



Apache Kafka Security Overview

Sven Erik Knop, Senior Solutions Architect



Who am I?

- Sven Erik Knop
- Senior Solutions Architect
- Architect, Developer, Presenter
- @sven_eric_knop



Agenda



Why security?



Implementation

Encryption

Authentication

Authorization



RBAC

Why RBAC?

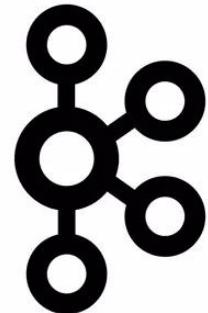
Implementation

Impact

Reminder: What is Apache Kafka?



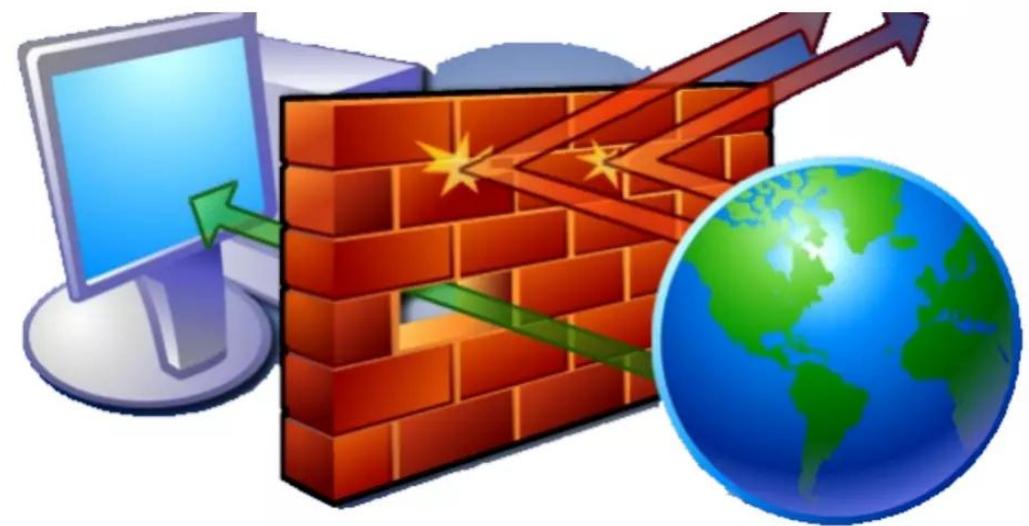
- Distributed, fault-tolerant and immutable log
- Use cases
 - Publish-subscribe queue
 - ETL (Extract, Transform, Load)
 - Event streaming platform
- Rich infrastructure through connectors and language APIs
- Open Source, with enterprise support and enhancements through companies like Confluent





Apache Kafka before security (<0.9)

- Kafka is infrastructure hidden behind the firewall
- Servers within the network are implicitly trusted
- No governance required (or possible)





New challenges

Multi-tenancy

Topics with
sensitive or
confidential data

Regulatory
requirements



Security Features added in Kafka 0.9



Encryption



Authentication



Authorization



Connection to
ZooKeeper

Encryption



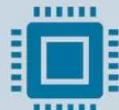
Encryption in flight

TLS = Transport Layer Encryption



Encryption at rest

OS or file system feature



End-to-End encryption

Needs to be built into the
producers and consumers



Encryption via TLS (Transport Layer Security)

Prevents eavesdropping

Prevents man-in-the-middle attacks

Industry-standard encryption

Broker certificate signed by either

- External, well-known Certificate Authority (CA) – GlobalSign, Let's Encrypt ...
- Internal CA, with clients using a truststore to verify its identity

`security.protocol=SSL` (instead of the default `PLAINTEXT`)

`ssl.endpoint.identification.algorithm=https`



TLS Mutual Authentication

If enabled,
provides trust into
client connections

Principal is used
for authorization

Certificates signed
by trusted
authority (CA)

Clients additionally use a keystore to provide their own identity

SASL (Simple Authentication and Security Layer)



Framework for authentication and data security



Decouples authentication mechanisms from the application protocols



Common interface for different implementations

SASL Implementations



Implementation	Class	Comments
PLAIN	org.apache.kafka.common.security.plain.PlainLoginModule	Username/Password
SCRAM	org.apache.kafka.common.security.scram.ScramLoginModule	SCRAM_SHA_(256 512)
GSSAPI	com.sun.security.auth.module.Krb5LoginModule	Kerberos
OAUTHBEARER	org.apache.kafka.common.security.oauthbearer.OAuthBearerLoginModule	Token based used for RBAC

security.protocol = SASL_PLAINTEXT | SASL_SSL

sasl.mechanism = PLAIN | SCRAM_SHA_256 | SCRAM_SHA_512 | GSSAPI | OAUTHBEARER

Configured in the brokers and the clients through a JAAS file

JAAS



- Java Authentication and Authorization Service
- For services (Kafka, ZooKeeper) grouped by *section names*
 - KafkaServer
 - KafkaClient
 - Server (for ZooKeeper Server)
 - Client (for ZooKeeper Client, e.g. Kafka)
- Alternatively, specify as a property

```
KafkaServer {
    org.apache.kafka.common.security.scram.ScramLoginModule required
        username="kafka"
        password="kafka";
};

Client {
    org.apache.zookeeper.server.auth.DigestLoginModule required
        username="admin"
        password="admin-secret";
};
```

```
sasl.mechanism=SCRAM-SHA-256
security.protocol=SASL_PLAINTEXT
sasl.jaas.config=org.apache.kafka.common.security.scram.ScramLoginModule required \
    username="producer" \
    password="producer-secret";
```

sasl.mechanism=PLAIN



- Username/Password
- Broker
 - By default, holds a list of username/password in the SASL config file
 - Requires rolling restart when users are added or deleted
 - Can be overwritten, for example
 - `listener.name.sasl_plaintext.plain.sasl.server.callback.handler.class=io.confluent.security.auth.provider.ldap.LdapAuthenticateCallbackHandler`
- Client

```
sasl.mechanism=PLAIN
security.protocol=SASL_PLAINTEXT
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule
required \
    username="alice" \
    password="alice-secret" ;
```

SCRAM



- **Salted Challenge Response Authentication Mechanism**
- Users and passwords are stored in ZooKeeper
- Dynamic configuration through *kafka-configs*
- Different encryption levels (SHA-256, SHA-512)
- CAVEAT: Need to add any super user as an entity as well



```
kafka-configs --zookeeper zookeeper:2182 --alter  
--add-config 'SCRAM-SHA-256=[password=xxxxxx],SCRAM-SHA-512=[password=xxxxxx]'  
--entity-type users  
--entity-name producer
```

```
sasl.mechanism=SCRAM-SHA-256  
security.protocol=SASL_PLAINTEXT  
sasl.jaas.config=org.apache.kafka.common.security.scram.ScramLoginModule  
required \  
username="producer" \  
password="xxxxxx";
```



Kerberos



- Network authentication protocol created by the MIT
- `sasl.enabled.mechanisms=GSSAPI`
 - Generic Security Service Application Program Interface
- External Key Distribution Center (KDC)
- Service (Broker, ZooKeeper ...) registers itself
 - `sasl.kerberos.service.name` must match principal
 - Default service names are 'kafka', 'zookeeper'



Kerberos client

- Either: create a ticket with *kinit* (login)
- Or: provide the keytab location and the principal name
- List existing tickets with *klist*

```
sasl.jaas.config=com.sun.security.auth.module.Krb5LoginModule required \
    useKeyTab=true \
    storeKey=true \
    keyTab="/etc/security/keytabs/kafka_client.keytab" \
    principal="kafkaclient1@EXAMPLE.COM";
```



Authorization via ACLS (Access Control Lists)



Permission to access or
modify resources like

Topics
Clusters
Consumer Groups



All, Literal and Prefix



Allow and deny access



Stored in ZooKeeper



ACL for users

ACLs are stored in ZooKeeper

- kafka-acls
- Can check in ZooKeeper directly

ACLs defined for

- Topic, Cluster, Group, TransactionId, DelegationToken
- (Principal, PermissionType, Operation, Host)

**Principal for users is
User:<user-name>**

- Confluent extension to use Group:<group-name> set in LDAP

**Can use resource name '*' for
catch all**

**To specify a subset of resources
with the same prefix, use**

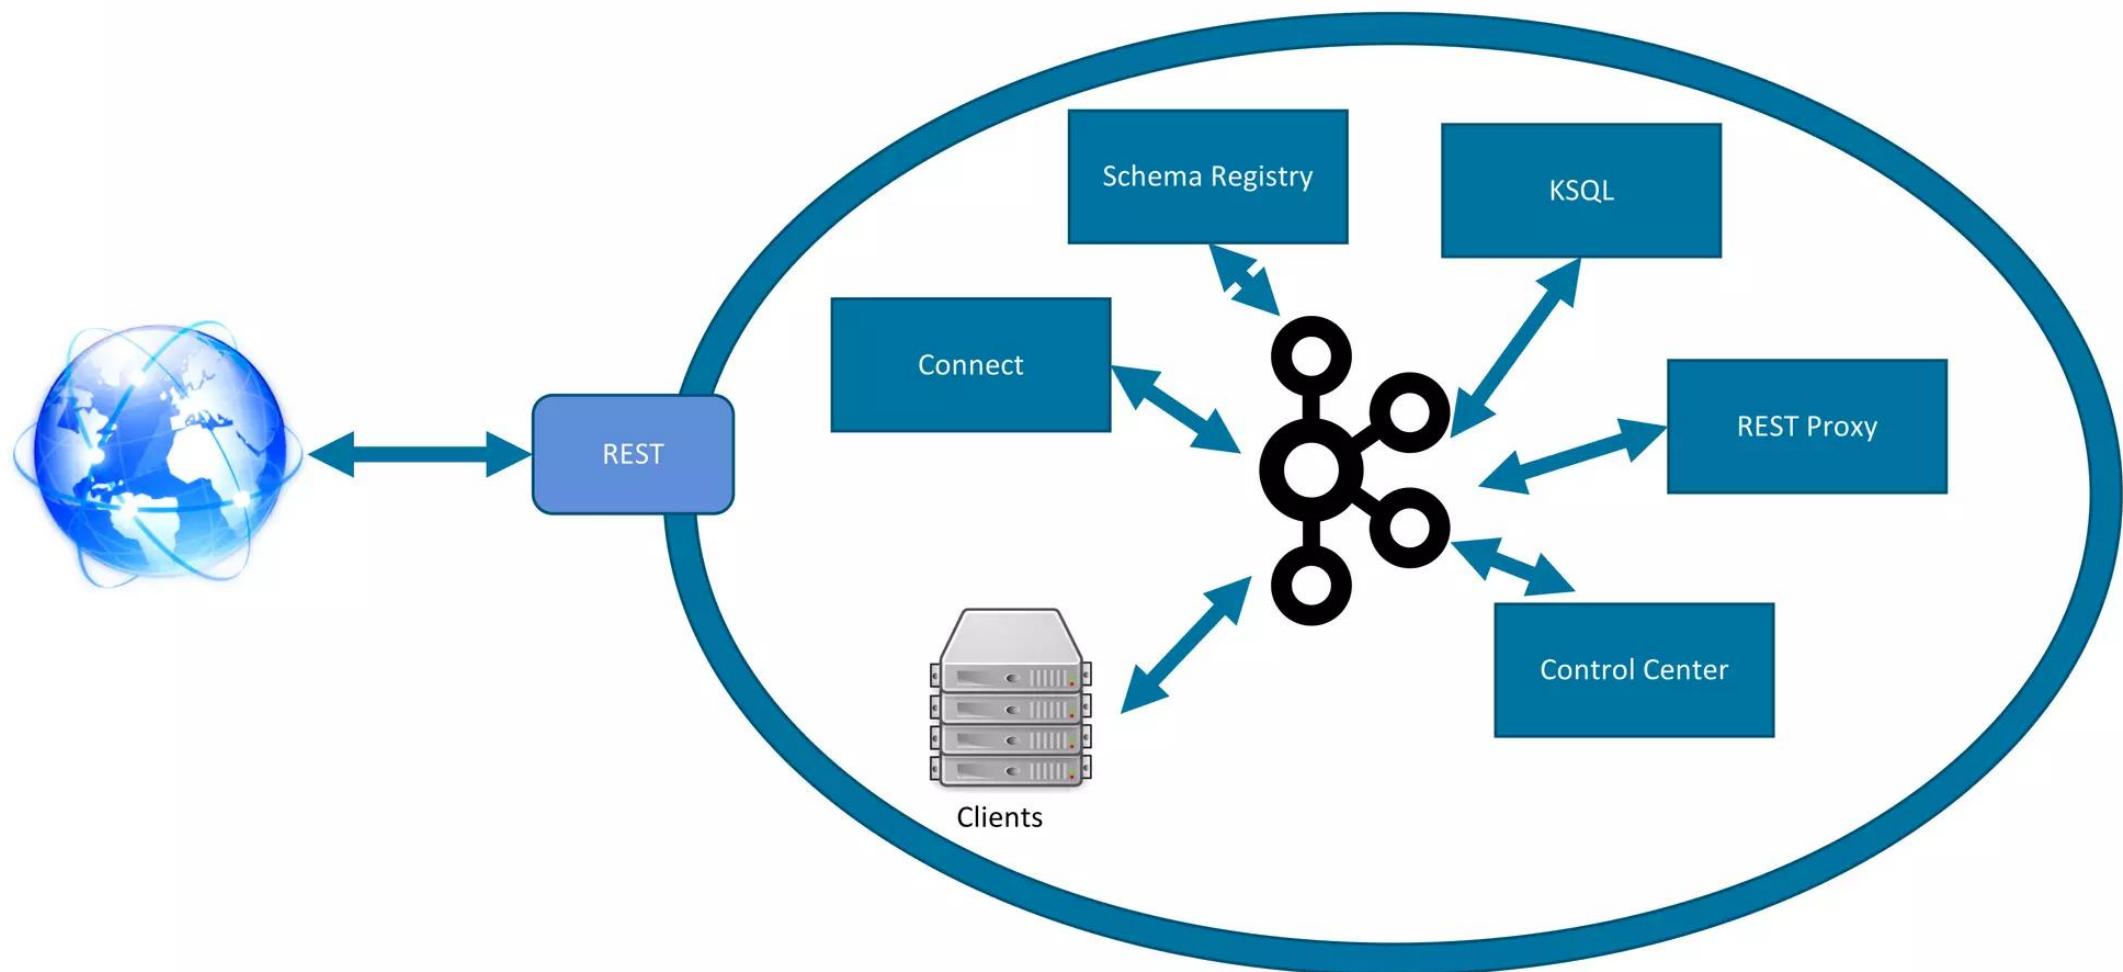
- --resource-pattern-type prefixed (default is literal)

Securing ZooKeeper

- Authentication
 - SASL
 - Digest
 - Kerberos
- Authorization
 - Securing your nodes through Kafka
 - `zookeeper.set.acl=true`
- Migrating from a non-secured node
 - `zookeeper-security-migration`



Kafka Infrastructure



Challenge



- Clients directly connected to the brokers use SASL or mutual auth for authentication
- Tools like KSQL or Connect are connected using a single user with *that user's* permissions
- ASK: Users of these tools need their own authentication and propagate their credentials



The Need for A Unified Approach

- Requirements
 - Services connect to Kafka using TLS Mutual Auth or SASL
 - External access via REST needs to impersonate an internal user
- Want
 - Single service to assign roles
 - Changes need to be traceable (ZooKeeper does not keep a log)





RBAC – Role Based Access Control



(Pre-defined) roles can be assigned to users or groups for a resource



Stored in a Kafka topic accessed via a new Kafka interface, the MDS (Meta Data Service)

Requires Confluent Server instead of Apache Kafka



User authentication backed by LDAP

Authorized users are given an OATHBEARER Token to identify themselves

Roles



Administration roles (SystemAdmin, ClusterAdmin, UserAdmin, ...)

Operators

ResourceOwners

Developers



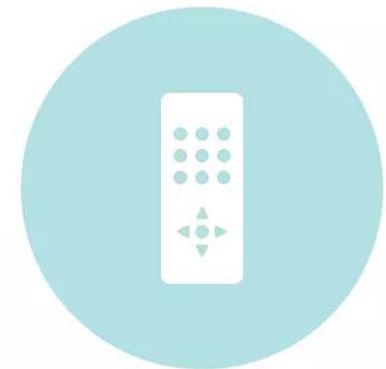
Roles Assignment to User/Group



CLI



API



CONTROL CENTER

Group membership management via LDAP

Confluent Control Center



Welcome to Confluent Platform

Login with your username and password

Log in



Password

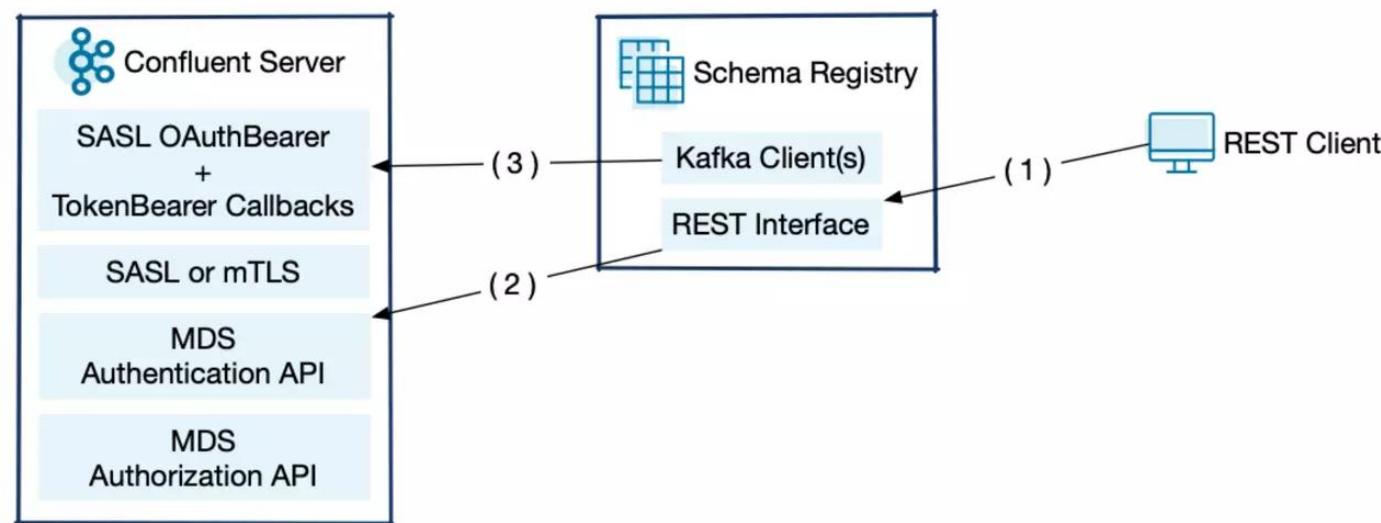


Log in

Schema Registry, Connect, KSQL, REST Proxy



- Standard REST username/password security (from LDAP)
- MDS Server provides OAUTHBEARER token
- Credentials are propagated

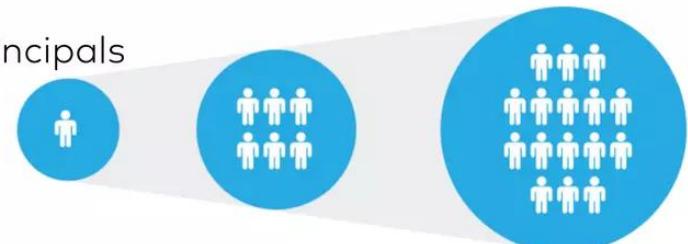
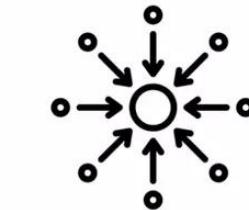




RBAC Key Benefits

Key Benefits Over Traditional ZK ACLs:

- Efficiency with Scale
 - Confluent Platform-wide security
 - Authorization at scale - with predefined Roles and LDAP Groups as Principals
 - Manage multiple clusters from 1 central location
- Increased Granular Policies
 - Connector level access control.
 - Granular access for Control Center users
 - On top of the Topic and SR Subjects
- Improved Multi-Tenancy support
 - Items without access do not show
 - Delegation of Administration Responsibilities



Requirements



- Confluent Server >= CP 5.4
- LDAP (Usually Microsoft Active Directory) for
 - User **authentication** (usually Username/Password SIMPLE BIND)
 - Group membership for **authorization**
- All services configured with
 - sasl.enabled.mechanism=OAUTHBEARER
 - sasl.login.callback.handler.class=
io.confluent.kafka.clients.plugins.auth.token.TokenUserLoginCallbackHandler
 - metadata.bootstrap.server.urls=http://mds-server:8090
- Do not configure standard producers/consumerstreams with OAUTHBEARER
 - Unsupported
 - Would require access to MDS server

Summary

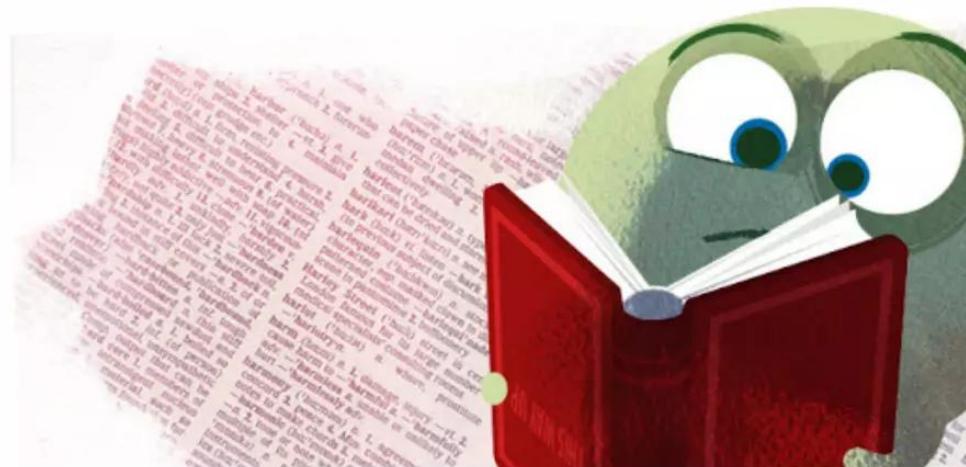


- Security is an important feature of a streaming platform
- Encryption and ACLs provide security for clients directly accessing the Kafka brokers
- RBAC adds a unified authentication and authorization interface for services around Kafka



Further reading

- https://docs.confluent.io/current/security/security_tutorial.html
- <https://docs.confluent.io/current/security/rbac/index.html>





Any questions?





CONFLUENT