



Getting Started with Confluent Schema Registry

Patrick Druley

Senior Solution Engineer @confluentinc



Agenda



<https://www.linkedin.com/in/patrickdruley/>

twitter: @PatrickLovesAK

1. Schema-less Kafka
2. Schema Registry Basics
3. 3 Good Habits
4. Schema Validation Demo



Schema-less Kafka

“Someone changed the date field from unix timestamp to datetime and broke all our reporting dashboards”

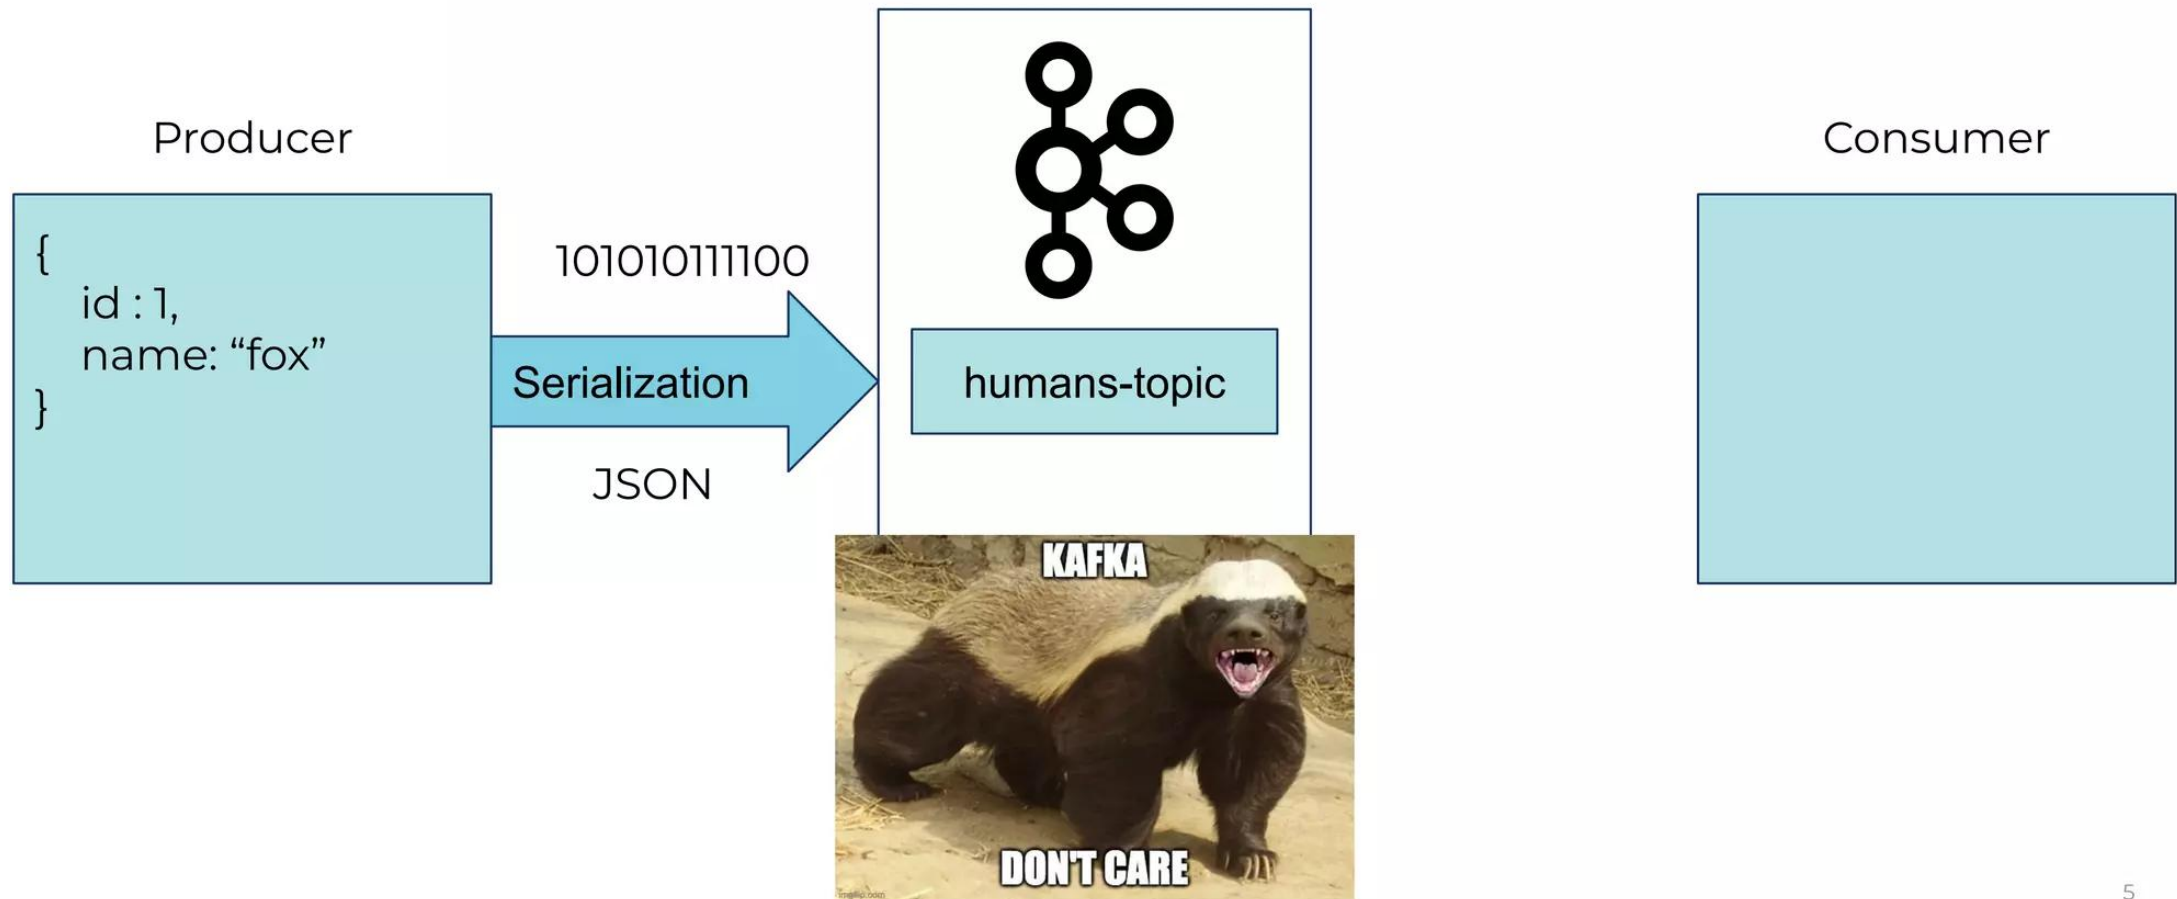
“I need to add a new field to support new functionality, but I don’t want to update ALL the potential downstream consumers of this data”

“My compliance team wants an audit to ensure there is no PII data in Kafka, there’s no easy way to get all the message metadata”

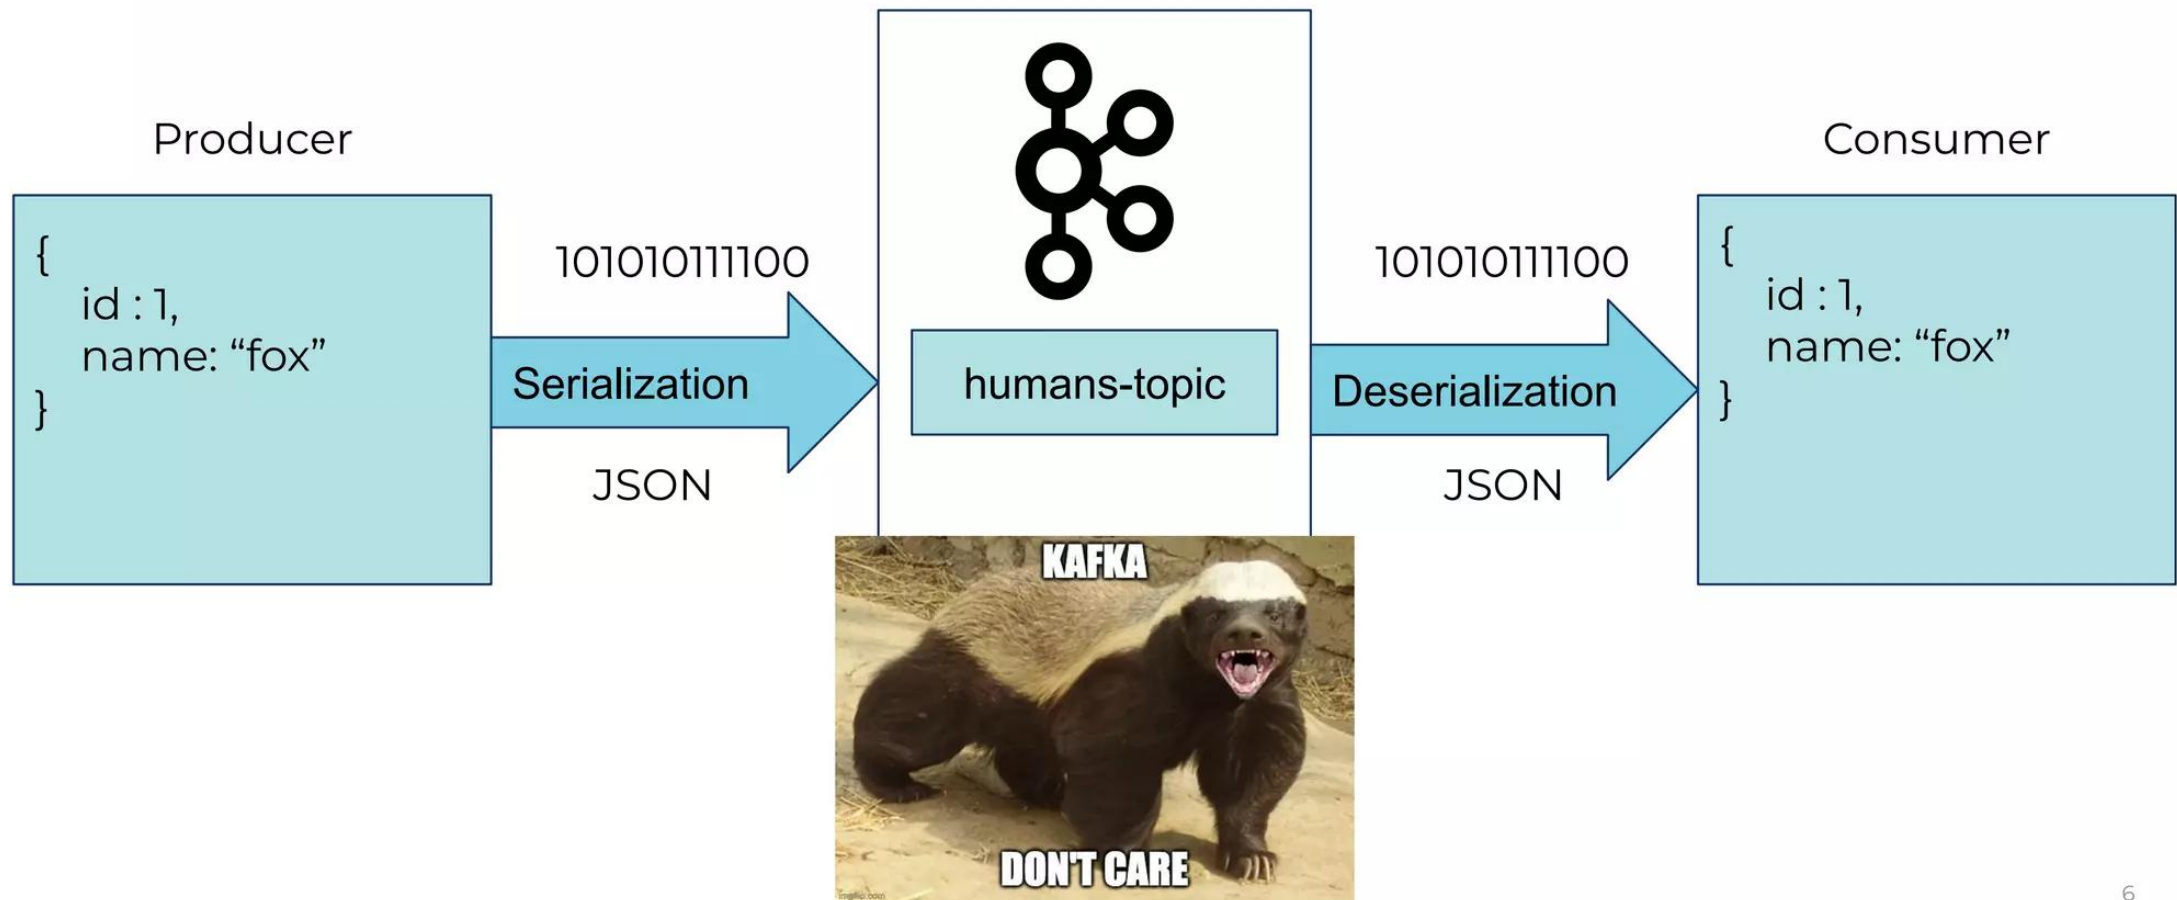
“My data is ‘democratized’ in Kafka, but my dev teams have no way to know what data is *in* Kafka without domain knowledge or coming to my platform team”



“My event is JSON” - Producer



“Sounds great” - Consumer



Time Passes



humans-topic

```
{  
  id: 1,  
  name: "fox"  
}
```

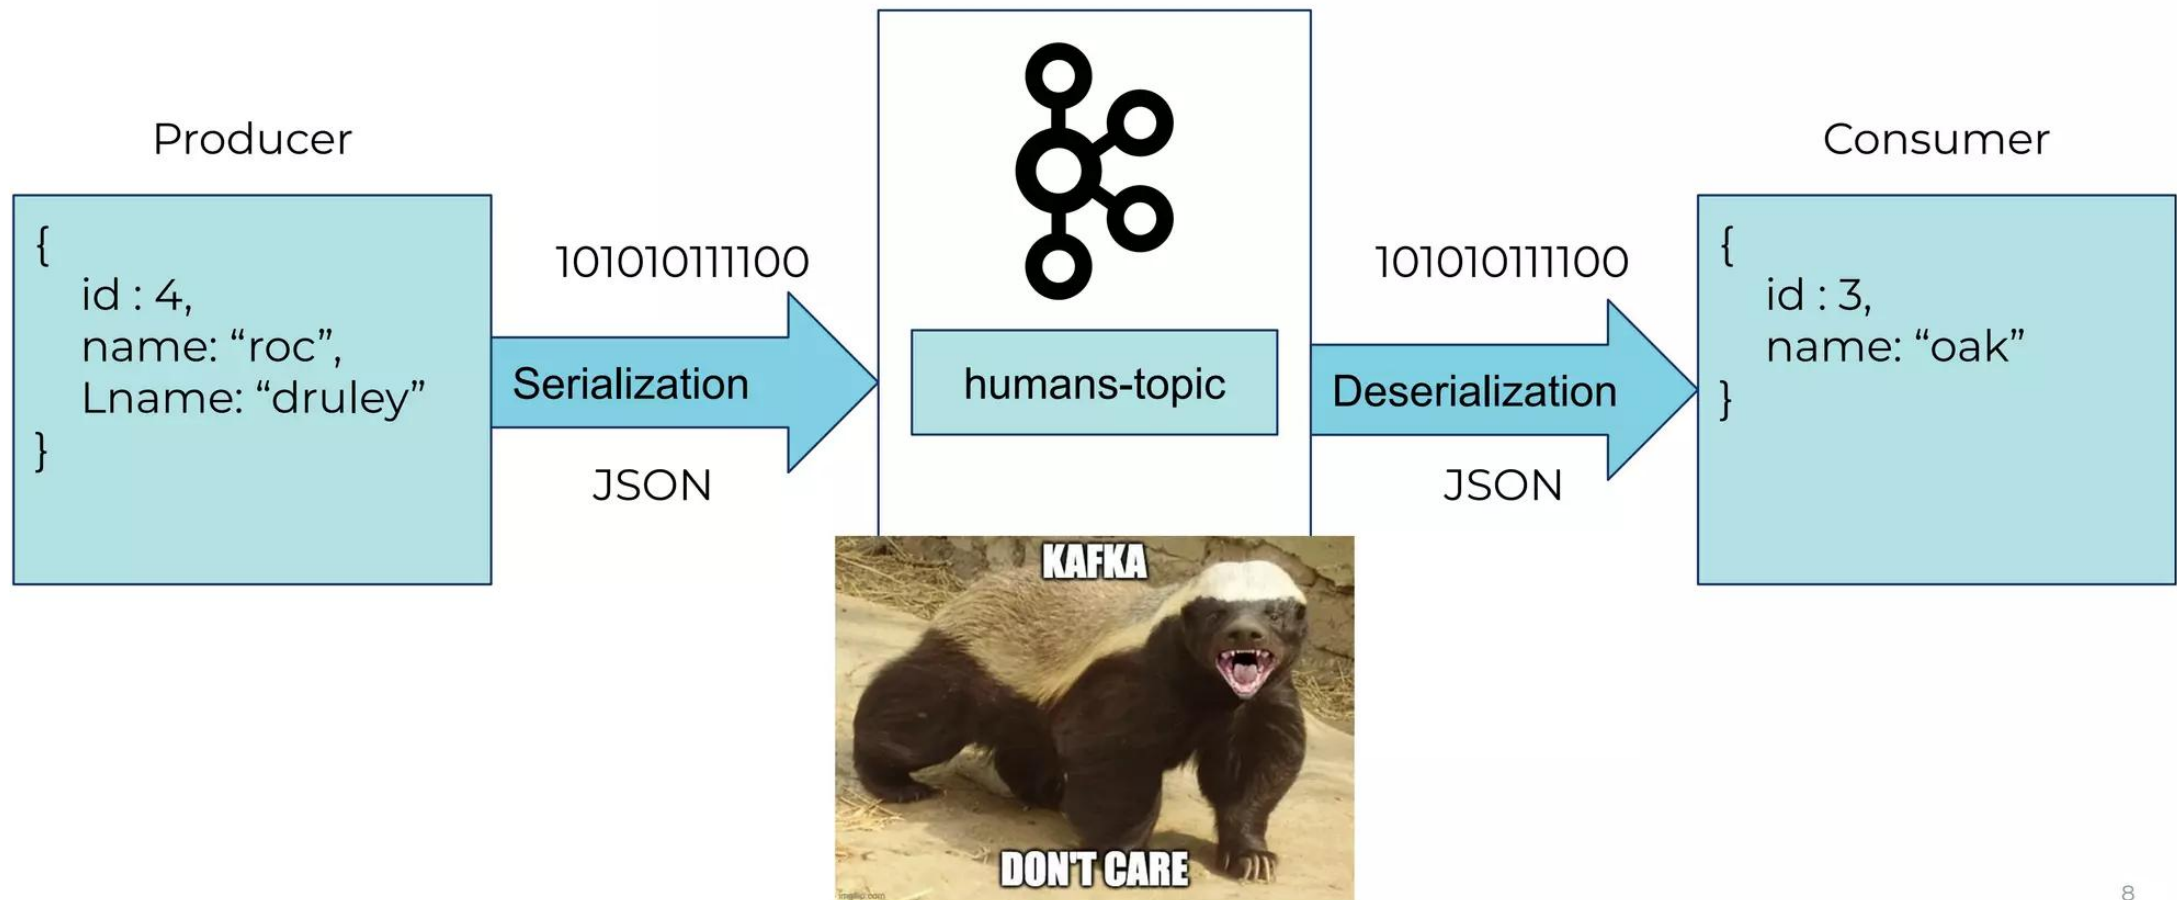
```
{  
  id: 2,  
  name: "jet"  
}
```

```
{  
  id: 3,  
  name: "oak"  
}
```

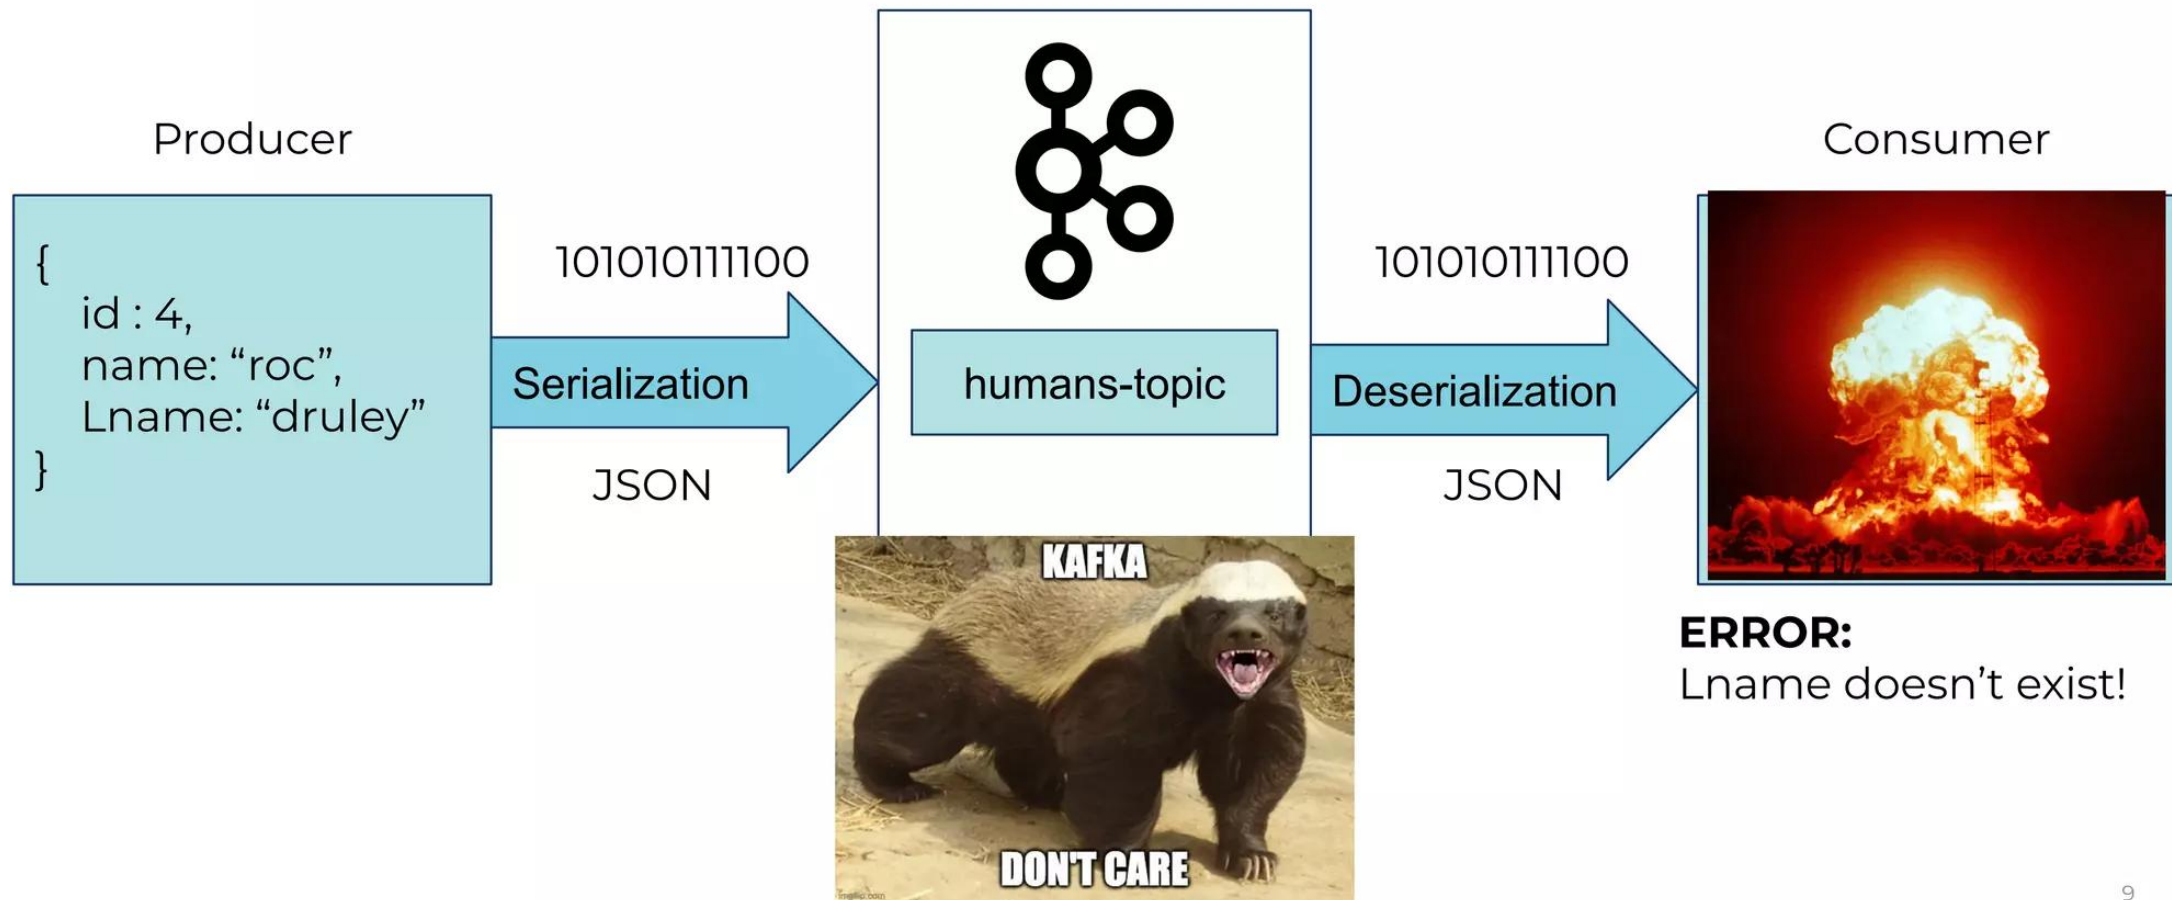
Tail

Head

“I added a field” - Producer



“You did what?” - Consumer



What Happened?

Consumer code breaks:

```
#new field Lname  
Lname = humans["Lname"]
```

KeyError: 'Lname'



humans-topic

<pre>{ id:1, name:"patrick" }</pre>	<pre>{ id:2, name:"roc" }</pre>	<pre>{ id:3, name:"oak" }</pre>	<pre>{ id:4, name:"roc", Lname:"druley" }</pre>
---	---	---	---

Tail

**But there's no "Lname"
in event 3.**

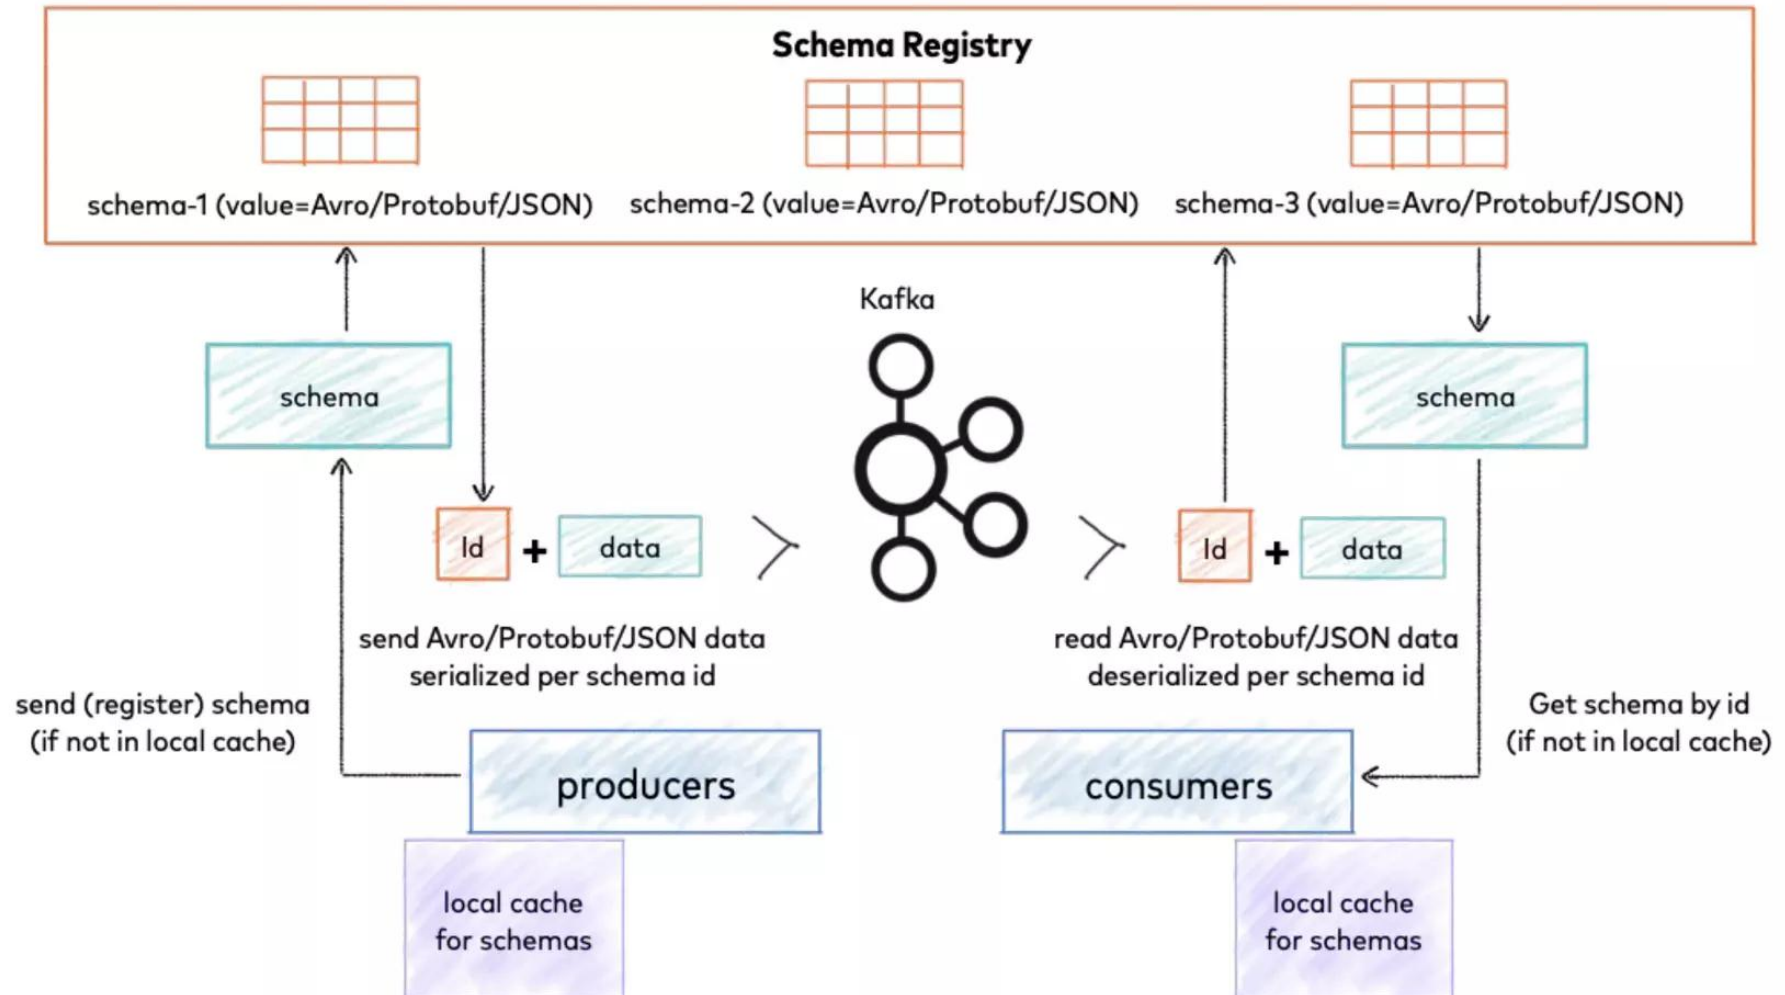
Consumer

Head



Schema Registry Basics

Confluent Schema Registry



Schema Compatibility - Avro



Compatibility Type	Changes Allowed	Check against which schemas	Upgrade First
BACKWARD	<ul style="list-style-type: none">• Delete Fields• Add optional fields	Last Version	Consumers
BACKWARD_TRANSITIVE	<ul style="list-style-type: none">• Delete Fields• Add optional fields	All previous versions	Consumers
FORWARD	<ul style="list-style-type: none">• Add Fields• Delete Optional Fields	Last Version	Producers
FORWARD_TRANSITIVE	<ul style="list-style-type: none">• Add Fields• Delete Optional Fields	All previous versions	Producers
FULL	<ul style="list-style-type: none">• Add Optional Fields• Delete Optional Fields	Last Version	Any order
FULL_TRANSITIVE	<ul style="list-style-type: none">• Add Optional Fields• Delete Optional Fields	All previous versions	Any order
NONE	<ul style="list-style-type: none">• All Changes Accepted	Compatibility checking disabled	Depends

Subject Naming Strategies



TopicNameStrategy

Subject Name =
Topic Name + [key|value]

Example:

Topic Name = mytopic
Value Subject Name = mytopic-value

Default

RecordNameStrategy

Subject Name =
Record Name + [key|value]

Example:

```
{ "type": "record",  
  "name": "myrecord",  
  "fields":  
    [ { "name": "f1",  
        "type": "string" } ]  
}
```

Value Subject Name = myrecord-value

Time Ordered Events

TopicRecordNameStrategy

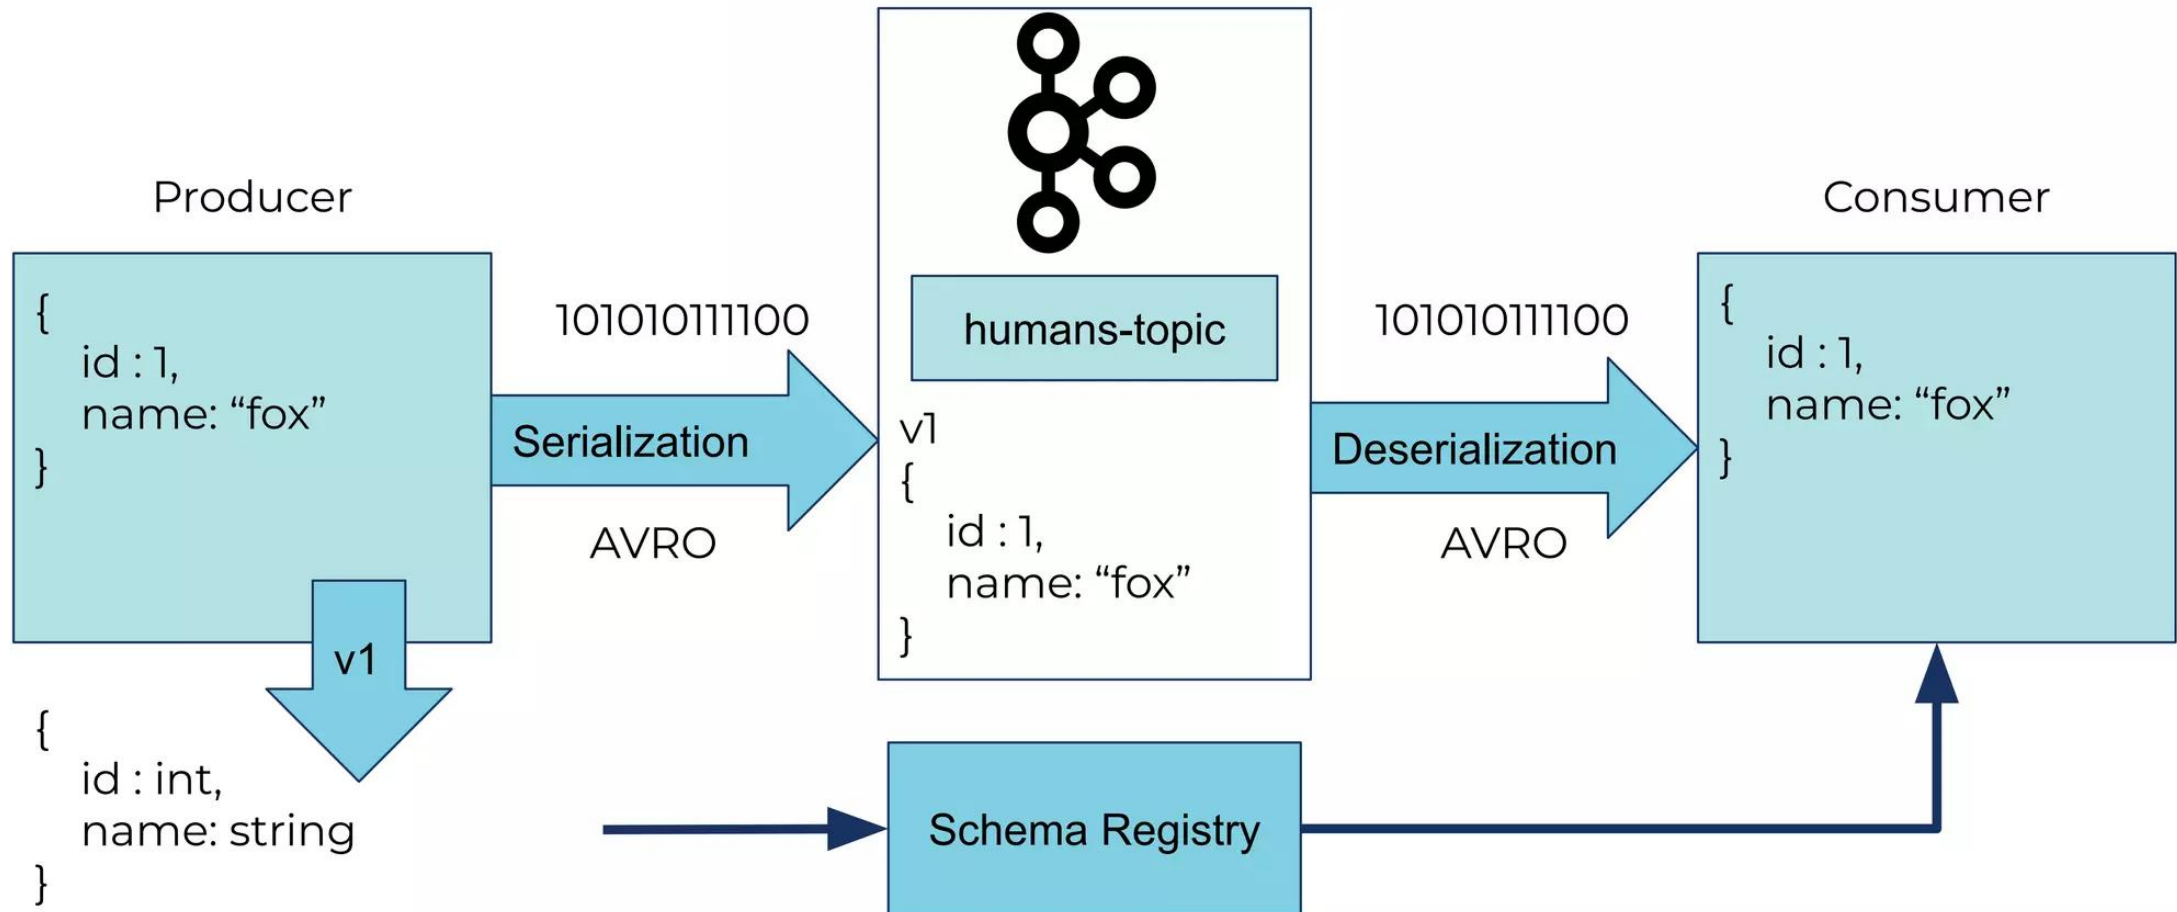
Subject Name =
Topic Name + Record Name + [key|value]

Example:

Topic Name = mytopic
Record Name = myrecord
Value Subject Name = mytopic-myrecord-value

Most Granular

Kafka with Schema Registry



Time Passes, Again



humans-topic

```
v1  
{  
  id:1,  
  name:"fox"  
}
```

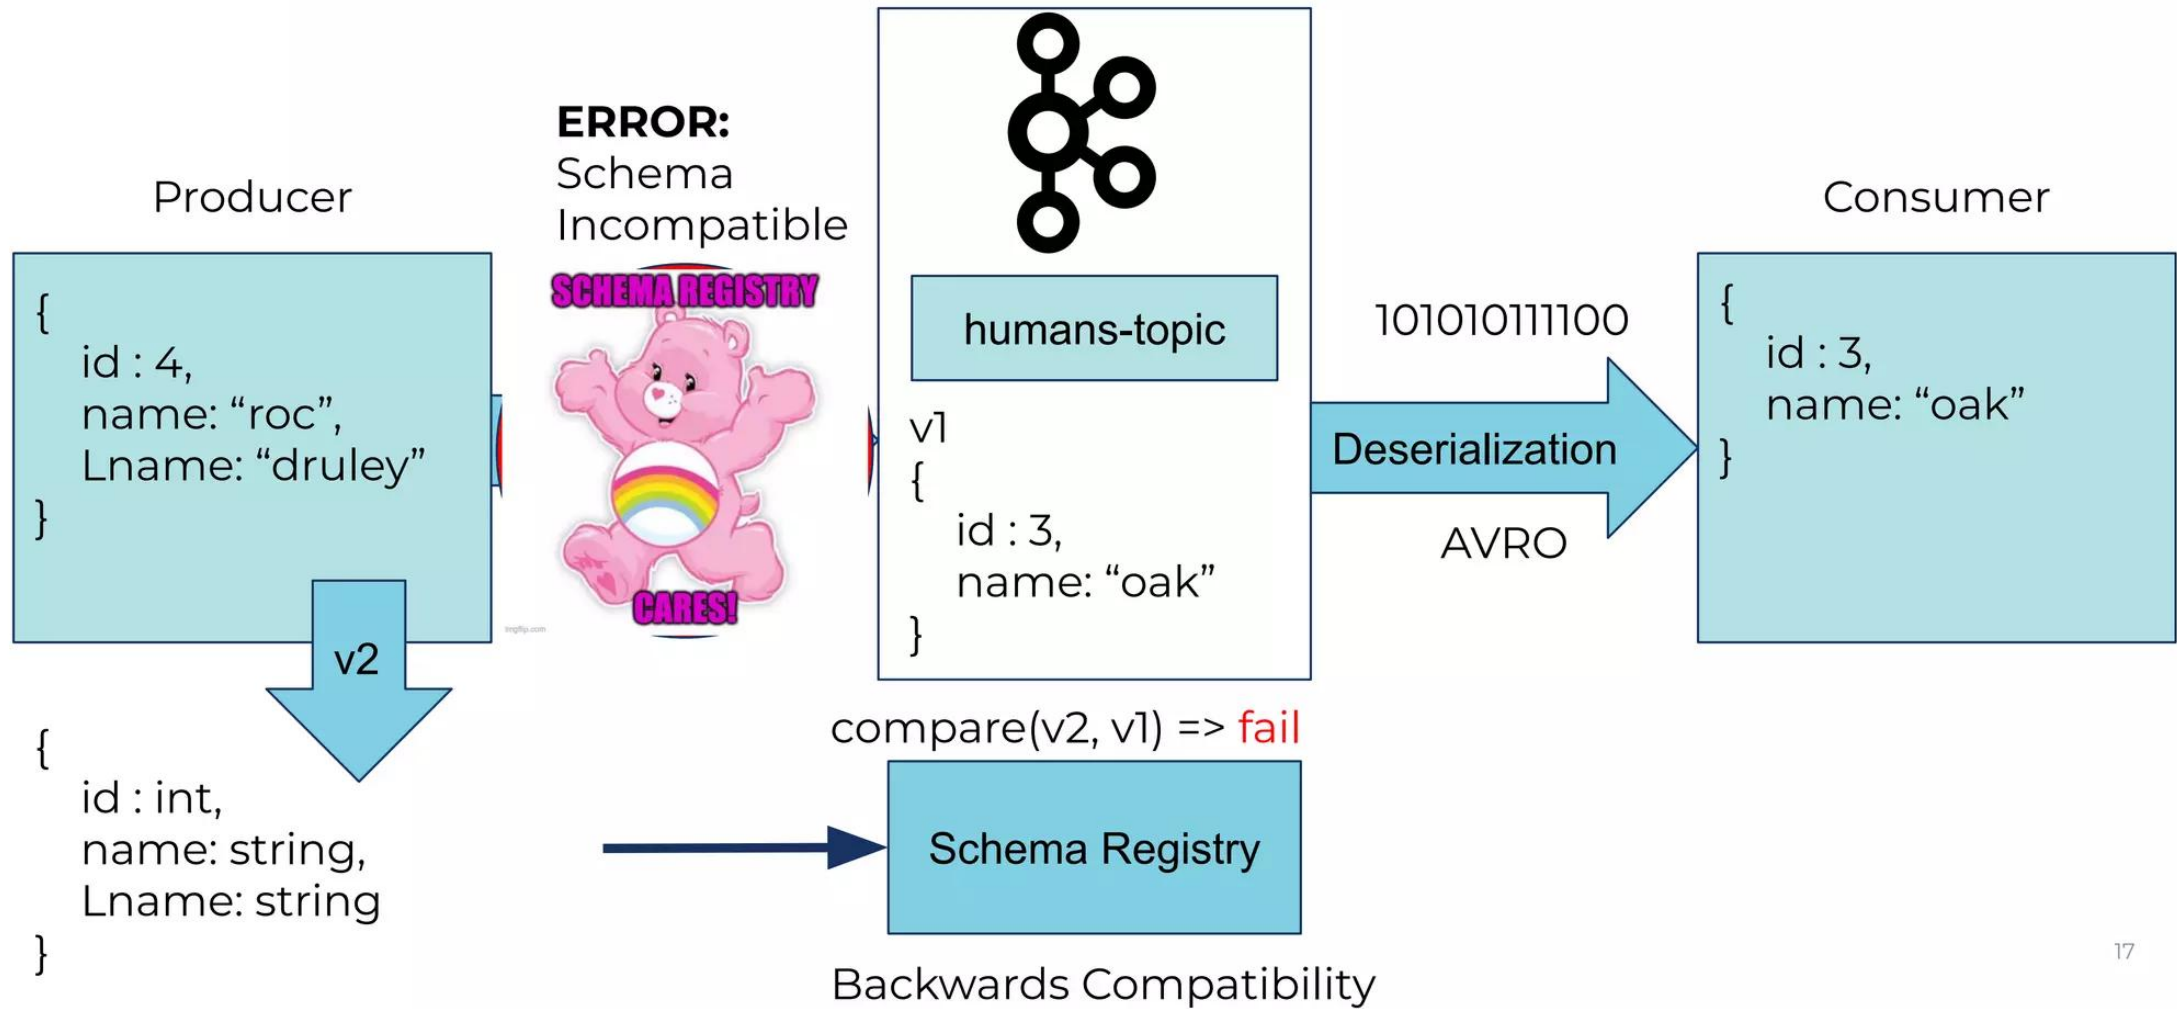
```
v1  
{  
  id:2,  
  name:"jet"  
}
```

```
v1  
{  
  id:3,  
  name:"oak"  
}
```

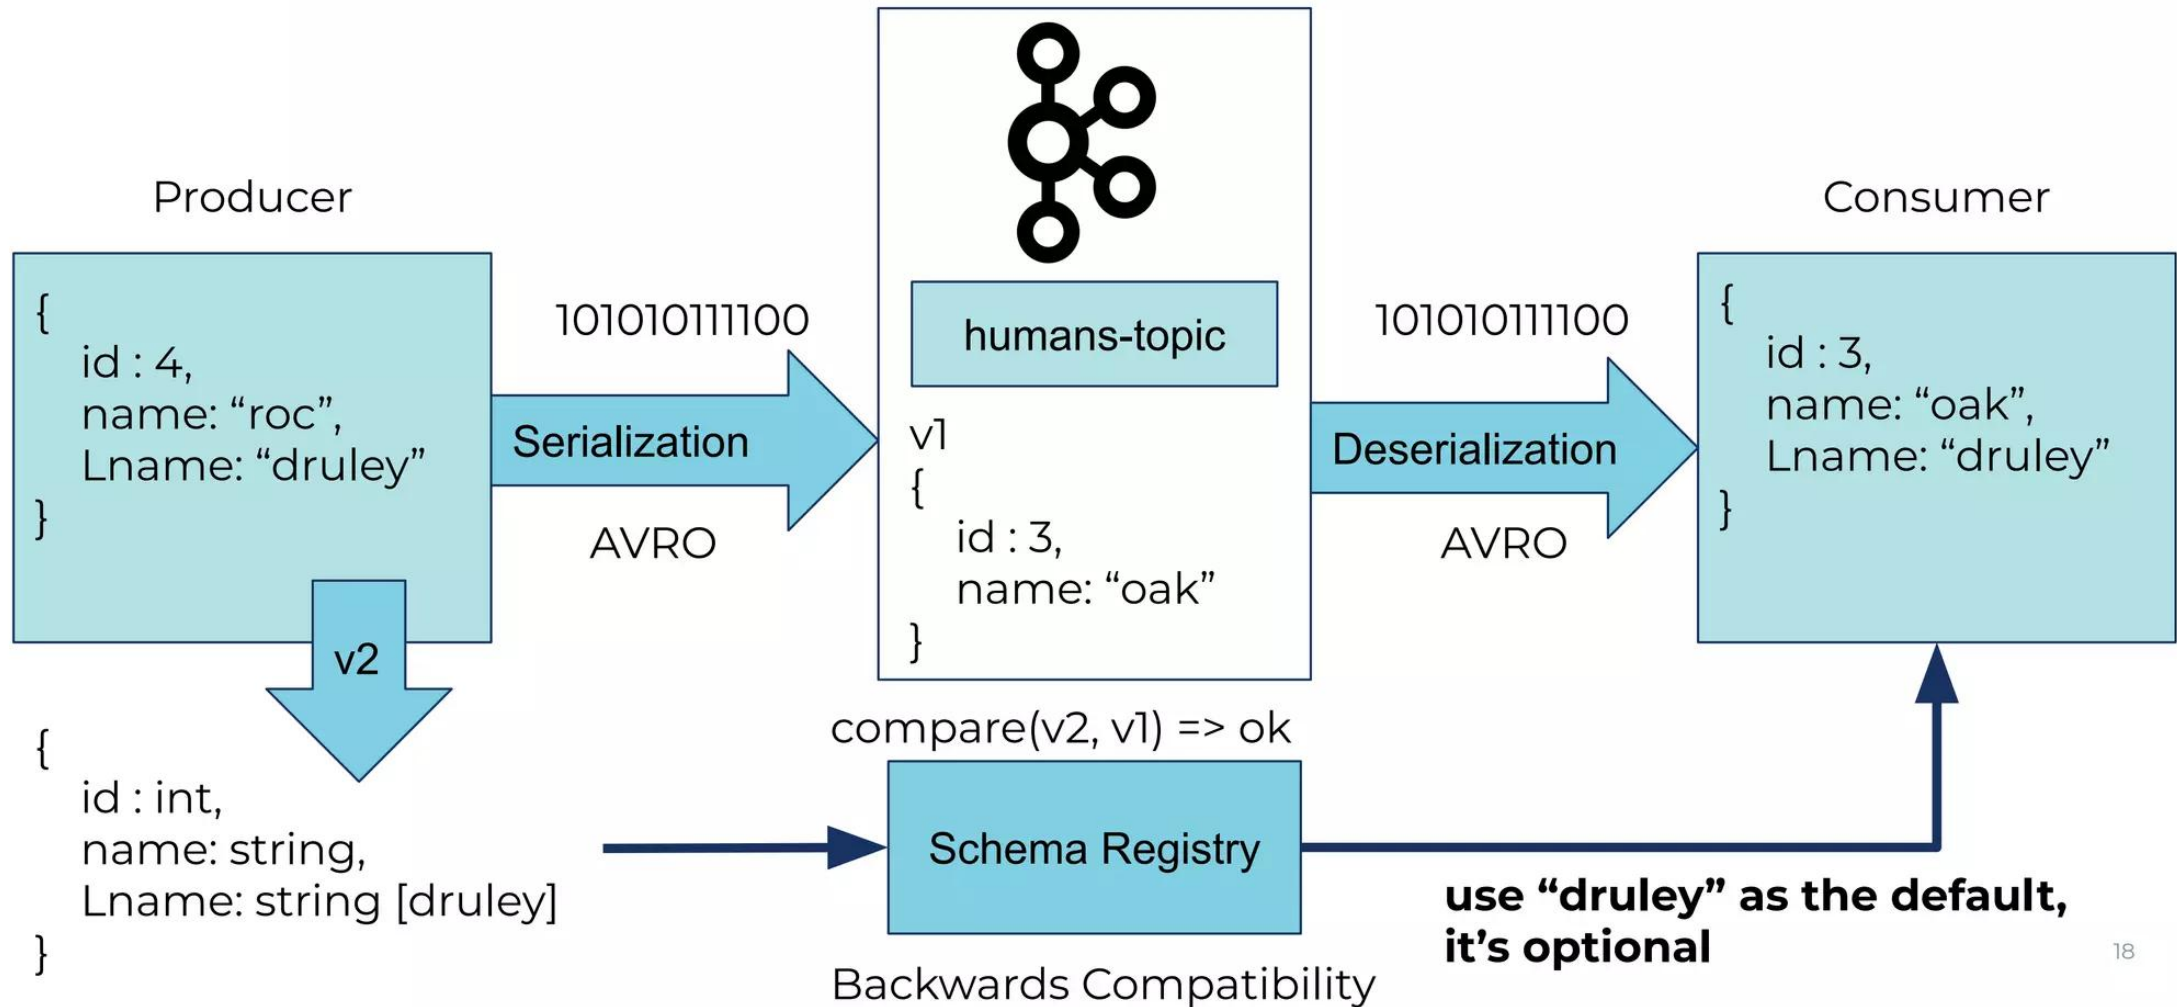
Tail

Head

“I added a field” - Producer



“I added an optional field” - Producer





3 Good Habits



1. Set auto register to false and add Schema Validation to CI/CD Pipeline.

Habit:

In prod and near-prod, clients should not automatically register new schemas. Ideally, this should be done in a CI/CD pipeline.

Producer Setting:

`auto.register.schema = false`

Exceptions:

1. Dev environments
2. Schema Registry ACLs are enabled using Confluent's security plugin
<https://docs.confluent.io/current/confluent-security-plugins/schema-registry/authorization/index.html#authorization-for-sr-operations-and-resources>

Schema Registry Maven Plugin - pom.xml



```
<plugin>
  <groupId>io.confluent</groupId>
  <artifactId>kafka-schema-registry-maven-plugin</artifactId>
  <version>${confluent.version}</version>
  <configuration>
    <schemaRegistryUrls>
      <param>${schemaRegistryUrl}</param>
    </schemaRegistryUrls>
    <userInfoConfig>${schemaRegistryBasicAuthUserInfo}</userInfoConfig>
    <subjects>
<transactions-value>src/main/resources/avro/io/confluent/examples/clients/basicavro/Payment2a.avsc</transactions-value>
      </subjects>
    </configuration>
    <goals>
      <goal>test-compatibility</goal>
    </goals>
  </plugin>
```

https://docs.confluent.io/current/schema-registry/schema_registry_onprem_tutorial.html#maven

2. Create new topics if you need to break compatibility.



Habit:

Figure out the right compatibility for you and create new topics in order to break it.

Don't retrofit your compatibility just for one time major schema changes.

Exceptions:

-Dev

3. Use Confluent Schema Validation.

Scale schemas reliably

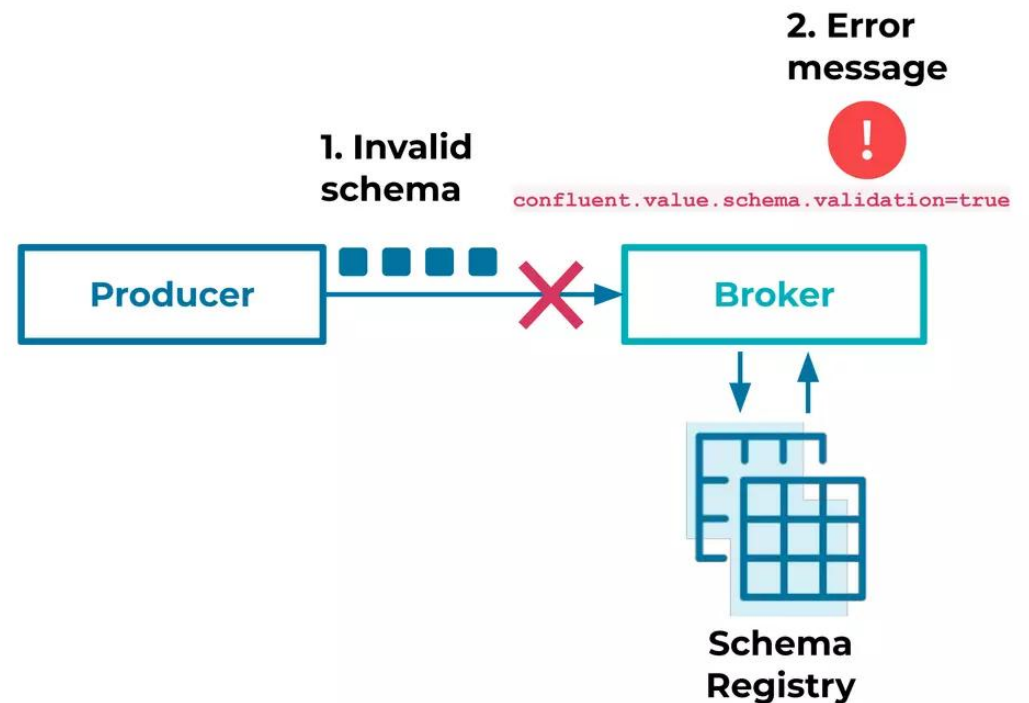
- **Automated broker-side** schema validation and enforcement
- **Direct interface** from the broker to Confluent Schema Registry

Granular control

- **Enabled validation** at the topic level
- Set **subject naming strategy** at the topic level



Broker Side Schema Validation





Schema Validation Demo

<https://docs.confluent.io/current/schema-registry/schema-validation.html#sv>

Learn Kafka.

Start building with
Apache Kafka at
Confluent Developer.



Confluent Developer
developer.confluent.io



Project Metamorphosis

Unveiling the next-gen event streaming platform

For Updates Visit
cnfl.io/pm



Jay Kreps
Co-founder and CEO
Confluent



Q&A



<https://www.linkedin.com/in/patrickdruley/>

patrick@confluent.io

<https://github.com/confluentinc/demo-scene/tree/master/industry-themes>



cnfl.io/meetups



cnfl.io/blog



cnfl.io/slack