



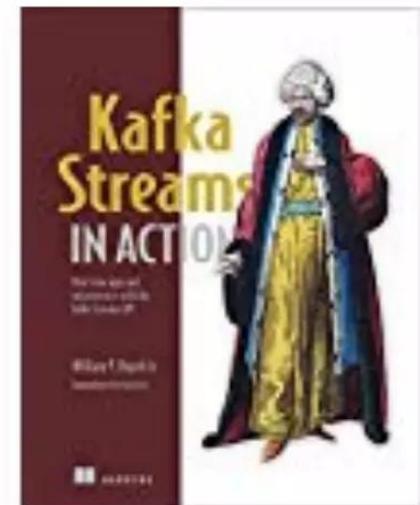
Introduction to Kafka Streams

Bill Bejeck, Integration Architect

Nice to meet you!



- Integration Architect on DevX team
- Prior to DevX ~3 years as engineer on Kafka Streams team
- Apache Kafka® Committer
- Author of “Kafka Streams in Action”



@bbejeck

Agenda



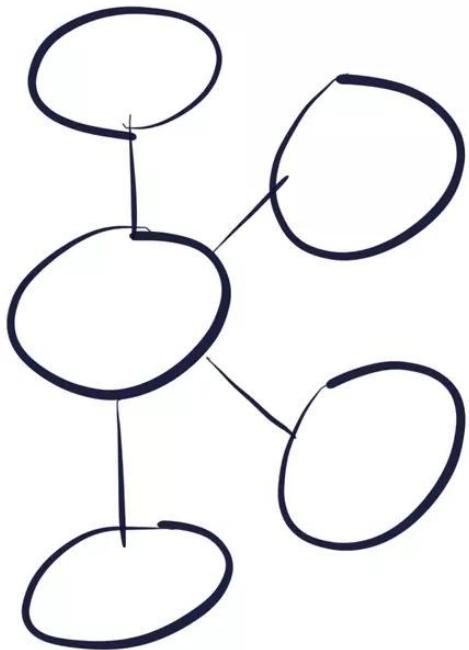
- Kafka Overview
- Life before Kafka Streams
- Kafka Streams Topology
- Kafka Streams Concepts/Architecture
- Kafka Streams API/Features

@bbejeck



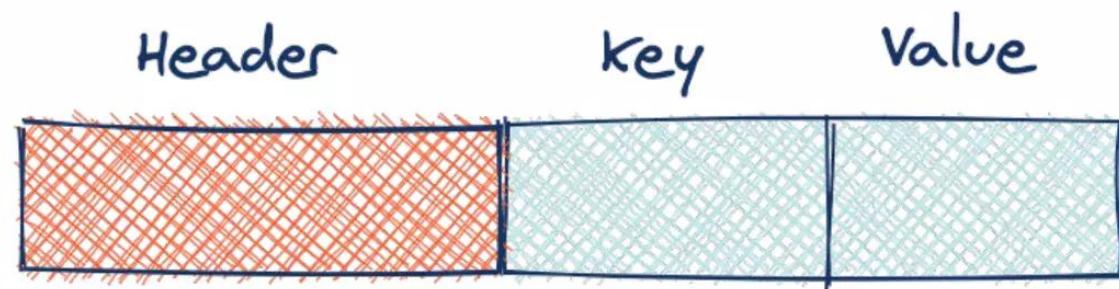
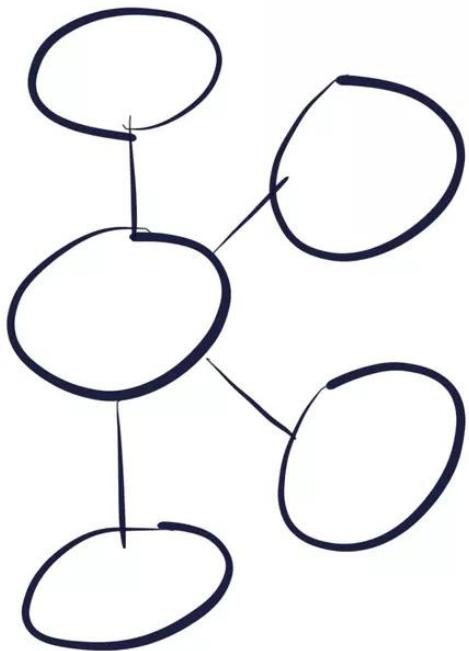
Kafka Overview

Kafka Overview



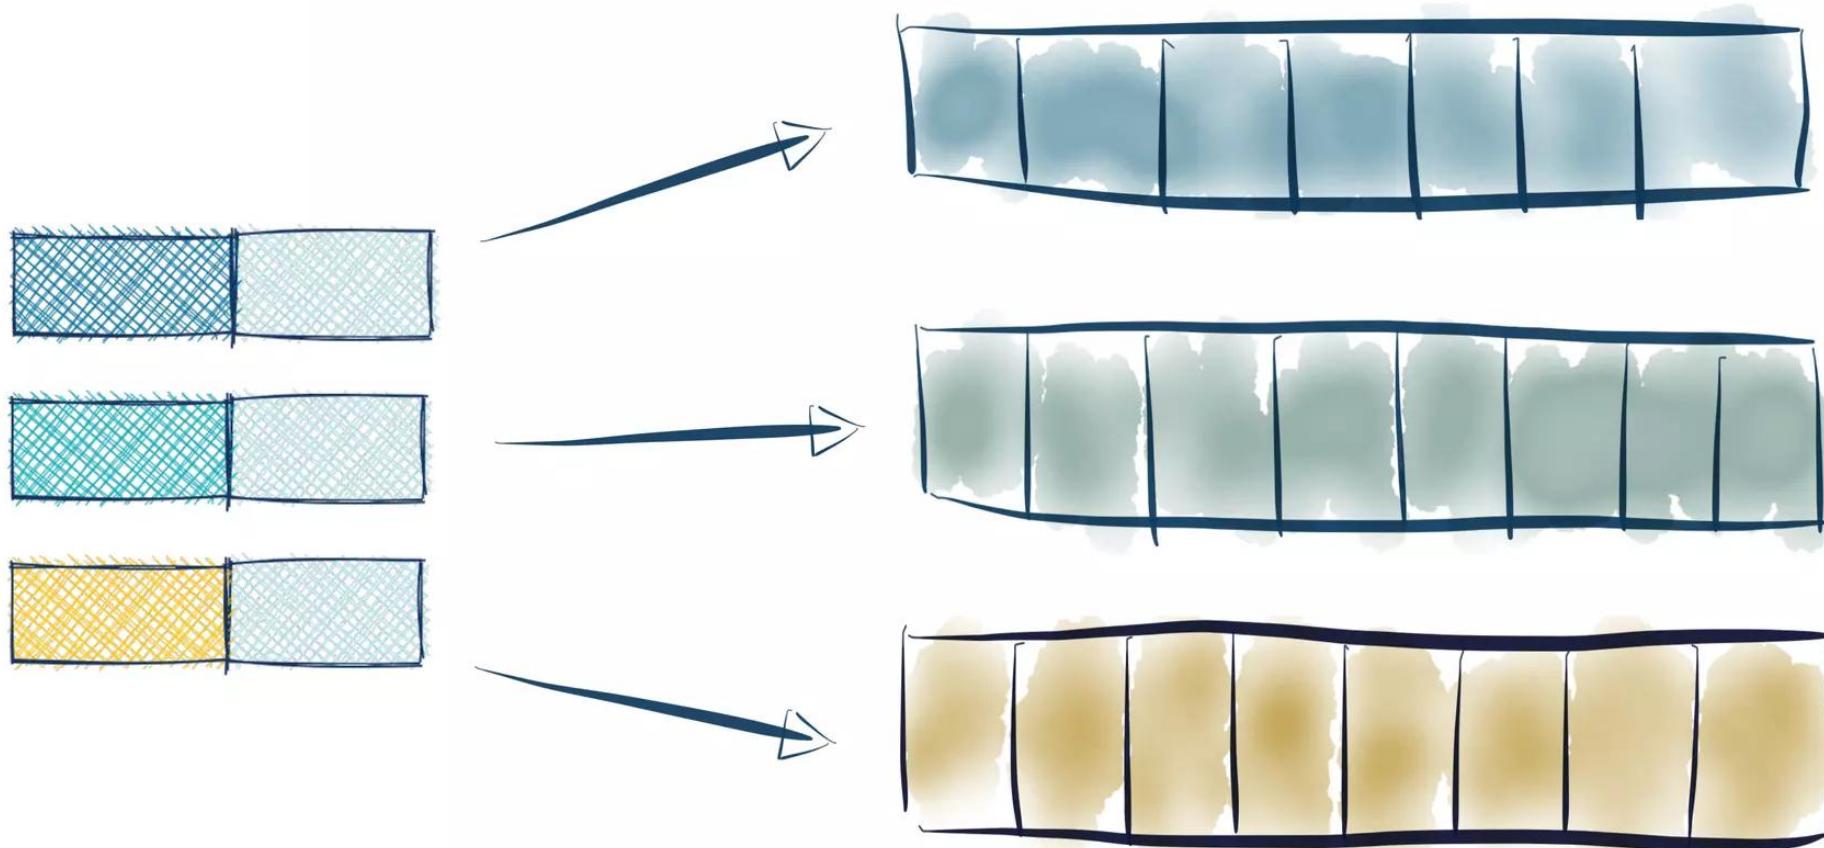
@bbejeck

Kafka Overview



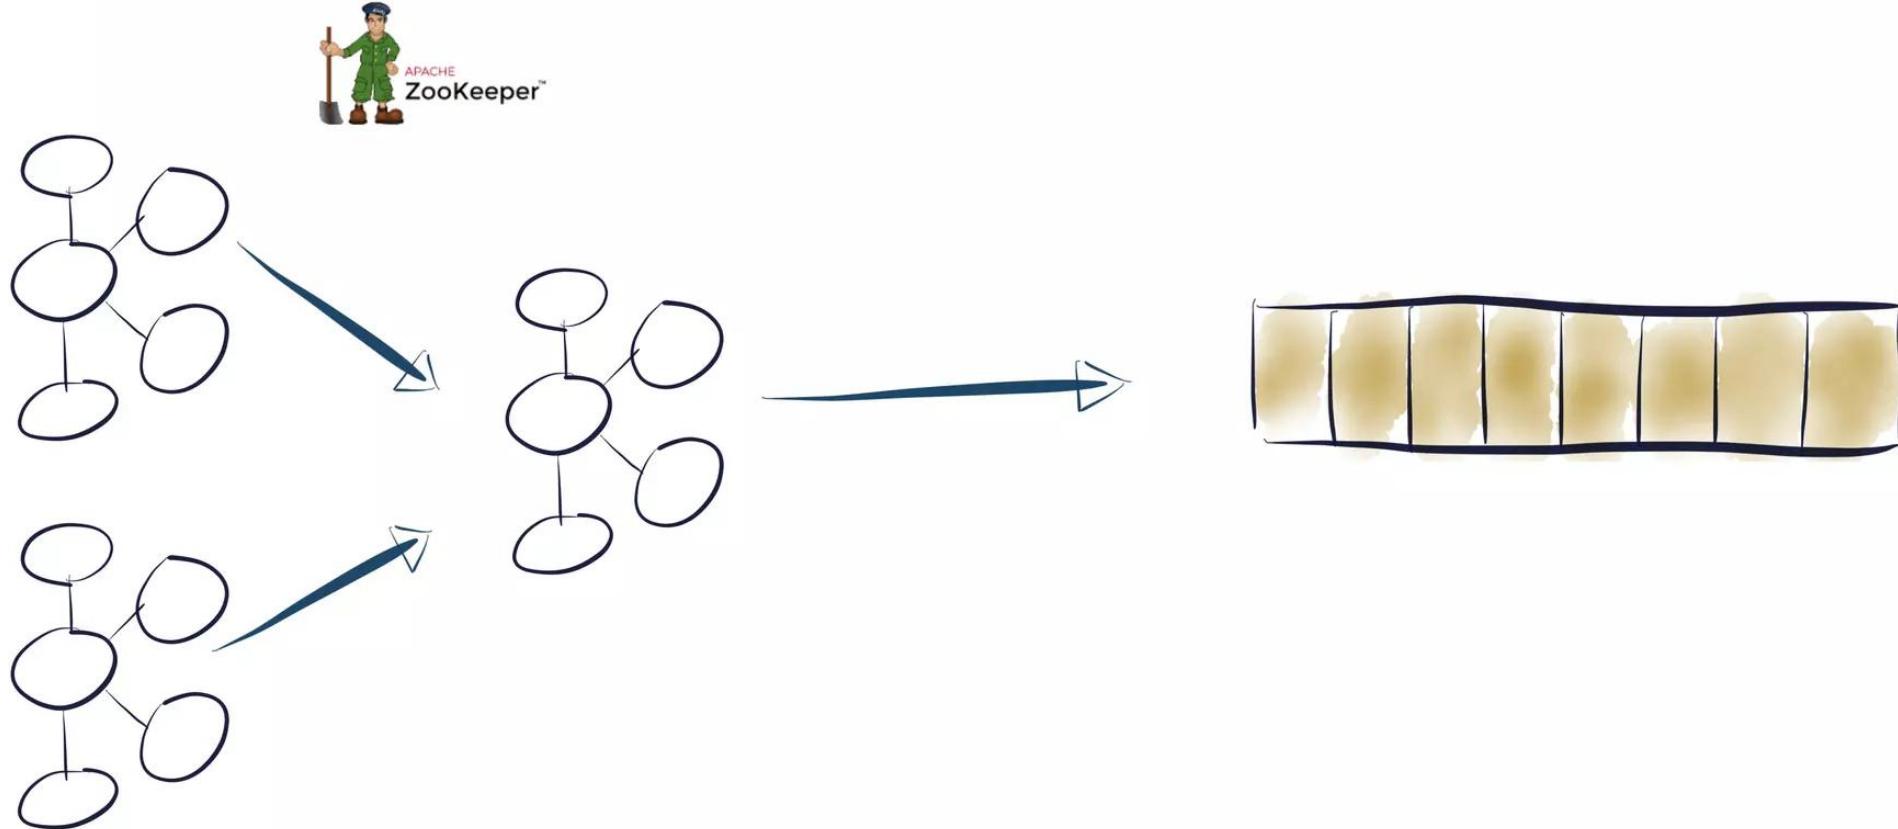
@bbejeck

Kafka Overview



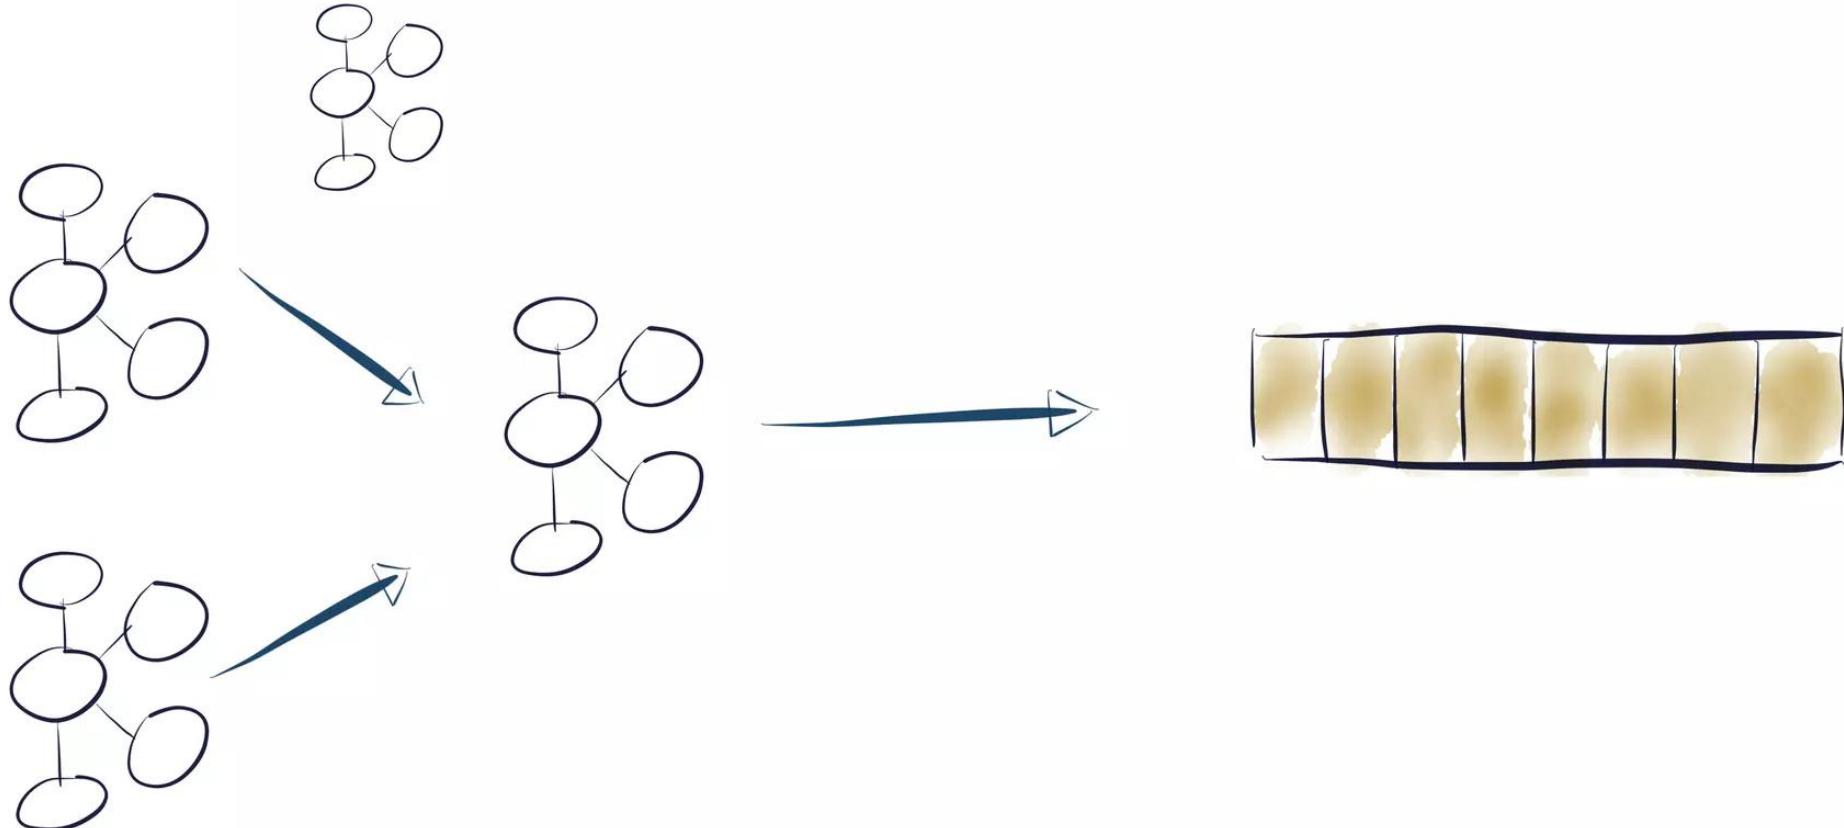
@bbejeck

Kafka Overview



@bbejeck

Kafka Overview

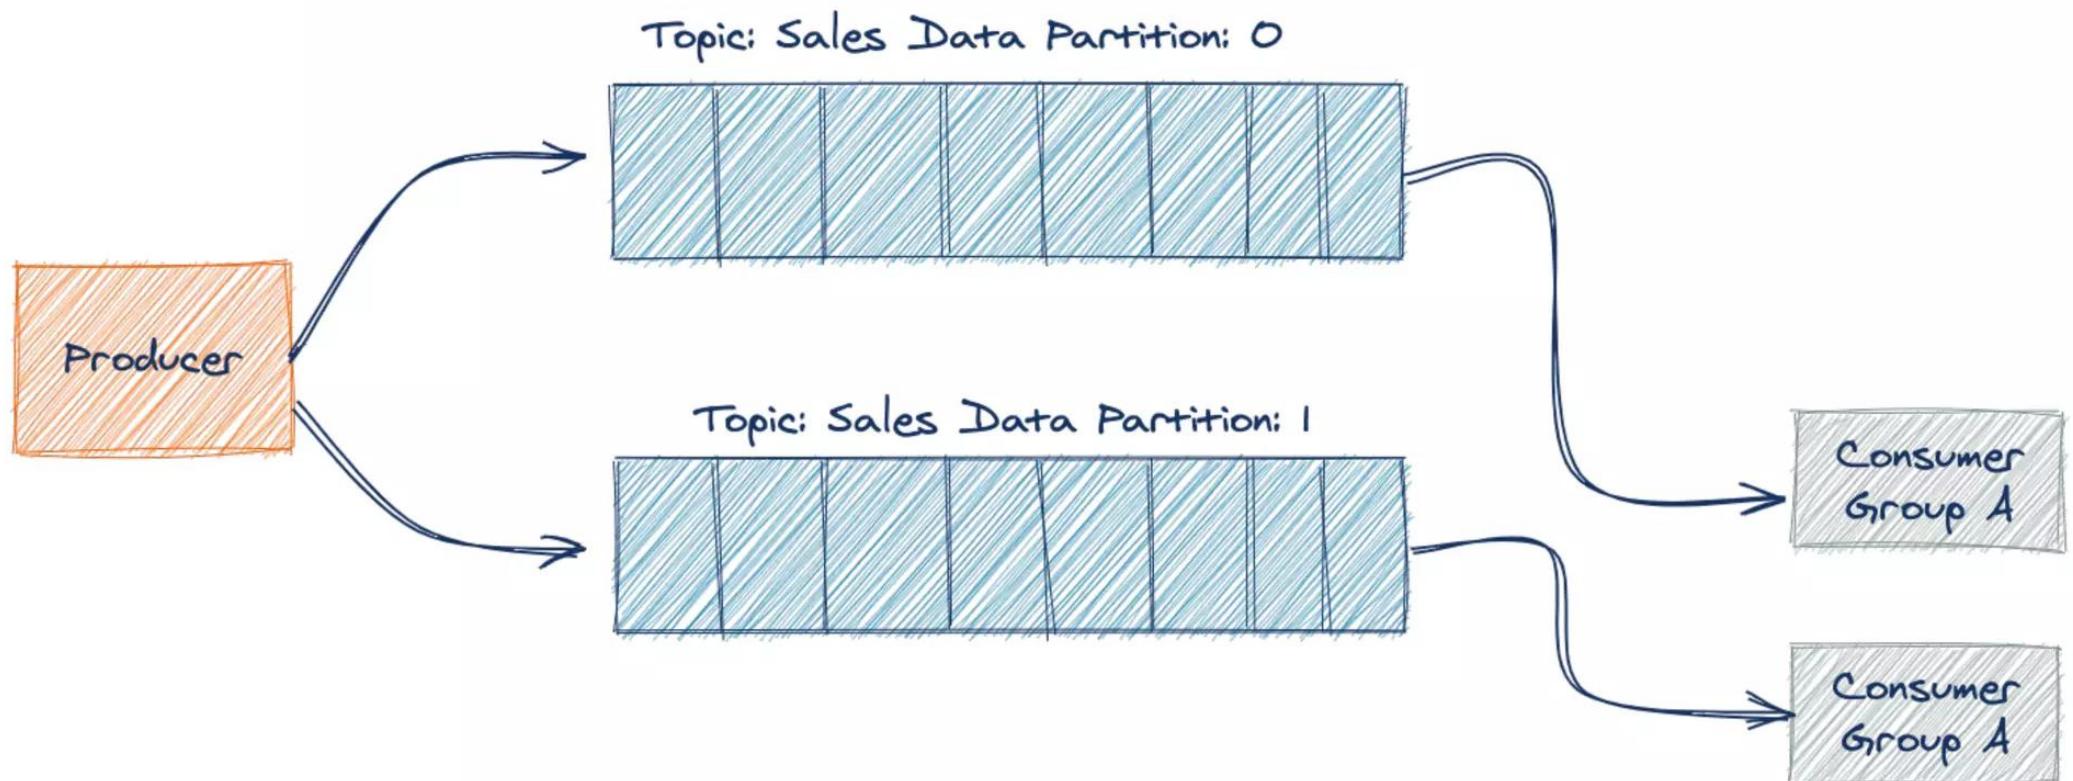


<https://cwiki.apache.org/confluence/display/KAFKA/KIP-500%3A+Replace+ZooKeeper+with+a+Self-Managed+Metadata+Quorum>

@bbejeck



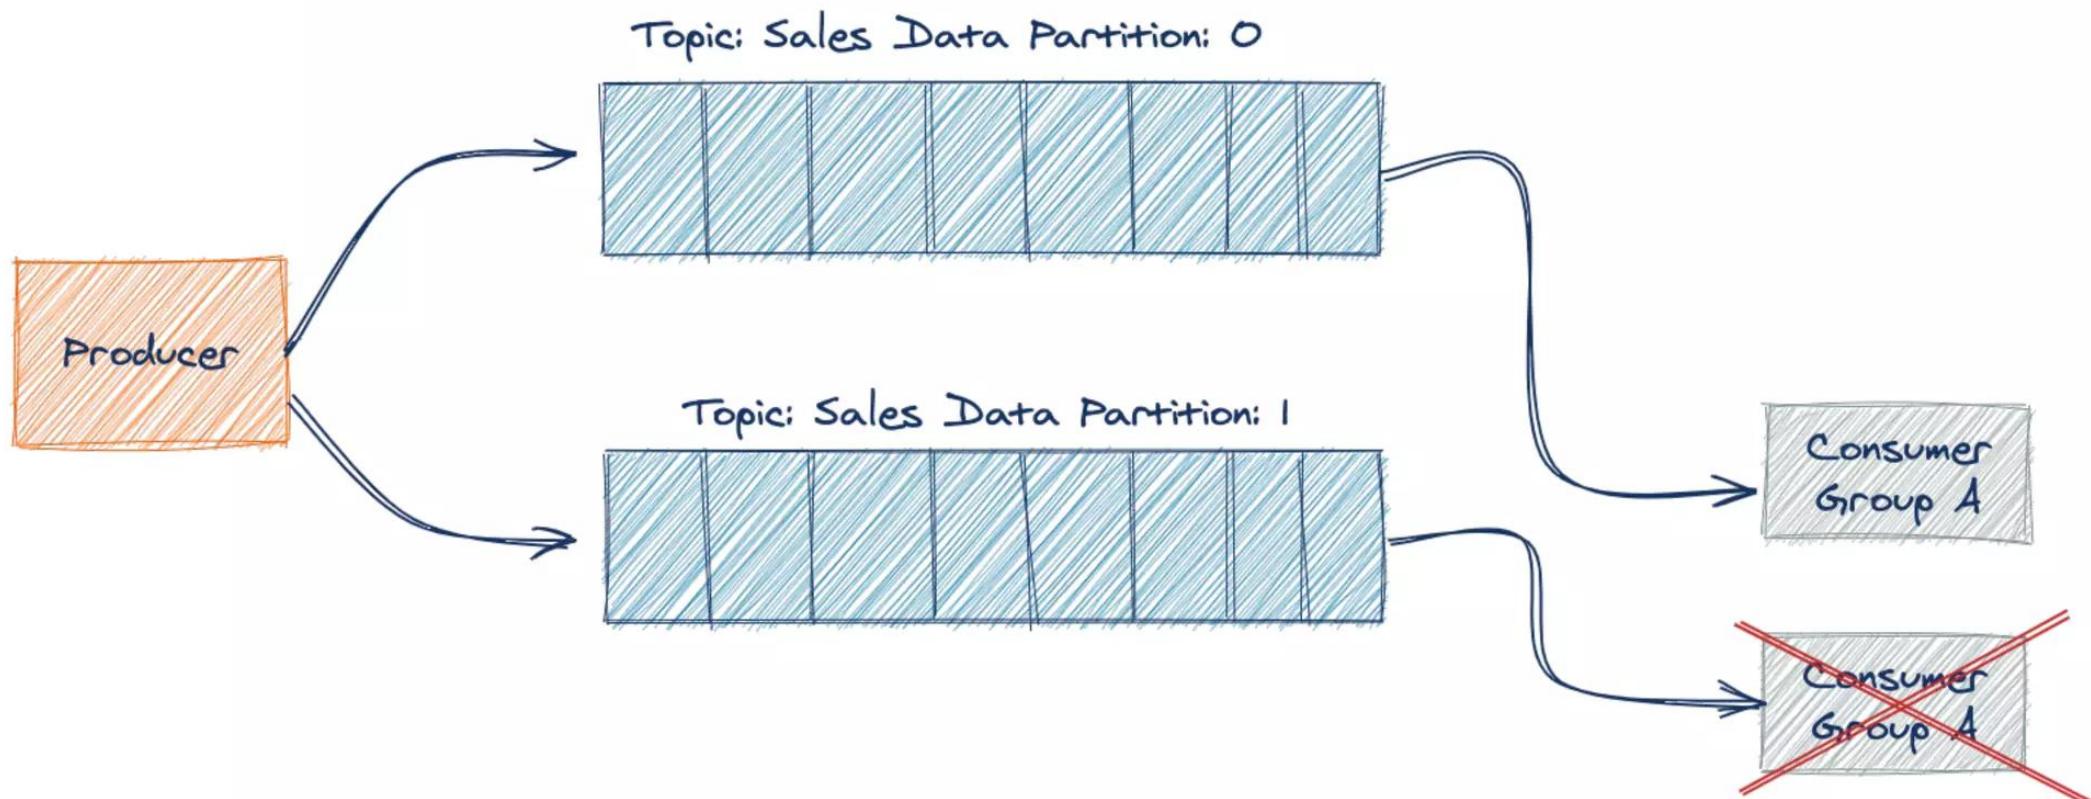
Kafka Overview - Clients



@bbejeck



Kafka Overview – Clients



@bbejeck

Kafka Overview – Clients



@bbejeck



Life before Kafka Streams

Life Before Kafka Streams – GroupBy Example



```
public static void main(String[] args) {  
    int counter = 0;  
    int sendInterval = 15;  
  
    Map<String, Integer> groupByCounts = new HashMap<>();  
  
    try(..consumer = new KafkaConsumer<>(consumerProperties());  
        ..producer = new KafkaProducer<>(producerProperties())){  
  
        consumer.subscribe(Arrays.asList("A", "B"));  
    }  
}
```

@bbejeck

Life Before Kafka Streams – GroupBy Example



```
while (true) {  
    ConsumerRecords<String, String> records = consumer.poll(Duration.ofSeconds(5));  
  
    for (ConsumerRecord<String, String> record : records) {  
        String key = record.key();  
        groupByCounts.compute(key, (k, v) -> v == null ? 1 : v + 1);  
    }  
}
```

@bbejeck

15

Life Before Kafka Streams – GroupBy Example



```
if(counter++ % sendInterval == 0) {  
    for(Entry<String, Integer> groupedEntry:groupByCounts.entrySet()) {  
        ProducerRecord<String, Integer> producerRecord =  
            new ProducerRecord<>("group-by-counts",  
                groupedEntry.getKey(),  
                groupedEntry.getValue());  
        producer.send(producerRecord);  
    }  
    consumer.commitSync();  
}
```

@bbejeck

Life Before Kafka Streams – GroupBy Example



```
if(counter++ % sendInterval == 0) {  
    for(Entry<String, Integer> groupedEntry:groupByCounts.entrySet()) {  
        ProducerRecord<String, Integer> producerRecord =  
            new ProducerRecord<>("group-by-counts",  
                groupedEntry.getKey(),  
                groupedEntry.getValue());  
        producer.send(producerRecord);  
    }  
    consumer.commitSync();  
}
```

@bbejeck



Sample Streams Application

```
...
stream = streamBuilder.stream(Arrays.asList("A", "B"))

stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .to("output-topic",
        Produced.with(Serdes.String(), Serdes.Long()))
```



Kafka Streams Topology



View the Topology

```
Topology topology = streamBuilder.build()
```

```
topology.describe().toString()
```

View the Topology



Topologies:

Sub-topology: 0

Source: KSTREAM-SOURCE-0000000000 (topics: [A, B])

--> KSTREAM-AGGREGATE-0000000001

Processor: KSTREAM-AGGREGATE-0000000001 (stores: [count-store])

--> KTABLE-TOSTREAM-0000000002

<-- KSTREAM-SOURCE-0000000000

Processor: KTABLE-TOSTREAM-0000000002 (stores: [])

--> KSTREAM-SINK-0000000003

<-- KSTREAM-AGGREGATE-0000000001

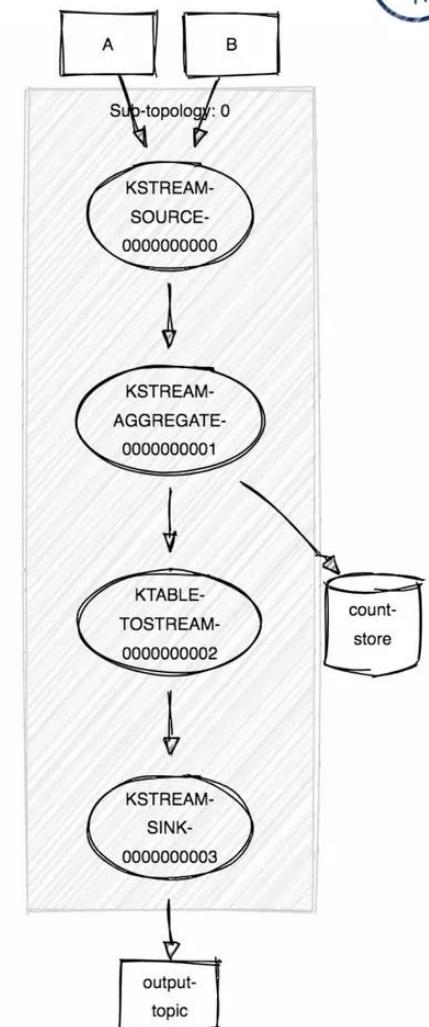
Sink: KSTREAM-SINK-0000000003 (topic: output-topic)

<-- KTABLE-TOSTREAM-0000000002

View Graphic Topology



- <https://github.com/zz85/kafka-streams-viz>



Kafka Streams Topology

View the Topology

```
stream = streamBuilder.stream(Arrays.asList("A", "B"));
```

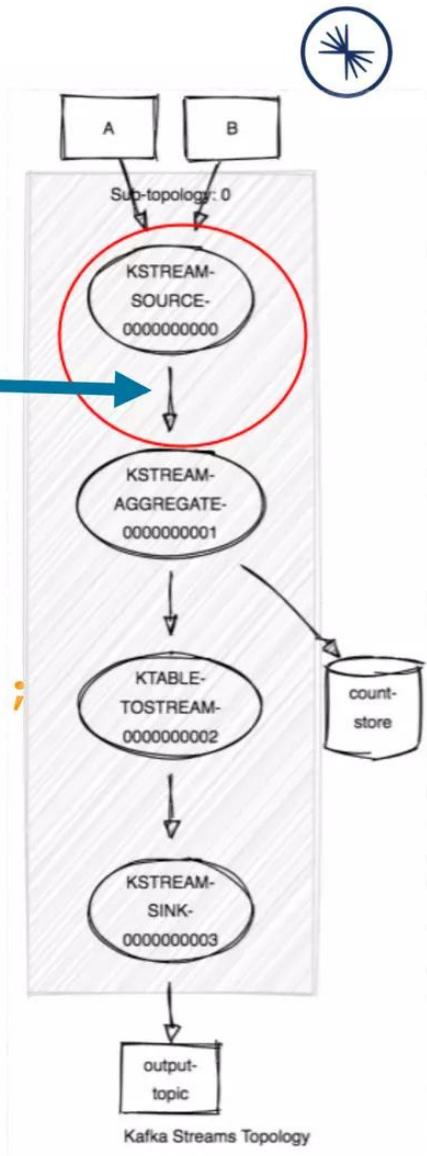
```
stream.groupByKey()
```

```
.count(Materialized.as("count-store"))
```

```
.toStream()
```

```
.to("output-topic",
```

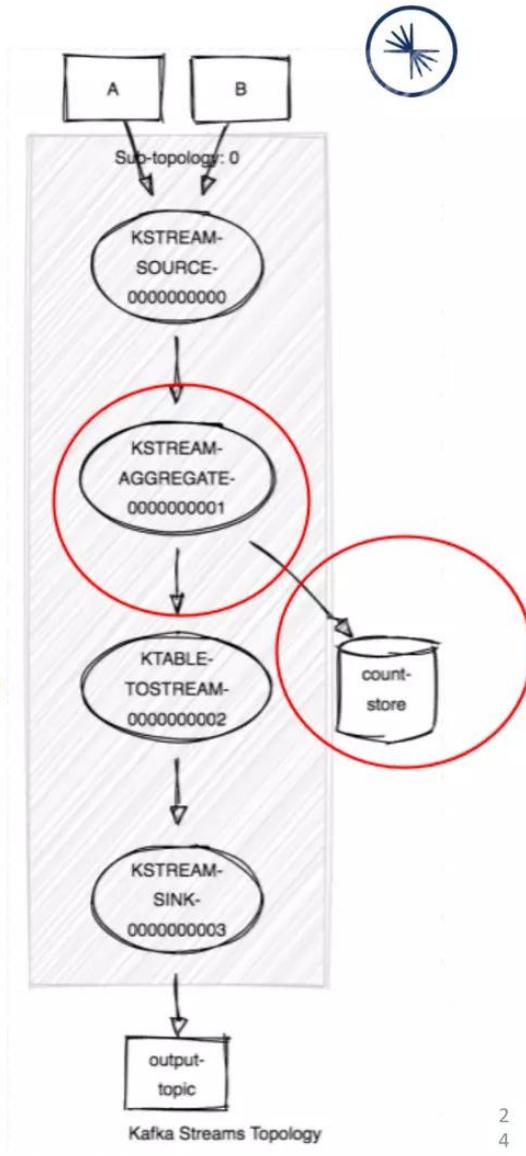
```
Produced.with(Serdes.String(), Serdes.Long())));
```



View the Topology

```
stream = builder.stream(Arrays.asList("A", "B"));

stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .to("output-topic",
        Produced.with(Serdes.String(), Serdes.Long()));
```

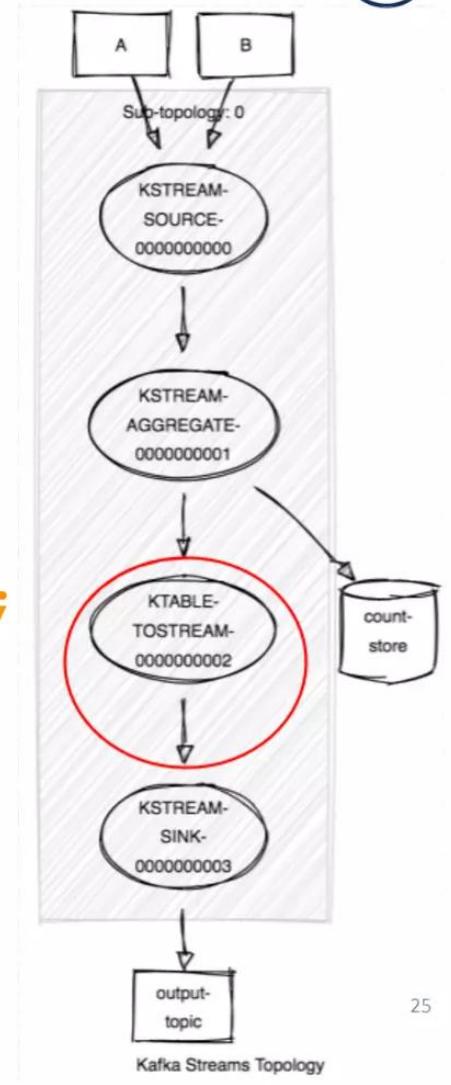




View the Topology

```
stream = builder.stream(Arrays.asList("A", "B"));

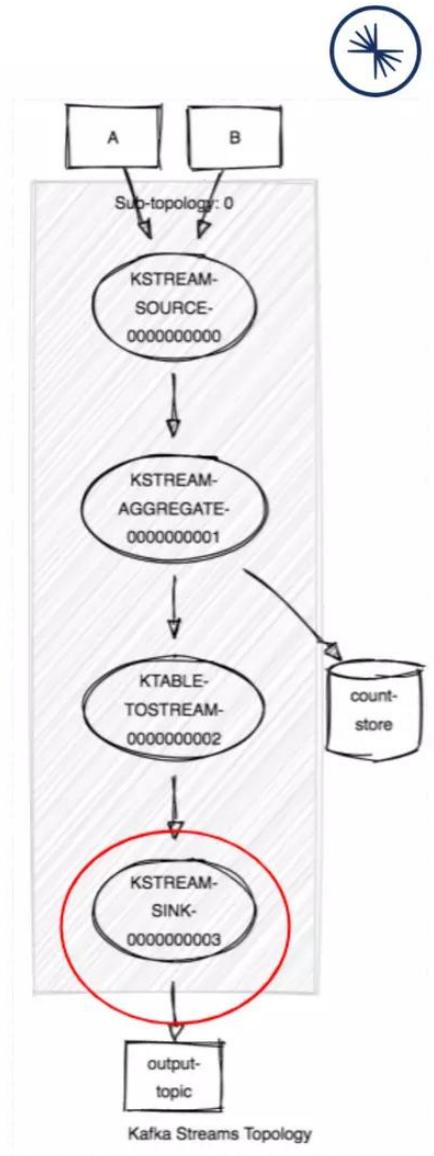
stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .to("output-topic",
        Produced.with(Serdes.String(), Serdes.Long()));
```



View the Topology

```
stream = builder.stream(Arrays.asList("A", "B"));

stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .to("output-topic",
Produced.with(Serdes.String(), Serdes.Long()));
```

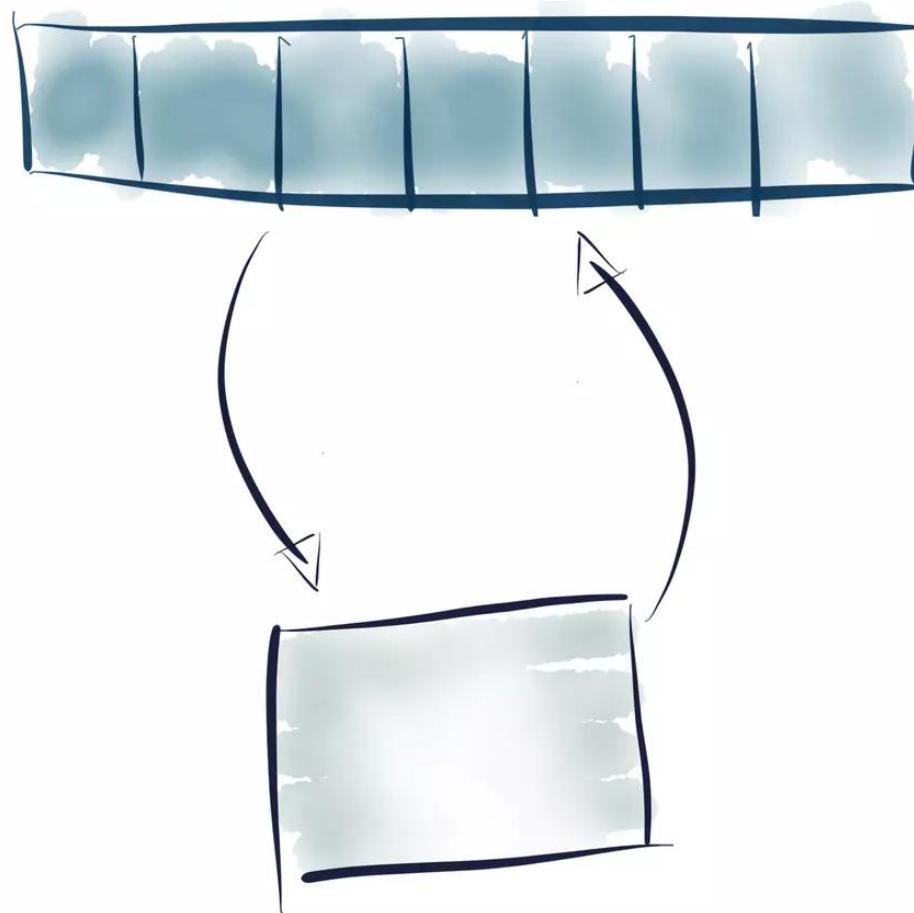
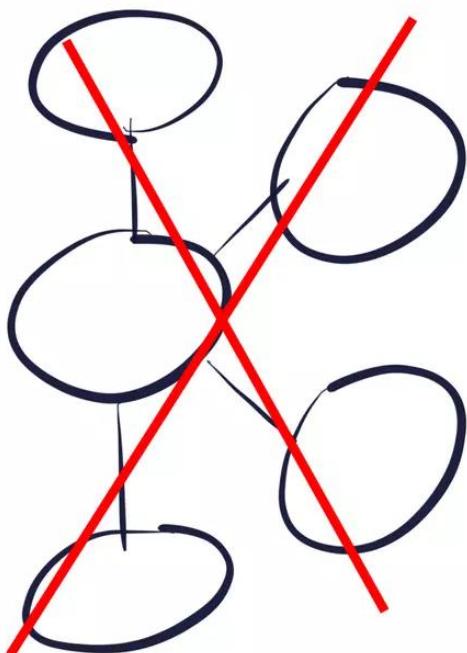




Kafka Streams Concepts/Architecture



Where does Kafka Streams run?



@bbejeck



Record Streams - KStream

Record Stream - Each key-value pair independent regardless of key



Update Stream - KTable



Update Stream - key-value pairs with the same key as previous records are updates to previous records

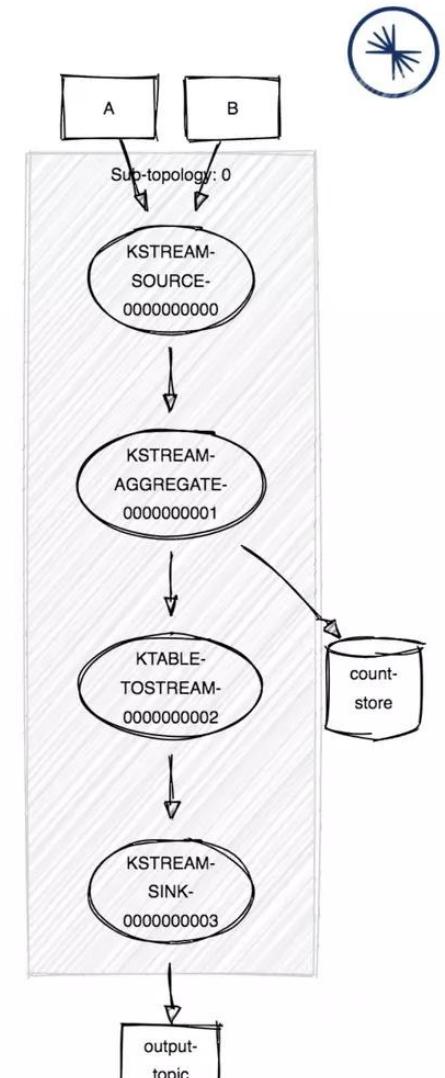
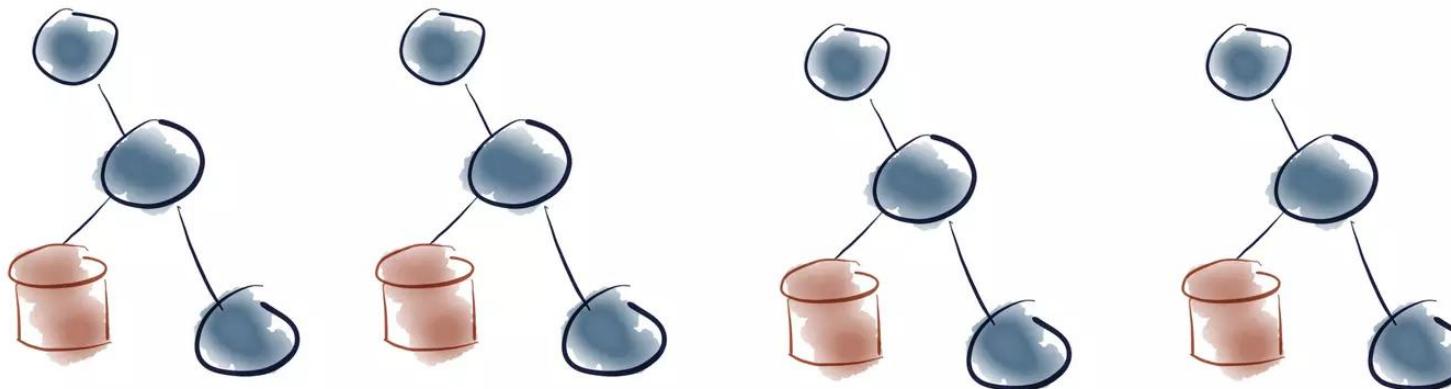


Kafka Streams Tasks

Topic A
A0 A1 A2 A3

Topic B
B0 B1

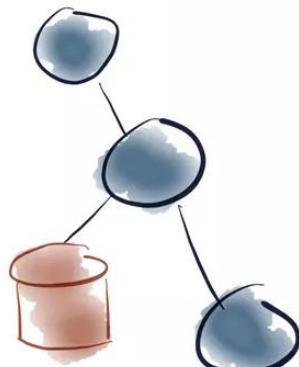
Number tasks = $\max(4, 2)$



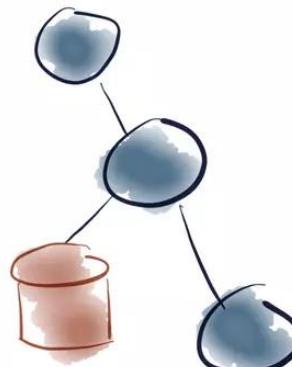
Kafka Streams Tasks

Topic A
A0 A1 A2 A3

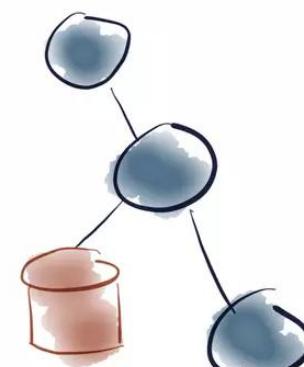
O_0
A0 B0



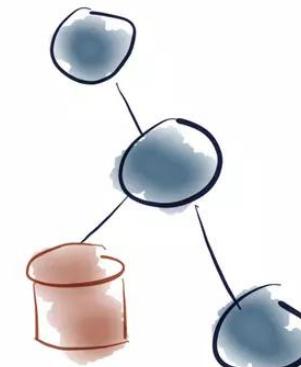
O_1
A1 B1



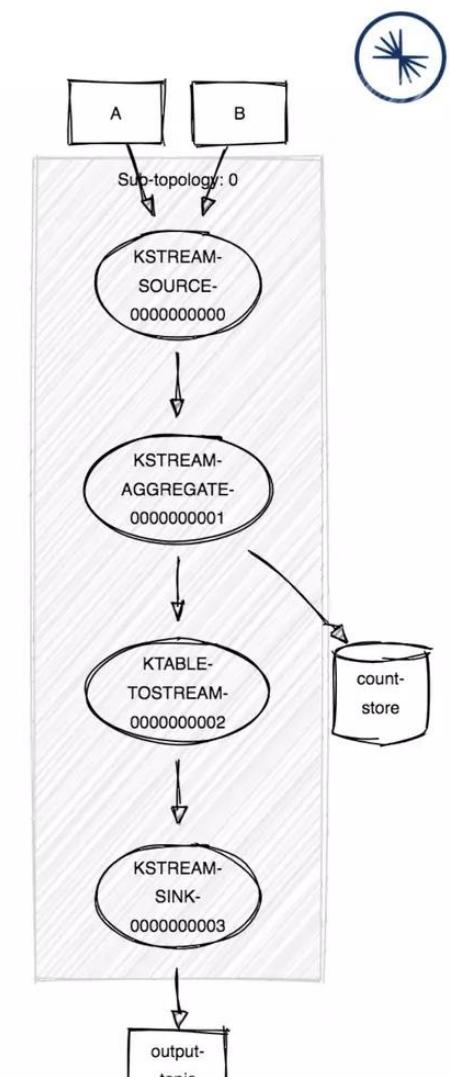
O_2
A2



O_3
A3

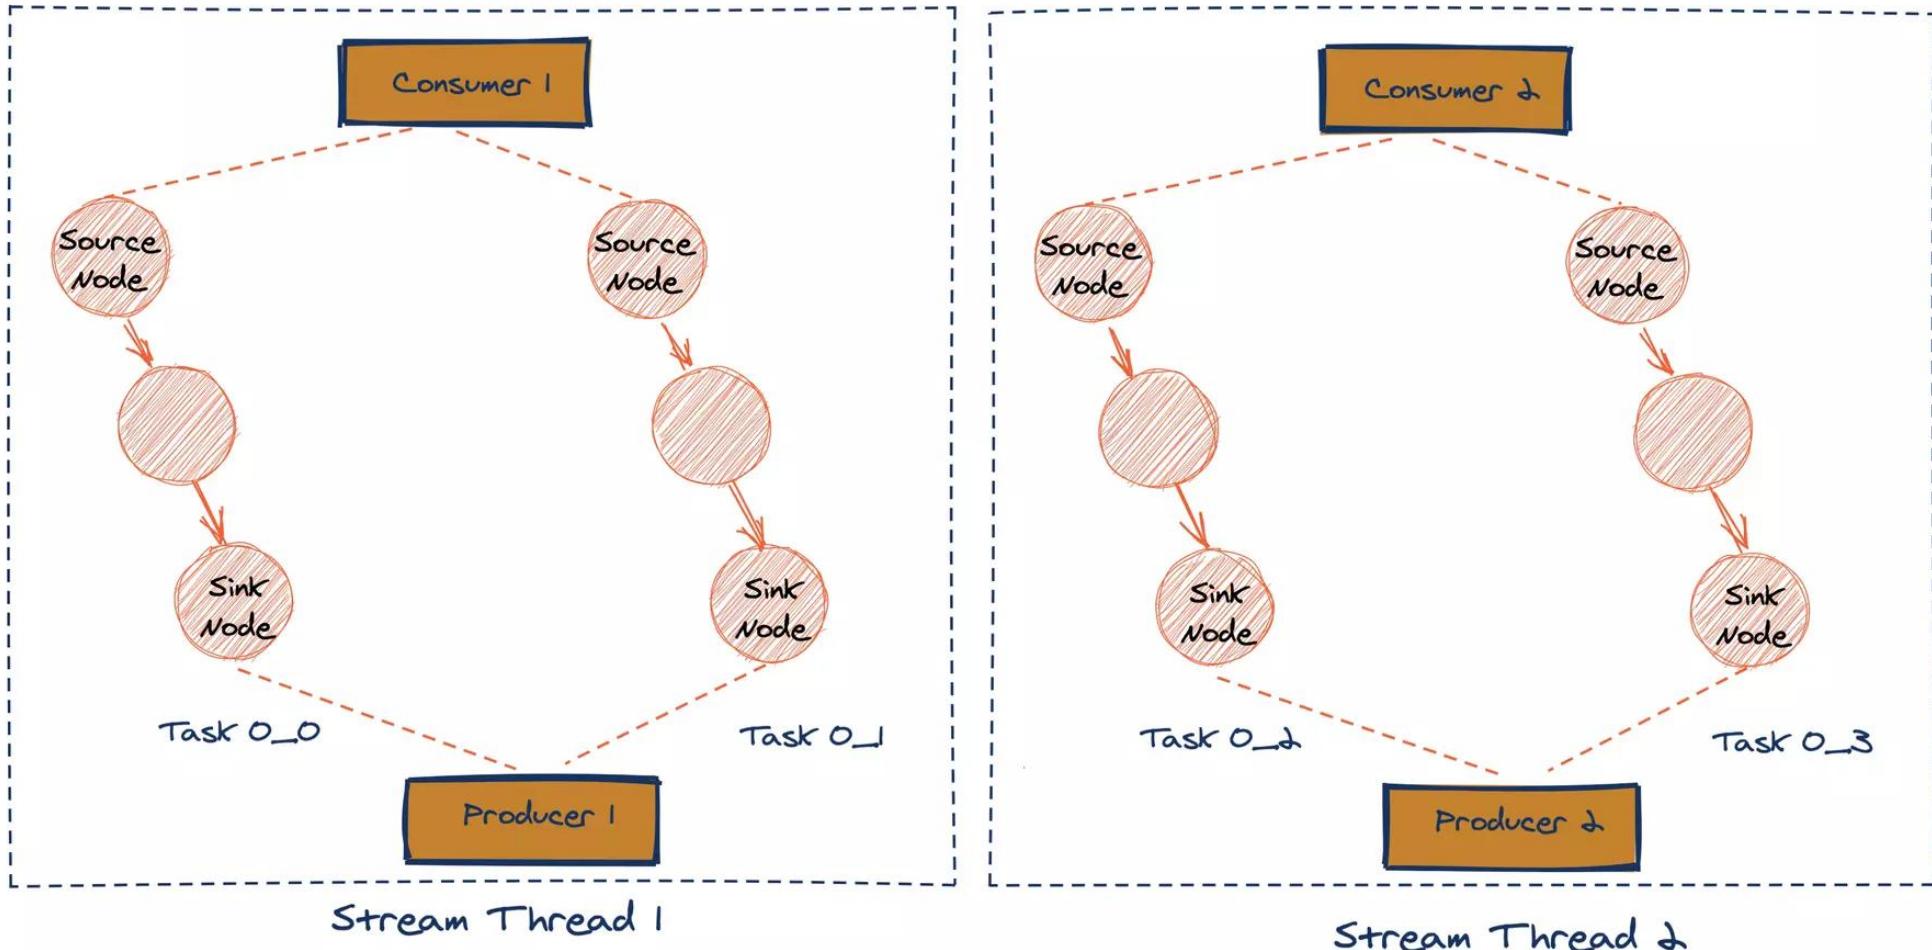


Topic B
B0 B1





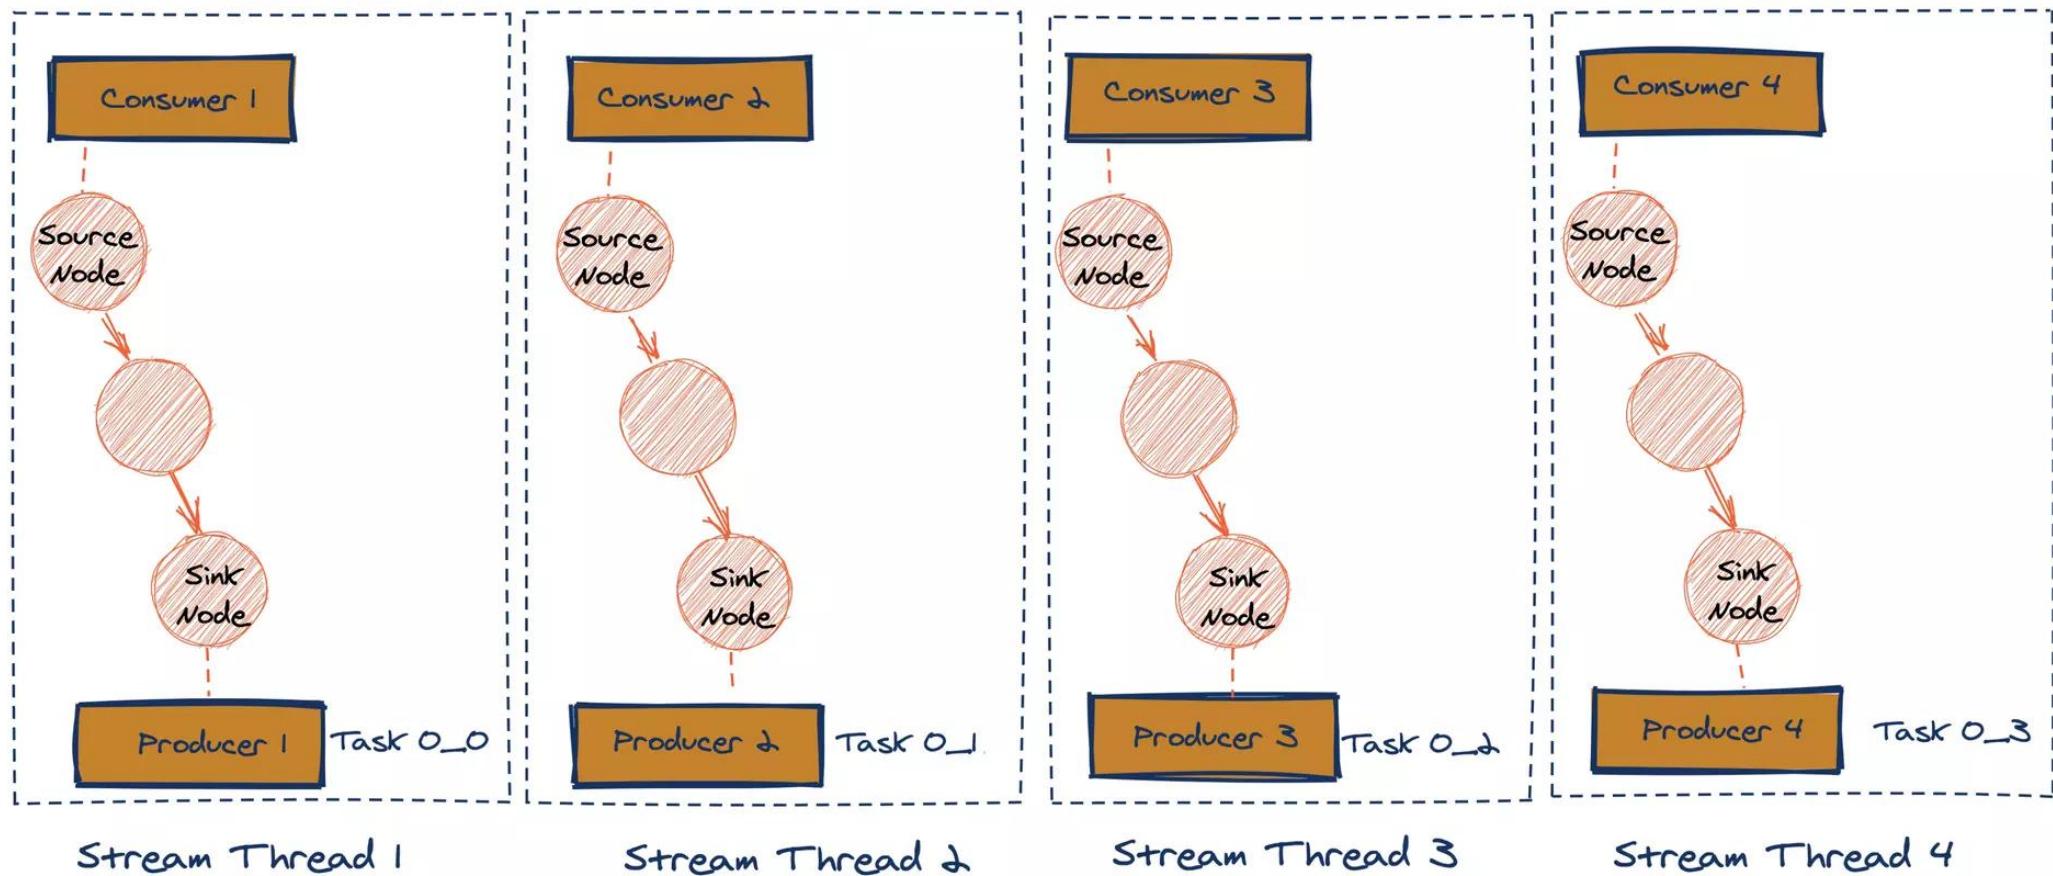
Kafka Streams Architecture - Threading



@bbejeck



Kafka Streams Architecture - Threading



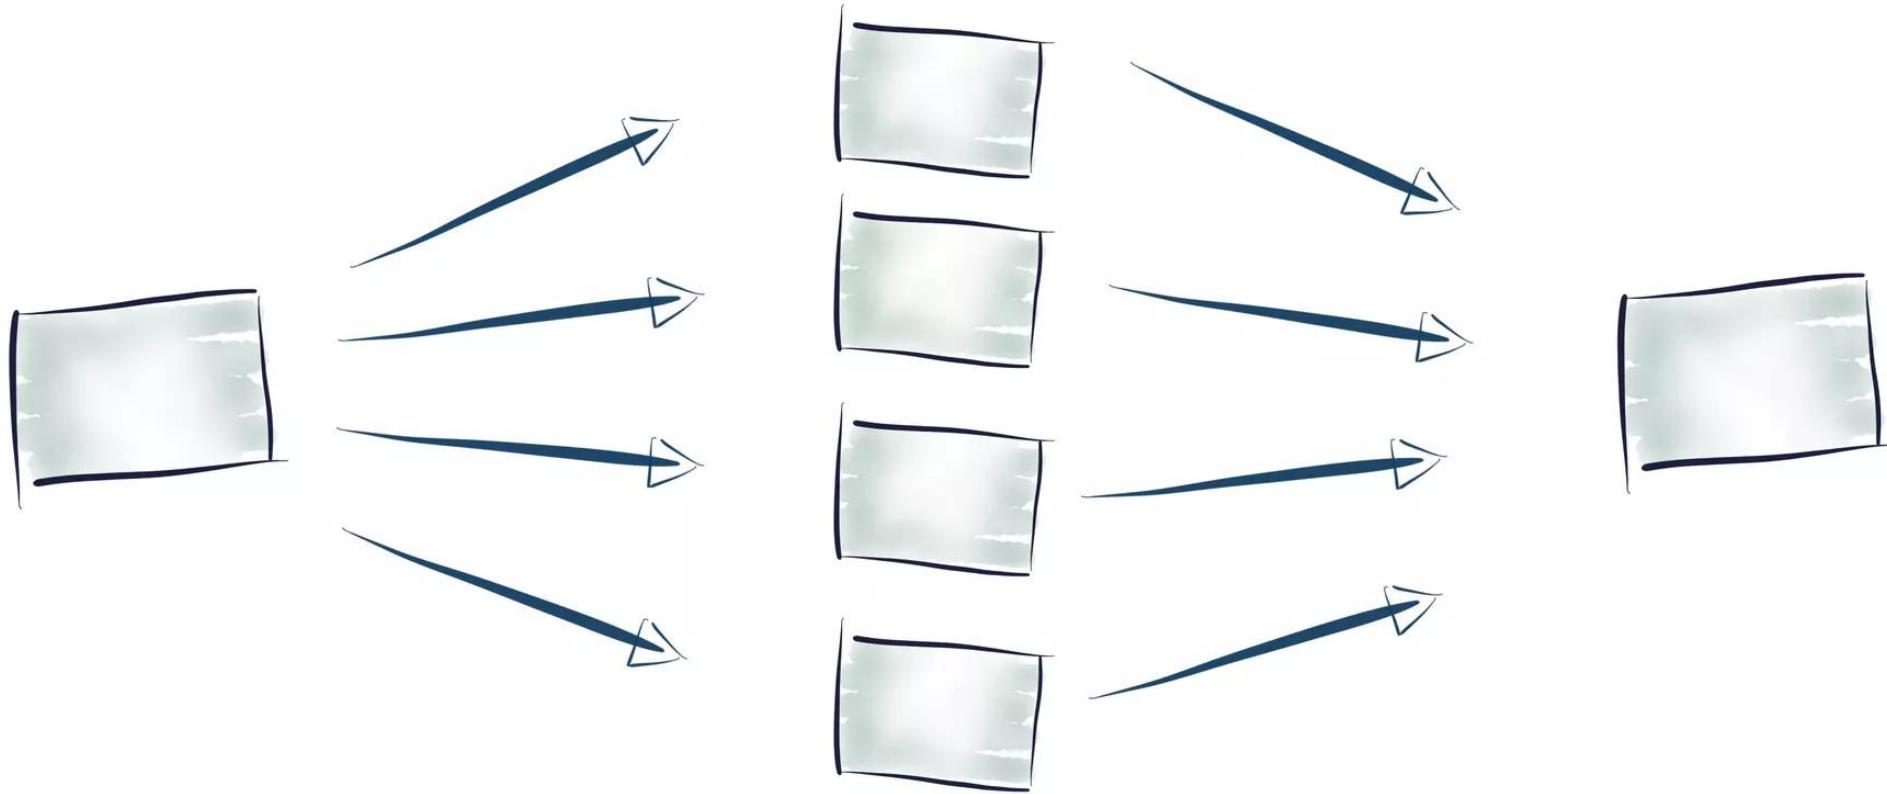
@bbejeck

Kafka Streams – Deployment Consideration



@bbejeck

Kafka Streams Dynamic Expansion/Contraction



@bbejeck



Kafka Streams API/Features

Kafka Streams – DSL API



```
...
stream = streamBuilder.stream(Arrays.asList("A", "B"))

stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .to("output-topic",
        Produced.with(Serdes.String(), Serdes.Long()))
```

@bbejeck

38

Processor API - WordCountProcessorDemo



```
final Topology builder = new Topology();

builder.addSource("Source", "streams-plaintext-input");

builder.addProcessor("Process", new MyProcessorSupplier(), "Source");
builder.addStateStore(Stores.keyValueStoreBuilder(
    Stores.inMemoryKeyValueStore("Counts"),
    Serdes.String(),
    Serdes.Integer()),
    "Process");

builder.addSink("Sink", "streams-wordcount-processor-output", "Process");
```

@bbejeck

Processor API - WordCountProcessorDemo



```
final Topology builder = new Topology();

builder.addSource("Source", "streams-plaintext-input");

builder.addProcessor("Process", new MyProcessorSupplier(), "Source");
builder.addStateStore(Stores.keyValueStoreBuilder(
    Stores.inMemoryKeyValueStore("Counts"),
    Serdes.String(),
    Serdes.Integer()),
    "Process");

builder.addSink("Sink", "streams-wordcount-processor-output", "Process");
```

@bbejeck

Processor API - WordCountProcessorDemo



```
final Topology builder = new Topology();

builder.addSource("Source", "streams-plaintext-input");

builder.addProcessor("Process", new MyProcessorSupplier(), "Source");
builder.addStateStore(Stores.keyValueStoreBuilder(
    Stores.inMemoryKeyValueStore("Counts"),
    Serdes.String(),
    Serdes.Integer()),
    "Process");

builder.addSink("Sink", "streams-wordcount-processor-output", "Process");
```

@bbejeck

Processor API - WordCountProcessorDemo



```
final Topology builder = new Topology();

builder.addSource("Source", "streams-plaintext-input");

builder.addProcessor("Process", new MyProcessorSupplier(), "Source");
builder.addStateStore(Stores.keyValueStoreBuilder(
    Stores.inMemoryKeyValueStore("Counts"),
    Serdes.String(),
    Serdes.Integer()),
    "Process");

builder.addSink("Sink", "streams-wordcount-processor-output", "Process");
```

@bbejeck

Kafka Streams – Combining DSL and Processor API



```
stream.groupByKey()
    .count(Materialized.as("count-store"))
    .toStream()
    .transform(MyTransformSupplier())
    .to("output-topic",
        Produced.with(Serdes.String(), Serdes.Long())))
  
  
transformValues, flatTransform, process (terminal node)
```

Kafka Streams – Stateless Operations



```
KStream.filter(...)  
KStream.map(...)  
KStream.mapValues(...)  
KStream.transform(...)  
..  
KTable.filter(...)  
KTable.groupBy(...)  
KTable.mapValues(...)
```

<https://docs.confluent.io/currentstreams/developer-guide/dsl-api.html#stateless-transformations>

@bbejeck



Stateful Operations

KStream.groupByKey().count()

KStream.groupByKey().reduce(..)

KStream.groupByKey().aggregate(..)

..

KTable.groupBy(..).count(..)

KTable.groupBy(..).reduce(..)

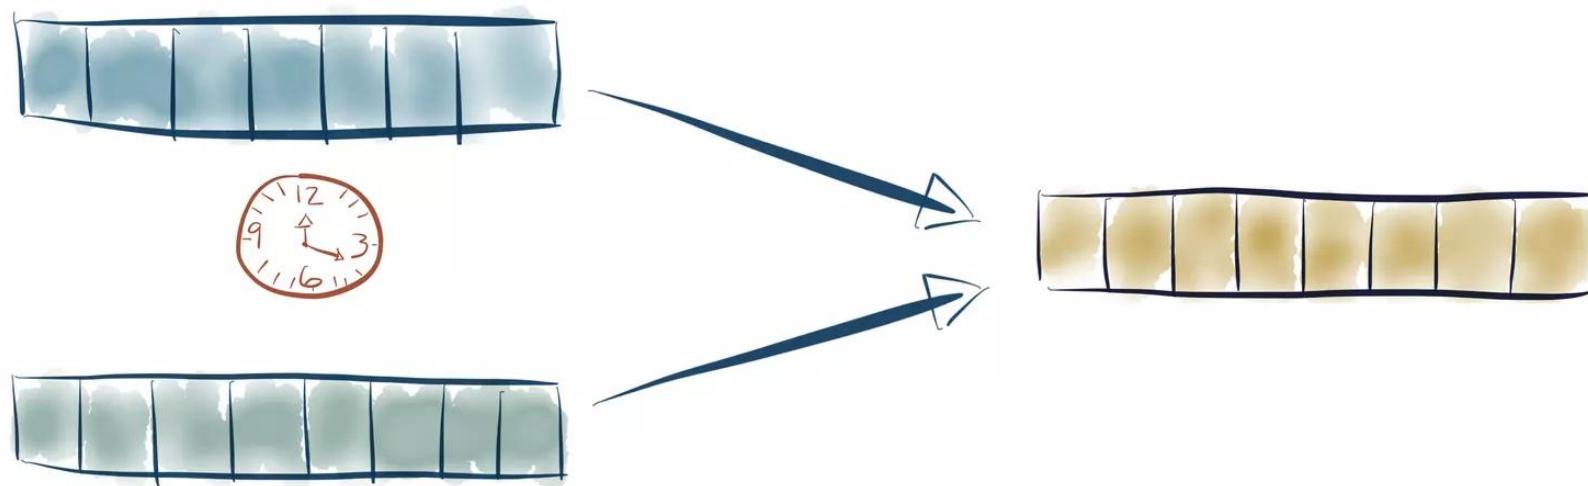
KTable.groupBy(..).aggregate(..)

<https://docs.confluent.io/currentstreams/developer-guide/dsl-api.html#stateful-transformations>

@bbejeck



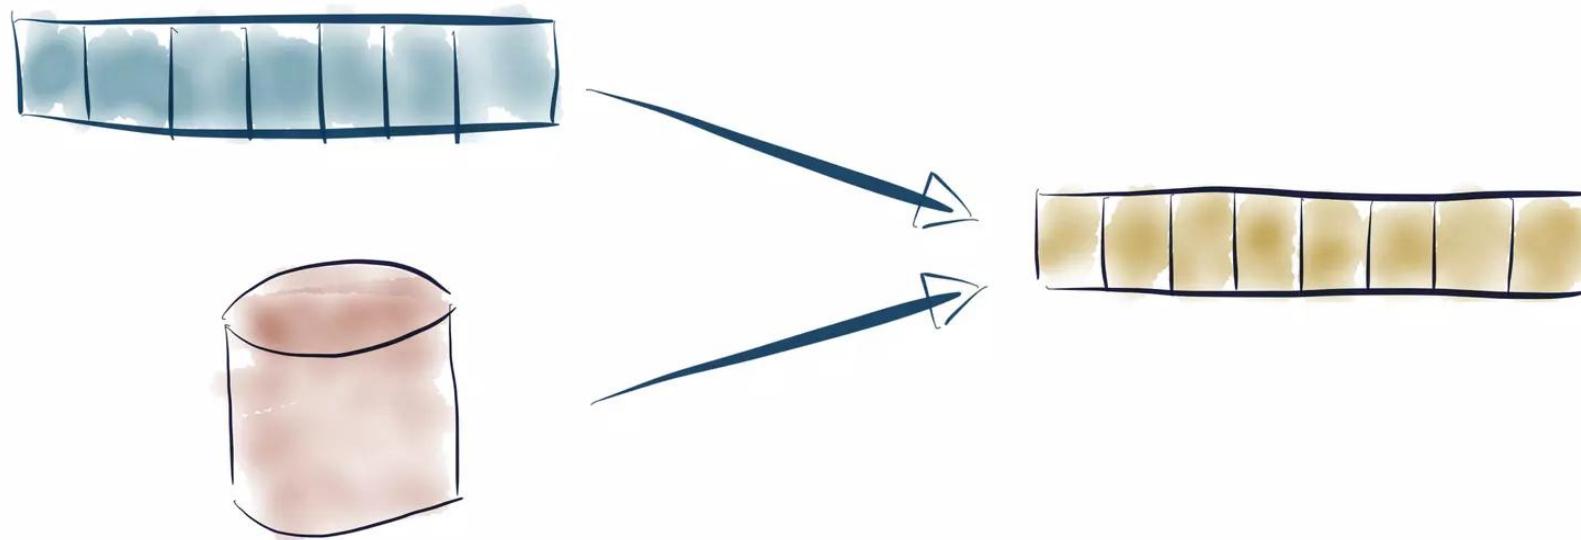
Stateful Operations – Stream-Stream Join



@bbejeck



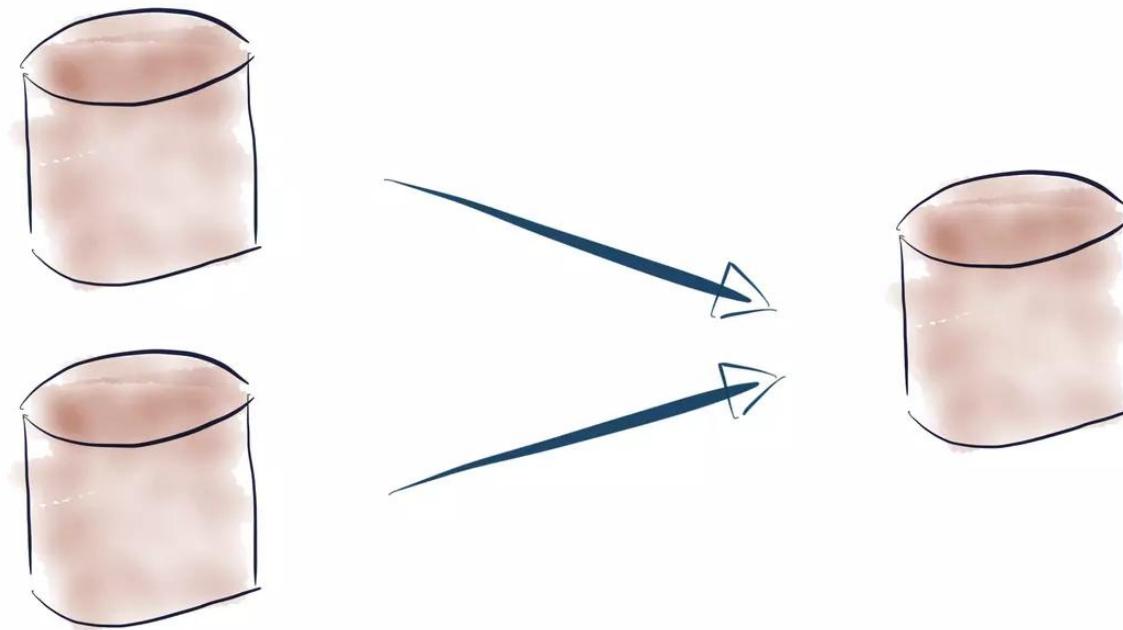
Stateful Operations – Stream-Table Join



@bbejeck



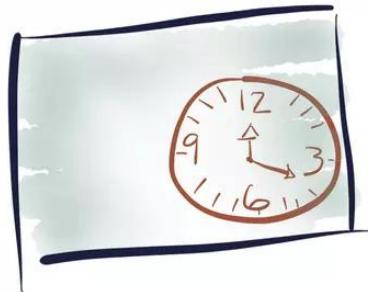
Stateful Operations – Table -Table Join



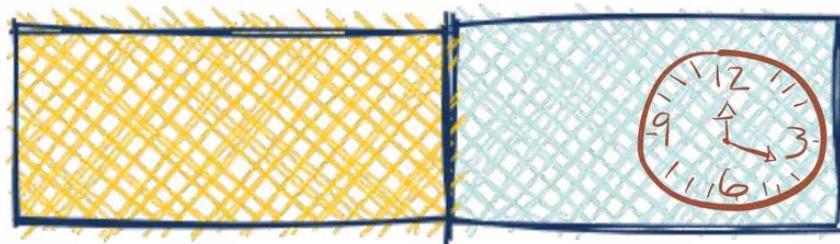
@bbejeck



Time in Kafka Streams



Wall-clock time is the time the application is processing the record



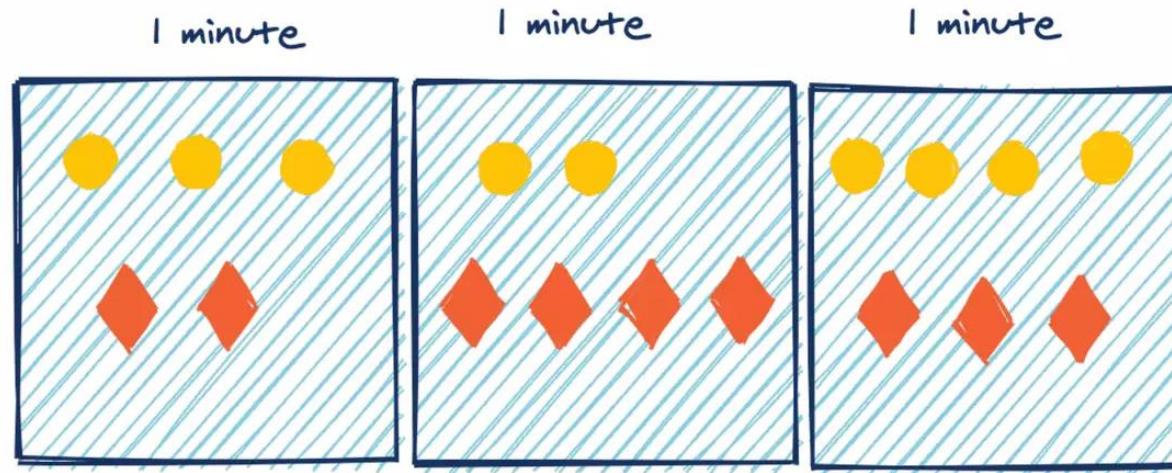
Time used by Kafka Streams is the timestamp of the record – pulled out by the `TimestampExtractor`.



Windowing – Tumbling Windows

```
Duration windowSize = Duration.ofMinutes(1);  
stream.groupByKey().count().windowedBy(TimeWindows.of(windowSize));
```

Tumbling windows - Non-overlapping events bounded by start and end time



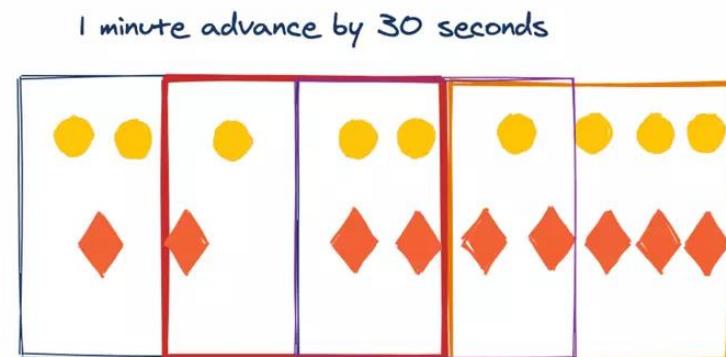
@bbejeck



Windowing – Hopping Windows

```
Duration windowSize = Duration.ofMinutes(1);  
Duration advance = Duration.ofSeconds(30);  
stream.groupByKey().count().windowedBy(TimeWindows.of(windowSize)).advanceBy(advance);
```

Hopping windows - overlapping fixed sized window bounded by start end time with an incremental update

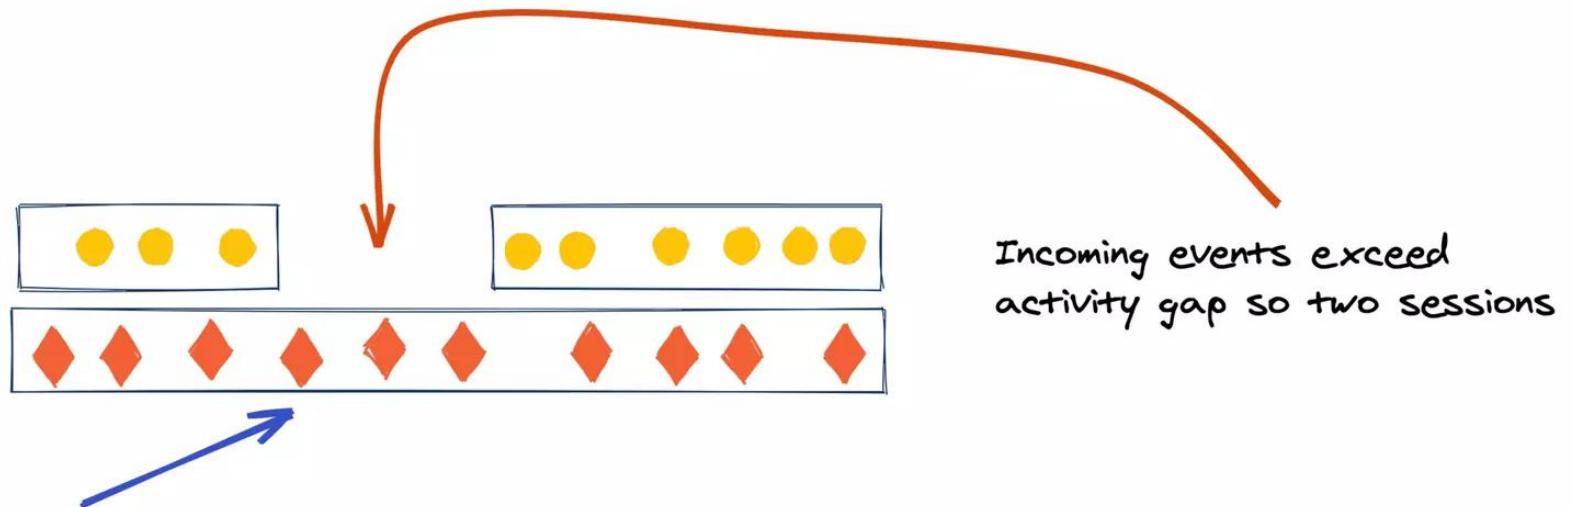




Windowing – Session Windows

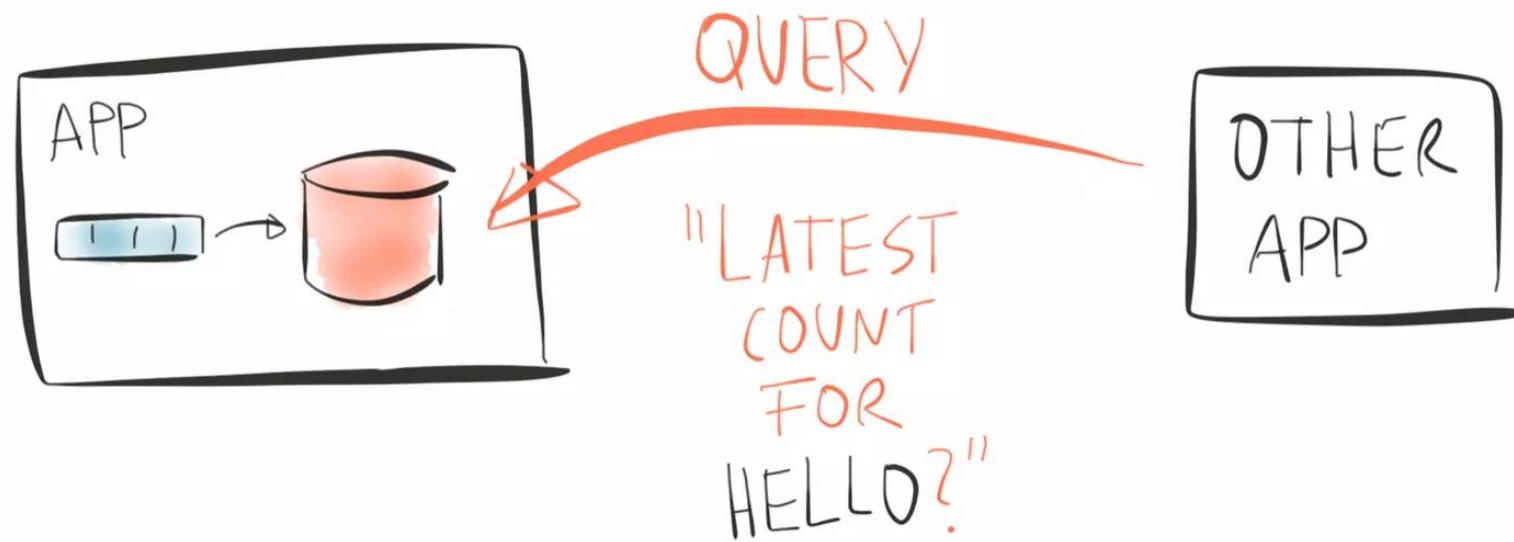
```
Duration inactivityGap = Duration.ofMinutes(3);  
stream.groupByKey().count().windowedBy(SessionWindows.of(inactivityGap));
```

Session Windows - Dynamic sized windows based on behavior inactivity defines a session



@bbejeck

Interactive Queries



Resources



- **Kafka Tutorials** - <https://kafka-tutorials.confluent.io/>
- **Confluent Developer** - developer.confluent.io
- **Examples** - <https://github.com/confluentinc/kafka-streams-examples/>
- **Documentation** - <https://docs.confluent.io/currentstreams>

@bbejeck

Thank you!

@bbejeck
bill@confluent.io



cnfl.io/meetups



cnfl.io/blog



cnfl.io/slack



Learn Kafka!

Confluent Developer
developer.confluent.io

