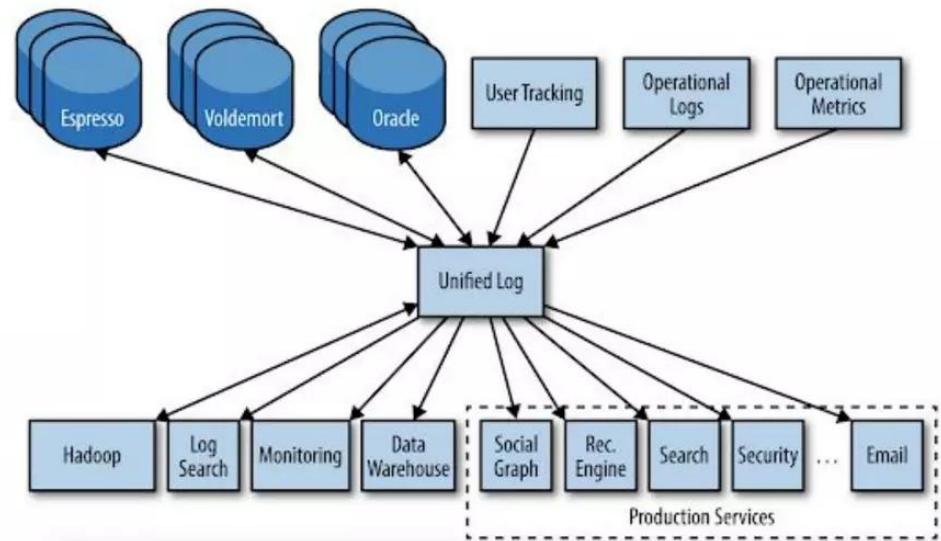
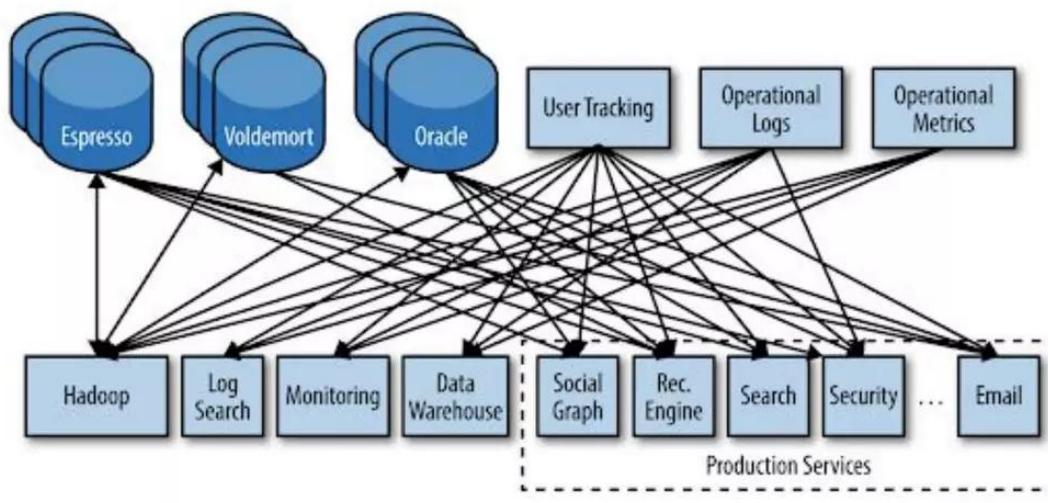


Kafka 102 - Applied

Alex Woolford, Systems engineer

why Kafka was created



Source: 10 logs, Jay Kreps, 2014

Cuatro Commas Club , , , ,

LinkedIn
4.5 trillion messages/day

Uber
1+ trillion messages/day

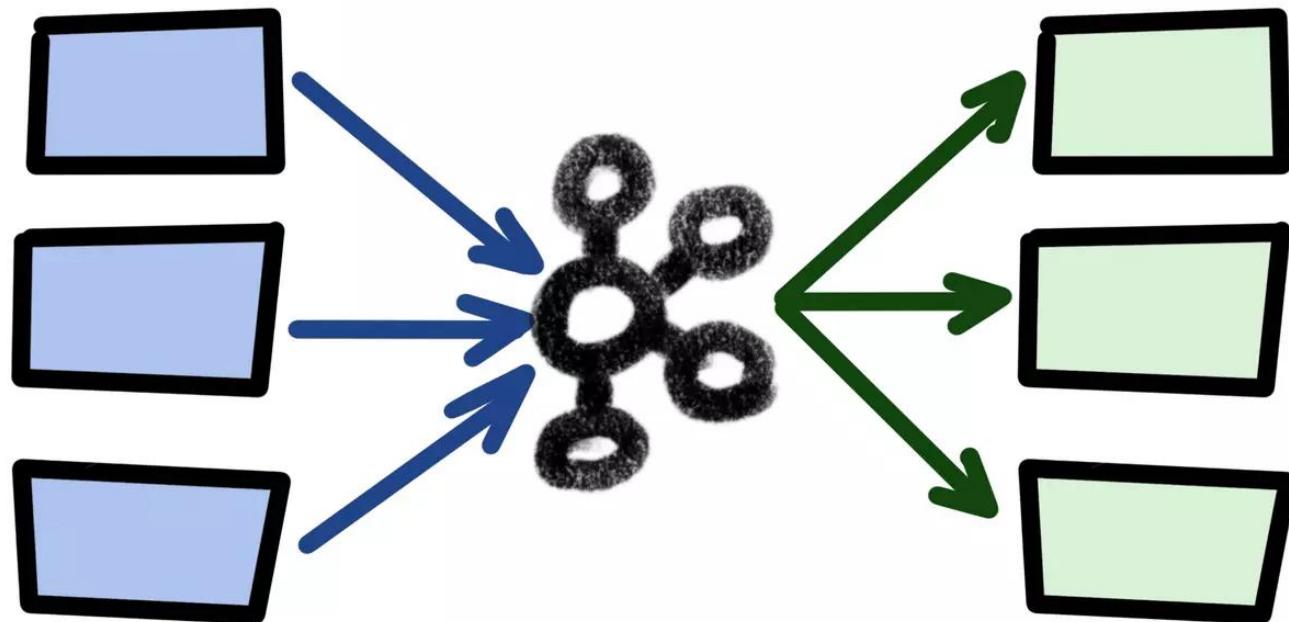
NETFLIX
2 trillion messages/day



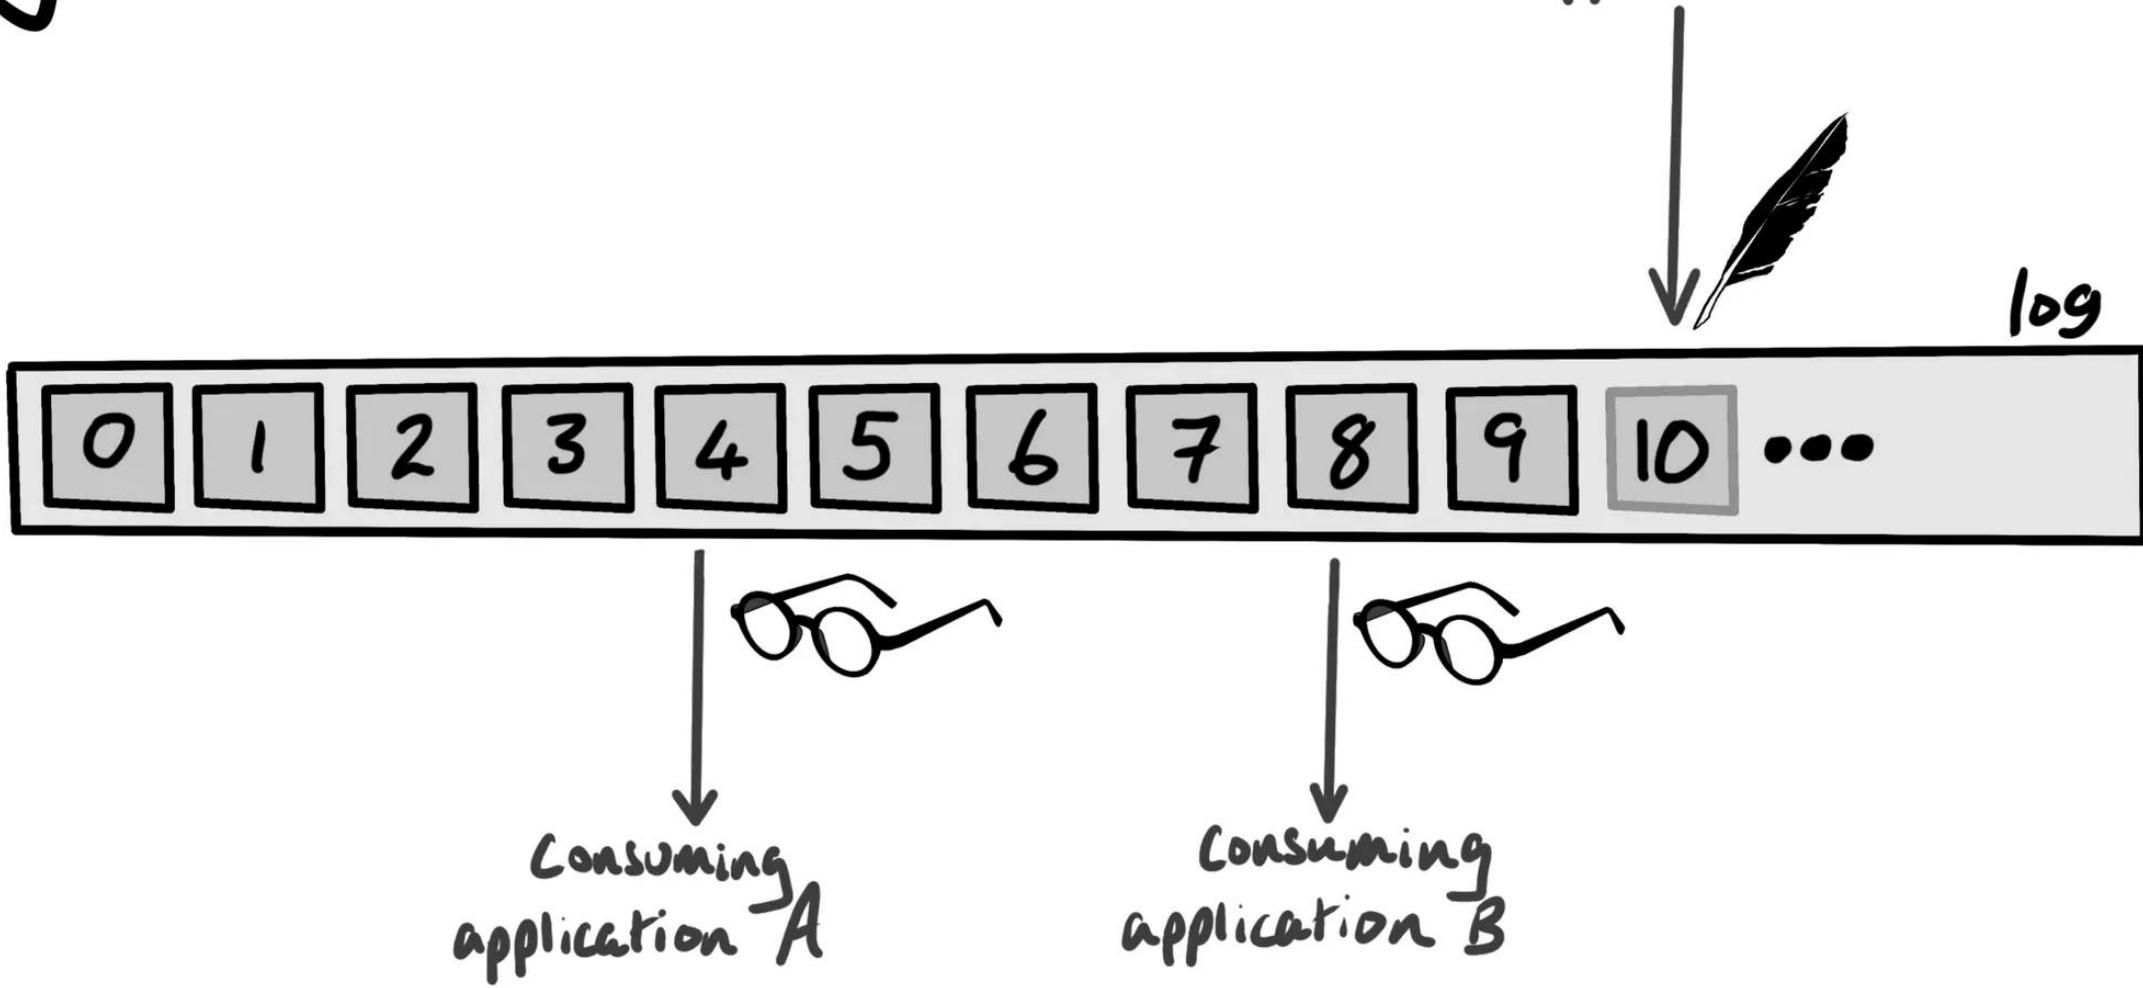
Microsoft
3 trillion messages/day

Kafka ❤️'s fan out

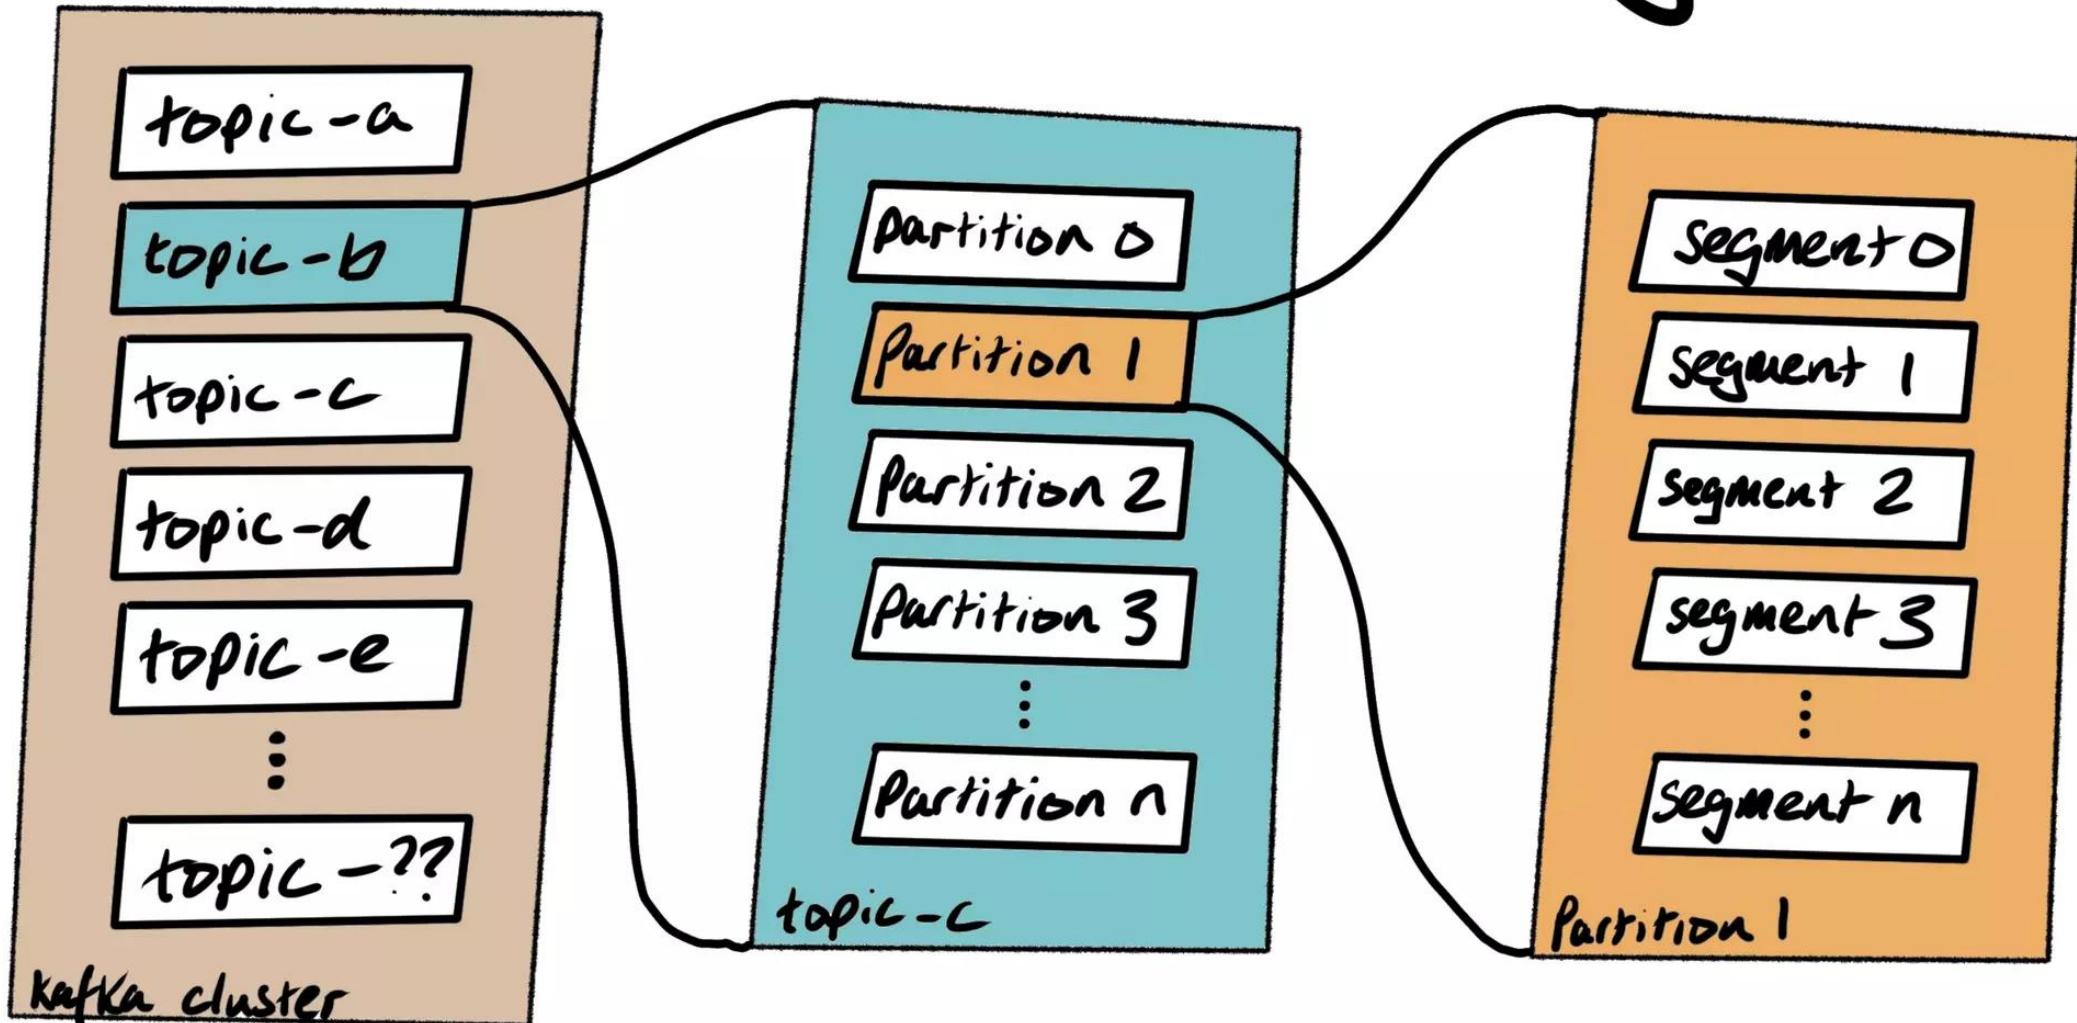
fan in fan out



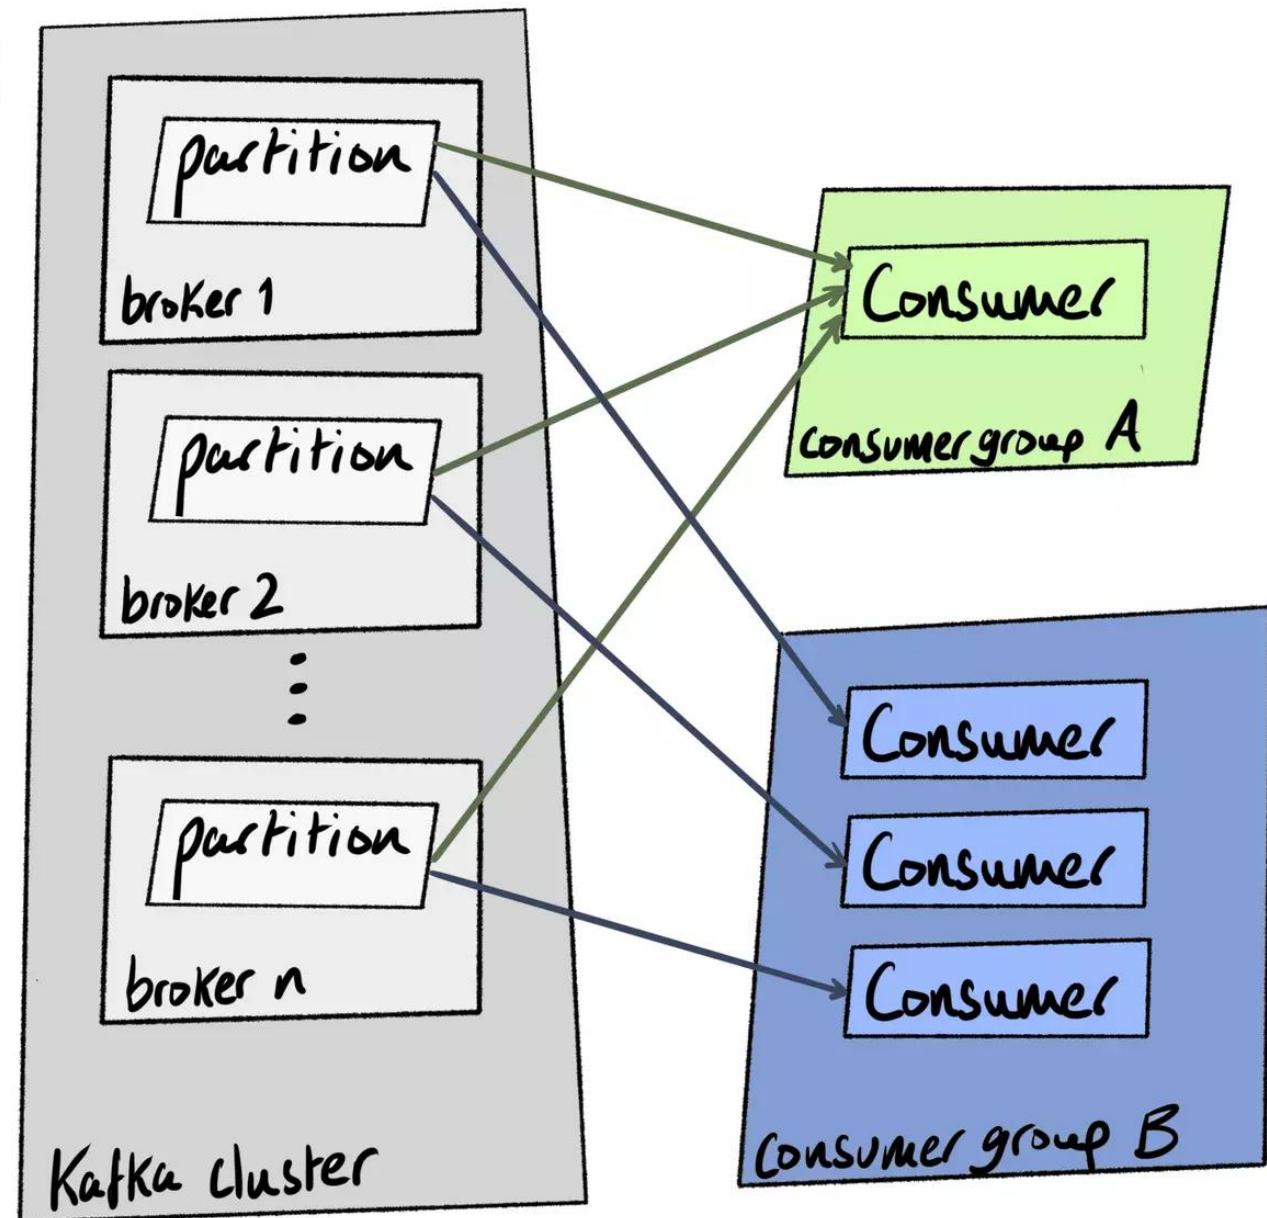
log structured data flow



topics, partitions, and segments



Consumer groups



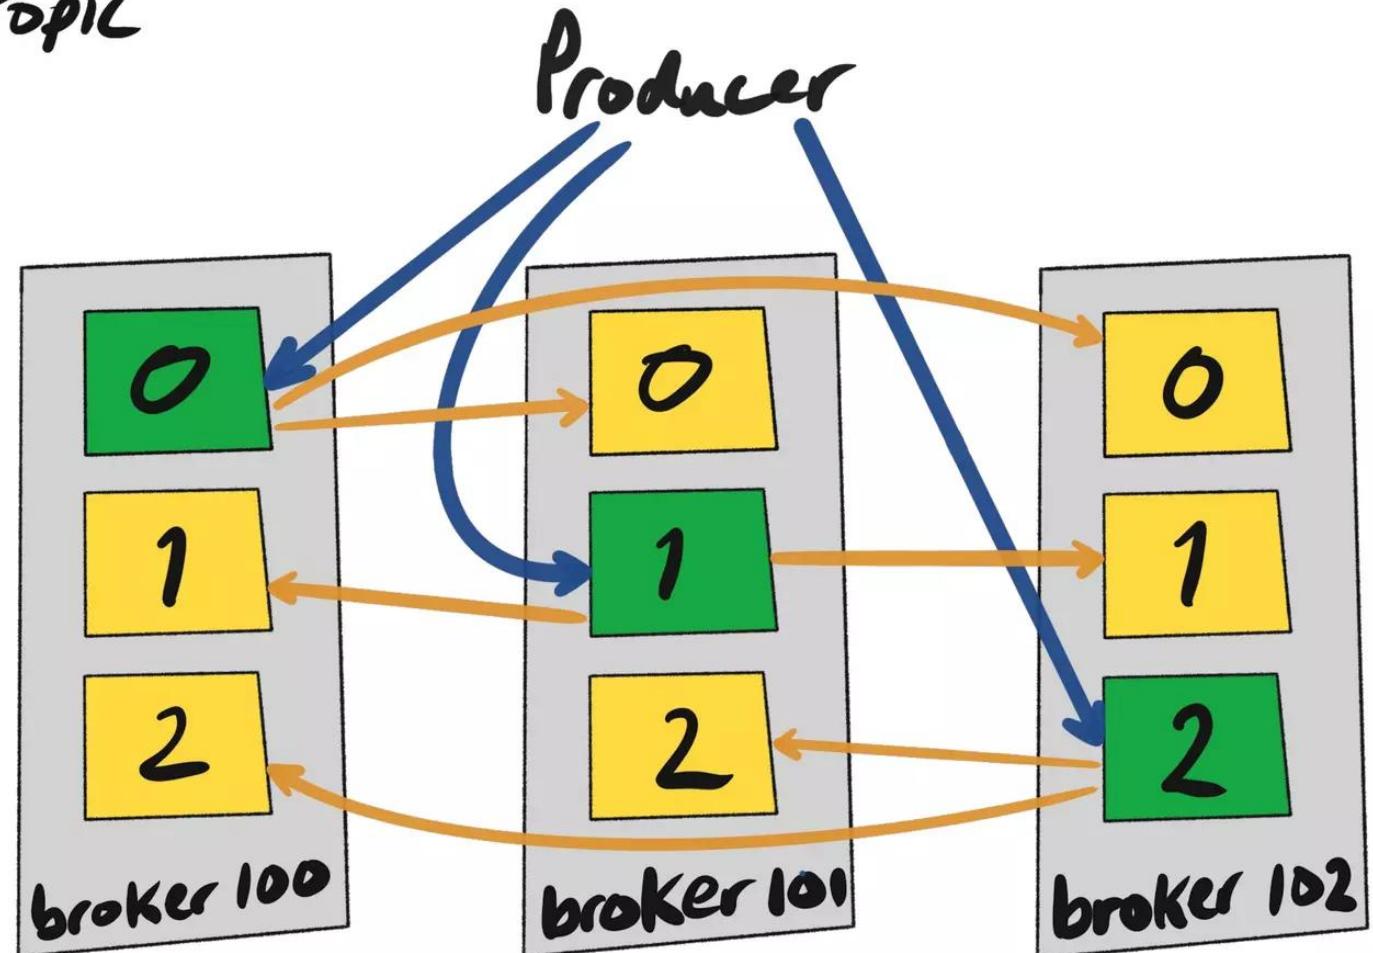
replication and in-Sync replicas (ISR)

topic : my-topic
partitions: 3
replicas: 3

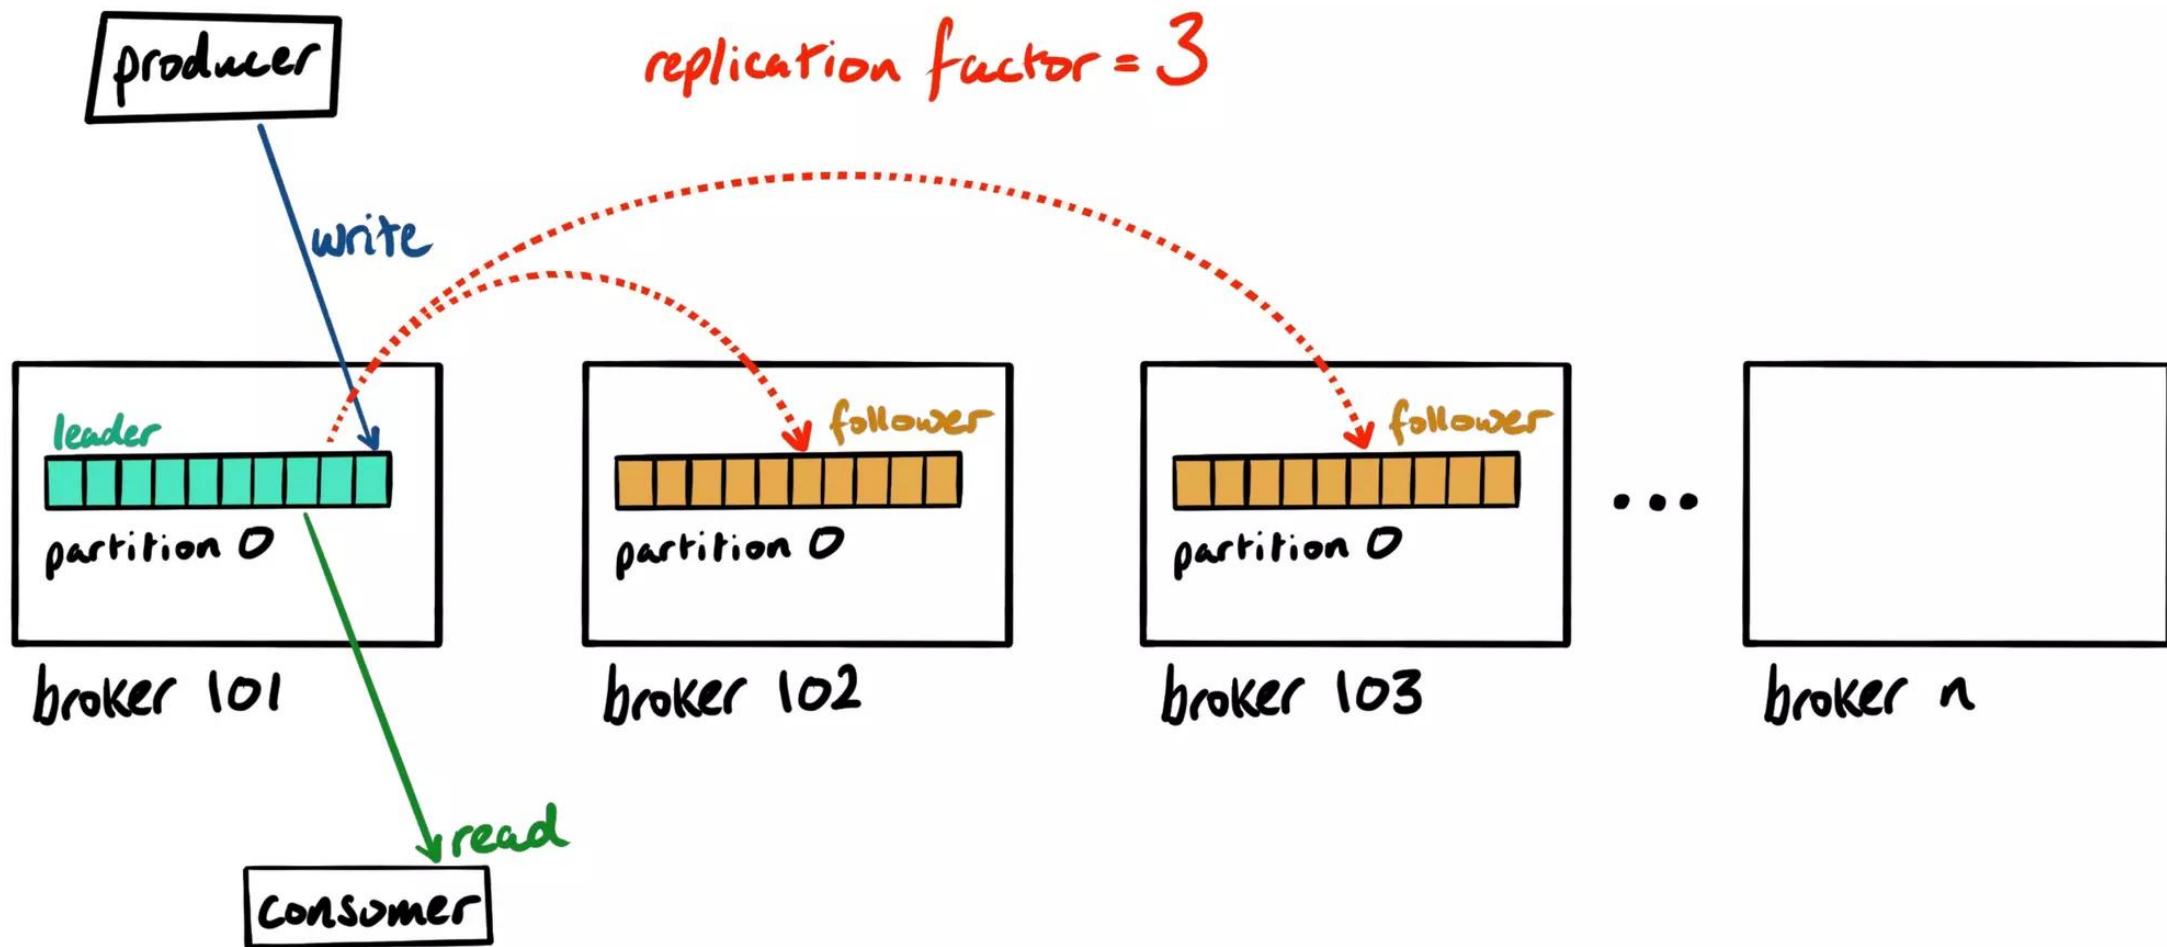
partition: 0
leader: 100
ISR: 101, 102

partition: 1
leader: 101
ISR: 100, 102

partition: 2
leader: 102
ISR: 100, 101



replication factor = 3



partitioning

why partition?

- load balancing

- semantic partitioning

out of the box behavior:

- With key: $\text{hash}(\text{key}) \% \text{ number of partitions}$
- no key: round robin

or write your own custom partitioner.

(you might do this if you have a composite key)

Kafka configs

Typical configs:

replication.factor : 3
min.insync.replicas : 2

} Topic (brokers)

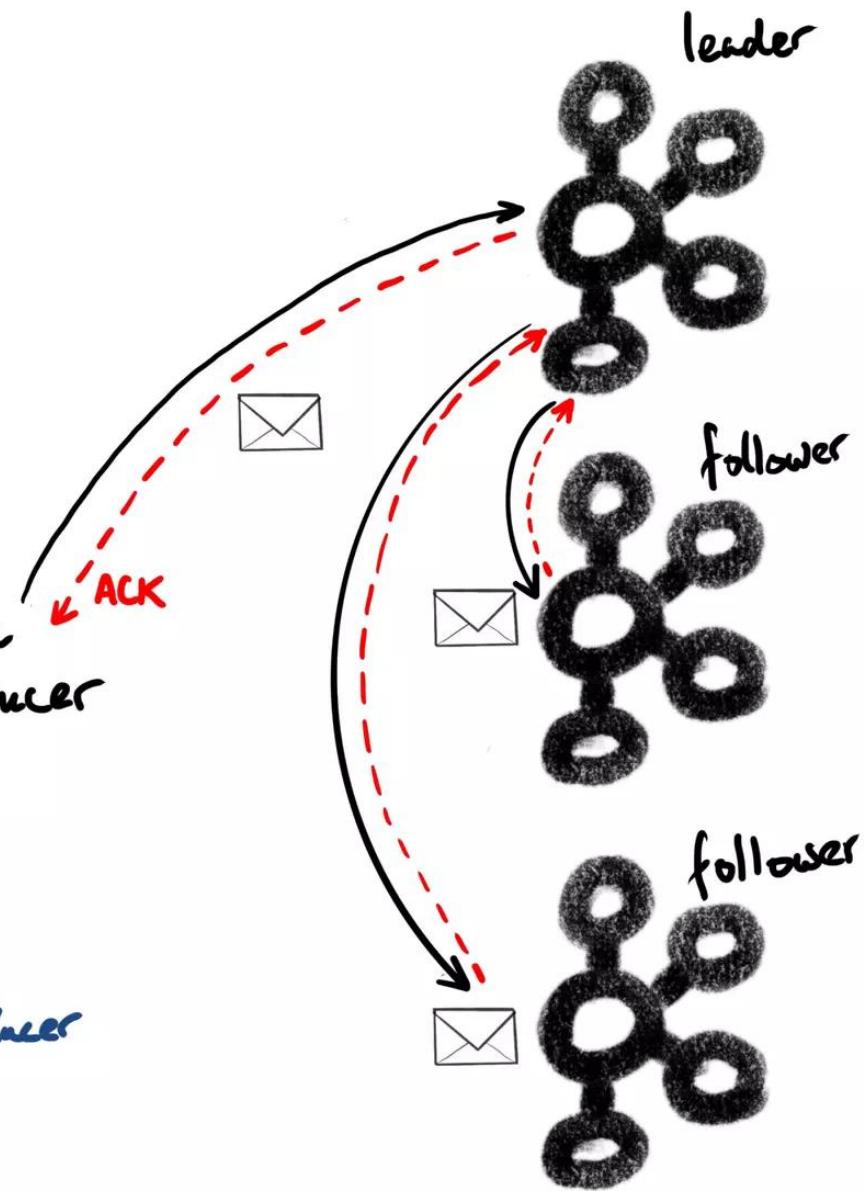
acks: 0 (fire and forget)

1 (data loss unlikely)

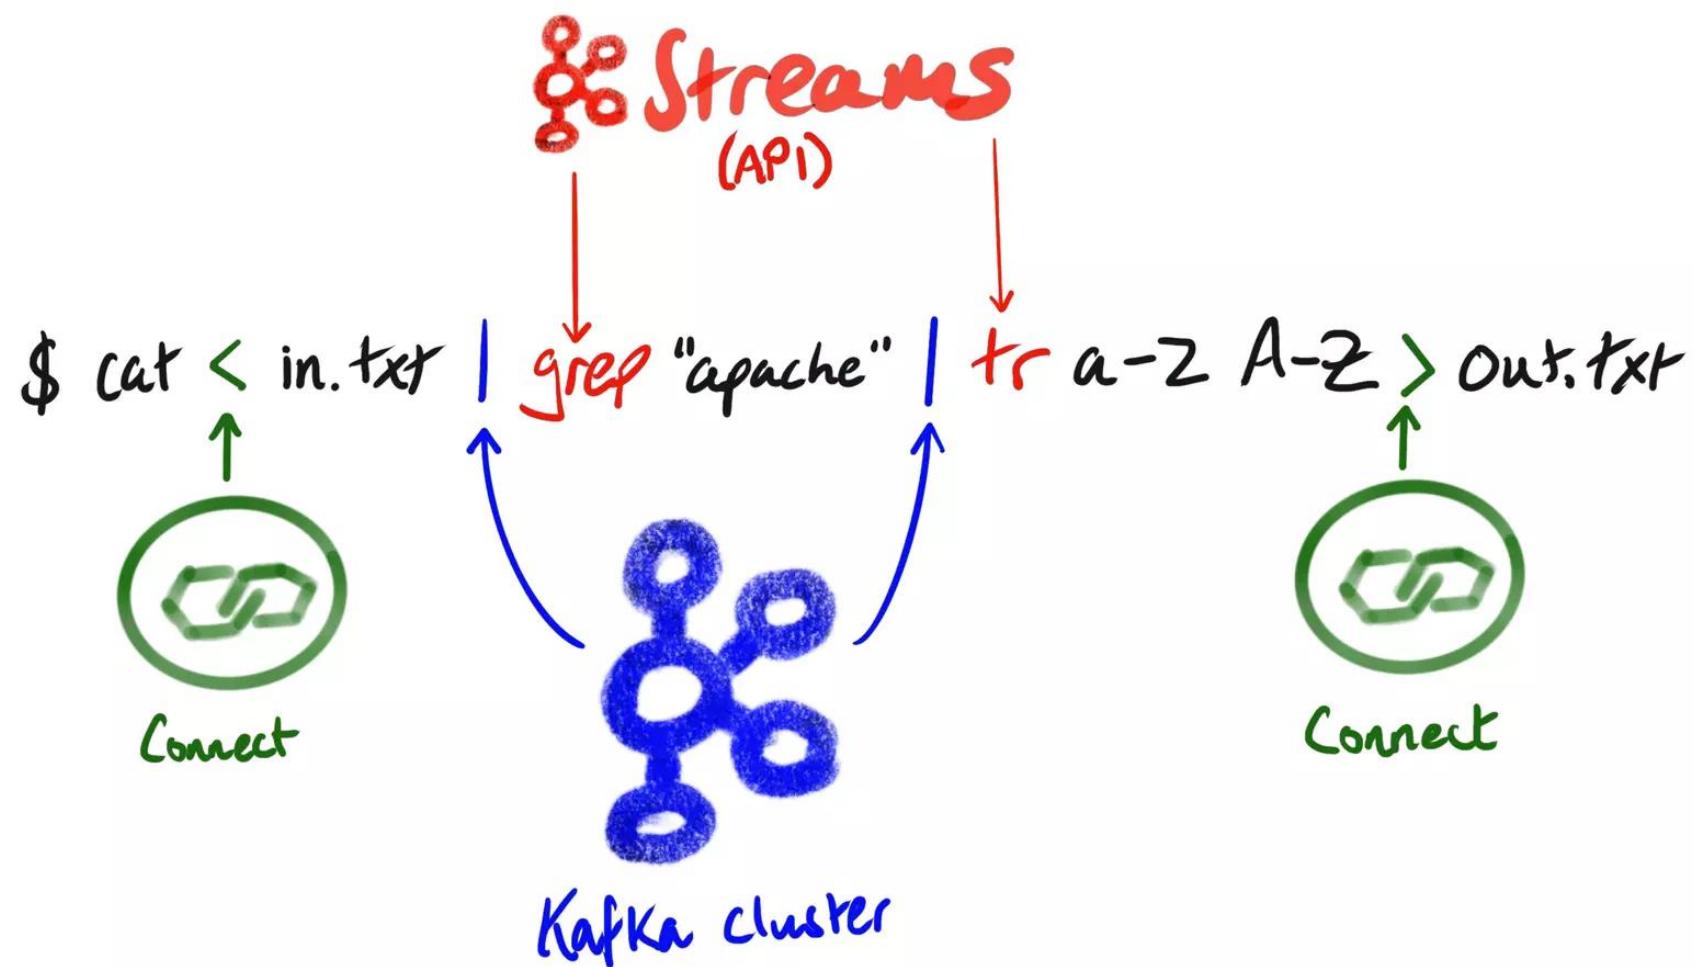
all (no data loss)

} Kafka producer

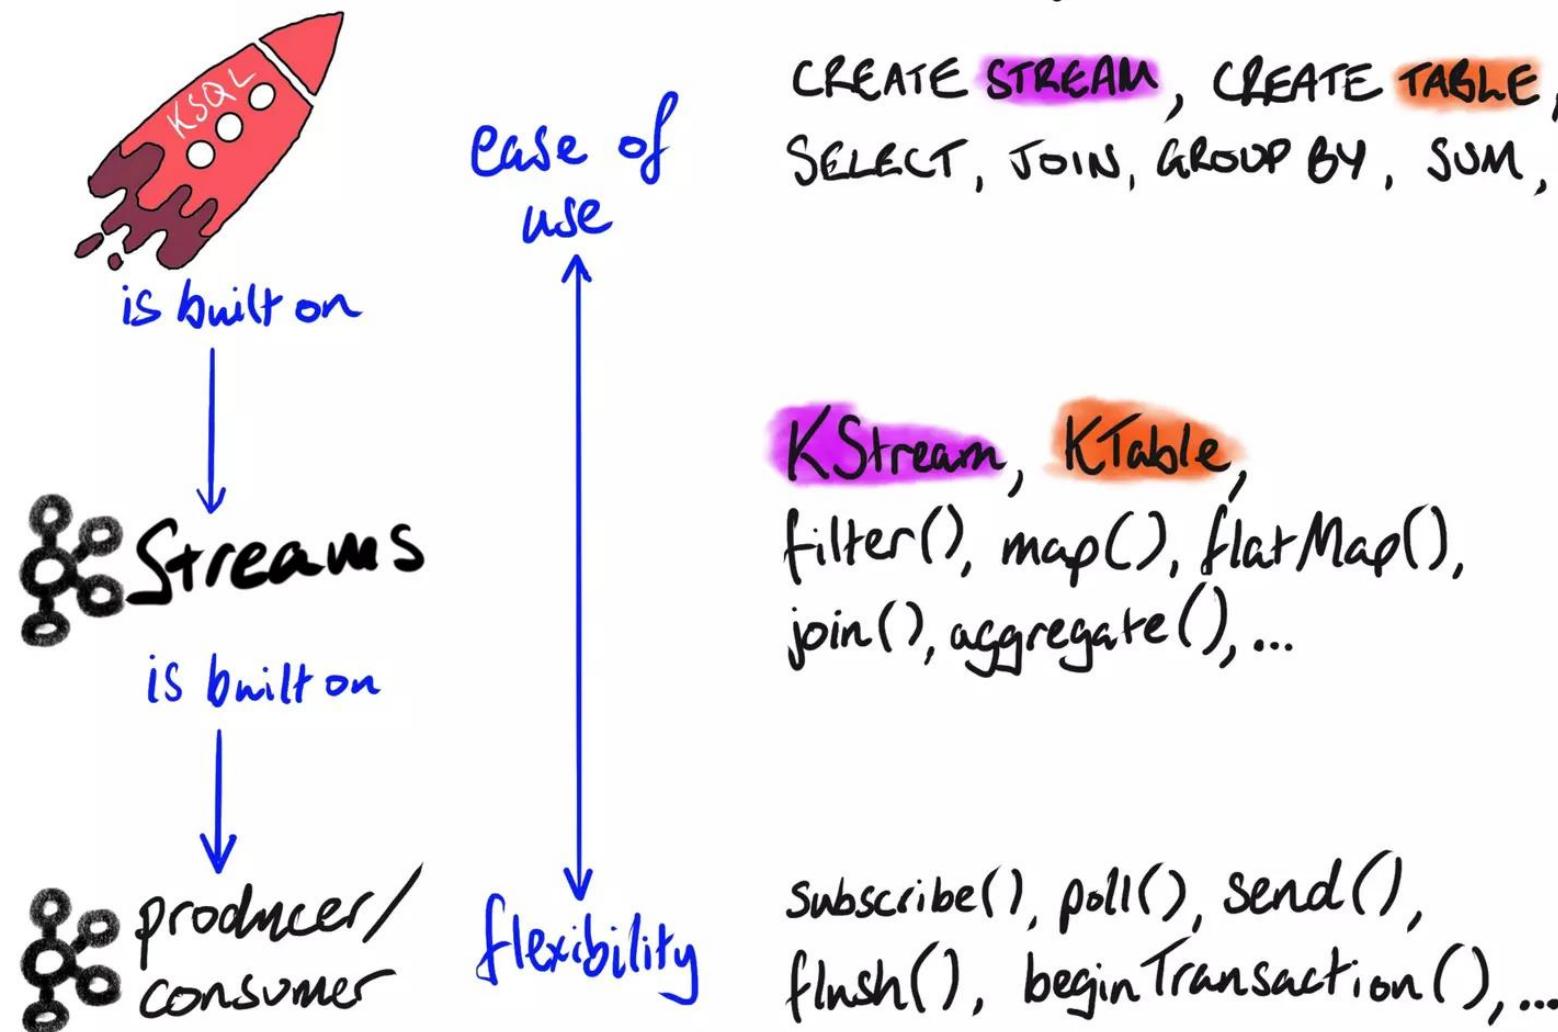
} producer



Kafka as a Linux Command



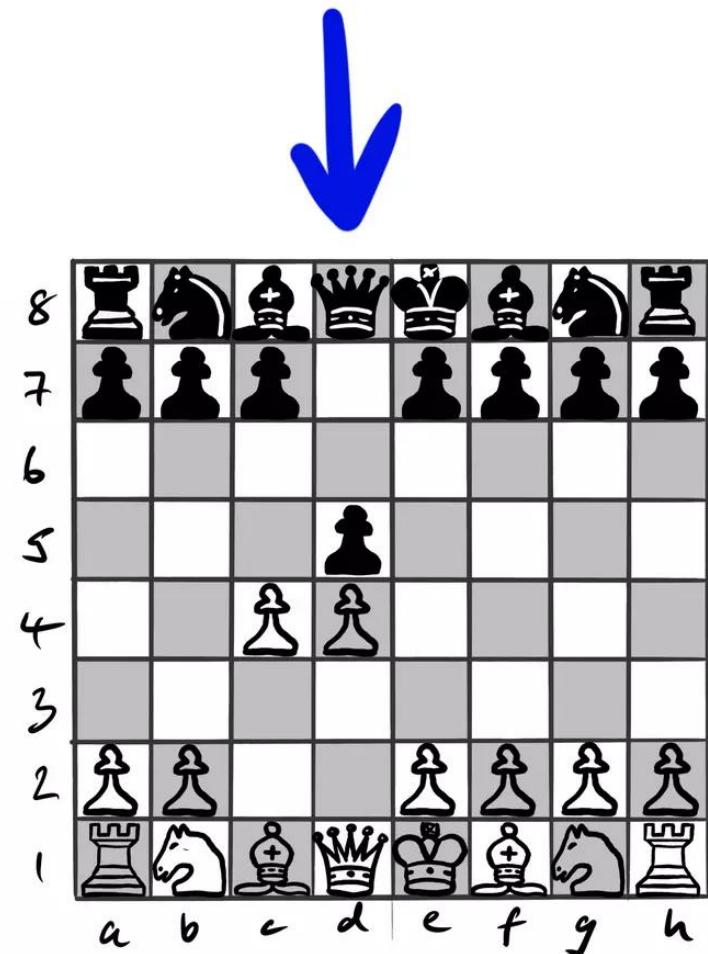
ksqlDB vs Kafka Streams vs producer/consumer



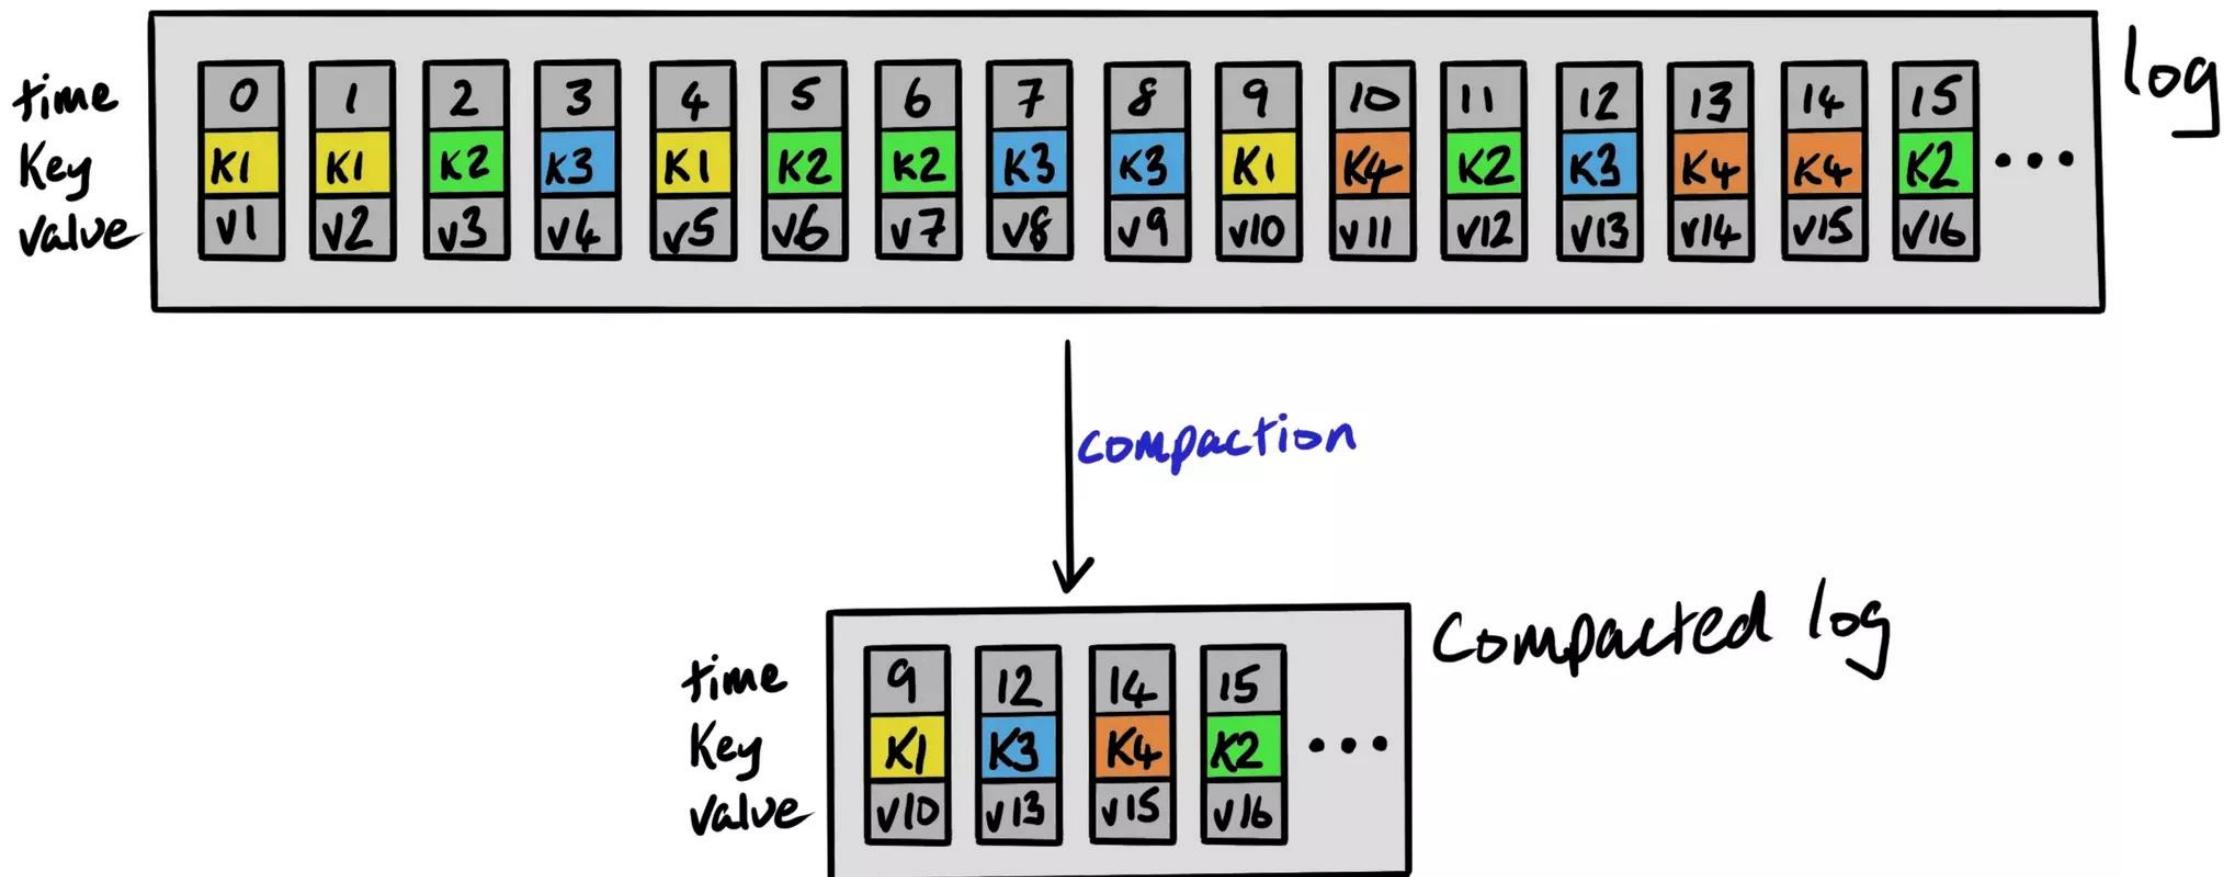
streams and tables



1. move white Queen's pawn forward two squares
2. move black Queen's pawn forward two squares
3. move white Queen-side bishop pawn forward two squares



Compacted topic



Stateless Operations

Branch

Filter

FilterNot

Flatmap

FlatmapValues

ForEach

GroupByKey

GroupBy

Map

MapValues

Peek

Print

SelectKey

toStream

* marks stream for re-partitioning

* always causes re-partitioning

* requires repartitioning if stream
is marked for it

Stateful Operations

Aggregations

Joins

Windows

Custom (PAPI)

Event-time processing

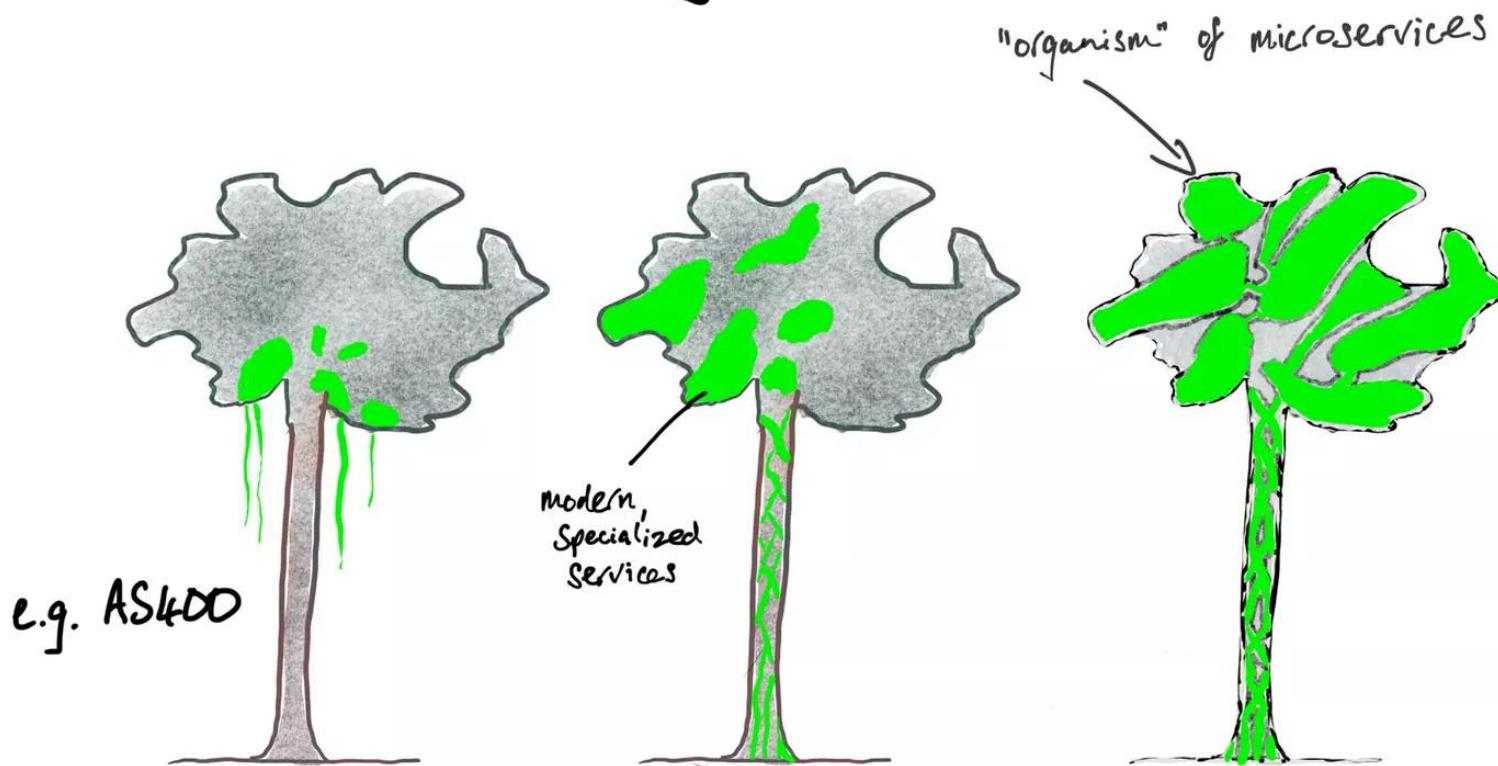
Stream time:

ingest time: when the record is written to a partition

event time: when the event occurred 

processing time: when the event was processed by
the streaming app. 

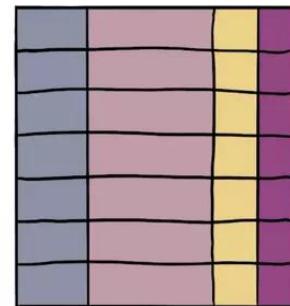
the Strangler pattern



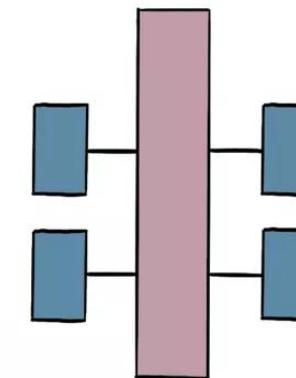
data access patterns

SQL

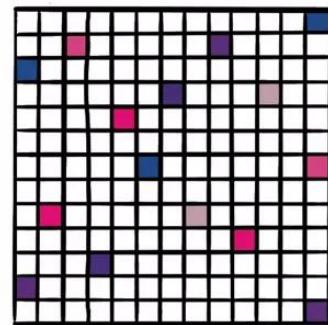
Ye olde RDBMS



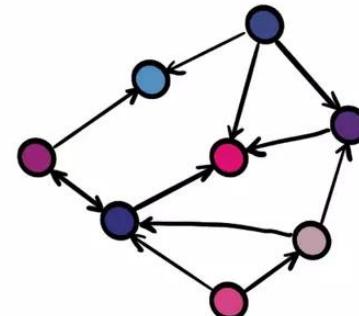
OLAP



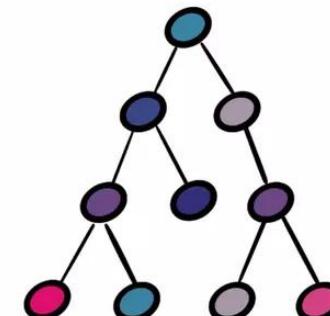
column - store



graph

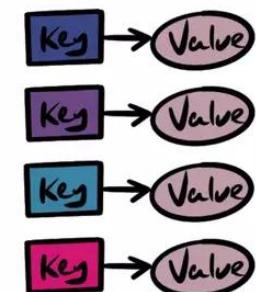


document

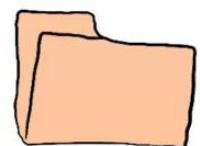
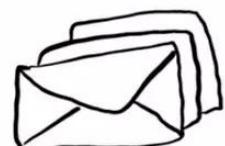


NoSQL

key / value



Sources



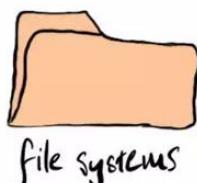
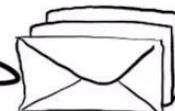
Connect

SMT

KSQL



Sinks

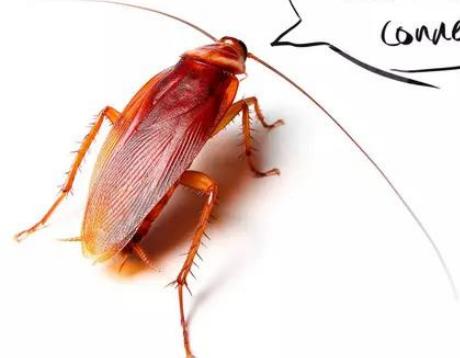


Connect

SMT

Streams

Kafka Connect



Wow! Over 100 connectors!



columnar DB's
DATASTAX



APACHE HBASE
APACHE PHOENIX
MAPR-DB

search/document DB's
splunk > Solr
Couchbase



Storage
hadoop

timeseries DB's
druid

change data capture (CDC)
ORACLE GOLDEN GATE
ATTUNITY SQDATA HVR



SQL
Java JDBC
SQL VERTICA
YugaByte DB
NETEZZA [ROCKSET]

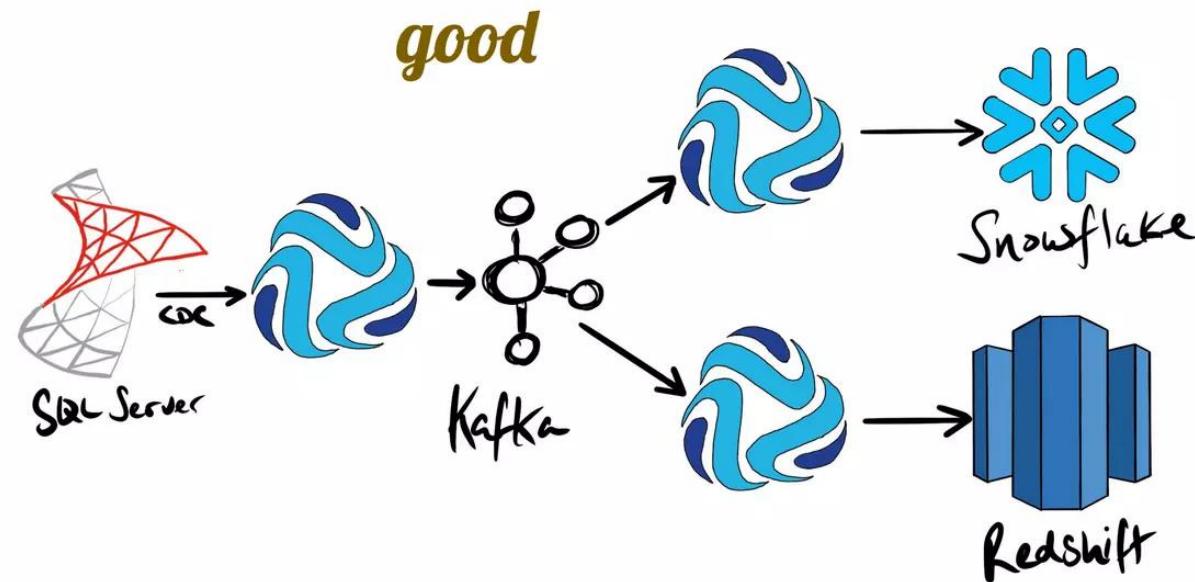
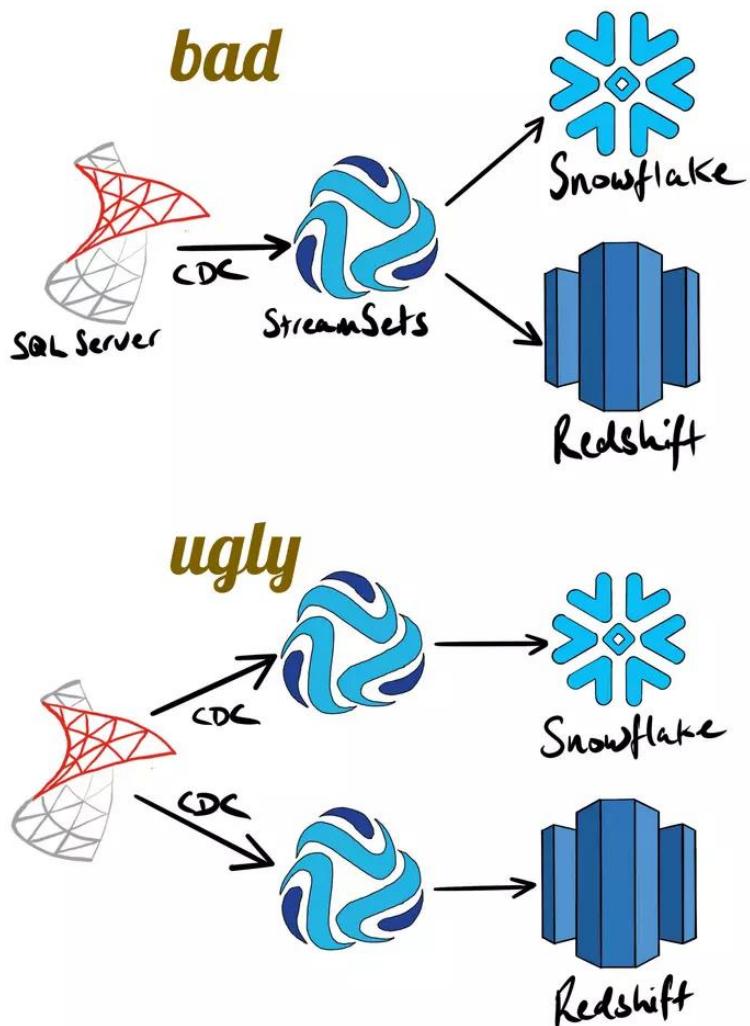
messaging
TIBCO®
JMS solace.
MQTT.org
Humio IBM MQ
ACTIVE MQ
PUSH TECHNOLOGY

GPU
kinetica
omnisci

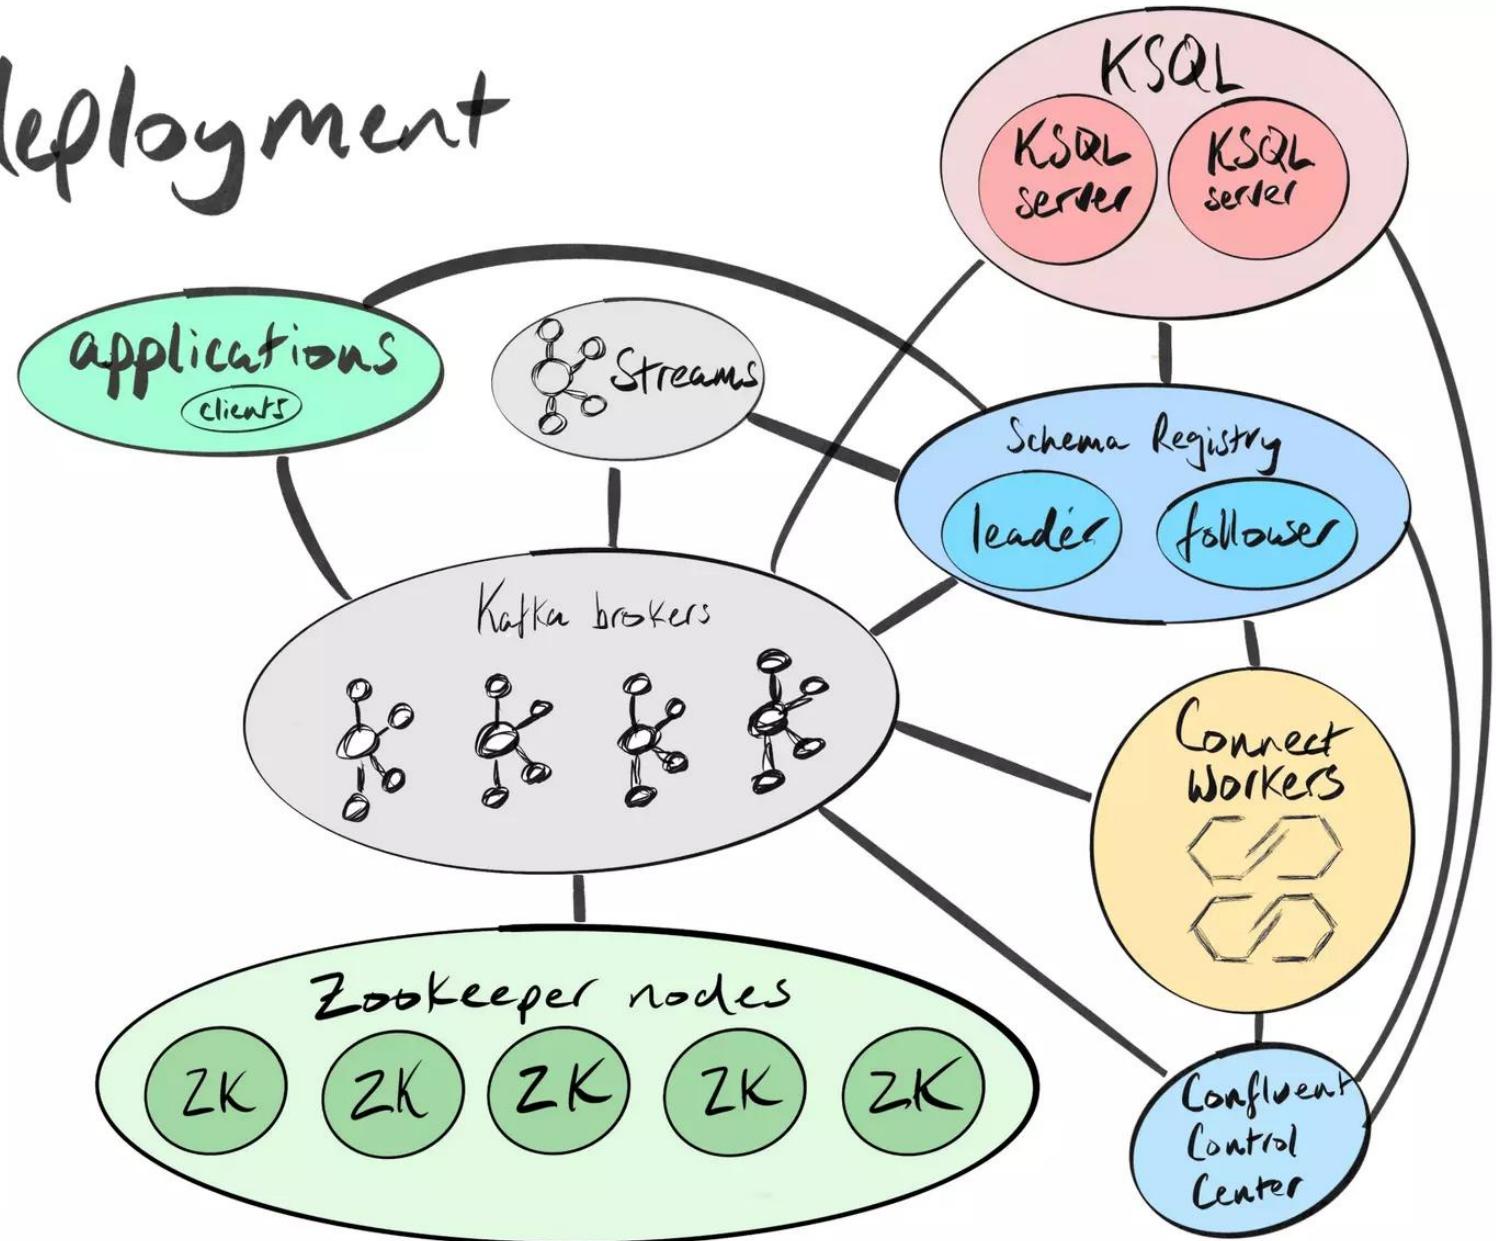
other stuff
salesforce
servicenow.
PRIVITAR VENAFI
StreamSets

confluent.io/hub

StreamSets: don't tightly couple pipelines



a typical deployment



Schema Evolution

```
person: {  
    lastName: string,  
    firstName: string,  
    age: int,  
    gender: [male, female]  
}
```



V1

```
person: {  
    lastName: string,  
    firstName: string,  
    ssn: string.  
    age: int,  
    gender: [male, female]  
}
```



V2

```
person: {  
    lastName: string,  
    firstName: string,  
    ssn: string,  
    age: int,  
    gender: [male, female, other]  
}
```

V3

PushTopic events from Salesforce

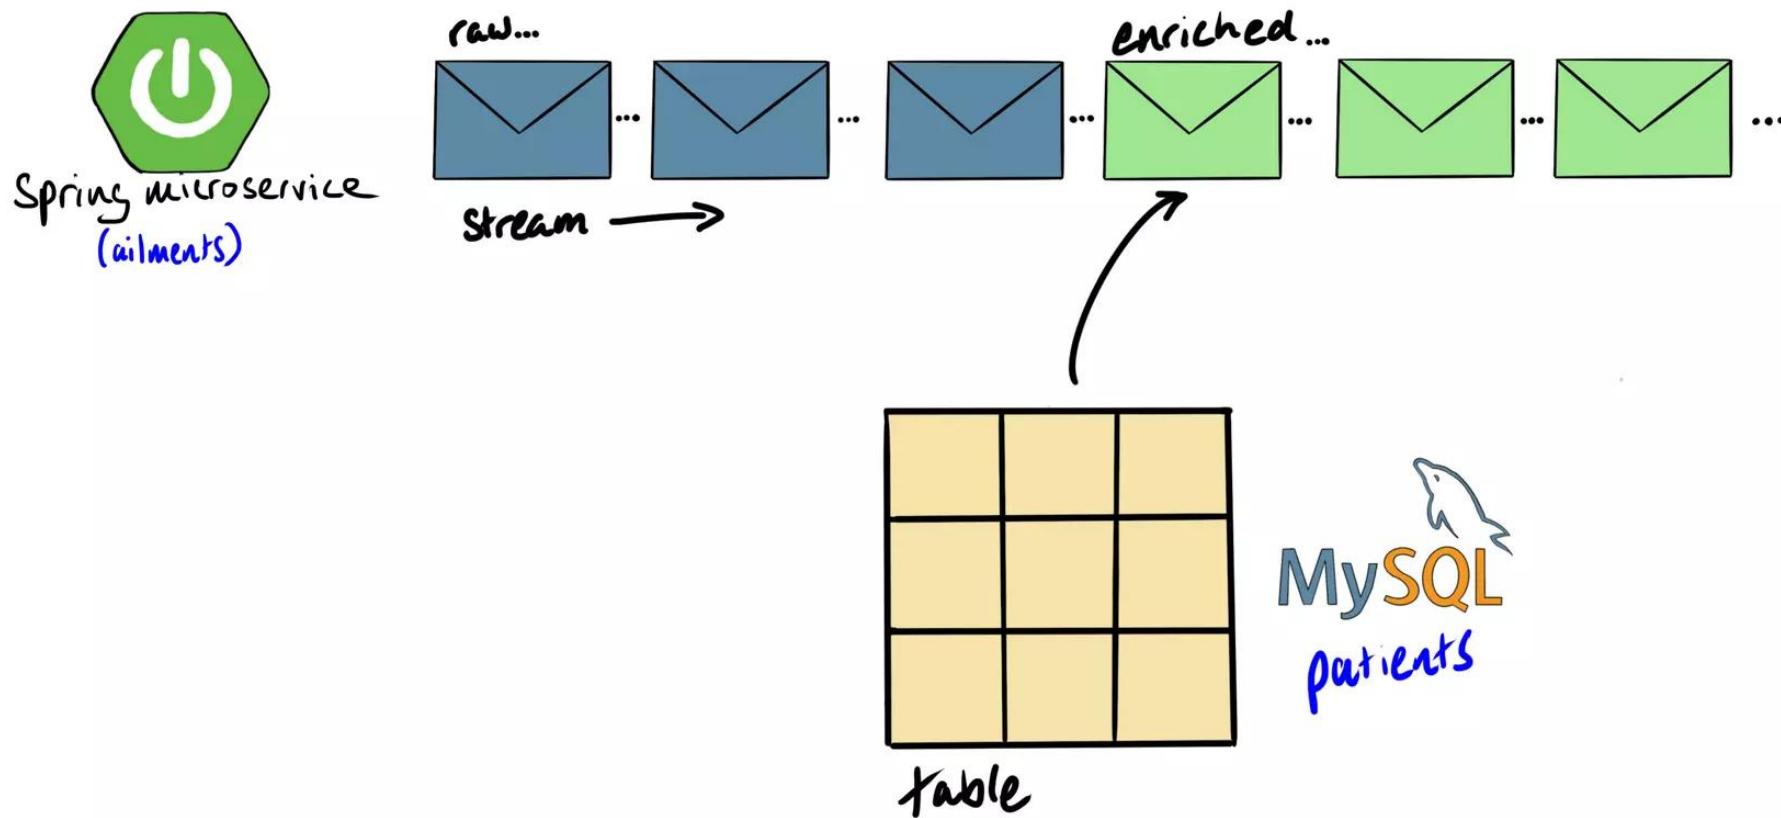
create PushTopic

```
PushTopic pushTopic = new PushTopic();
pushTopic.Name = 'ContactUpdates';
pushTopic.Query = 'SELECT Id, [...] FROM Contact';
pushTopic.ApiVersion = 50.0;
pushTopic.NotifyForOperationCreate = true;
pushTopic.NotifyForOperationUpdate = true;
pushTopic.NotifyForOperationUndelete = true;
pushTopic.NotifyForOperationDelete = true;
pushTopic.NotifyForFields = 'All';
insert pushTopic;
```

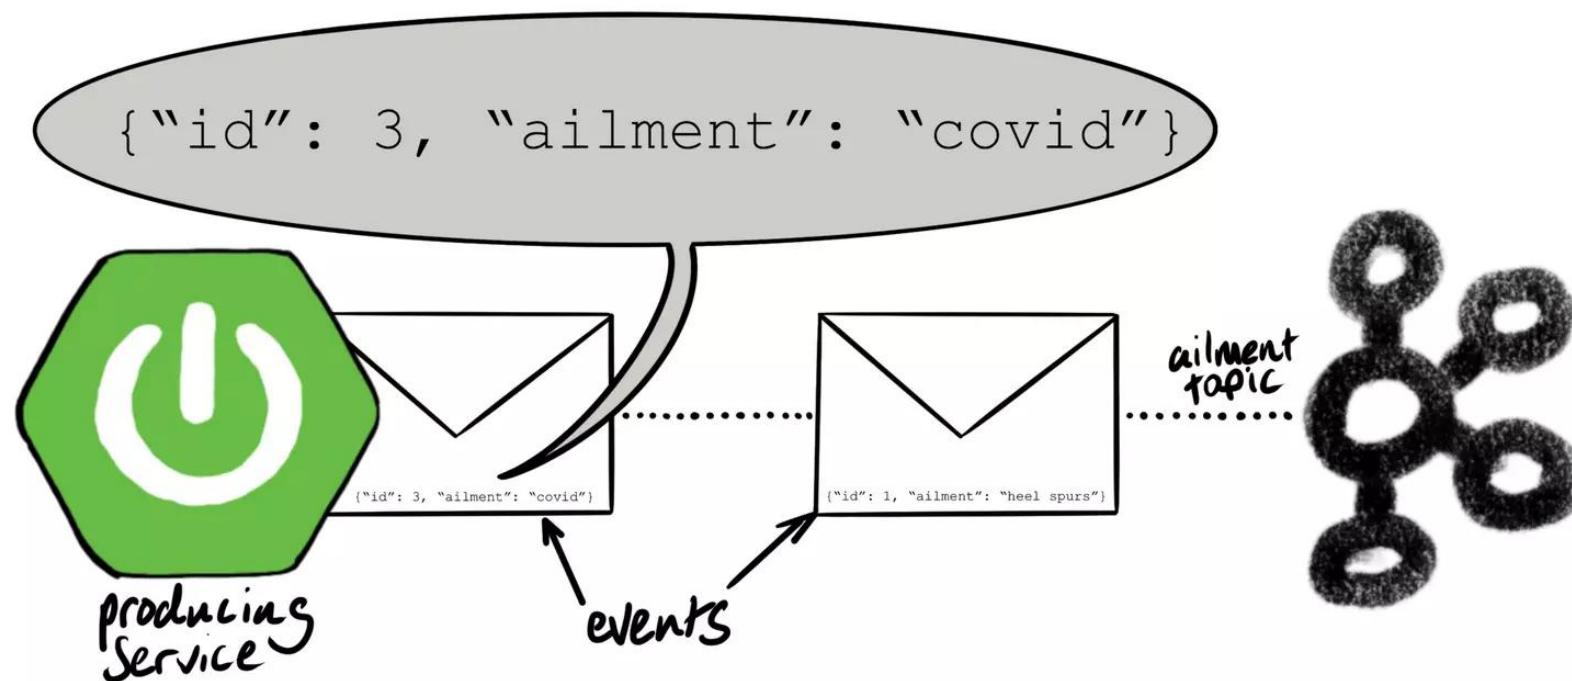
create connector

```
http PUT cp01.woolford.io:8083/connectors/sfdc-contact-updates/config <<< '
{
    "confluent.topic.bootstrap.servers": "cp01.woolford.io:9092,cp02...",
    "confluent.topic.replication.factor": "1",
    "connector.class":
    "io.confluent.salesforce.SalesforcePushTopicSourceConnector",
    "kafka.topic": "sfdc-contact-updates",
    "name": "sfdc-contact-updates",
    "salesforce.consumer.key": "*****",
    "salesforce.consumer.secret": "*****",
    "salesforce.initial.start": "all",
    "salesforce.object": "Contact",
    "salesforce.password": "cxI3o3R$****",
    "salesforce.password.token": "iggX9UEPizeaizvpg02Vj****",
    "salesforce.push.topic.name": "ContactUpdates",
    "salesforce.username": "alex@woolford.io",
    "tasks.max": "1"
}
```

Stream/table join



The ailments



The patients.



```
mysql> use healthcare;
```

```
Database changed
```

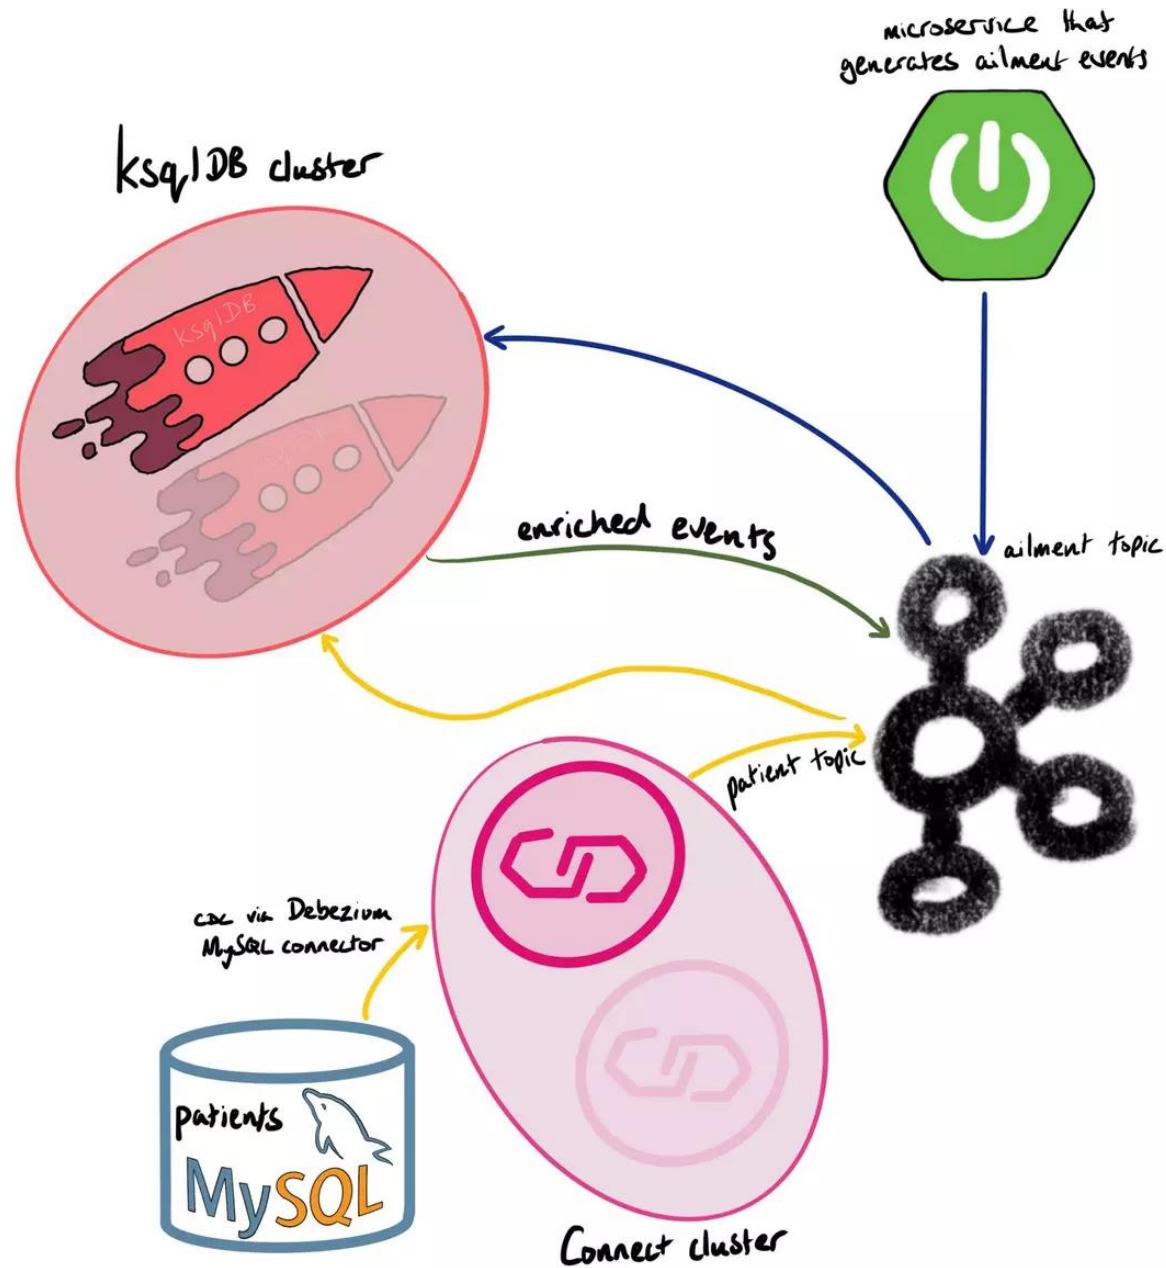
```
mysql> select * from patient;
```

id	firstname	lastname	enroll_start	enroll_end
1 Donald	Trump		2019-04-14	2020-04-13
2 Vladimir	Putin		2019-05-20	2020-05-19
3 Boris	Johnson		2019-01-24	2020-01-24
4 Mark	Zuckerberg		2019-03-03	2020-03-02
5 Benjamin	Netanyahu		2019-08-24	2020-08-23
6 Jair	Bolsonaro		2019-09-19	2020-09-18
7 Farty	McFartface		2019-08-23	2020-08-22

```
7 rows in set (0.00 sec)
```

change data capture from MySQL with the Debezium connector

```
http PUT http://cp01.woolford.io:8083/connectors/mysql-cdc-healthcare/config <<< '
{
    "name": "mysql-cdc-healthcare",
    "connector.class": "io.debezium.connector.mysql.MySqlConnector",
    ...
    "include.schema.changes": "false",
    "snapshot.mode": "when_needed",
    "table.whitelist": "healthcare.patient",
    "transforms": "unwrap,HoistField",
    "transforms.unwrap.type": "io.debezium.transforms.ExtractNewRecordState",
    "transforms.HoistField.type": "org.apache.kafka.connect.transforms.ExtractField$Key",
    "transforms.HoistField.field": "id"
}'
```



create the stream and table in ksqlDB

```
CREATE TABLE PATIENT (
    id STRING,
    firstname VARCHAR,
    lastname VARCHAR,
    enroll_start BIGINT,
    enroll_end BIGINT
) WITH (
    kafka_topic='deept thought.healthcare.patient',
    value_format='JSON',
    KEY='id'
);
```

```
CREATE STREAM AILMENT (
    id STRING,
    ailment VARCHAR
) WITH (
    kafka_topic='ailment',
    value_format='JSON',
    KEY='id'
);
```

Streaming join in ksqlDB

```
select
    ailment.id as id,
    ailment,
    firstname,
    lastname,
    case
        when ailment.rowtime > enroll_start * 86400 * 1000
            and ailment.rowtime <= enroll_end * 86400 * 1000
        then true
        else false
    end as covered
from ailment inner join patient
on ailment.id = patient.id
emit changes;
```

ksqlDB

Typical ksqlDB use cases:

- streaming ETL
- real-time monitoring and analytics
- data exploration and discovery
- anomaly detection
- personalization
- sensor and IoT
- customer 360

events from video



KSQL user defined function

enrichment with MaxMind's geoIP

```
select getgeoforip(ip) from mystream emit changes;
```

github.com/alexwoolford/ksql-udf-geoip

enrichment with EXIF from an image

NYC released 173 million anonymized taxi journeys
... which was easily de-anonymized thanks to EXIF

```
SELECT D.dropoff_latitude, D.dropoff_longitude, F.total_amount, F.tip_amount
FROM tripData AS D, tripFare AS F
WHERE D.hack_license = F.hack_license AND D.pickup_datetime = F.pickup_datetime
AND pickup_datetime > "2013-07-08 19:33:00" AND pickup_datetime < "2013-07-08 19:37:00"
AND pickup_latitude > 40.719 AND pickup_latitude < 40.7204
AND pickup_longitude > -74.0106 AND pickup_longitude < -74.01;
```



```
from exif import Image

with open('IMG_20181103_200537.jpg', 'rb') as
image_file:
    my_image = Image(image_file)
```

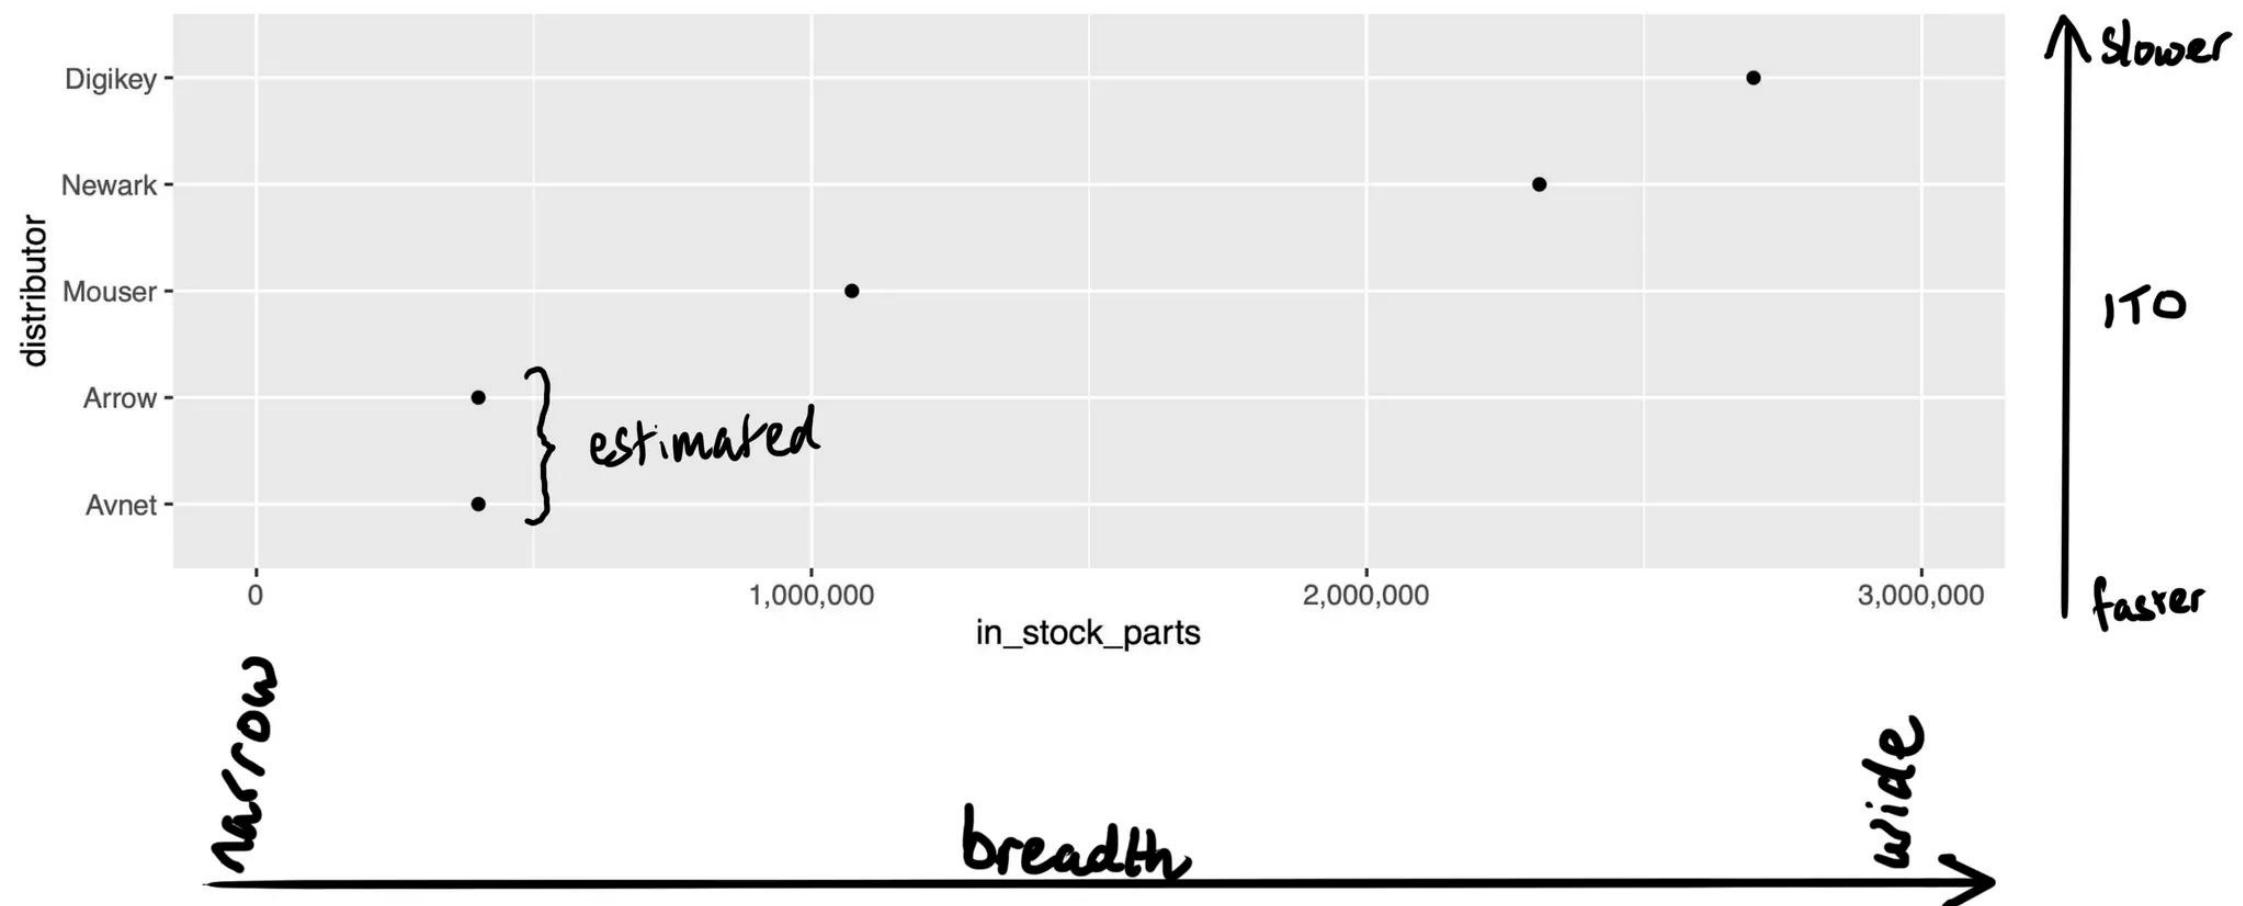
lat/lon
timestamp
camera settings
copyright

Data-driven inventory expansion

The challenge: add SKU's without
decreasing inventory turns

$$\text{inventory turnover} = \frac{\text{cost of goods sold}}{\text{average inventory}}$$

more SKUs == slower inventory turns



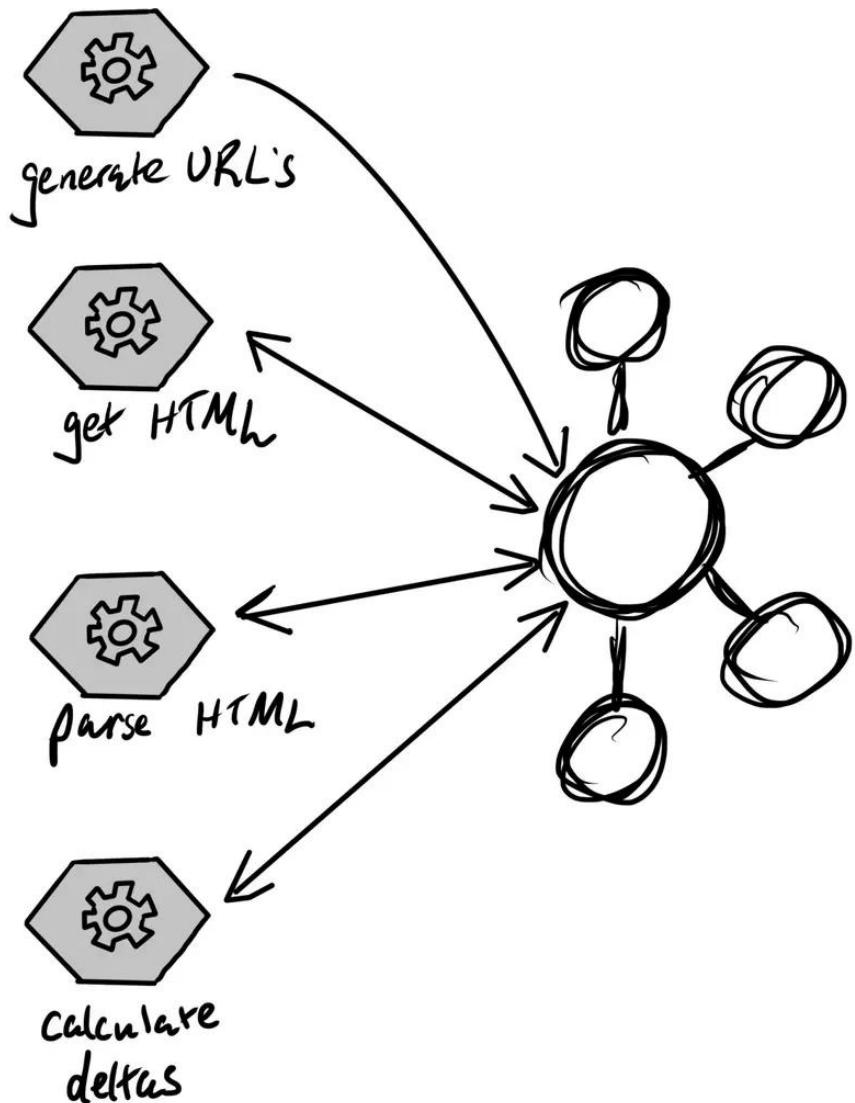
Information distortion: the bullwhip effect

order variability up the Pampers supply chain



source: Lee, Padmanabhan and Whang, 1997

Supply chain visibility



products/en/development-boards-kits-programmers/evaluation-and-demonstration-boards-and-

Image	Part Number	Manufacturer Part Number	Manufacturer	Description	Quantity Available	Unit Price USD	Minimum Quantity
	1568-1209-ND	BOB-12009	SparkFun Electronics	LOGIC LEVEL CONVERTER - BI-DIREC	1,638 - Immediate	\$2.95000	1
	1568-1279-ND	BOB-09118	SparkFun Electronics	SPARKFUN OPTO-ISOLATOR BREAKOUT	279 - Immediate	\$4.95000	1
	497-10828-ND	EVAL6235PD	STMicroelectronics	EVAL BOARD FOR THE L6235PD	0	Active	1 Non-Stock
	1568-1193-ND	BOB-11189	SparkFun Electronics	TRANSCEIVER BREAKOUT - MAX3232	1,387 - Immediate	\$5.95000	1
	1460-1201-ND	TMC2208 SILENTSTEPSTICK	Trinamic Motion Control GmbH	TMC2208 STEPPER DRIVER BOARD	1,188 - Immediate	\$7.32000	1
	768-1317-ND	LC234X	FTDI, Future Technology Devices International Ltd	MOD USB UART DEV FT234XD	284 - Immediate	\$7.78000	1

Audio Products - [222 New Products](#)

- [Accessories](#) (659 items)
- [Alarms, Buzzers, and Sirens](#) (5607 items)
- [Amplifiers](#) (23 items)
- [Buzzer Elements, Piezo Benders](#) (143 items)
- [Guitar Parts, Accessories](#) (213 items)
- [Microphones](#) (1767 items)
- [Speakers](#) (2608 items)
- [Vacuum Tubes](#) (1408 items)

Battery Products - [344 New Products](#)

- [Accessories](#) (274 items)
- [Batteries Non-Rechargeable \(Primary\)](#) (698 items)
- [Batteries Rechargeable \(Secondary\)](#) (837 items)
- [Battery Chargers](#) (415 items)
- [Battery Holders, Clips, Contacts](#) (1704 items)
- [Battery Packs](#) (2832 items)
- [Cigarette Lighter Assemblies](#) (69 items)

<https://www.digikey.com/products/en/audio-products/alarms-buzzers-and-sirens/157?&quantity=0&pageSize=500&page=1>

Predictable URL's

The # of pages, for each category, can be determined by the # of parts.

E.g. ceiling($5,607 / 500$) = 12

lots of parts on
each page

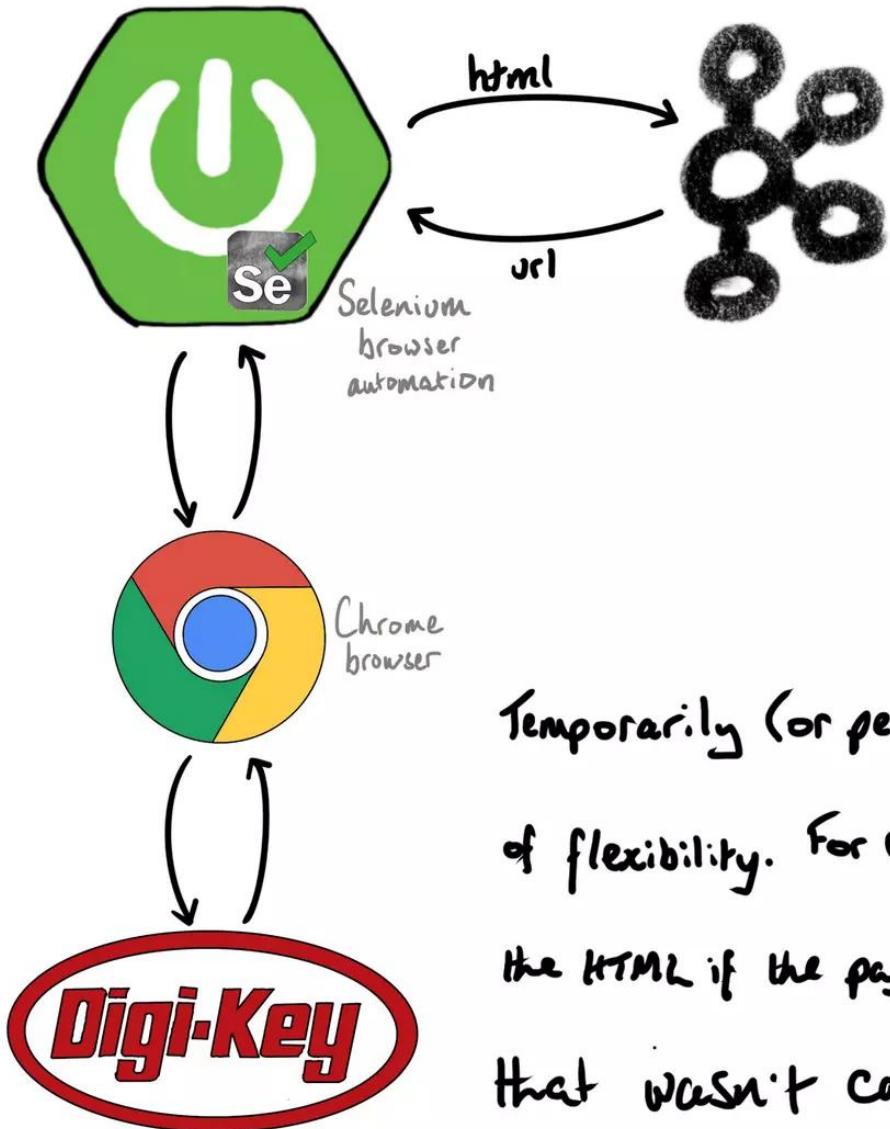
Show in-stock
items only

easy pagination



Although it's common practice, web scraping is ethically dubious and can lead to legal problems (e.g. trespass to chattels).

Always spread the requests across a pool of residential IP addresses and go slow, so it doesn't look like a denial of service attack.



Spring Boot microservice reads from URL topic, gets the HTML, and writes that back to another topic

use Selenium when the pages have to execute
JavaScript to display the content

Temporarily (or permanently) storing the HTML gives a lot of flexibility. For example, it's possible to re-parse information from the HTML if the page format changes, or if there's some useful information that wasn't captured.

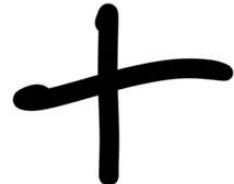
Compare Parts		Image	Digi-Key Part Number	Manufacturer Part Number	Manufacturer	Description	Quantity Available ?	Unit Price USD	Minimum Quantity	Packaging	Series	Part Status
			▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼	▲ ▼
<input type="checkbox"/>			102-1153-ND	CEM-1203(42)	CUI Devices	AUDIO MAGNETIC XDCR 3-5V TH	57,928 - Immediate	\$0.74000	1	Tray ?	CEM	Active
<input type="checkbox"/>			102-1286-ND	CSQG703BP	CUI Devices	AUDIO MAGNETIC XDCR 2-4V TH	31,760 - Immediate	\$0.83000	1	Tray ?	CSQ	Active
<input type="checkbox"/>			102-1156-ND	CEM-1212S	CUI Devices	AUDIO MAGNETIC XDCR 6-16V TH	40,718 - Immediate	\$0.86000	1	Tray ?	CEM	Active
<input type="checkbox"/>			102-1458-ND	CST-931AP	CUI Devices	AUDIO MAGNETIC XDCR 2-4V TH	69,146 - Immediate	\$1.06000	1	Tray ?	CST	Active
<input type="checkbox"/>			102-1199-2-ND	CT-1205-SMT-TR	CUI Devices	AUDIO MAGNETIC XDCR 3-8V SMD	36,000 - Immediate	\$1.18678	450	Tape & Reel (TR) ?	CT	Active



```
{"timestamp":1598492553789,"dpn":"DM330028-2-ND","mpn":"DM330028-2","mfg":"Microchip Technology","qoh":3}
{"timestamp":1598492553789,"dpn":"296-28530-ND","mpn":"MSP-EXP430FR5739","mfg":"Texas Instruments","qoh":166}
{"timestamp":1598492553789,"dpn":"497-STM32H7B3I-DK-ND","mpn":"STM32H7B3I-DK","mfg":"STMicroelectronics","qoh":1}
 {"timestamp":1598492553789,"dpn":"TMDSLCDK6748-ND","mpn":"TMDSLCDK6748","mfg":"Texas Instruments","qoh":10}
 {"timestamp":1598492553789,"dpn":"497-12073-ND","mpn":"STM8SVLDISCOVERY","mfg":"STMicroelectronics","qoh":5}
 {"timestamp":1598492553789,"dpn":"497-13930-ND","mpn":"STM32L100C-DISCO","mfg":"STMicroelectronics","qoh":12}
 {"timestamp":1598492553789,"dpn":"428-4367-ND","mpn":"CY8CKIT-146","mfg":"Cypress Semiconductor Corp","qoh":20}
 {"timestamp":1598492553789,"dpn":"428-4510-ND","mpn":"CY8CKIT-147","mfg":"Cypress Semiconductor Corp","qoh":10}
 {"timestamp":1598492553789,"dpn":"497-10660-ND","mpn":"STM8L-DISCOVERY","mfg":"STMicroelectronics","qoh":19}
 {"timestamp":1598492553789,"dpn":"497-18395-ND","mpn":"NUCLEO-G070RB","mfg":"STMicroelectronics","qoh":22}
 {"timestamp":1598492553789,"dpn":"497-15095-ND","mpn":"NUCLEO-F070RB","mfg":"STMicroelectronics","qoh":1}
 {"timestamp":1598492553789,"dpn":"497-15979-ND","mpn":"NUCLEO-F031K6","mfg":"STMicroelectronics","qoh":10}
 {"timestamp":1598492553789,"dpn":"497-16284-ND","mpn":"NUCLEO-L011K4","mfg":"STMicroelectronics","qoh":16}
 {"timestamp":1598492553789,"dpn":"497-18155-ND","mpn":"NUCLEO-8L152R8","mfg":"STMicroelectronics","qoh":16}
 {"timestamp":1598492553789,"dpn":"497-18239-ND","mpn":"NUCLEO-L412KB","mfg":"STMicroelectronics","qoh":20}
```

Calculate the deltas

```
{"timestamp":1598492567811,"dpn":"1568-DEV-16996-ND","mpn":"DEV-16996","mfg":"SparkFun Electronics","qoh":196}
```

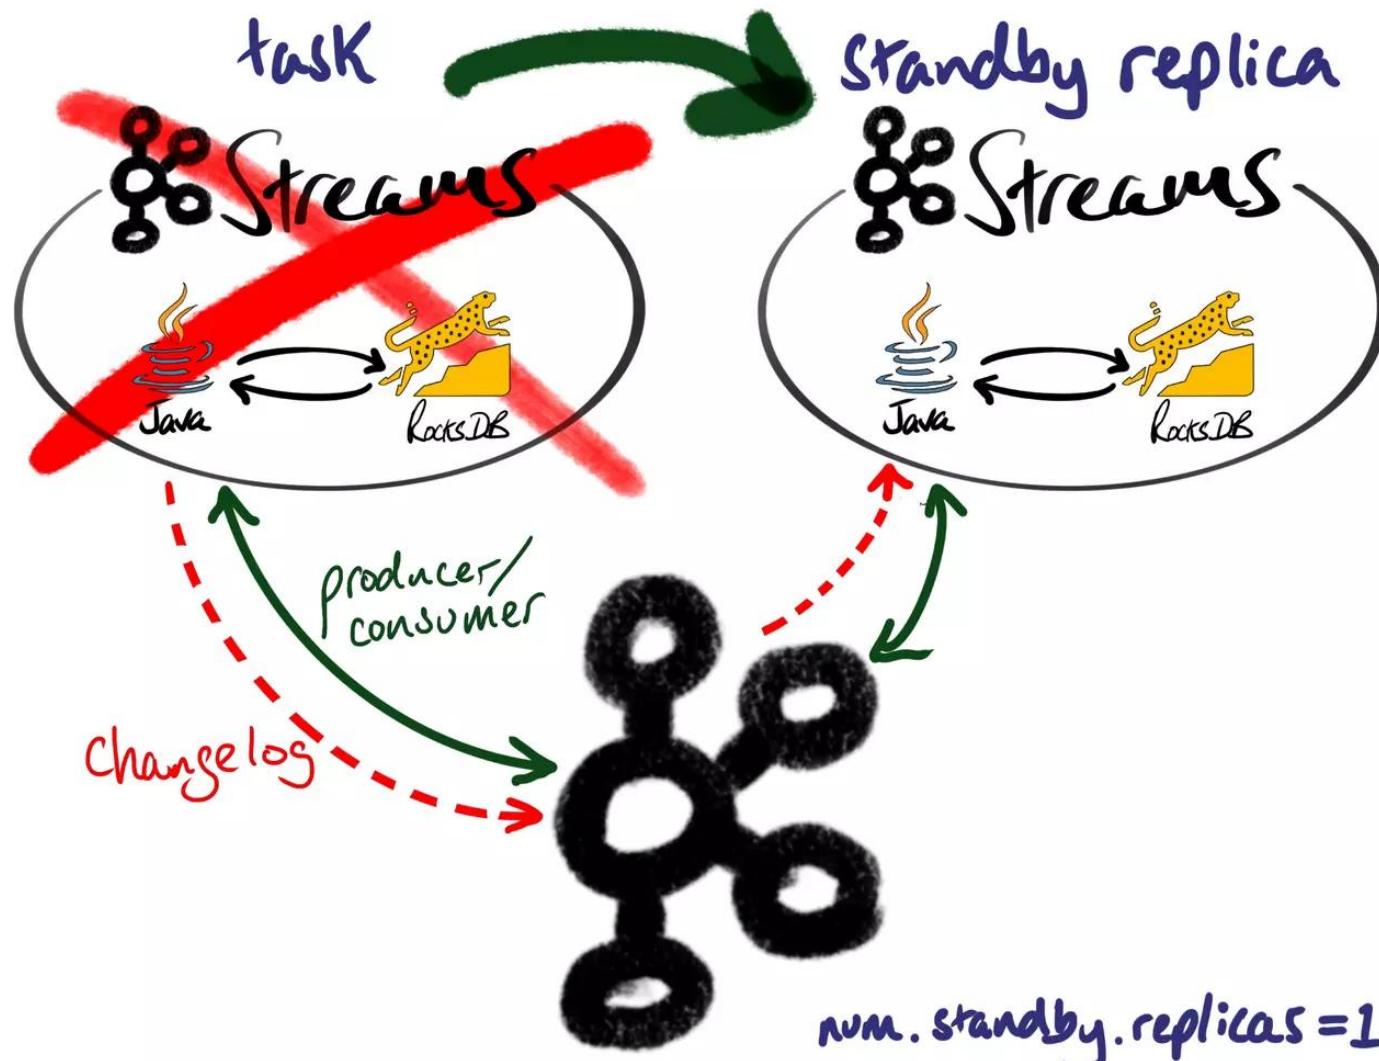


```
{"timestamp":1598492654204,"dpn":"1568-DEV-16996-ND","mpn":"DEV-16996","mfg":"SparkFun Electronics","qoh":194}
```

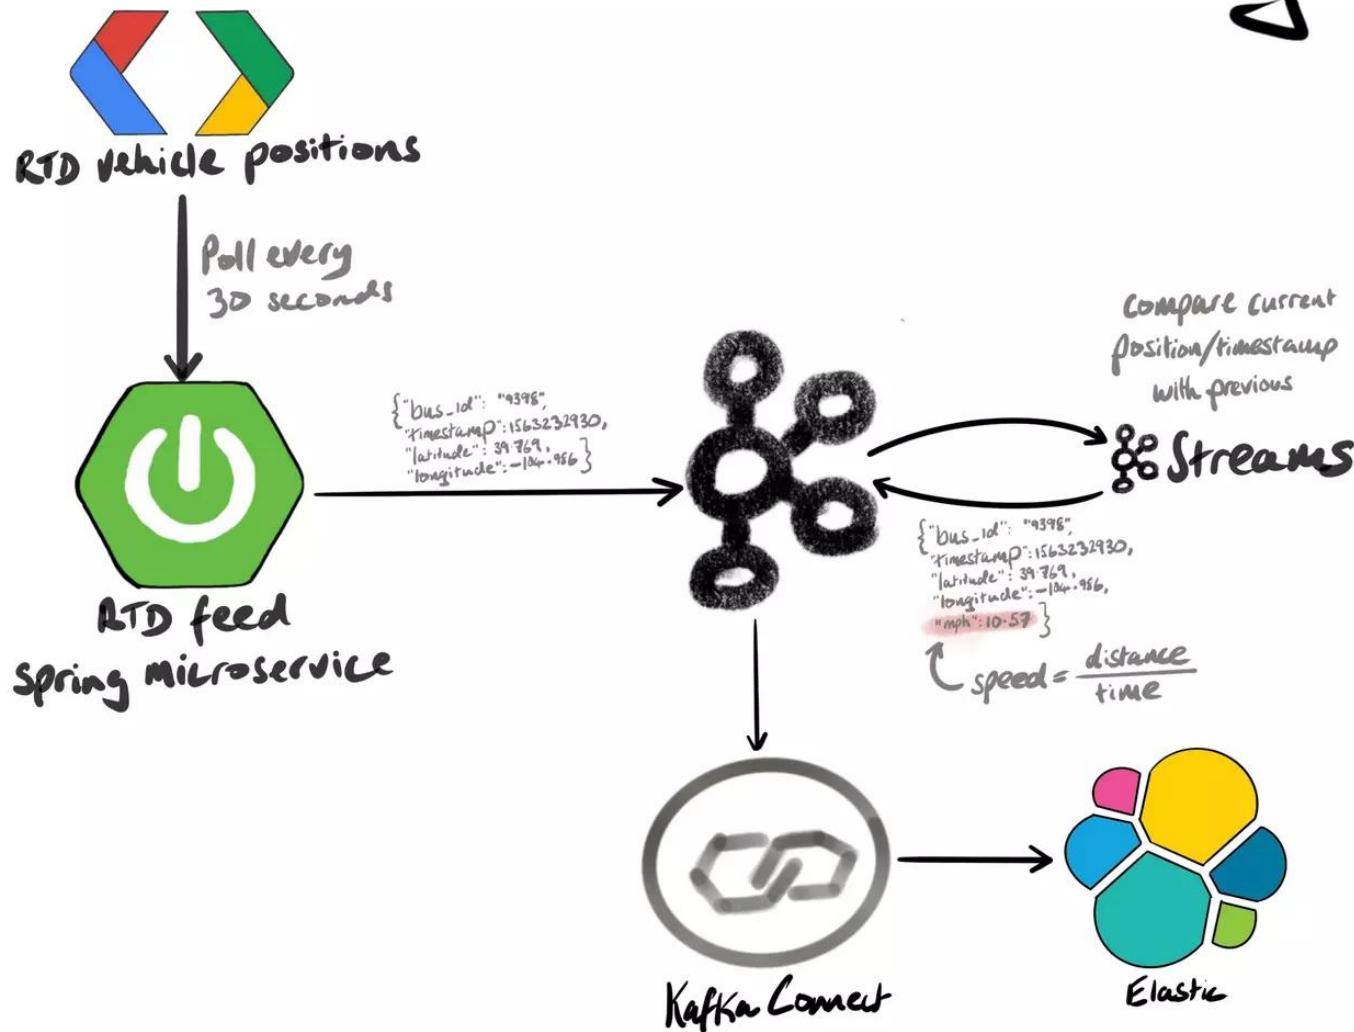


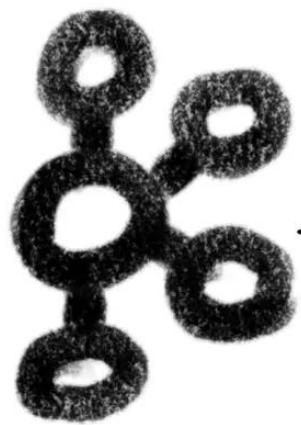
```
{"timestamp":1598492654204,"dpn":"1568-DEV-16996-ND","mpn":"DEV-16996","mfg":"SparkFun Electronics","delta":2}
```

Kafka Streams' state store



how fast are RTD buses going?





```
{  
    "_id": "291541",  
    "courseID": 495420,  
    "courseName": "Physical Science",  
    "SectionID": 2408179,  
    "taskID": 2851,  
    "taskName": "Final"  
    "progressScore": "A",  
    "progressPercent": 97.85  
}
```

PS4 Killswitch

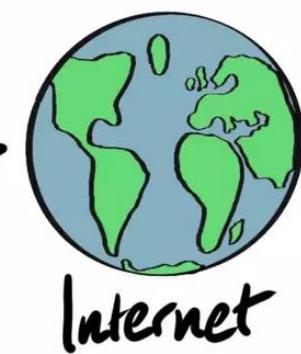


```
if (badGrade == true) {  
    //add PS4 to "deny-internet" group  
    ....  
}
```

PUT --data '{ "name": "deny-internet", "member": [{ "name": "ps4-host" }] }'
<https://10.0.1.1:443/api/v2/cmdb/firewall/addgrp/deny-internet>



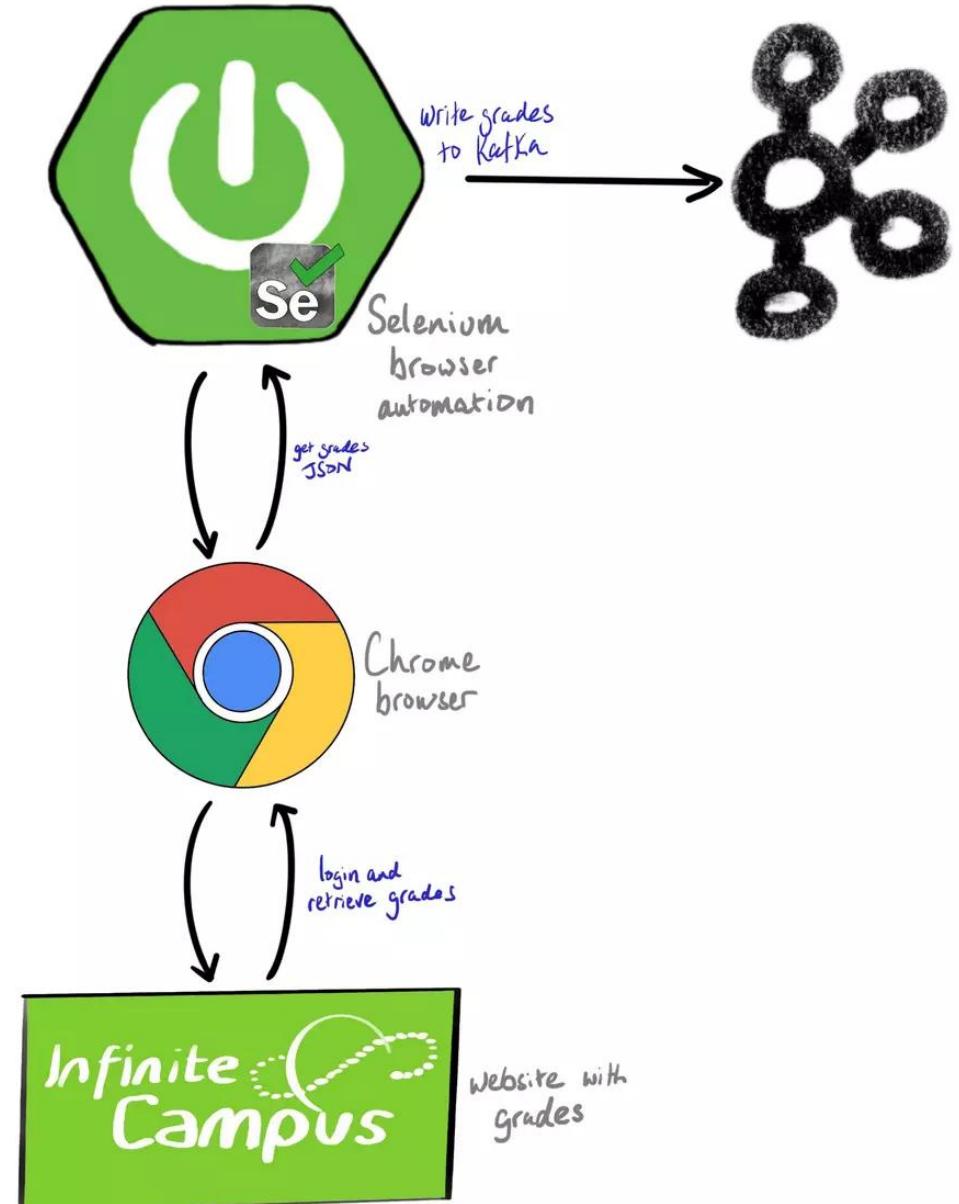
Fortinet firewall



Internet

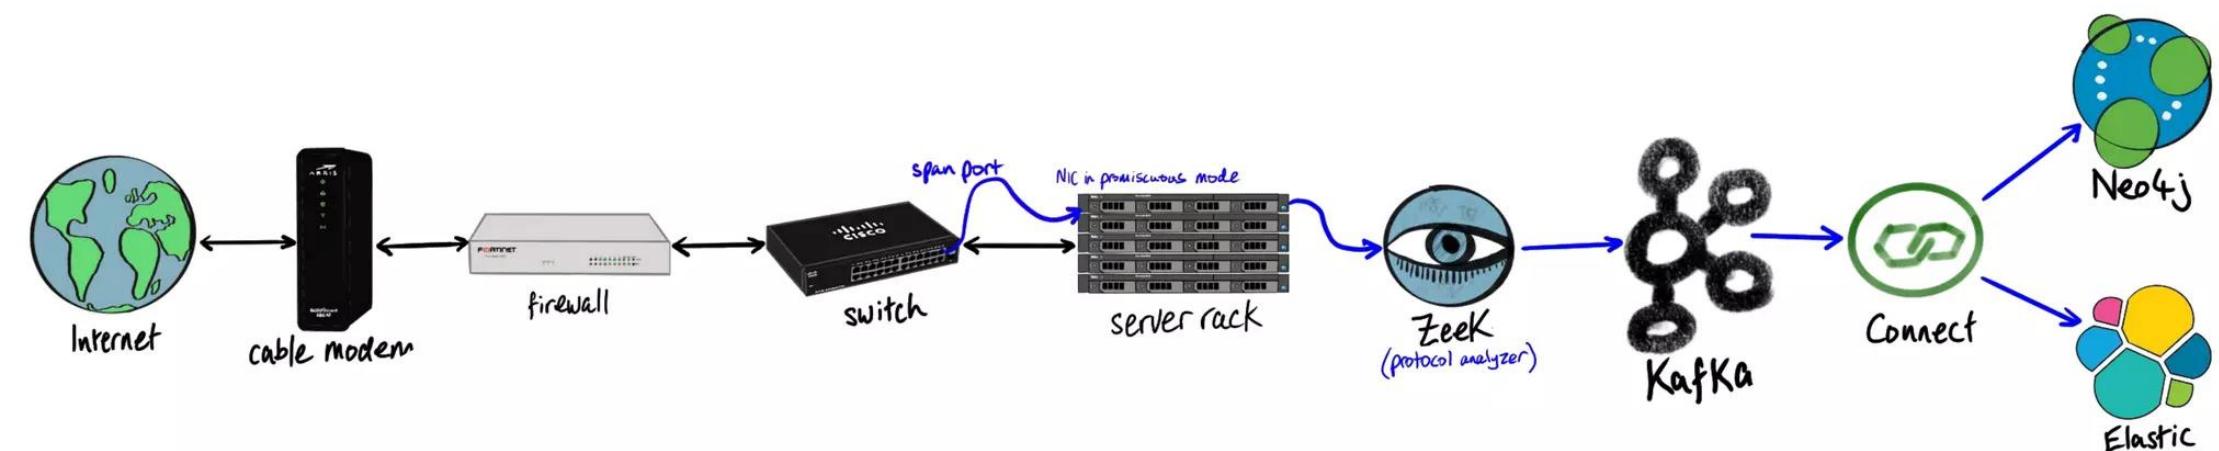
Bad grades?

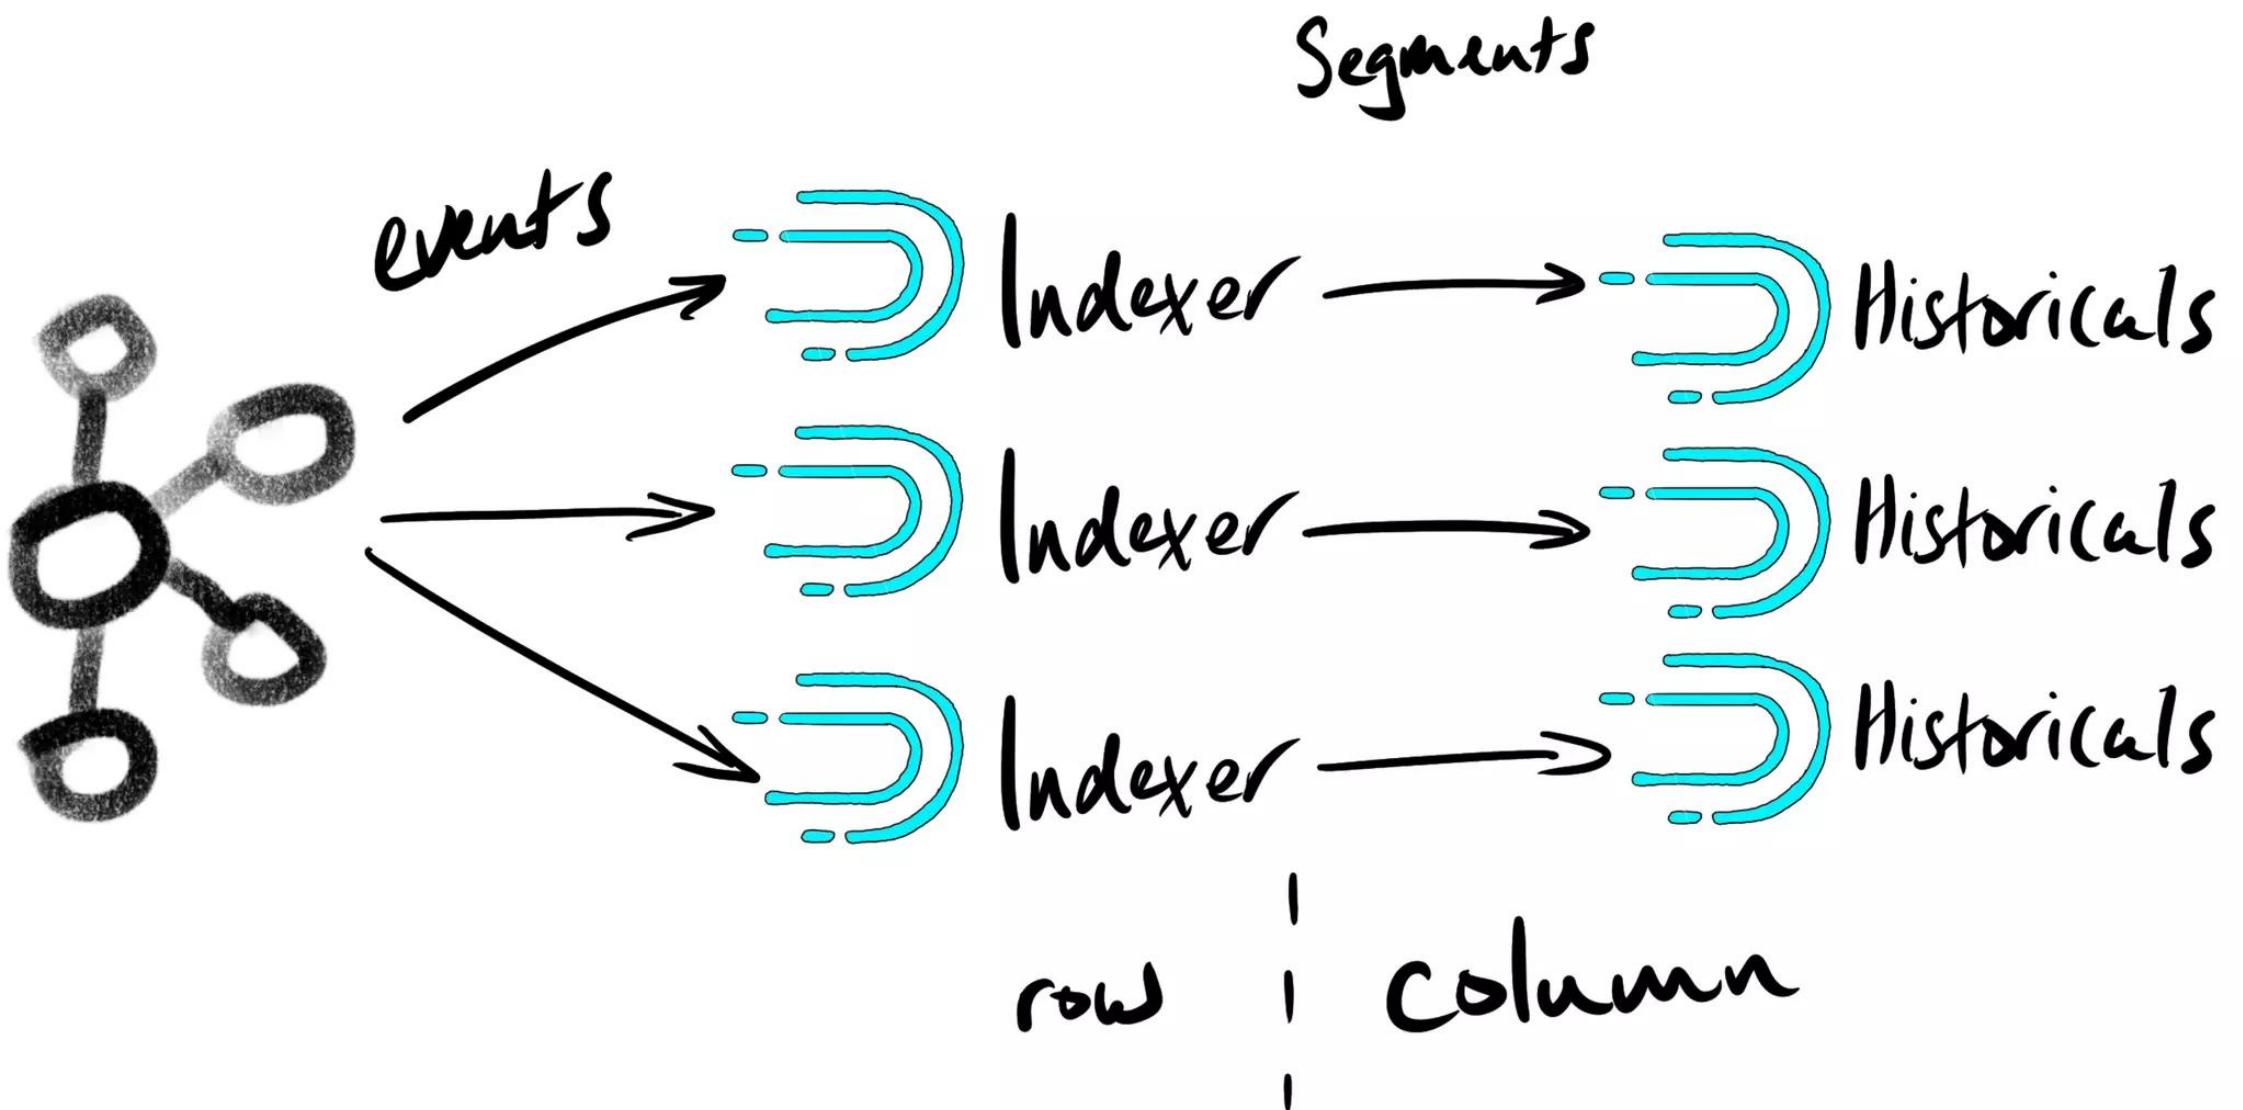
No PS4.



Real-time network traffic analysis

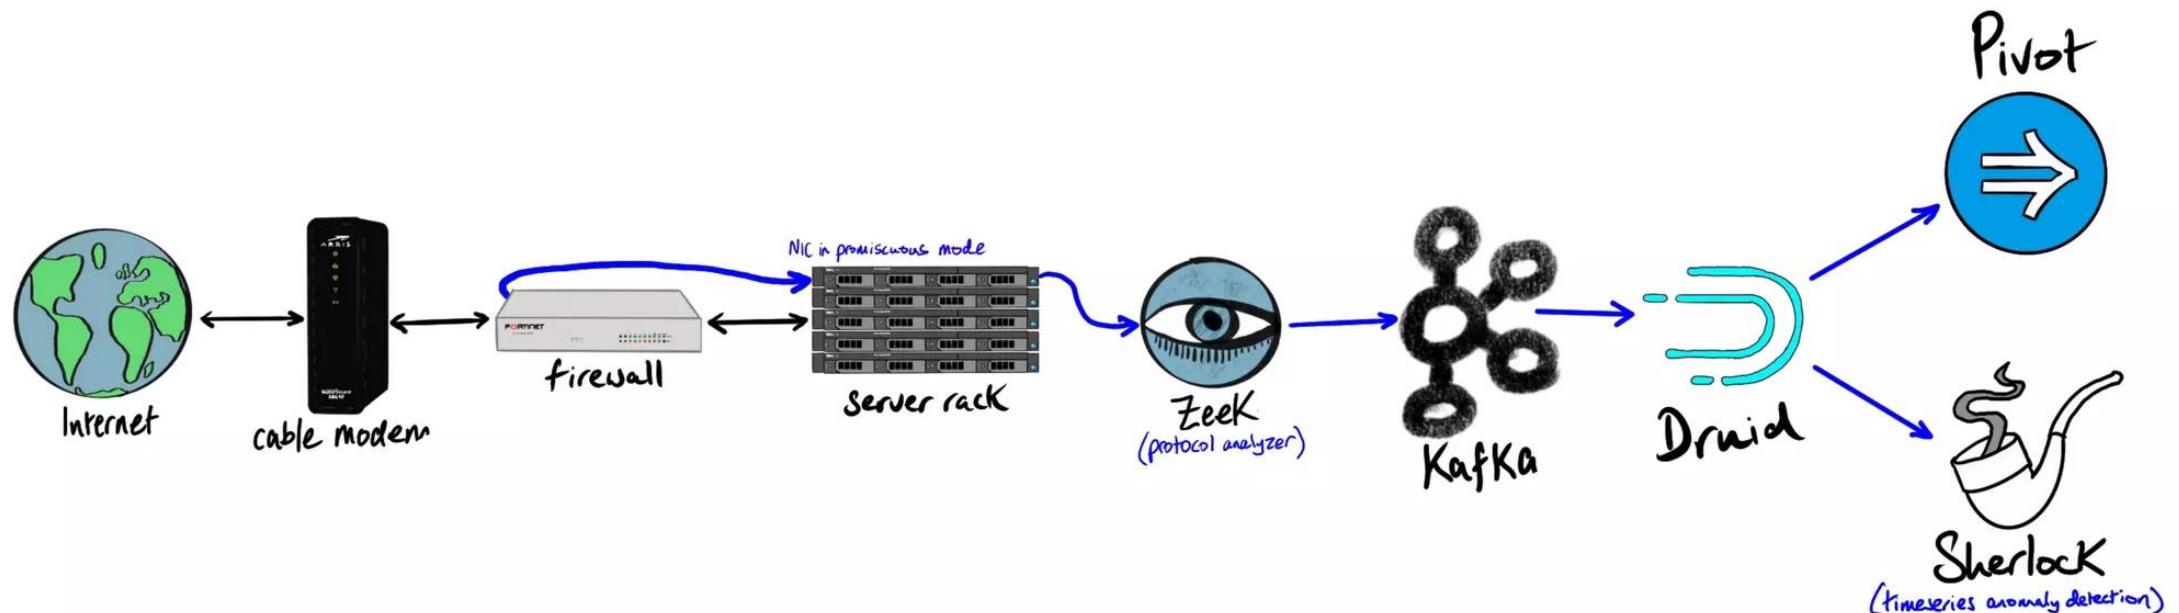
...with Zeek, Kafka, Neo4j, and Elastic





Real-time network traffic analysis

... with Zeek, Kafka, Druid, and Pivot



timeseries anomaly detection with Sherlock

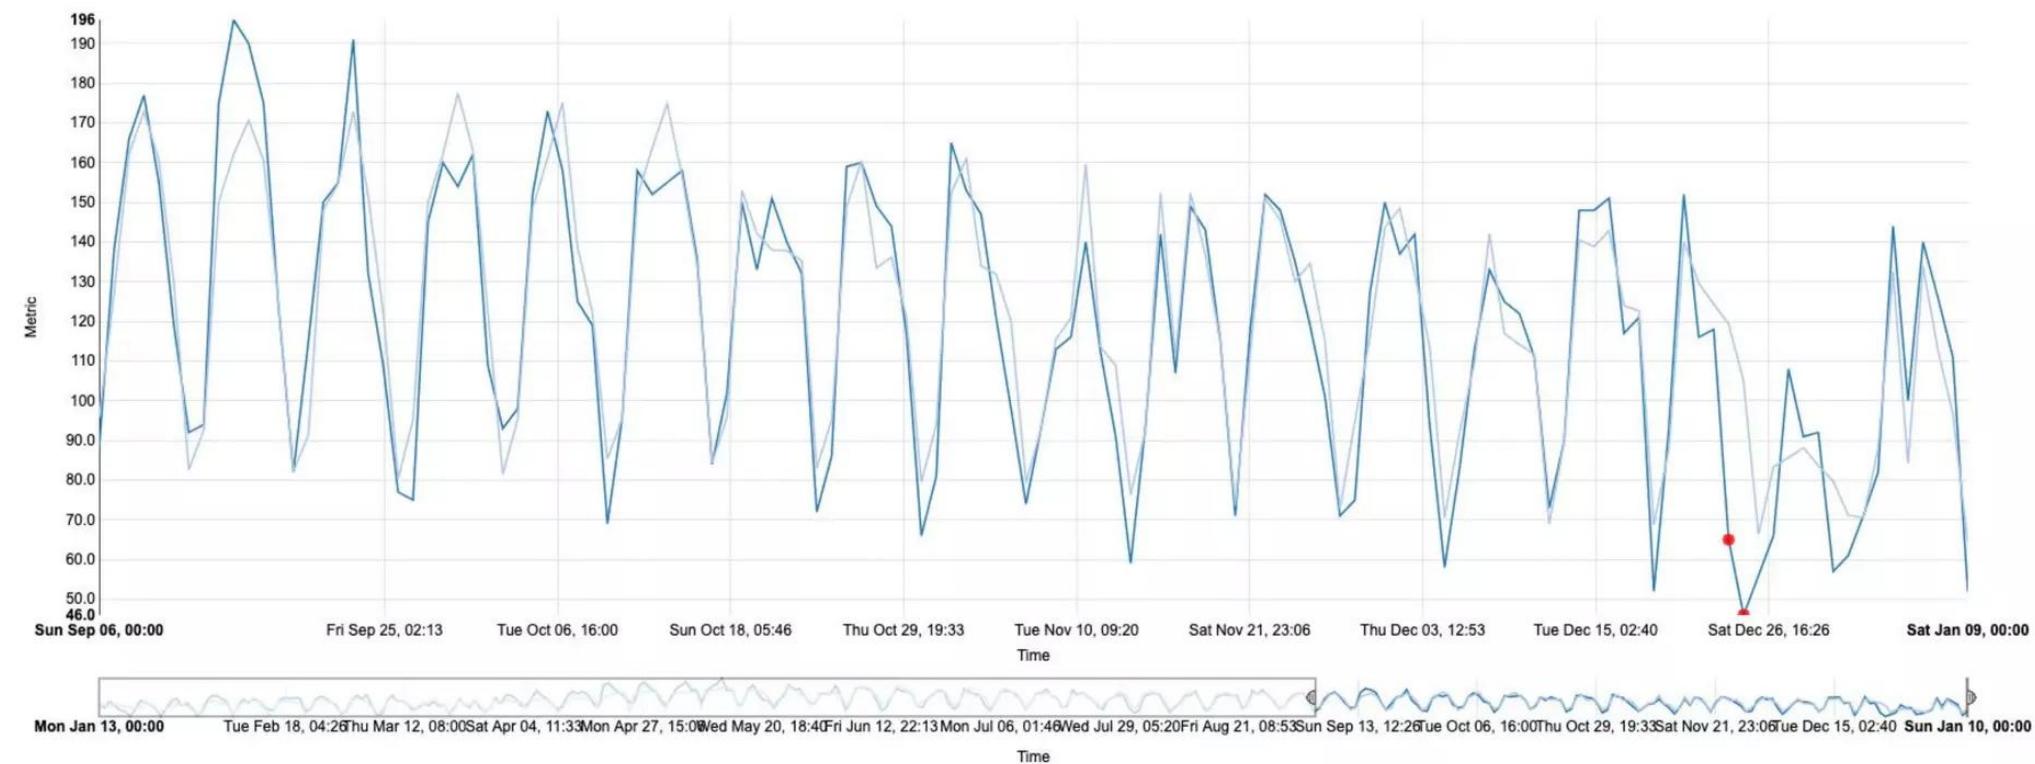
Sherlock



Sherlock

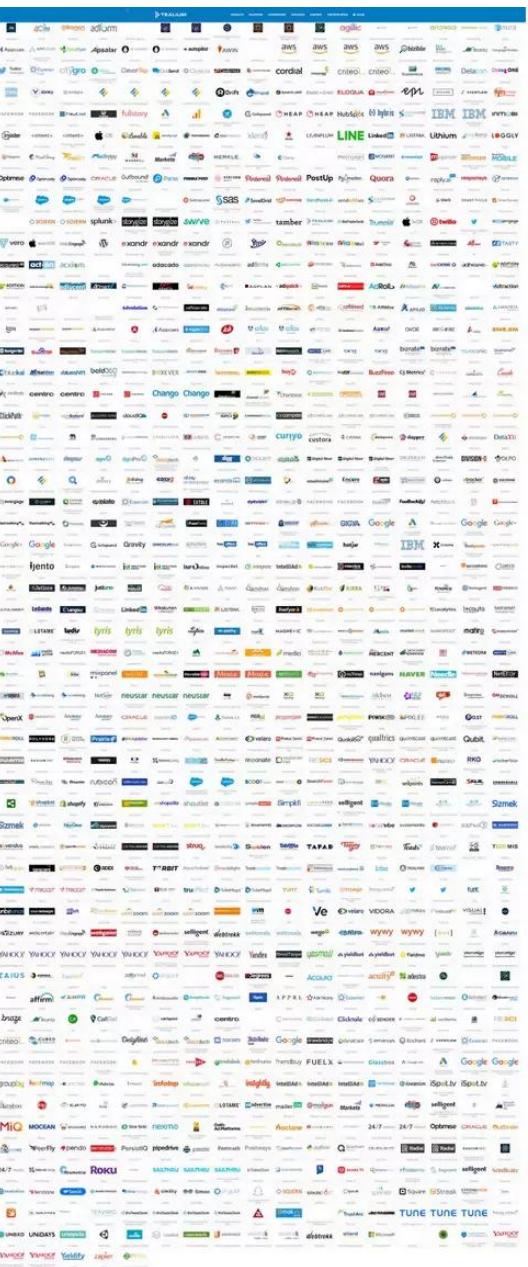
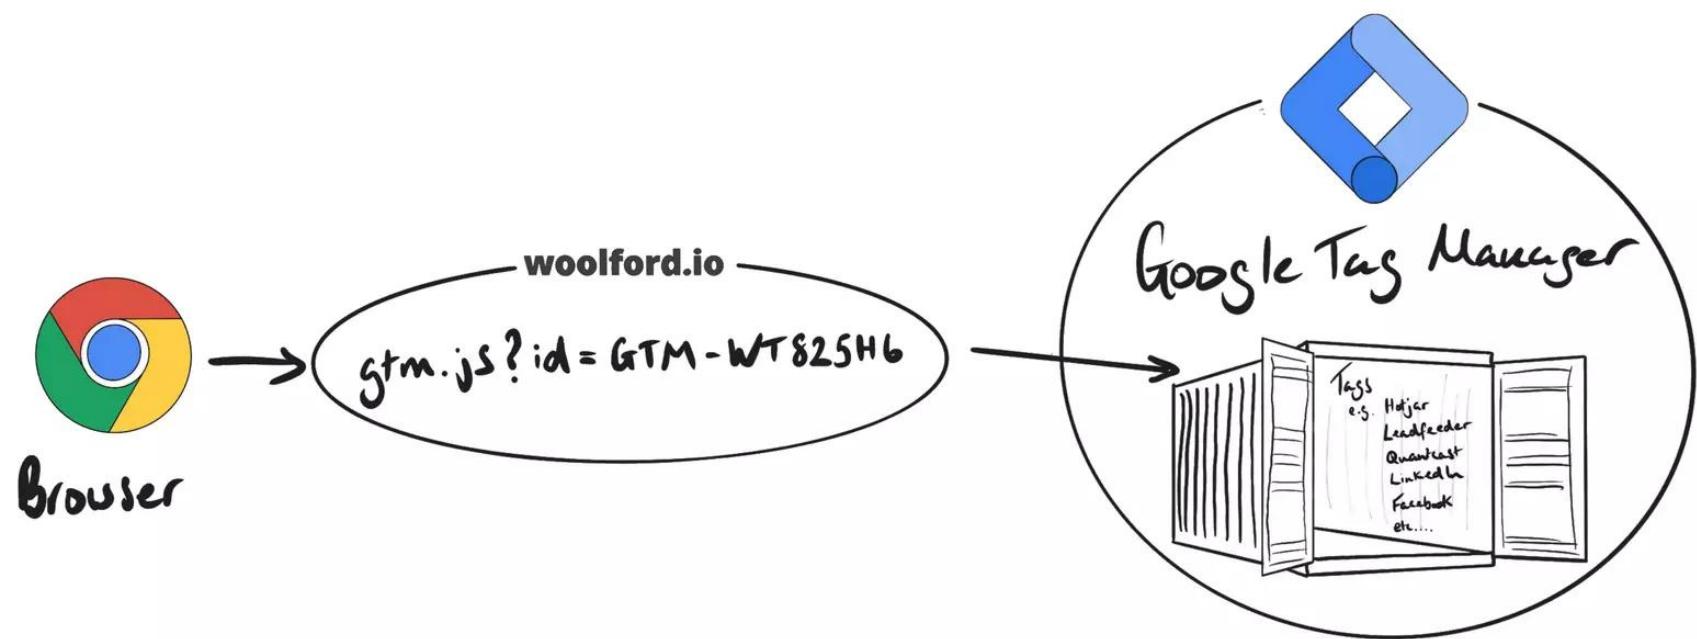
Timeseries List:

views; ▾

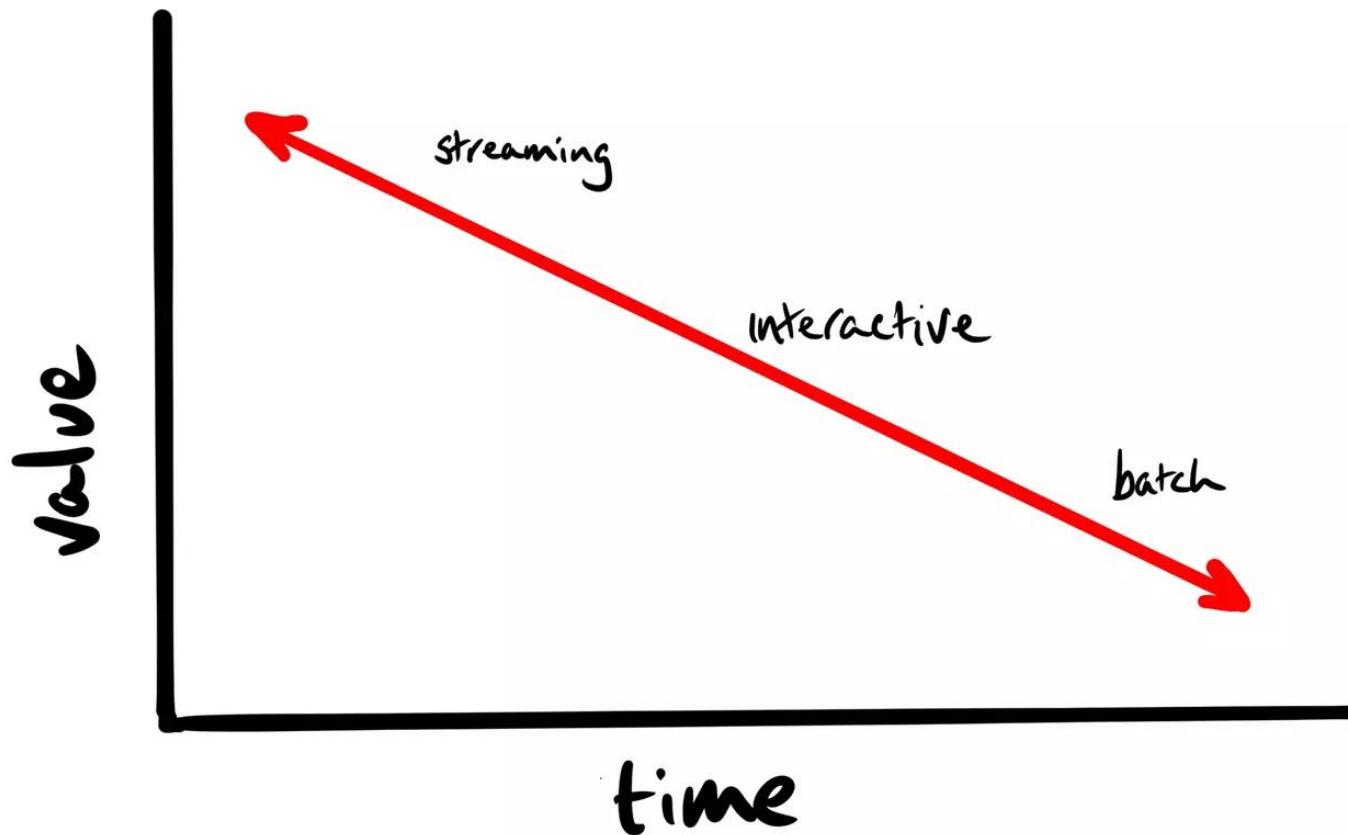


Metric	Group By Dimensions	Anomaly Info	Metric Deviation	Model Info
Metric: views		Thu 24 Dec 2020 00:00:00 UTC Fri 25 Dec 2020 00:00:00 UTC	▼-45% ▼-56%	Model: KSigmaModel Params: 3.0

Javascript Tags

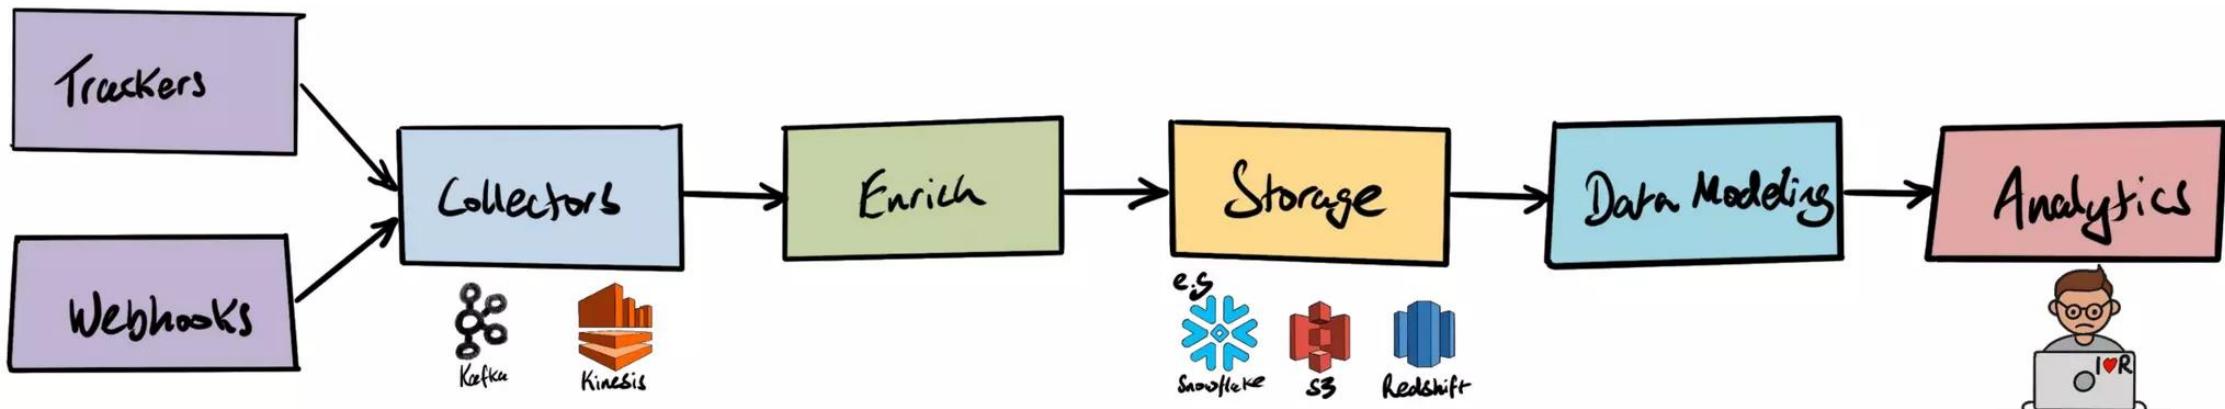


time-value of data

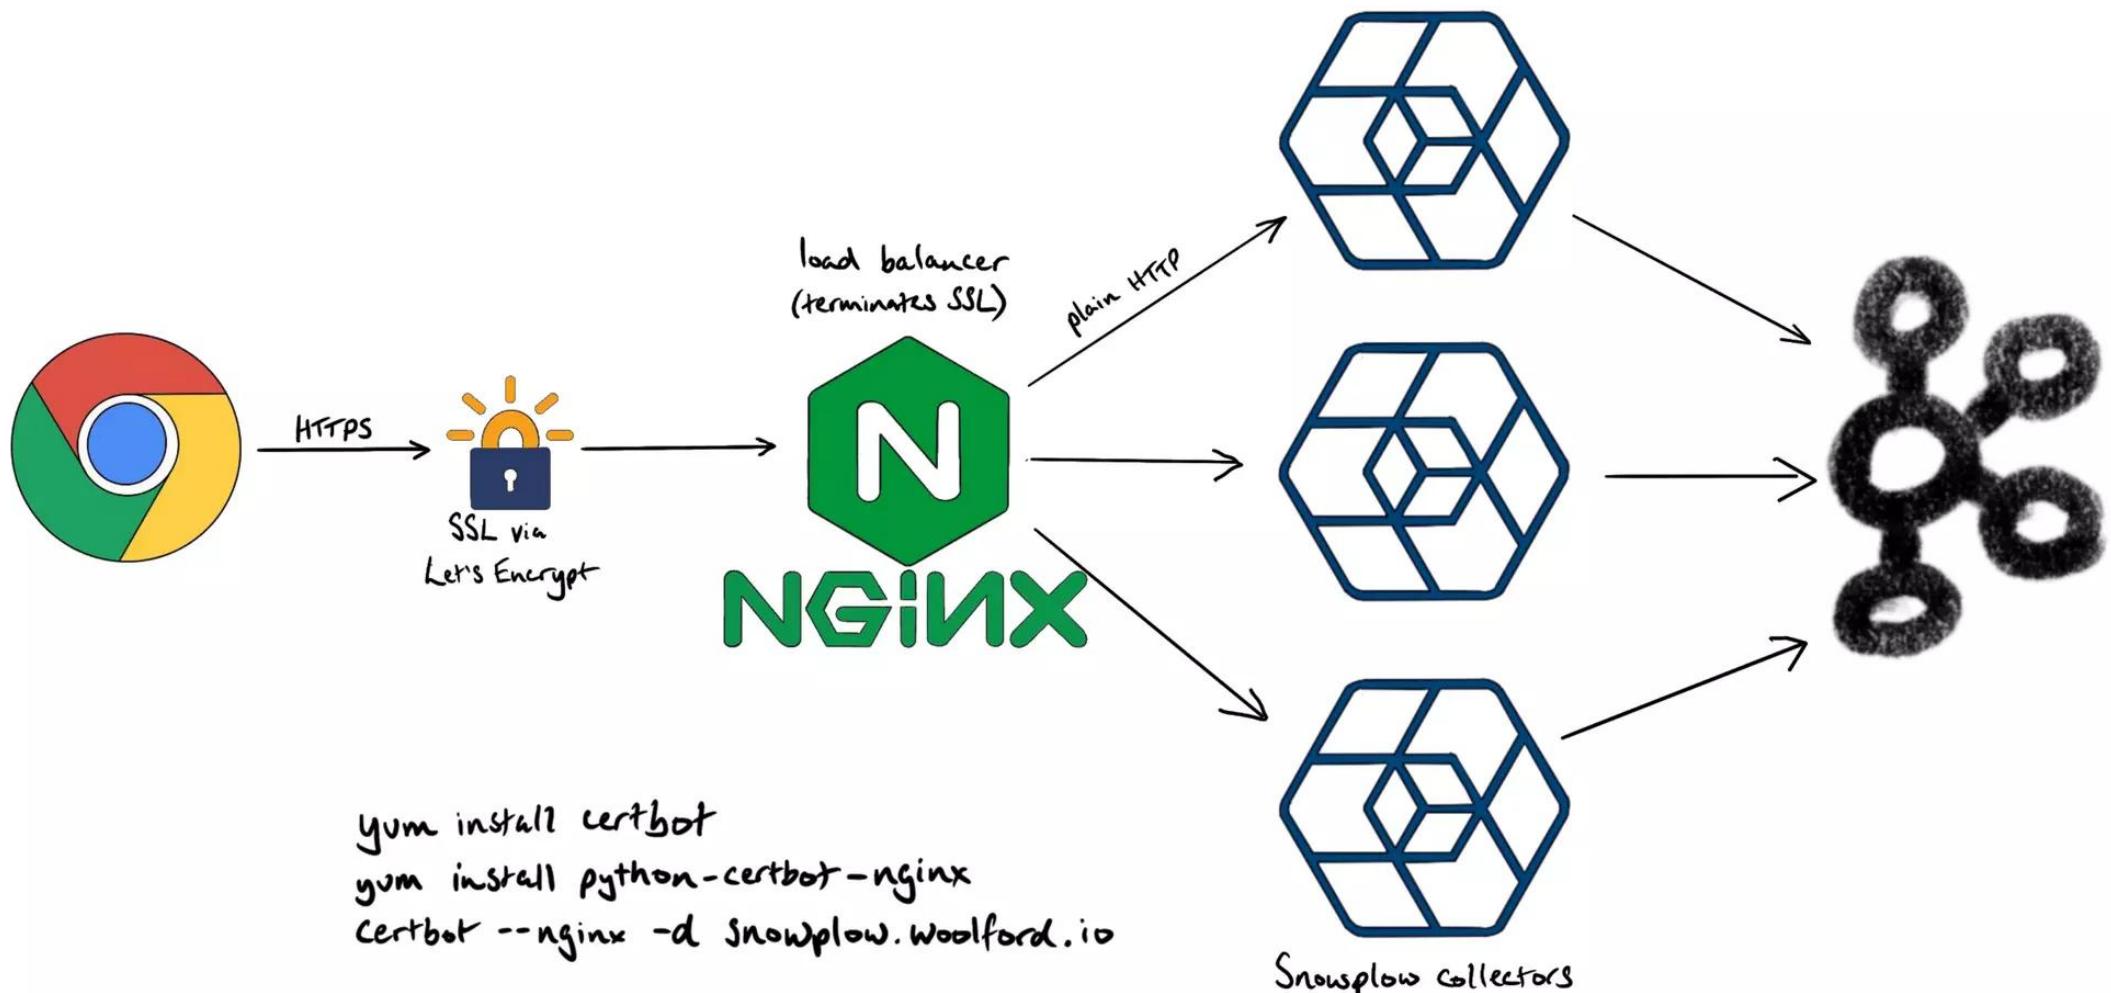


Source: Mike Franklin, UC Berkeley

Snowplow Analytics



SSL with Let's Encrypt; terminated by Nginx



Visualizing Kafka Streams topologies

...

final Topology = builder.build;

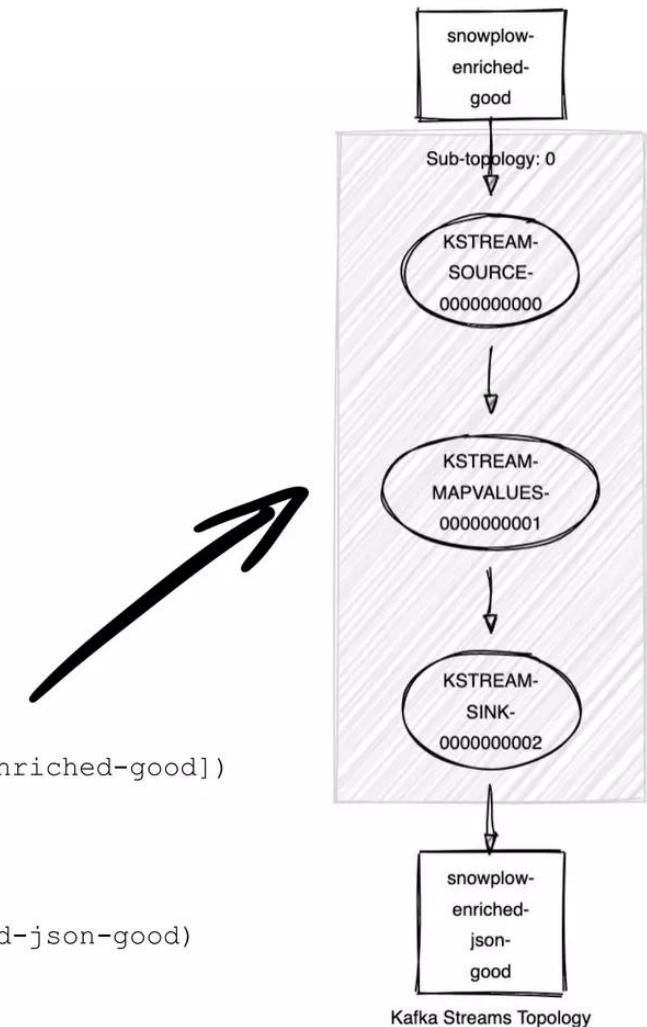


paste the output from topology.describe() into

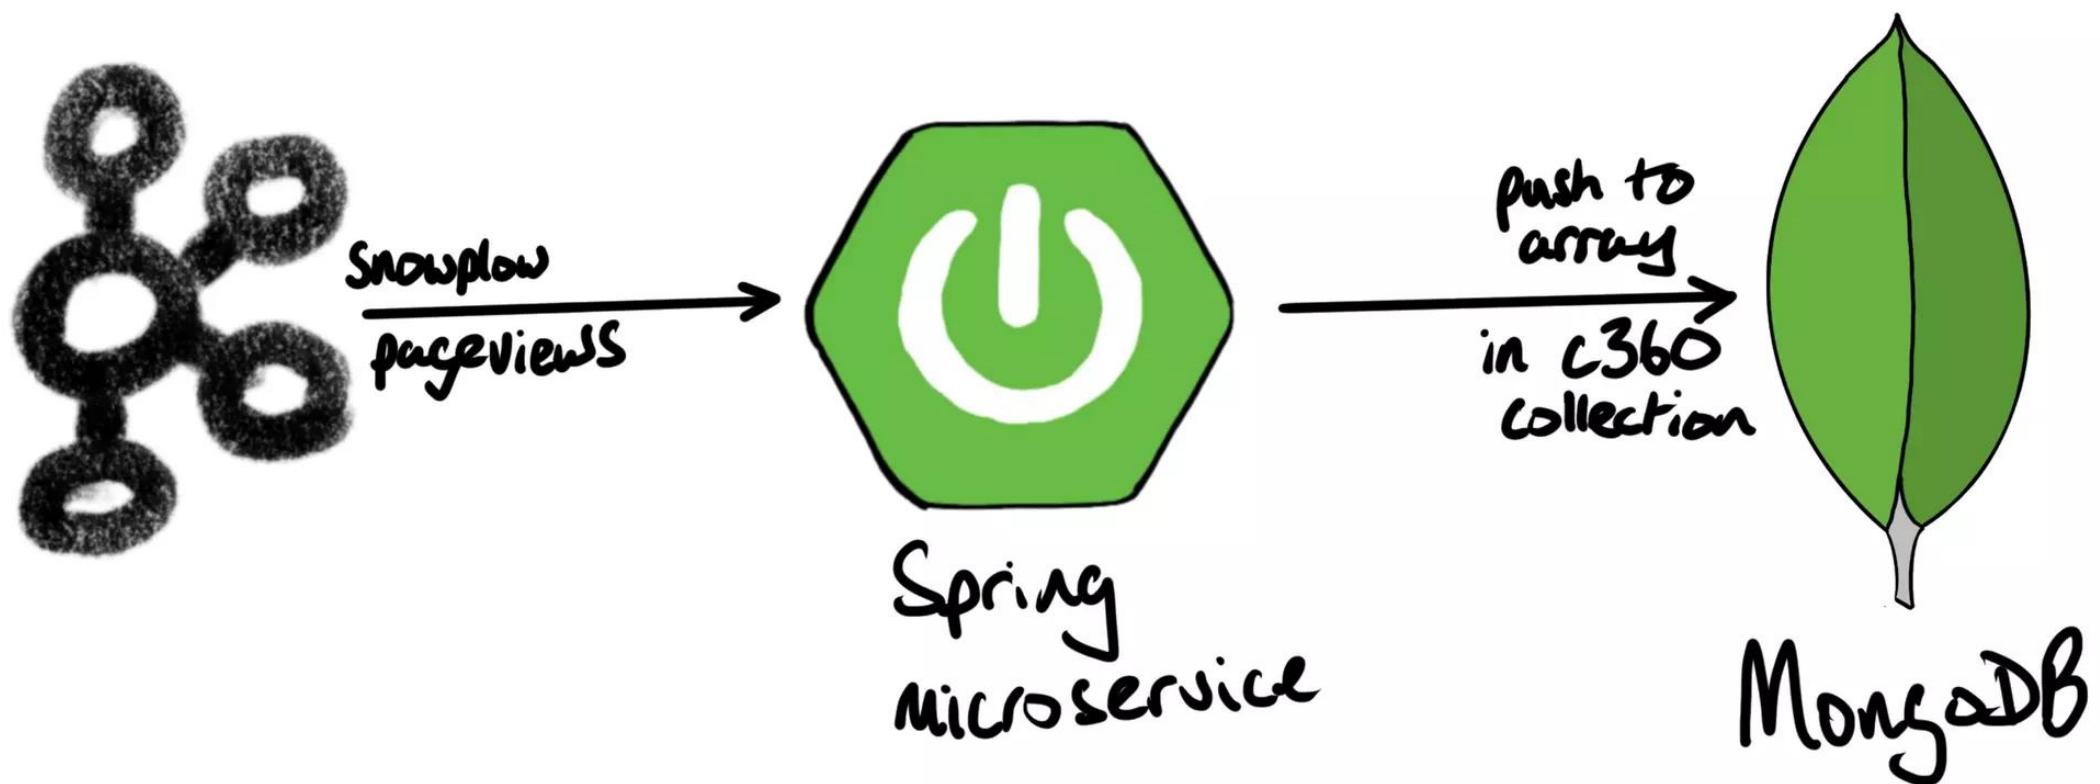


zz85.github.io/Kafka-streams-viz

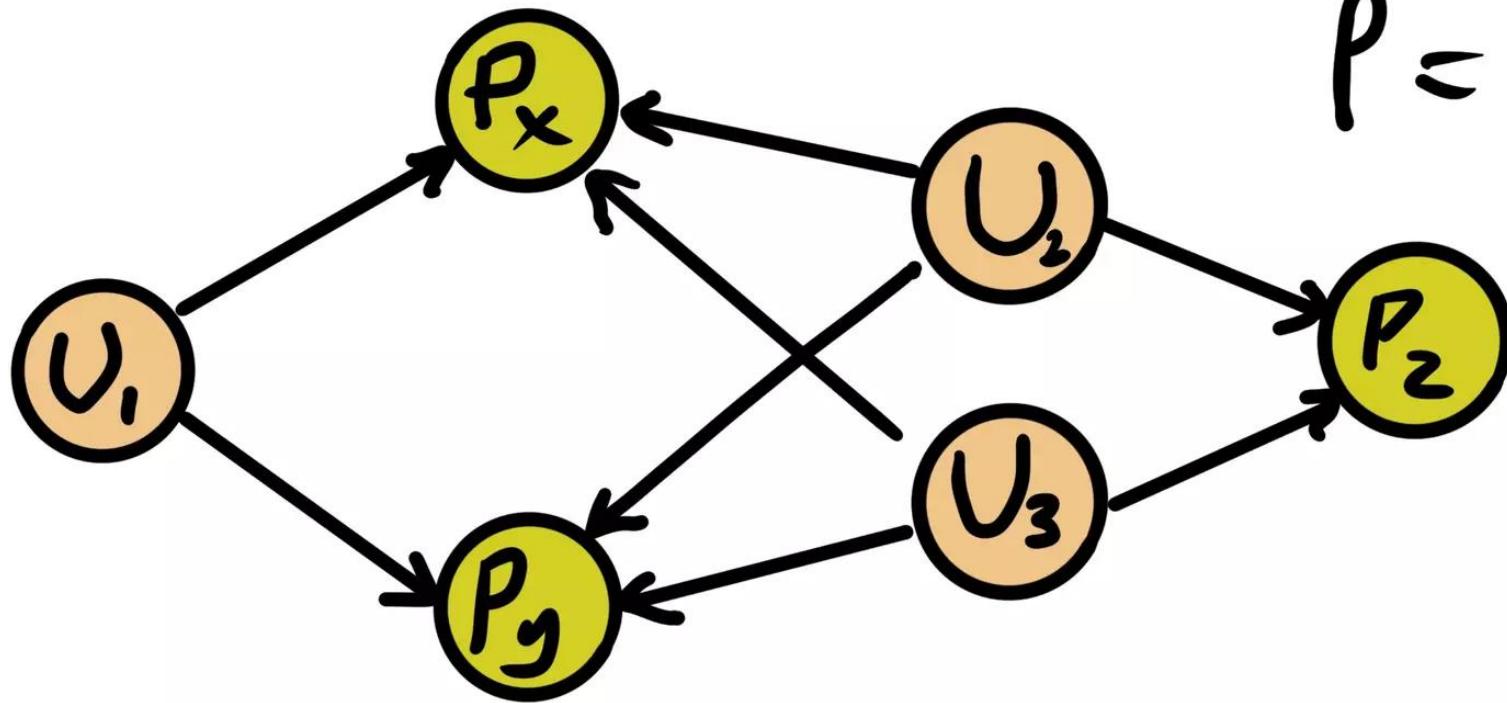
```
Sub-topology: 0
Source: KSTREAM-SOURCE-0000000000 (topics: [snowplow-enriched-good])
--> KSTREAM-MAPVALUES-0000000001
Processor: KSTREAM-MAPVALUES-0000000001 (stores: [])
--> KSTREAM-SINK-0000000002
<-- KSTREAM-SOURCE-0000000000
Sink: KSTREAM-SINK-0000000002 (topic: snowplow-enriched-json-good)
--> KSTREAM-MAPVALUES-0000000001
```



build profiles



Kraph-based recommender



$U = \text{user}$

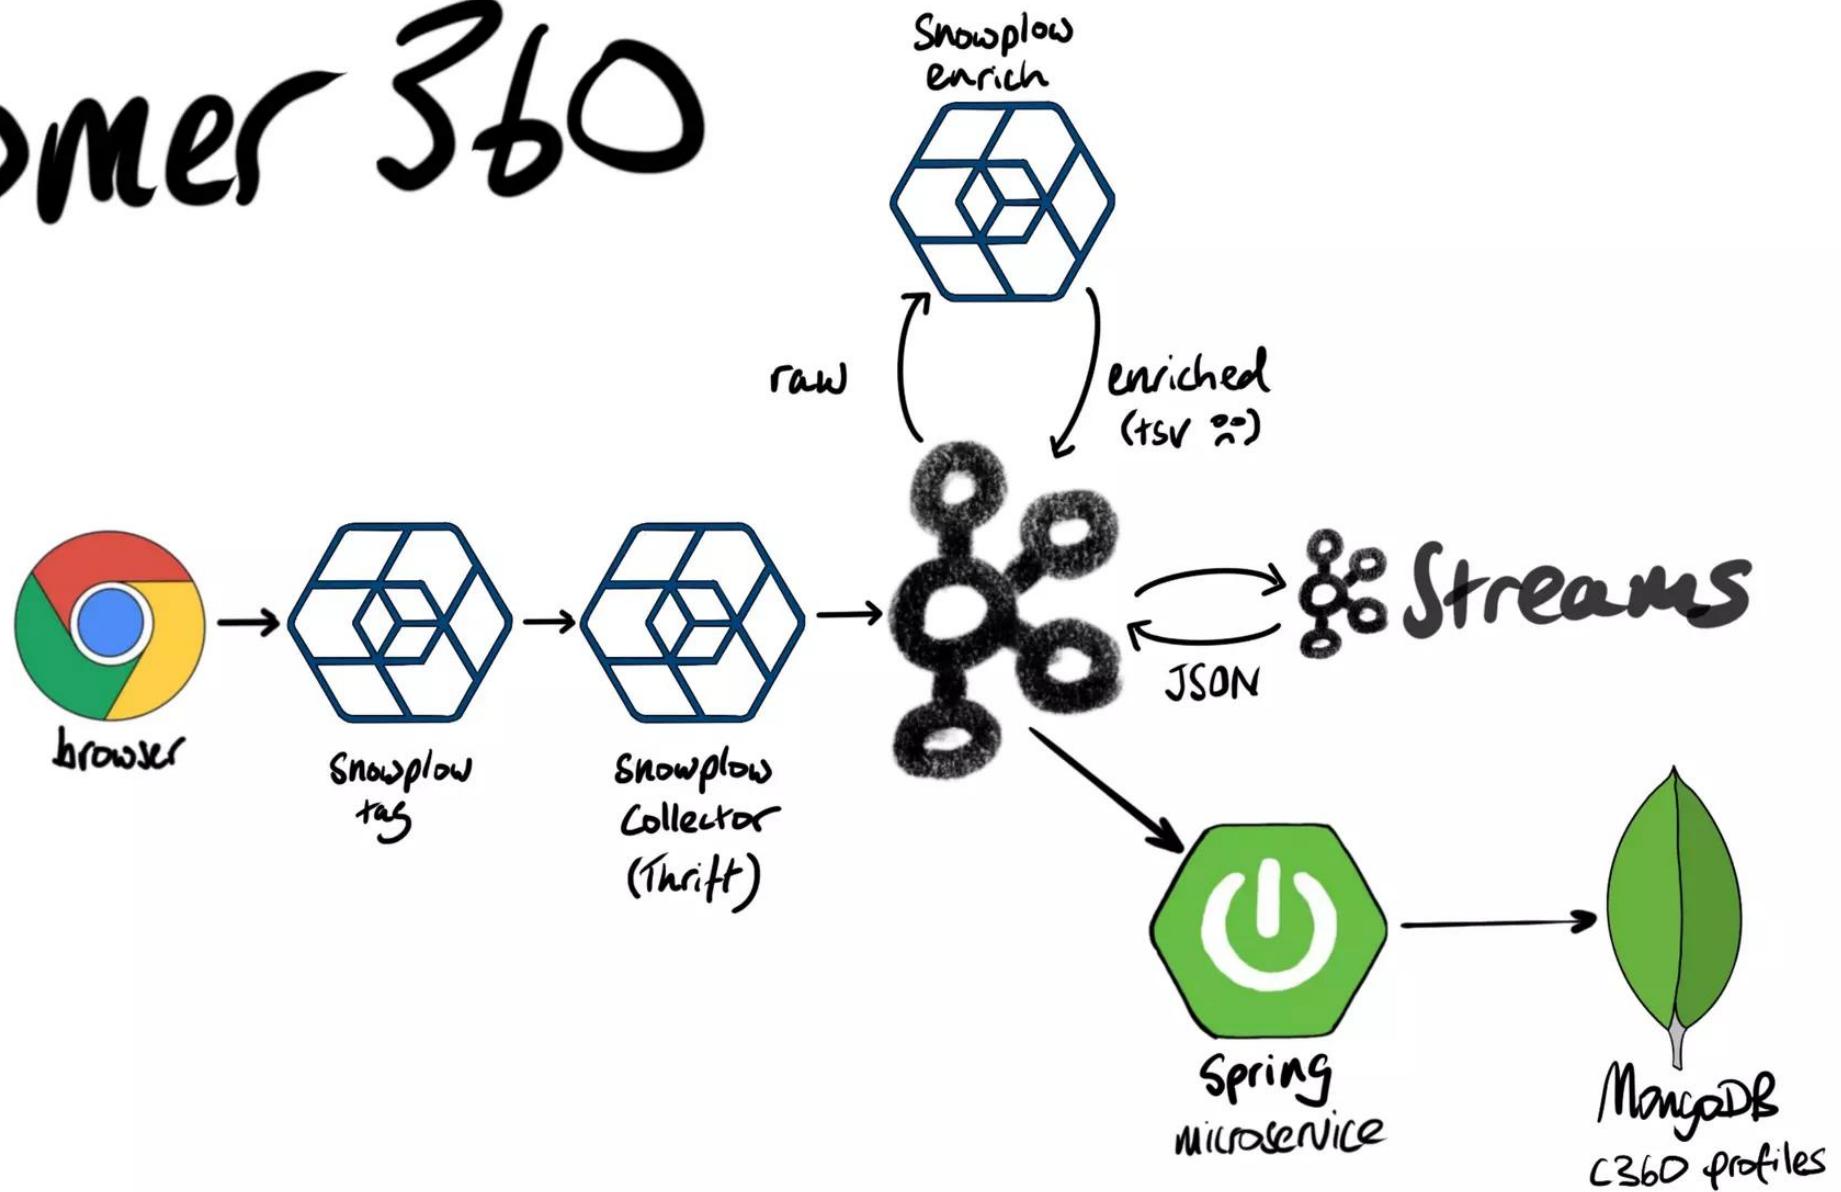
$P = \text{page}$

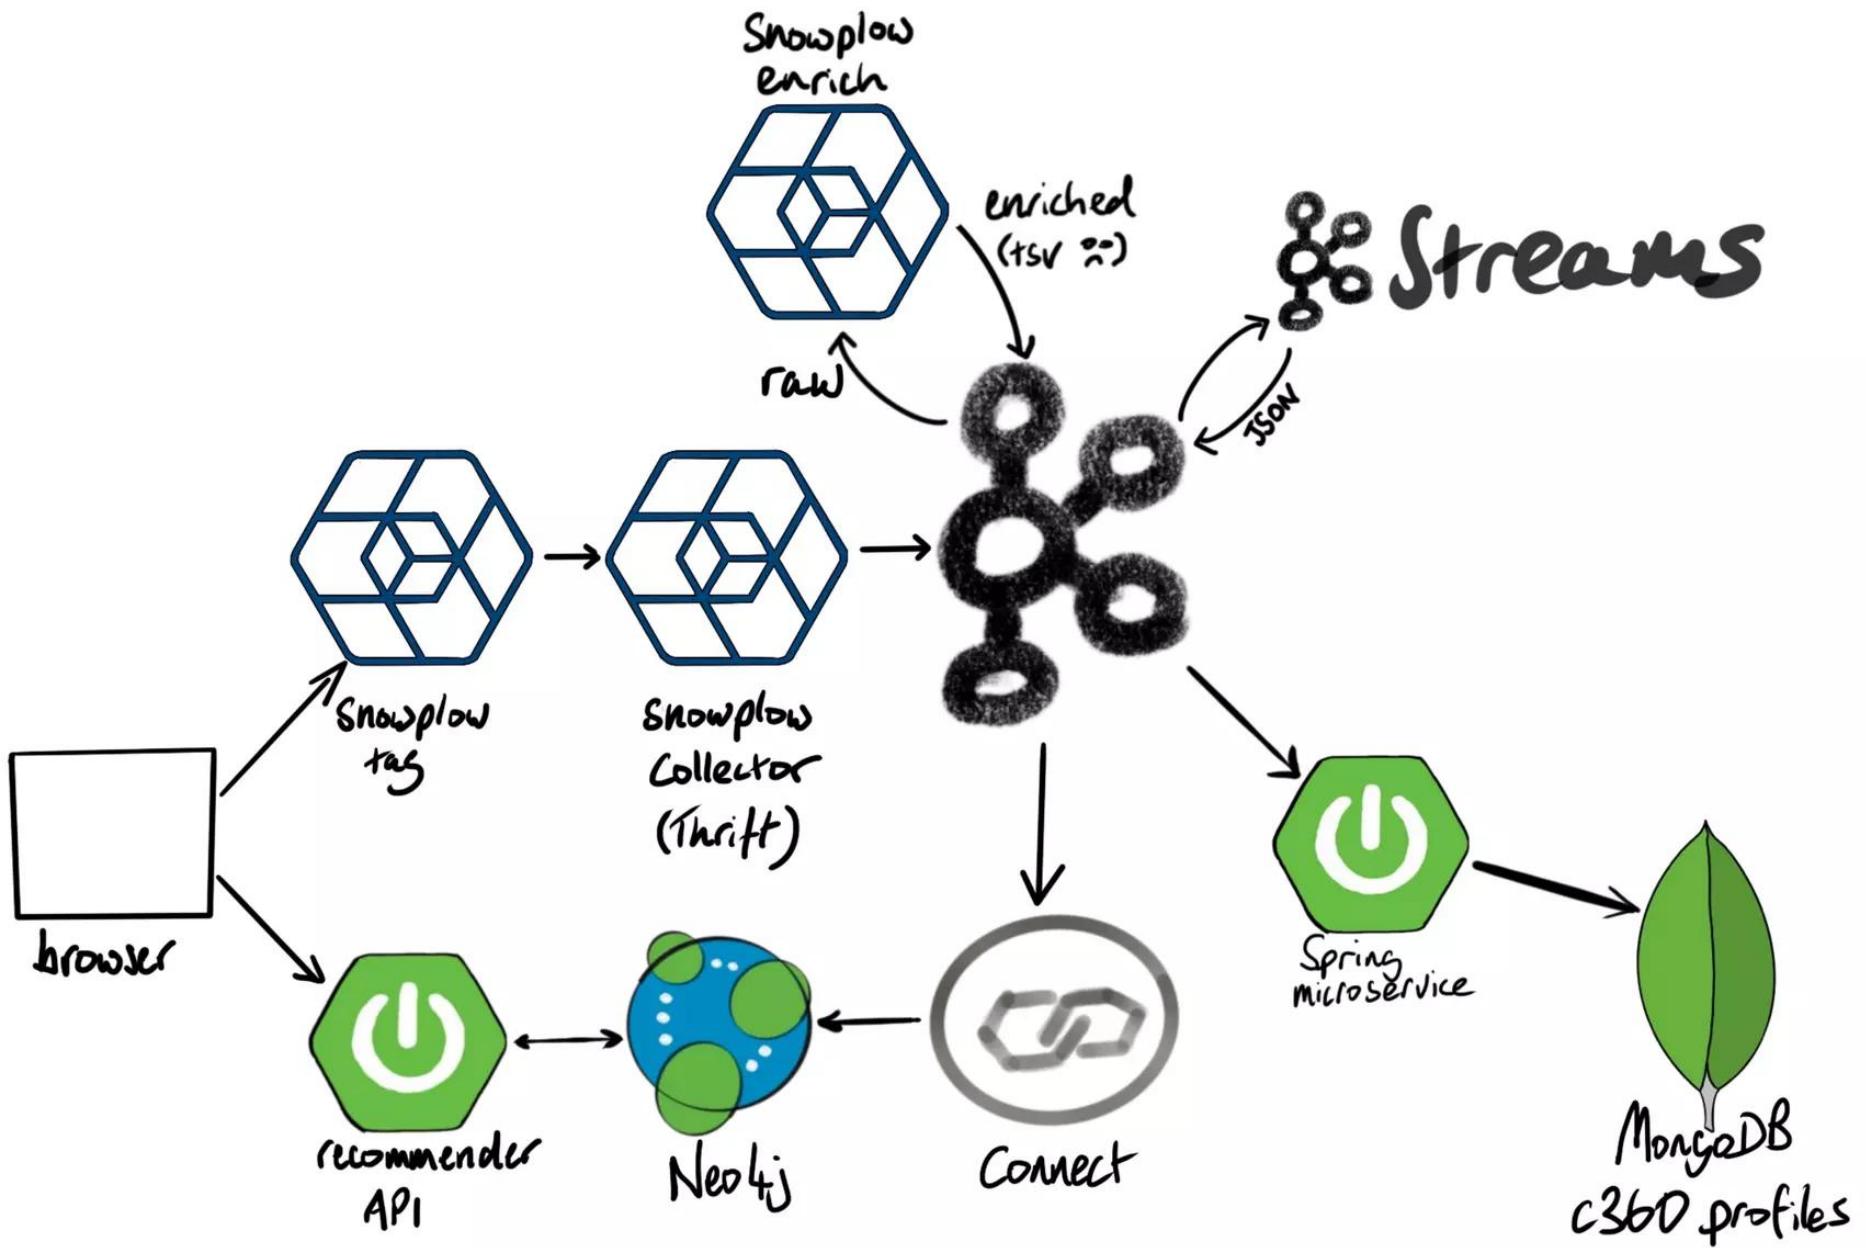
Cypher statement to build the graph

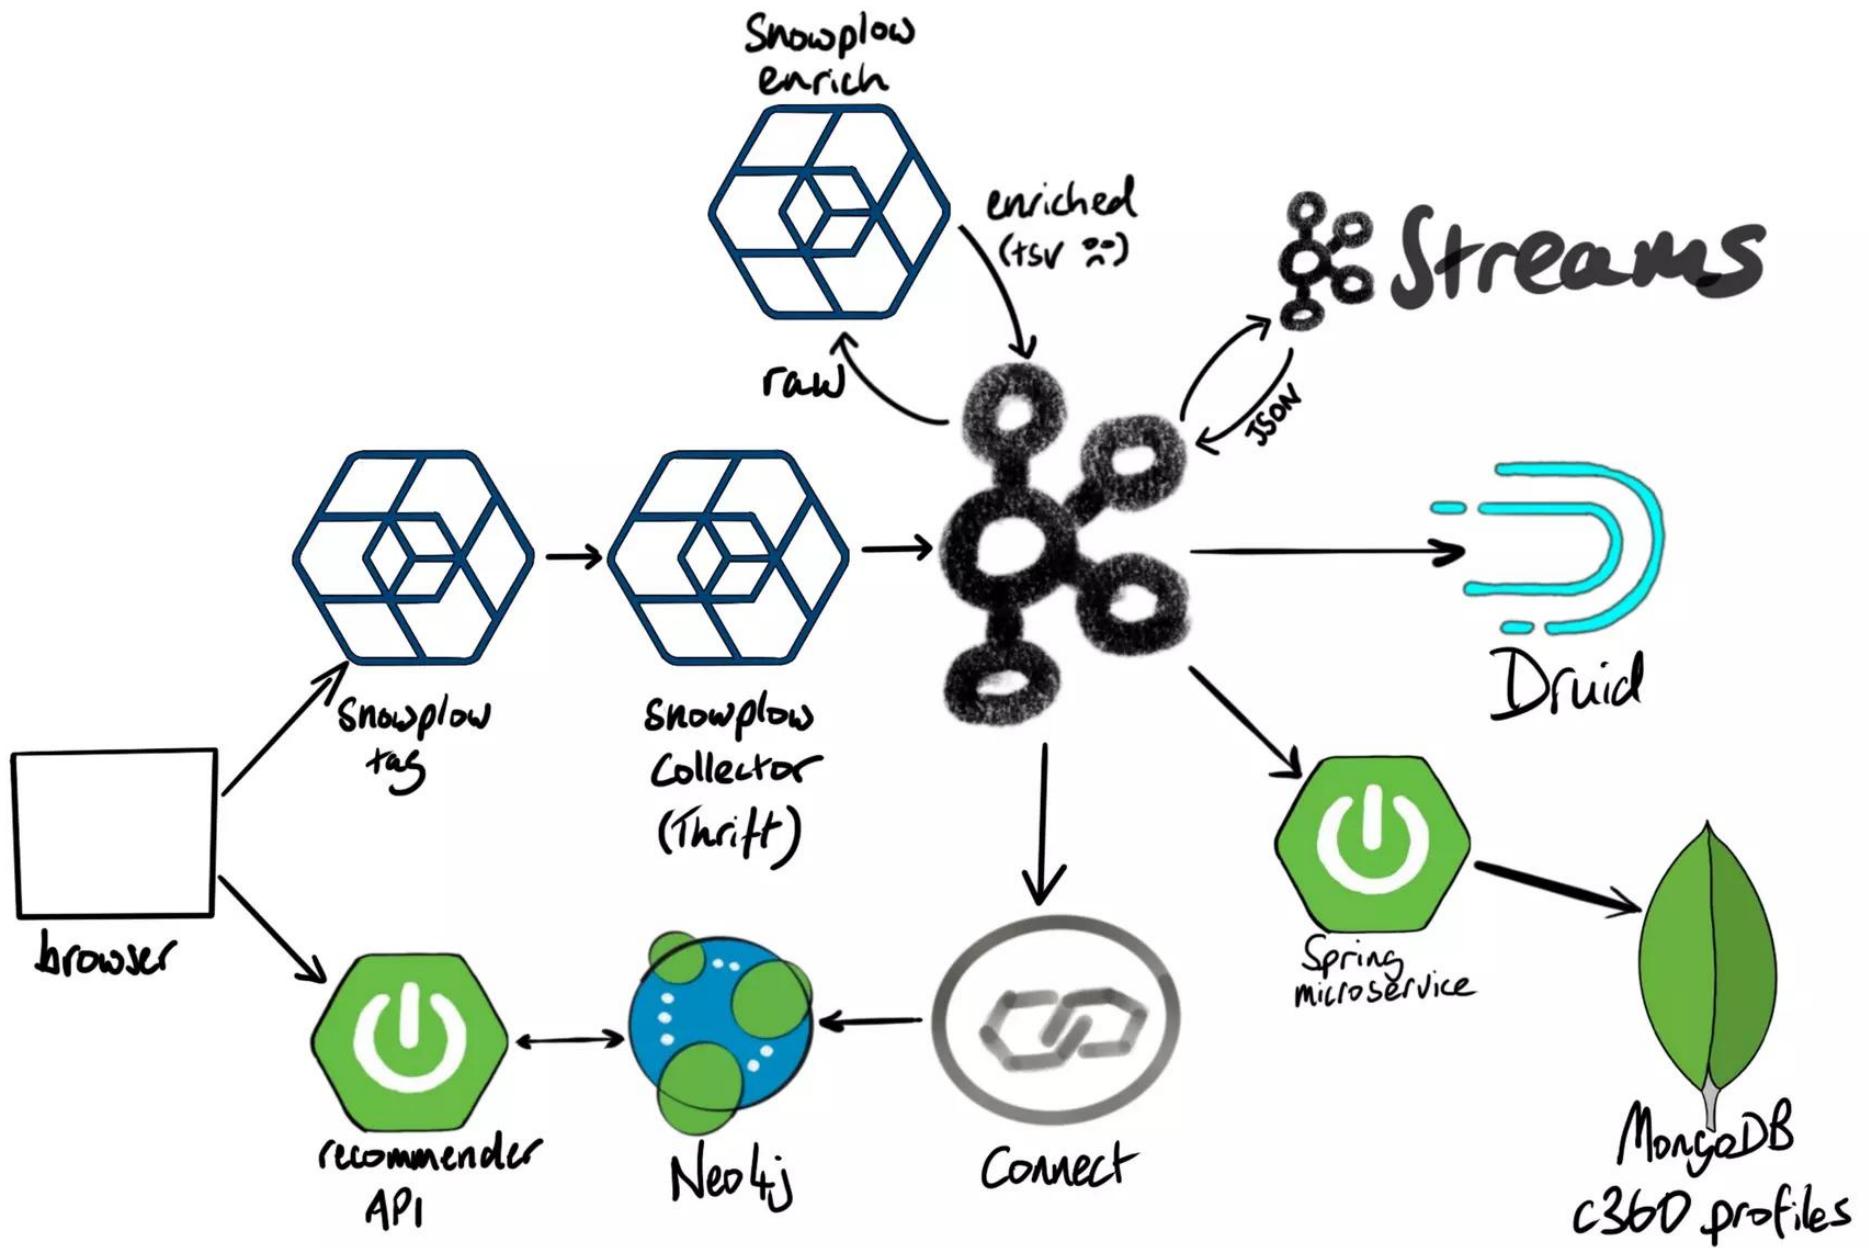
```
merge (network_userid:network_userid { id: event.network_userid })  
merge (page_url:page_url { id: event.page_url })  
merge (network_userid)-[:VIEWED]-(page_url)
```

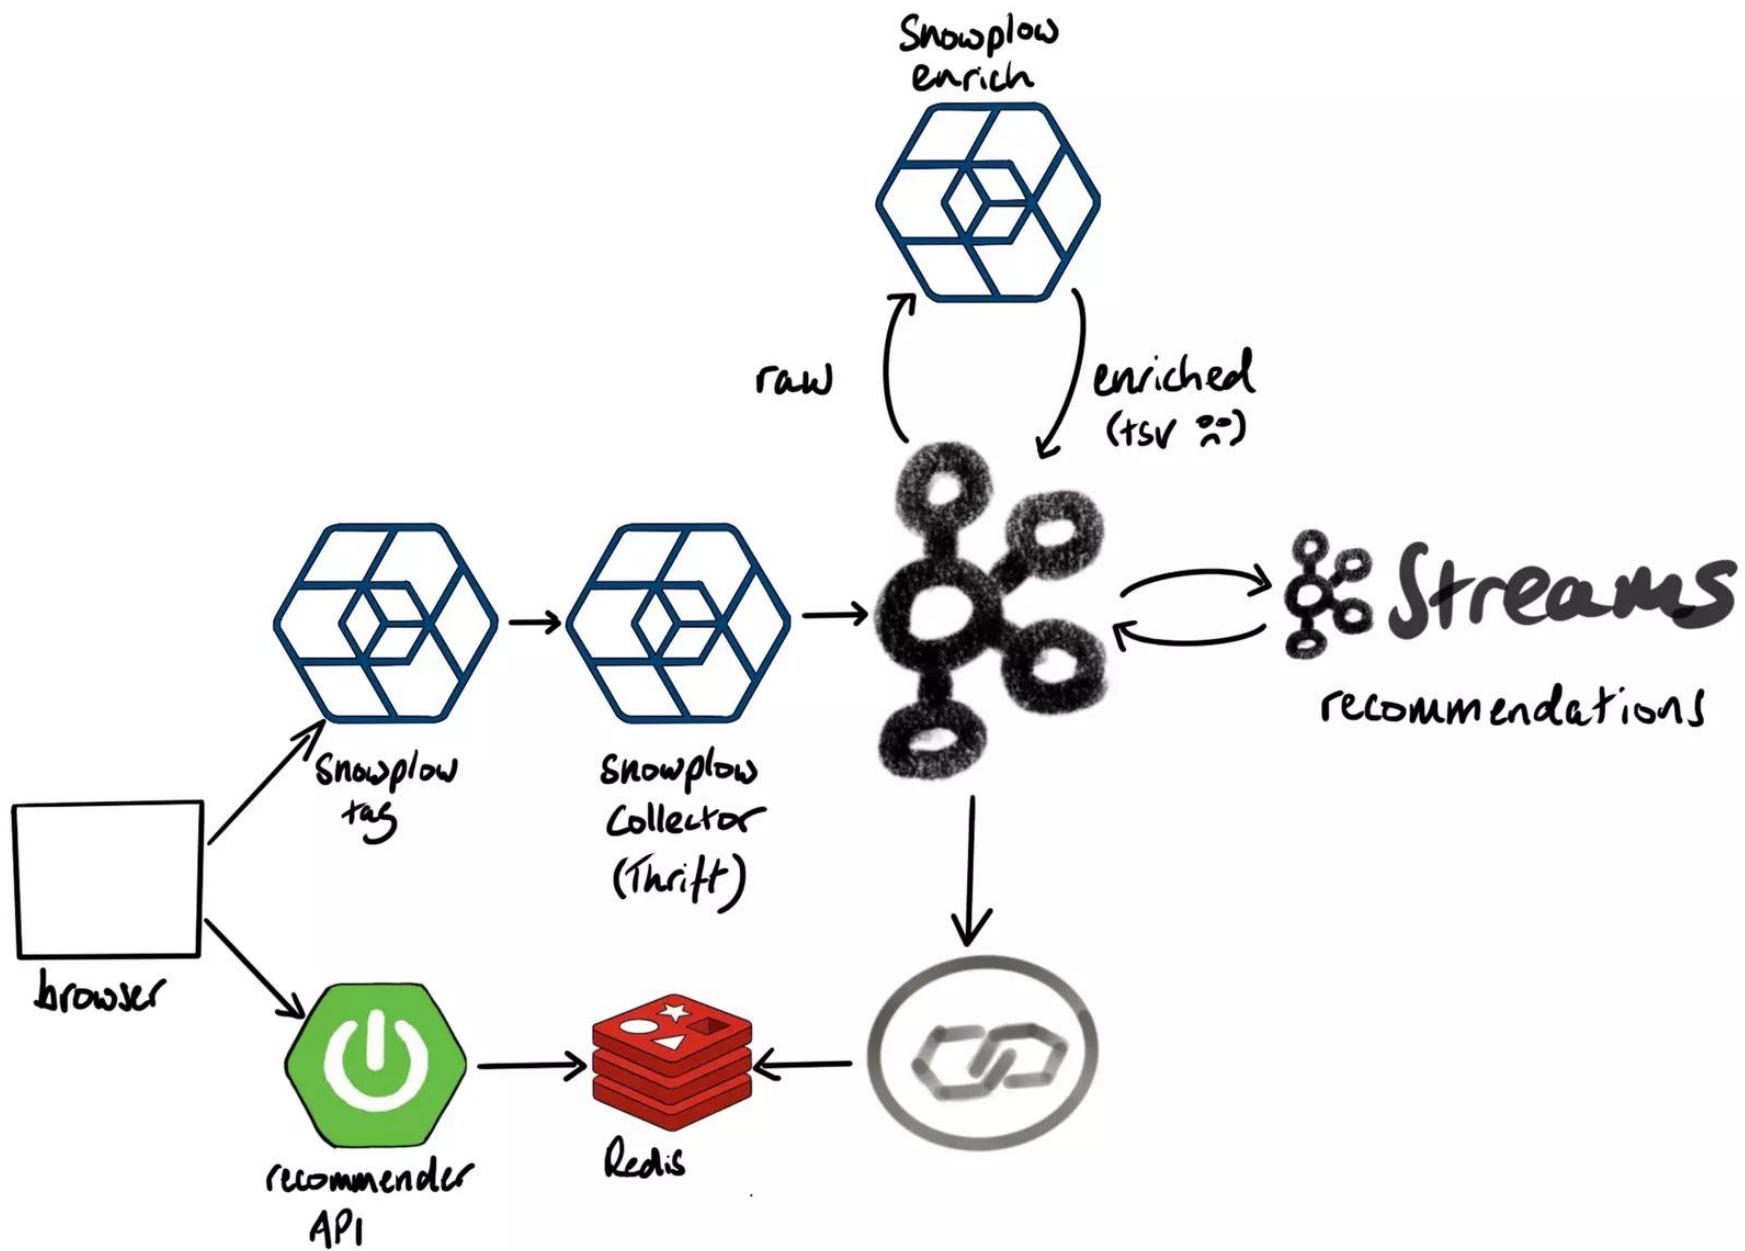
① get or create user
② get or create page
③ create relationship

Customer 360

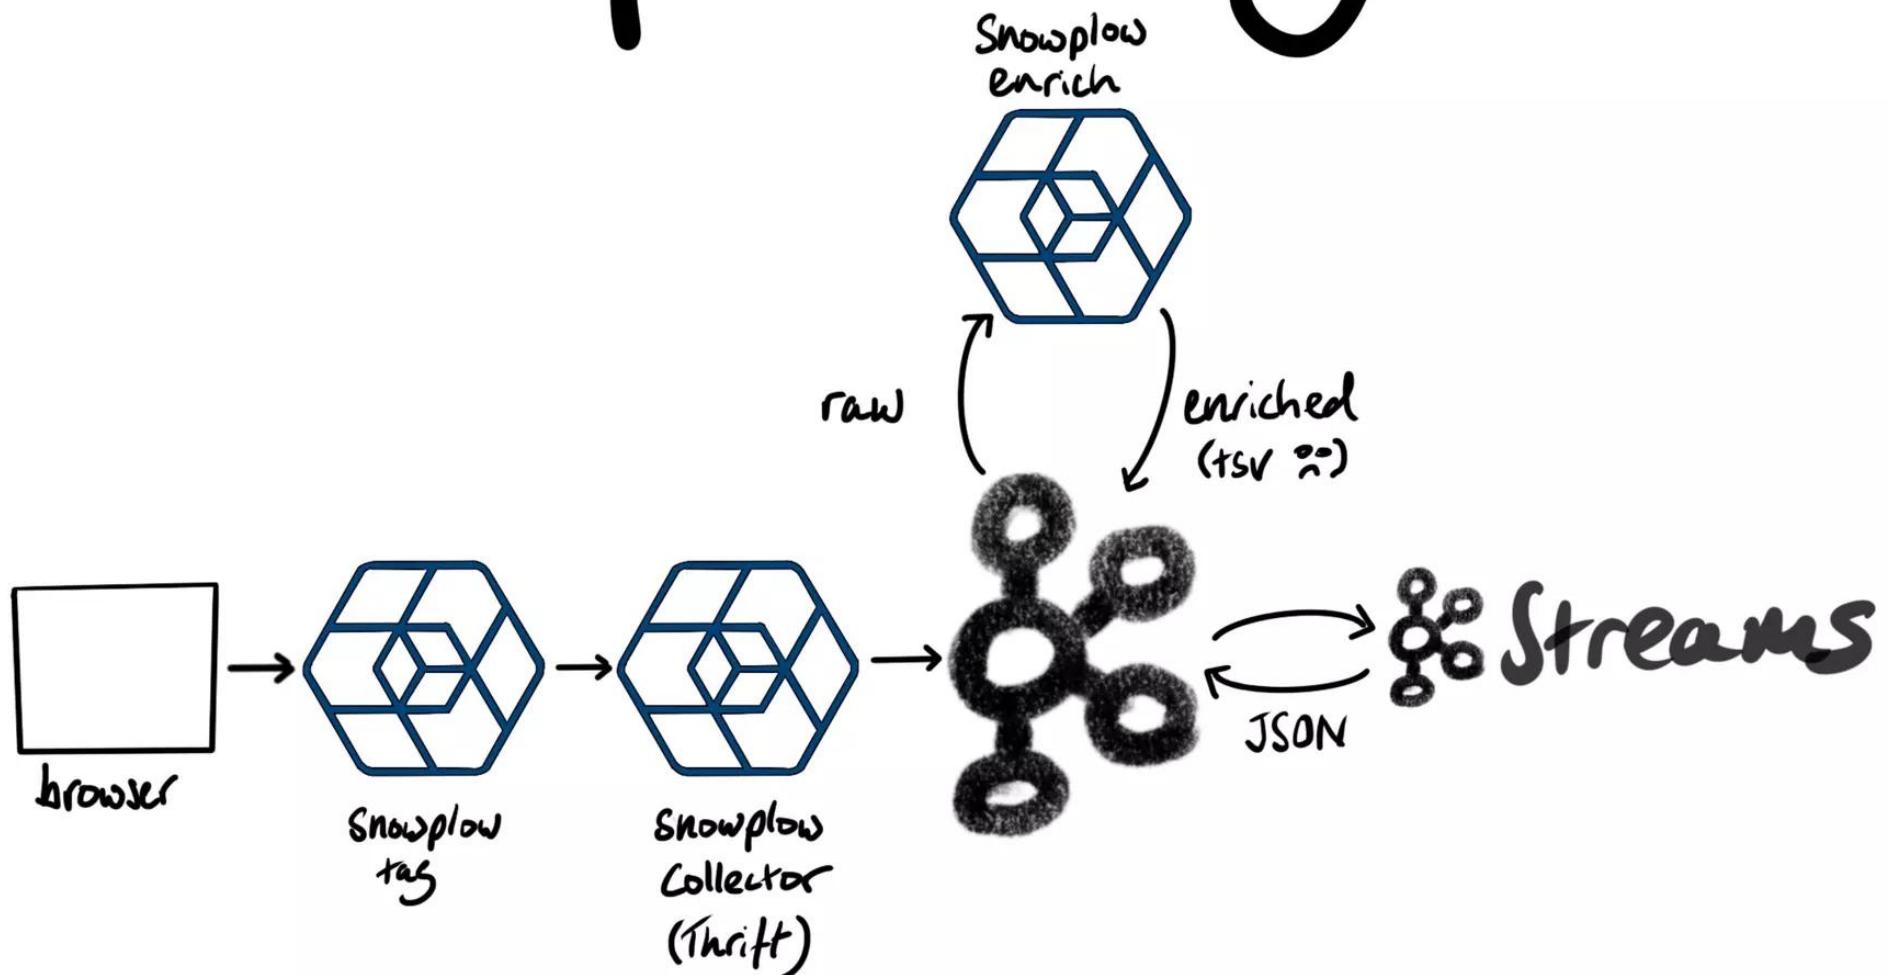




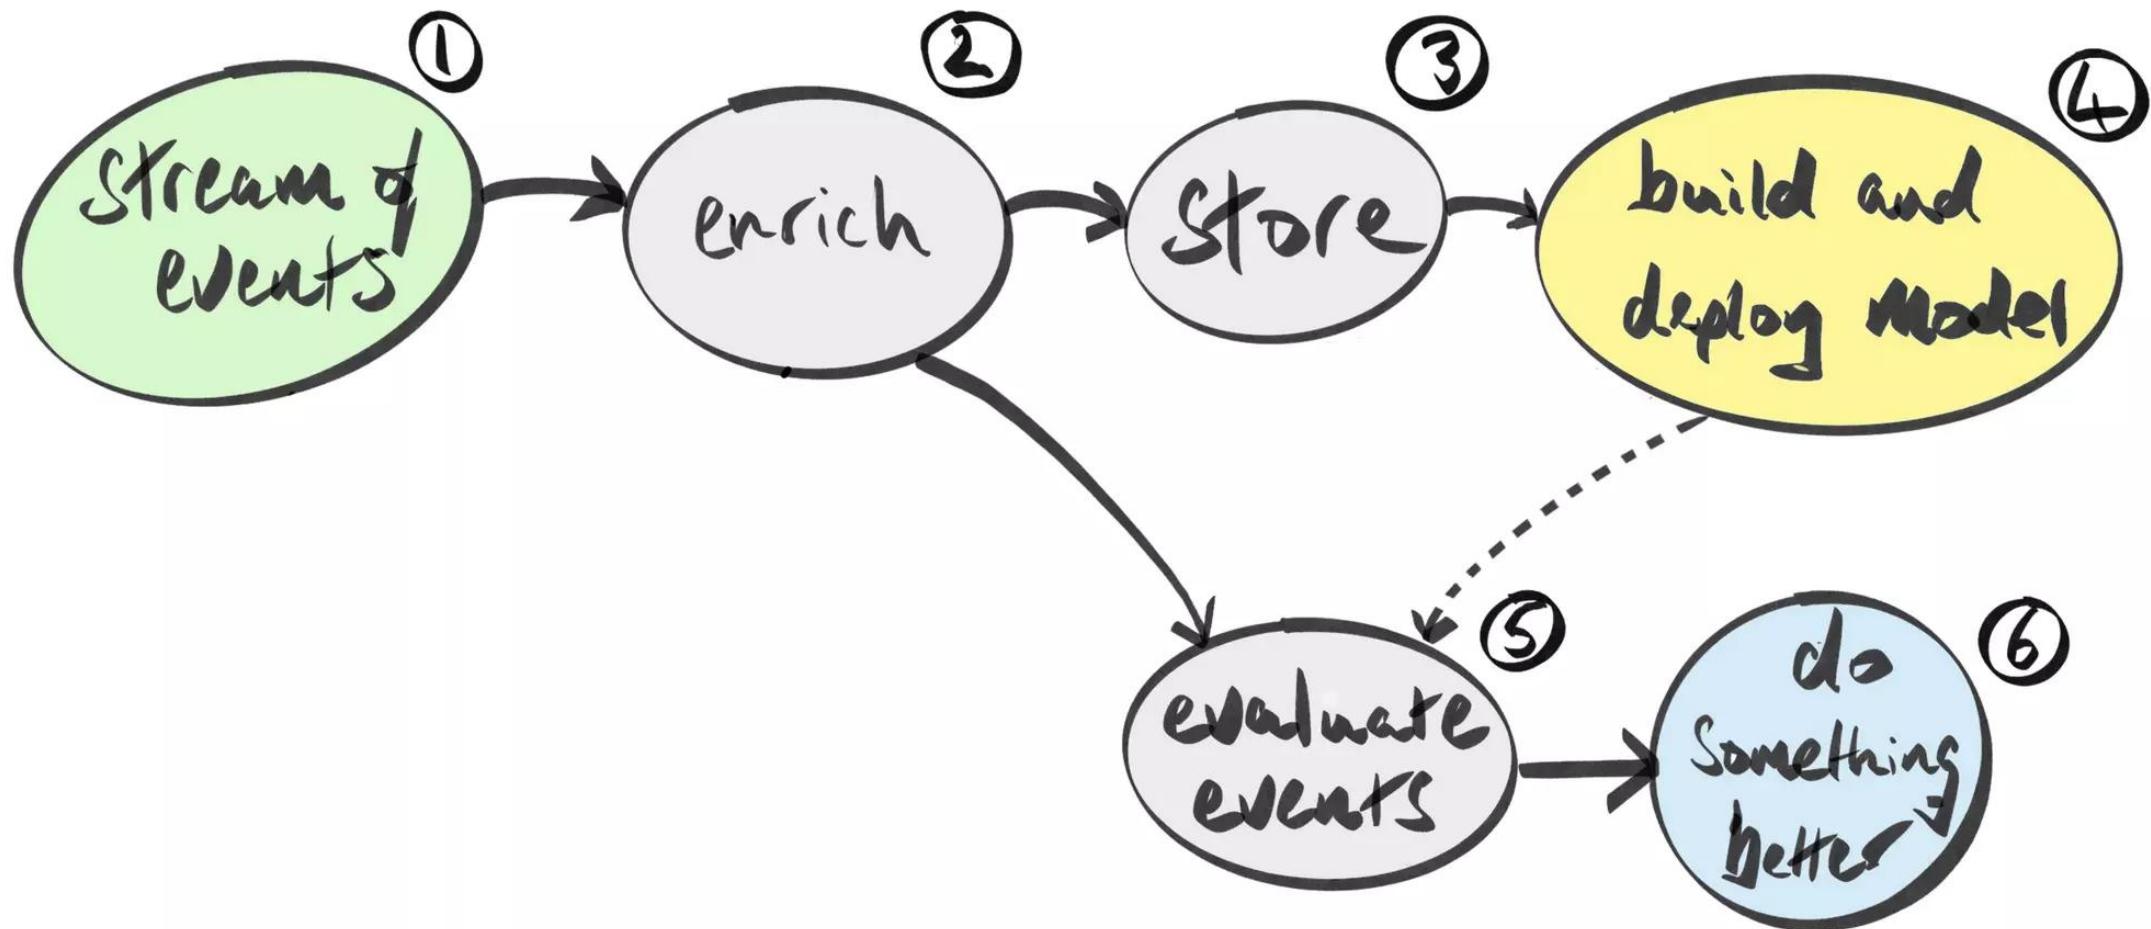




Snowplow ingest



Putting AI into streams



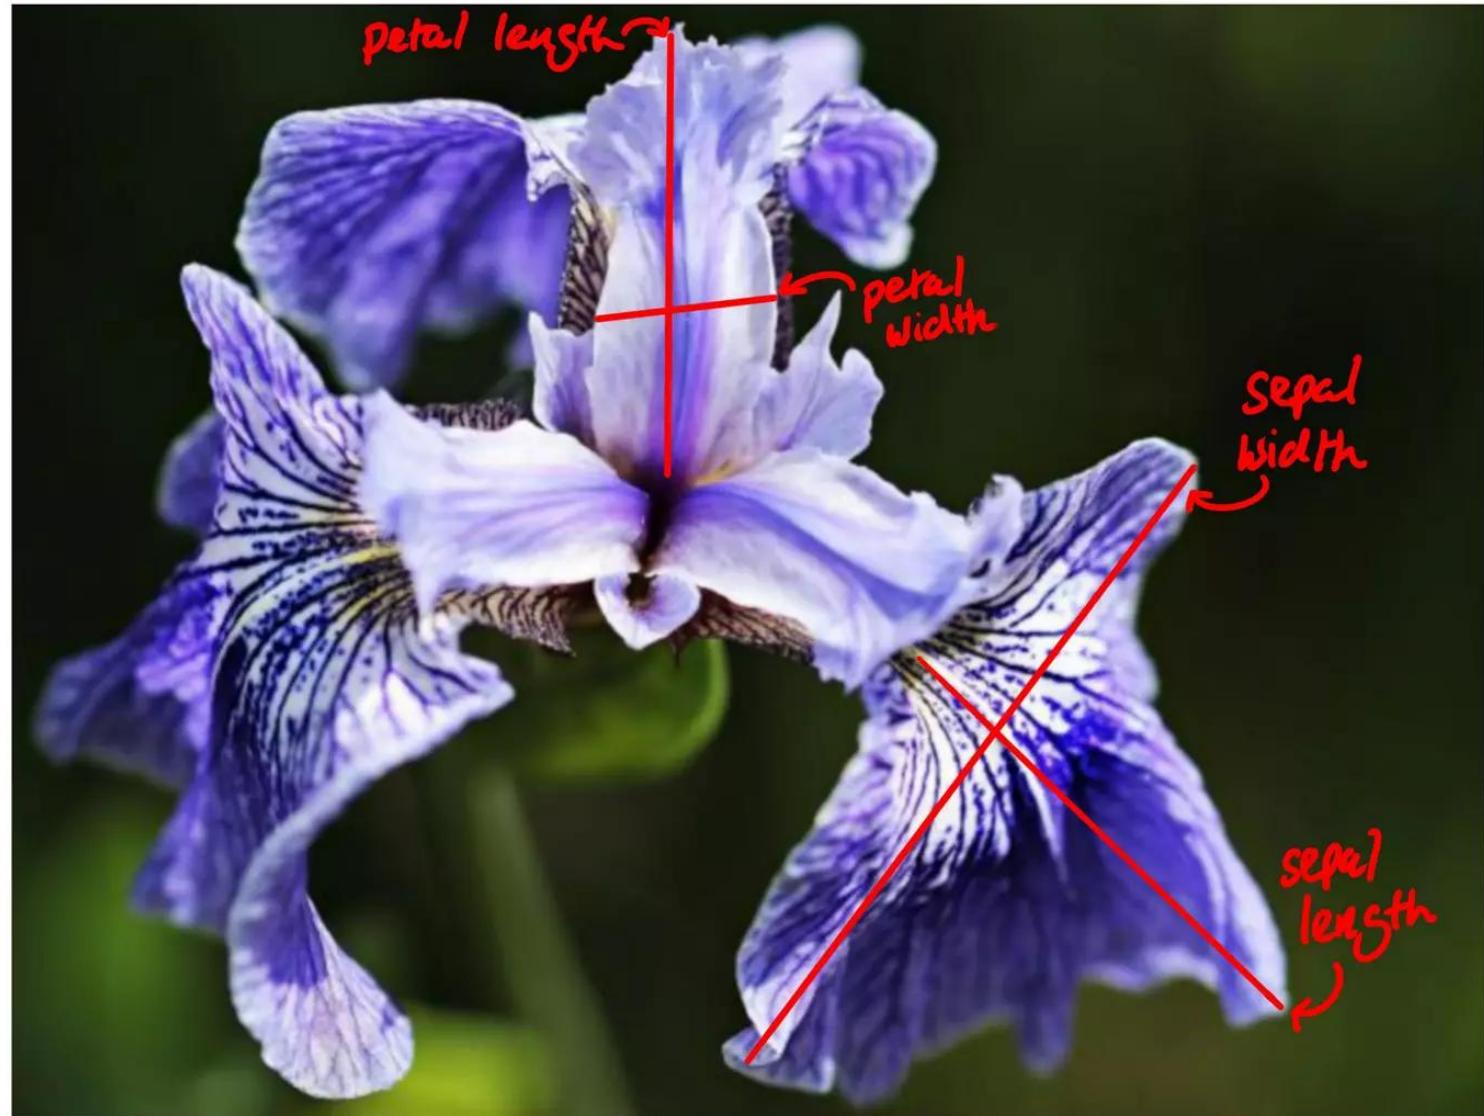
$$y = f(x_1, x_2, x_3 \dots)$$

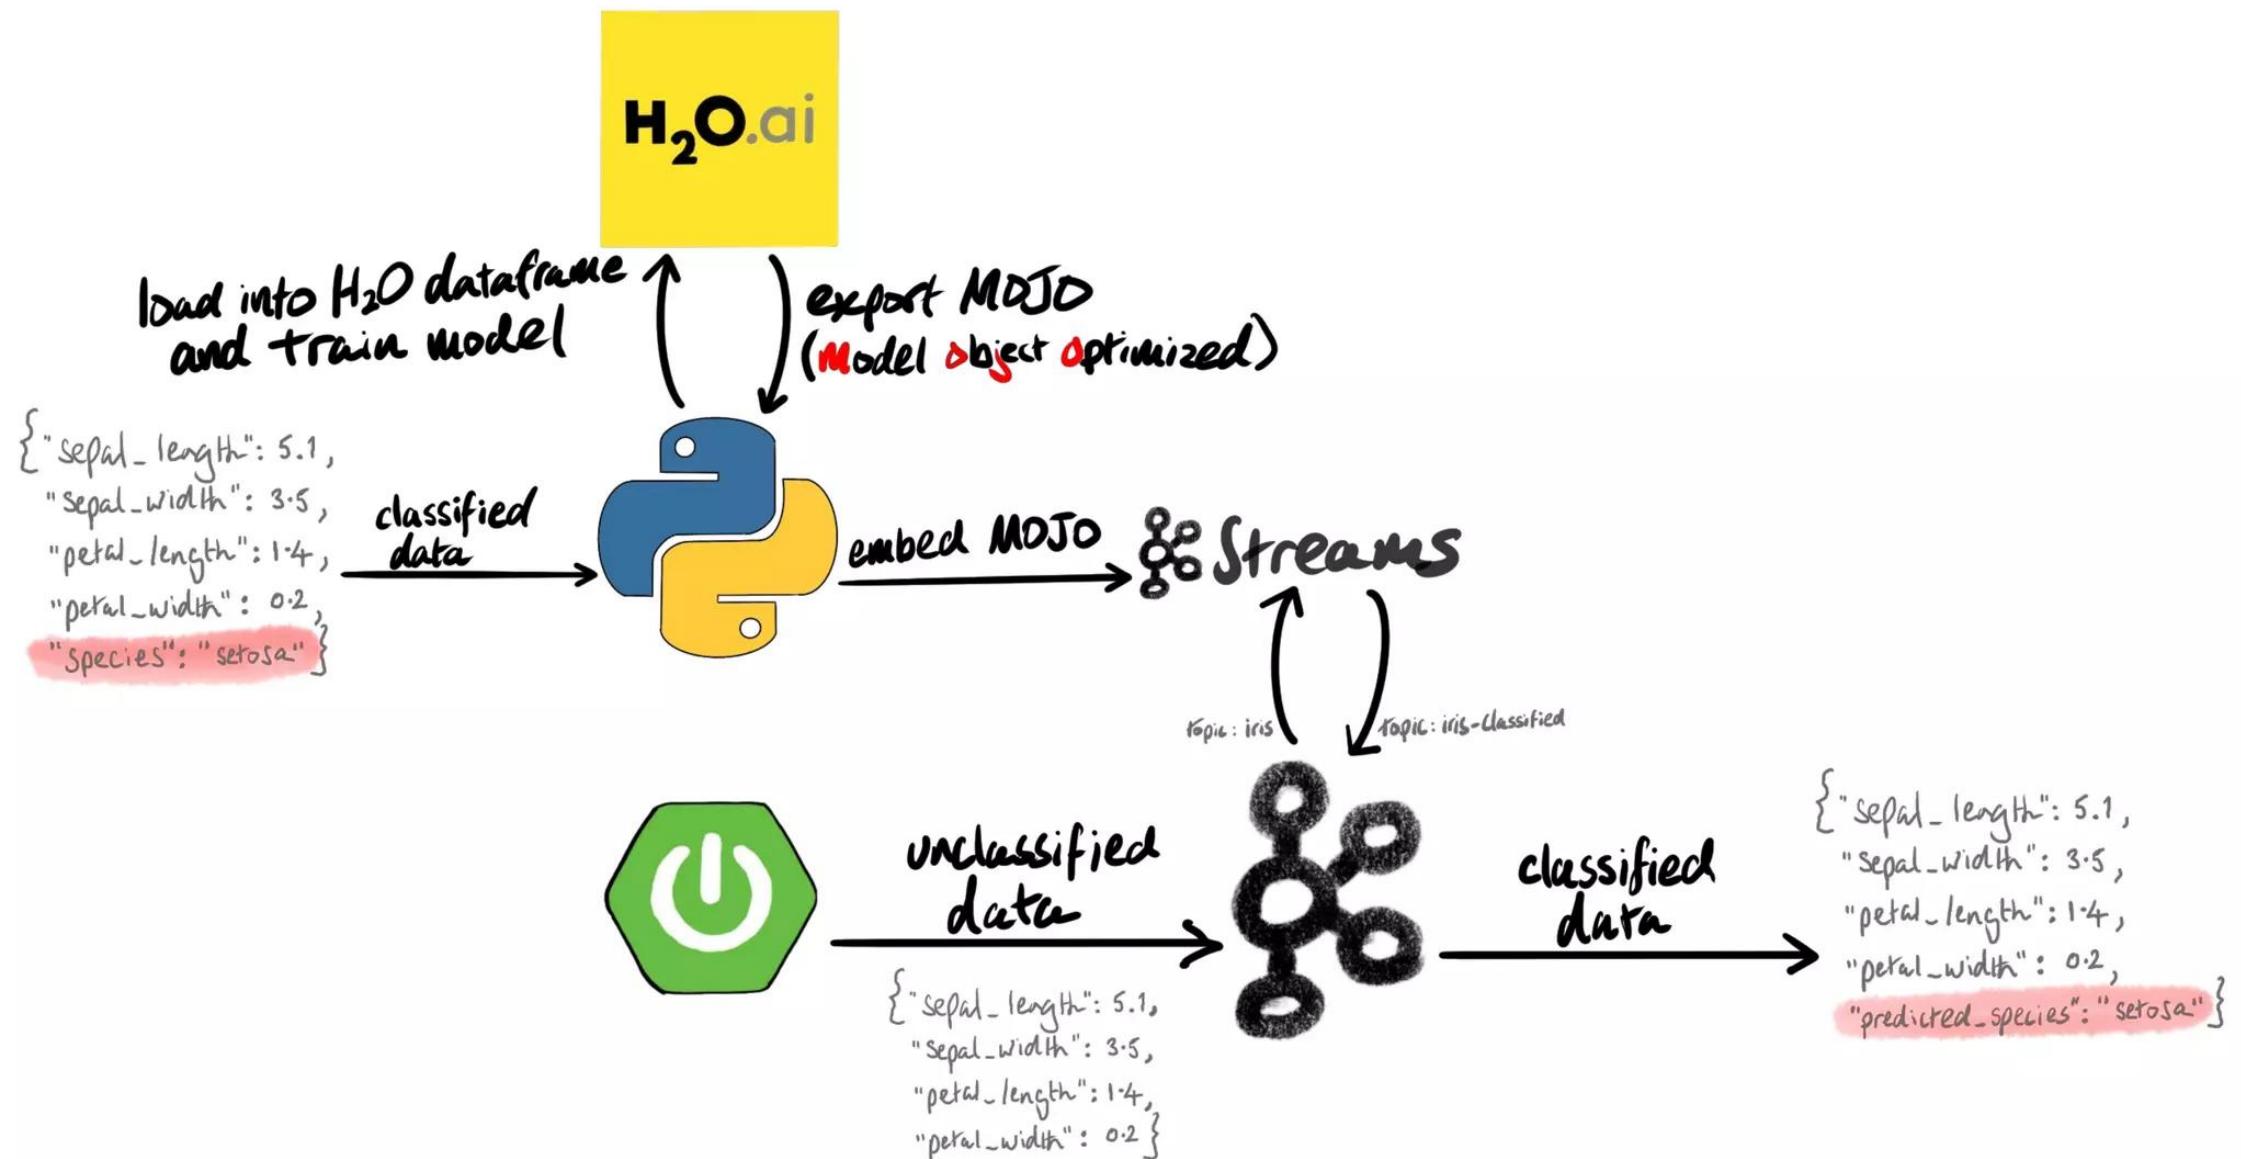
↑
thing we're trying to predict

↑
things that might influence the thing we're trying to predict

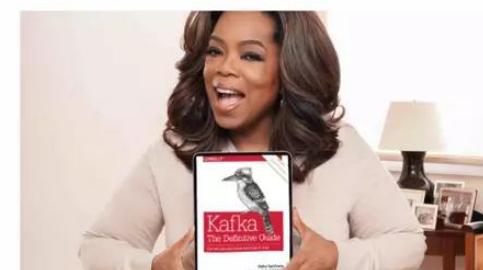
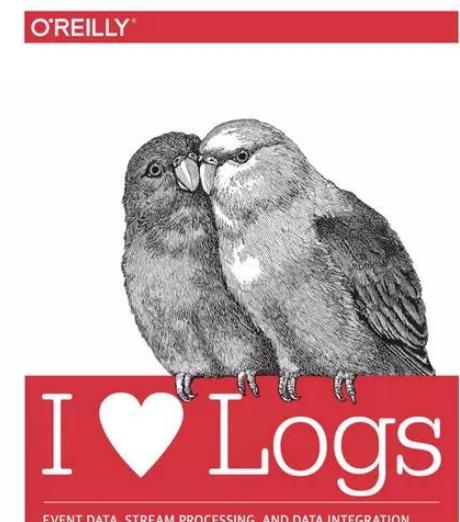
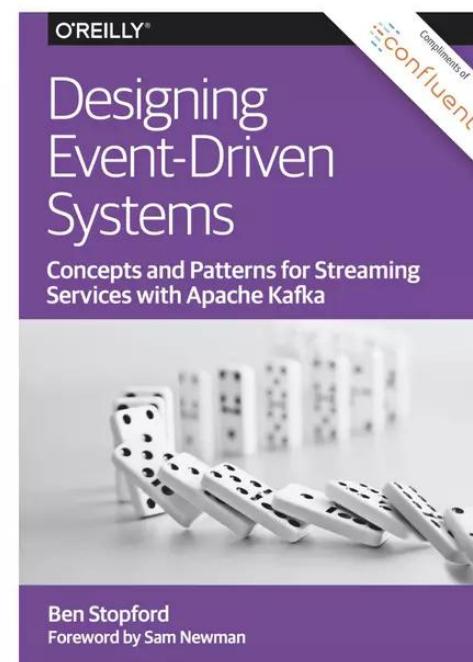
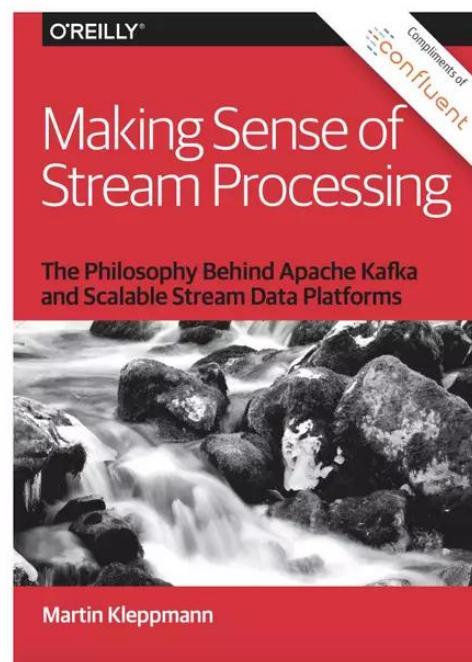
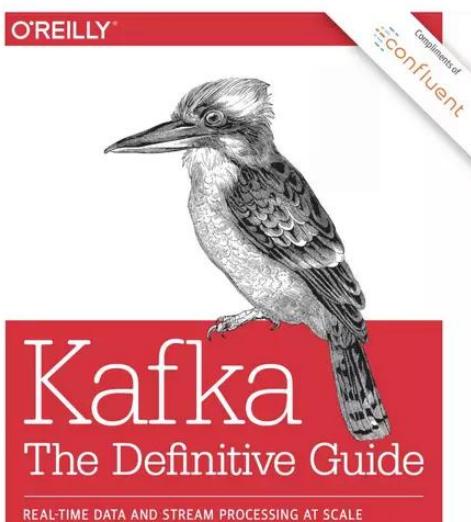
- Step 1) create lots of candidate models.
- Step 2) choose the best one.
- Step 3) stick it in the stream.

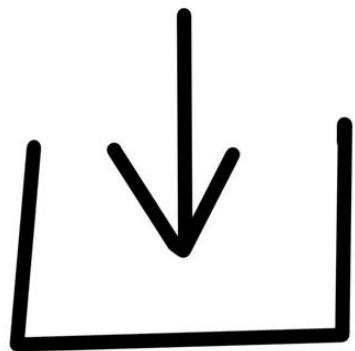
the iris dataset



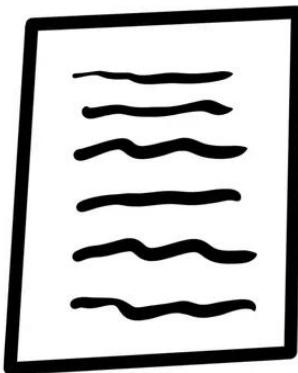


confluent.io/apache-kafka-stream-processing-book-bundle

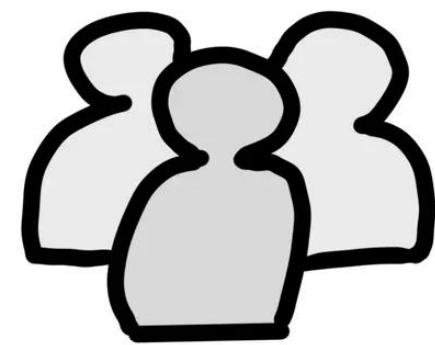




cnfl.io/download



cnfl.io/blog



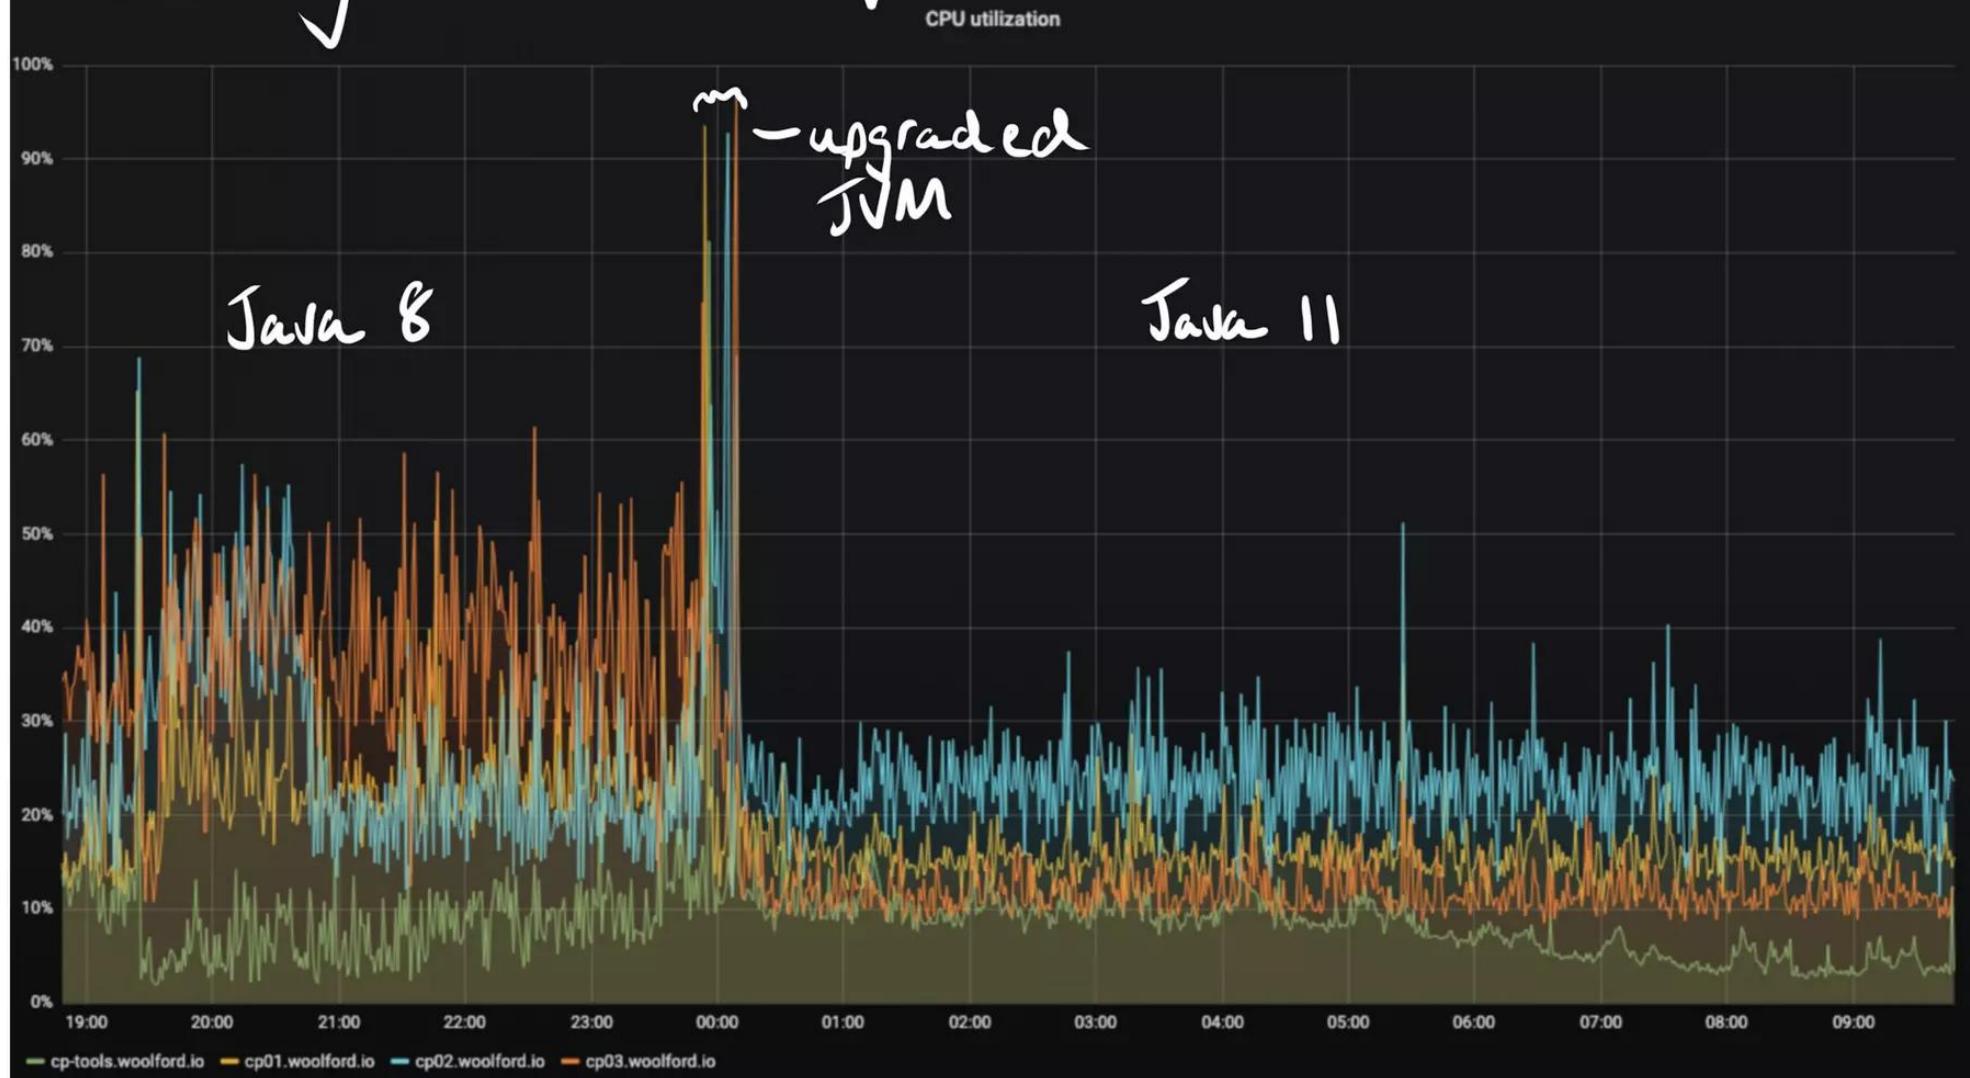
cnfl.io/meetups



LinkedIn: alexwoolford

Twitter: @bissgus_data

Using SSL? Upgrade to Java 11



Avro and Schema Registry

strict ordering in Kafka

exactly once

transactions

monitoring with Prometheus, Datadog, Burrow, Lightbend Lag Exporter

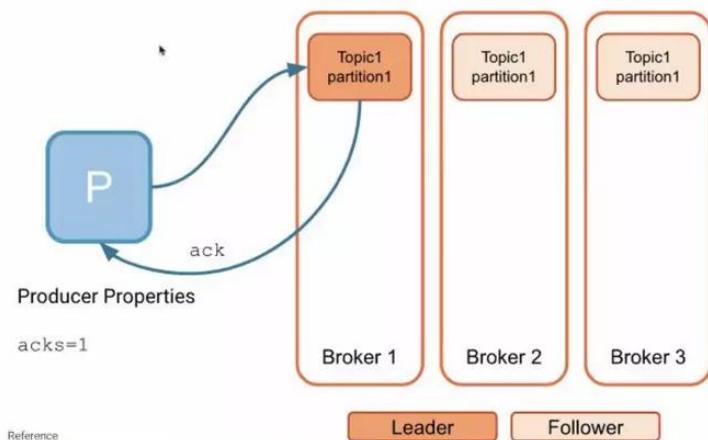
tailing log files with FileBeat

consumer offsets

zero copy, sequential reads, and the page cache

producer request configuration properties

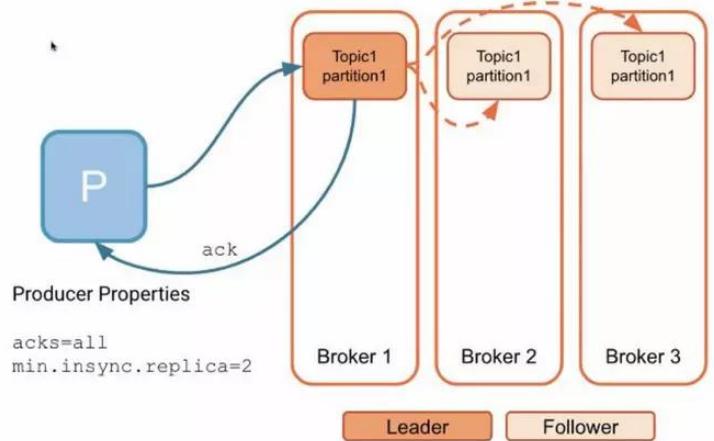
Producer Guarantees



confluent

Reference
<https://www.confluent.io/blog/exactly-once-semantics-are-easy-here-how-apache-kafka-does-it/>

Producer Guarantees



confluent

consumer request configuration properties

bridge to cloud example

OSQuery

Persistent stores, punctuate, grace period, suppress

timeseries analysis

How to run KSQL

CLI
REST
C3
headless?

Terminology

- **Stream** - unbounded, continuously updating data set
- **Stream partition** - ordered, replayable, and fault-tolerant sequence of immutable data records
- **Data record** - key-value pair
- **Topology** - computational logic of the data processing (typically expressed by a DAG)
- **Stream processor** - node in the topology (i.e. processing step). It can be declarative (DSL) or imperative (PAPI)
- **KStream** - abstraction of a record stream
- **KTable** - abstraction of a changelog stream
- **KGlobalTable** - KTable with the data populated from all partitions



Agenda

- Developing Streams Apps
 - Streams and Tables
 - Topology
 - Serdes
 - Joins and windows
 - Error handling
 - Suppress
 - Exposing interactive queries
 - Testing
 - Tricks
- How Streams works
 - Threading model
 - State store
 - Key configuration
 - Processing guarantees
- Running Streams Apps
 - Monitoring
 - Security

DSL vs PAPI

<https://docs.confluent.io/current/streams/developer-guide/dsl-api.html>

<https://docs.confluent.io/current/streams/developer-guide/processor-api.html>



DSL API

- Preferred option for most of the use cases
- 9 out of 10 developers pick the DSL

```
KStream<String, String> textLines = ...;

KStream<String, Long> wordCounts = textLines
    // Split each text line, by whitespace, into words. The text lines are the record
    // values, i.e. you can ignore whatever data is in the record keys and thus invoke
    // `flatMapValues` instead of the more generic `flatMap`.
    .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\W+")))
    // Group the stream by word to ensure the key of the record is the word.
    .groupByKey((key, word) -> word)
    // Count the occurrences of each word (record key).
    //
    // This will change the stream type from `KGroupedStream<String, String>` to
    // `KTable<String, Long>` (word -> count).
    .count()
    // Convert the `KTable<String, Long>` into a `KStream<String, Long>`.
    .toStream();
```

Processor API

- Same logic as previous code
- Full flexibility but much more manual work
- DSL and Processor API can be mixed (i.e. use DSL whenever possible and Processor for exceptions)

```
public class WordCountProcessor implements Processor<String, String> {  
  
    private ProcessorContext context;  
    private KeyValueStore<String, Long> kvStore;  
  
    @Override  
    @SuppressWarnings("unchecked")  
    public void init(ProcessorContext context) {  
        // keep the processor context locally because we need it in punctuate() and commit()  
        this.context = context;  
  
        // retrieve the key-value store named "Counts"  
        kvStore = (KeyValueStore) context.getStateStore("Counts");  
  
        // schedule a punctuate() method every second based on event-time  
        this.context.schedule(Duration.ofSeconds(1), PunctuationType.STREAM_TIME, (timestamp) -> {  
            KeyValueIterator<String, Long> iter = this.kvStore.all();  
            while (iter.hasNext()) {  
                KeyValue<String, Long> entry = iter.next();  
                context.forward(entry.key, entry.value.toString());  
            }  
            iter.close();  
  
            // commit the current processing progress  
            context.commit();  
        });  
    }  
  
    @Override  
    public void process(String dummy, String line) {  
        String[] words = line.toLowerCase(Locale.getDefault()).split(" ");  
  
        for (String word : words) {  
            Integer oldValue = this.kvStore.get(word);  
  
            if (oldValue == null) {  
                this.kvStore.put(word, 1);  
            } else {  
                this.kvStore.put(word, oldValue + 1);  
            }  
        }  
    }  
  
    @Override  
    public void close() {  
        // nothing to do  
    }  
}
```



Should I use DSL or the Processor API ?

“**DSL** looks awesome but I feel limited, I need more control.”

“**PAPI** looks simpler, powerful, and I have more control on things”

“With great power comes great responsibility !”

Don't need to go one way or the other we can mix both the **DSL** and **PAPI**.

Before using **PAPI**, start thinking if there is no other way to implement it.

Most probably you are not thinking in a Event Streaming way yet !



Naming the topology

No naming

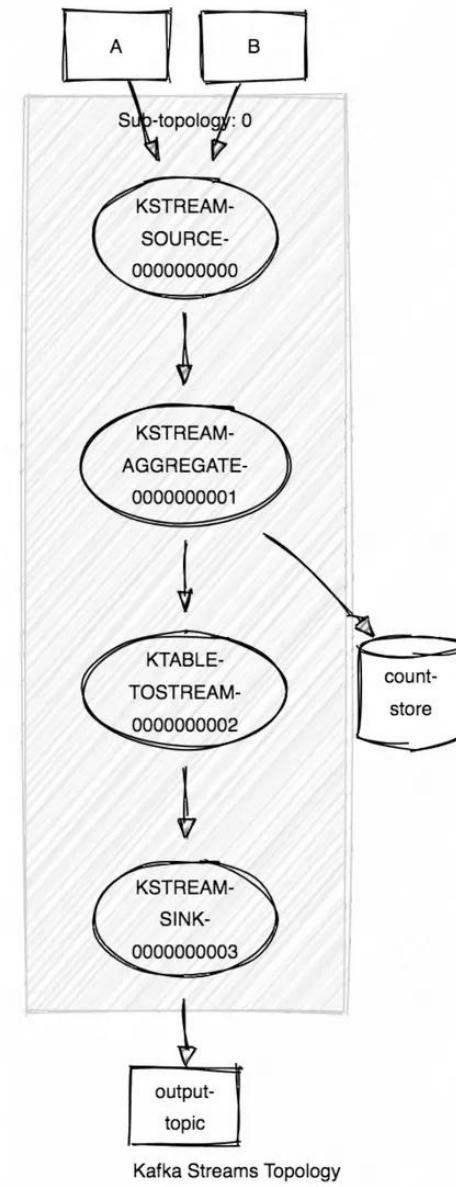
```
KStream<String, String> stream =  
builder.stream("input");  
  
stream.filter((k,v) -> !v.equals("invalid_txn"))  
    .mapValues((v) -> v.substring(0,5))  
    .to("output")
```

With naming

```
KStream<String, String> stream =  
builder.stream("input",  
Consumed.as("Customer_transactions_input_topic"));  
  
stream.filter((k,v) -> !v.equals("invalid_txn"),  
Named.as("filter_out_invalid_txns"))  
    .mapValues((v) -> v.substring(0,5),  
Named.as("Map_values_to_first_6_characters"))  
    .to("output",  
Produced.as("Mapped_transactions_output_topic"));
```

Visualize the topology!

<https://zz85.github.io/kafka-streams-viz/>





Map vs MapValues

Map vs Map Values



```
final StreamsBuilder builder = new StreamsBuilder();
final KTable<Long, Long> count = builder
    .stream(inputTopic, Consumed.with(Serdes.Long(), Serdes.String()))
    .map((key, val) -> new KeyValue<>(key, "new value for " + val))
    .groupByKey()
    .count();
count
    .toStream()
    .to(outputTopic, Produced.with(Serdes.Long(), Serdes.Long()));

Topology topology = builder.build();
```

Here, even though the key is unchanged, which will trigger a repartition because as Kafka Streams will flag the operation as if the key was modified.



Map vs Map Values

```
final StreamsBuilder builder = new StreamsBuilder();
final KTable<Long, Long> count = builder
    .stream(inputTopic, Consumed.with(Serdes.Long(), Serdes.String()))
    .mapValues(val -> "new value for " + val)
    .groupByKey()
    .count();
count
    .toStream()
    .to(outputTopic, Produced.with(Serdes.Long(), Serdes.Long()));

Topology topology = builder.build();
```

Now that we replaced the topology step with a mapValues, this won't trigger a repartition, leading to better performance.

Requirements



- Join on message keys
 - In Confluent 5.4 there is the support non-key joining on KTable ([KIP-213](#))
- Co-partitioned
- Streams to Streams are time windowed

In Confluent 5.5 we introduce CoGroup ([KIP-150](#)) !

KIP-150: Add Cogroup to the DSL



- In the past, aggregating multiple streams into one could be complicated and error-prone.
- It generally requires you to group + aggregate all of the streams into tables, then make multiple outer join calls.
- Using the new co-group operator will:
 - Clean up the syntax of your programs
 - Reduce the number of state store invocations
 - Increase performance

```
KGroupedStream<K, V1> grouped1 =  
    builder.stream("topic1").groupByKey();
```

```
KGroupedStream<K, V2> grouped2 =  
    builder.stream("topic2").groupByKey();
```

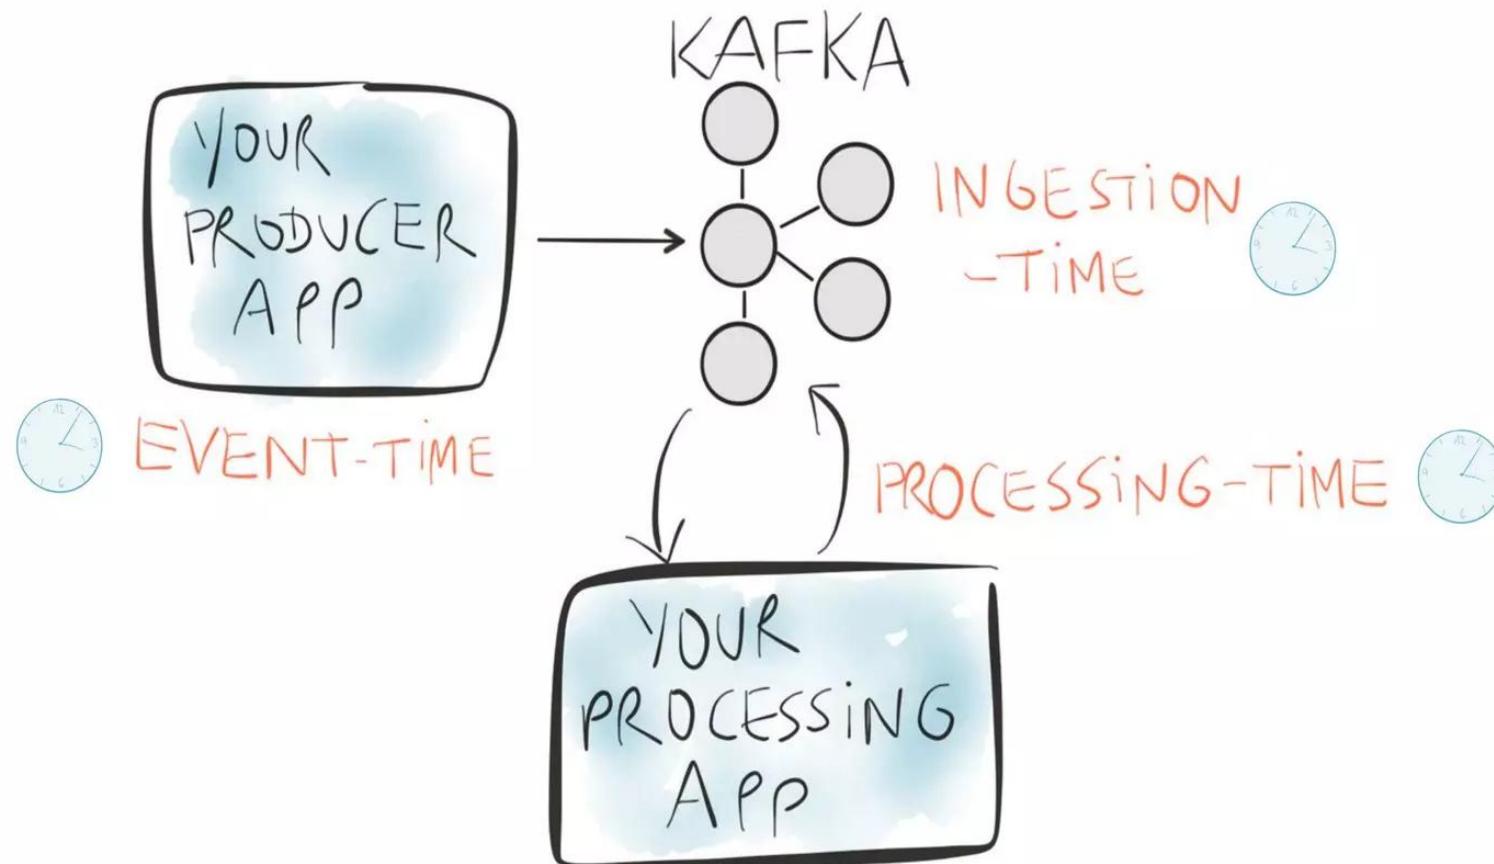
```
KGroupedStream<K, V3> grouped3 =  
    builder.stream("topic3").groupByKey();
```

```
KTable<K, CG> cogrouped =  
    grouped1  
        .cogroup(aggregator1)  
        .cogroup(grouped2, aggregator2)  
        .cogroup(grouped3, aggregator3)  
        .aggregate(initializer1, materialized1);
```

Time time time time !

<https://docs.confluent.io/current/streams/concepts.html#time>

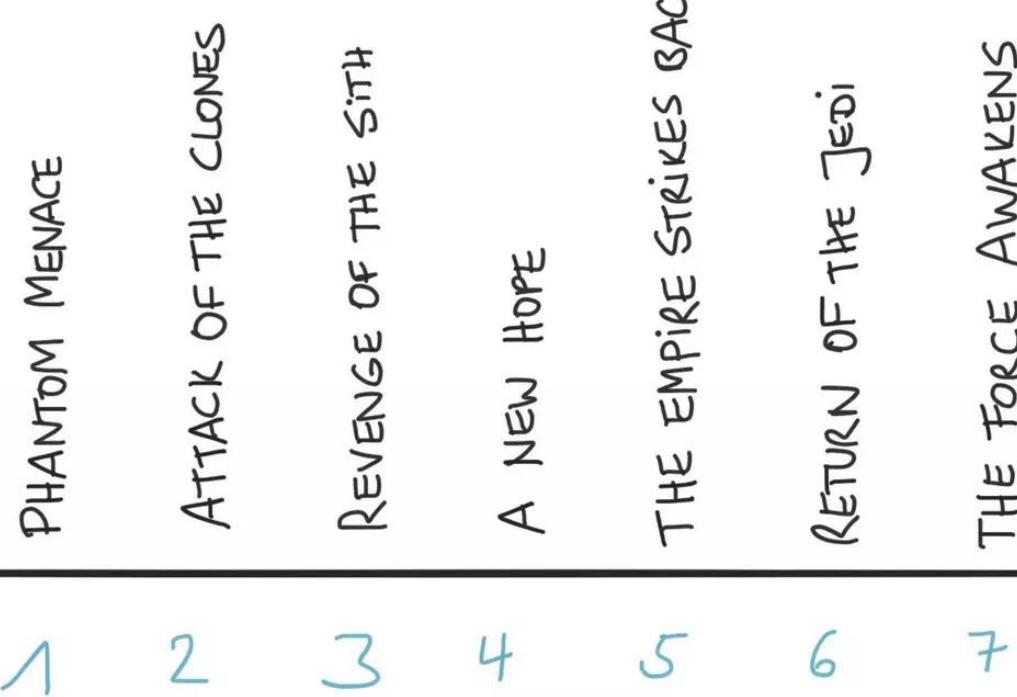
Time



Time



STAR WARS



PROCESSING-TIME 1999 2002 2005 1977 1980 1983 2015

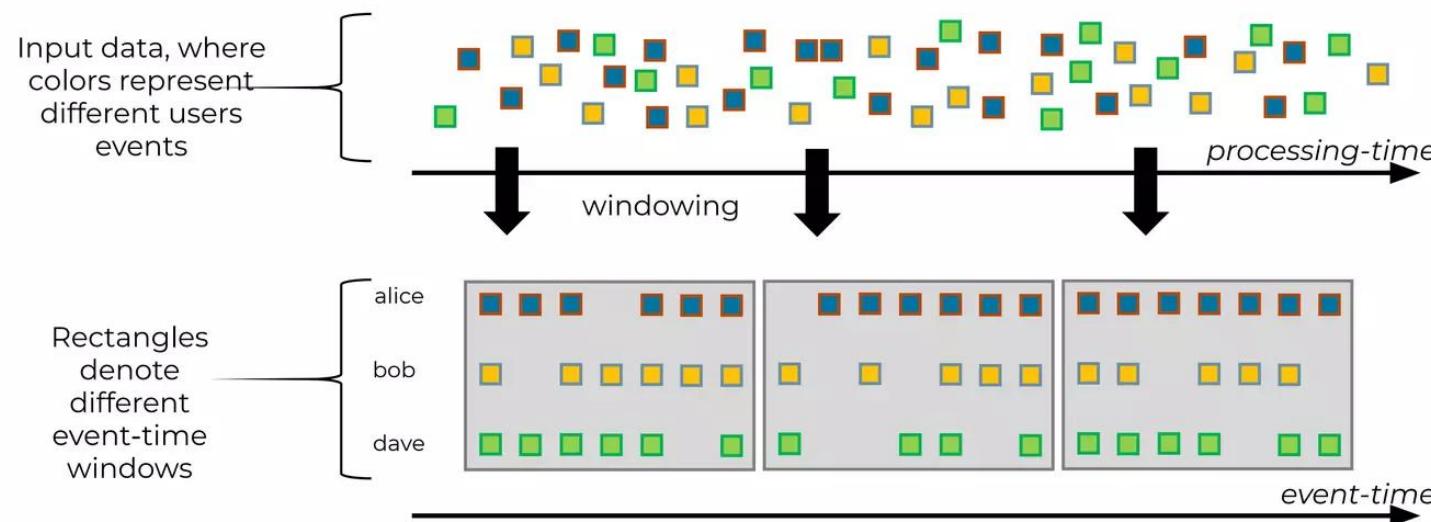
Windowing

<https://docs.confluent.io/currentstreams/concepts.html#windowing>



Windowing

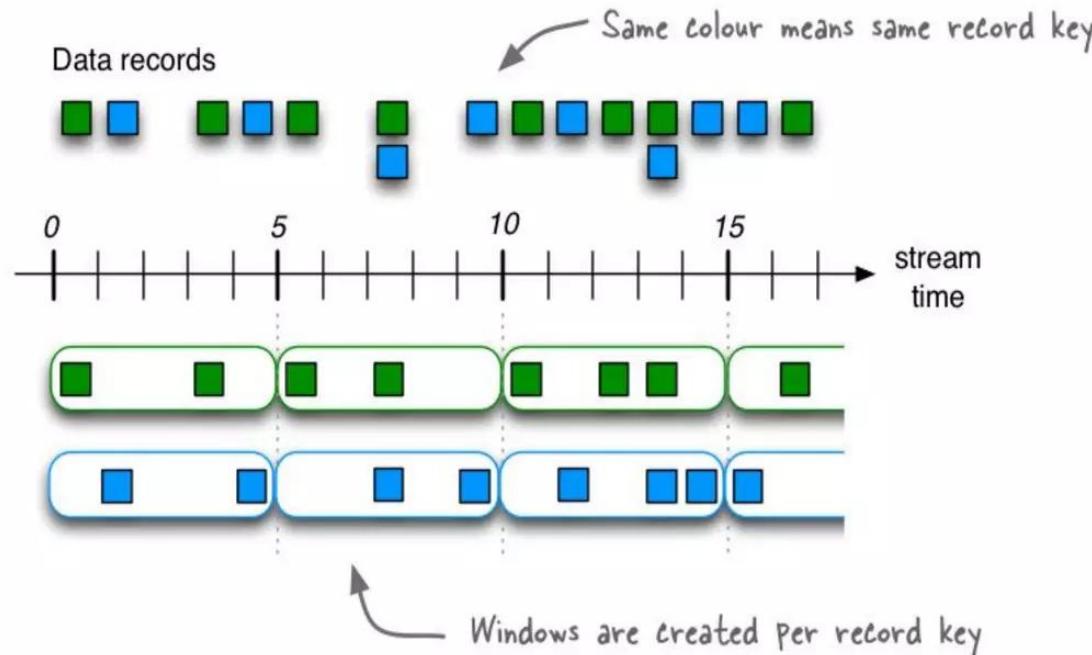
- **To group related events in a stream**
- Use case examples:
 - Time-based analysis of ad impressions ("number of ads clicked in the past hour")
 - Monitoring statistics of telemetry data ("1min/5min/15min averages")
 - Analyzing user browsing sessions on a news site





Tumbling window

A 5-min Tumbling Window

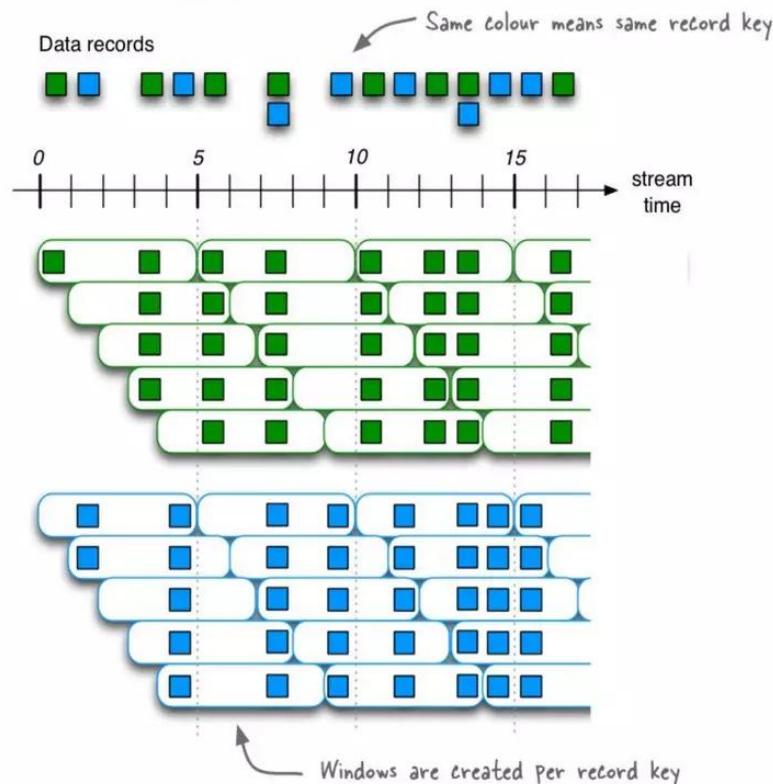


```
// Key (String) is user ID, value (Avro record) is the page view event for that user.  
// Such a data stream is often called a "clickstream".  
KStream<String, GenericRecord> pageViews = ...;  
  
// Count page views per window, per user, with tumbling windows of size 5 minutes  
KTable<Windowed<String>, Long> windowedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(TimeWindows.of(Duration.ofMinutes(5)))  
    .count();
```



Hopping window

A 5-min Hopping Window with a 1-min "hop"

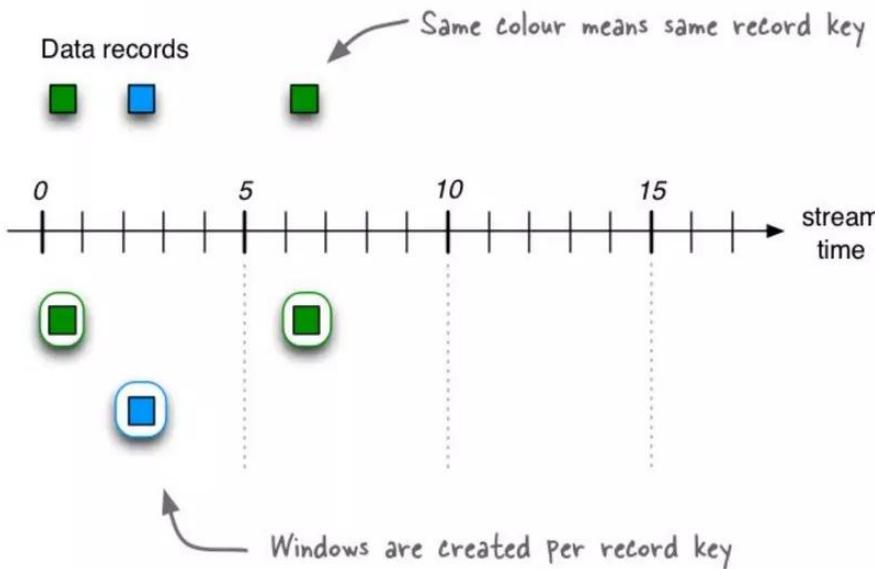


```
// Key (String) is user ID, value (Avro record) is the page view event for that user.  
// Such a data stream is often called a "clickstream".  
KStream<String, GenericRecord> pageViews = ...;  
  
// Count page views per window, per user, with hopping windows of size 5 minutes that advance every 1 minute  
KTable<Windowed<String>, Long> windowedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(TimeWindows.of(Duration.ofMinutes(5)).advanceBy(Duration.ofMinutes(1))))  
    .count()
```

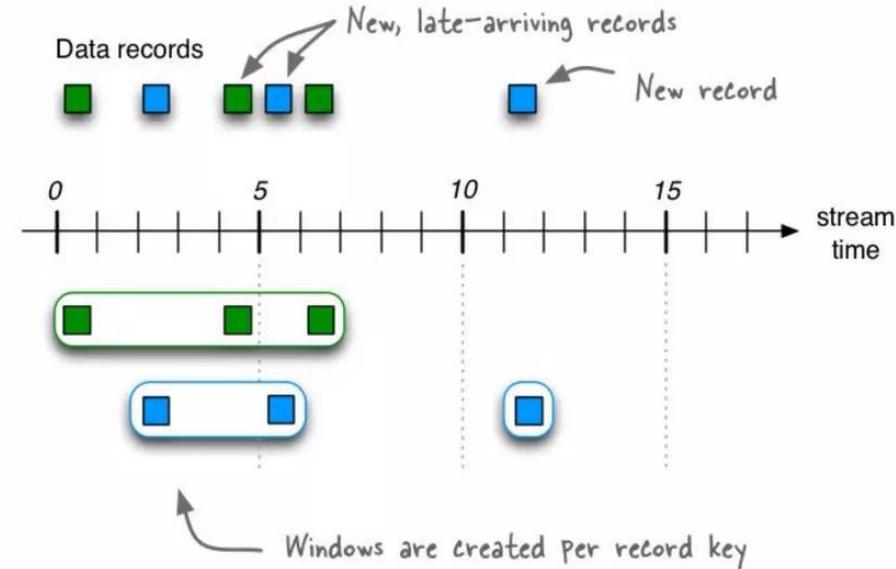


Session window

A Session Window with a 5-min inactivity gap



A Session Window with a 5-min inactivity gap



```
// Key (String) is user ID, value (Avro record) is the page view event for that user.  
// Such a data stream is often called a "clickstream".  
KStream<String, GenericRecord> pageViews = ...;  
  
// Count page views per session, per user, with session windows that have an inactivity gap of 5 minutes  
KTable<Windowed<String>, Long> sessionizedPageViewCounts = pageViews  
    .groupByKey(Grouped.with(Serdes.String(), genericAvroSerde))  
    .windowedBy(SessionWindows.with(Duration.ofMinutes(5)))  
    .count();
```



Kafka Streams: Write standard Java apps and microservices to process your data in real-time

- No separate processing cluster required
- Develop on Mac, Linux, Windows
- Deploy to containers, VMs, bare metal, cloud
- Powered by Kafka: elastic, scalable, distributed, battle-tested
- Perfect for small, medium, large use cases
- Fully integrated with Kafka security
- Exactly-once processing semantics
- Part of Apache Kafka

```
KStream<User, PageViewEvent> pageViews = builder.stream("pageviews-topic");
KTable<Windowed<User>, Long> viewsPerUserSession = pageViews
    .groupByKey()
    .count(SessionWindows.with(TimeUnit.MINUTES.toMillis(5)), "session-views");
```

<https://docs.confluent.io/currentstreams/>



Kafka Streams

Transform Data with Real-Time Applications



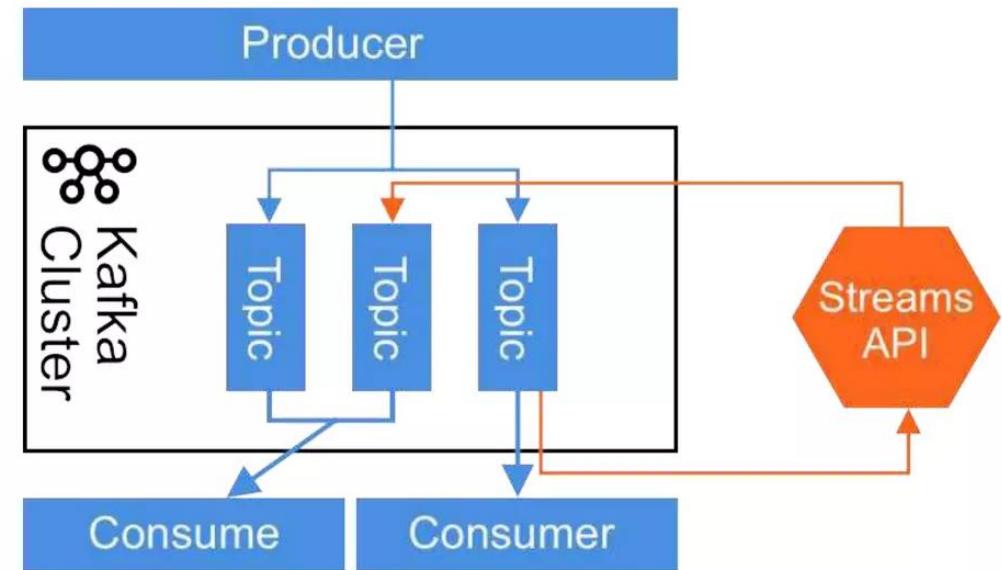
Overview

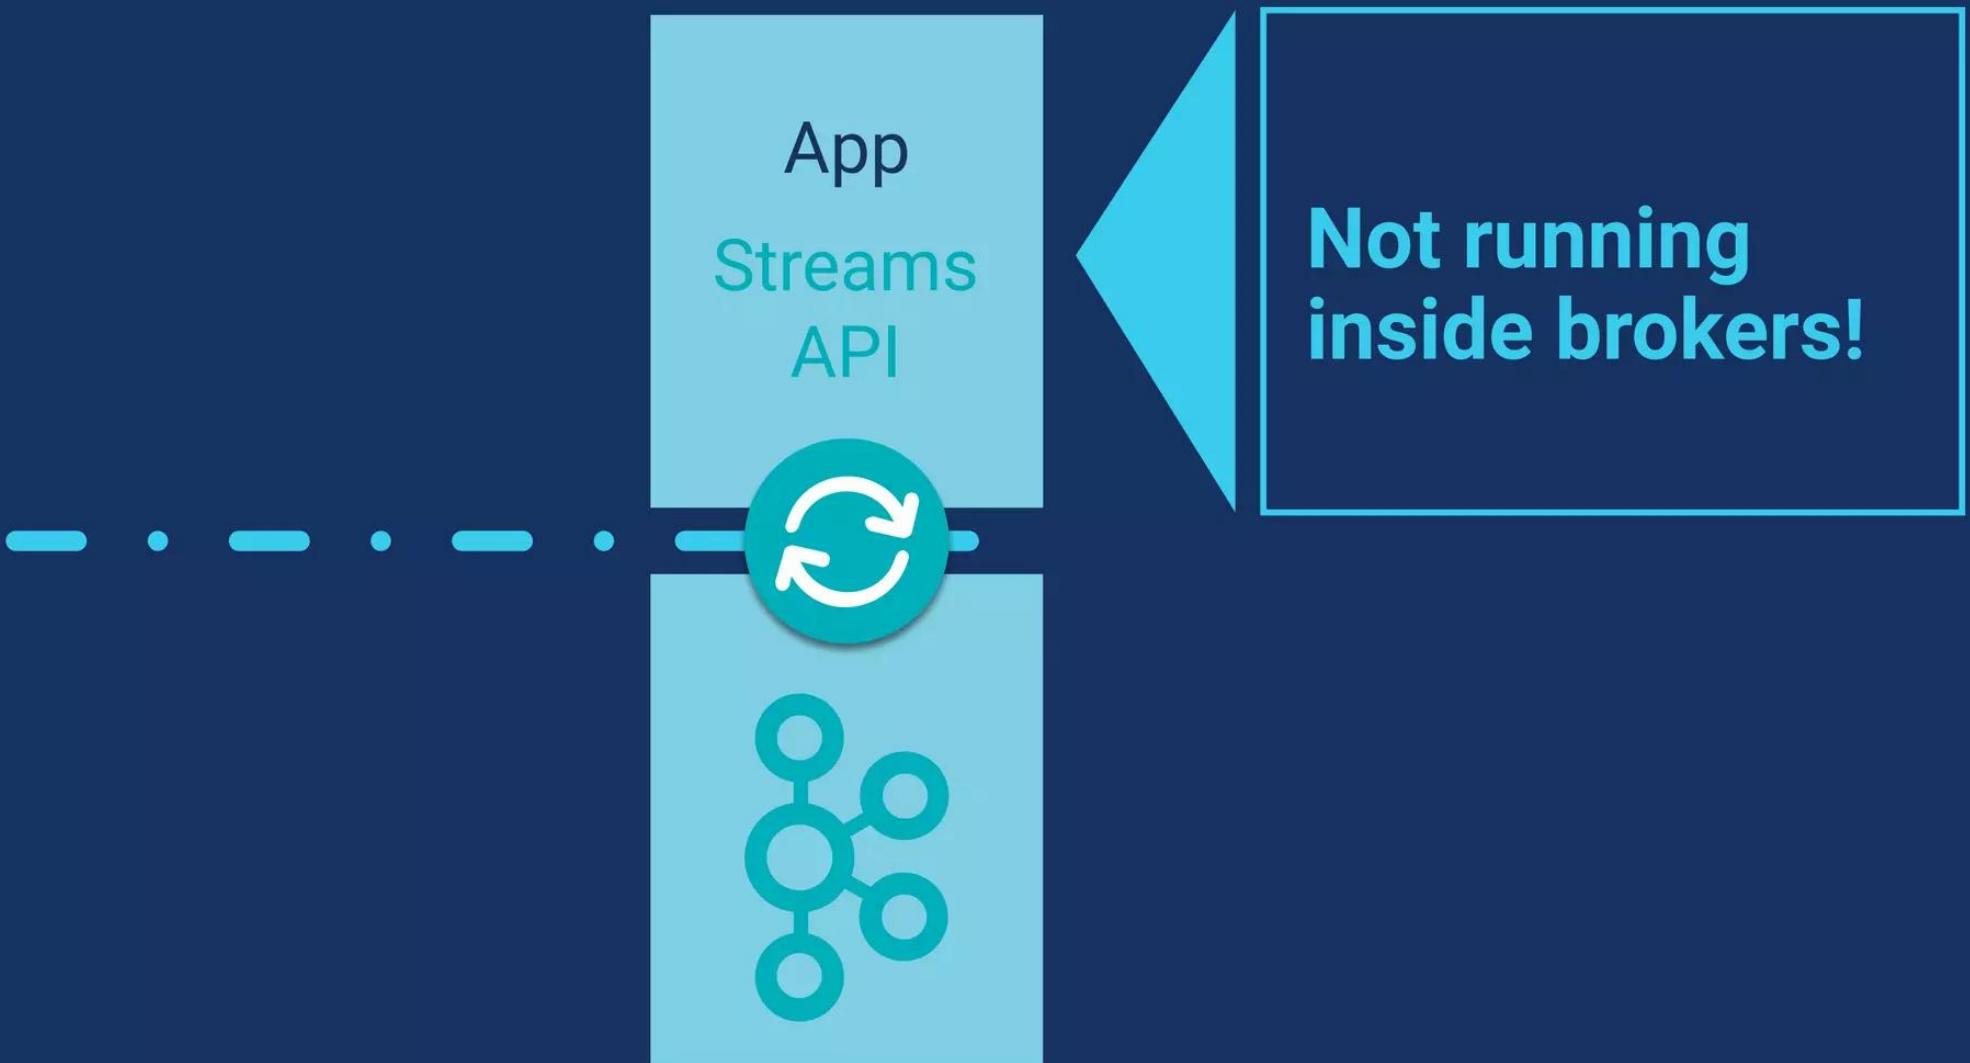
- Write standard Java applications
- No separate processing cluster required
- Exactly-once processing semantics
- Elastic, highly scalable, fault-tolerant
- Fully integrated with Kafka security



Example Use Cases

- Microservices
- Continuous queries
- Continuous transformations





Things Kafka Streams Does



Runs
everywhere



Clustering
done for you



Exactly once
processing



Event time
processing



Integrated
database

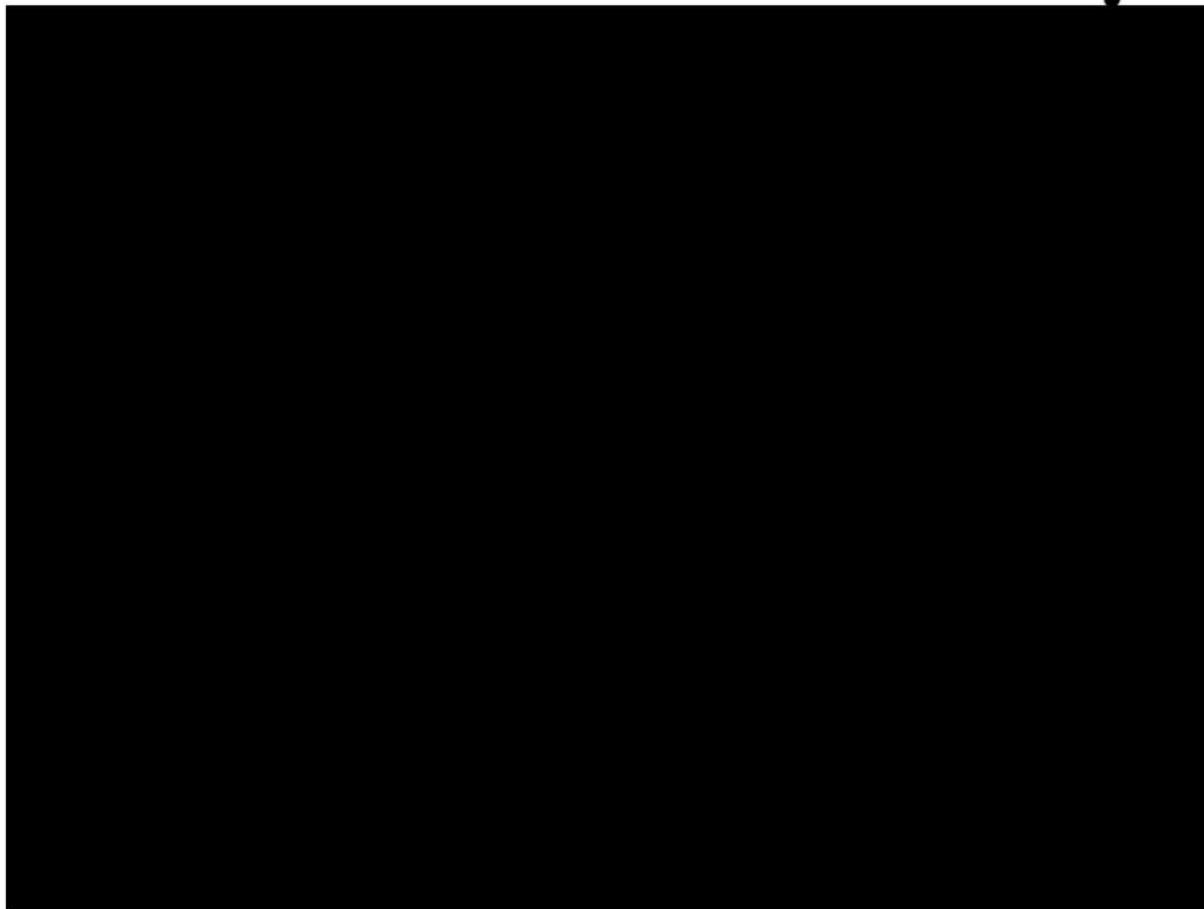


Joins, windowing,
aggregation



S/M/L/XL/XXL/XXXL
sizes

each consumer has its own position



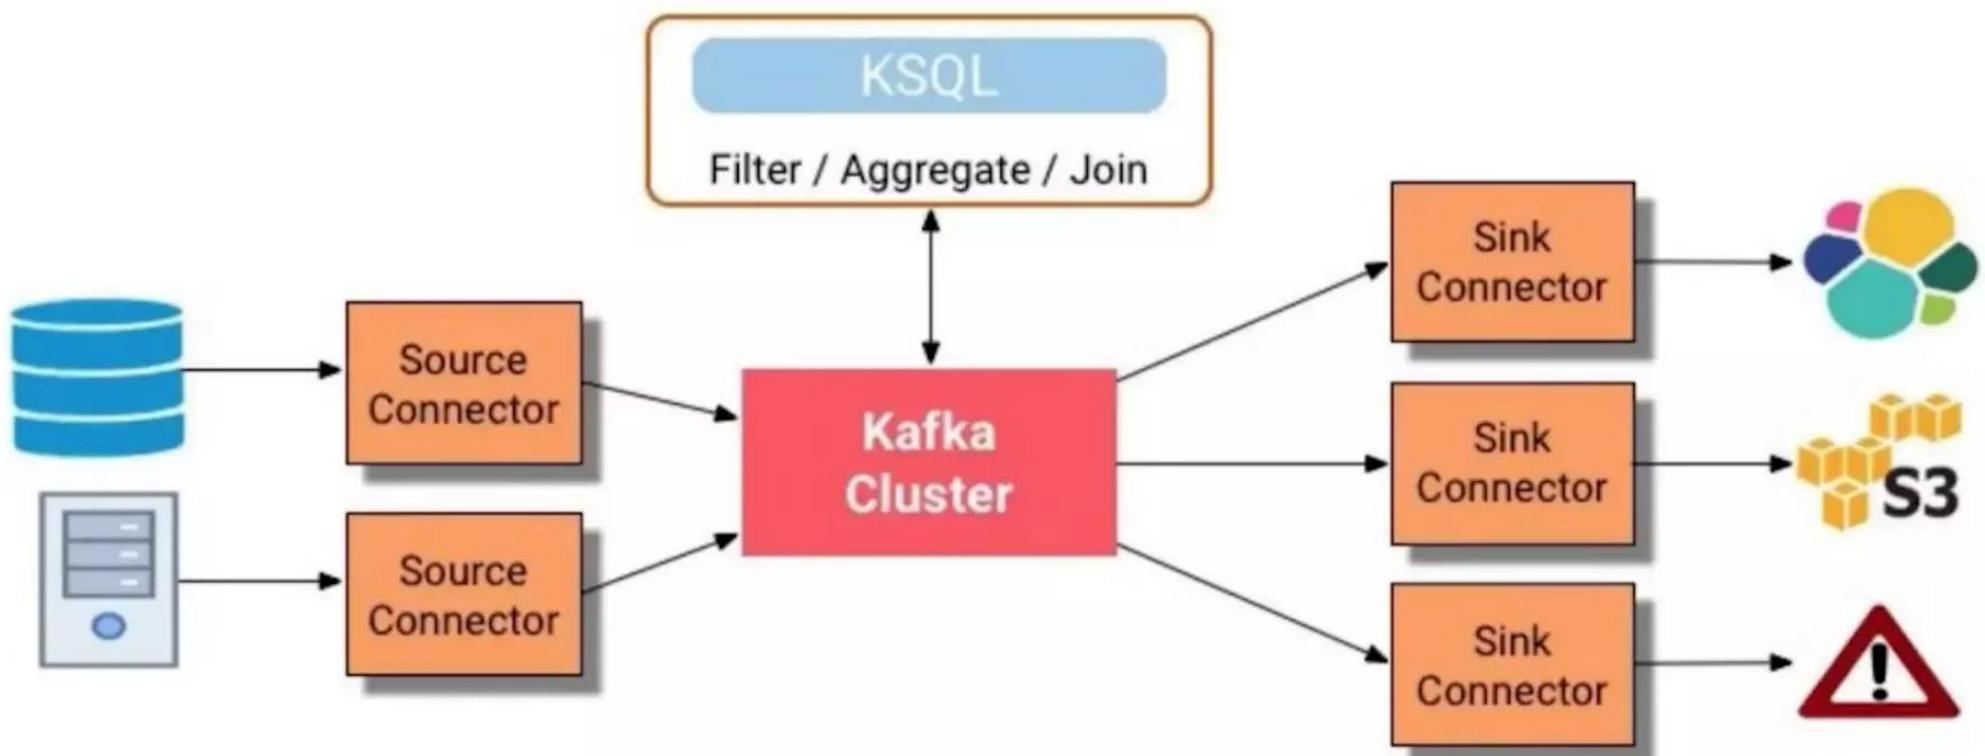
offsets are stored in Kafka topic -- consumer_offsets



Kafka Streams

- Java Library
- no out-of-the-box runnable application
- Build Jars to deploy
- No interactivity

KSQL Server



Types of KSQL statement - p1

Create Data Sources:

- CREATE TABLE ...
- CREATE STREAM ...
- SHOW TABLES; SHOW STREAMS;

Transforming persistent queries

Types of KSQL statement - p2

Metadata management:

- SHOW STREAMS / TABLES / QUERIES
- DESCRIBE STREAM / TABLE
- EXPLAIN QUERY
- DROP STREAM / TABLE
- TERMINATE
- PRINT

User Defined Functions

- Called inline in KSQL
- Types are: UDF, UDAF, UDTF
- Implemented in Java & deployed as JARs to ksql-server

