


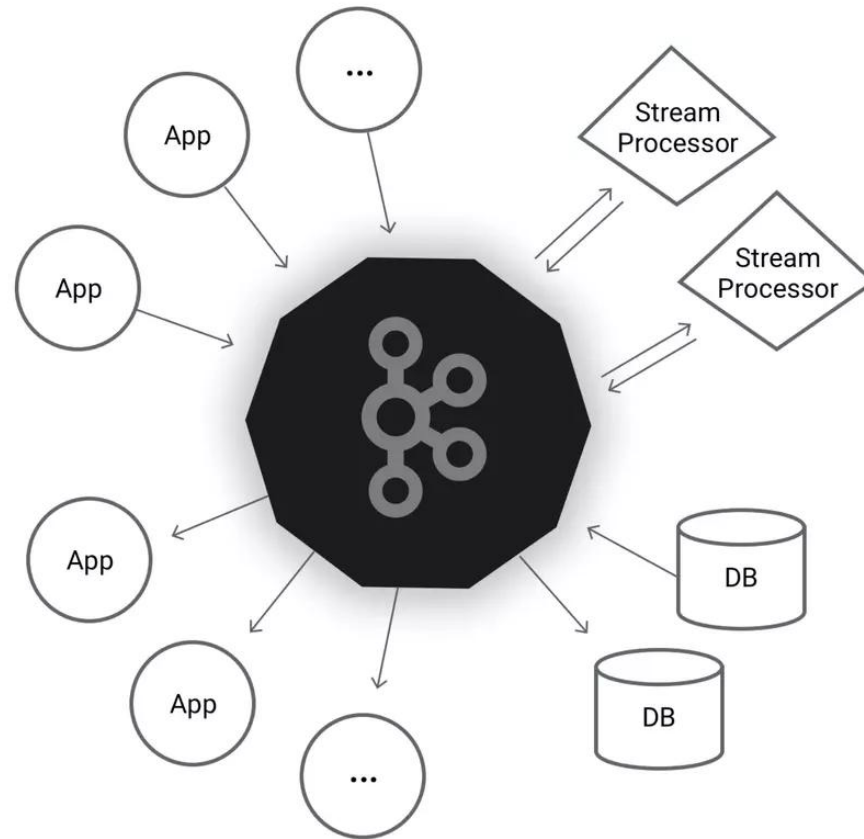


# Kafka Streams: What it is, and how to use it?

Matthias J. Sax | Software Engineer |  @MatthiasJSax  
Apache Kafka PMC member



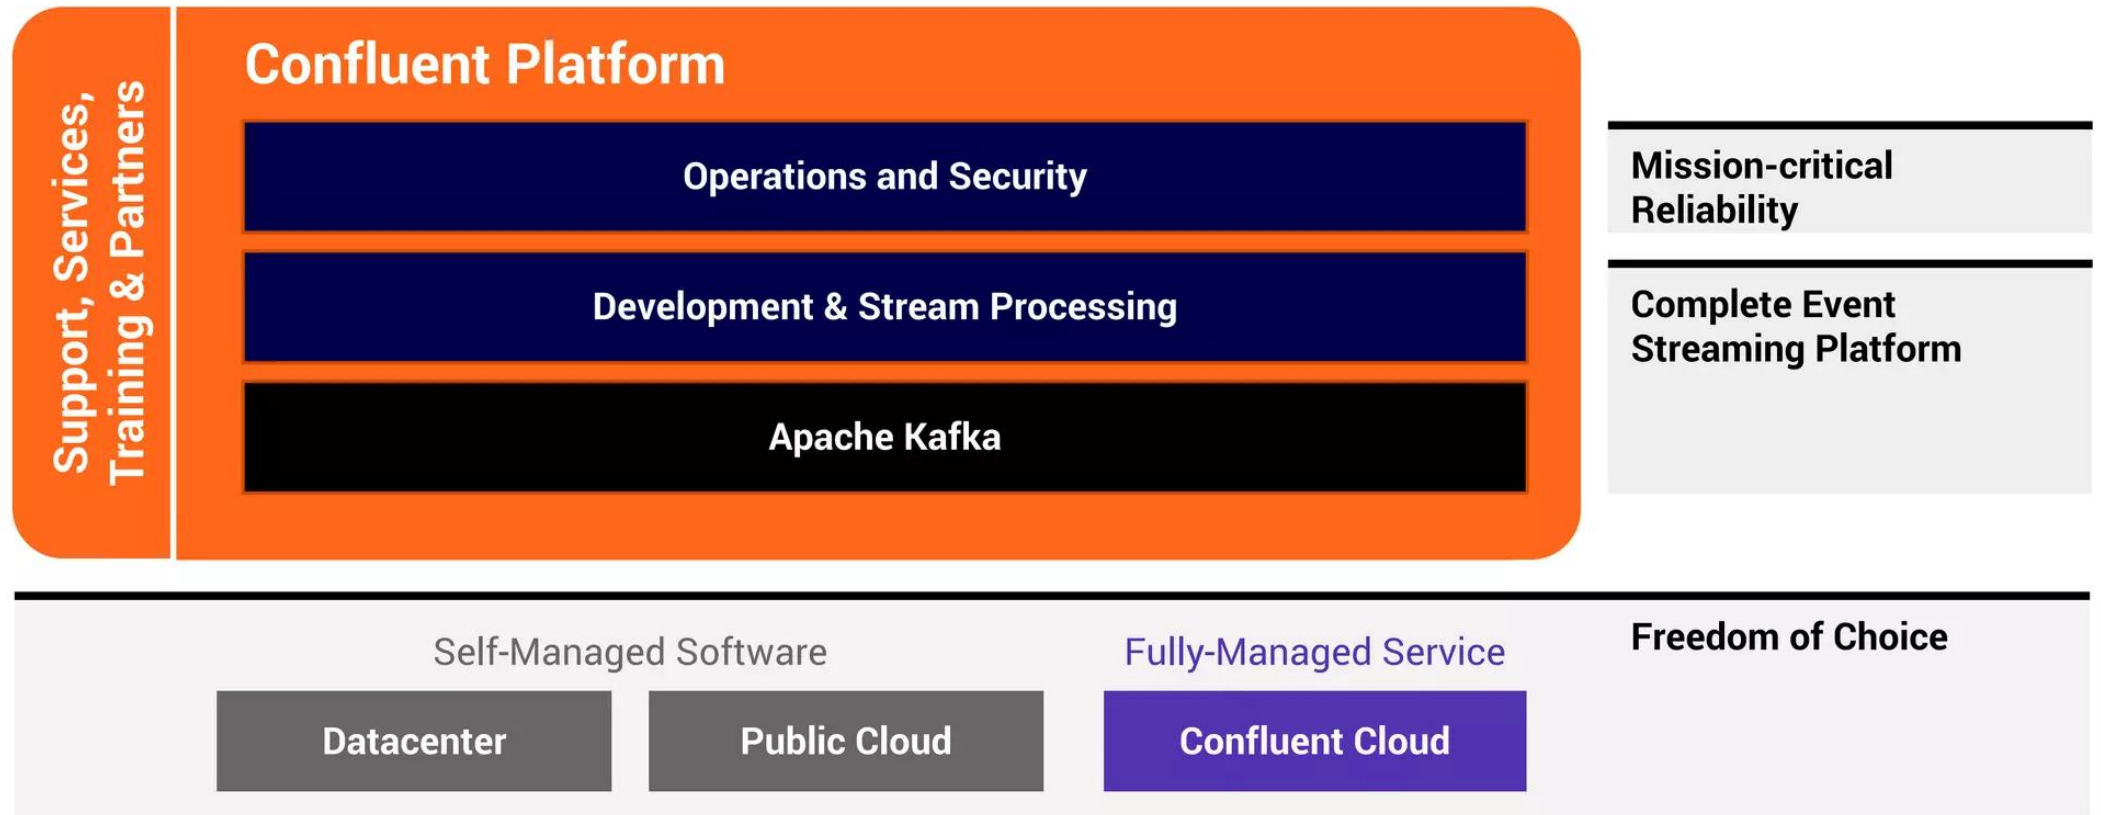
# Apache Kafka<sup>®</sup>: A Distributed Streaming Platform






# Confluent Platform

The Event Streaming Platform Built by the Original Creators of Apache Kafka®





# Kafka Streams

(aka Streams API)

# Streams API

—

**Java client library  
for building  
distributed apps**

—

**Part of the  
Apache Kafka  
project**

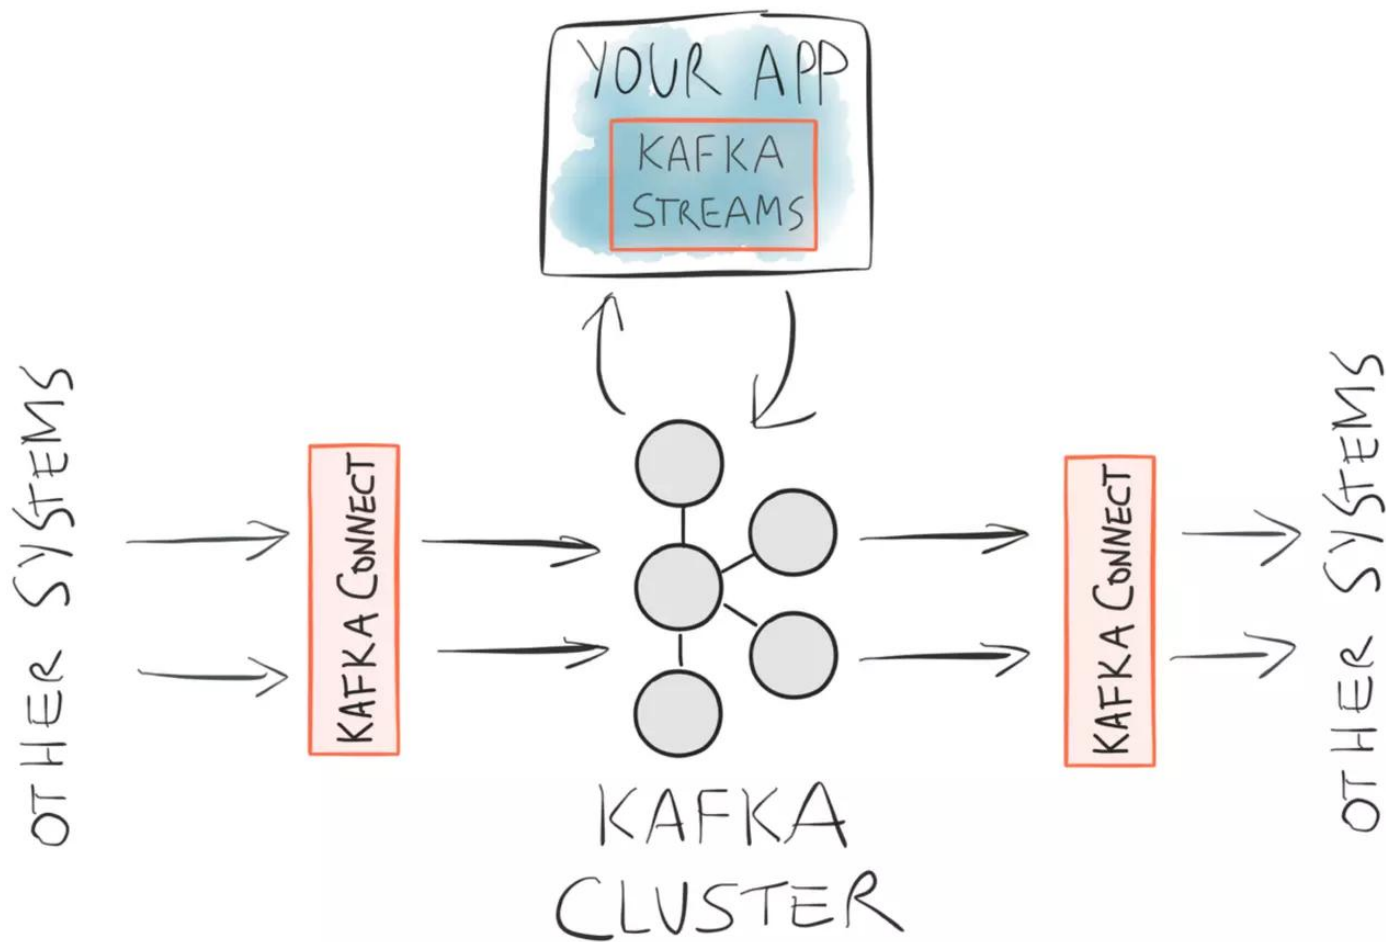
—

**Read-process-  
write pattern**

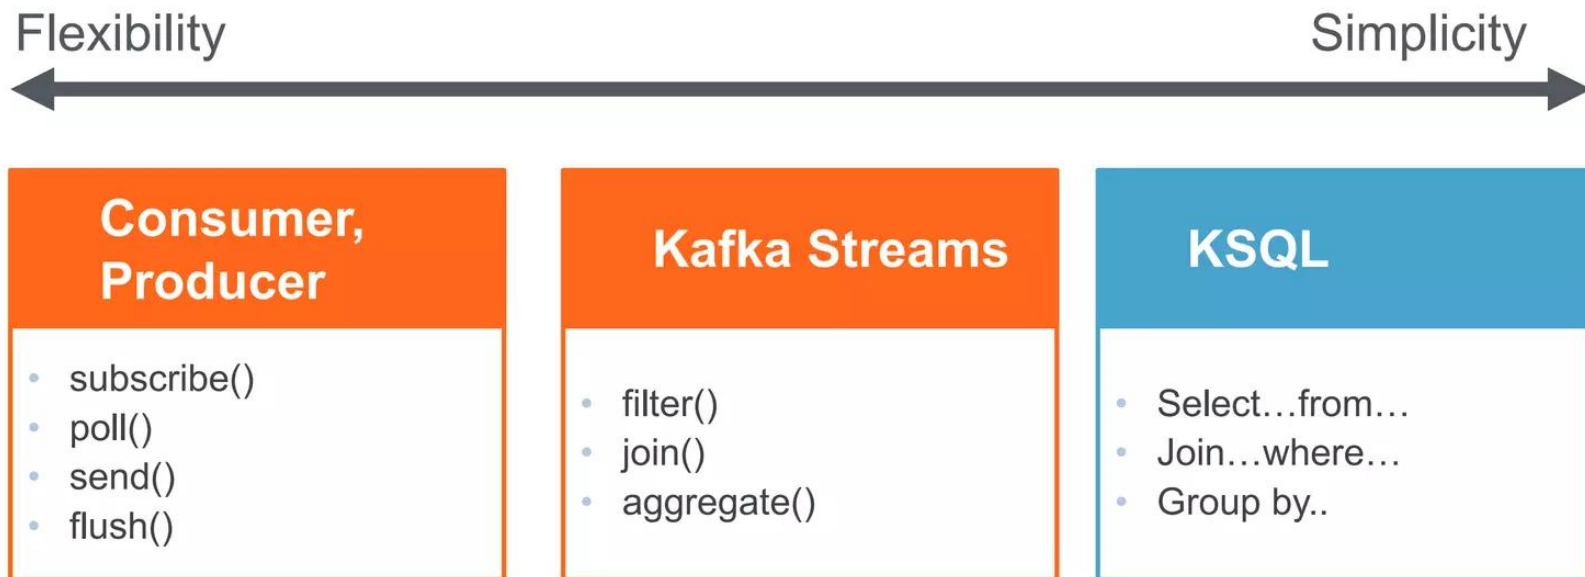
—

**State handling,  
fault-tolerance,  
scaling**






# Why should I use it?



# How can I use it?


```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-streams</artifactId>  
  <version>2.3.0</version>  
</dependency>
```

# How do I deploy my application?




# Scaling Your Application


EXPAND



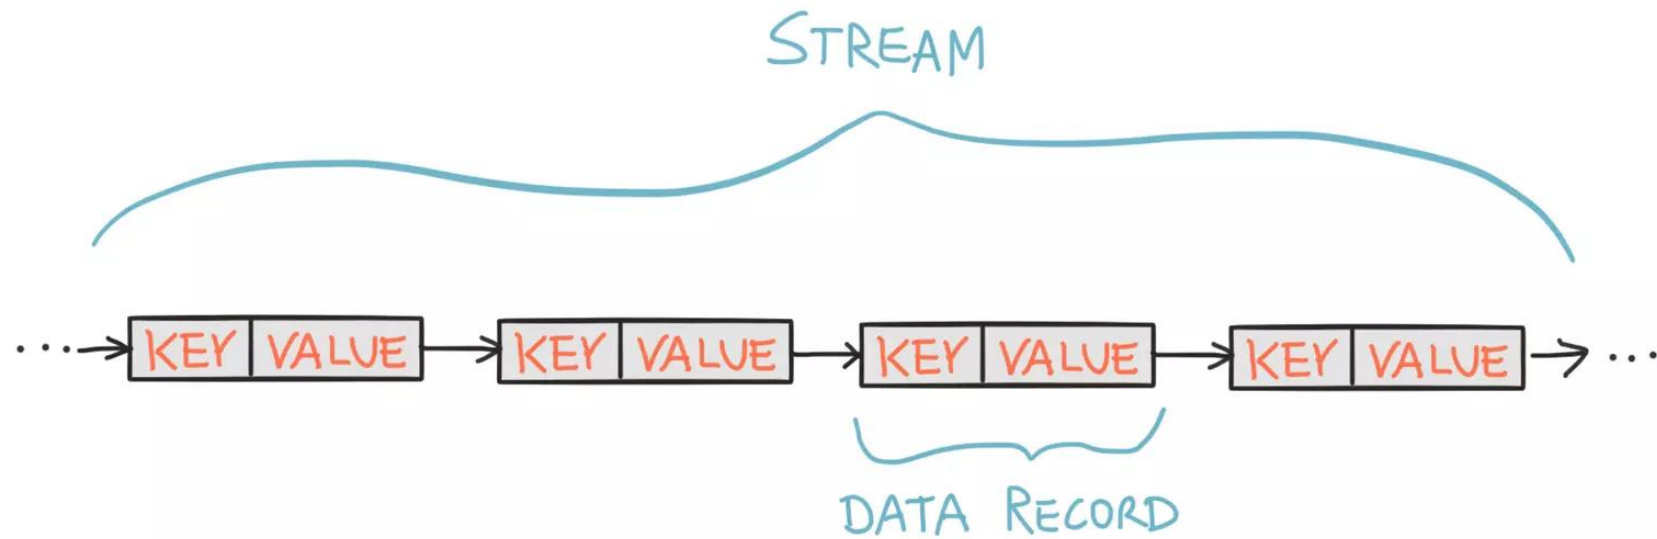
# Scaling Your Application



# Scaling Your Application




# Data Model





# Transformations (map, filter, etc.)



```
StreamsBuilder builder = new StreamsBuilder();
```

```
KStream<Long,String> inputStream =  
    builder.stream("input-topic");
```

```
KStream<Long,String> outputStream =  
    inputStream.mapValues(  
        value -> value.toLowerCase());
```

```
outputStream.to("output-topic");
```

# Stream Aggregations



TRANSFORM

MUTATIONS



```
KTable<Long,Long> countPerKey  
    = stream.groupByKey().count();
```

## Data Model (cont.)

STREAMS ARE EVERYWHERE!

TABLES ARE EVERYWHERE, TOO!

# Streams

  
**Facts/events**

  
**Infinite/unbounded**

  
**Append-only**

  
**Immutable**

# Tables

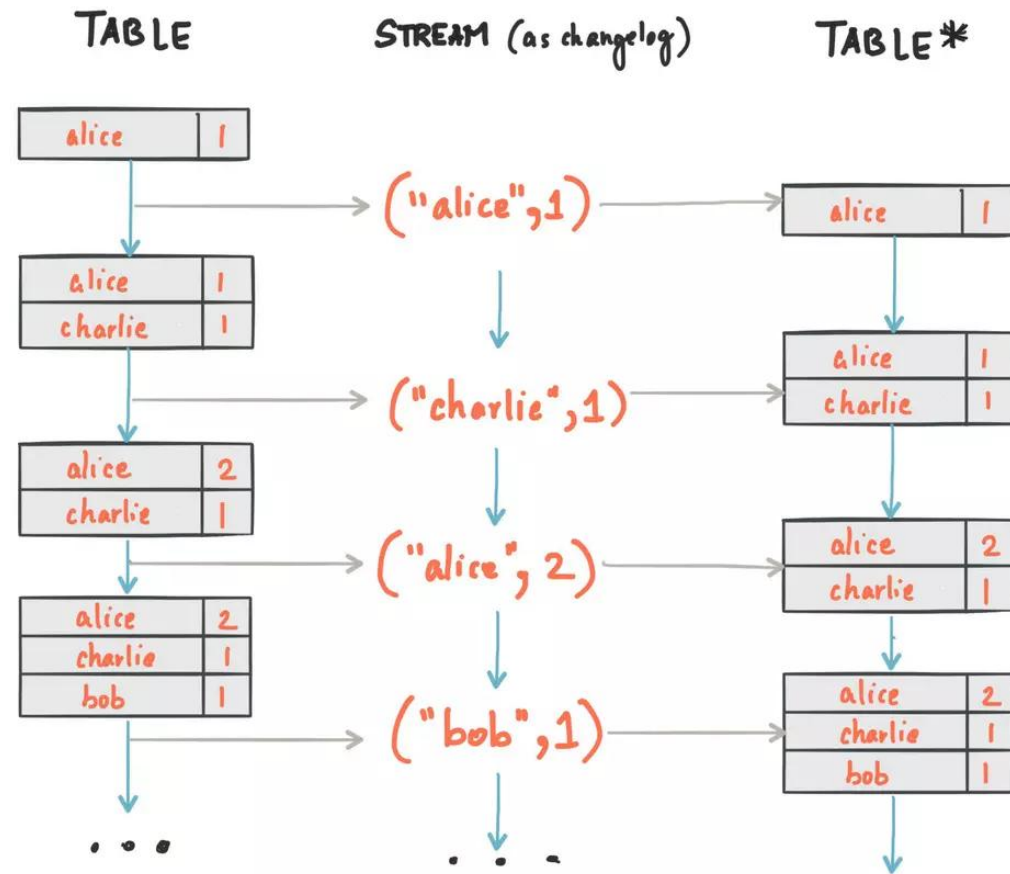
  
**Current state**

  
**Finite/Bounded**

  
**Insert,  
Update,  
Delete**

  
**Mutable**

# Stream-Table Duality




# Stream Aggregations: The Changelog



TRANSFORM


```
KTable<Long,Long> countPerKey  
    = stream.groupByKey().count();  
  
countPerKey.toStream().to("topic");
```

MUTATIONS



# Creating Tables

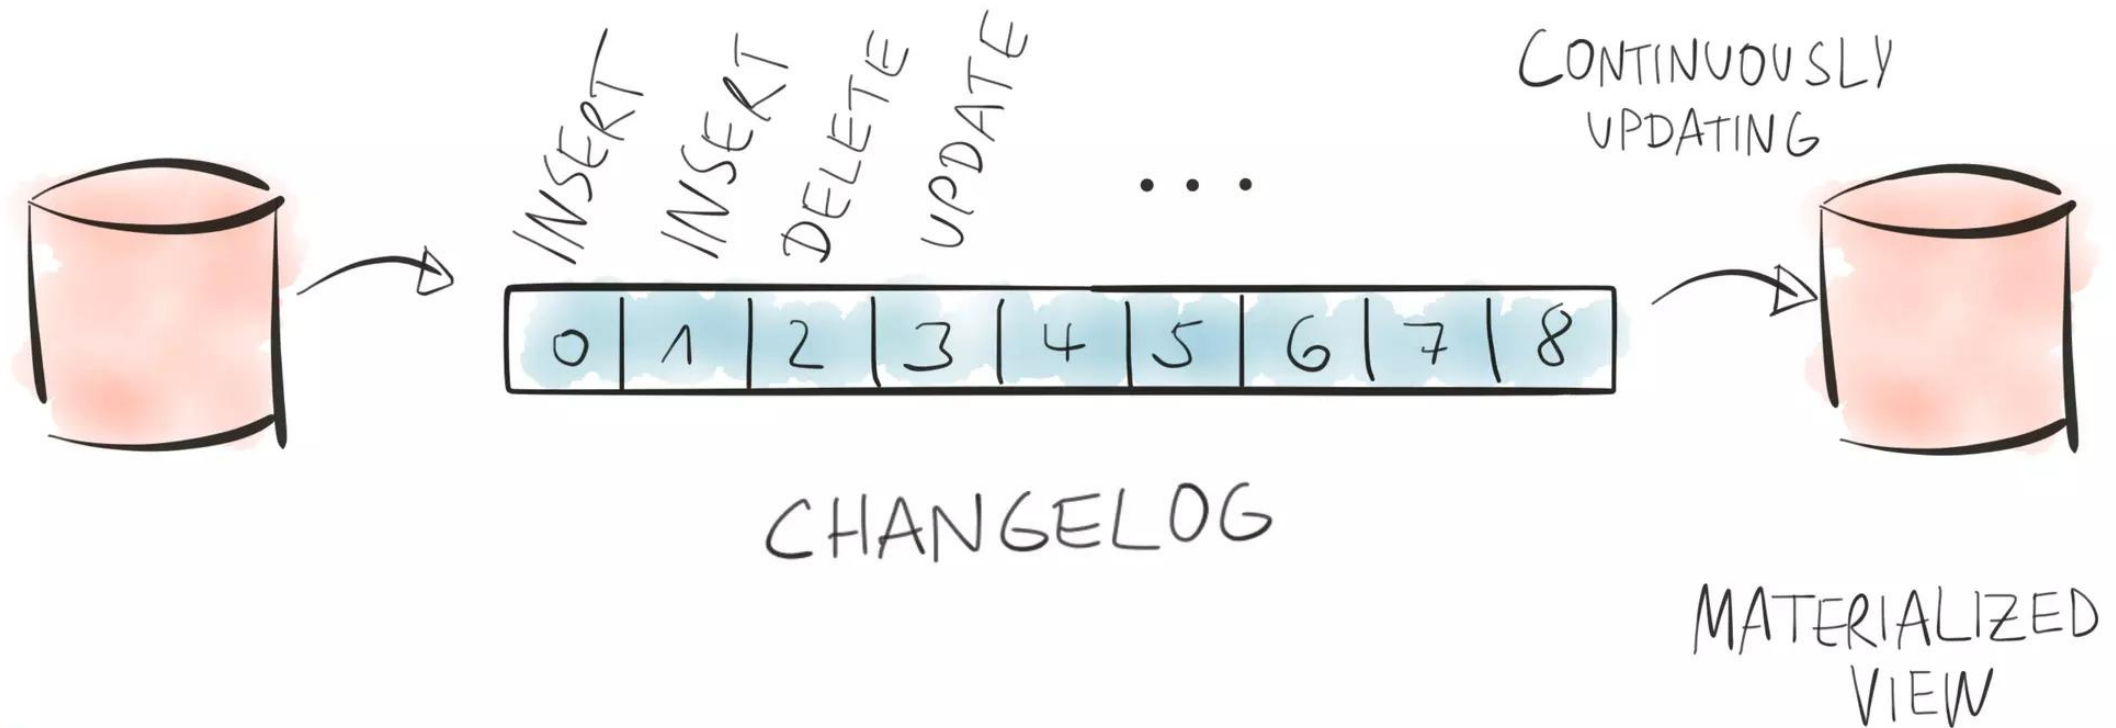
```
KTable<Long,String> inputTable = builder.table("changelog-topic");
```






# Transforming Tables

```
KTable<Long,String> outputTable = inputTable.filter();
```



# Joining Streams and Tables



```
KStream<Long,String> enrichedStream =  
    inputStream.join(inputTable, ...);
```

groupBy

windowBy  
(tumbling,  
hopping,  
session)

flatMap

filter


foreach

mapValues

join  
(stream-stream,  
stream-table,  
table-table)

aggregate

# Processor API



```
Topology topology = new Topology();

topology.addSource(...);
topology.addProcessor(...);
topology.addSink(...);

topology.addStateStore(...);

// DSL compiles down to a topology:
// StreamsBuilder builder = ...
// Topology topology = builder.build();

KafkaStreams instance =
    new Kafka Streams(topology, config);
instance.start();
```

# The Processor Interface

```
public class MyProcessor<K,V>
    implements Processor<K,V> {

    private ProcessorContext context;
    private KeyValueStore store;

    public void init(ProcessorContext c) {
        context = c;
        store = context.getStateStore(...);
    }

    public void process(K key, V value) {
        // business logic
        state.put(key, value);
        context.forward(key, value);
    }

    public void close() {}
}
```

- Full flexibility
- Emit zero or more output records
- Lookup or modify data in your state stores
- Access metadata (e.g., timestamp, headers, etc) via the context
- Schedule punctuations:
  - `context.schedule(...)`
  - Event or wall-clock time based
  - Regular callbacks



# Consumer, Producer

- subscribe()
- poll()
- send()
- flush()

- subscribe()
- poll()
- send()
- flush()

# Kafka Streams

- `filter()`
- `join()`
- `aggregate()`

- filter()
- join()
- aggregate()

# KSQL

- Select...from...
- Join...where...
- Group by..

- Select...from...
- Join...where...
- Group by..

Flexibility

Simplicity



### Consumer, Producer

- subscribe()
- poll()
- send()
- flush()

### Streams PAPI

- source / processor / sink
- state store
- process()
- punctuations

### Streams DSL

- filter()
- join()
- aggregate()

### KSQL

- Select...from...
- Join...where...
- Group by..



# How to get started...

- <https://docs.confluent.io/current/streams/index.html>
- <https://github.com/confluentinc/kafka-streams-examples>
- <https://www.confluent.io/resources/>
  - Kafka Summit talk recording
  - White papers
  - Books
- <https://www.manning.com/books/kafka-streams-in-action>
  - By Bill Bejeck

## Getting help:

- User mailing list: <https://kafka.apache.org/contact>
- Confluent Google Group: <https://groups.google.com/forum/#!topic/confluent-platform>
- Confluent Community Slack: <https://launchpass.com/confluentcommunity>
- StackOverflow



Until the end of 2019 Confluent are giving new users \$50 of free usage per month for their first 3 months

## Here's advice on how to use this promotion to try Confluent Cloud for free!

### Sign up for a Confluent Cloud account

Please bear in mind that you will be required to enter credit card information but will not be charged unless you go over the \$50 usage in any of the first 3 months or if you don't cancel your subscription before the end of your promotion.



### You won't be charged if you don't go over the limit!

Get the benefits of Confluent Cloud, but keep an eye on your account making sure that you have enough remaining free credits available for the rest of your subscription month!!



### Cancel before the 3 months end If you don't want to continue past the promotion

If you fail to cancel within your first three months you will start being charged full price. To cancel, immediately stop all streaming and storing data in Confluent Cloud and email [cloud-support@confluent.io](mailto:cloud-support@confluent.io)

[bit.ly/TryConfluentCloud](https://bit.ly/TryConfluentCloud)

Available on





We are hiring!

 @MatthiasJSax

matthias@confluent.io | mjsax@apache.org