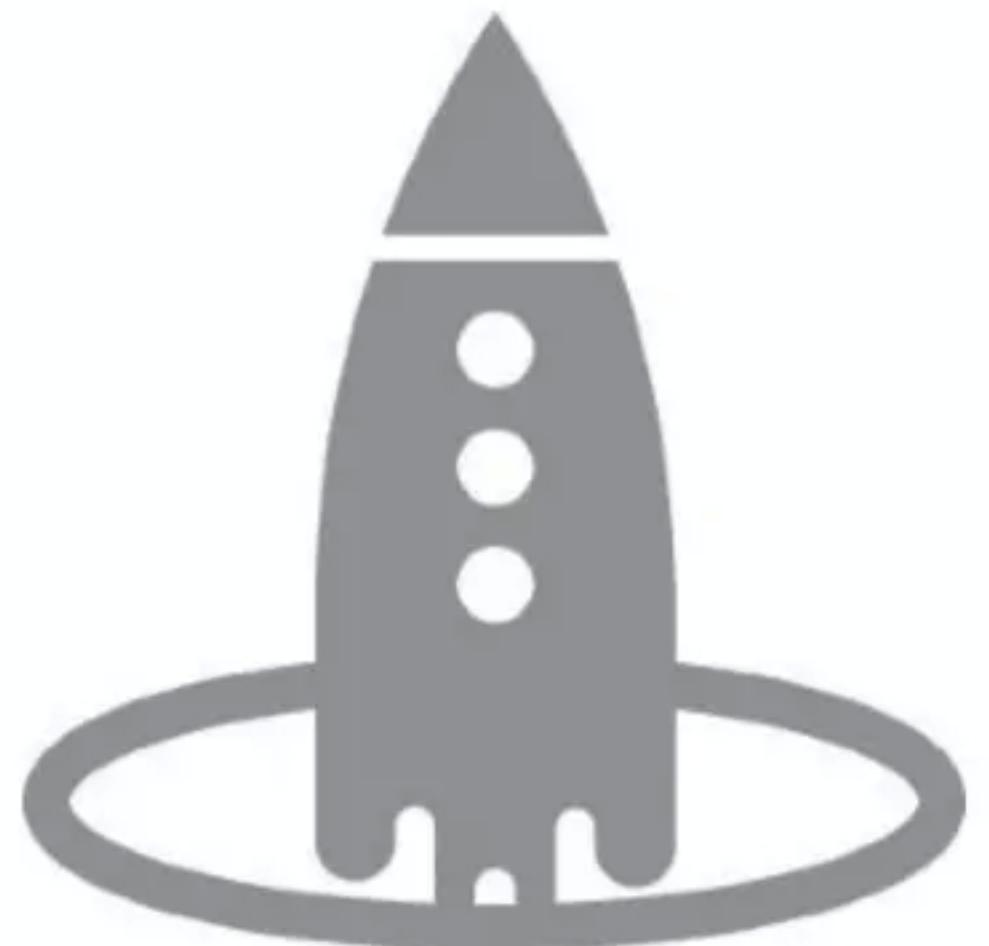




KSQL in Practice

Almog Gavra
Engineer, KSQL

- **Intro**
- Develop
- Deploy
- Operate
- Common Mistakes
- Q&A



Why KSQL?

Why KSQL?

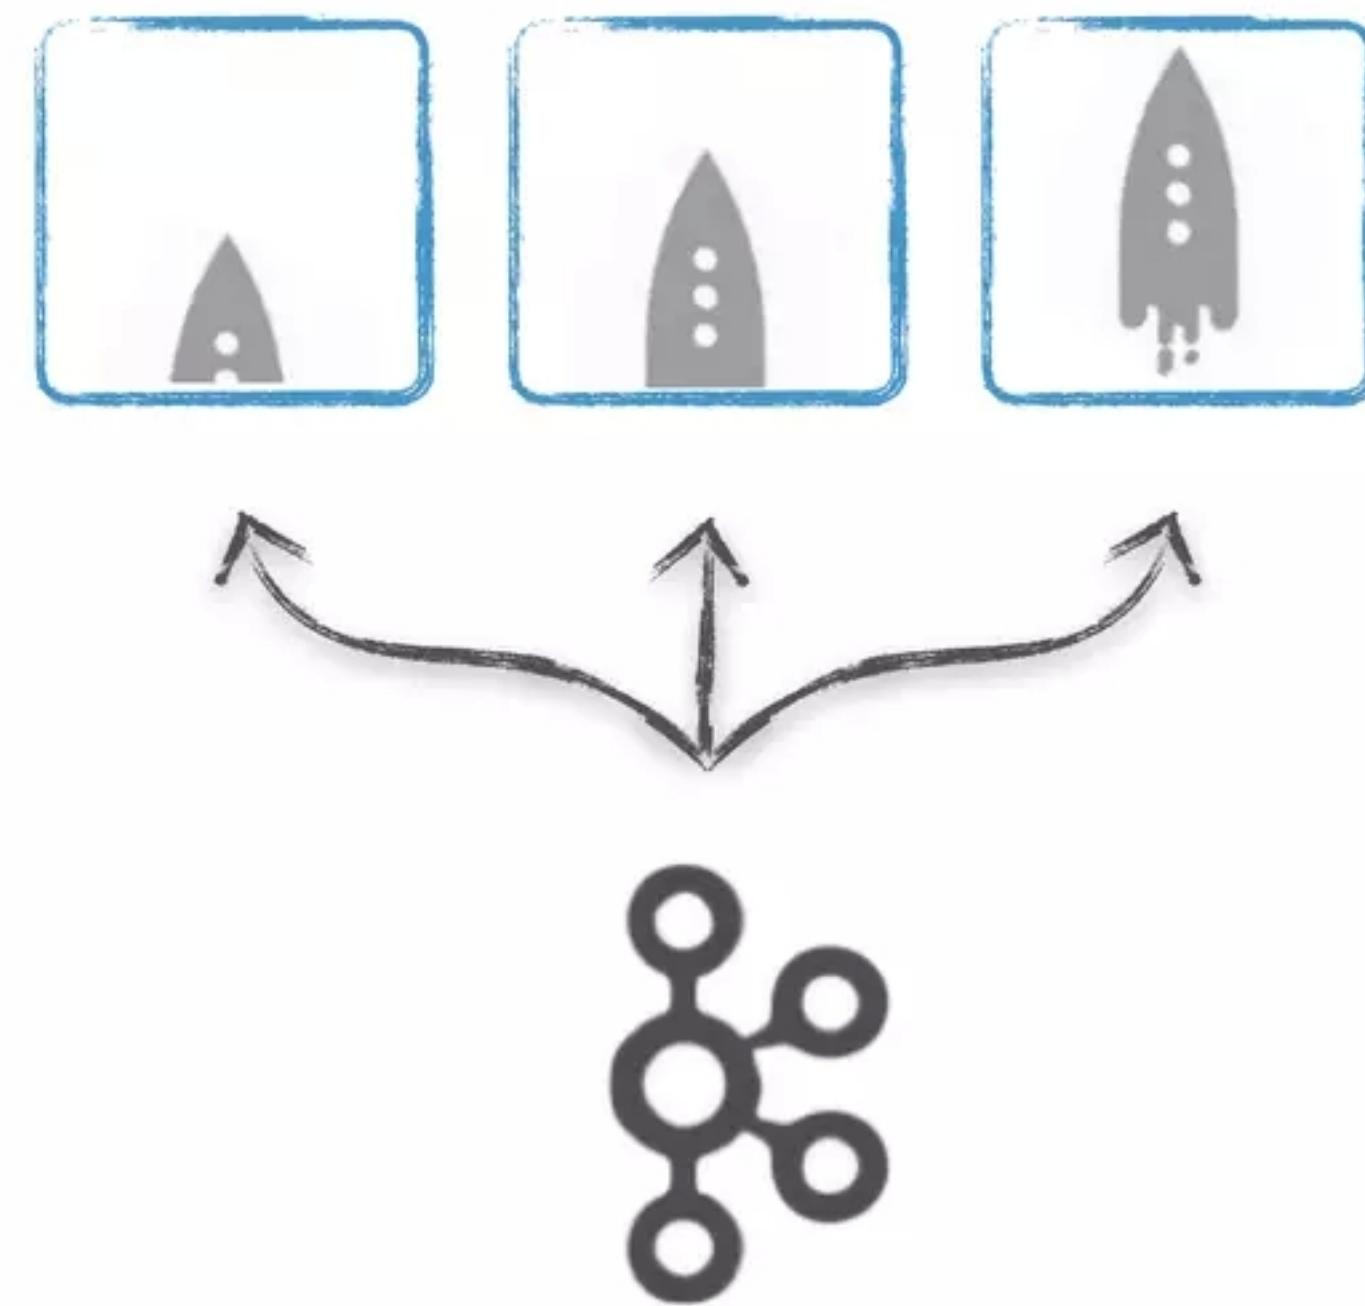
Performance and
Robustness from Apache
Kafka®



Why KSQL?

Simplified Stream
Processing

Performance and
Robustness from Apache
Kafka®

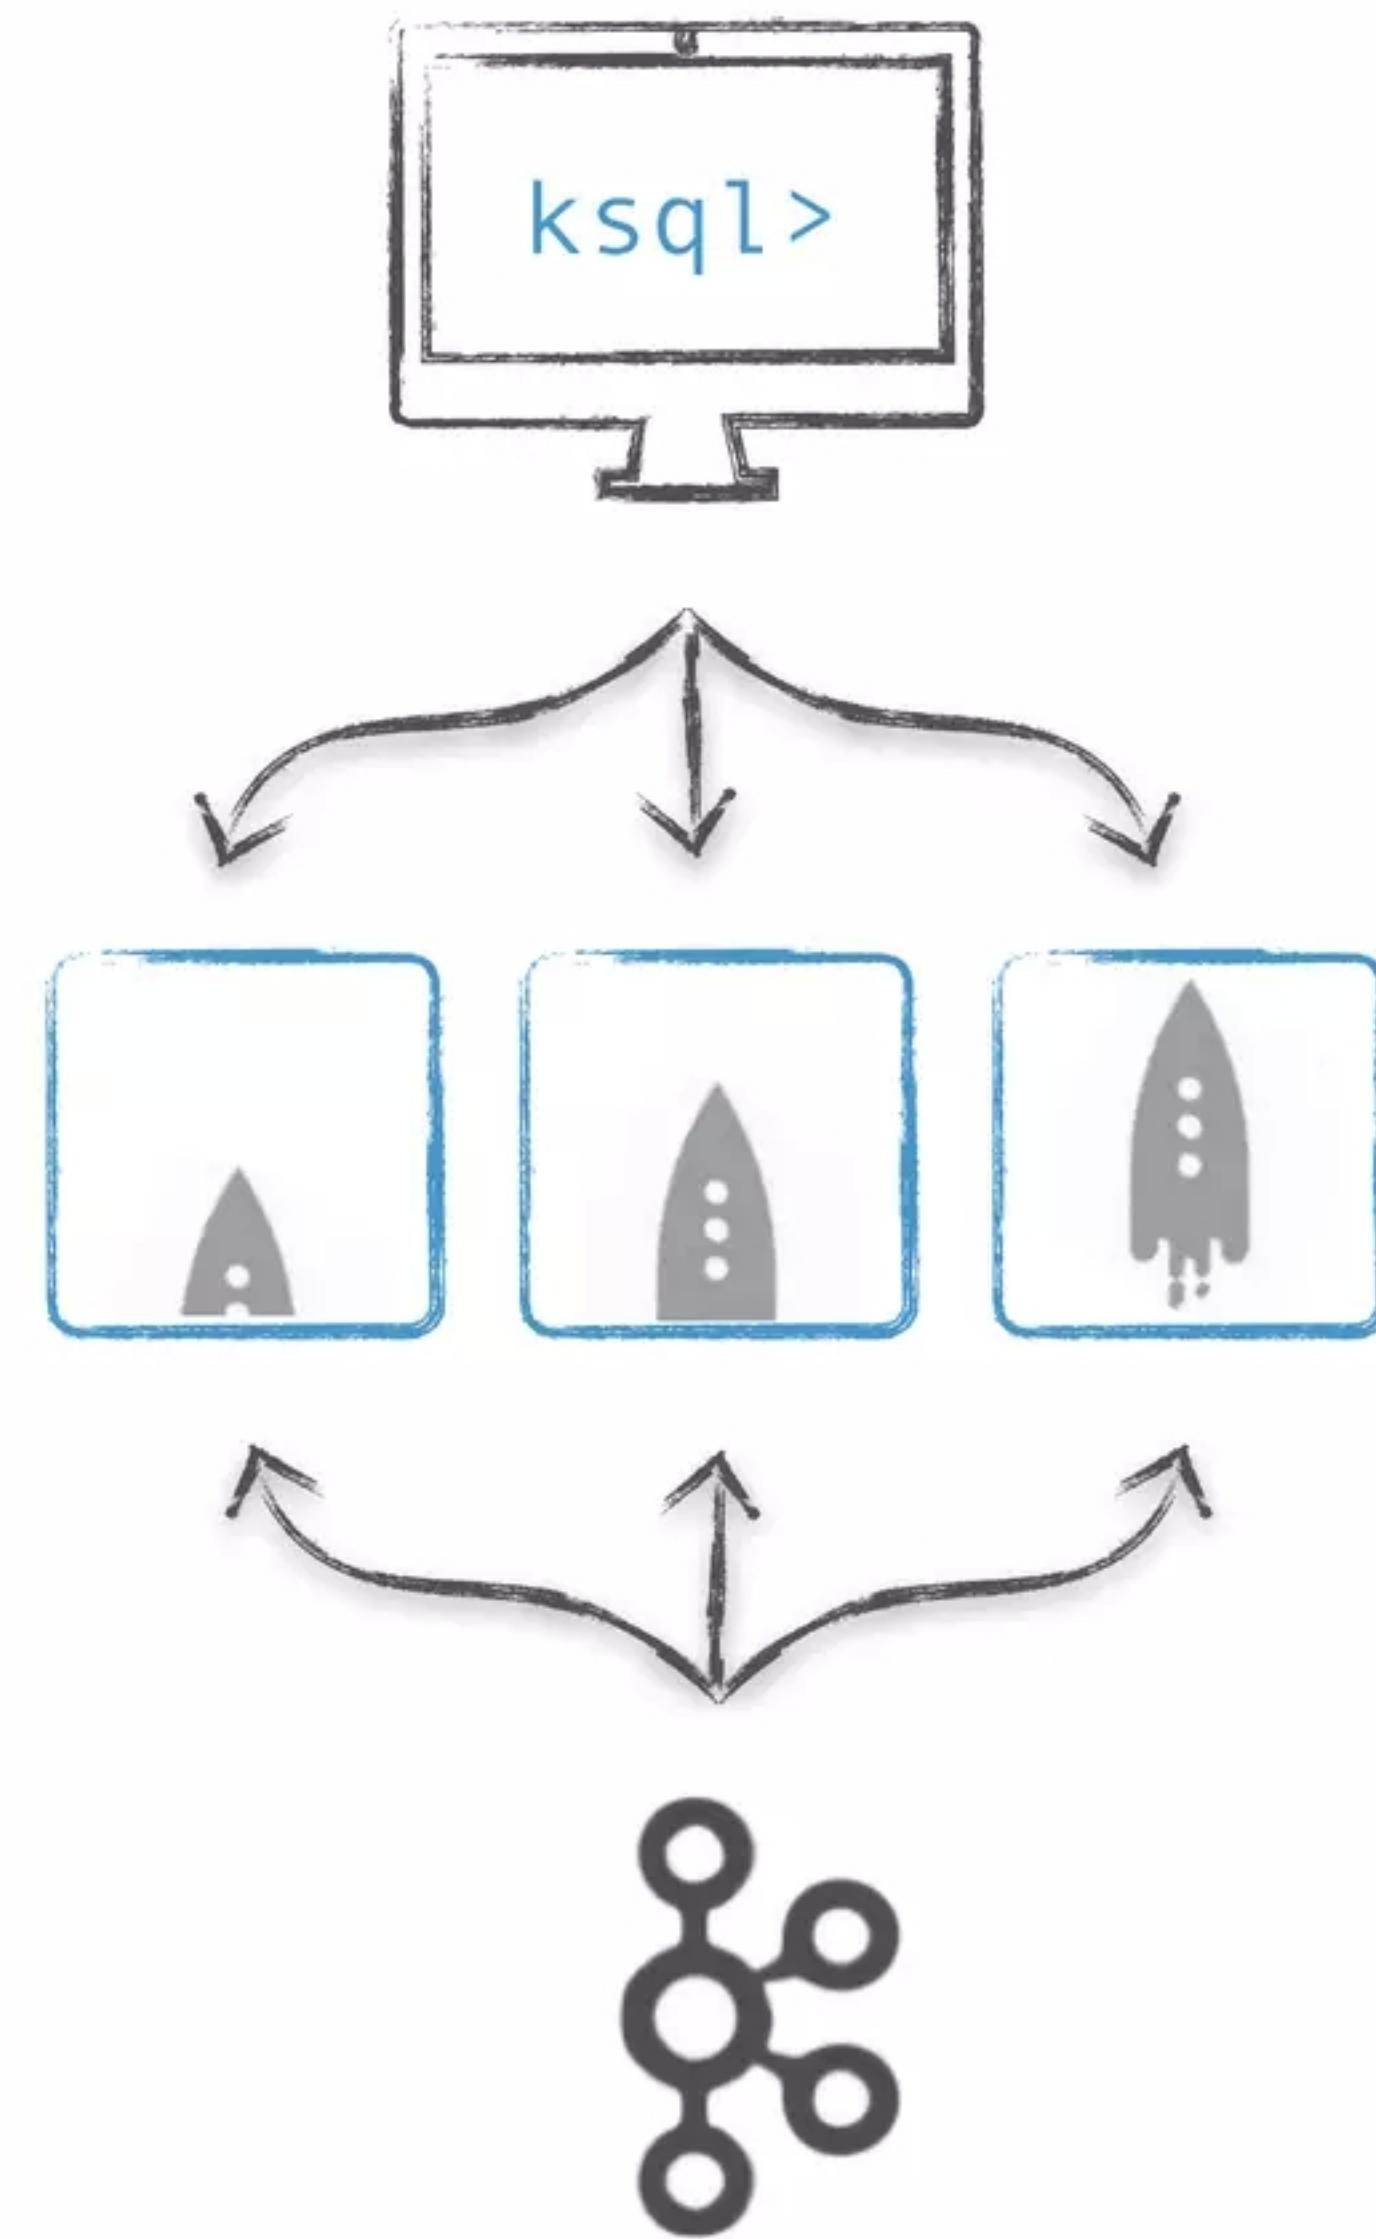


Why KSQL?

Familiar & Expressive
SQL-Like Language

Simplified Stream
Processing

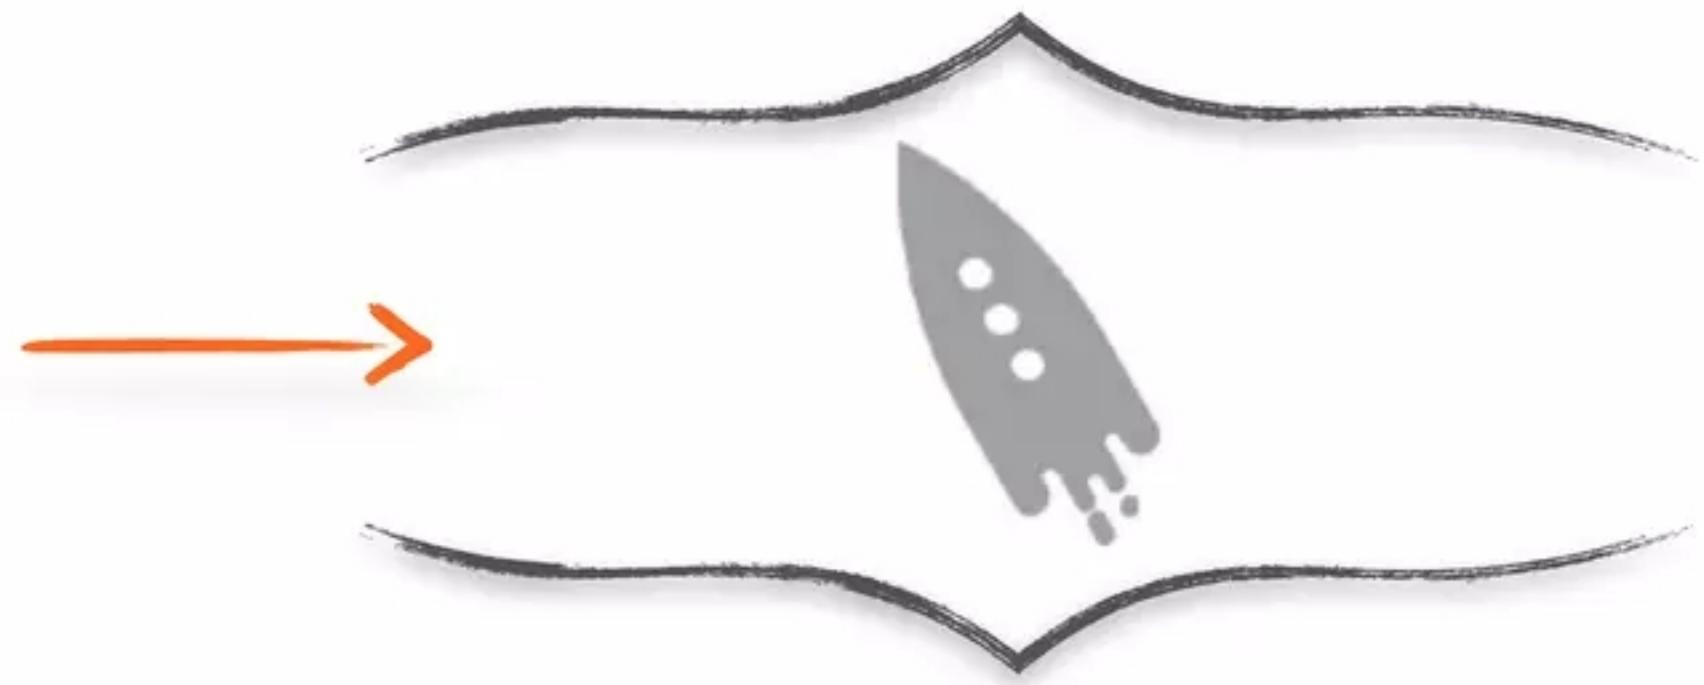
Performance and
Robustness from Apache
Kafka®



```
SELECT * FROM authorizations  
WHERE card_number = 123;
```

User submits SQL

```
SELECT * FROM authorizations  
WHERE card_number = 123;
```



User submits SQL

KSQL builds
KafkaStreams topology

```
SELECT * FROM authorizations  
WHERE card_number = 123;
```



User submits SQL

KSQL builds
KafkaStreams topology

KafkaStreams
executes the topology

Fraud Detection via KSQL



“find all card numbers that had more than 3 authorization attempts within 5 seconds”

Fraud Detection via KSQL

```
CREATE TABLE possible_frauds AS  
SELECT card_number, count(*)  
FROM authorizations  
GROUP BY card_number  
HAVING count(*) > 3;
```

“find all card numbers that had more than 3 authorization attempts within 5 seconds”

Fraud Detection via KSQL

```
CREATE TABLE possible_frauds AS
SELECT card_number, count(*)
FROM authorizations
WINDOW TUMBLING (SIZE 5 SECONDS)

GROUP BY card_number
HAVING count(*) > 3;
```

“find all card numbers that had more than 3 authorization attempts within 5 seconds”

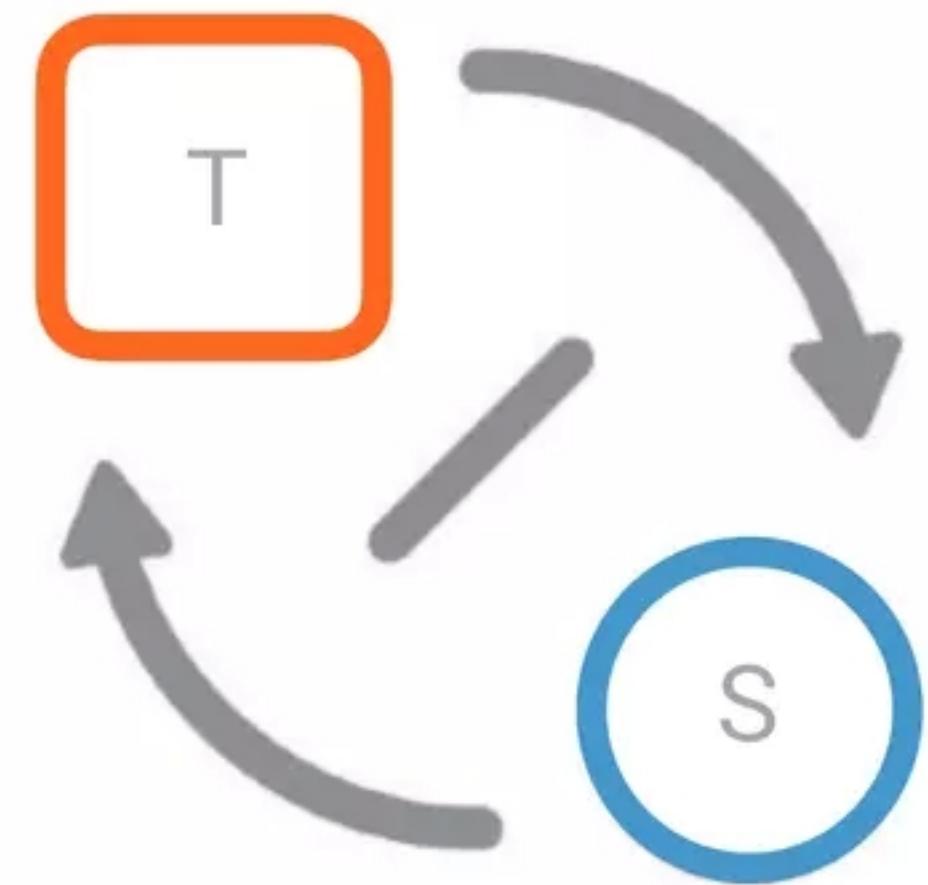
- Intro
- **Develop**
- Deploy
- Operate
- Common Mistakes
- Q&A

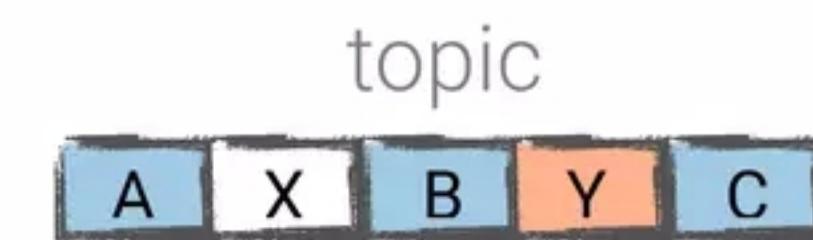
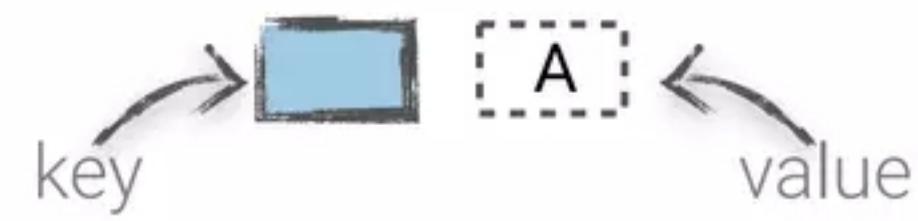


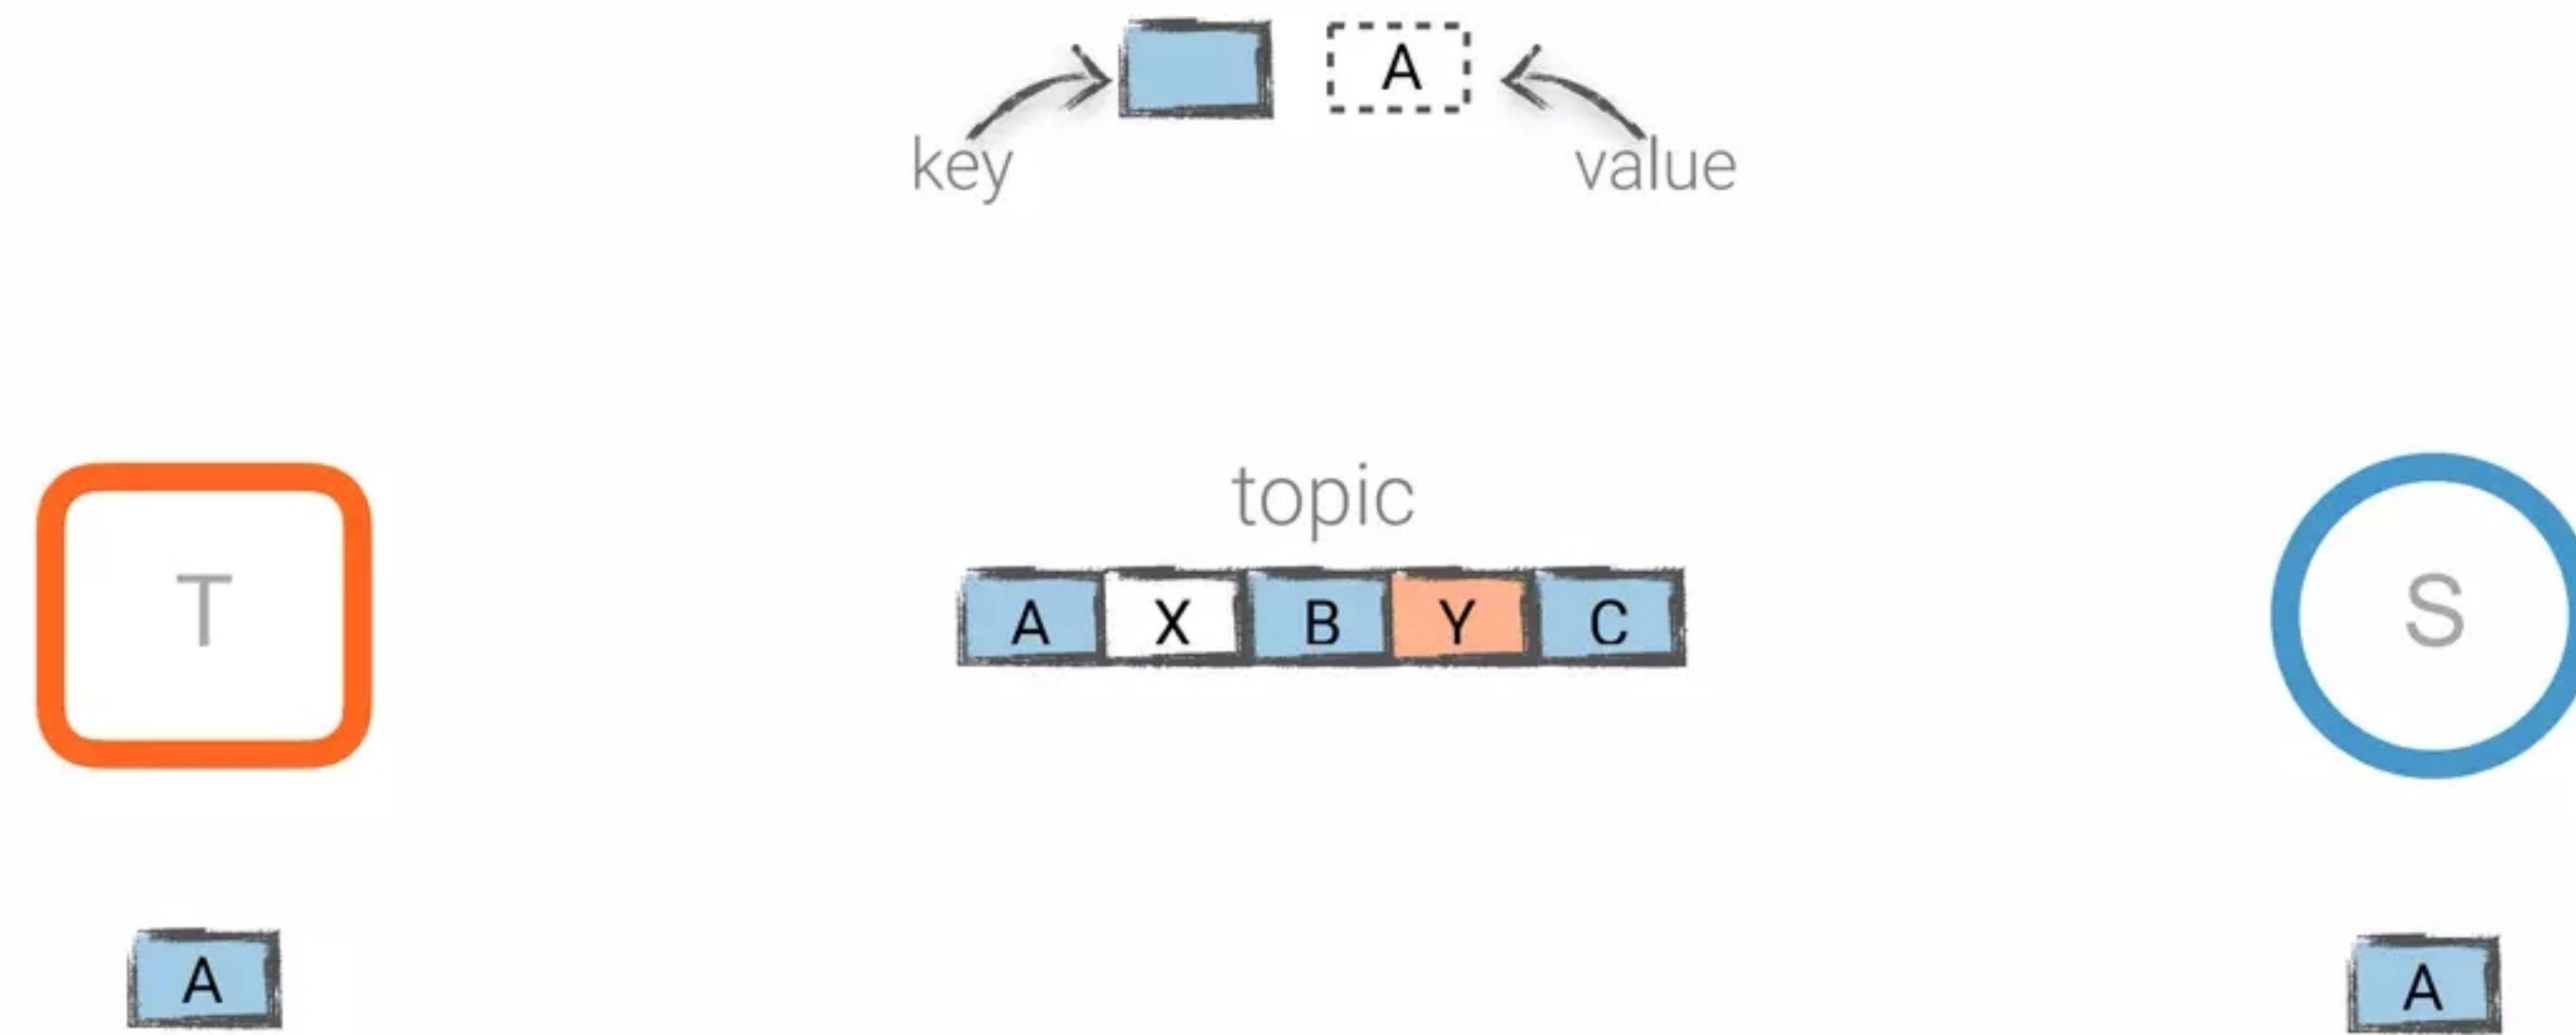
but first...

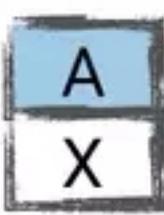
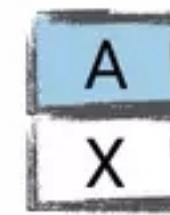
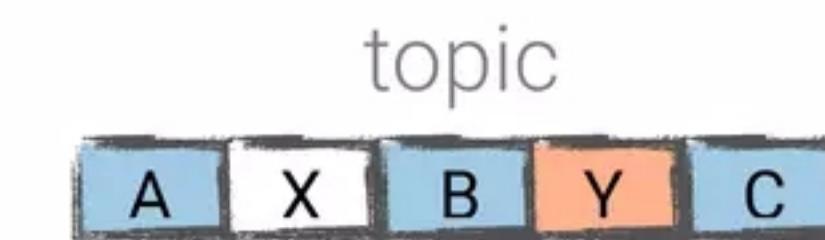
but first...

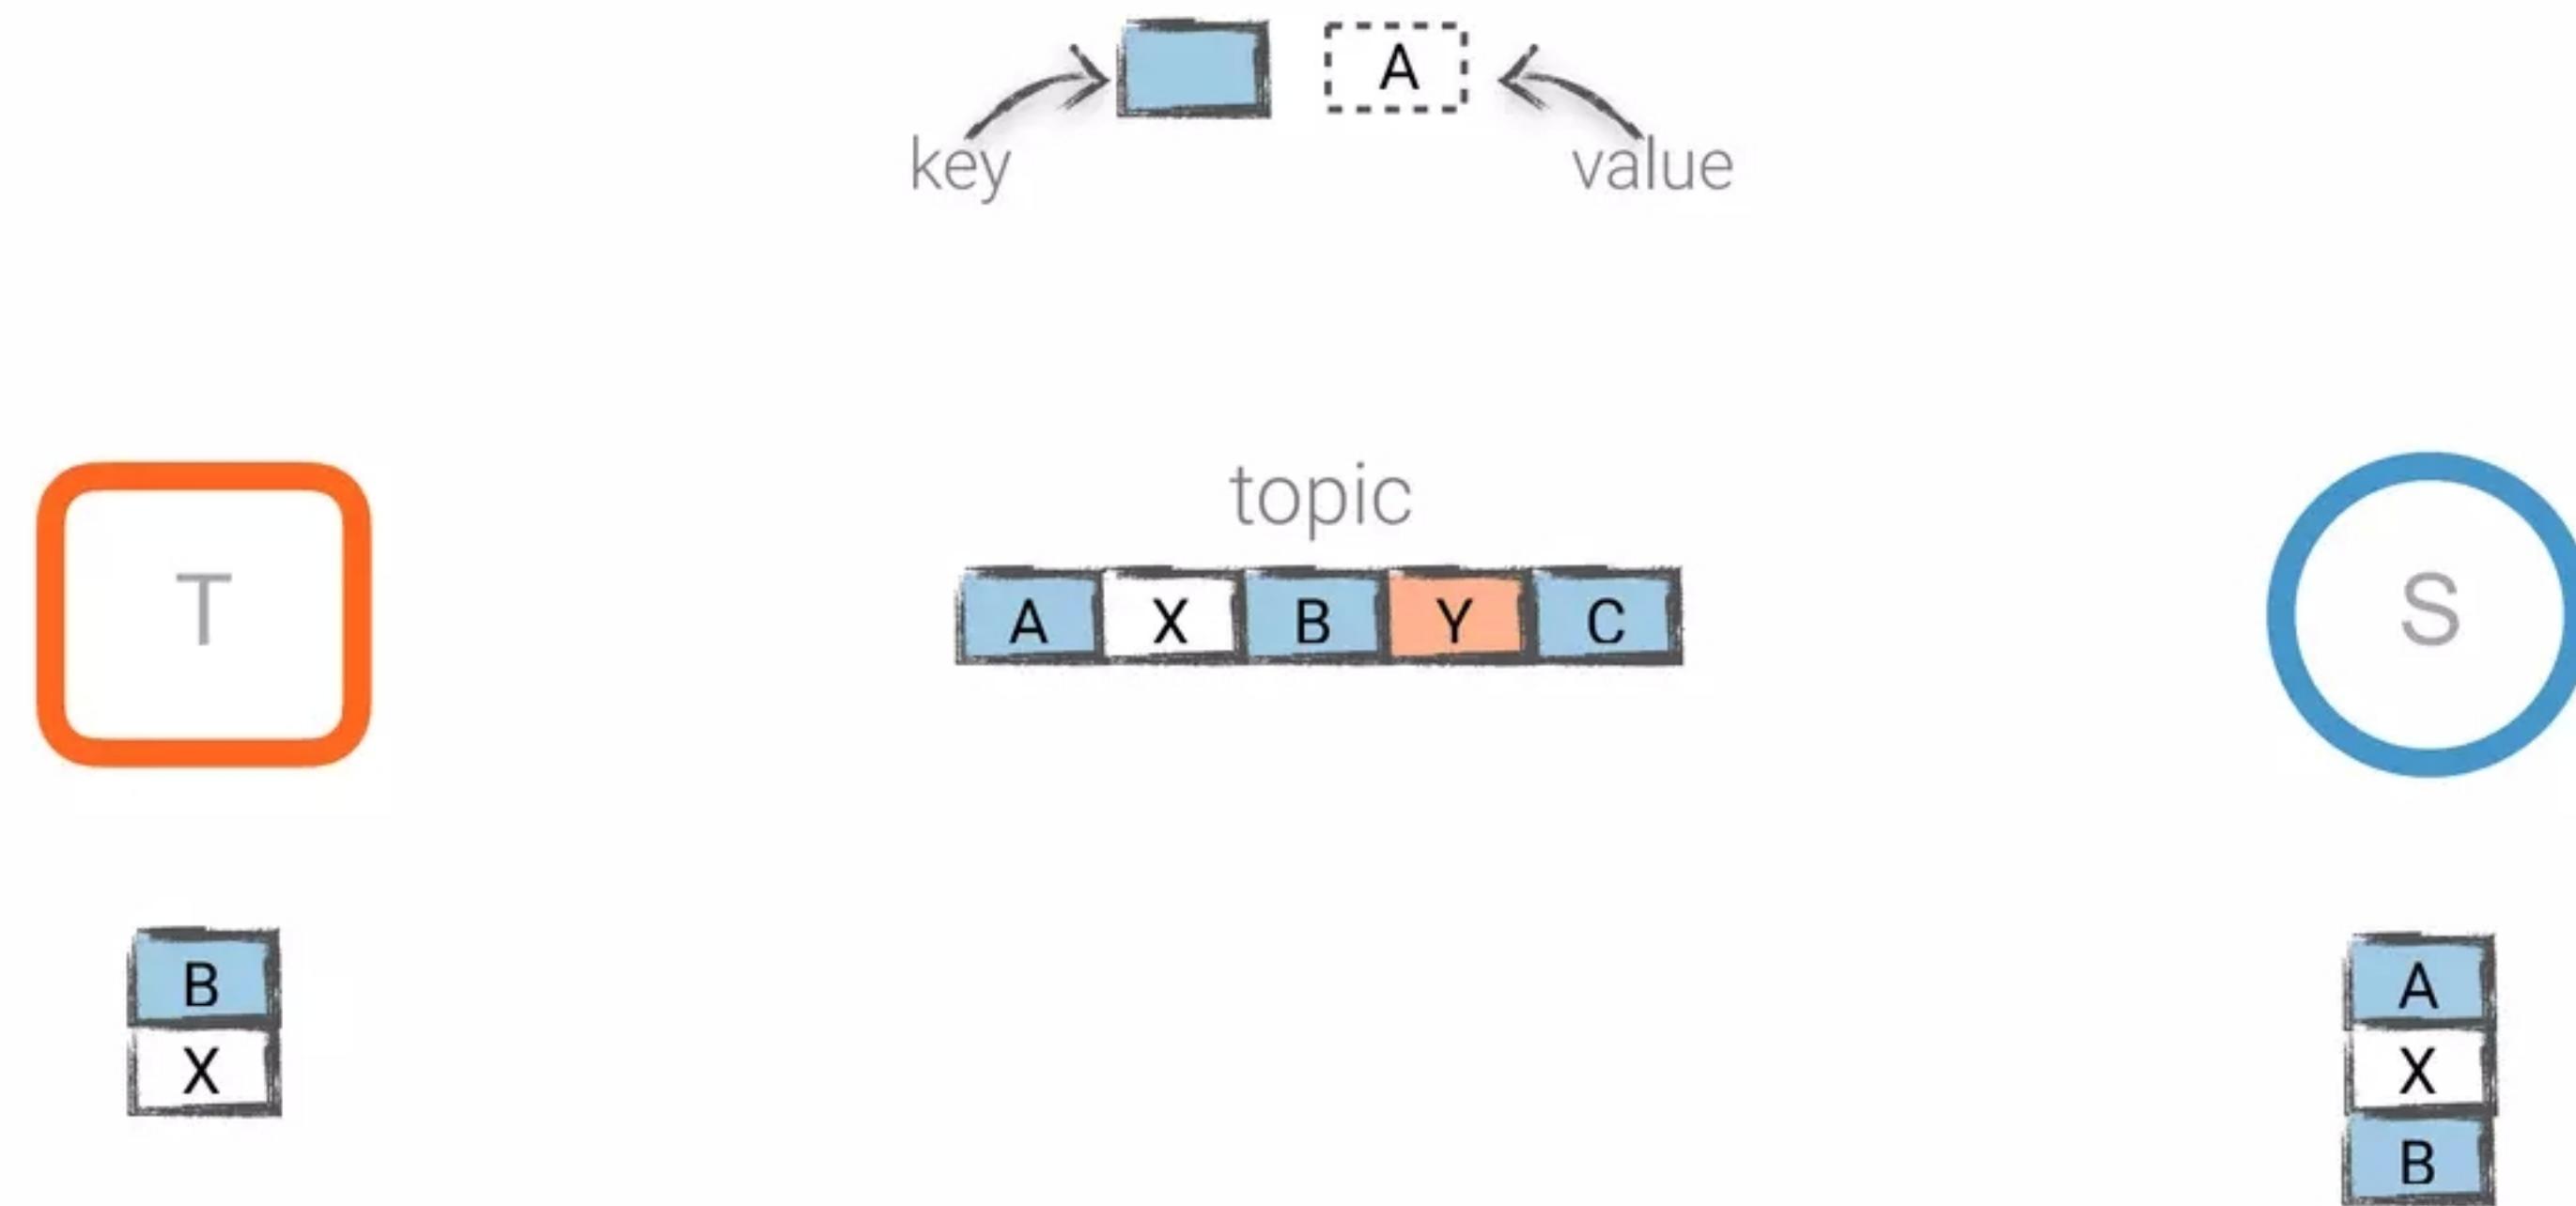
Stream/Table Duality

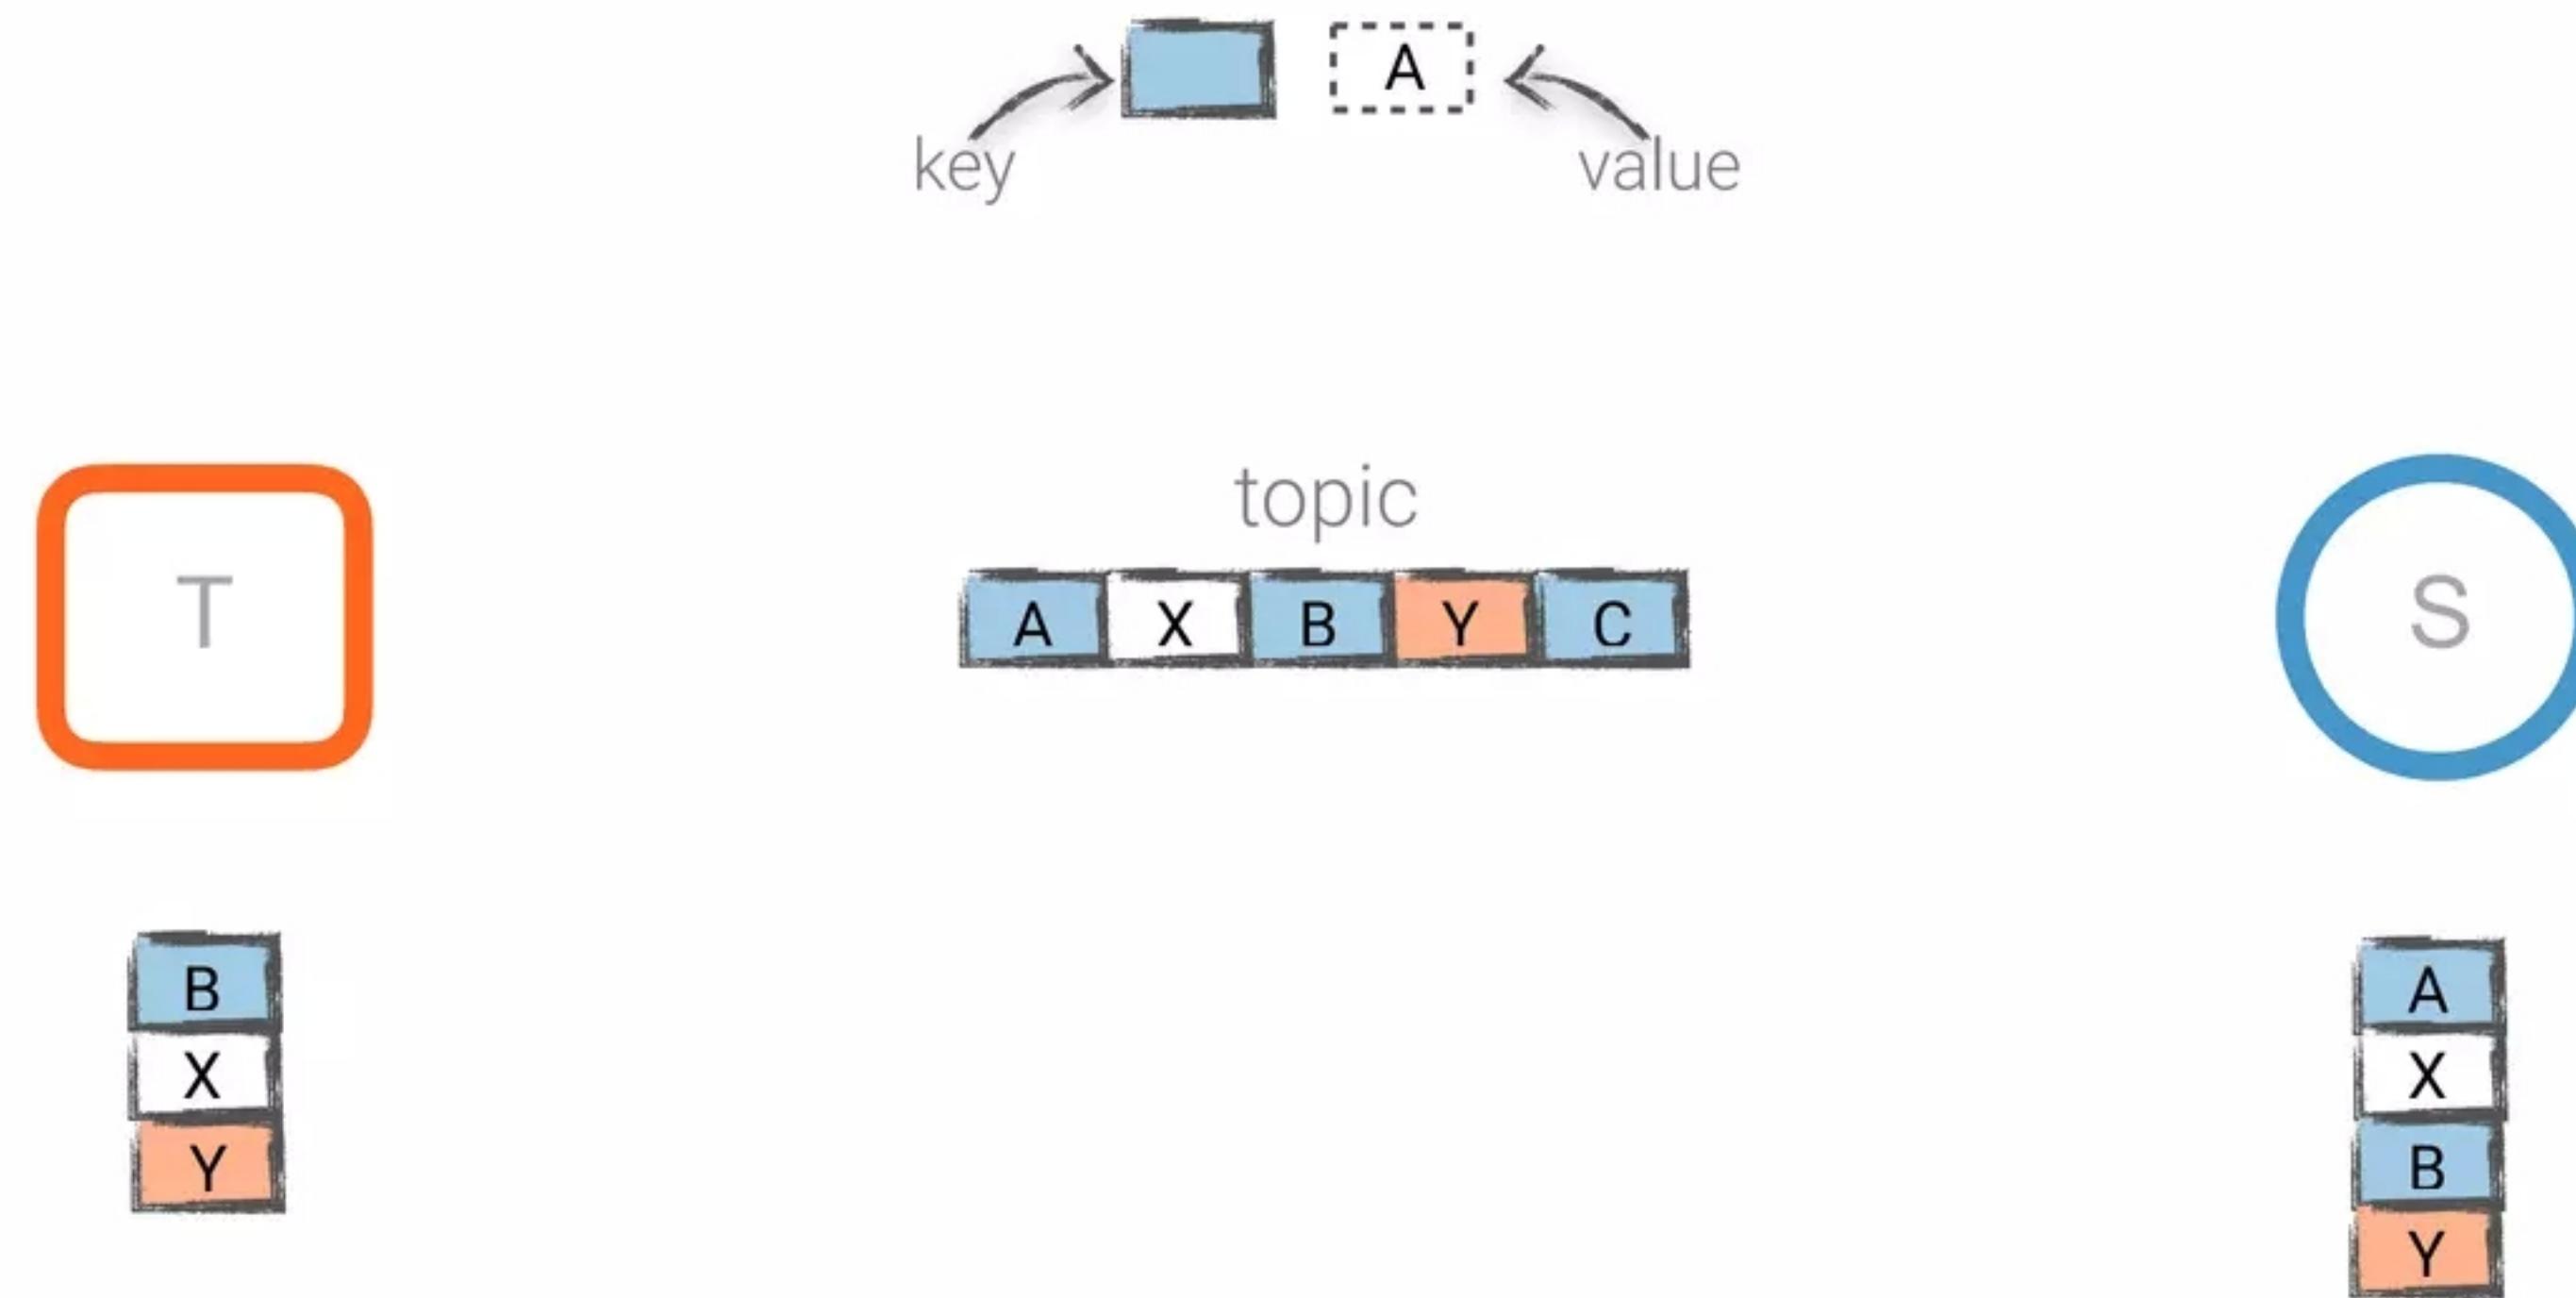


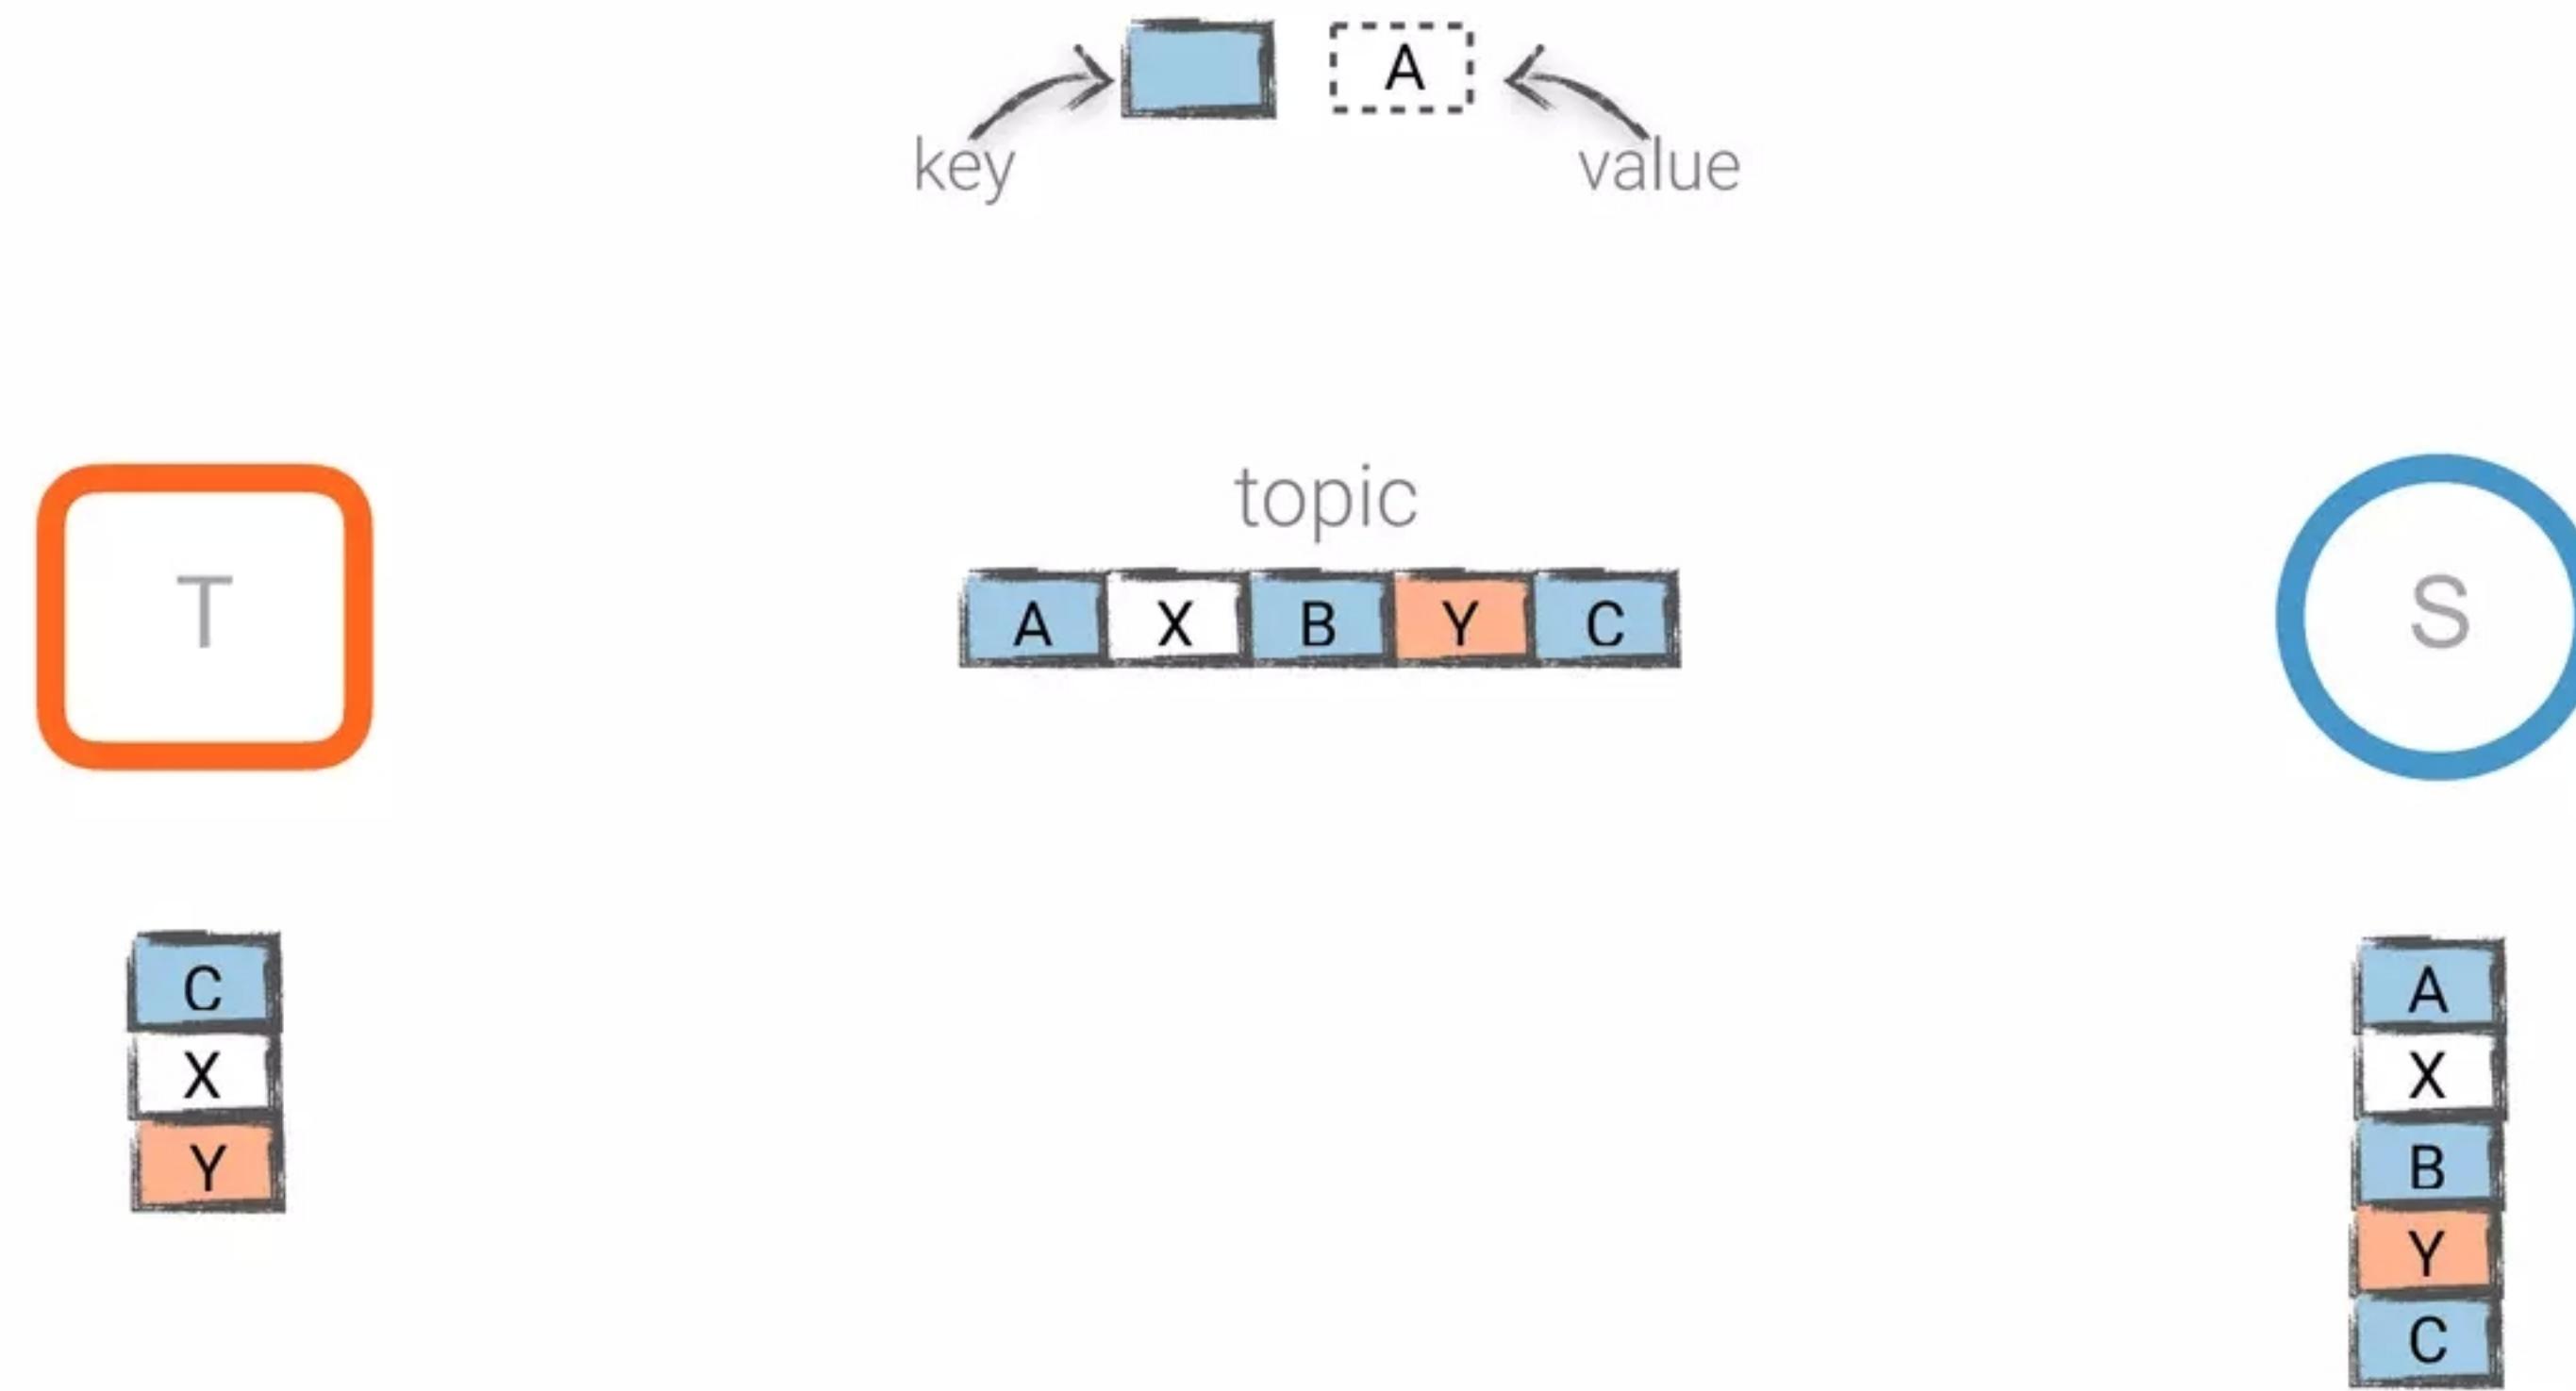


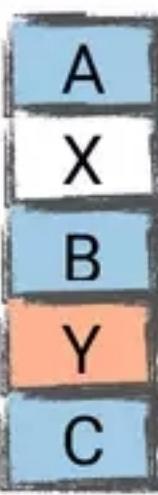
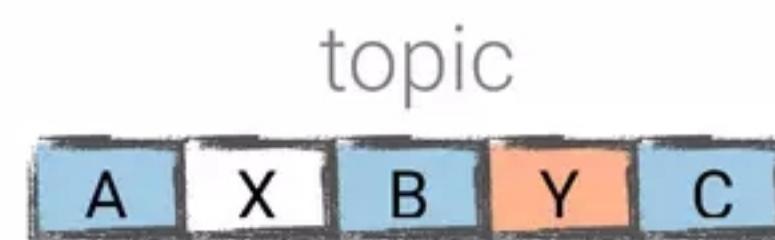




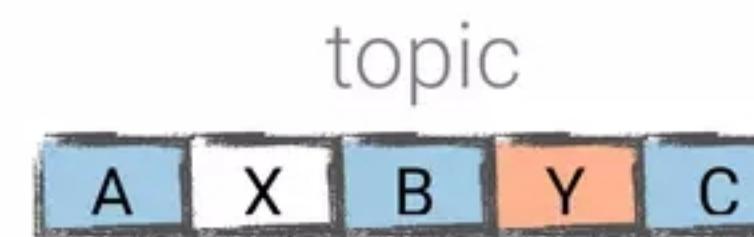








- ▶ a **table** is the *current state* of your data
- ▶ the topic is a *changelog*



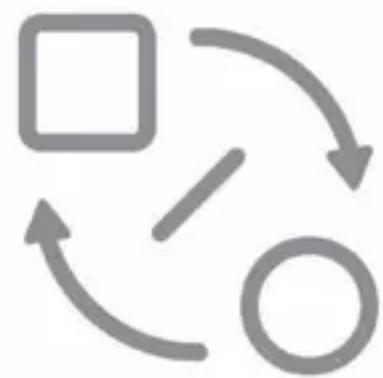
- ▶ a **table** is the *current state* of your data
- ▶ the topic is a *changelog*

- ▶ a **stream** is the *historical state* of your data
- ▶ the topic is a *sequence*

Development Lifecycle



Import Data



Create Stream/Table



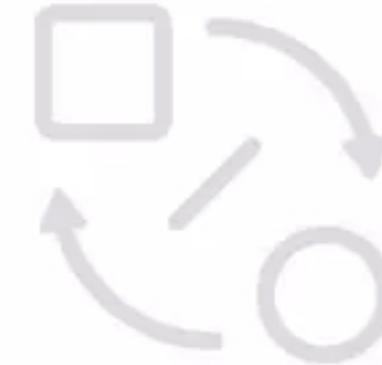
Write KSQL Queries



Verify Output



Import Data



Create Stream/Table



Write KSQL Queries

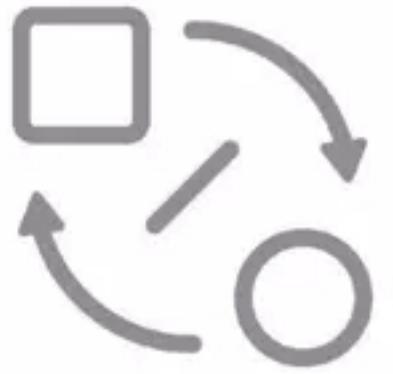


Verify Output

Kafka
Connect`INSERT INTO foo VALUES *``./ksql-datagen`



Import Data



Create Stream/Table



Write KSQL Queries



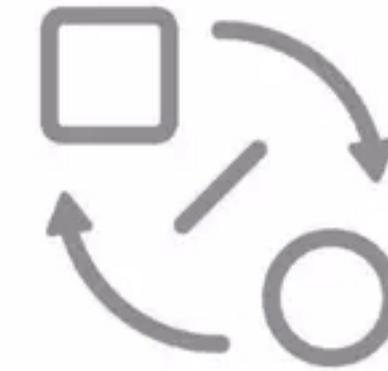
Verify Output

```
CREATE TABLE/STREAM name (schema ...)  
WITH (kafka_topic="topic");
```

DDL



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output

```
CREATE TABLE/STREAM name (schema ...)  
WITH (kafka_topic="topic");
```

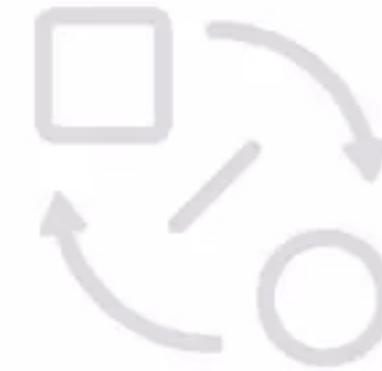
DDL

```
CREATE TABLE/STREAM name (schema ...)  
AS SELECT * FROM source;
```

DML



Import Data



Create Stream/Table



Write KSQL Queries



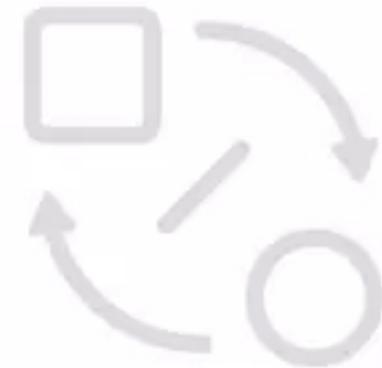
Verify Output



{ REST }



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output



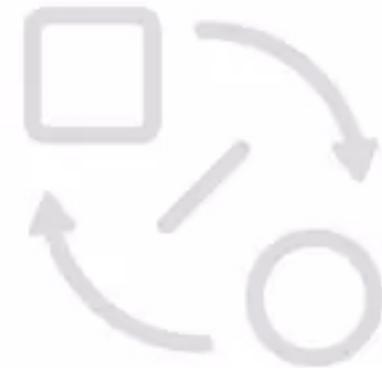
UDF



UDAF



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output

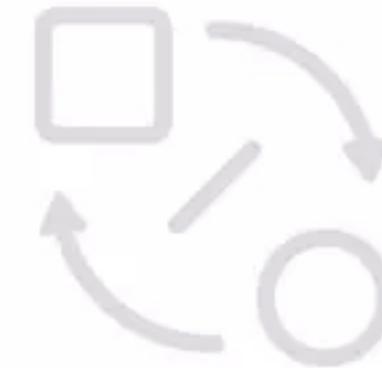


UDF

```
SELECT amount, OBFUSCATE(card_number) AS id  
FROM AUTHORIZATIONS;
```



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output

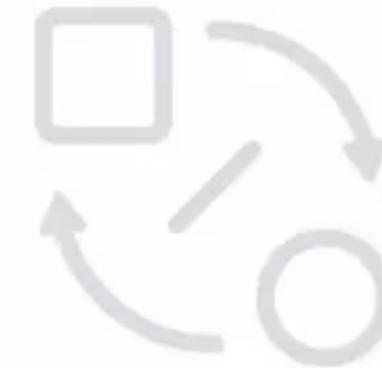


UDAF

```
SELECT card_number, SUM(cost) AS total_cost  
FROM AUTHORIZATIONS WINDOW TUMBLING (SIZE 7 DAYS)  
GROUP BY card_number;
```



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output



UDAF

```
SELECT card_number, SUM(cost) AS total_cost  
FROM AUTHORIZATIONS WINDOW TUMBLING (SIZE 7 DAYS)  
GROUP BY card_number;
```

(123, \$10)
(123, \$20)
(456, \$40)
(123, \$5)

S

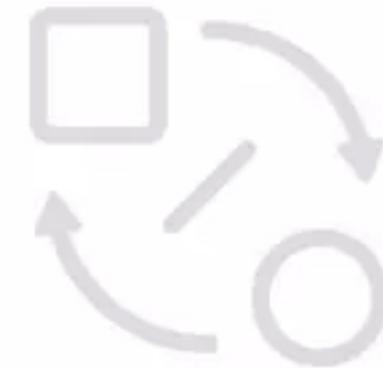


(123_w1, \$30)
(456_w1, \$40)
(123_w2, \$5)

T



Import Data



Create Stream/Table



Write KSQL Queries



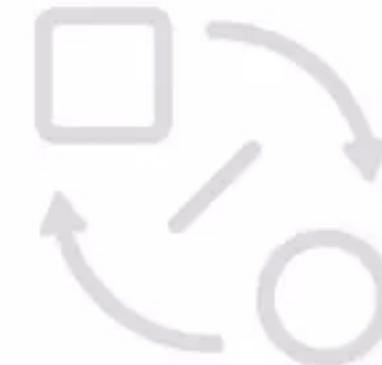
Verify Output

ksql-test-framework*

```
SELECT * FROM stream LIMIT 3
```



Import Data



Create Stream/Table



Write KSQL Queries



Verify Output

ksql-test-framework*

```
{  
  "statements": [  
    "CREATE STREAM TEST (source VARCHAR) WITH (kafka_topic='test_topic');",  
    "CREATE STREAM O AS SELECT CONCAT('prefix-', source) AS C FROM TEST;"],  
  "inputs": [  
    {"topic": "test_topic", "value": {"source": "s1"}},  
    {"topic": "test_topic", "value": {"source": "s2"}]},  
  "outputs": [  
    {"topic": "O", "value": {"C": "prefix-s1"}},  
    {"topic": "O", "value": {"C": "prefix-s2"}]}  
}
```

The screenshot shows a macOS terminal window titled "Terminalizer". The window has three colored status icons (red, yellow, green) in the top-left corner. The terminal output is as follows:

```
=  
=   [ \ / \ / \ / \ / \ ] =  
=   [ ' / | \ \ \ \ \ ] =  
=   [ < \ \ \ \ \ \ \ \ ] =  
=   [ . \ \ \ \ \ \ \ \ ] =  
=   [ \ \ \ \ \ \ \ \ \ ] =  
=   [ \ \ \ \ \ \ \ \ \ ] =  
= Streaming SQL Engine for Apache Kafka® =  
=====
```

Copyright 2017-2018 Confluent Inc.
CLI v5.3.0-SNAPSHOT, Server v5.3.0-SNAPSHOT located at <http://localhost:8088>
Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

* Feature Coming in CP 5.3

- Intro
- Develop
- **Deploy**
- Operate
- Common Mistakes
- Q&A



Deployment Options

Deployment Options



Interactive Mode

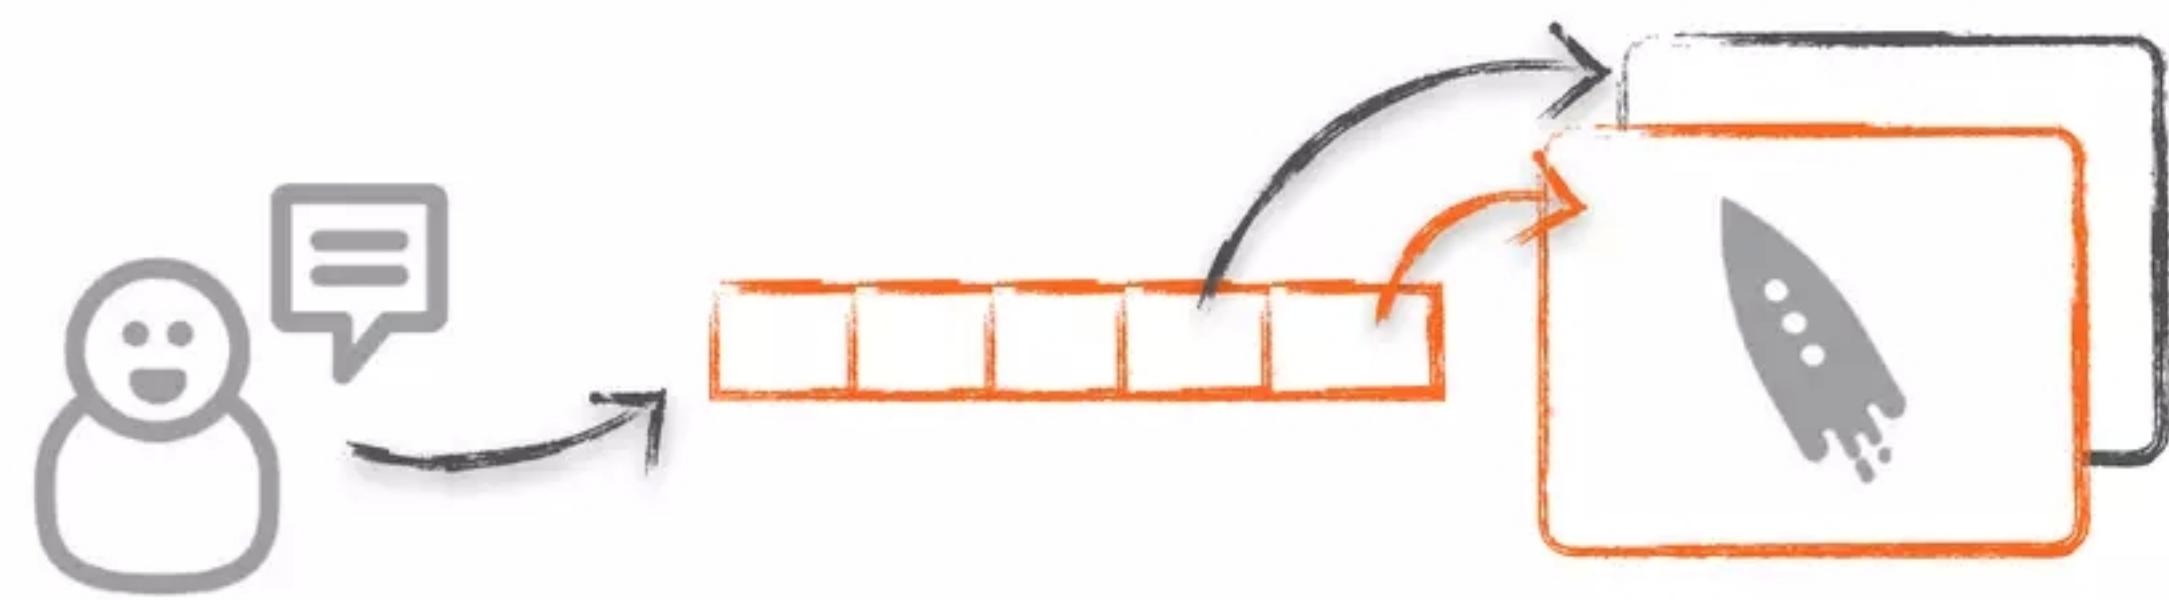
Deployment Options



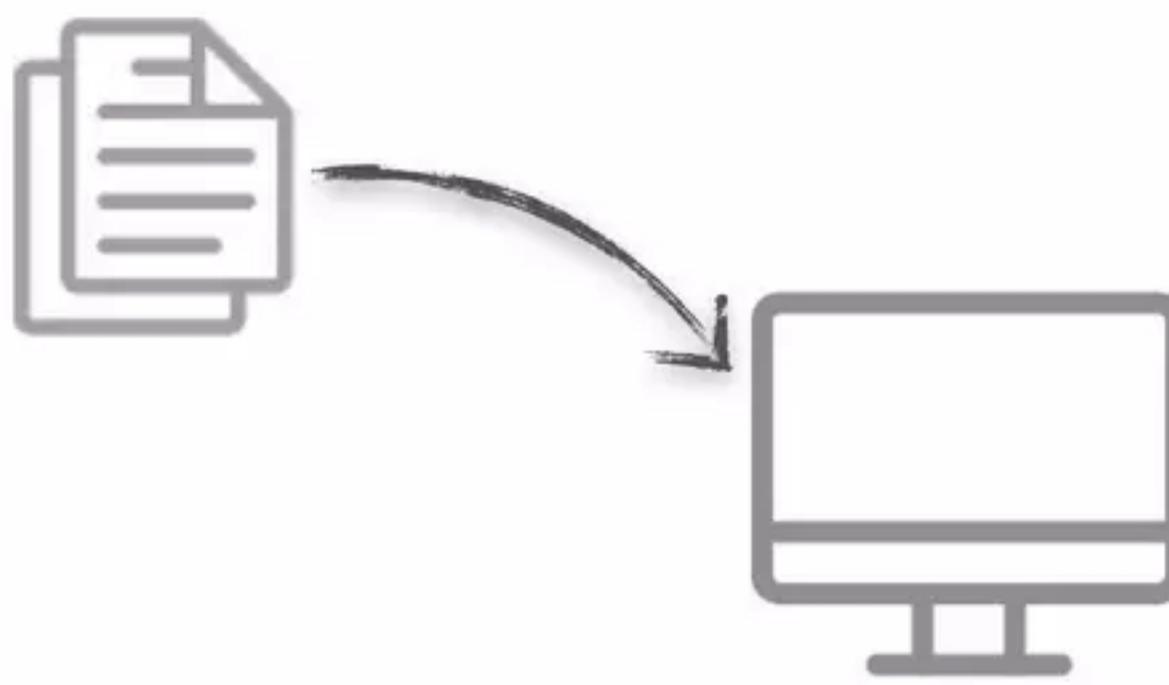
Interactive Mode



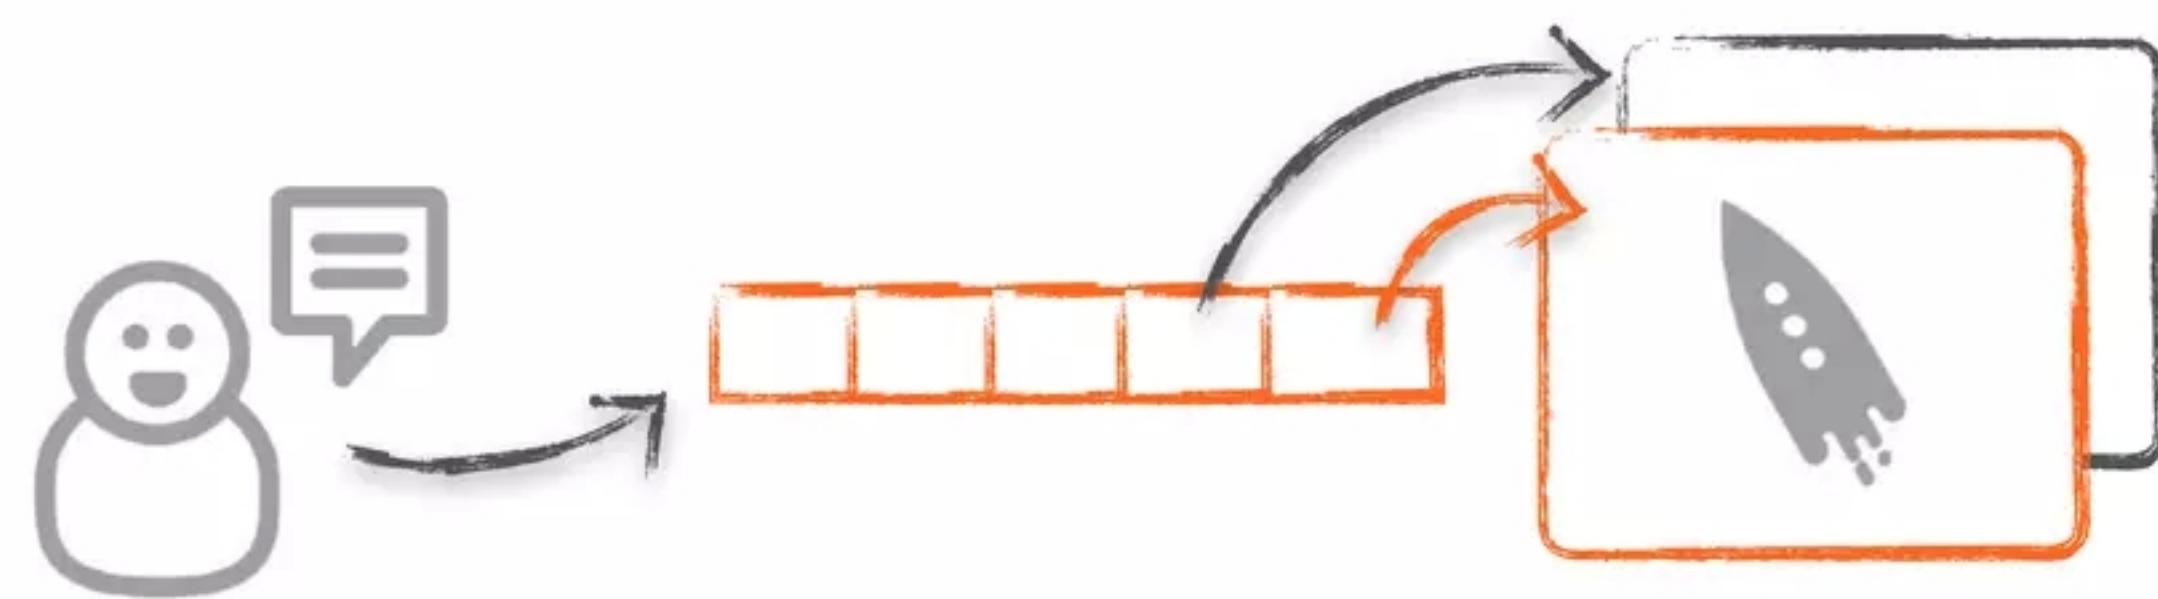
Headless Mode



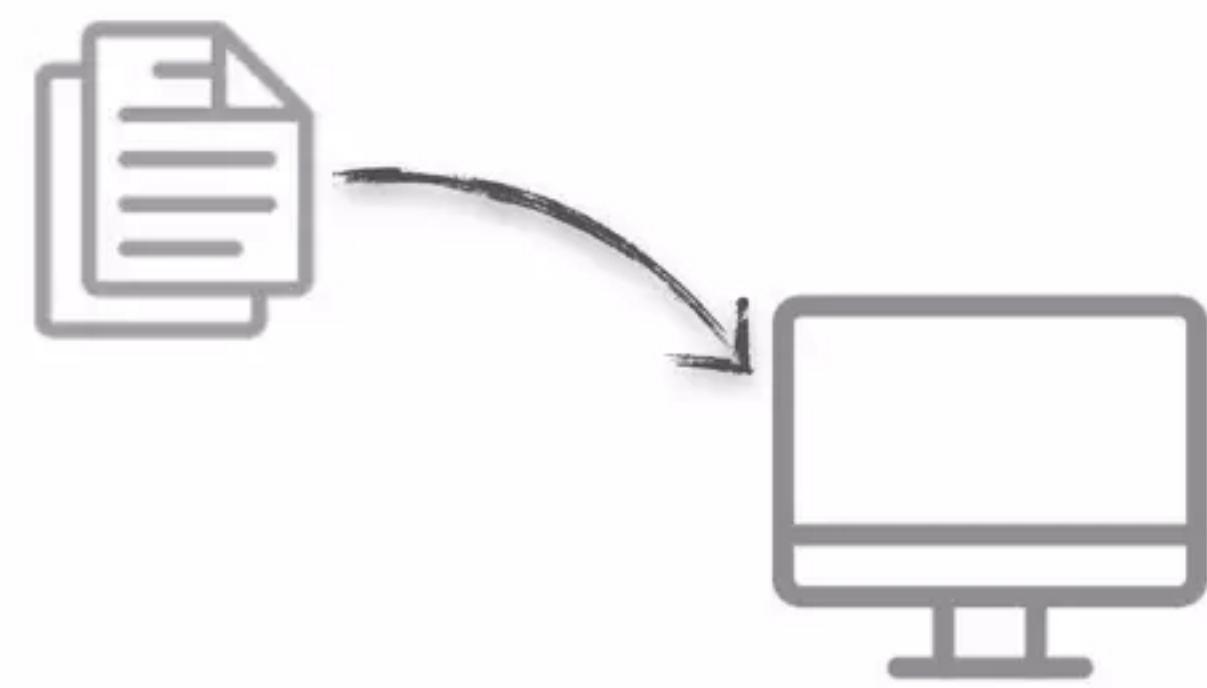
Interactive Mode



Headless Mode



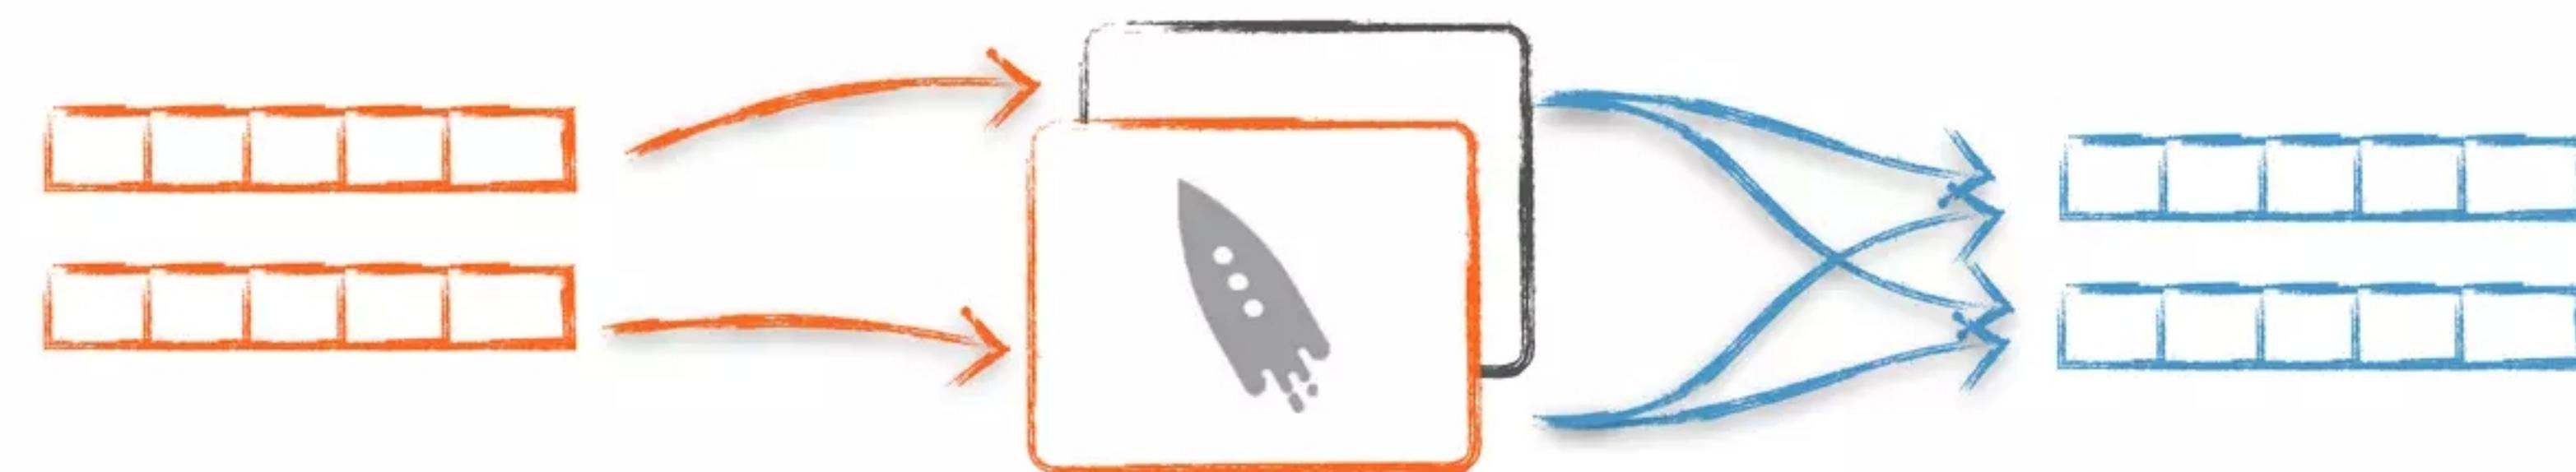
Interactive Mode



Headless Mode



KSQL scales relative
to partitions in Kafka



KSQL scales relative
to partitions in Kafka



... and rebalances if a
replica is lost

KSQL scales relative
to partitions in Kafka



✓ **tip:** initially over partitioning
your input topics can help
scale!

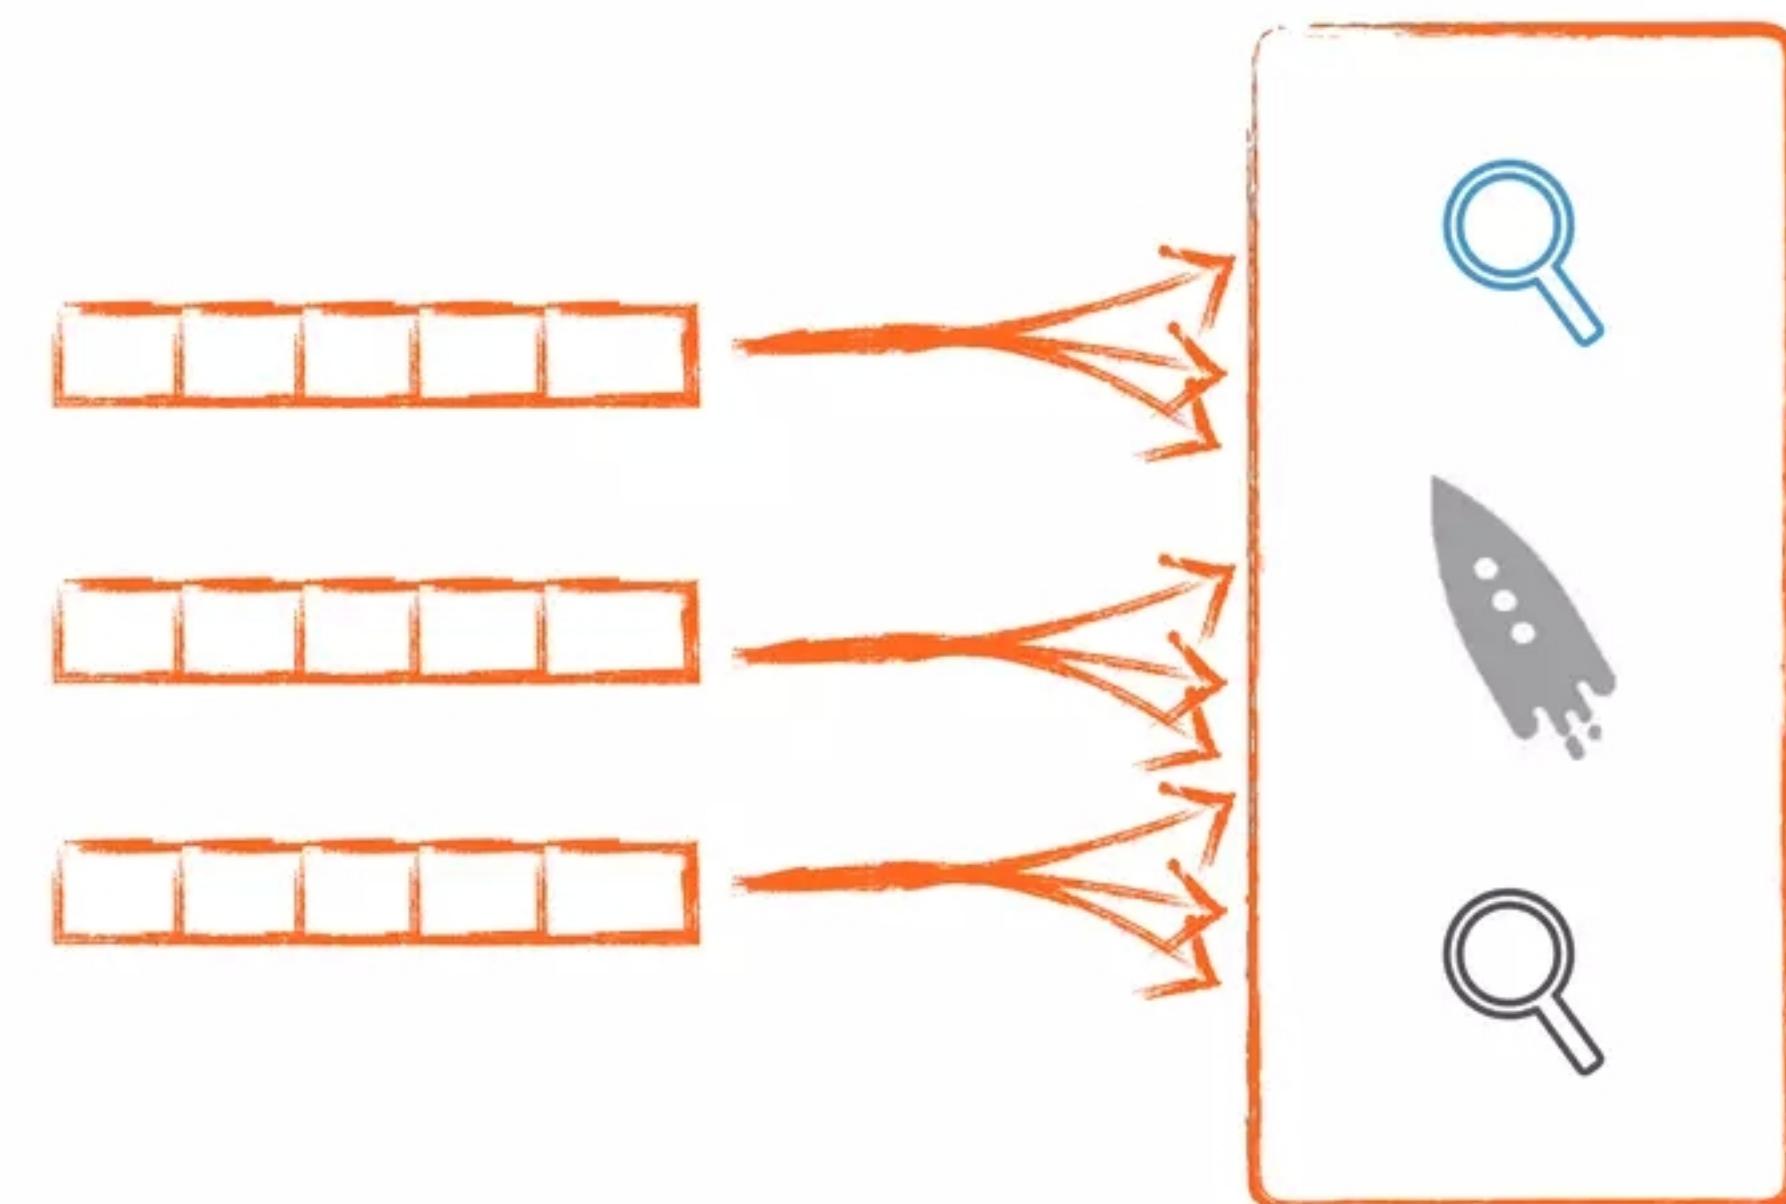
... and rebalances if a
replica is lost

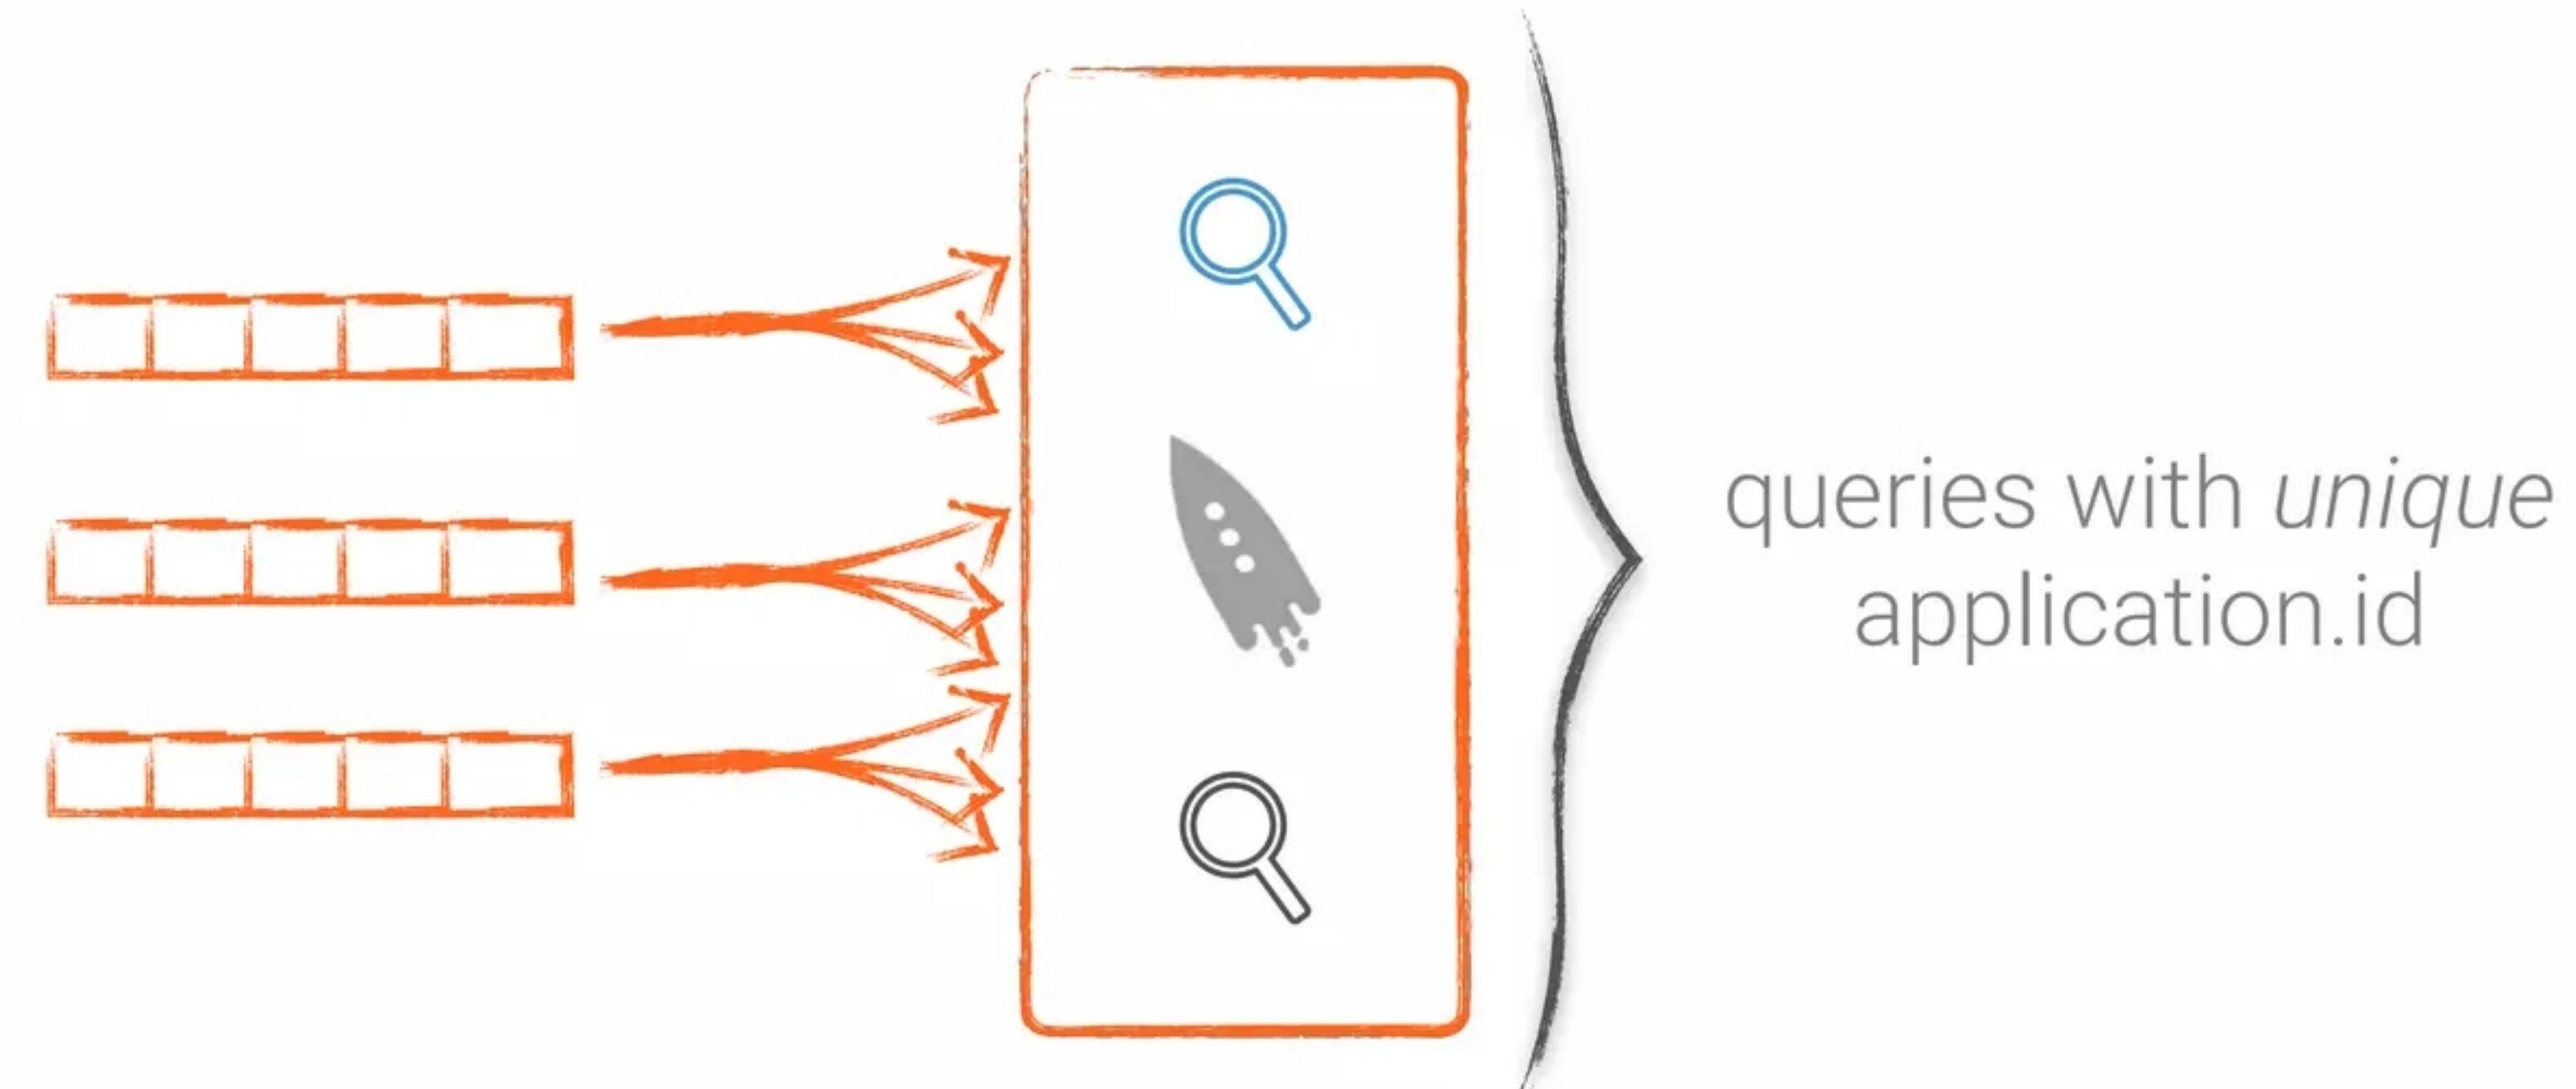


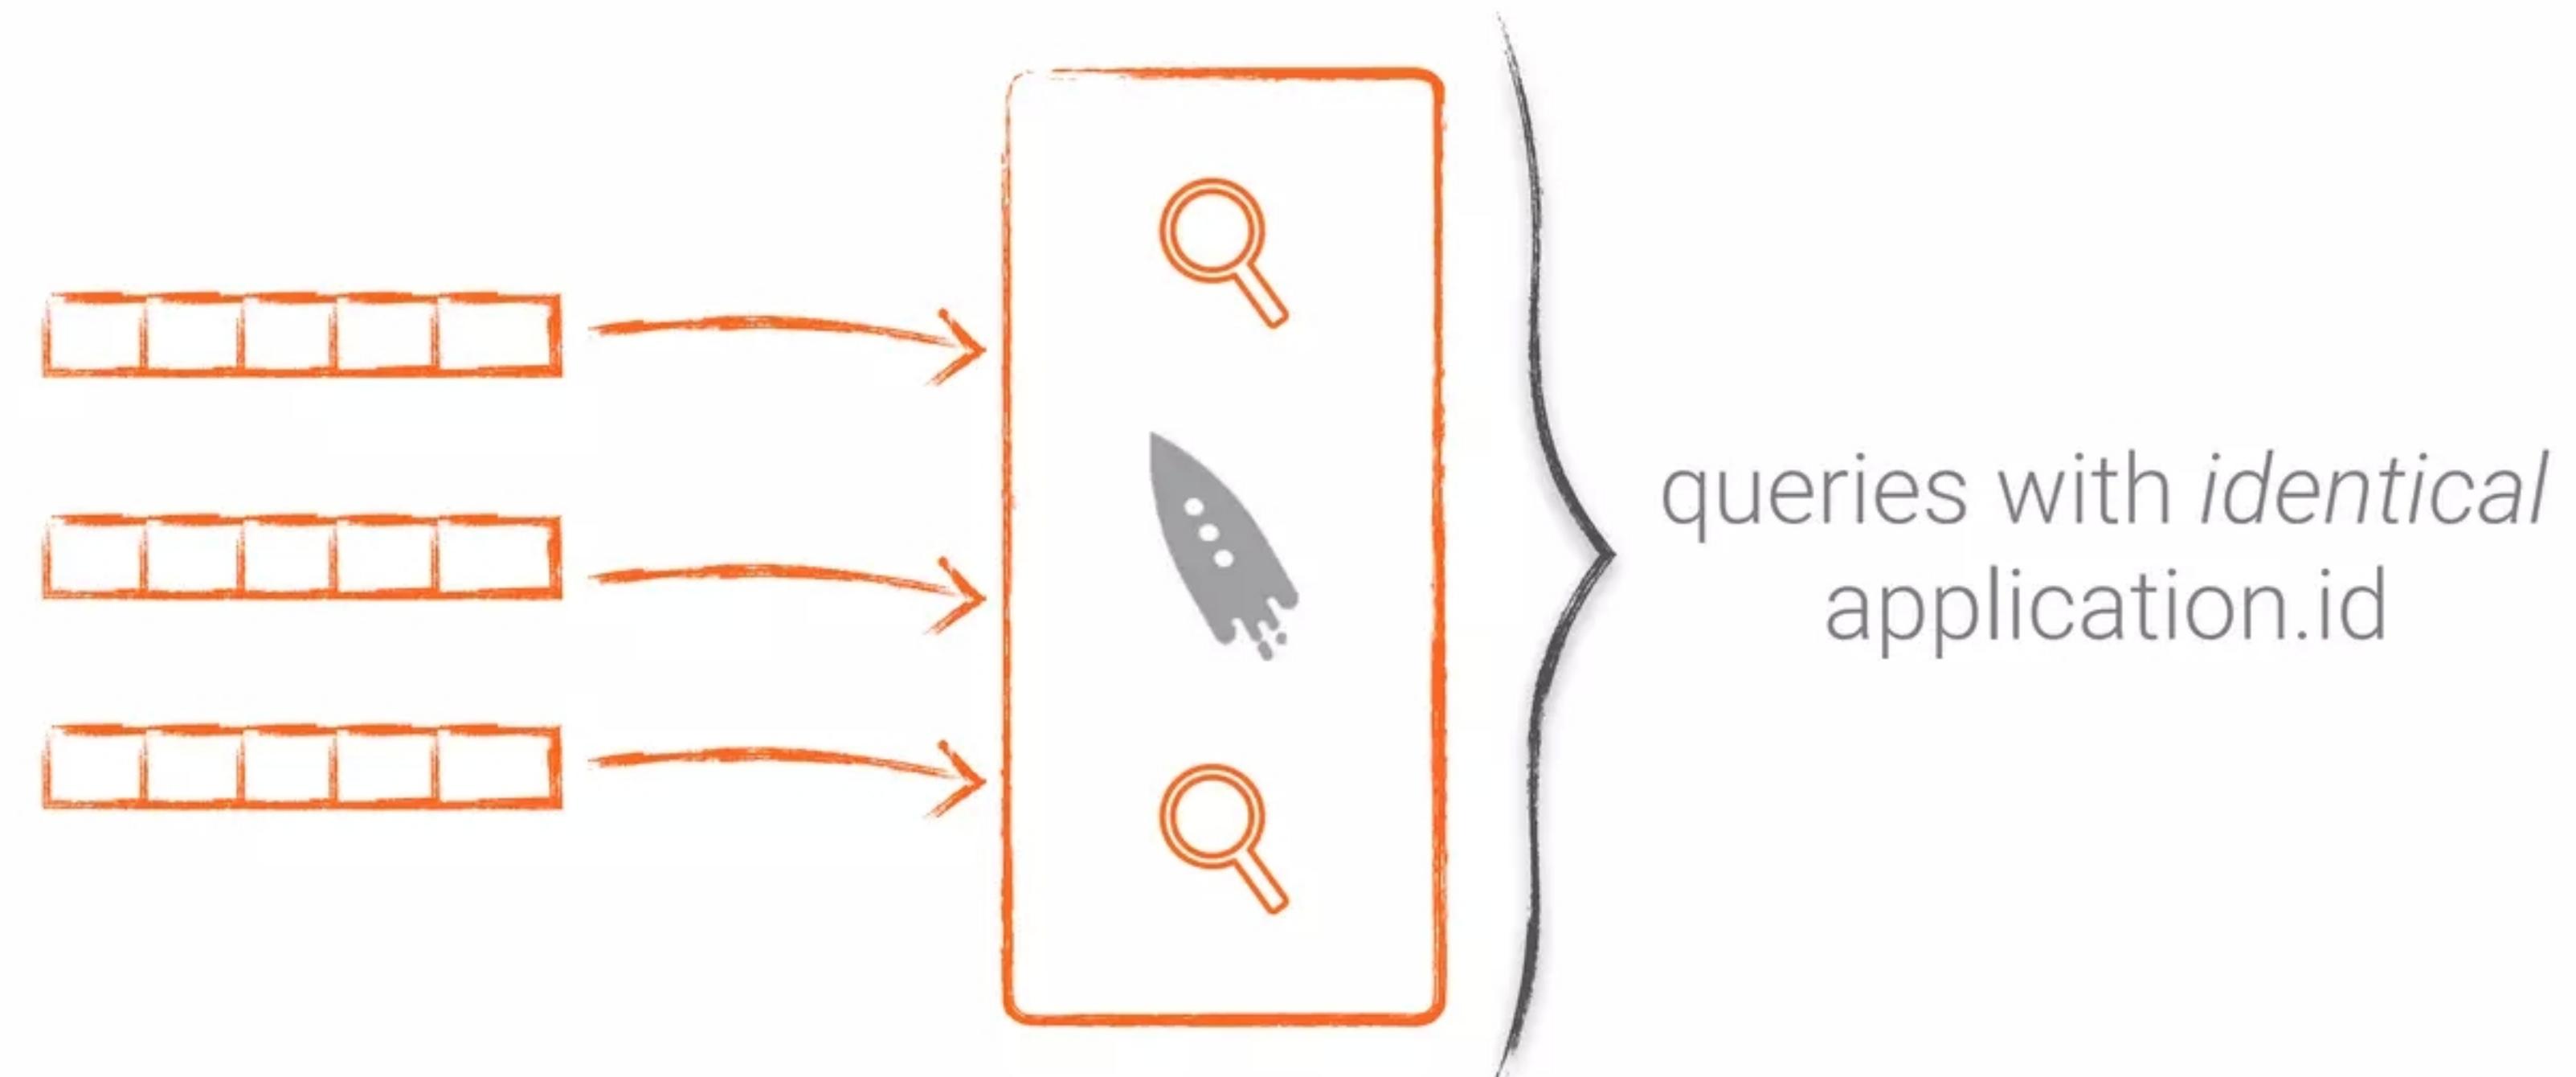
application.id

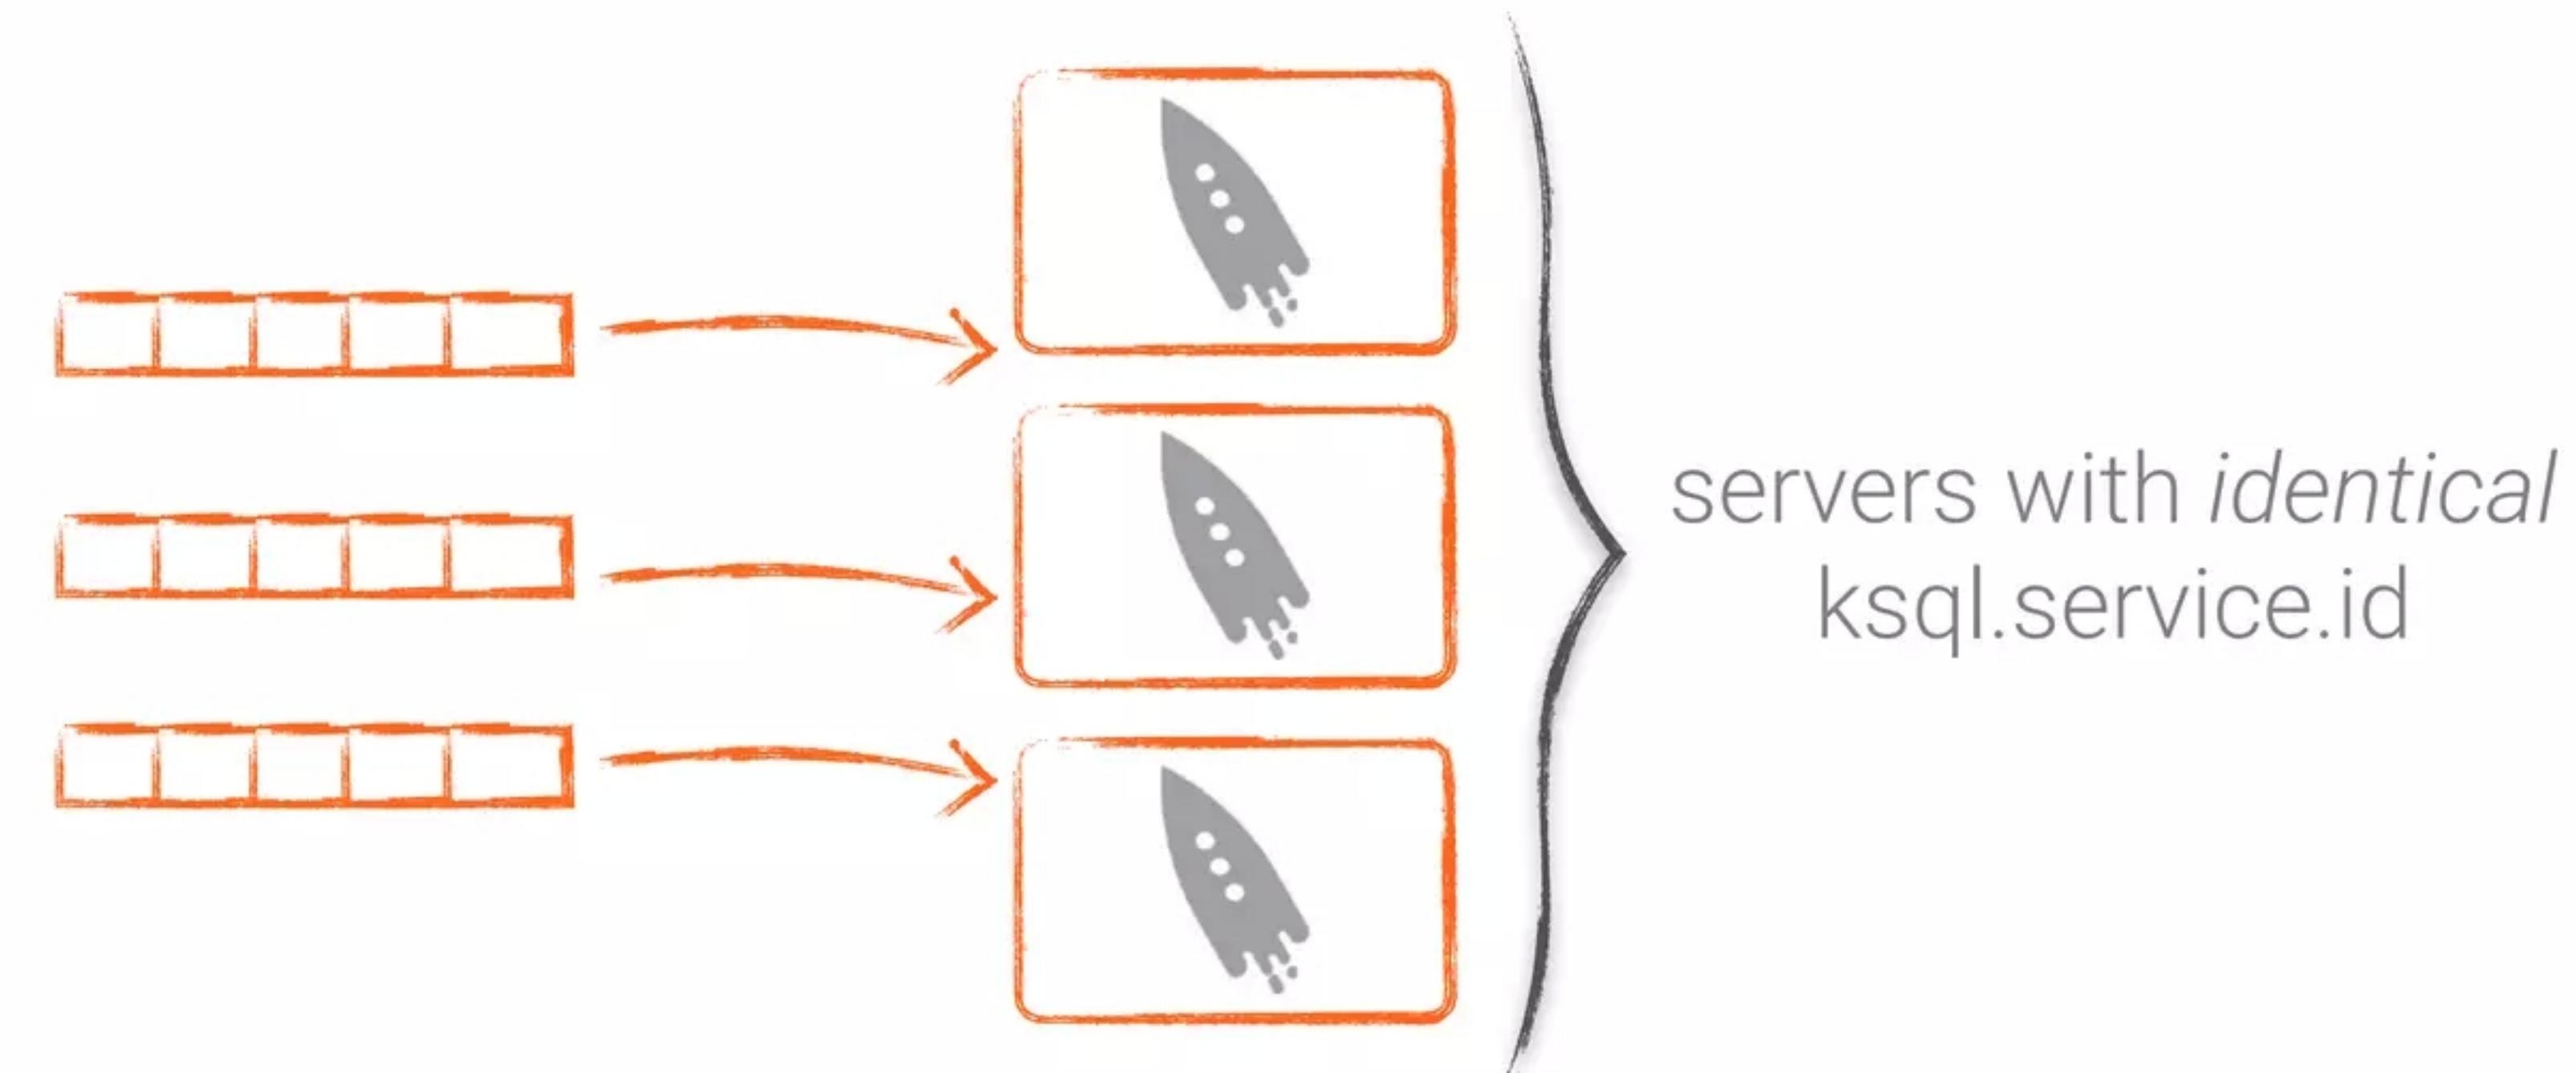


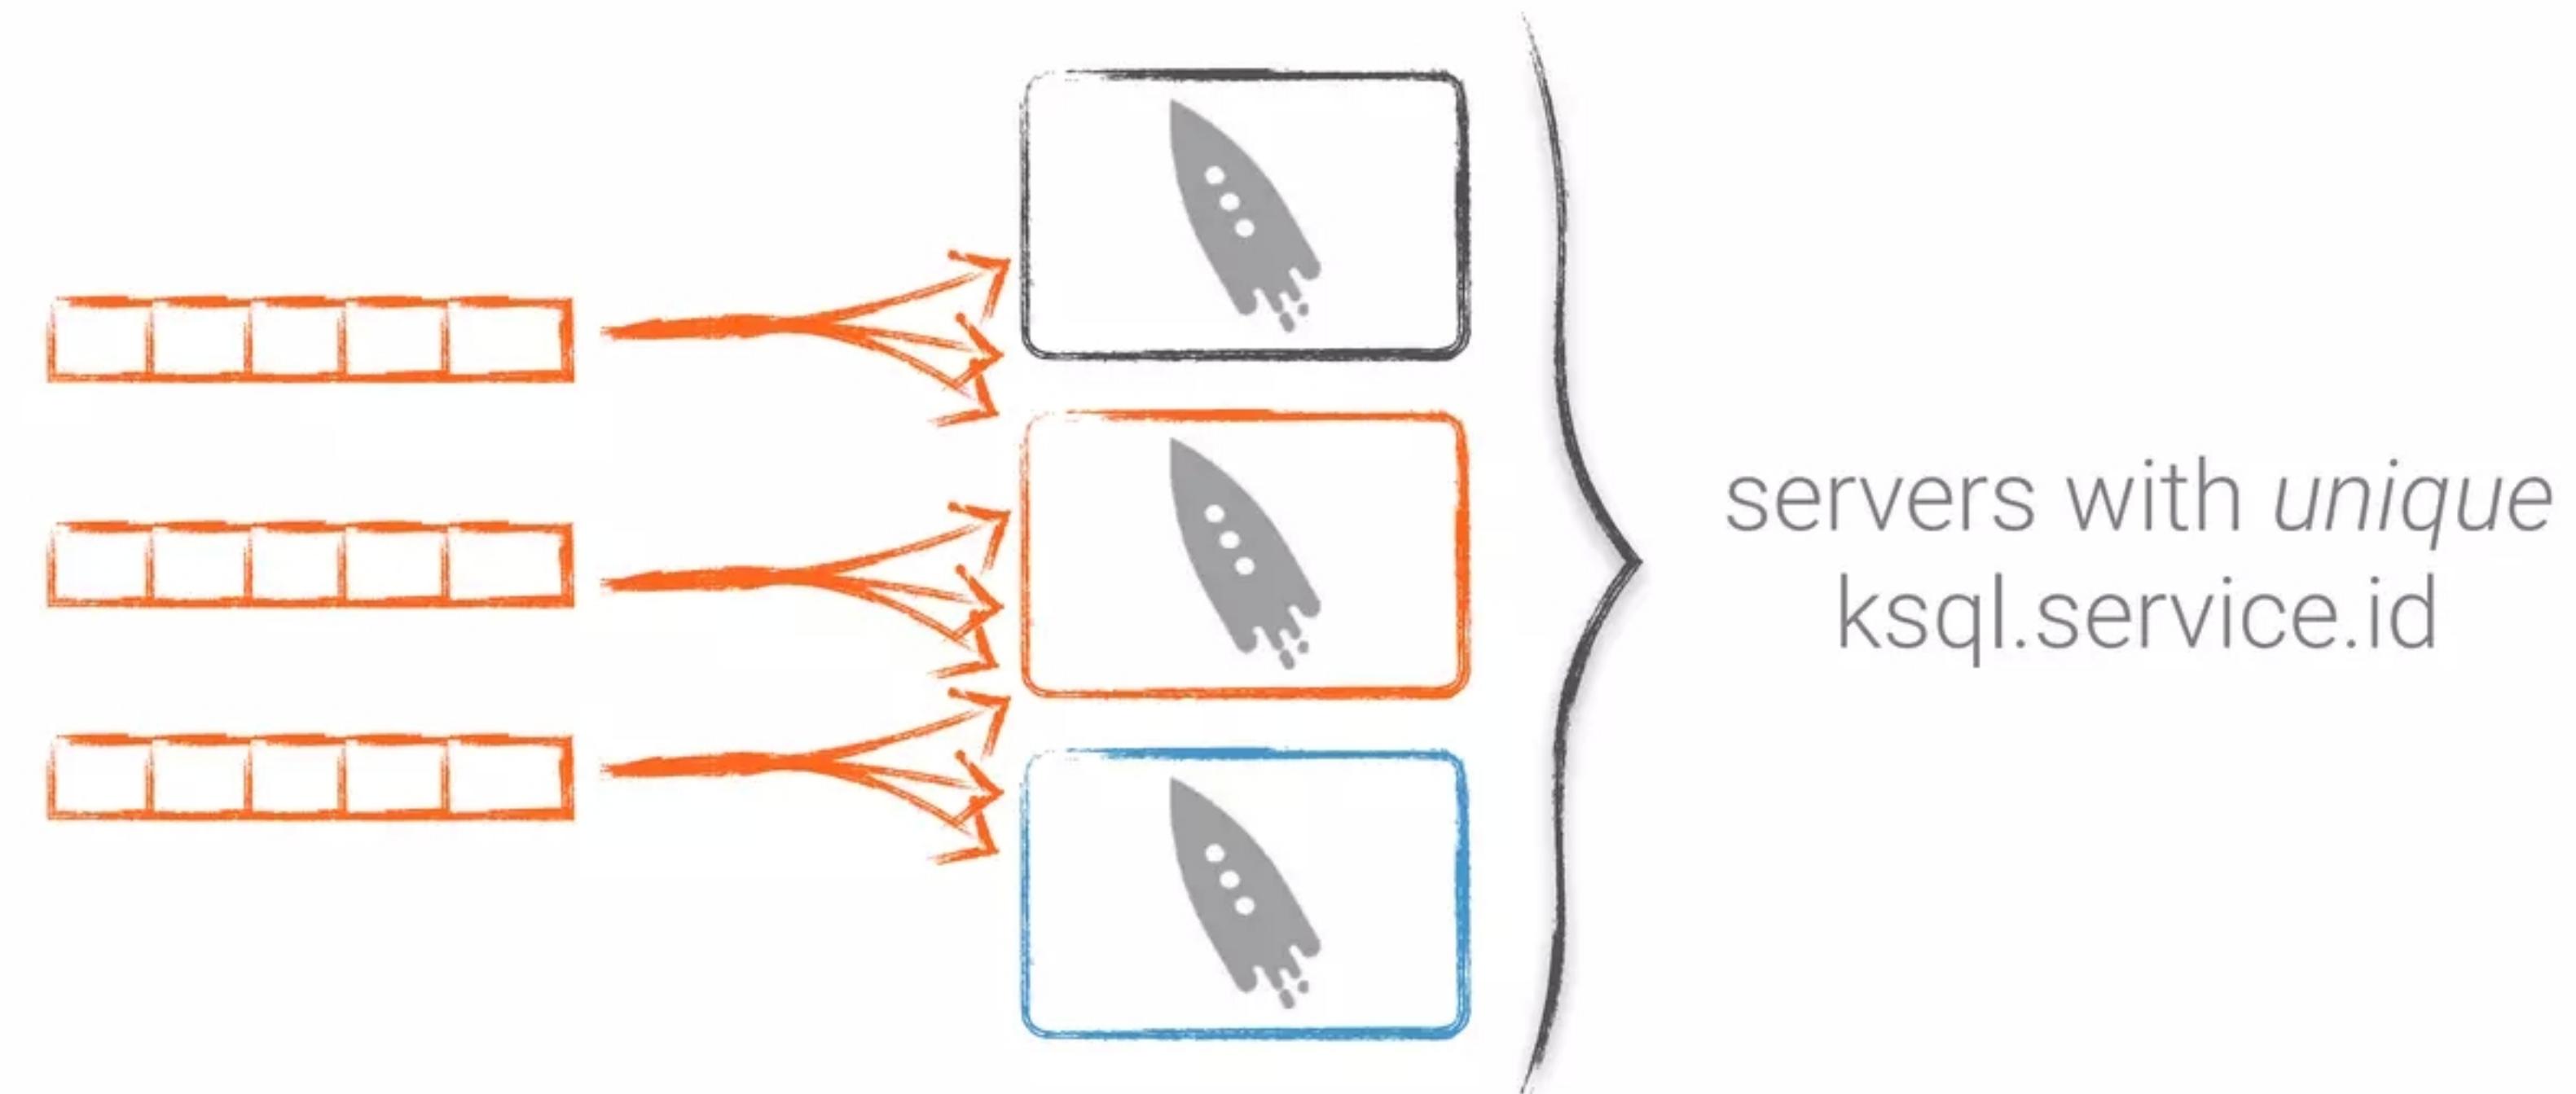
ksql.service.id



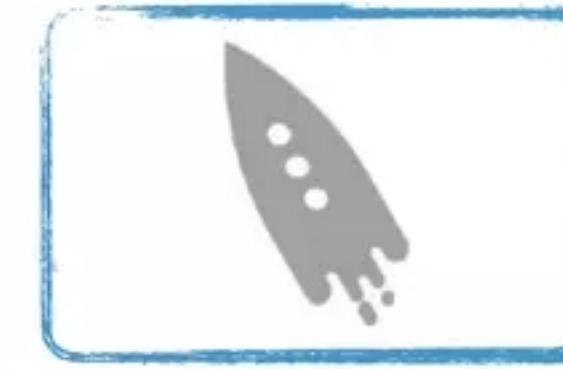




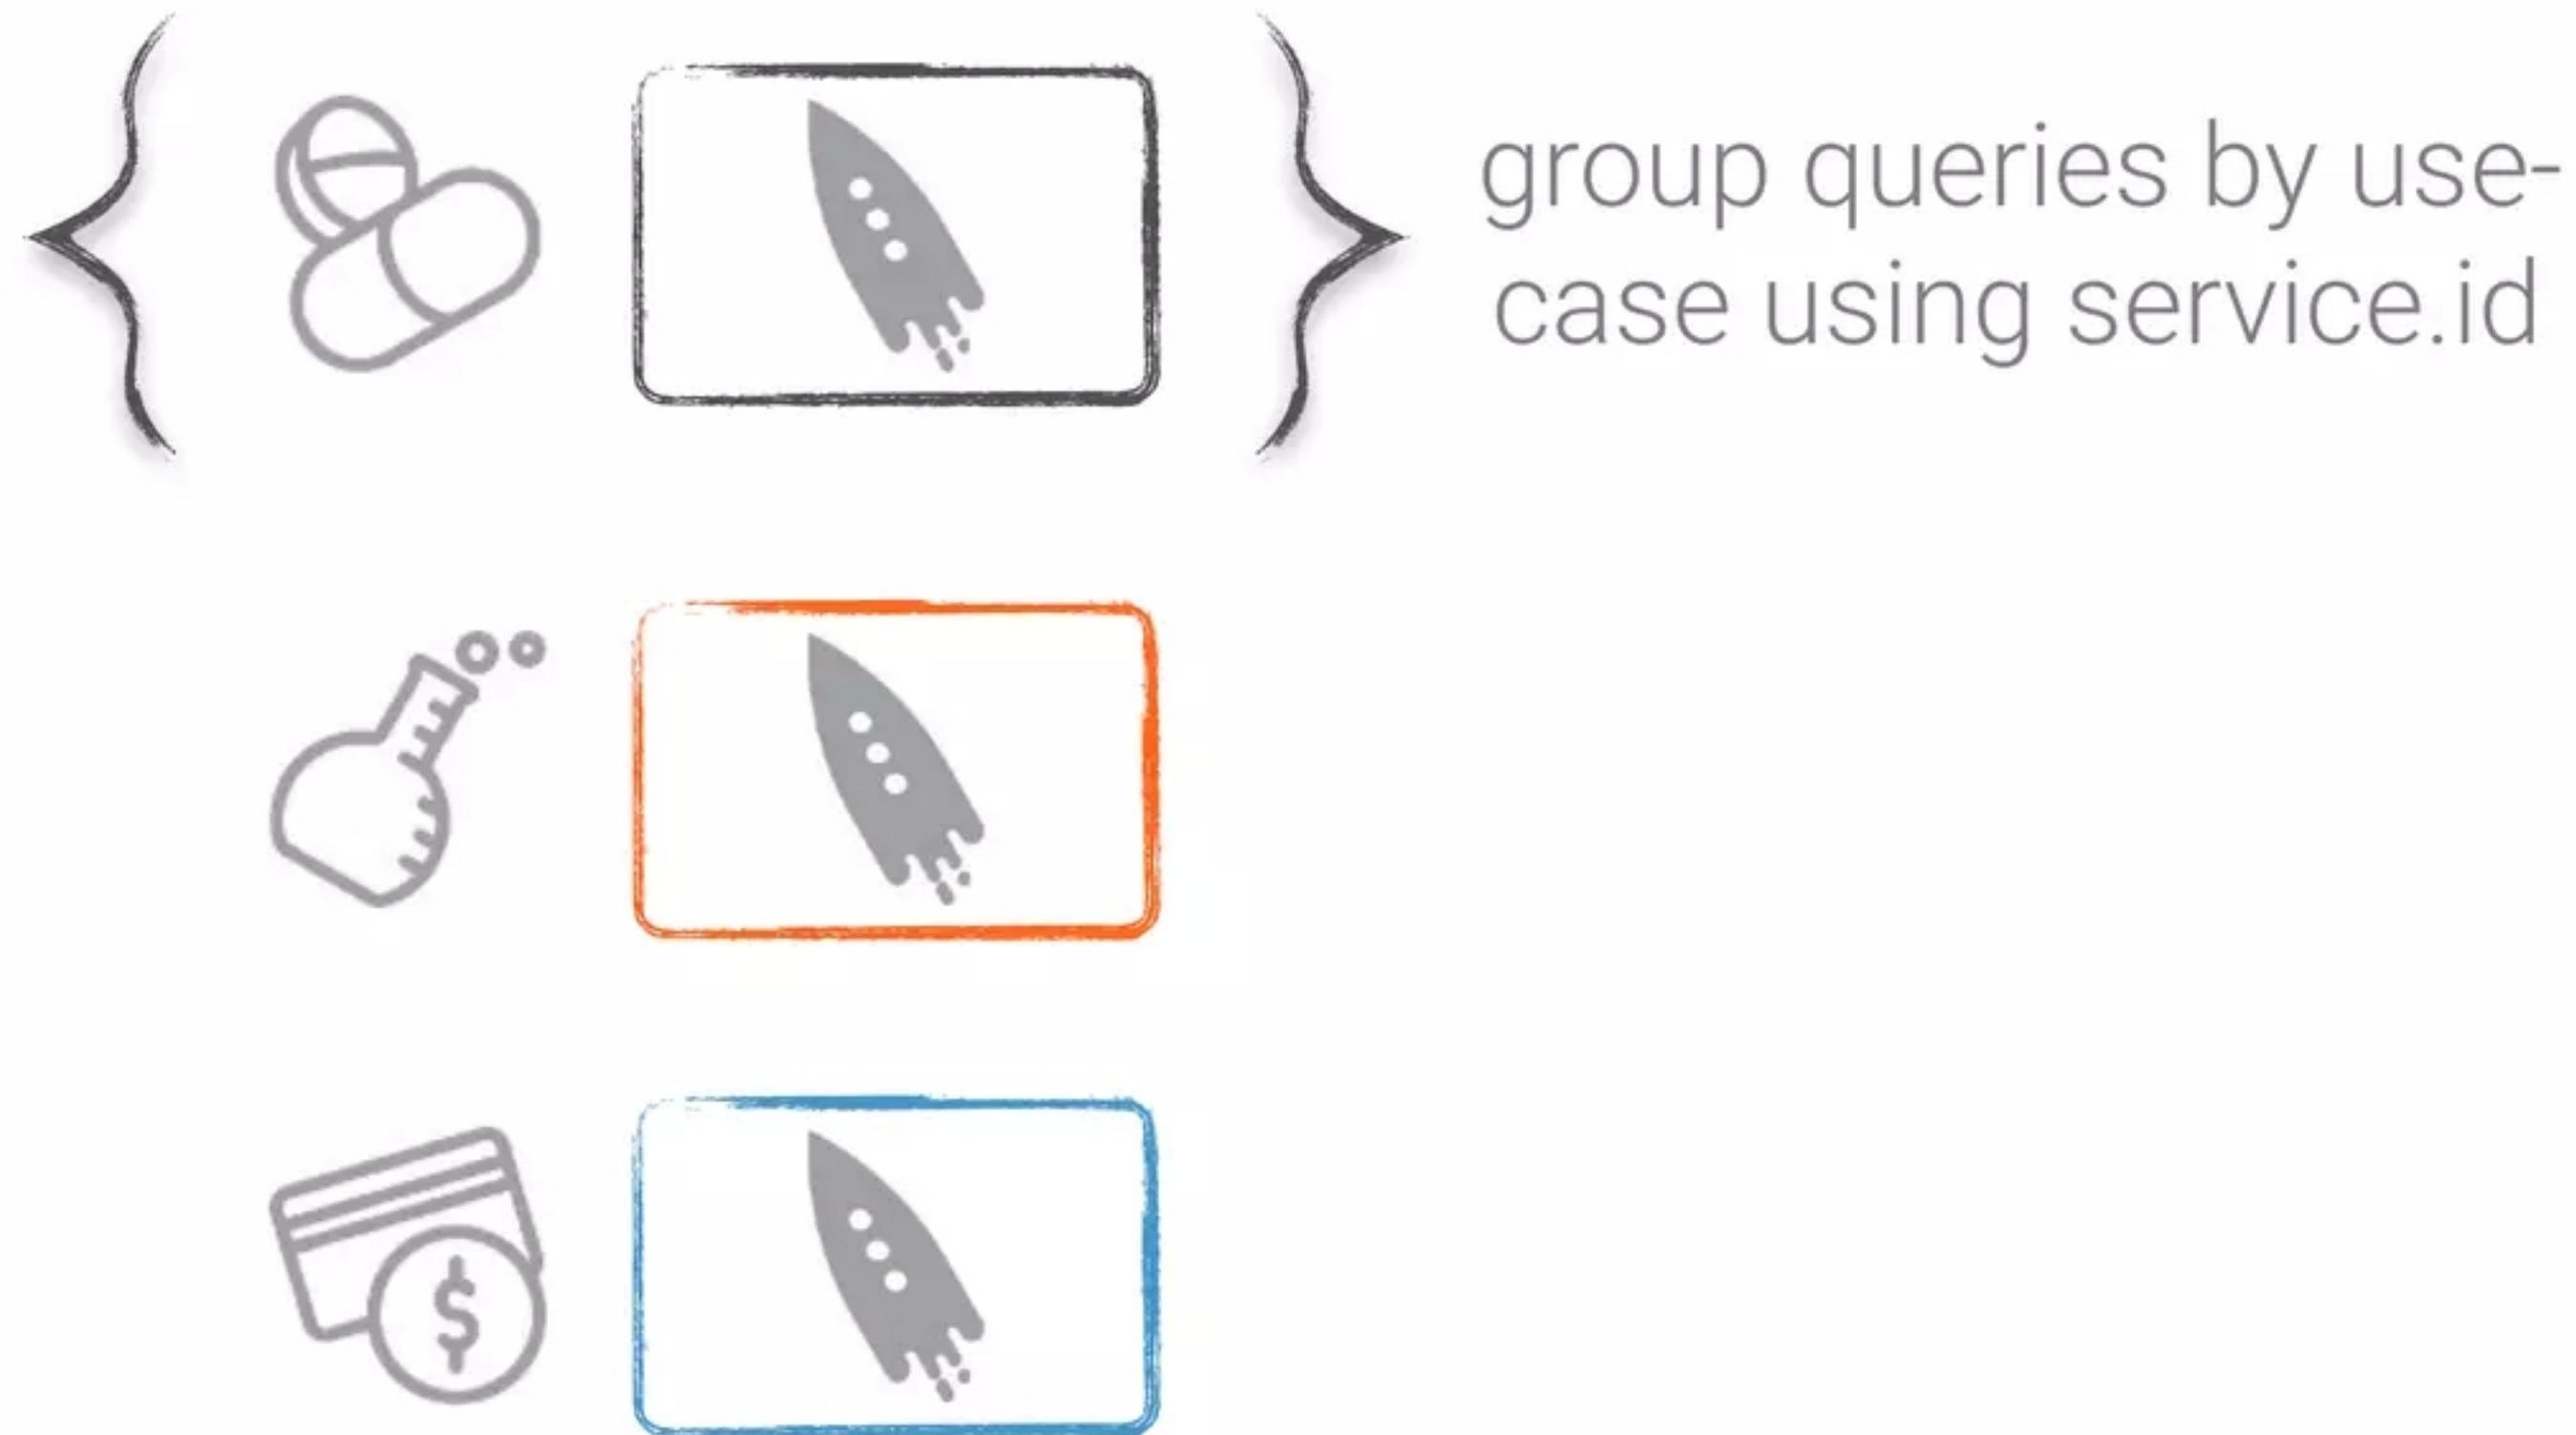




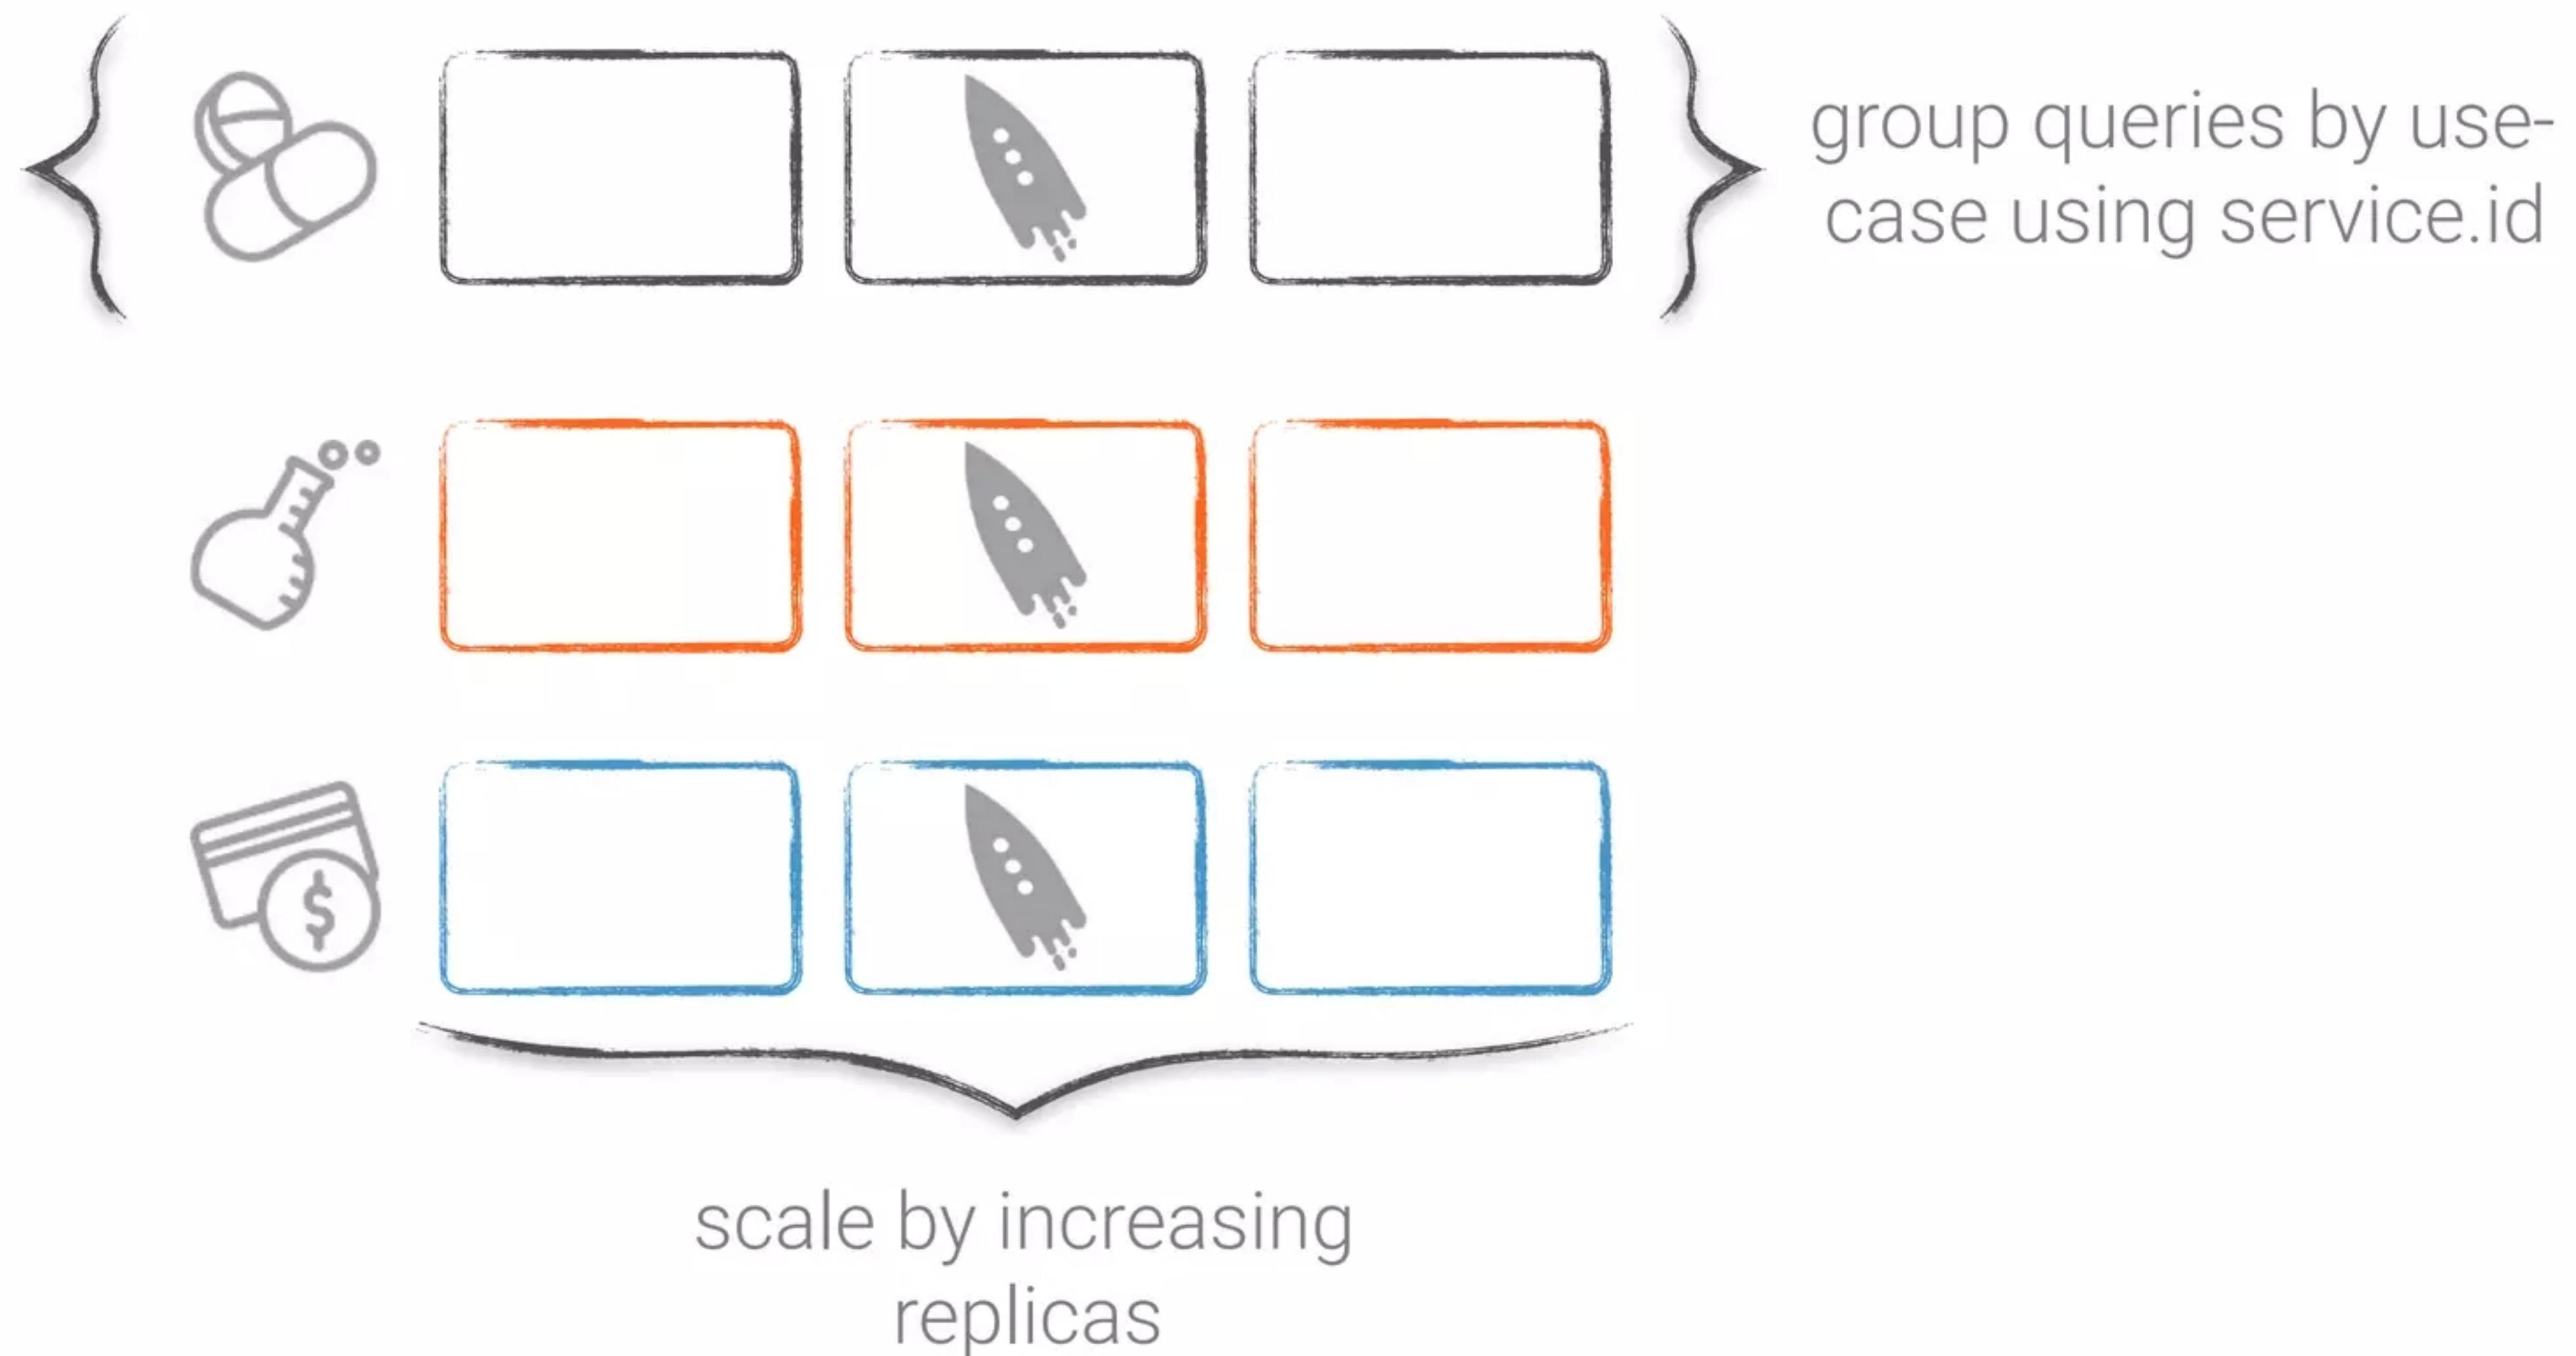
Recommended Deployment Topology



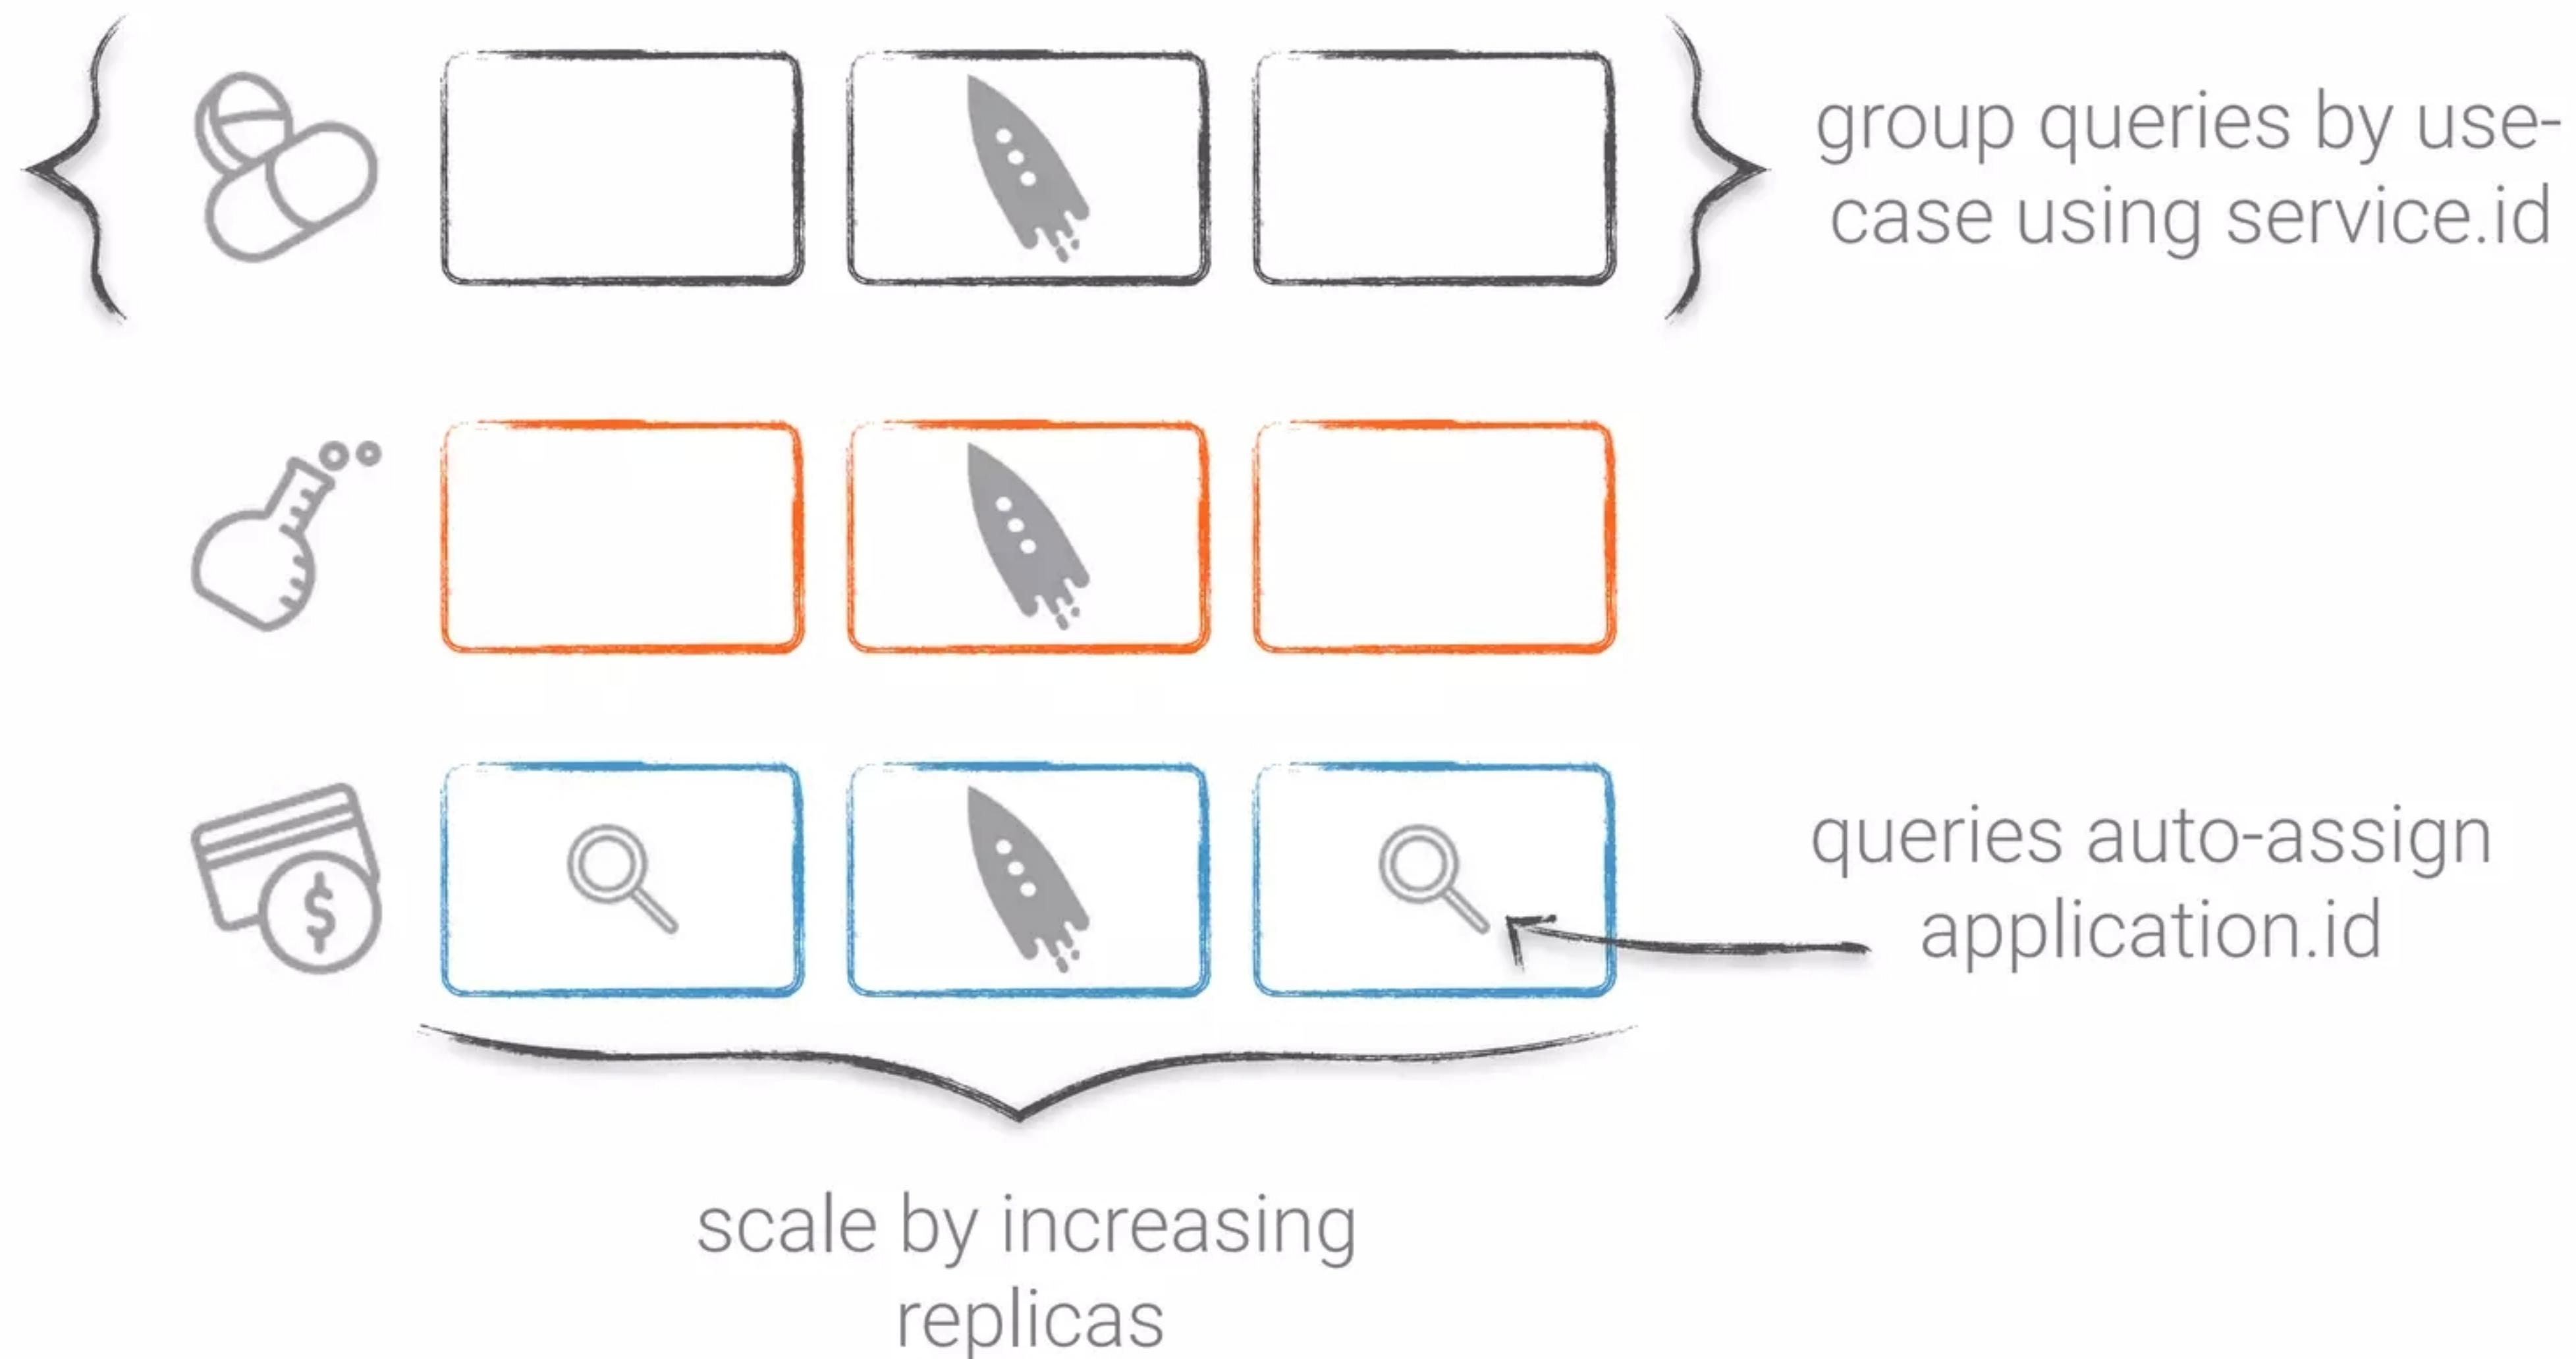
Recommended Deployment Topology



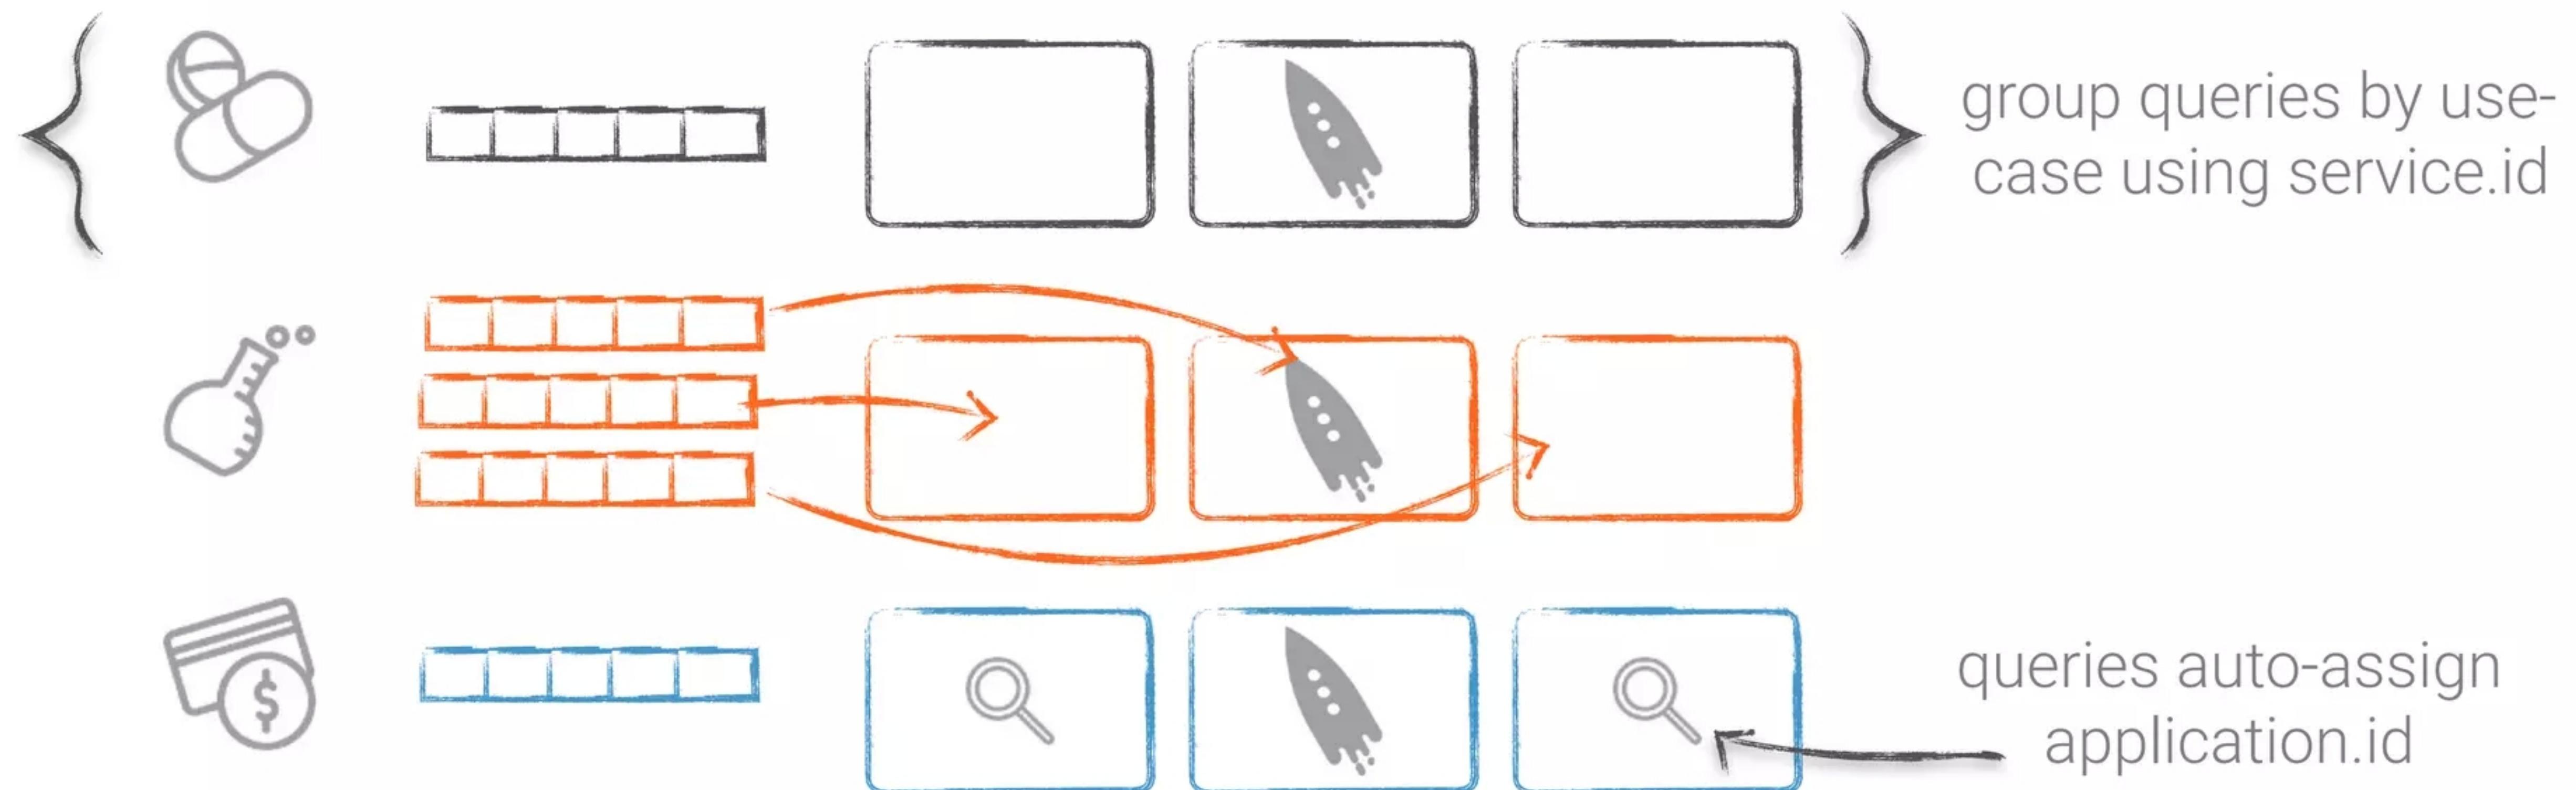
Recommended Deployment Topology



Recommended Deployment Topology

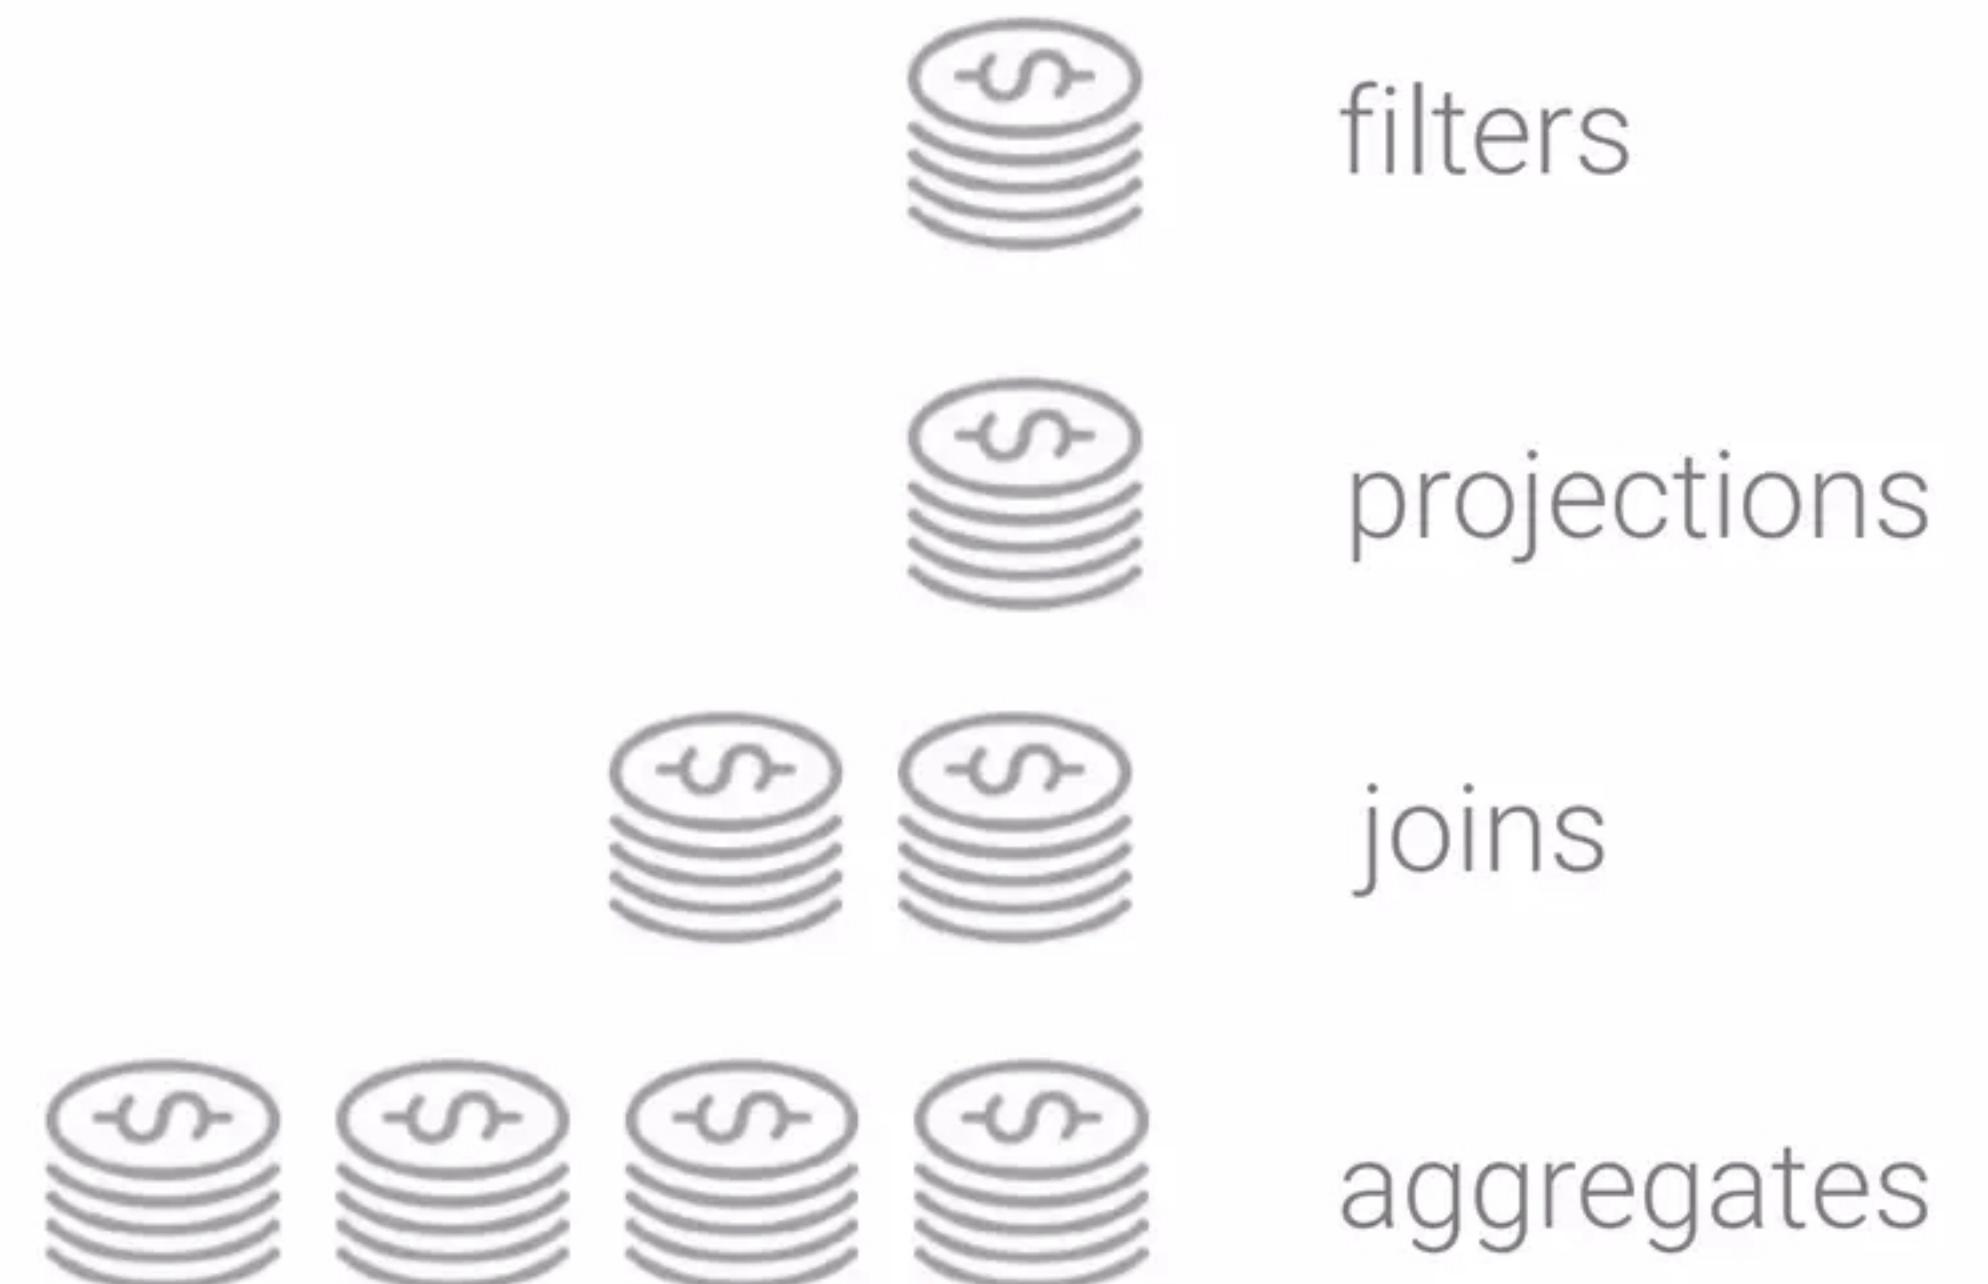


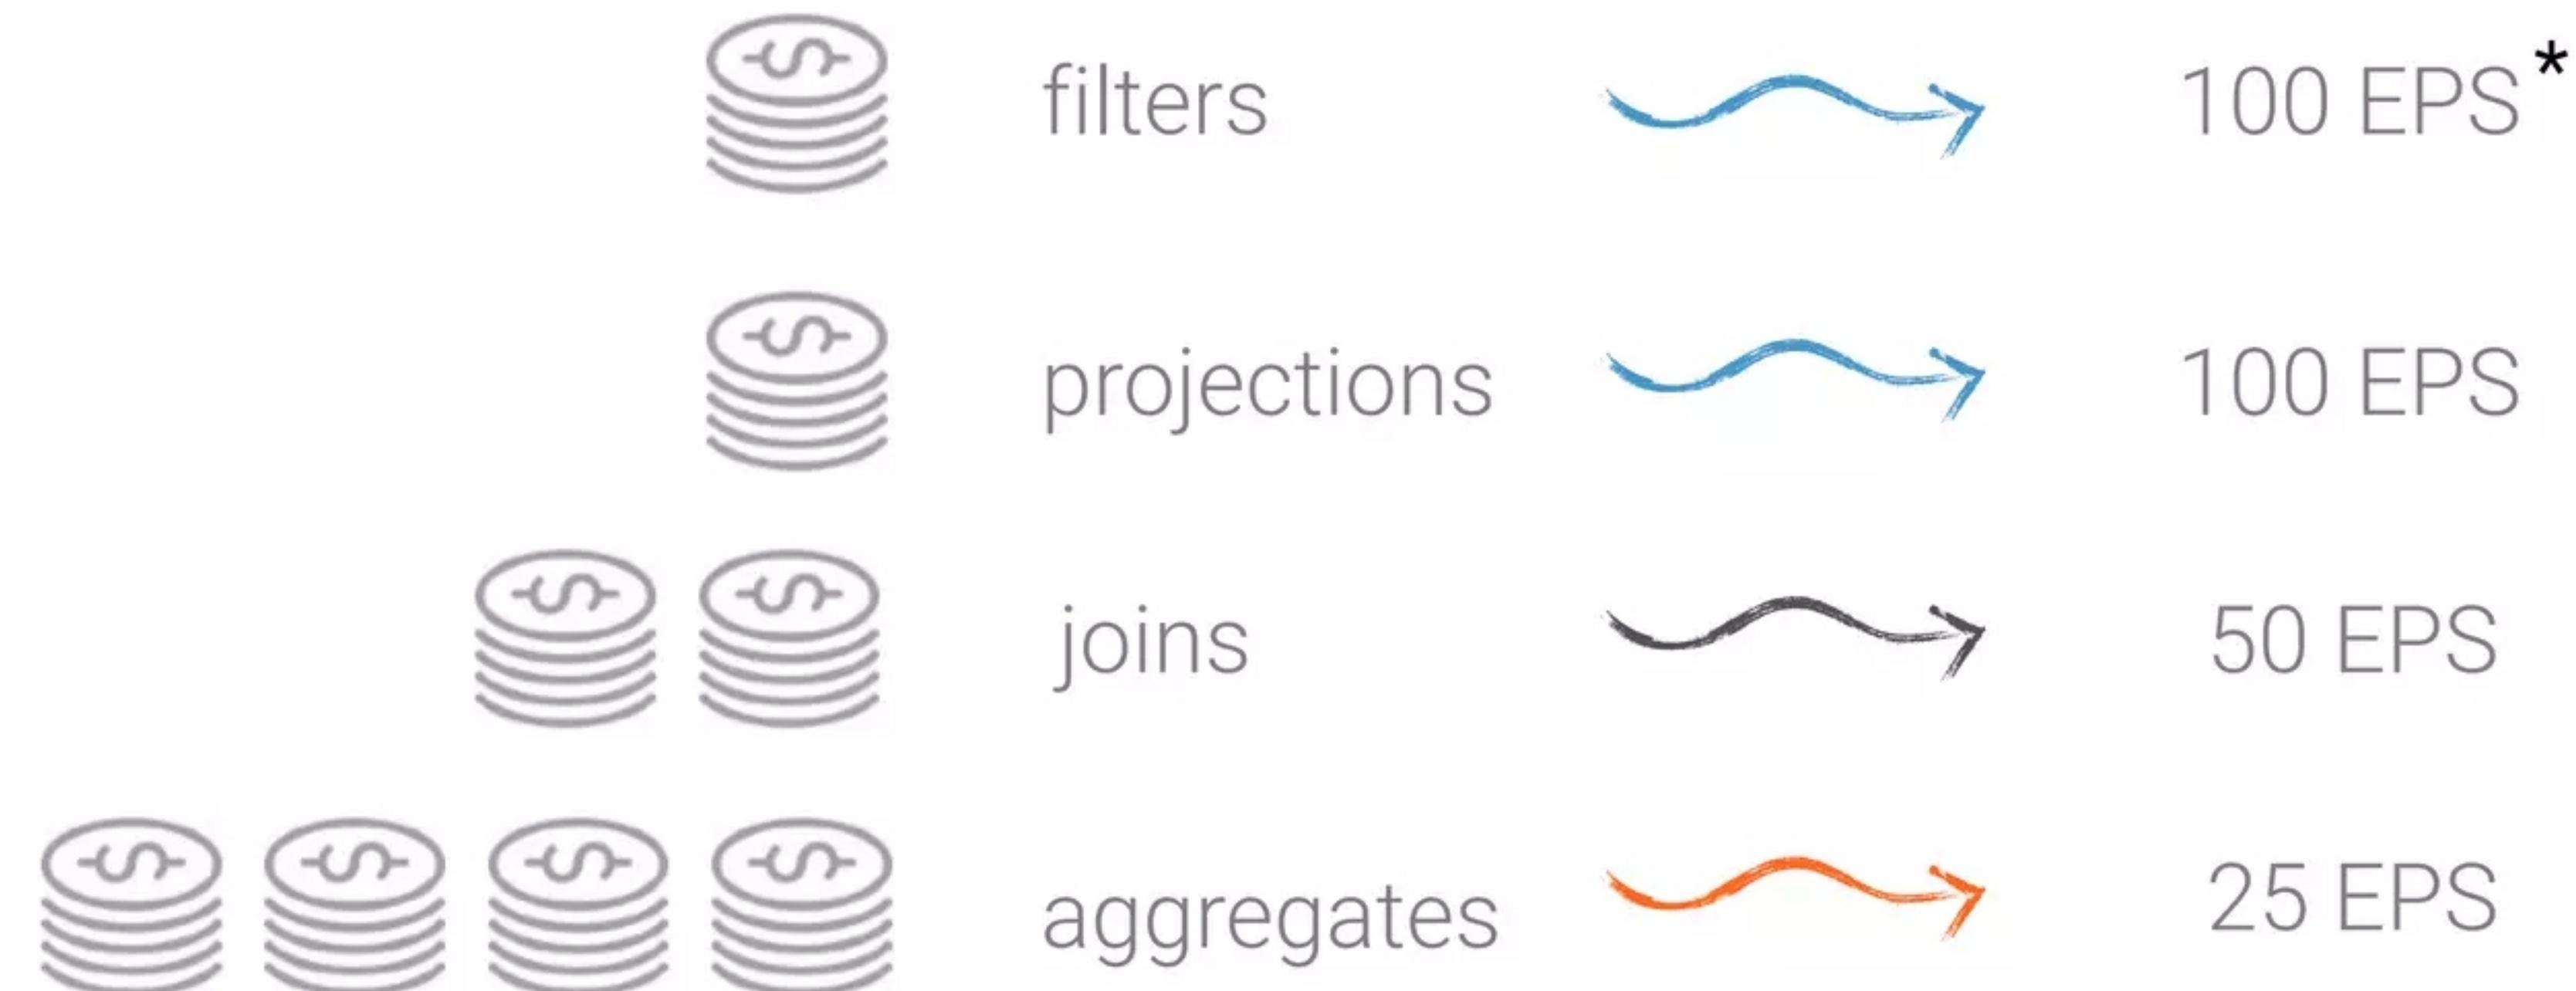
Recommended Deployment Topology

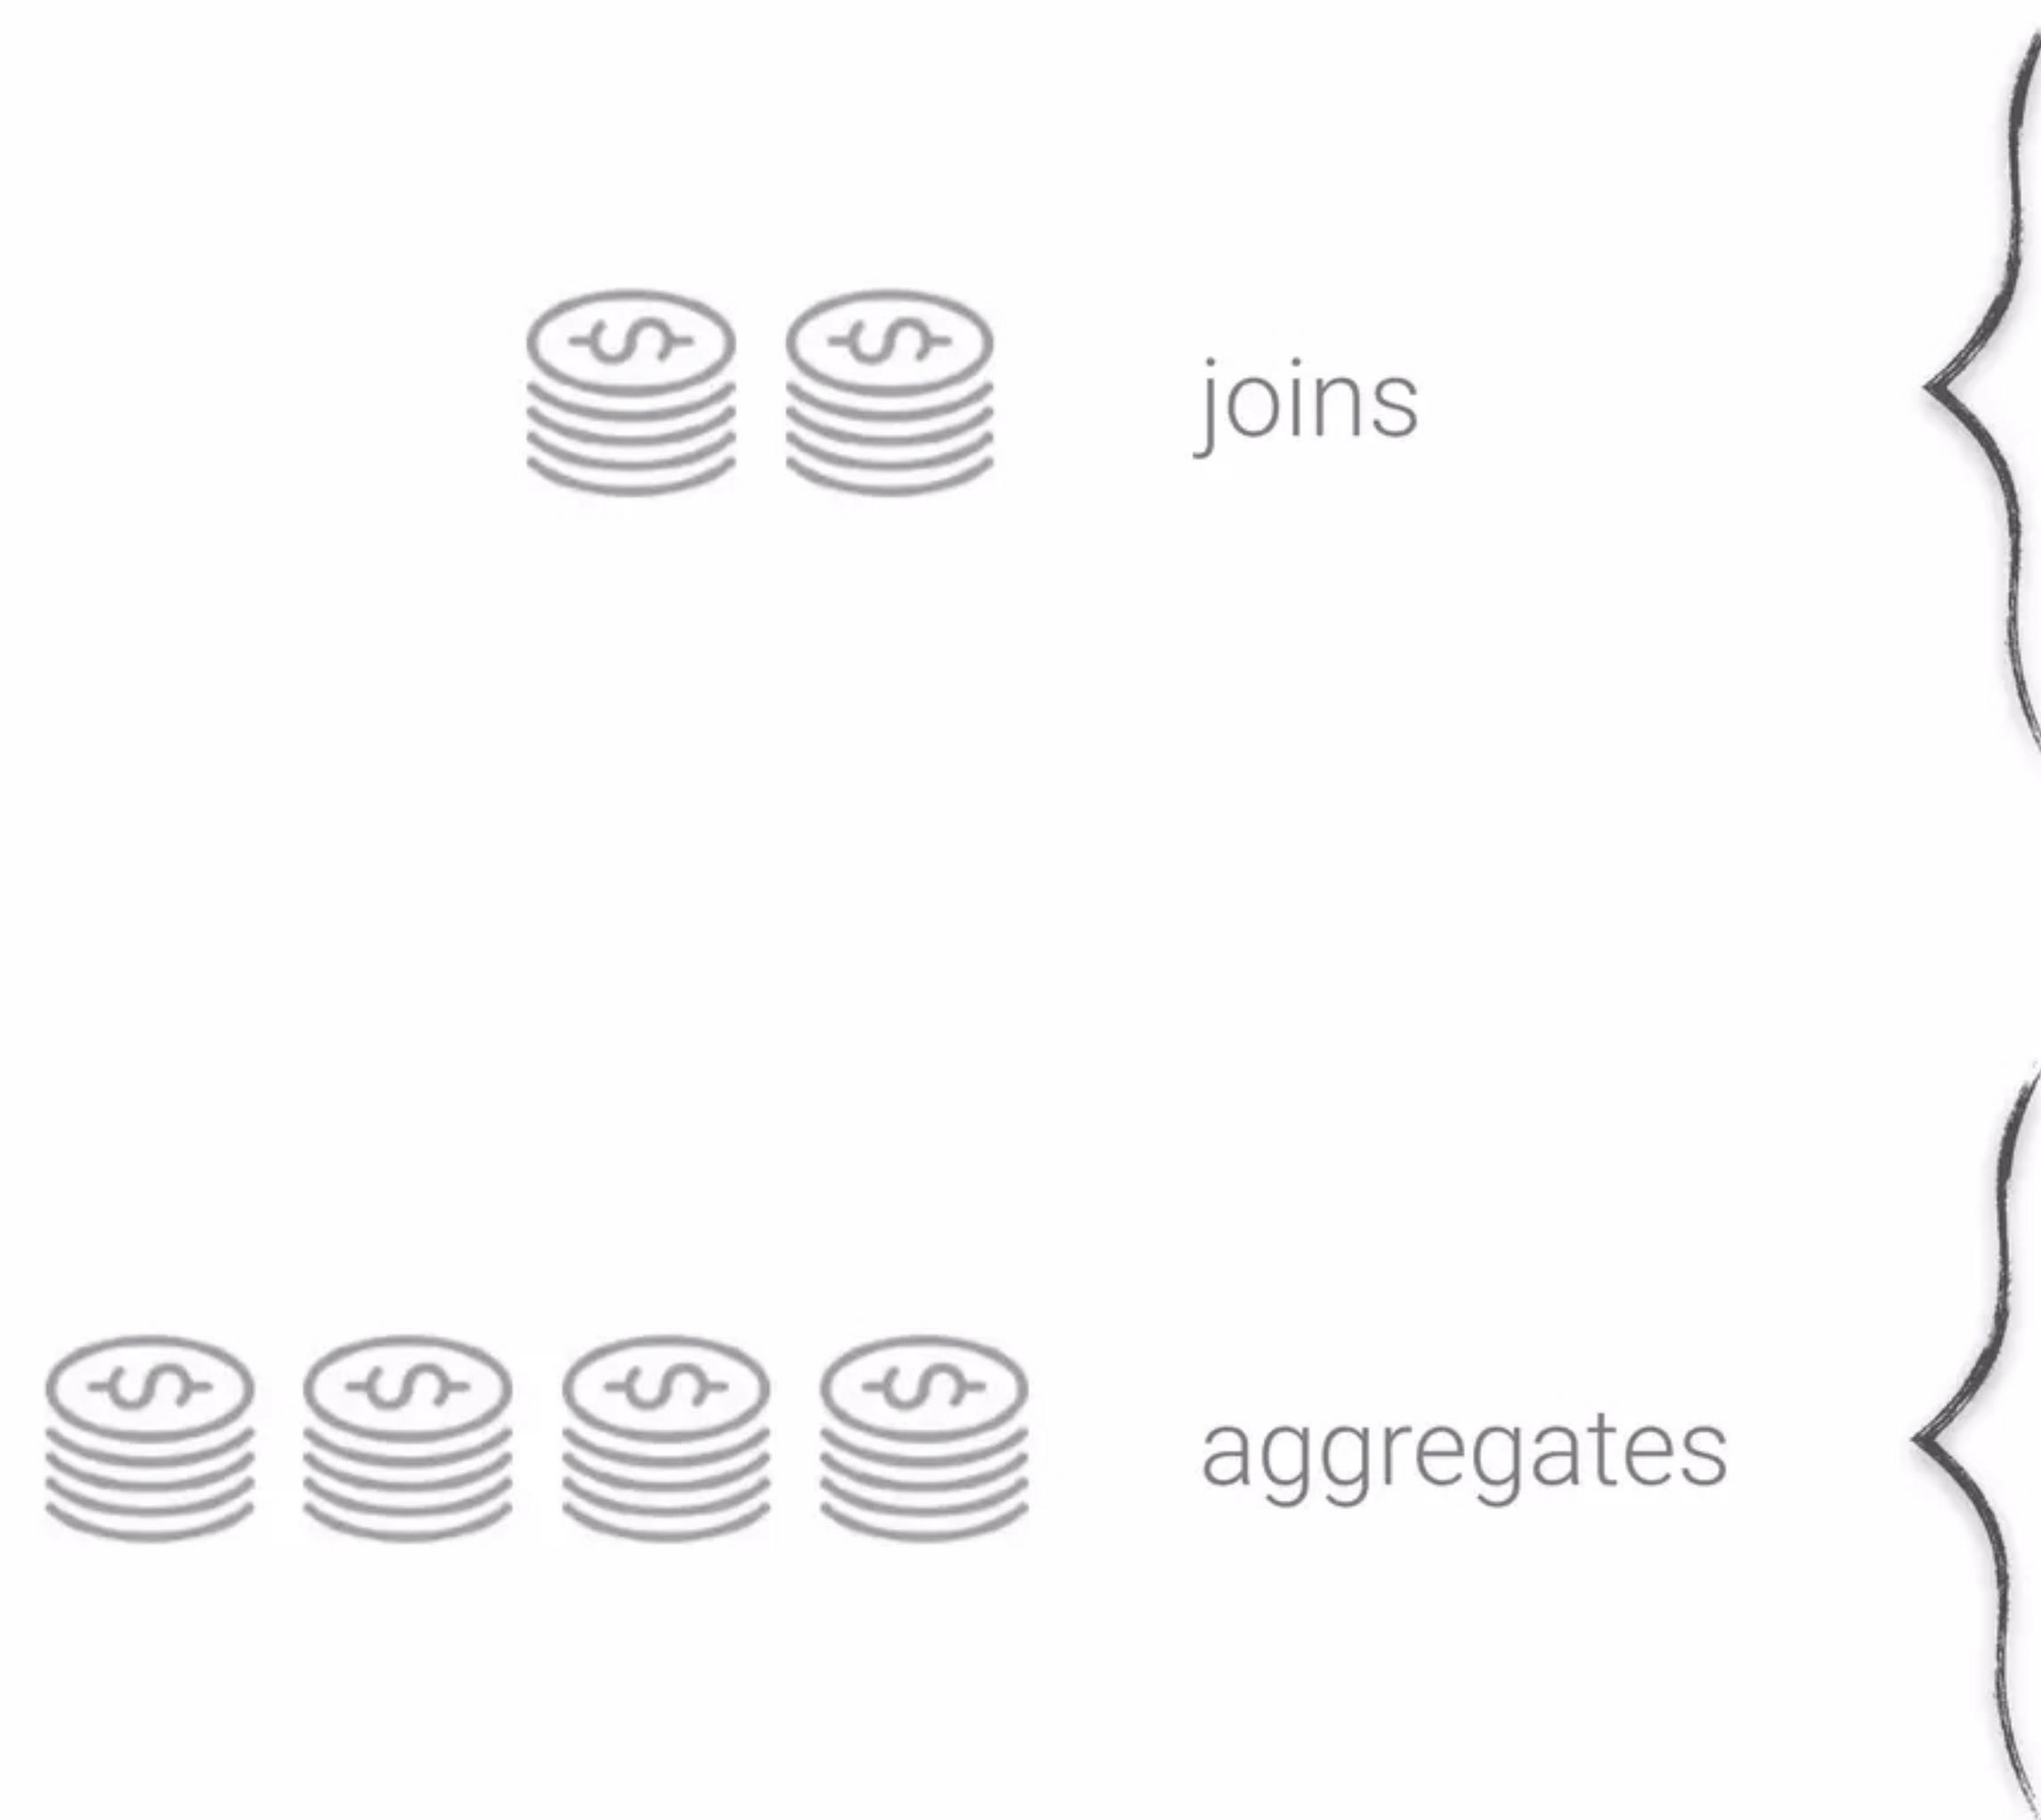


```
application.id = cluster-foo_123
```

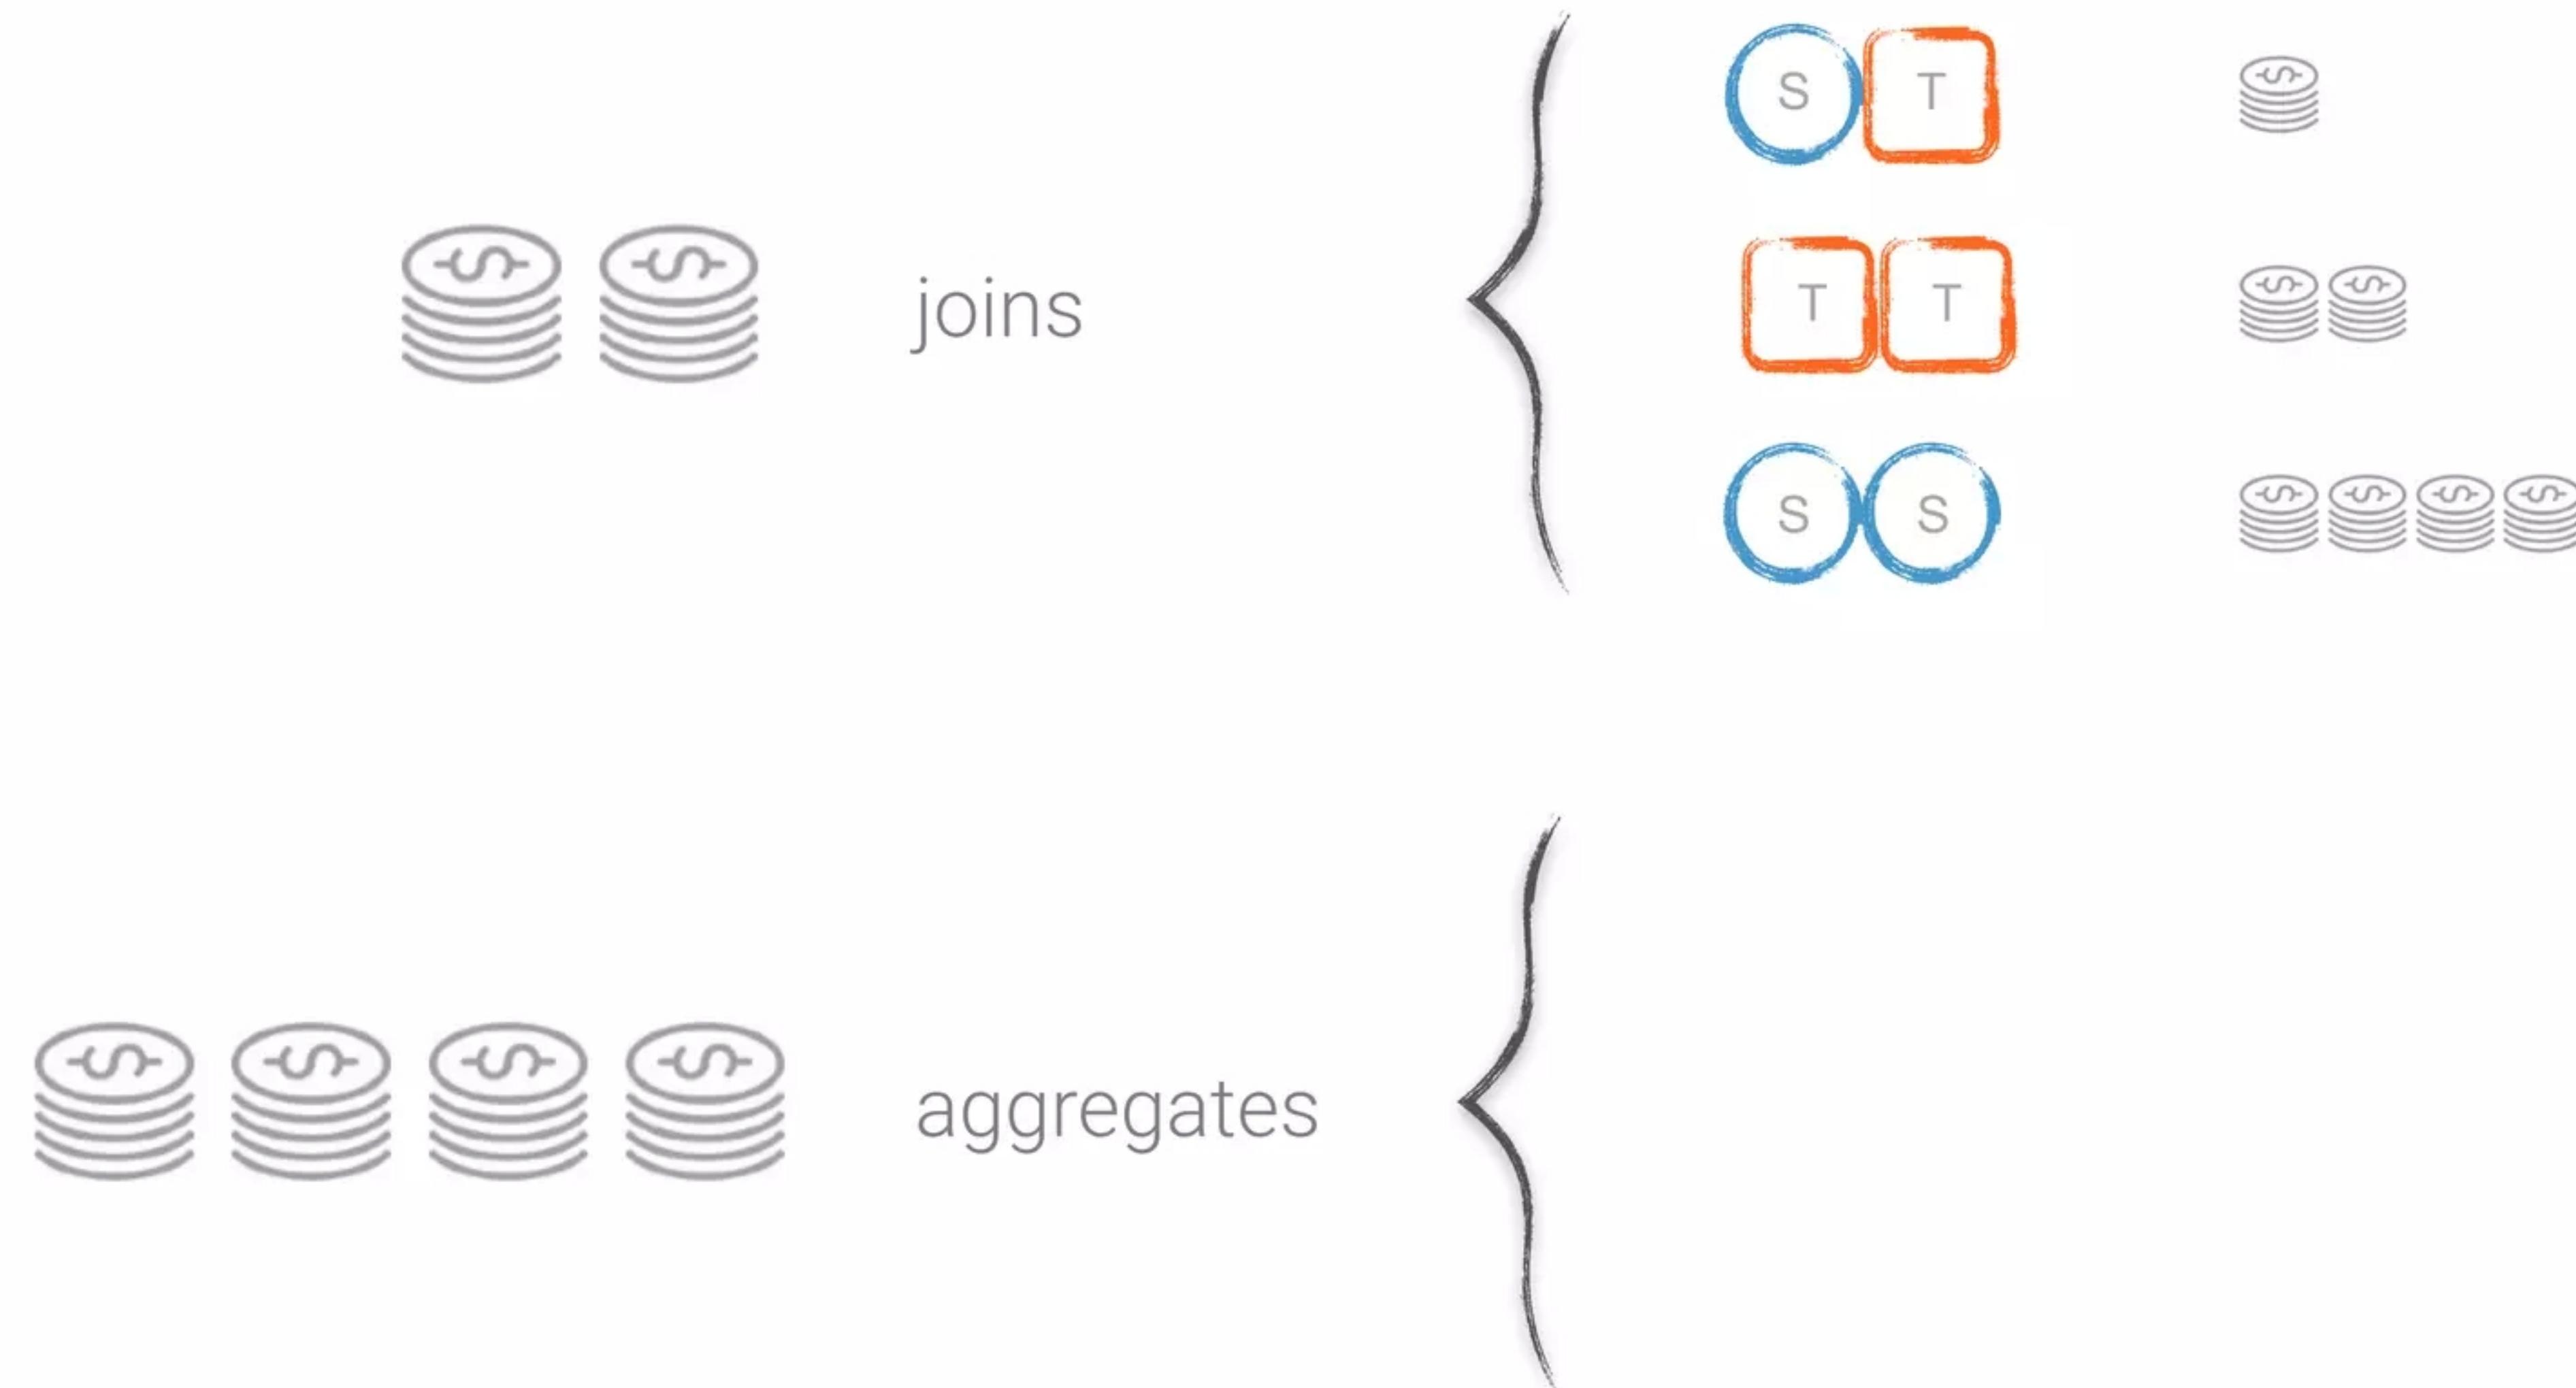
application.id = { cluster-foo_123
ksql.service.id + query_prefix + query_id }



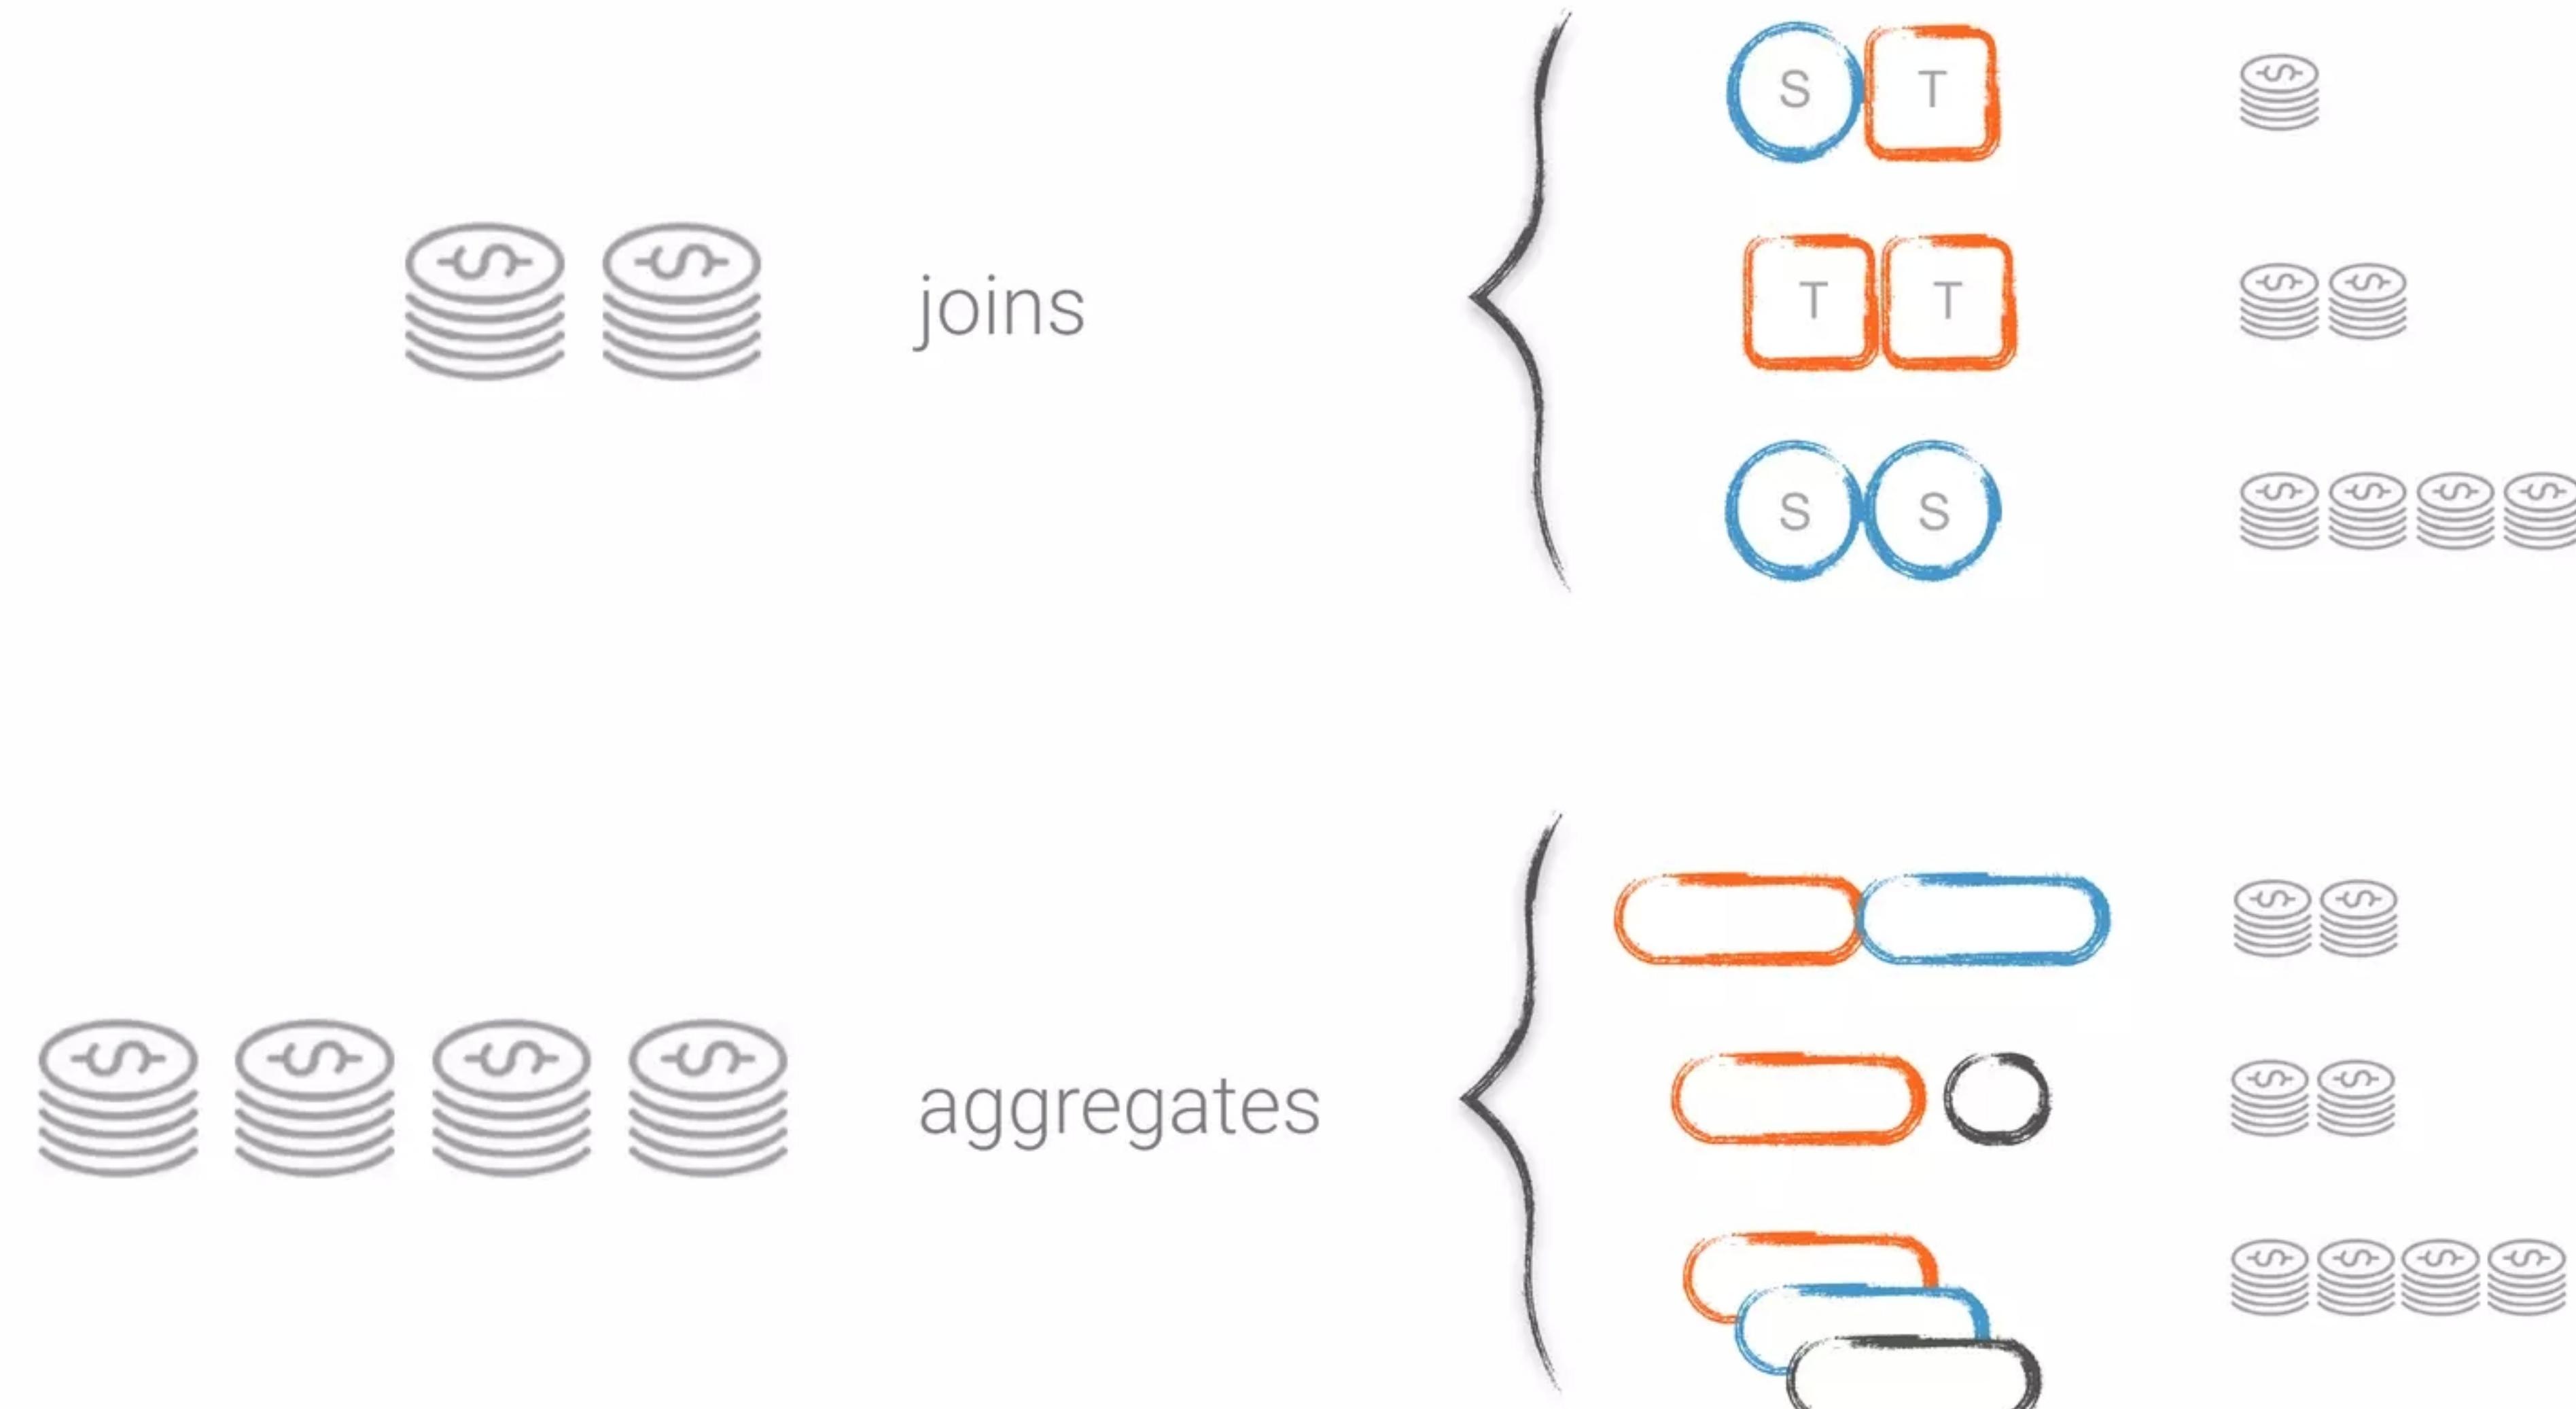




see: <https://docs.confluent.io/current/streams/developer-guide/dsl-api.html#stateful-transformations>



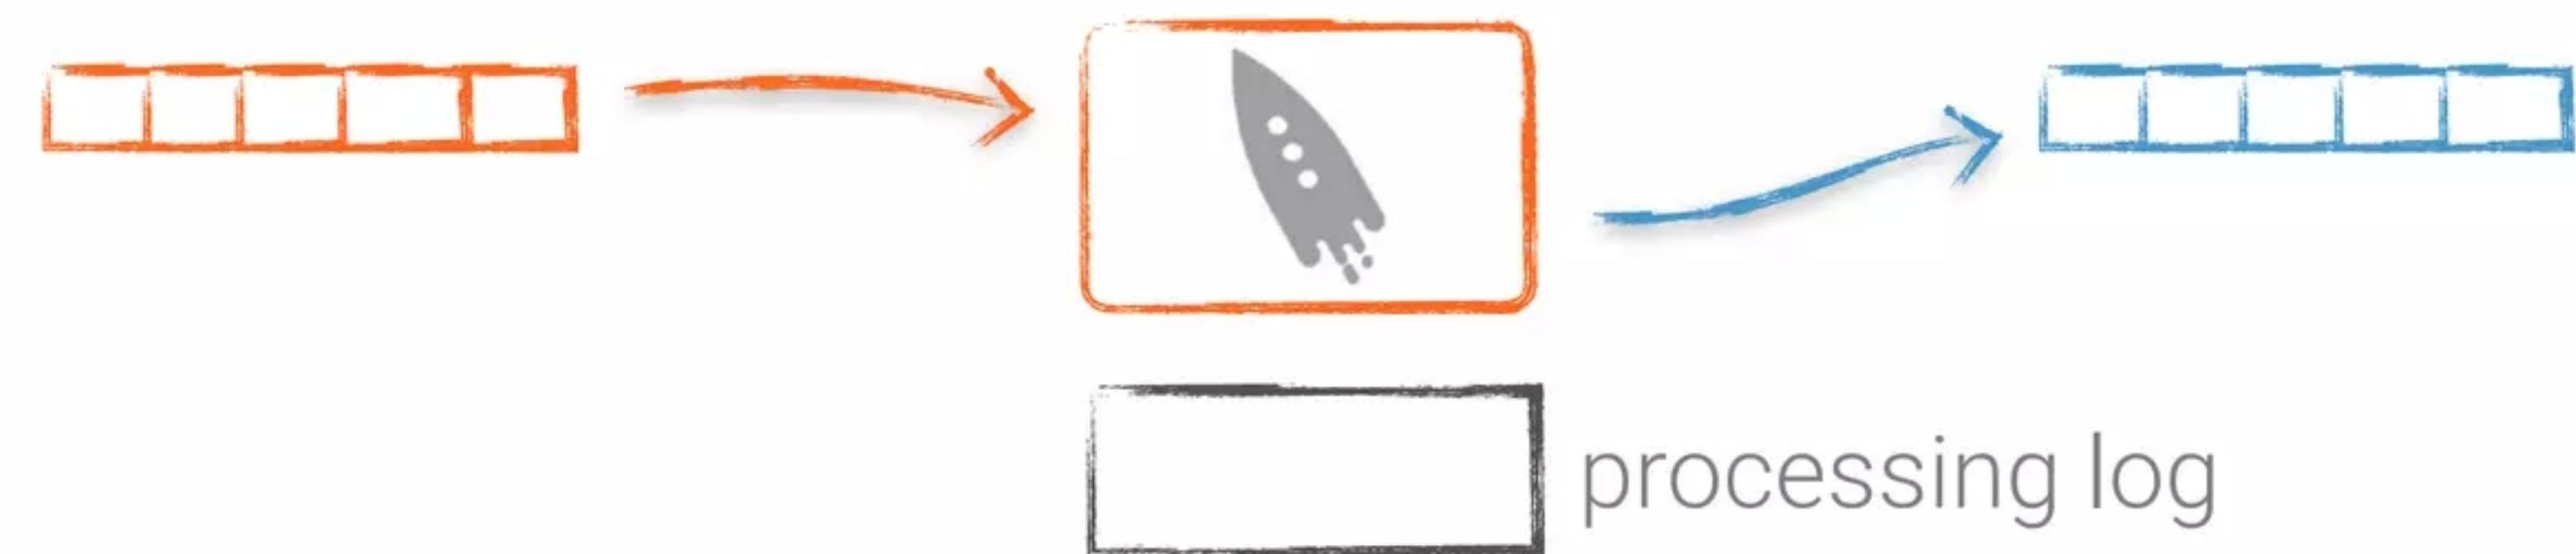
see: [https://docs.confluent.io/current\(streams\)/developer-guide/dsl-api.html#stateful-transformations](https://docs.confluent.io/current(streams)/developer-guide/dsl-api.html#stateful-transformations)

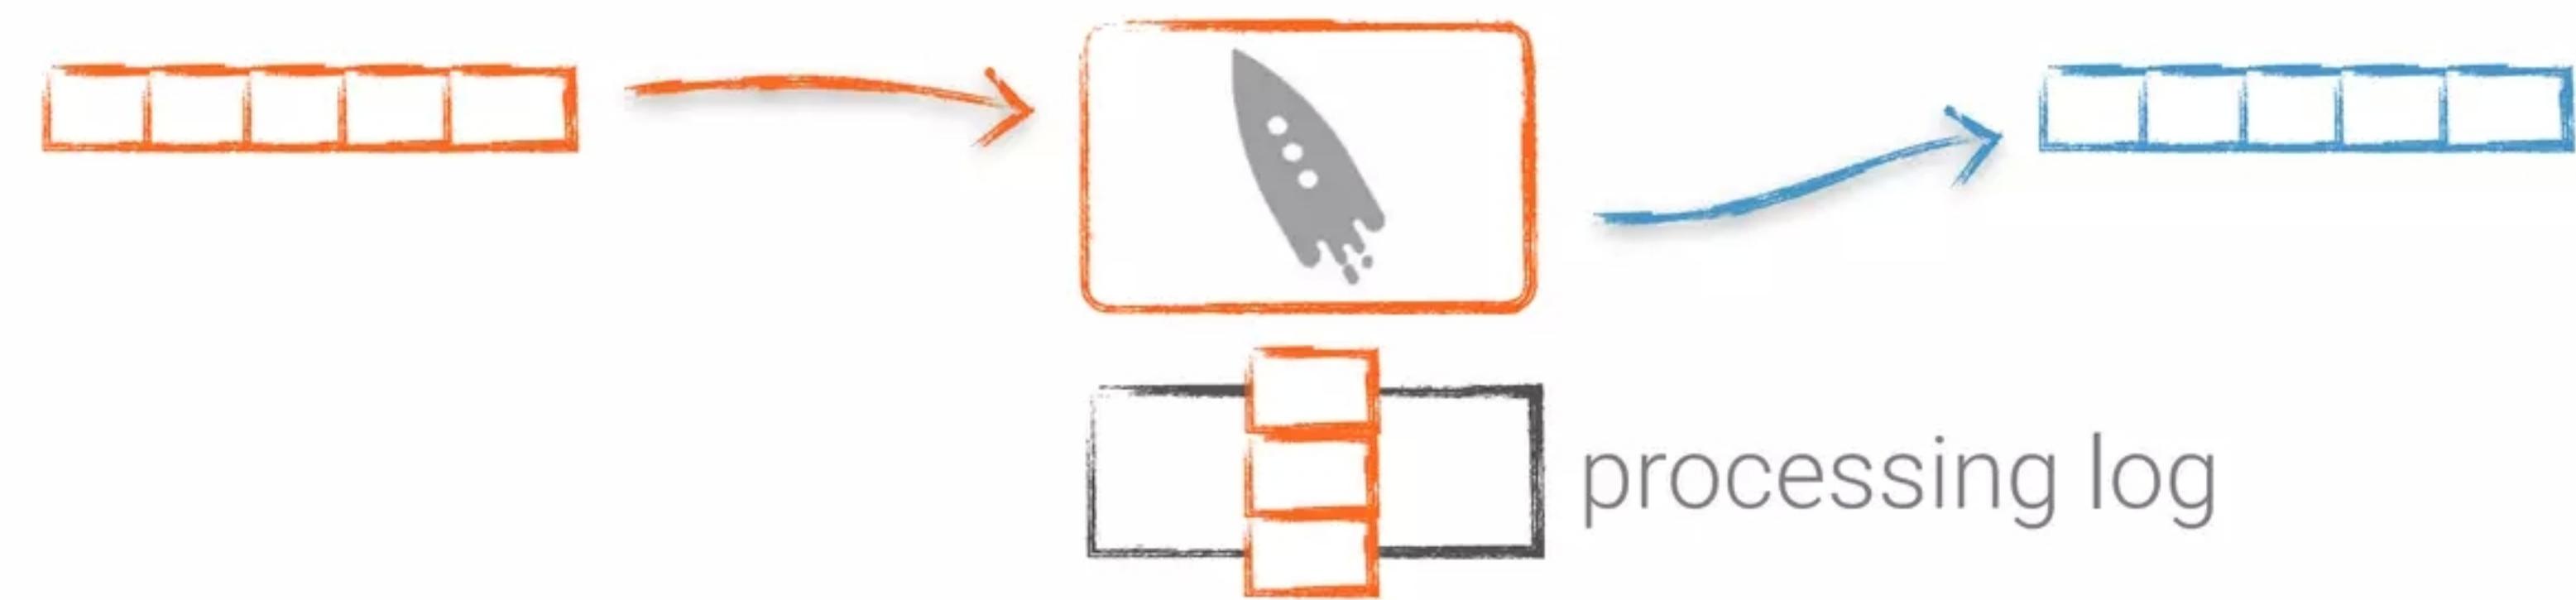


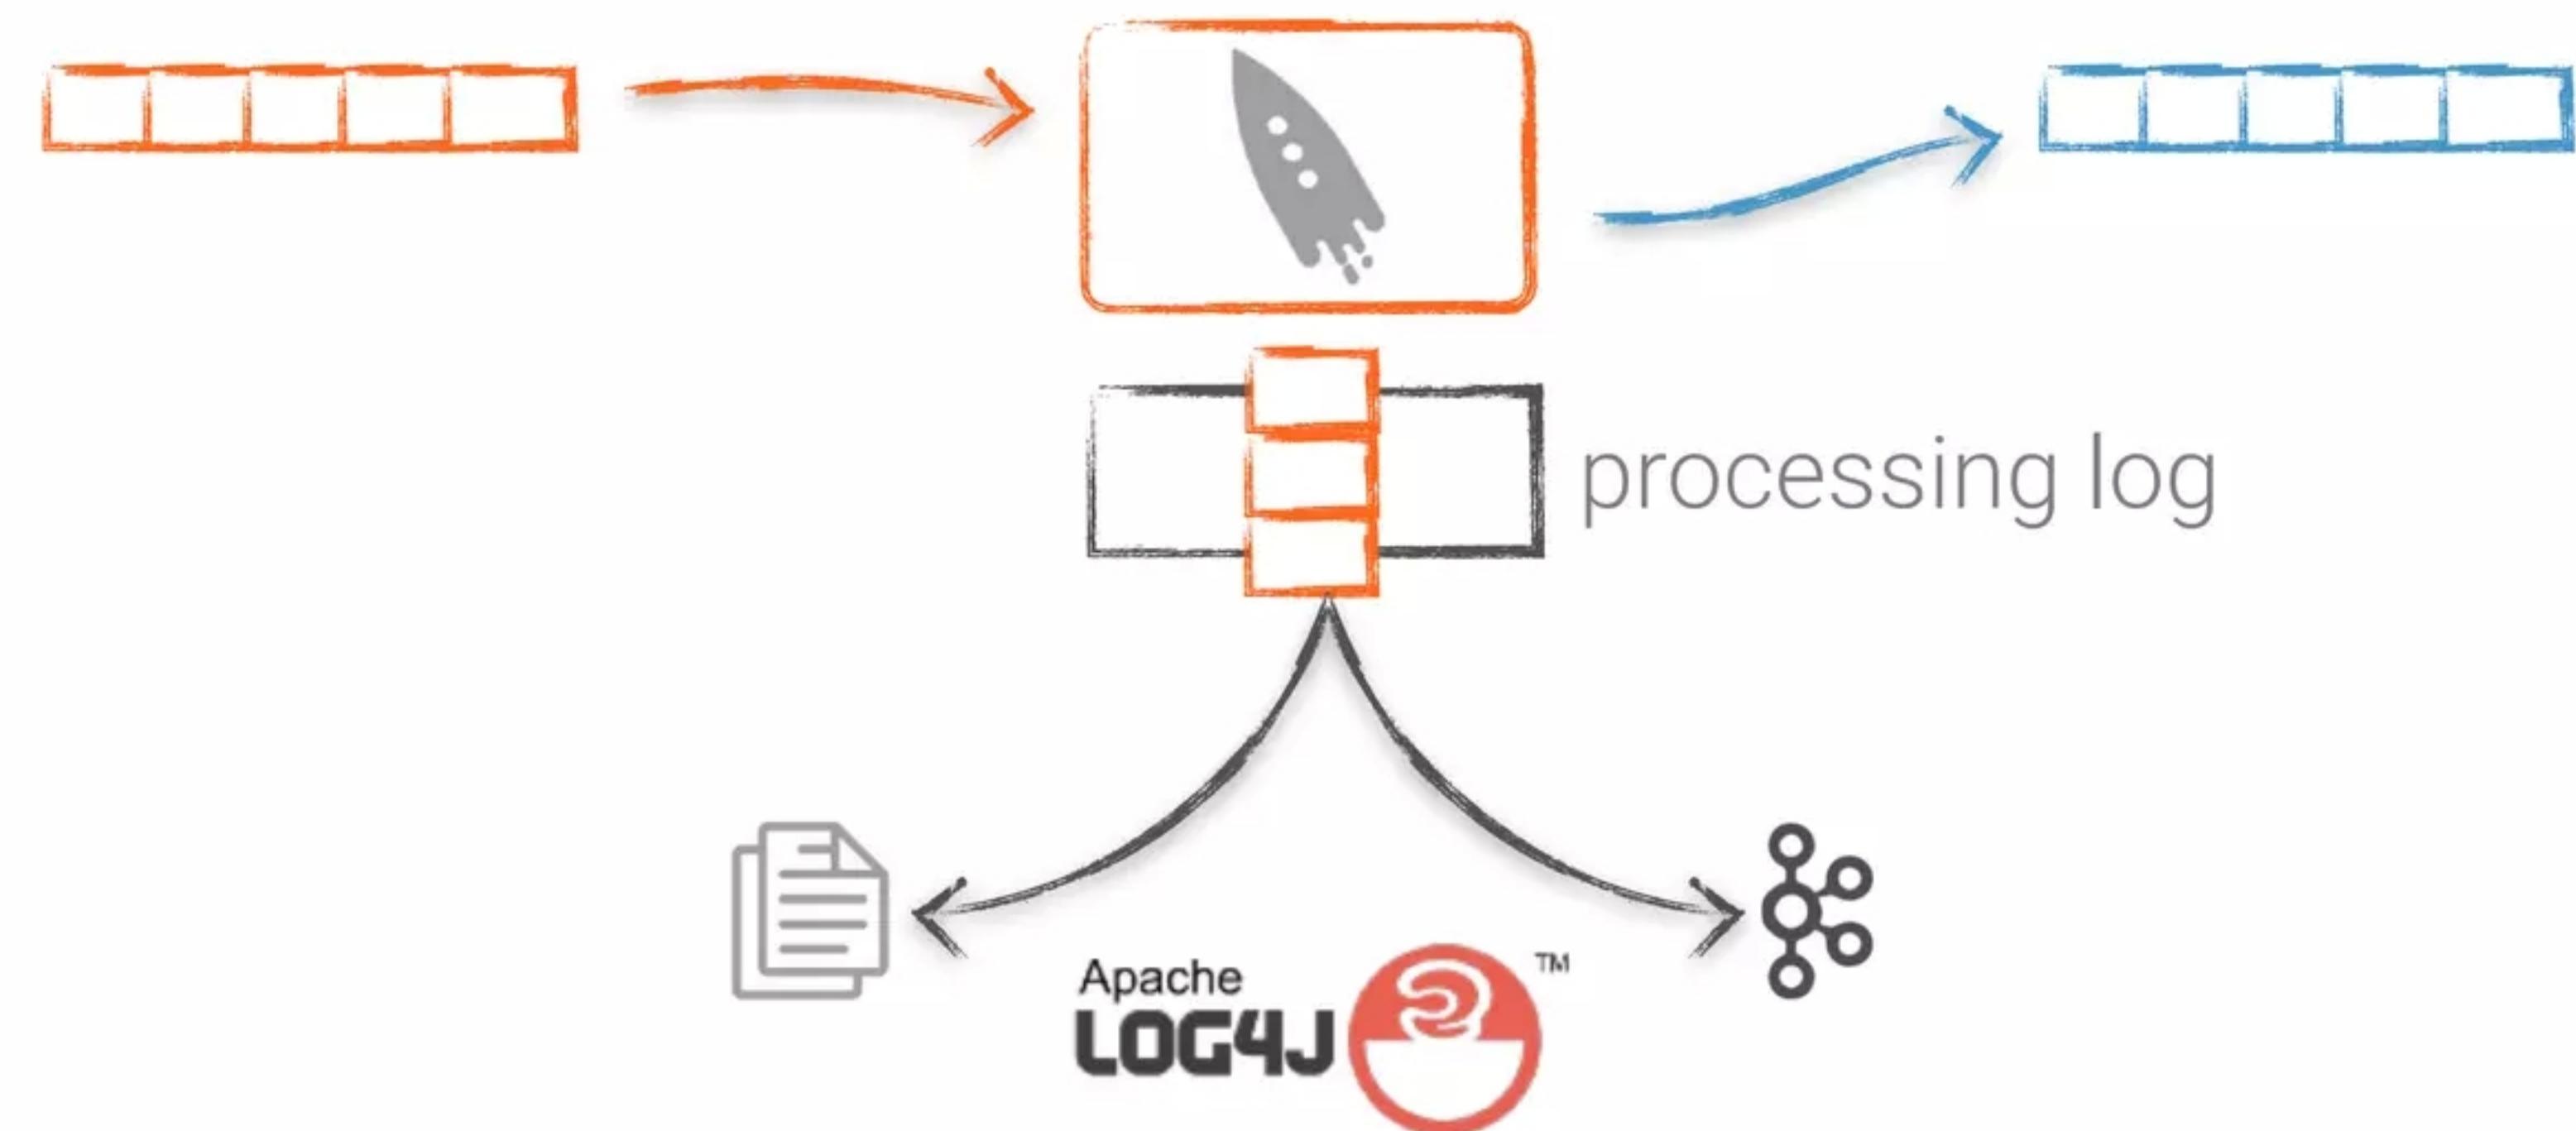
see: <https://docs.confluent.io/currentstreams/developer-guide/dsl-api.html#stateful-transformations>

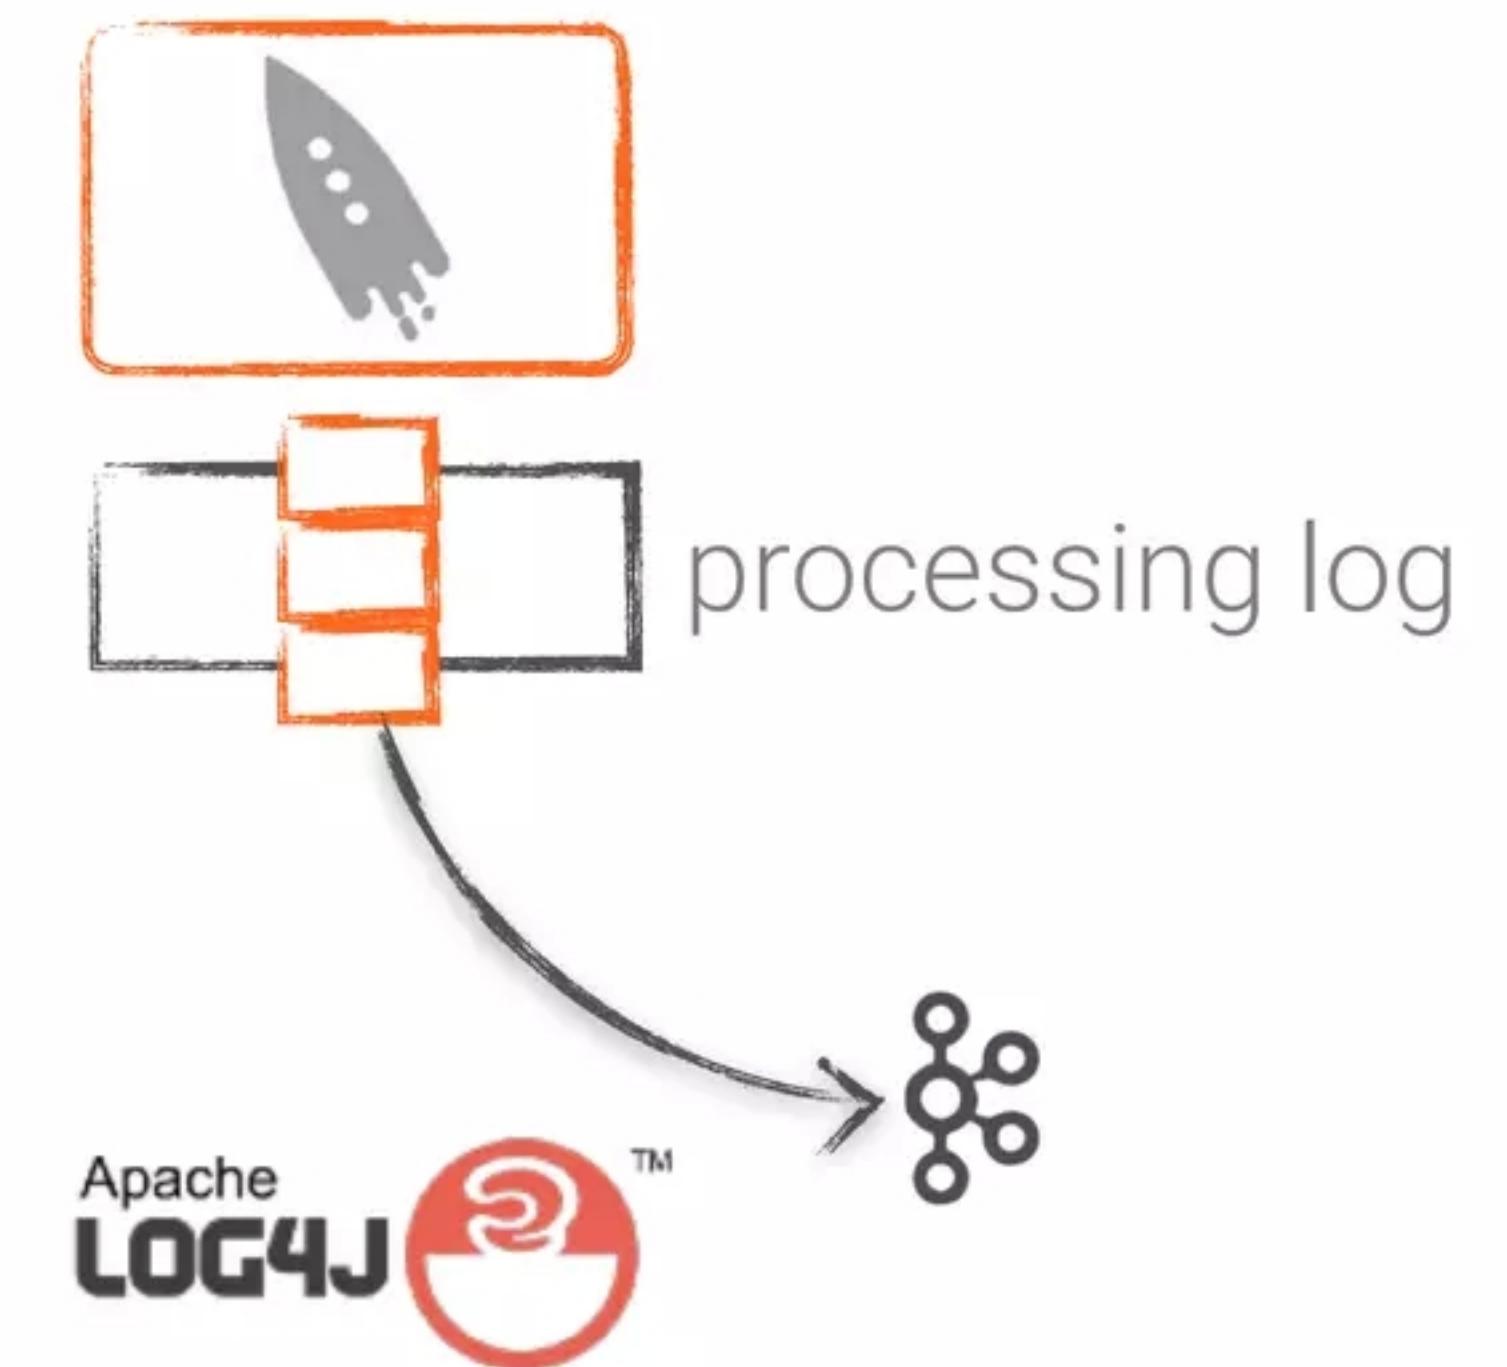
- Intro
- Develop
- Deploy
- **Operate**
- Common Mistakes
- Q&A

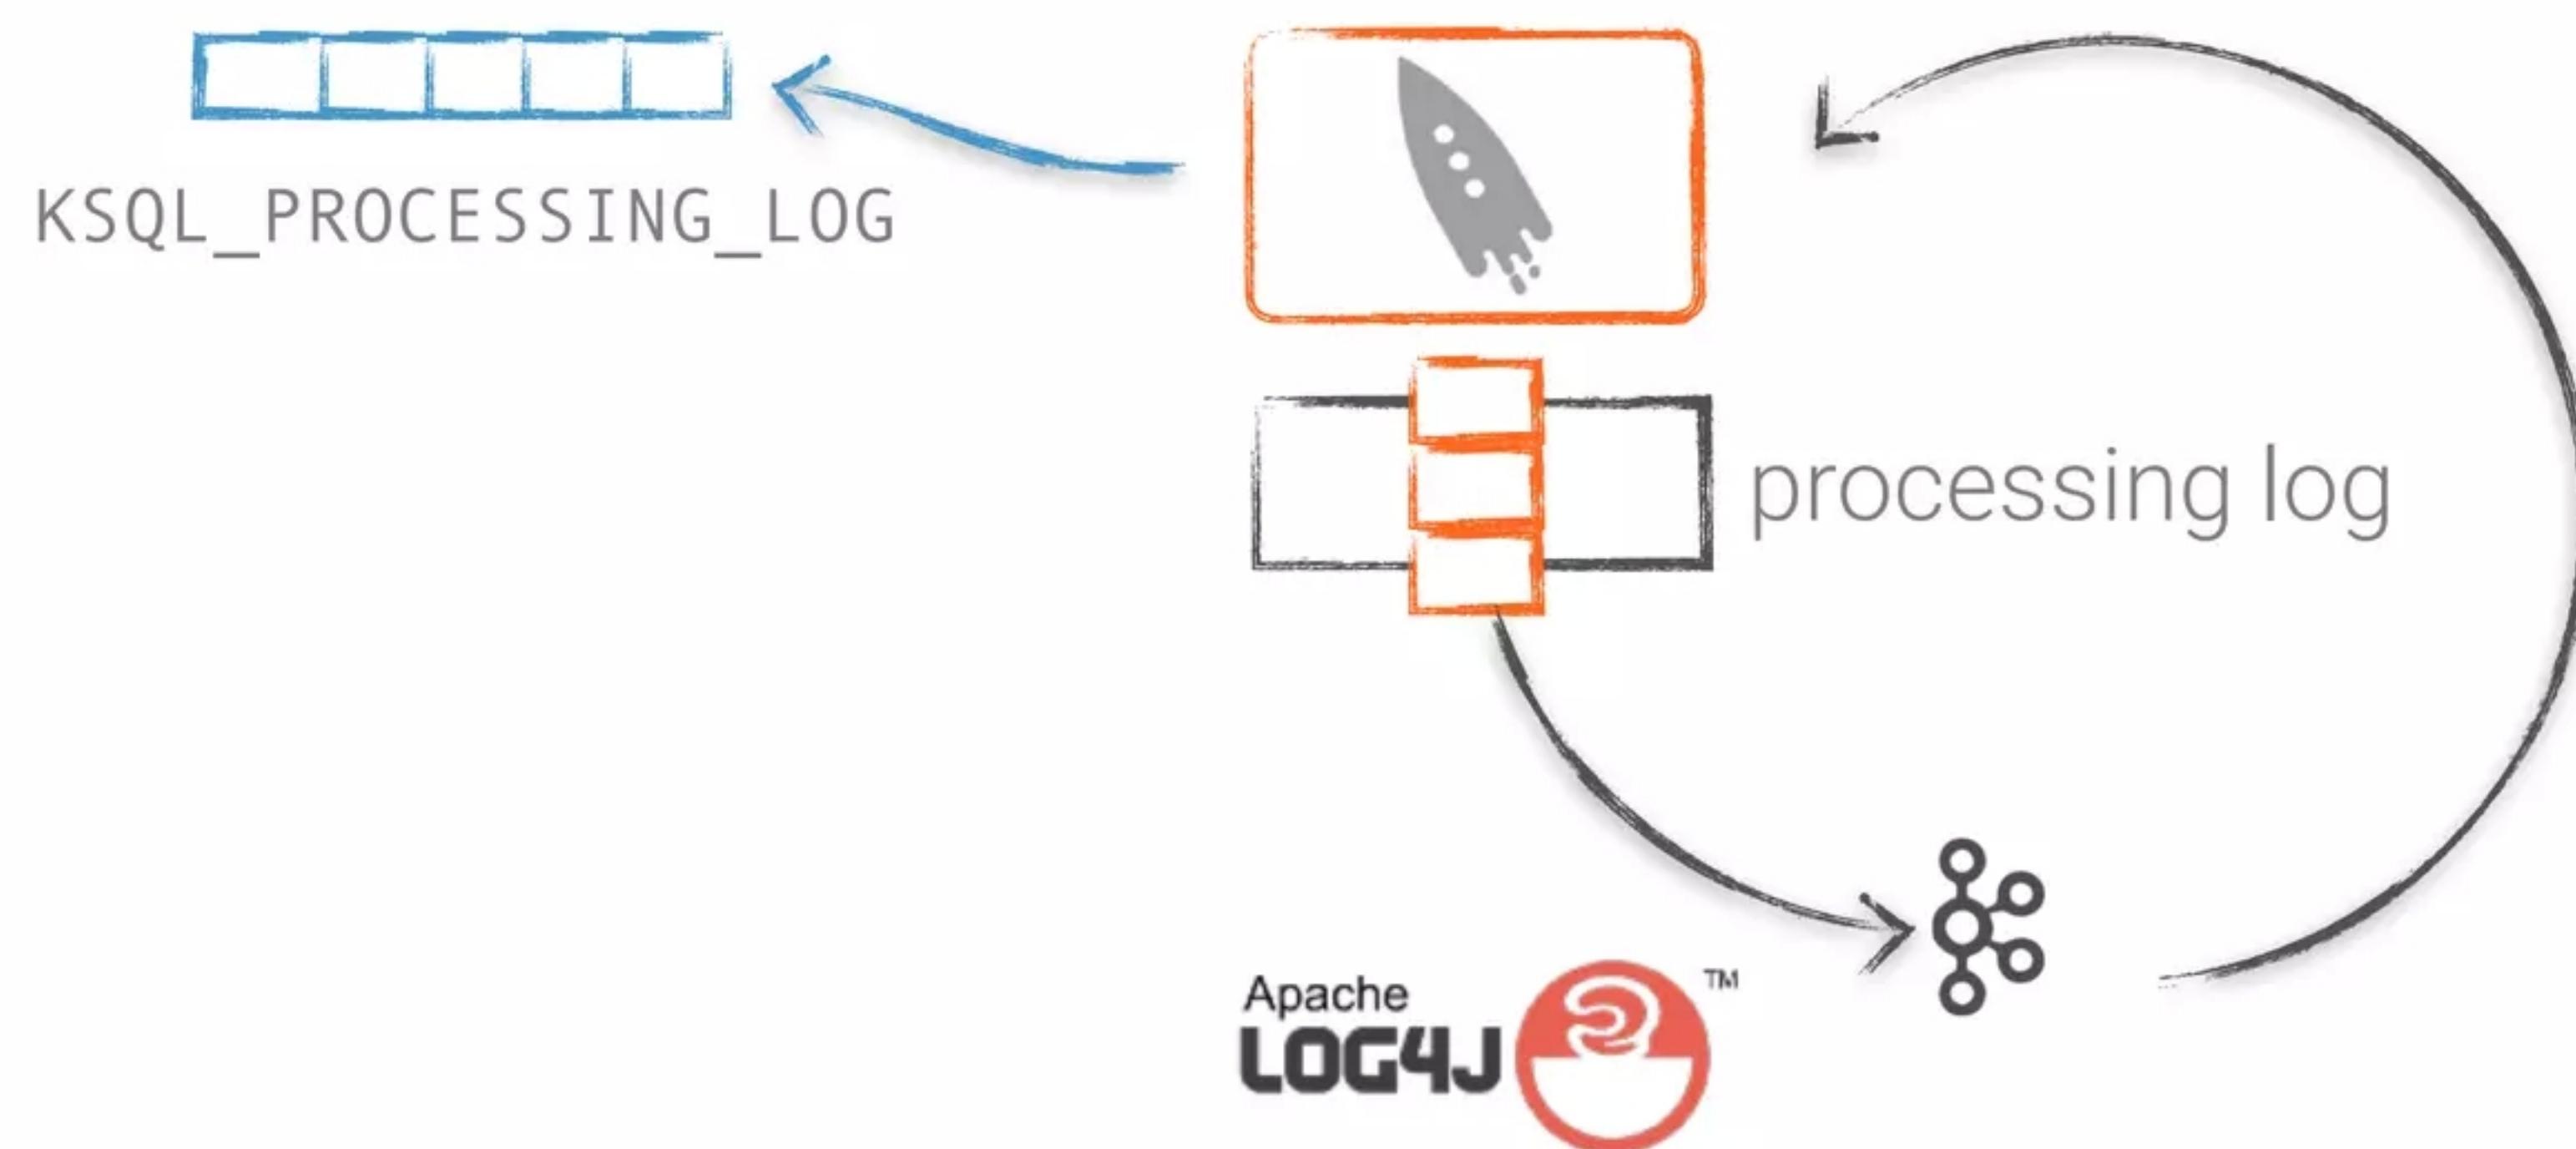


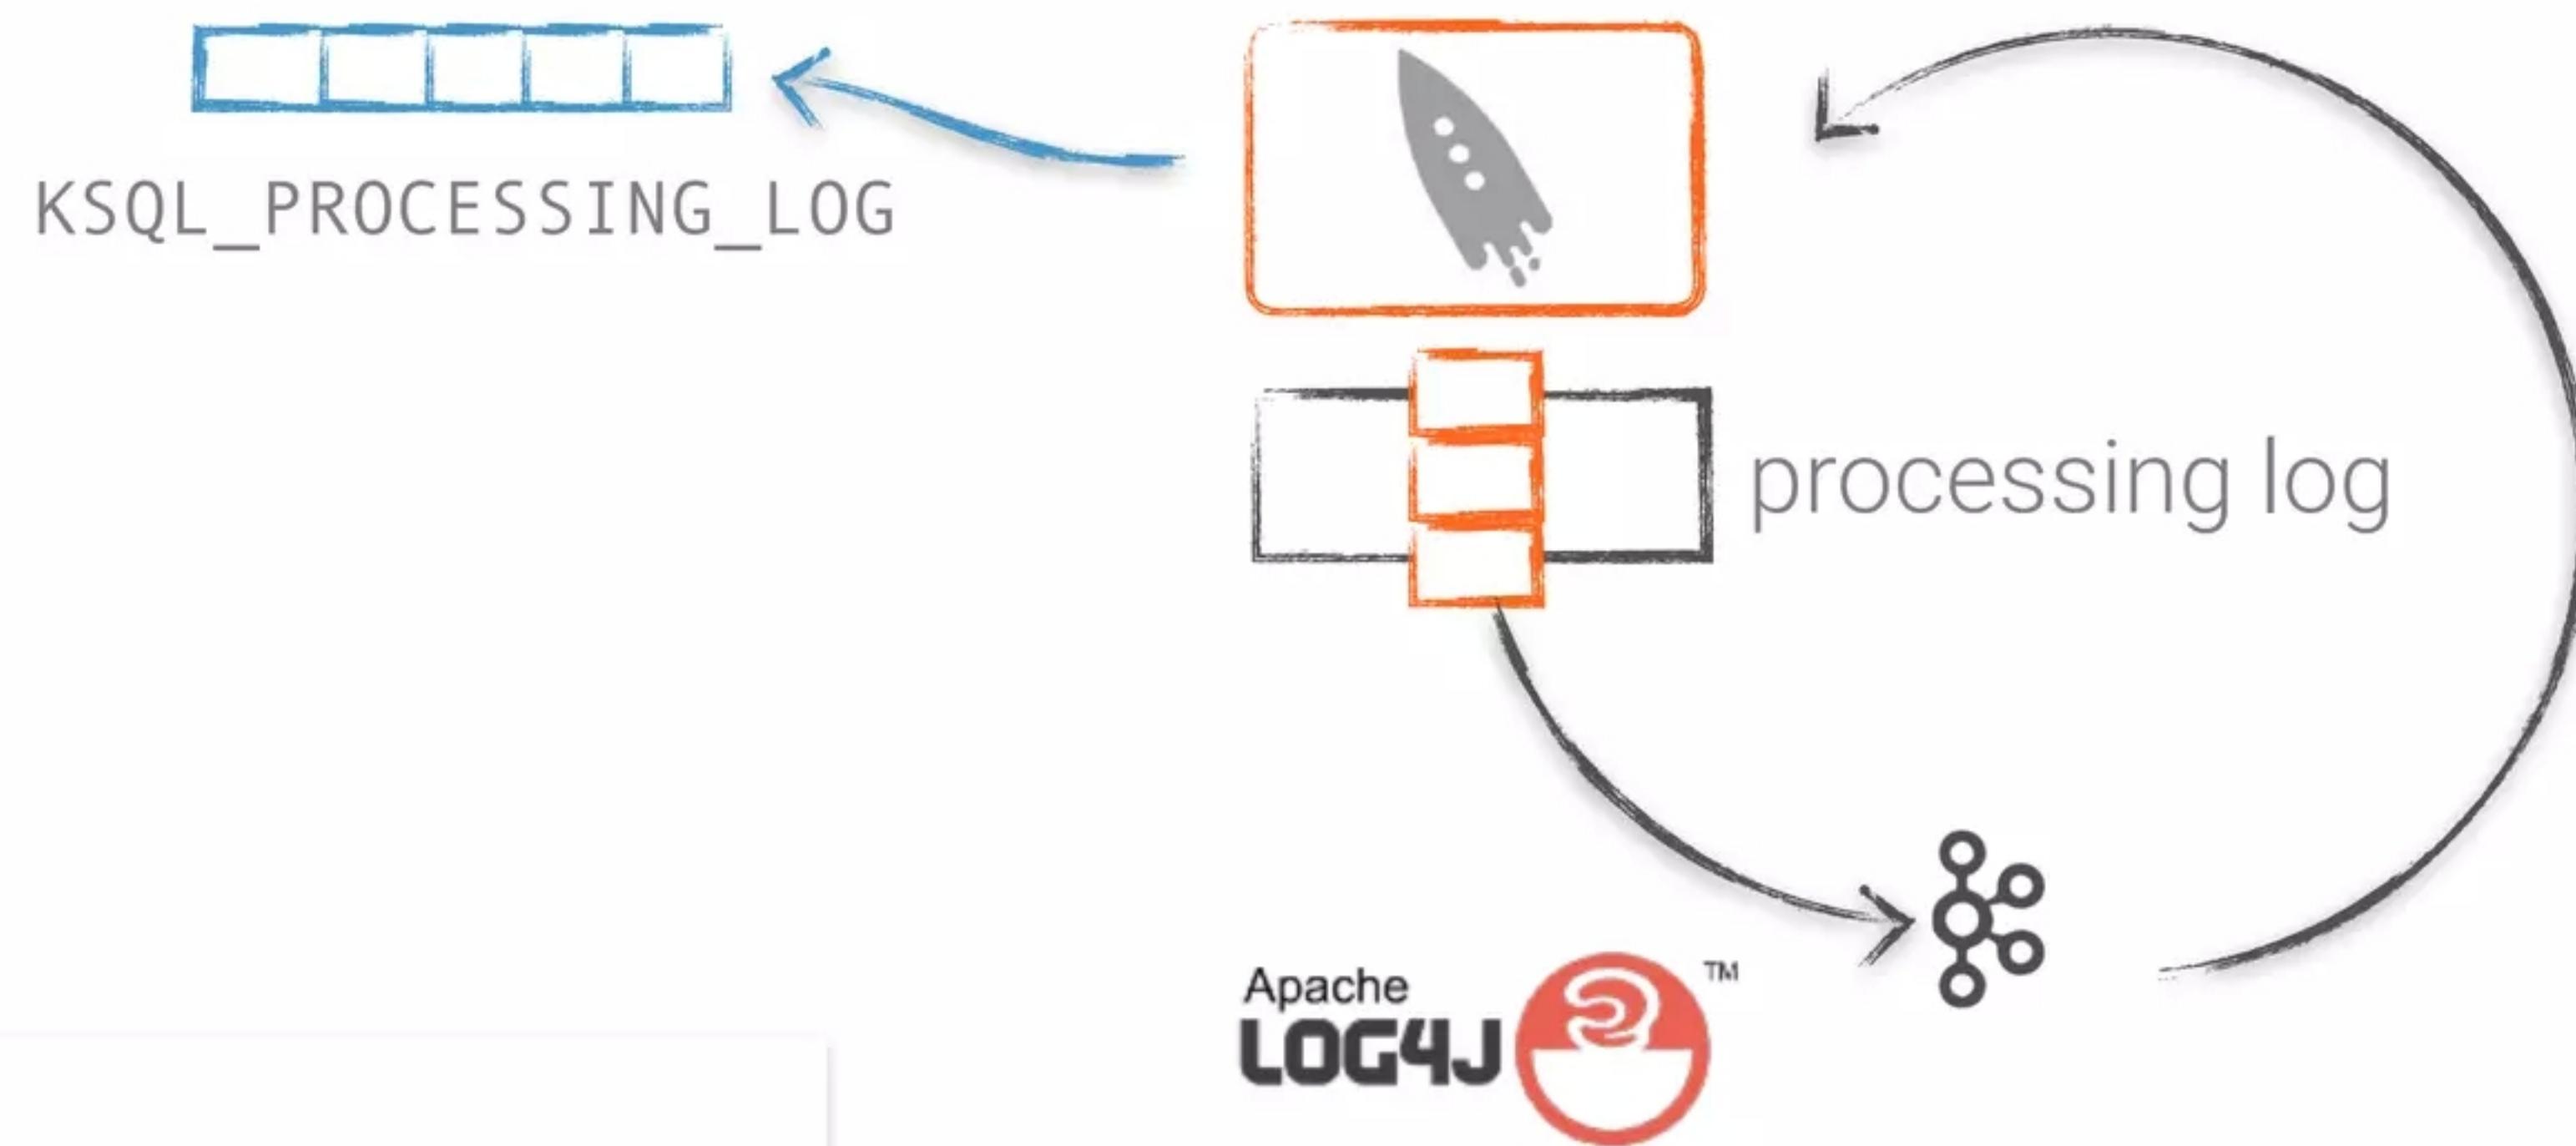




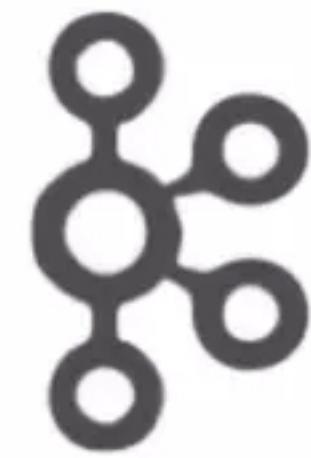


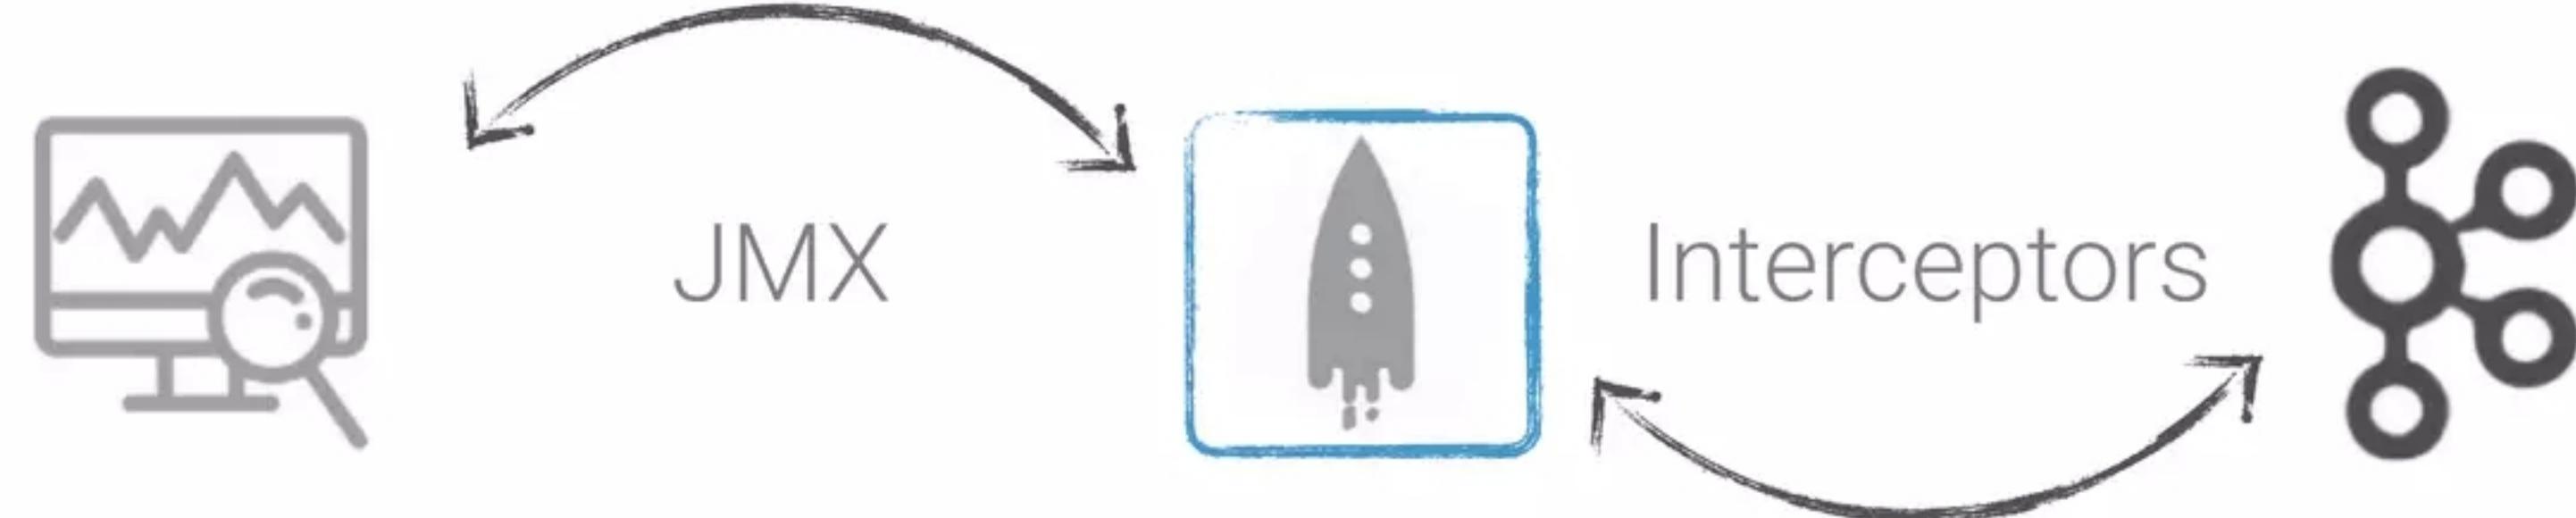


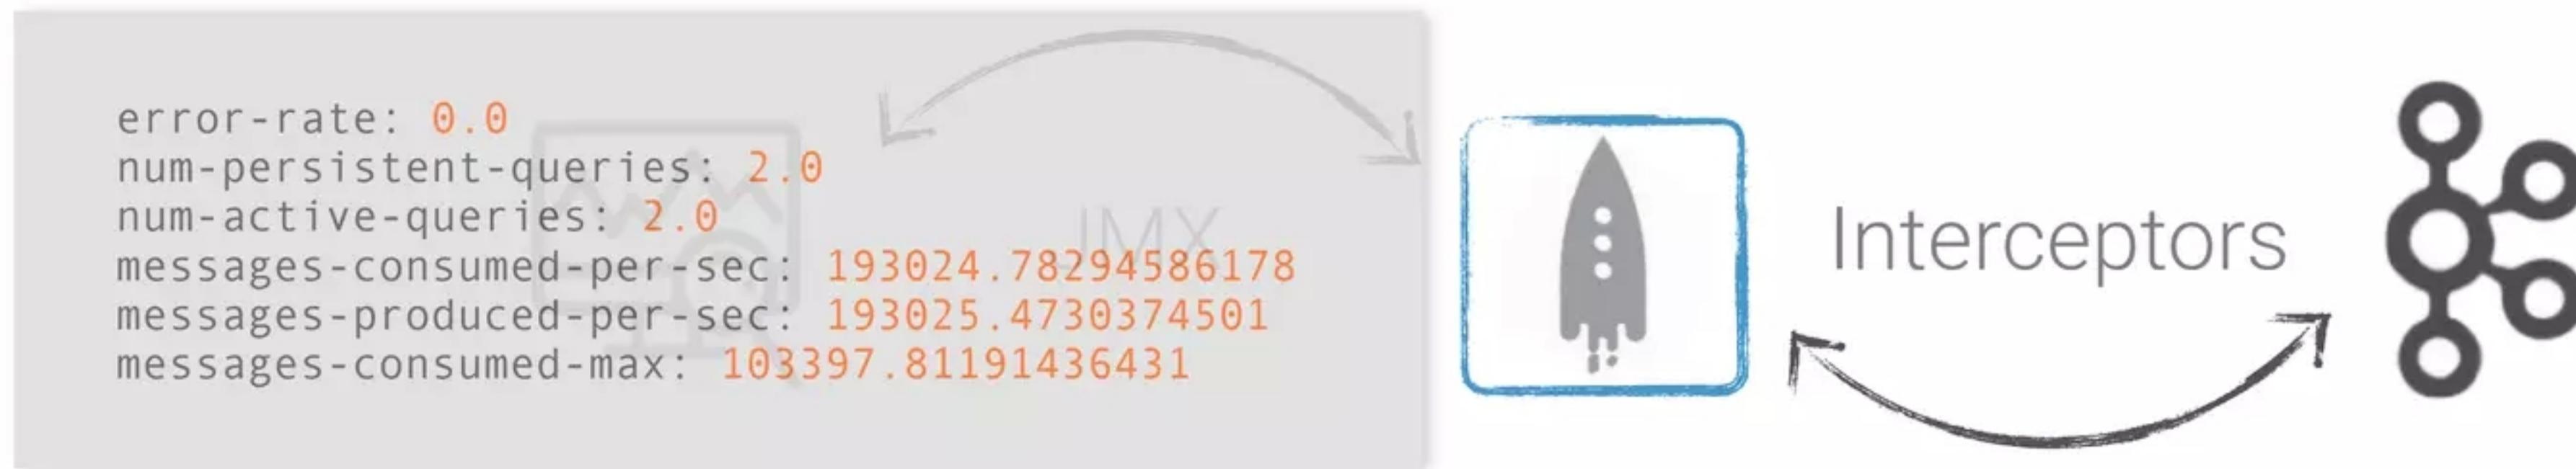


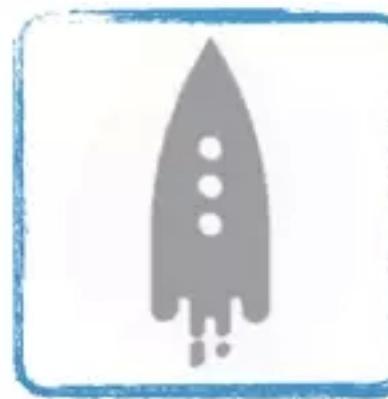


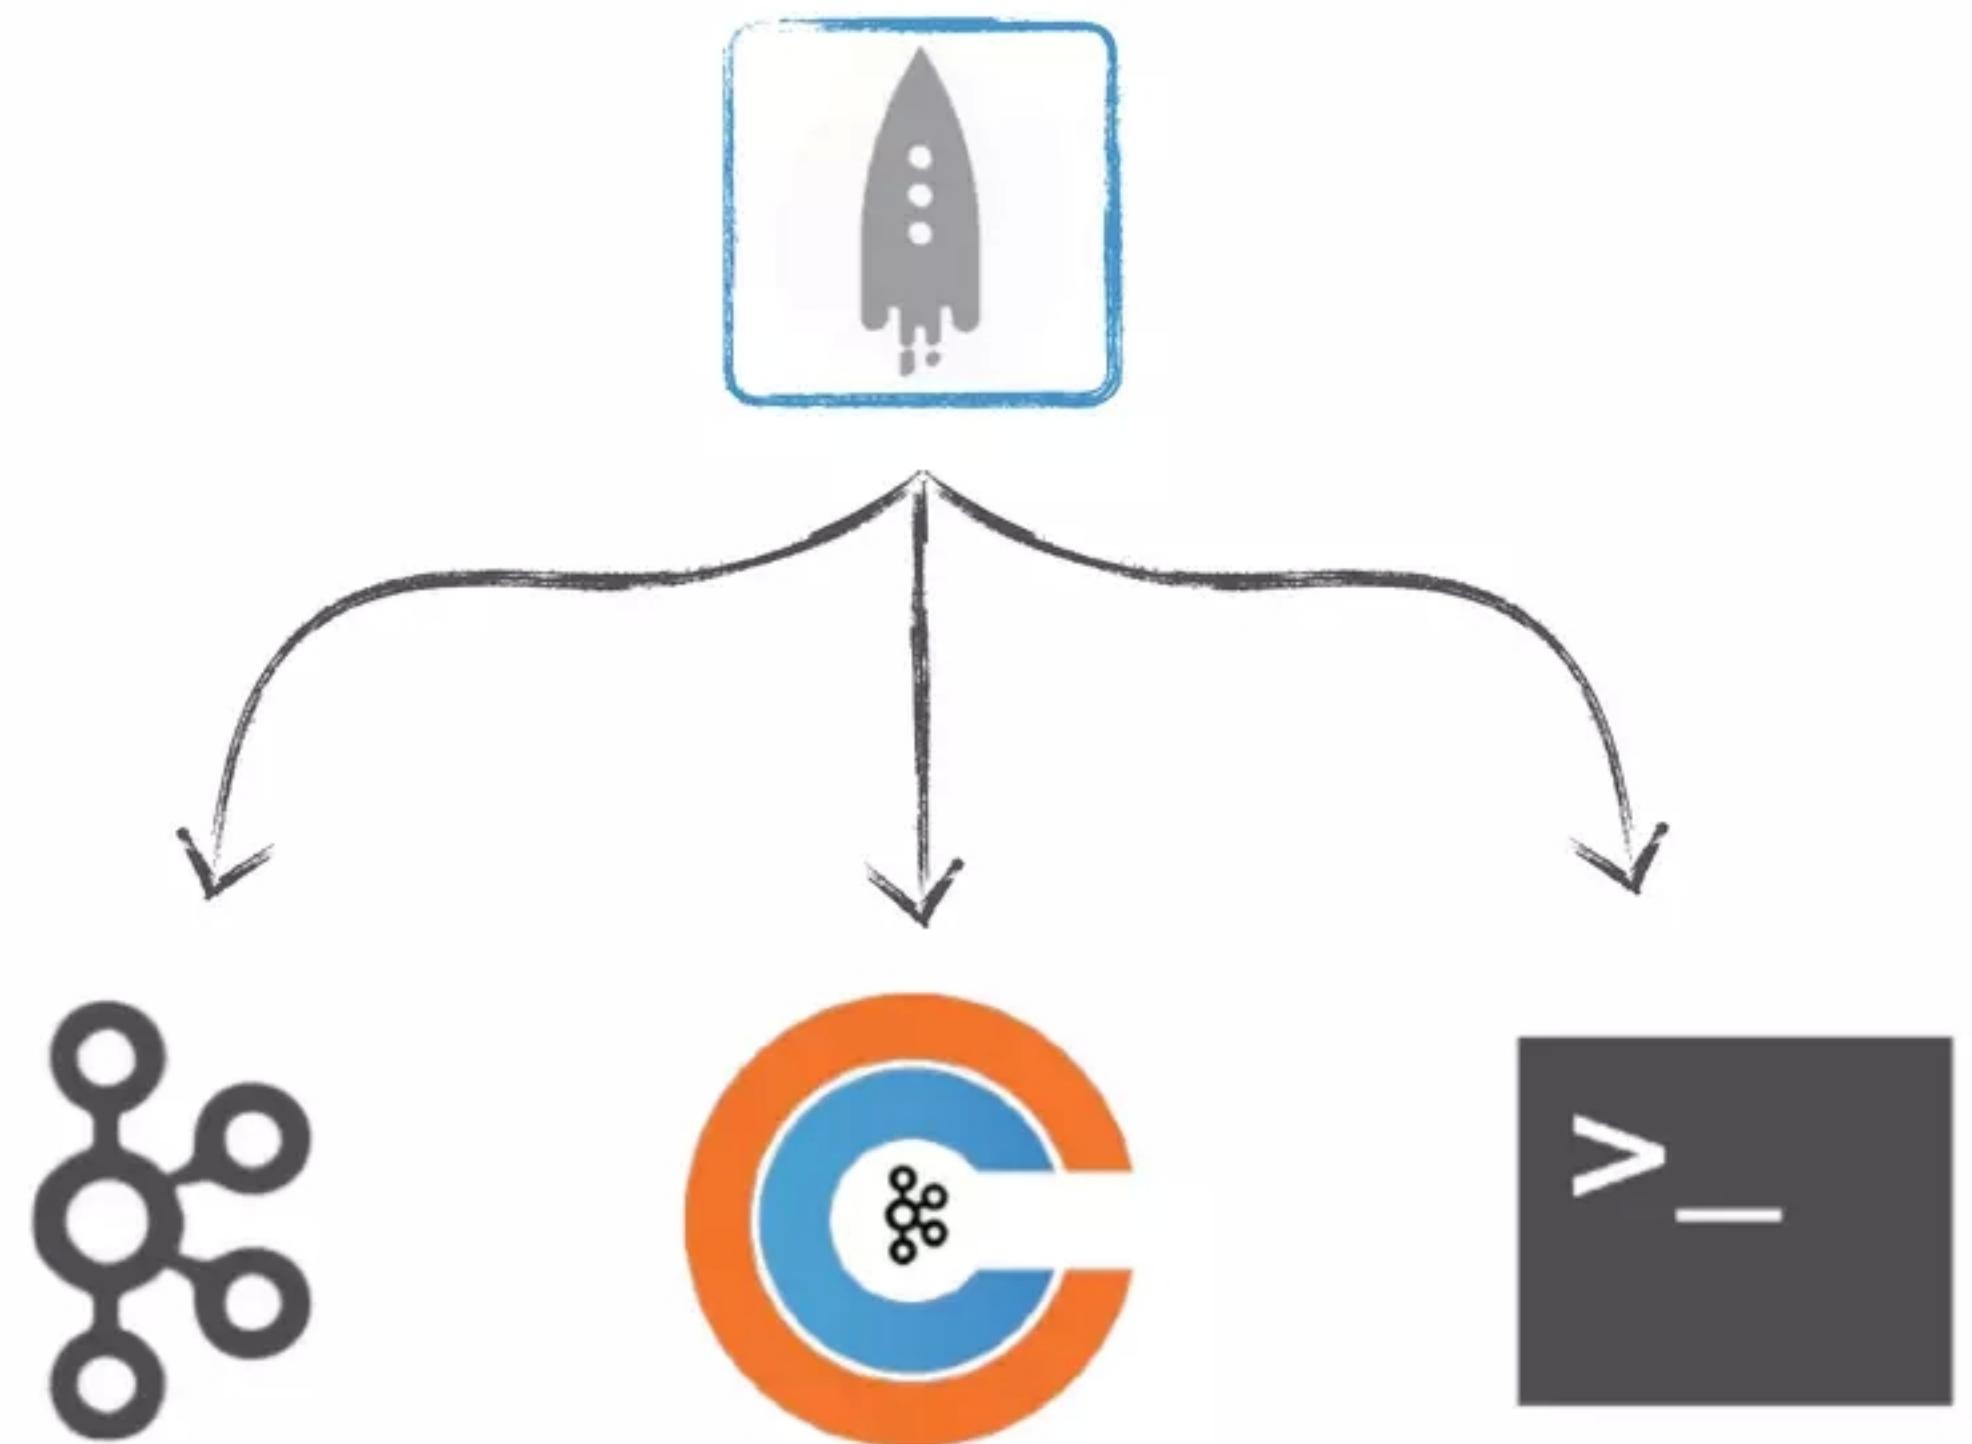
✓ **tip:** the best way to debug streaming is by streaming

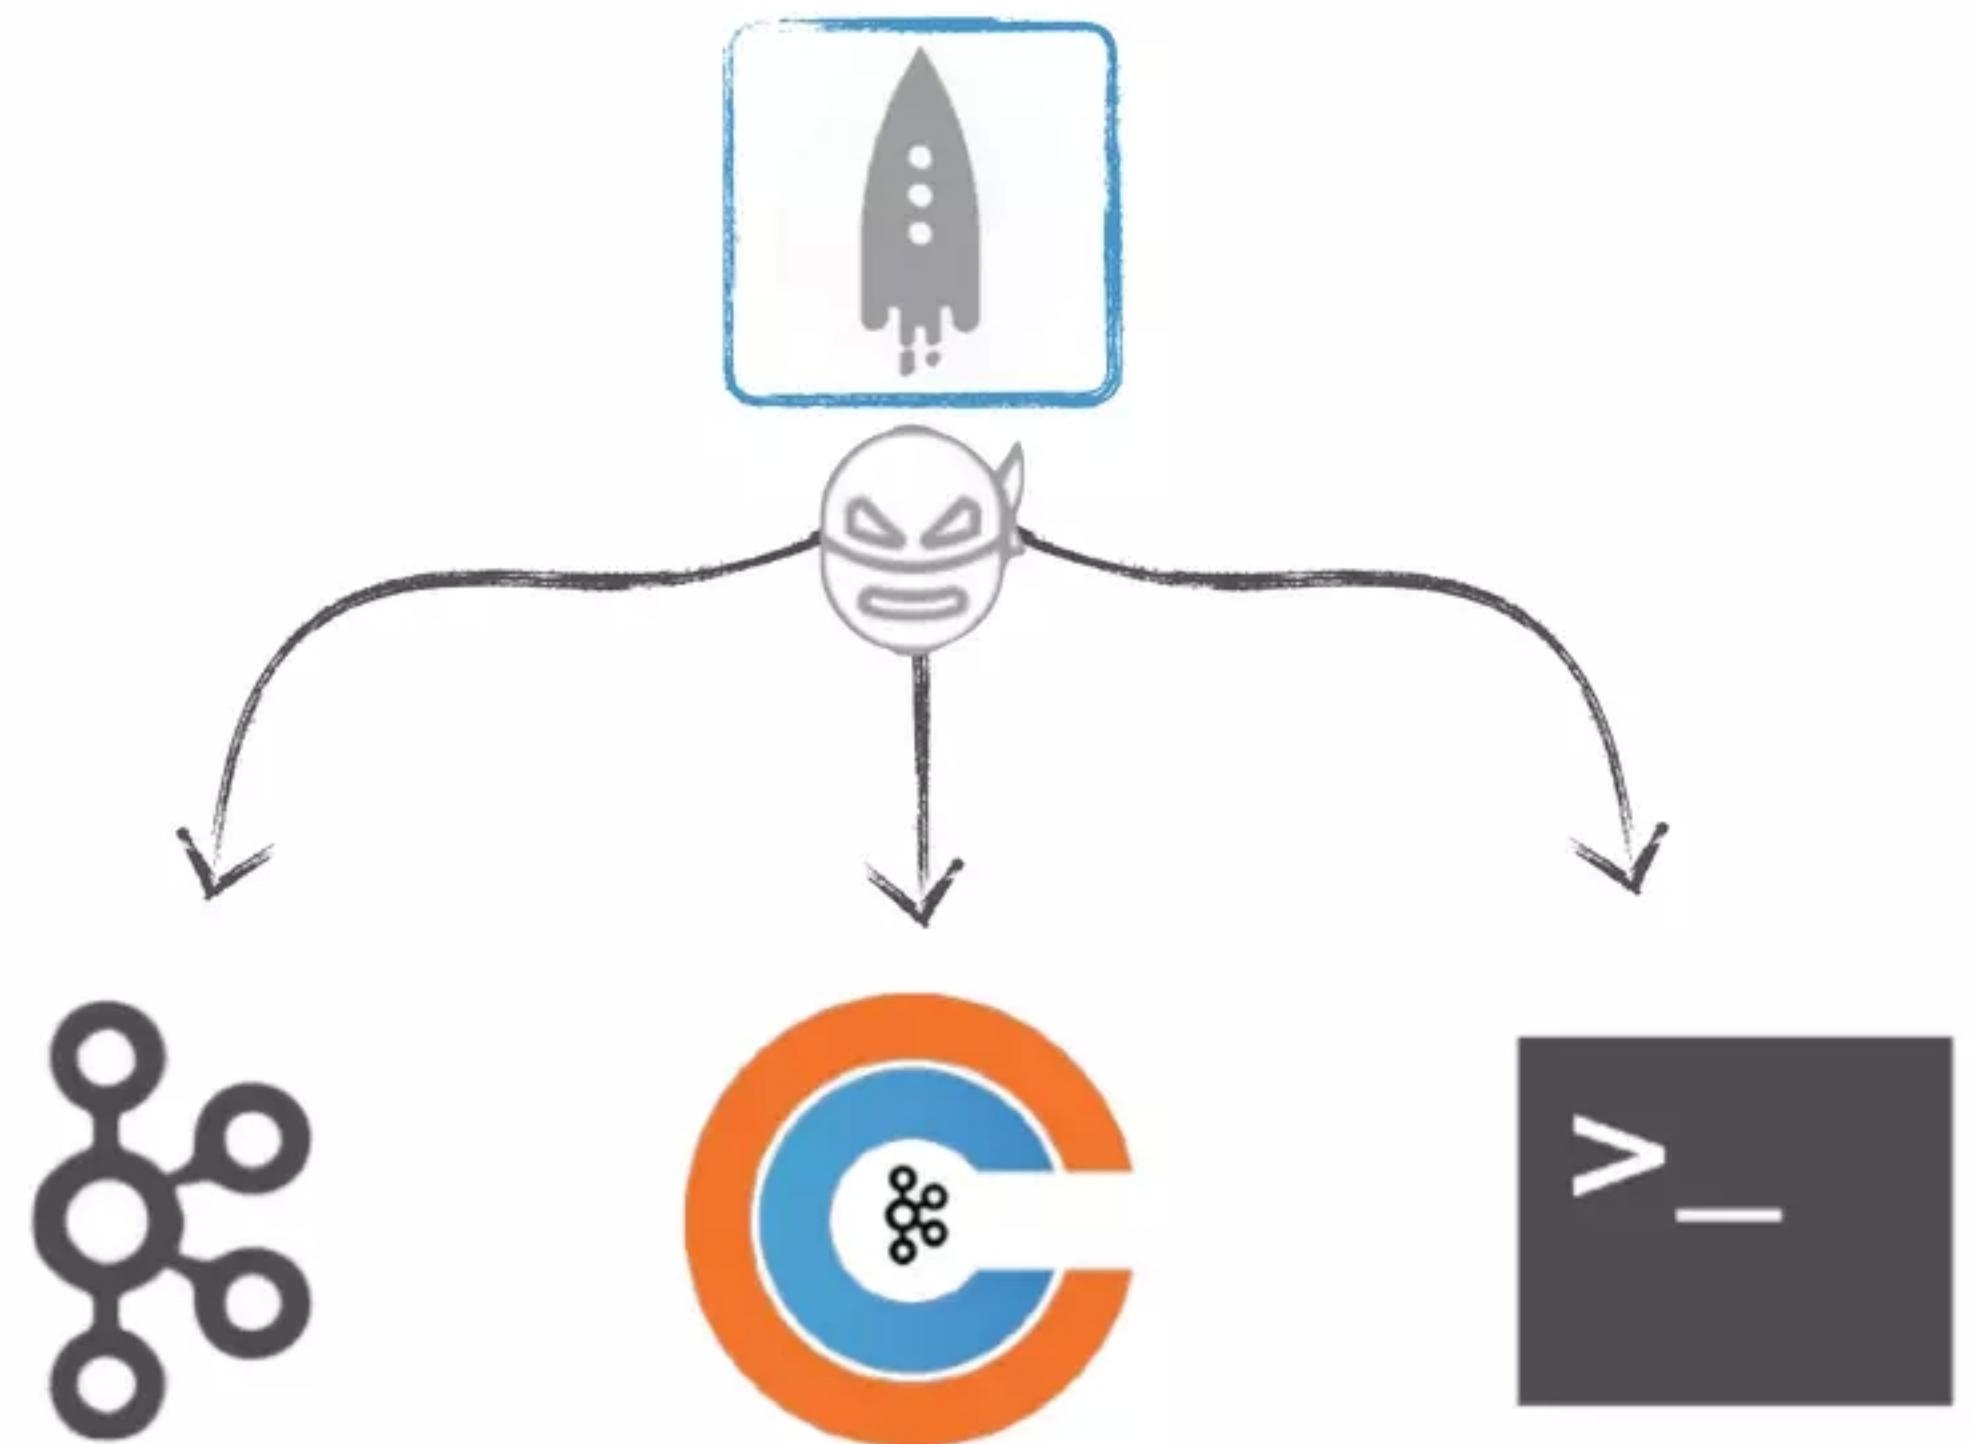












Encryption



see: <https://www.confluent.io/kafka-summit-ny19/ksql-and-security>

Encryption



Authentication



Encryption



Authentication



Authorization



- Intro
- Develop
- Deploy
- Operate
- **Common Mistakes**
- Q&A





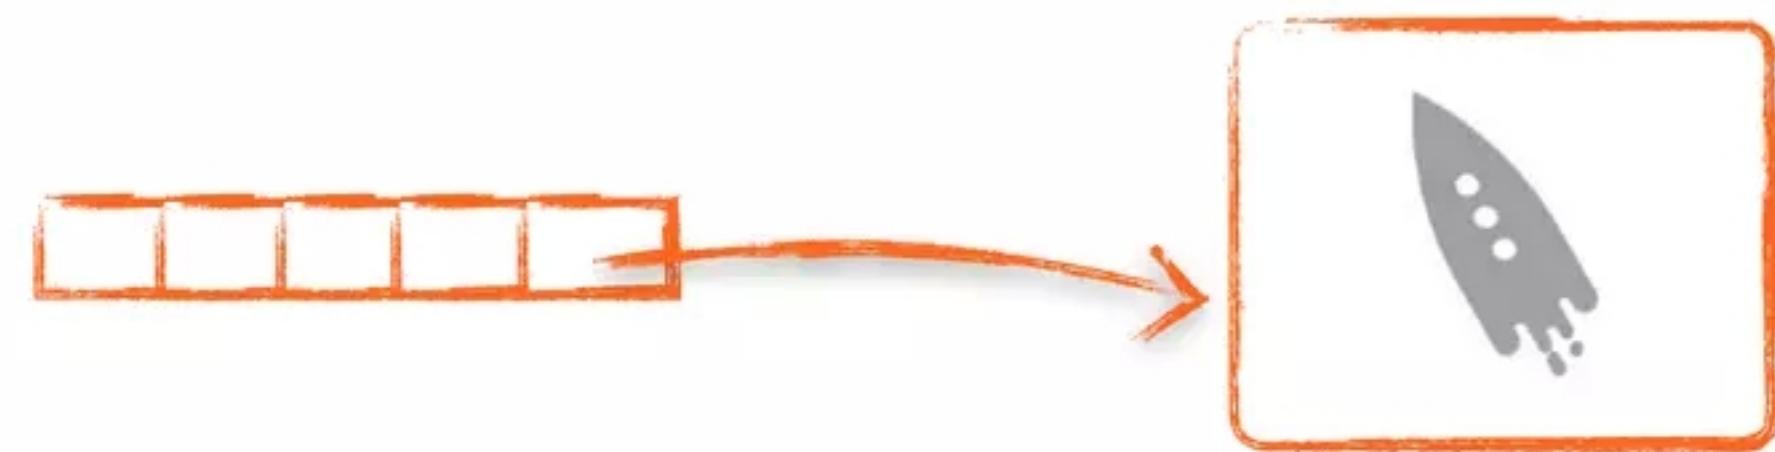
Queries Are Continuous



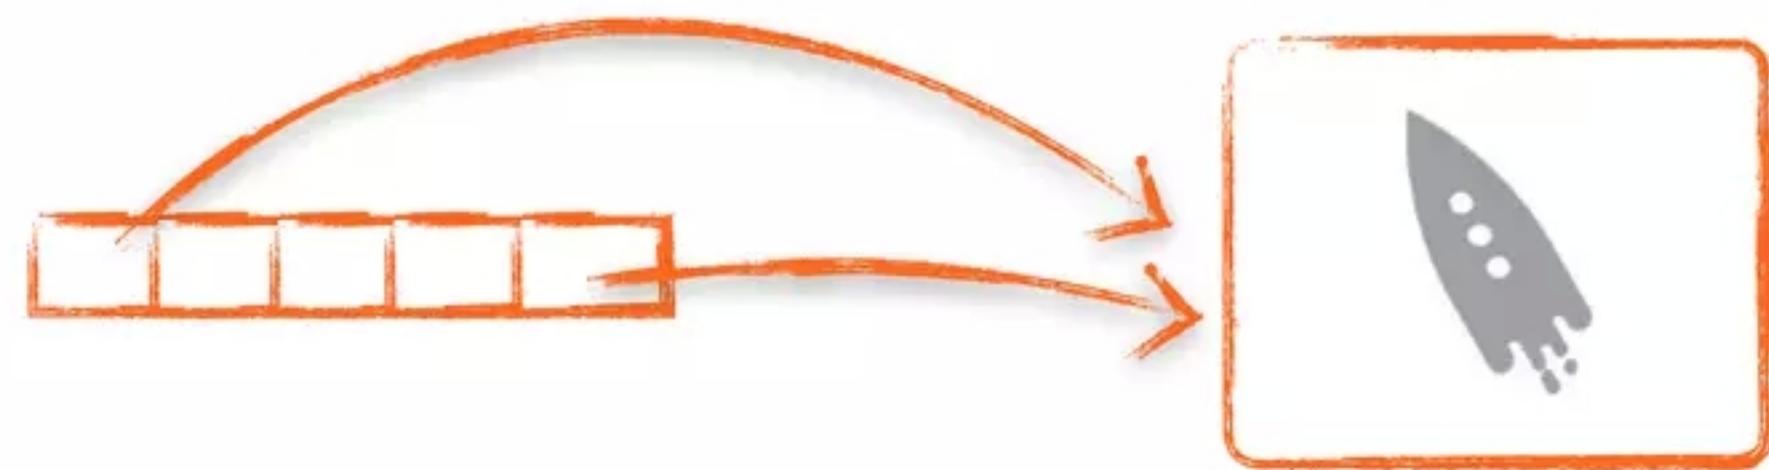
Queries Are Continuous



new queries read from most recent by default



new queries read from most recent by default



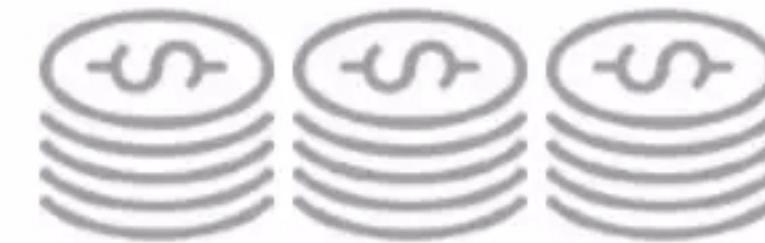
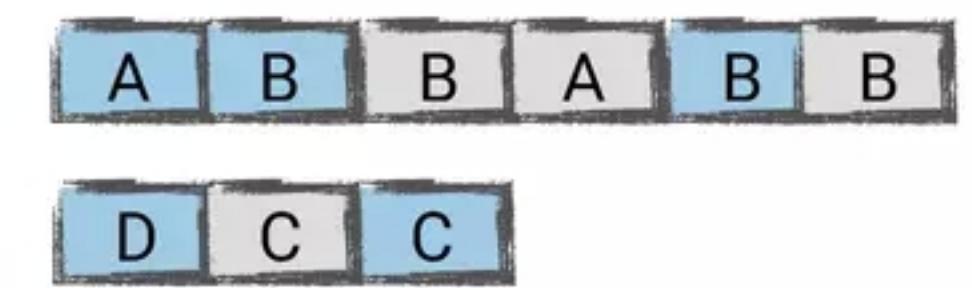
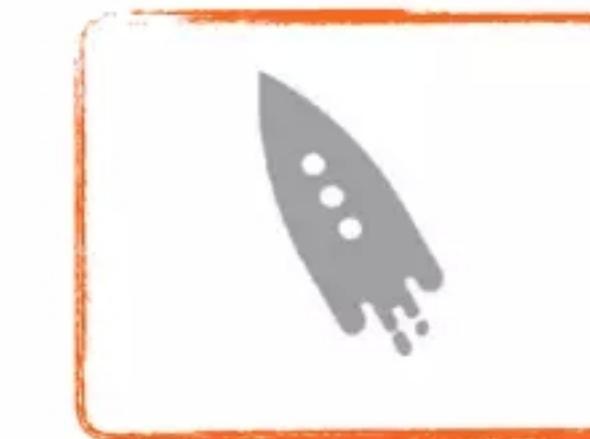
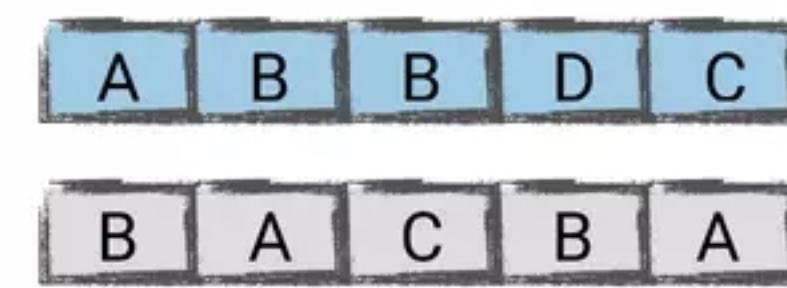
```
SET 'auto.offset.reset' = 'earliest'
```

joins may cause
repartition

A	B	B	D	C
B	A	C	B	A



joins may cause
repartition



repartition does not
preserve order across
partitions...

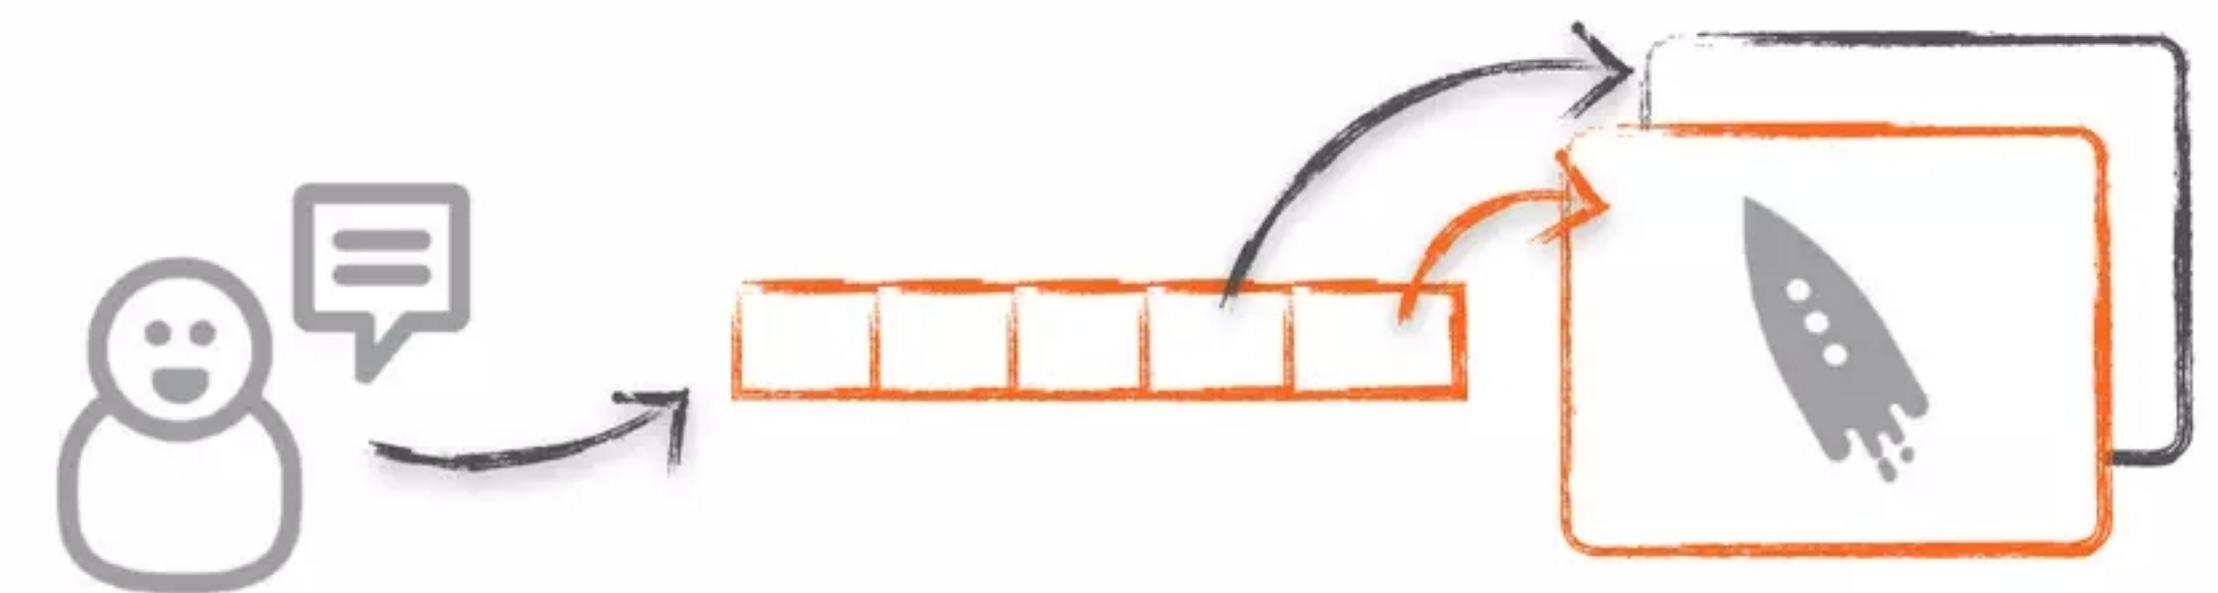
repartition does not
preserve order across
partitions...

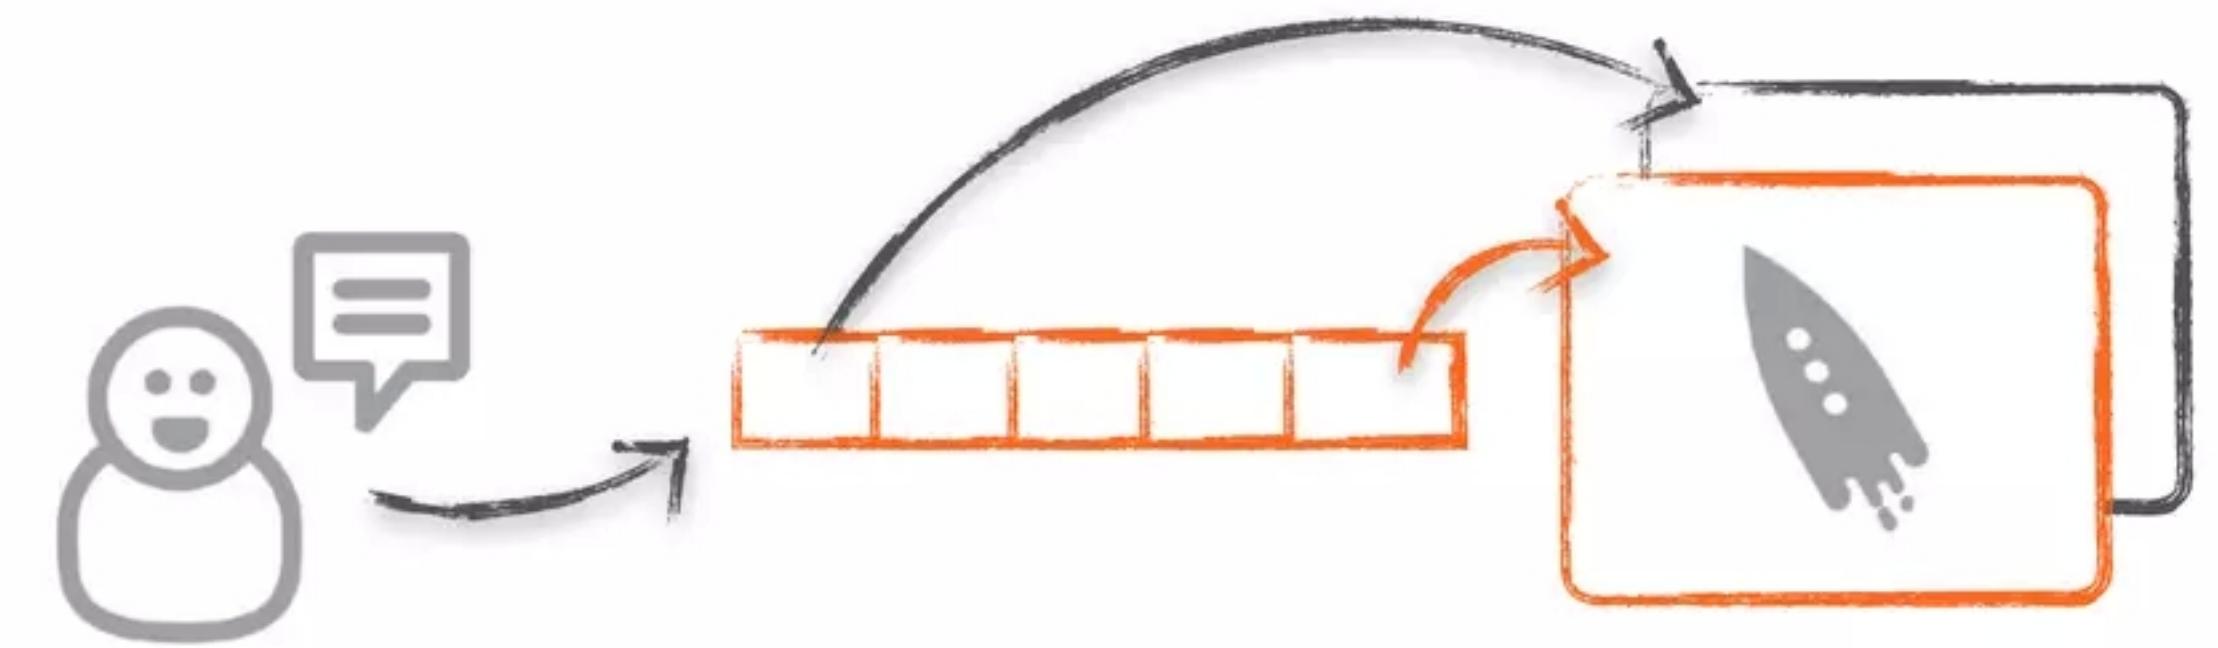


repartition does not
preserve order across
partitions...

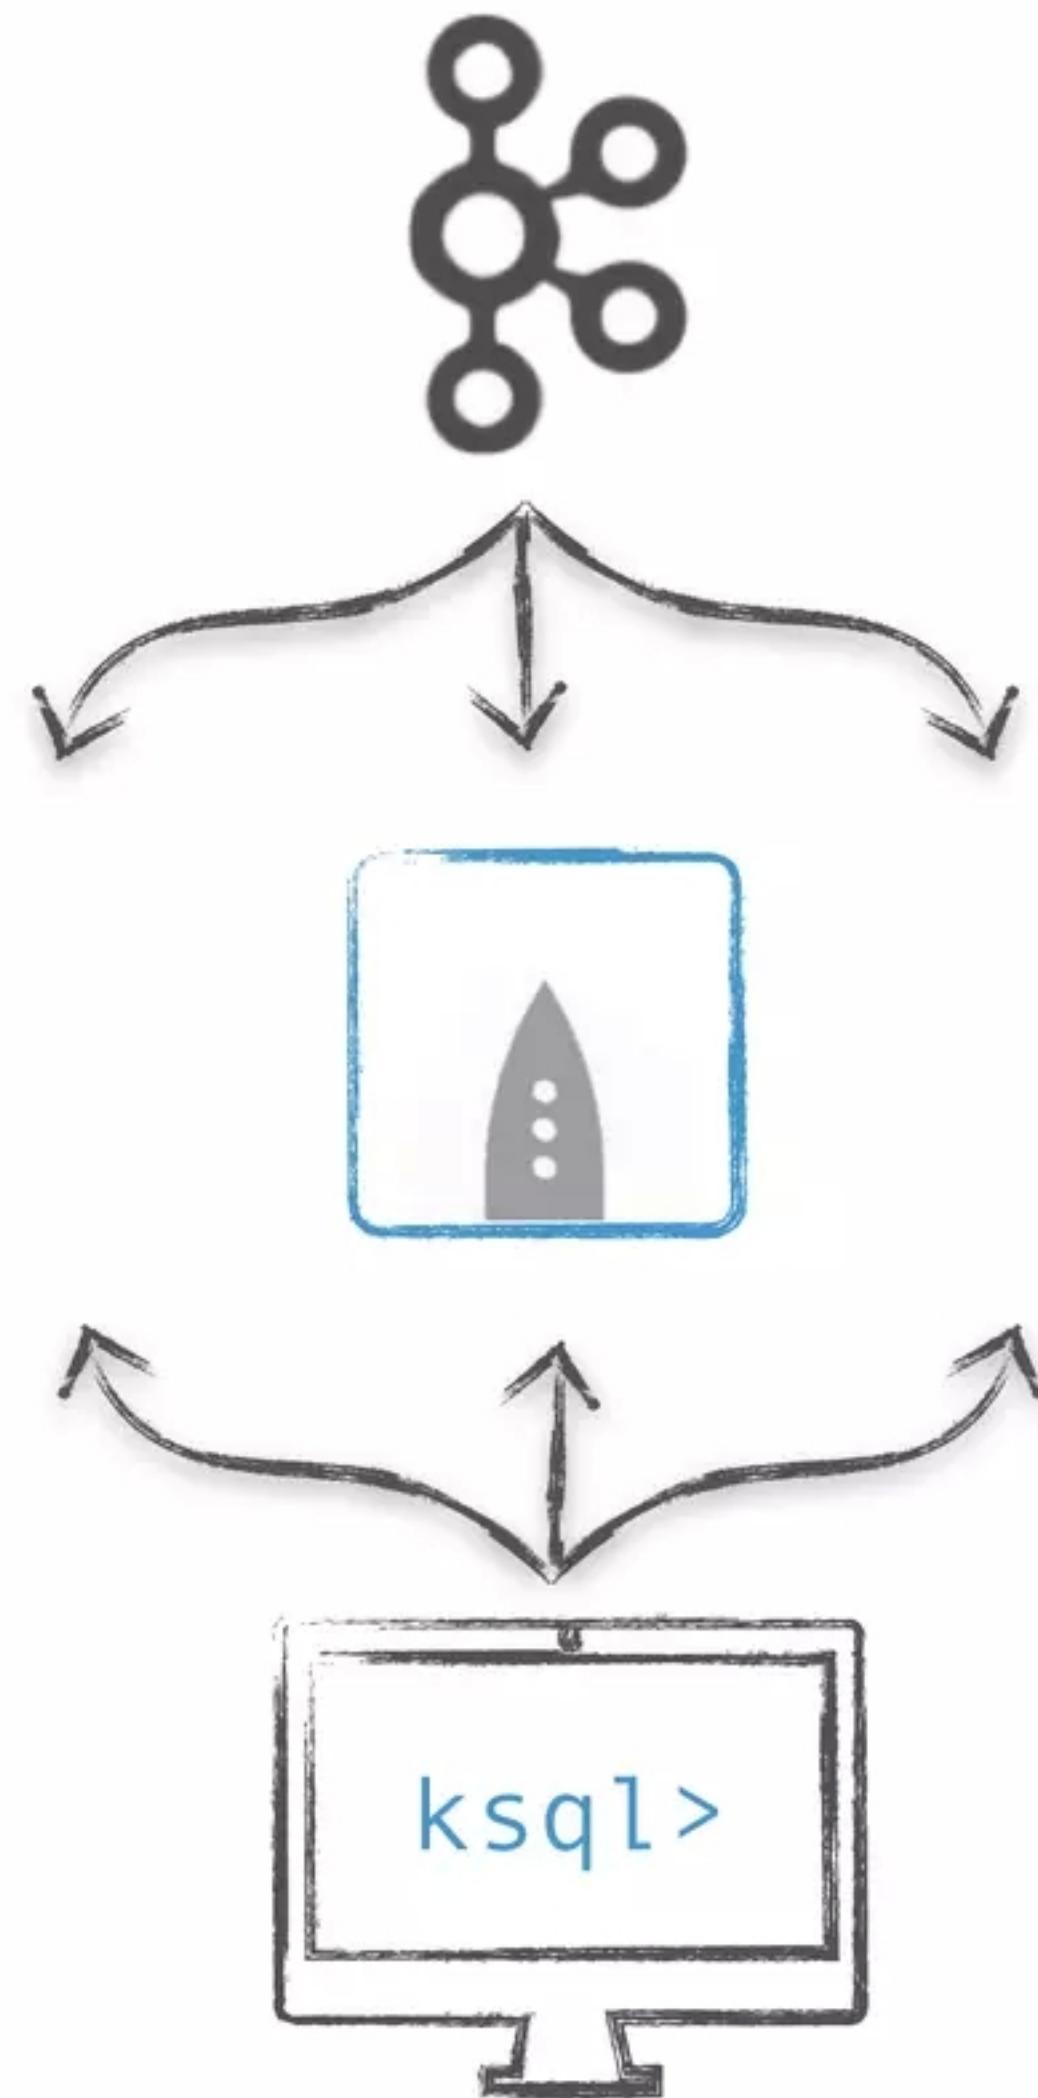


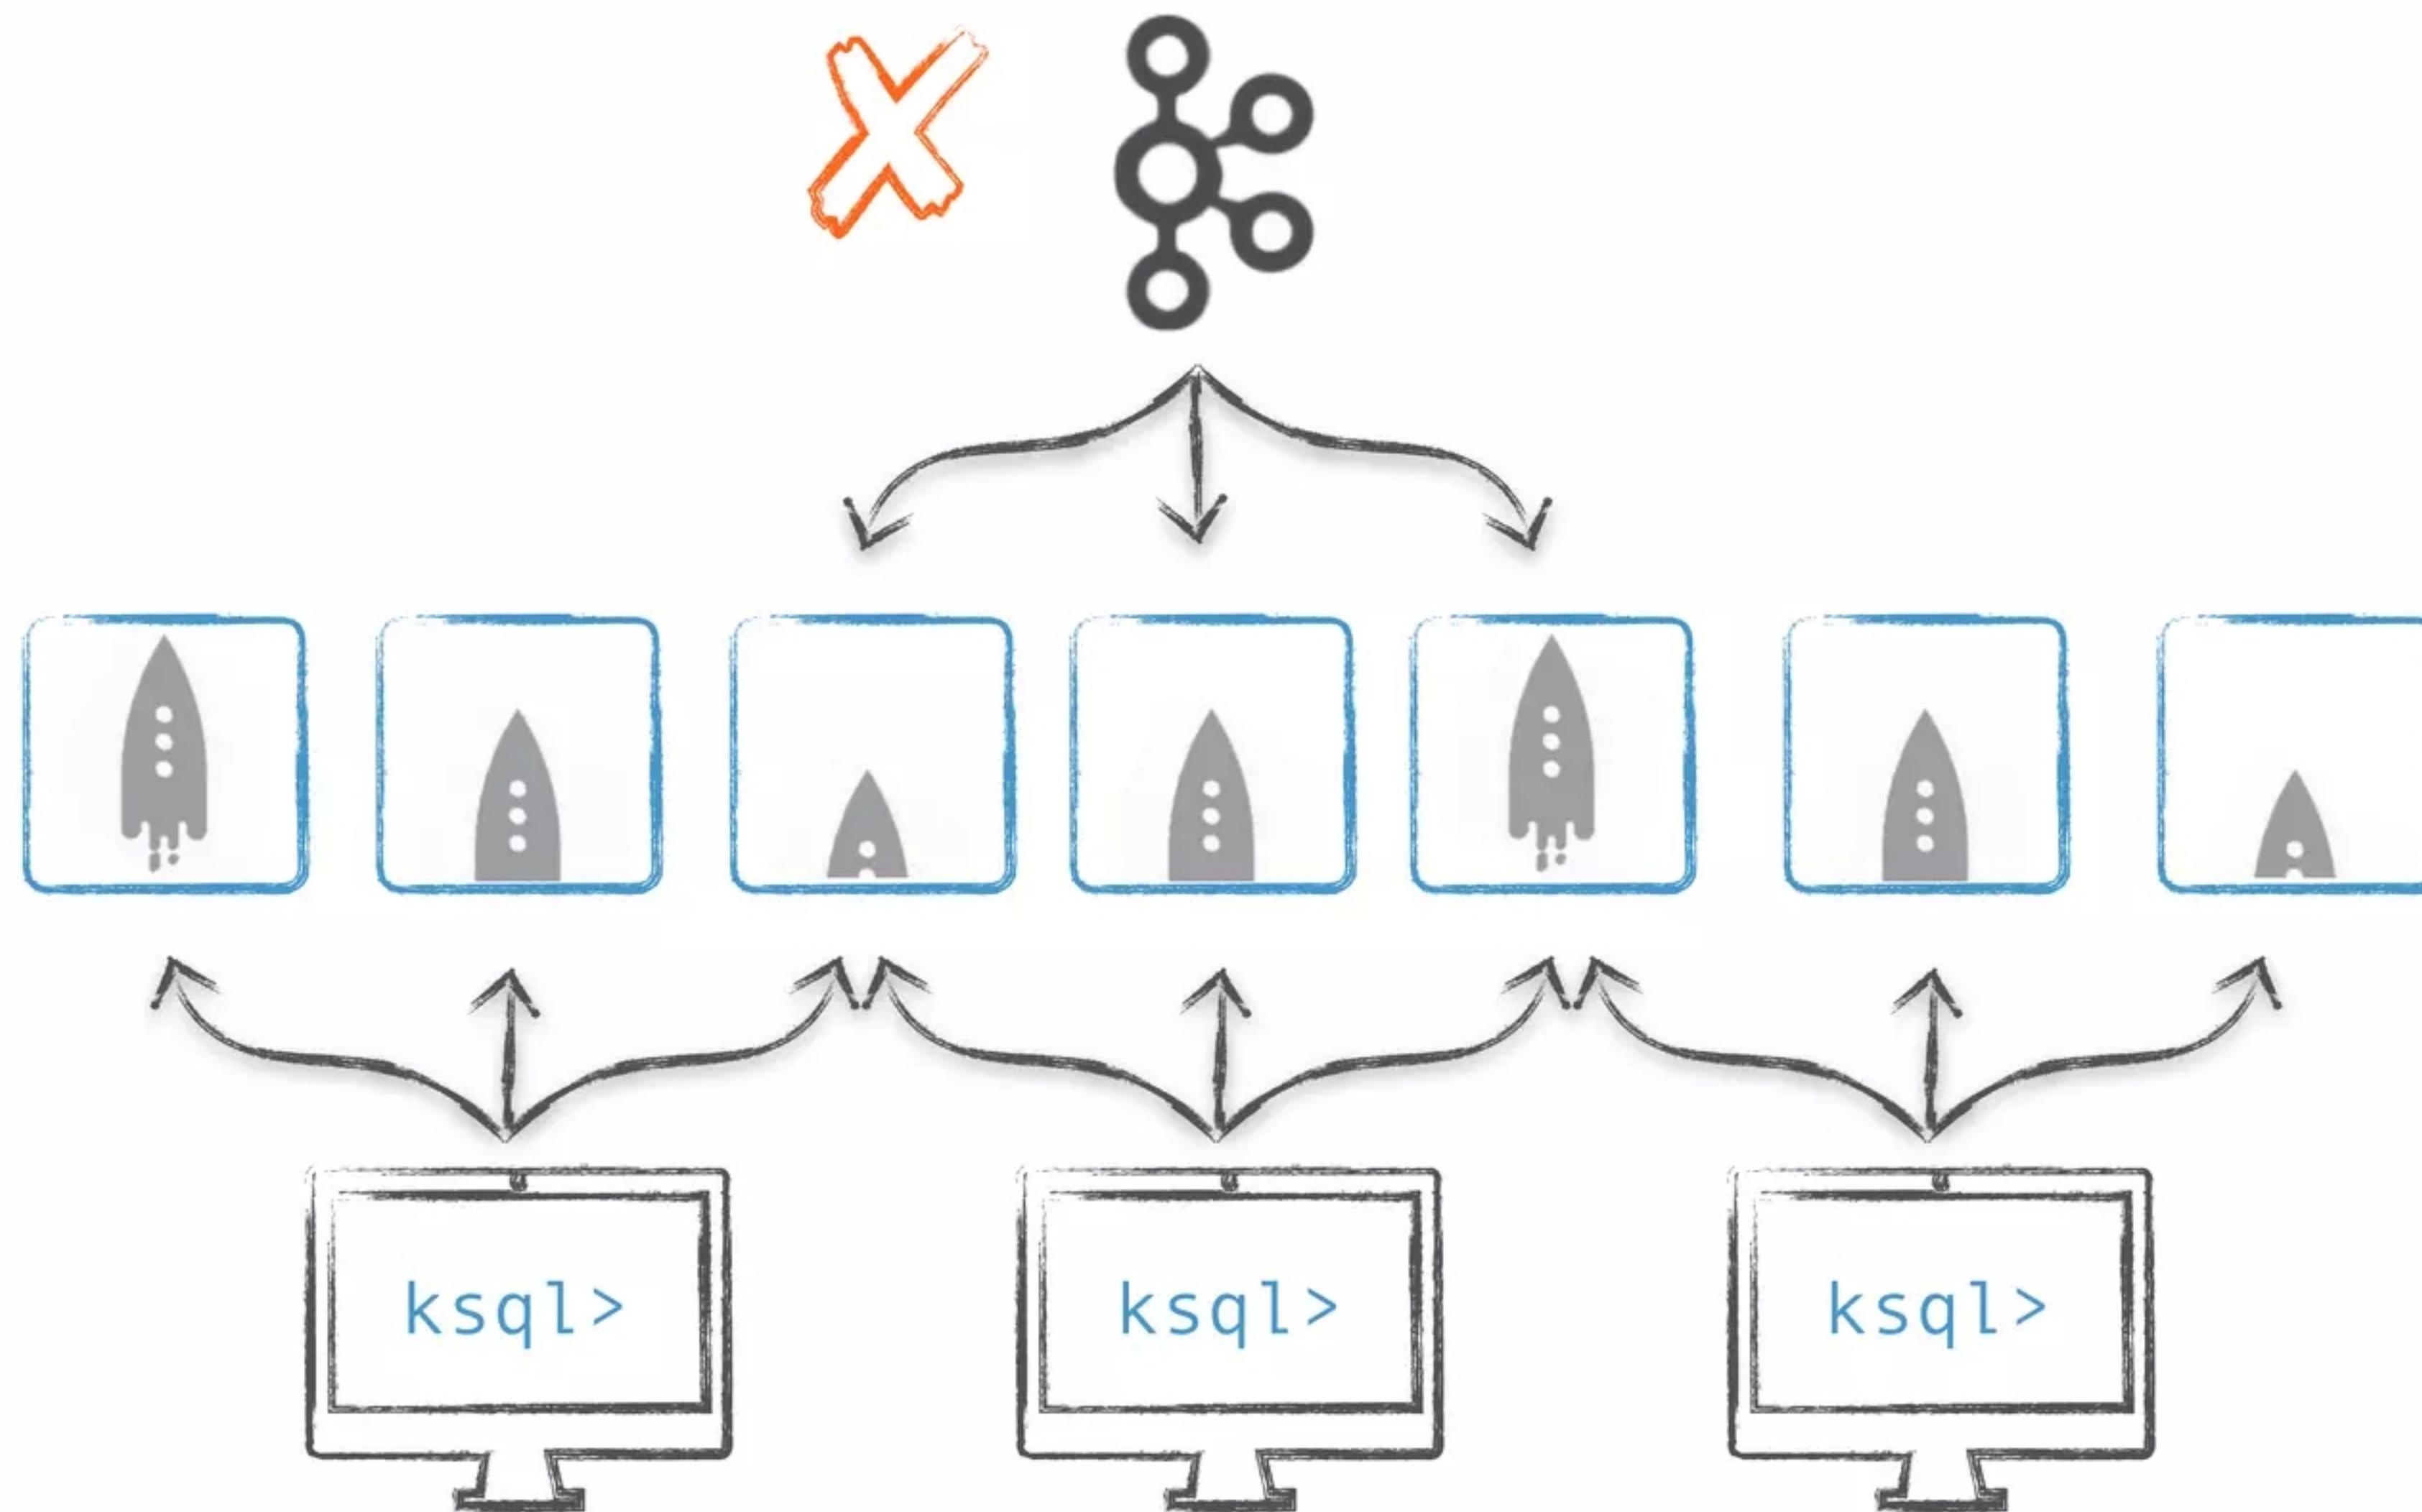
so tables cannot be
rekeyed

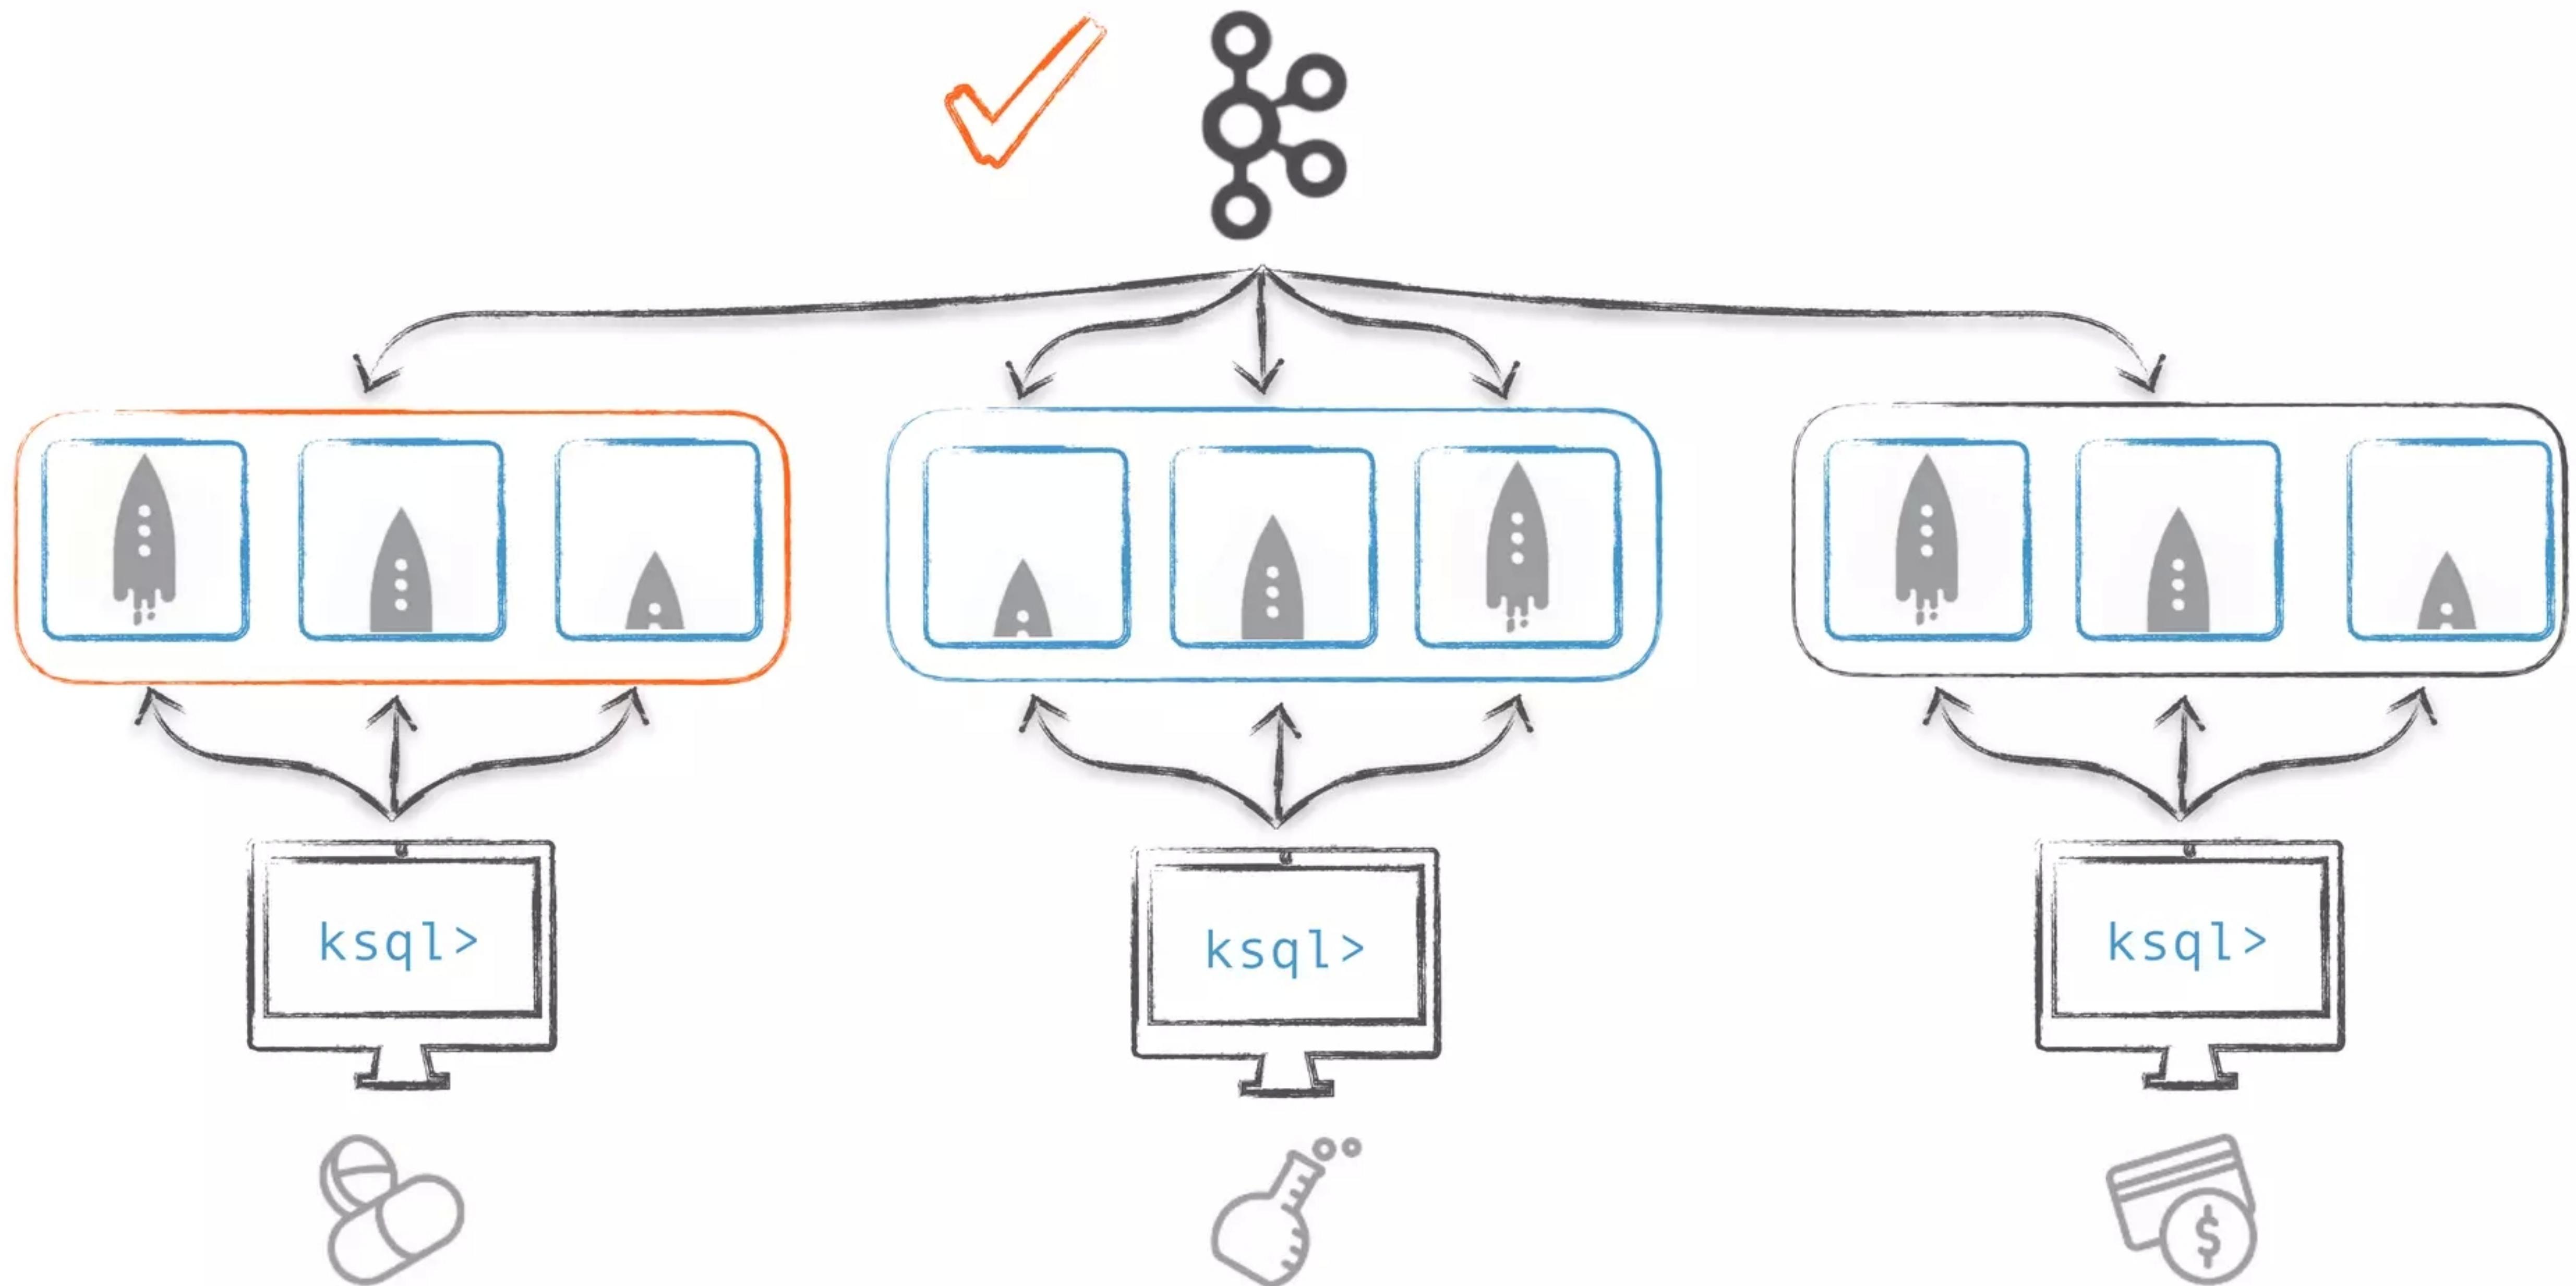




crashes replay (live query) state







- Intro
- Develop
- Deploy
- Operate
- Common Mistakes
- **Q&A**





 [in/agavra](https://www.linkedin.com/in/agavra)

 almog@confluent.io