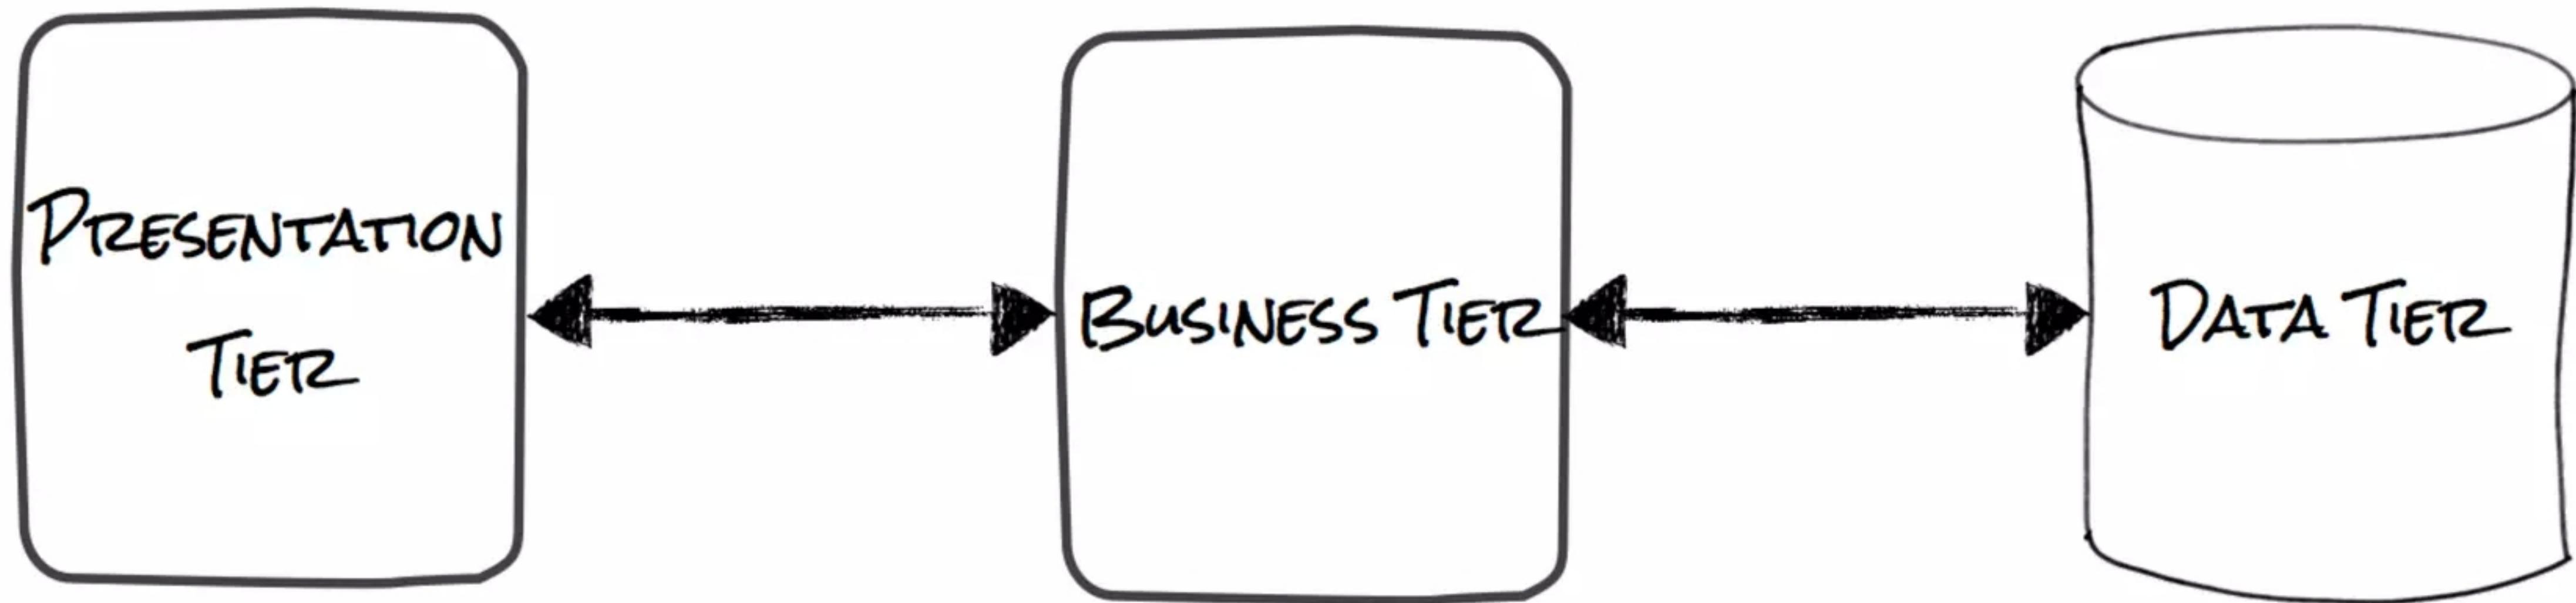


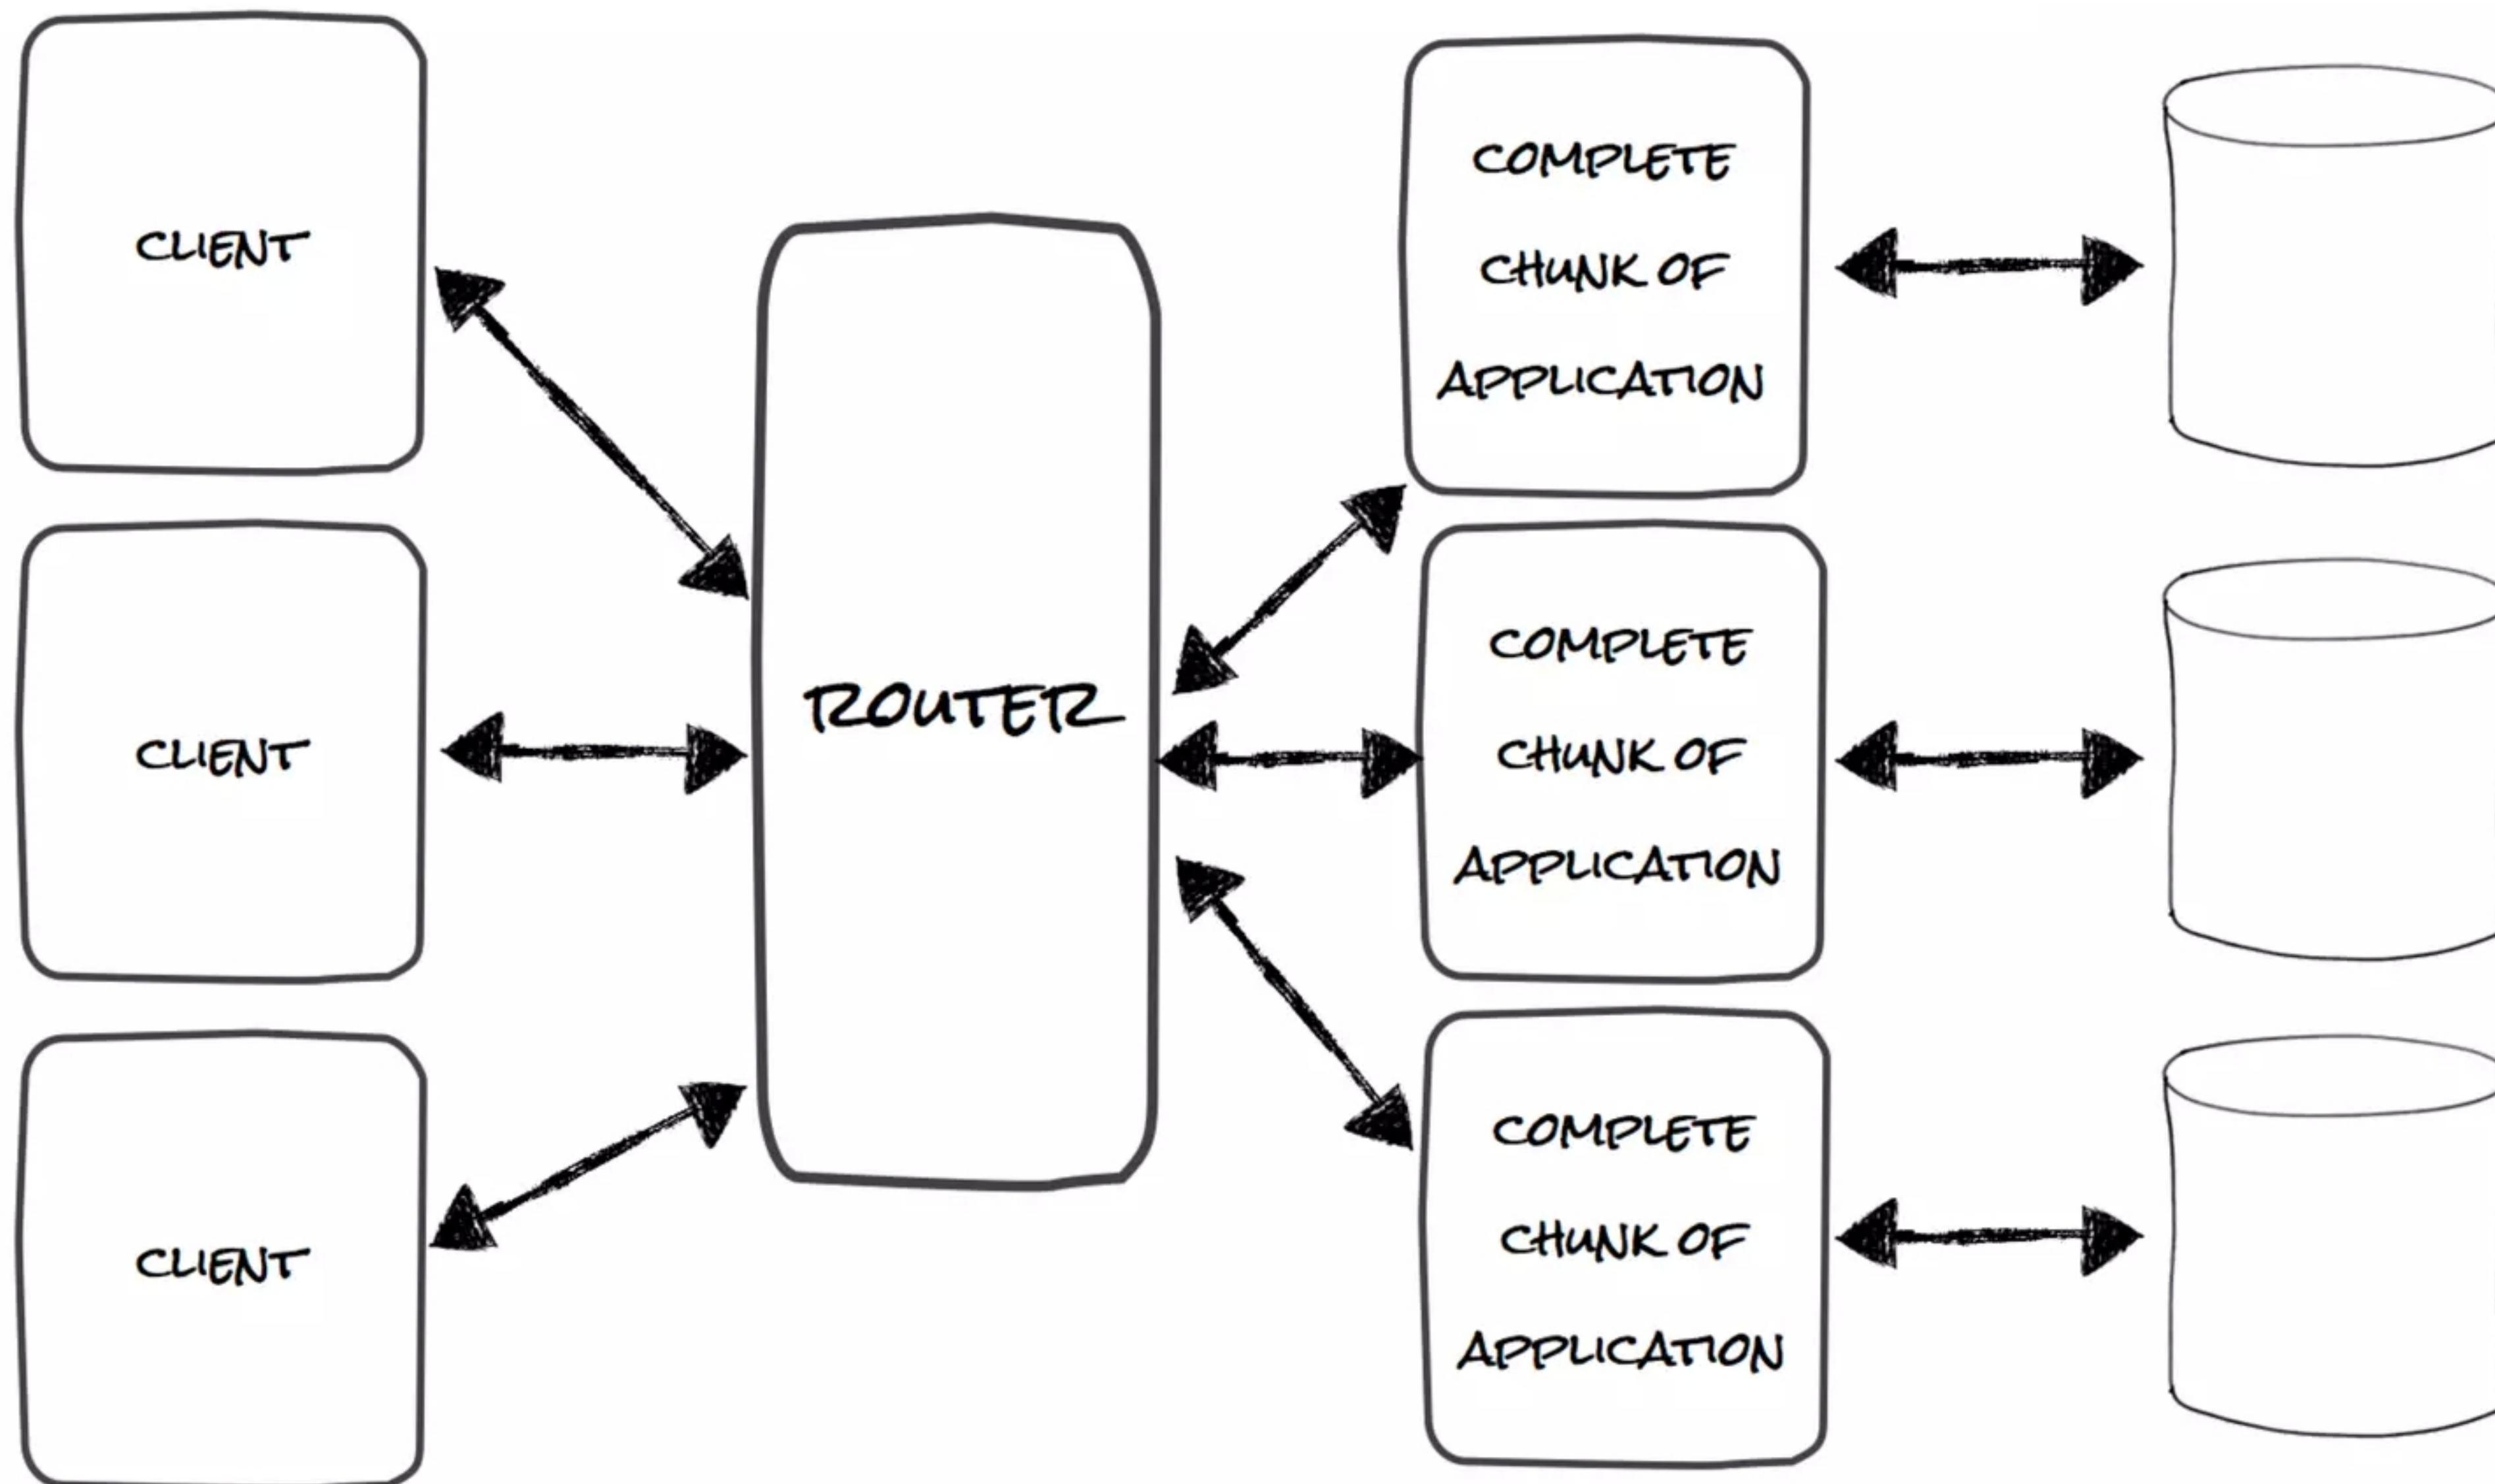
# Four Streaming Data Architectures

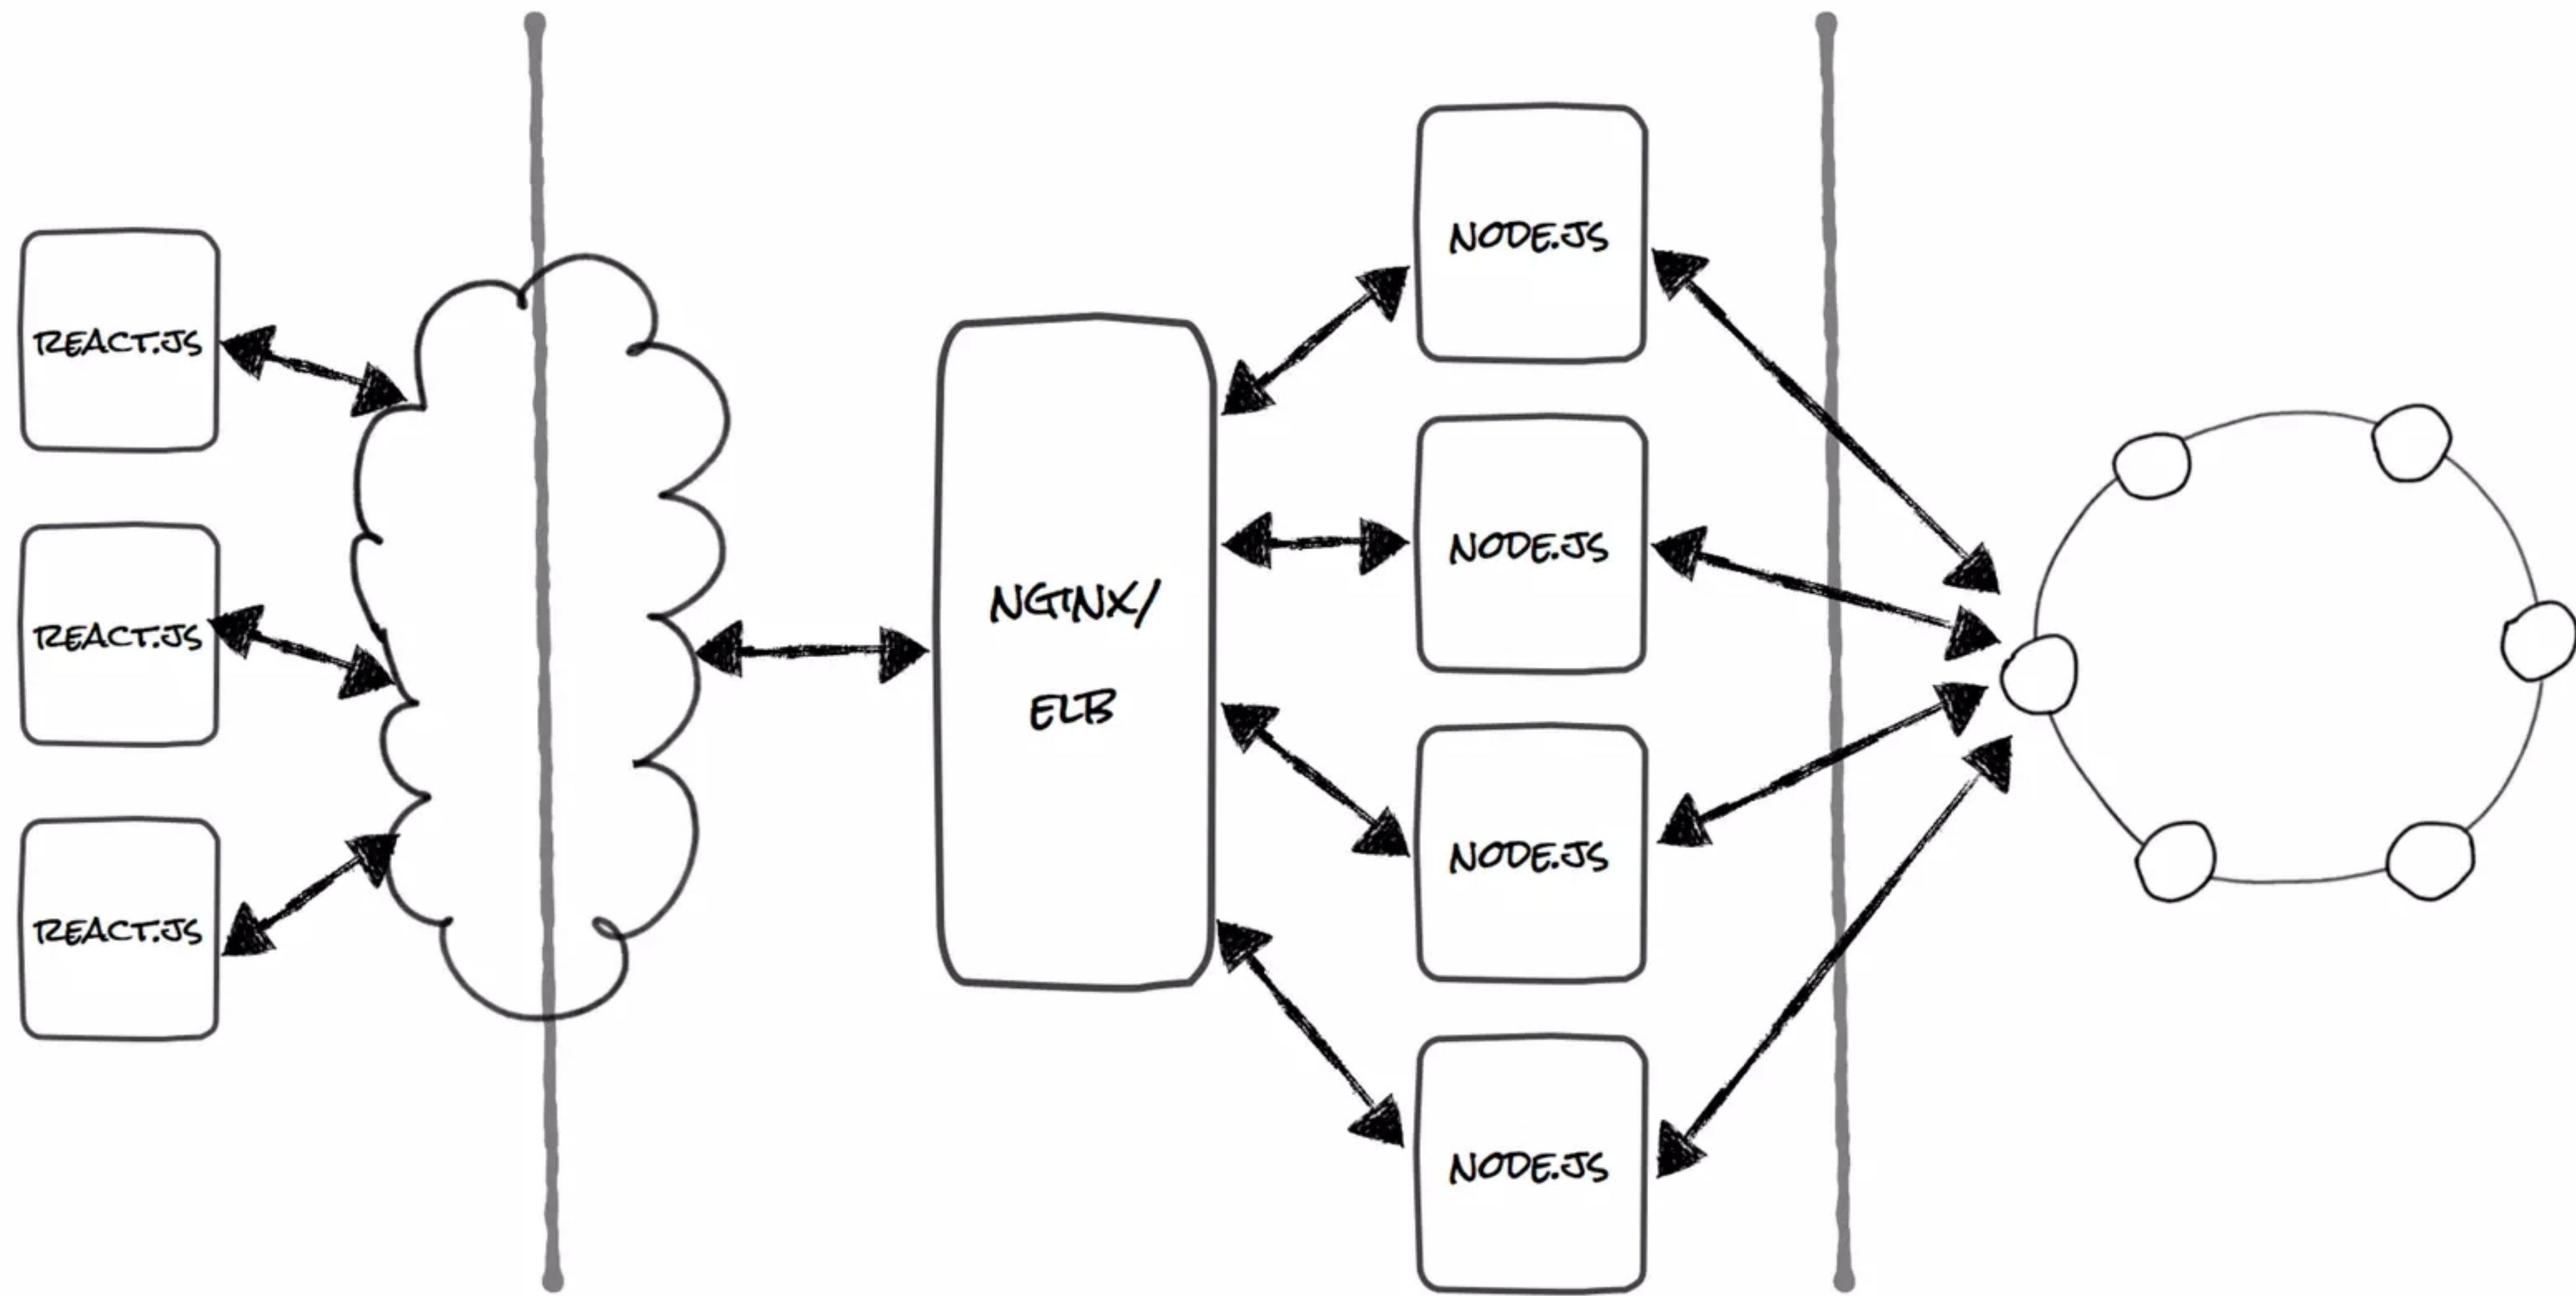
@tlberglund

---

# The Old Story









---

# The Streaming Story



Maria-José Viñas, NASA Earth Science News Team.



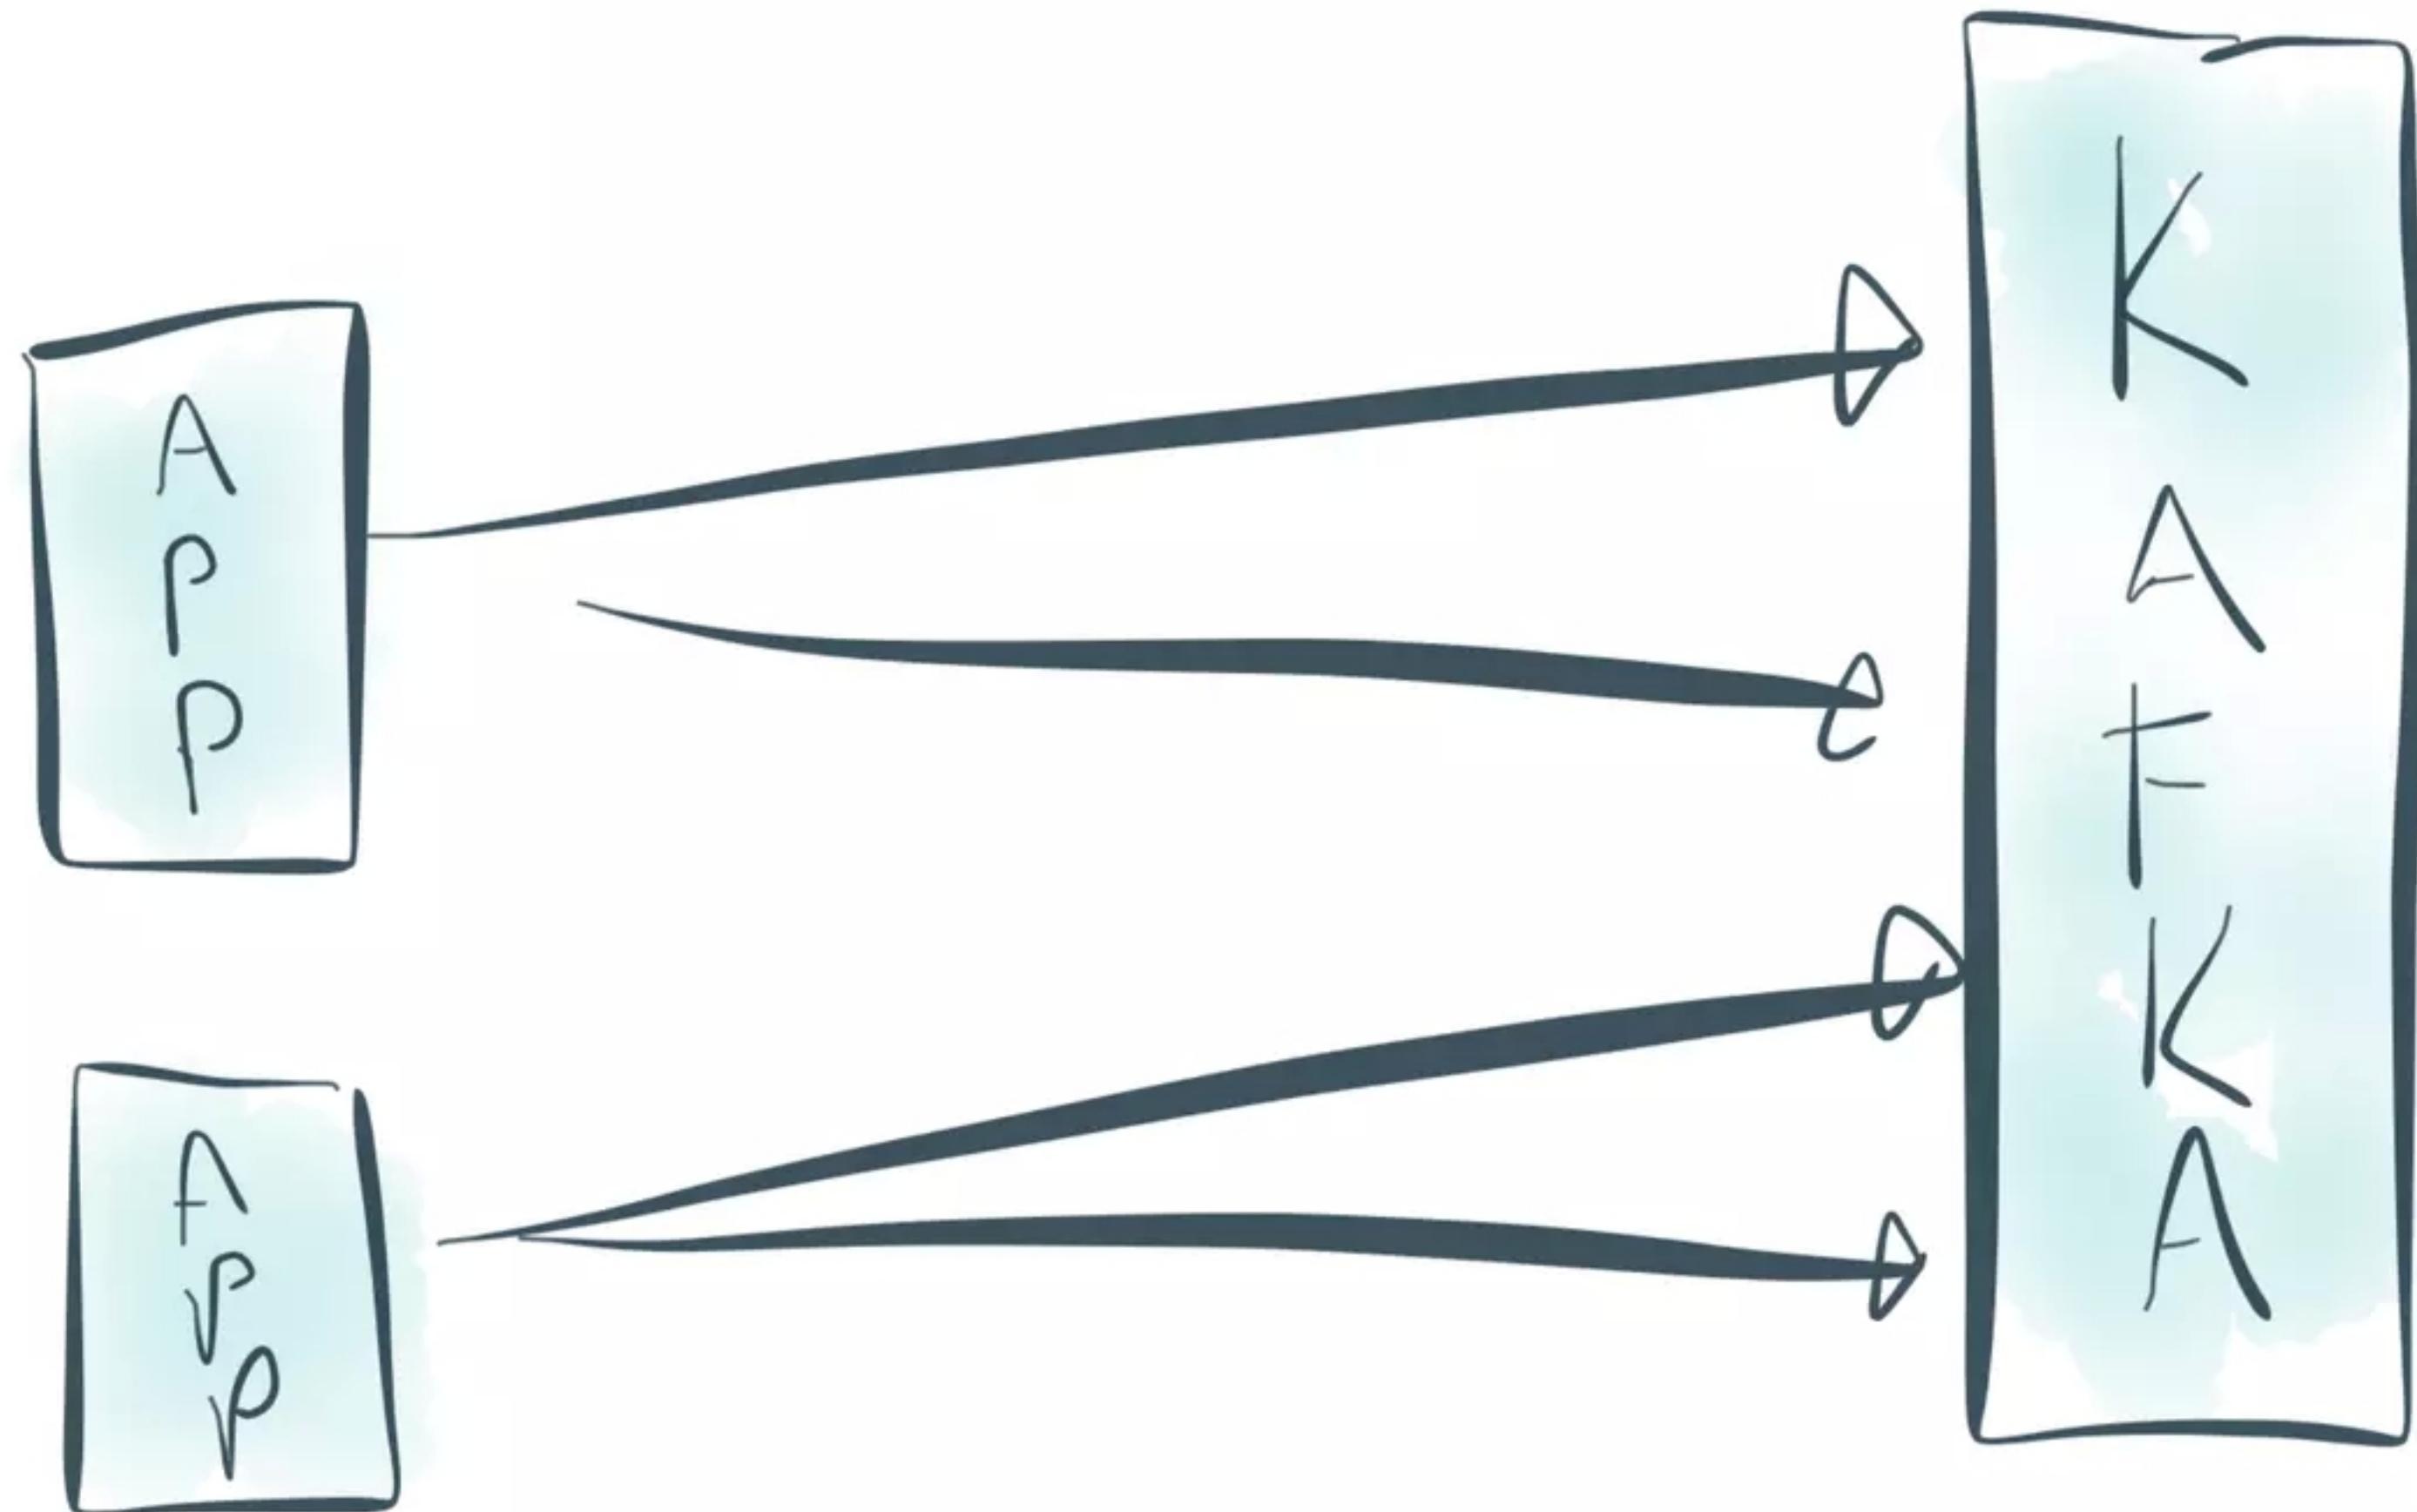
Image credit: NASA/JPL/UCSD/JSC



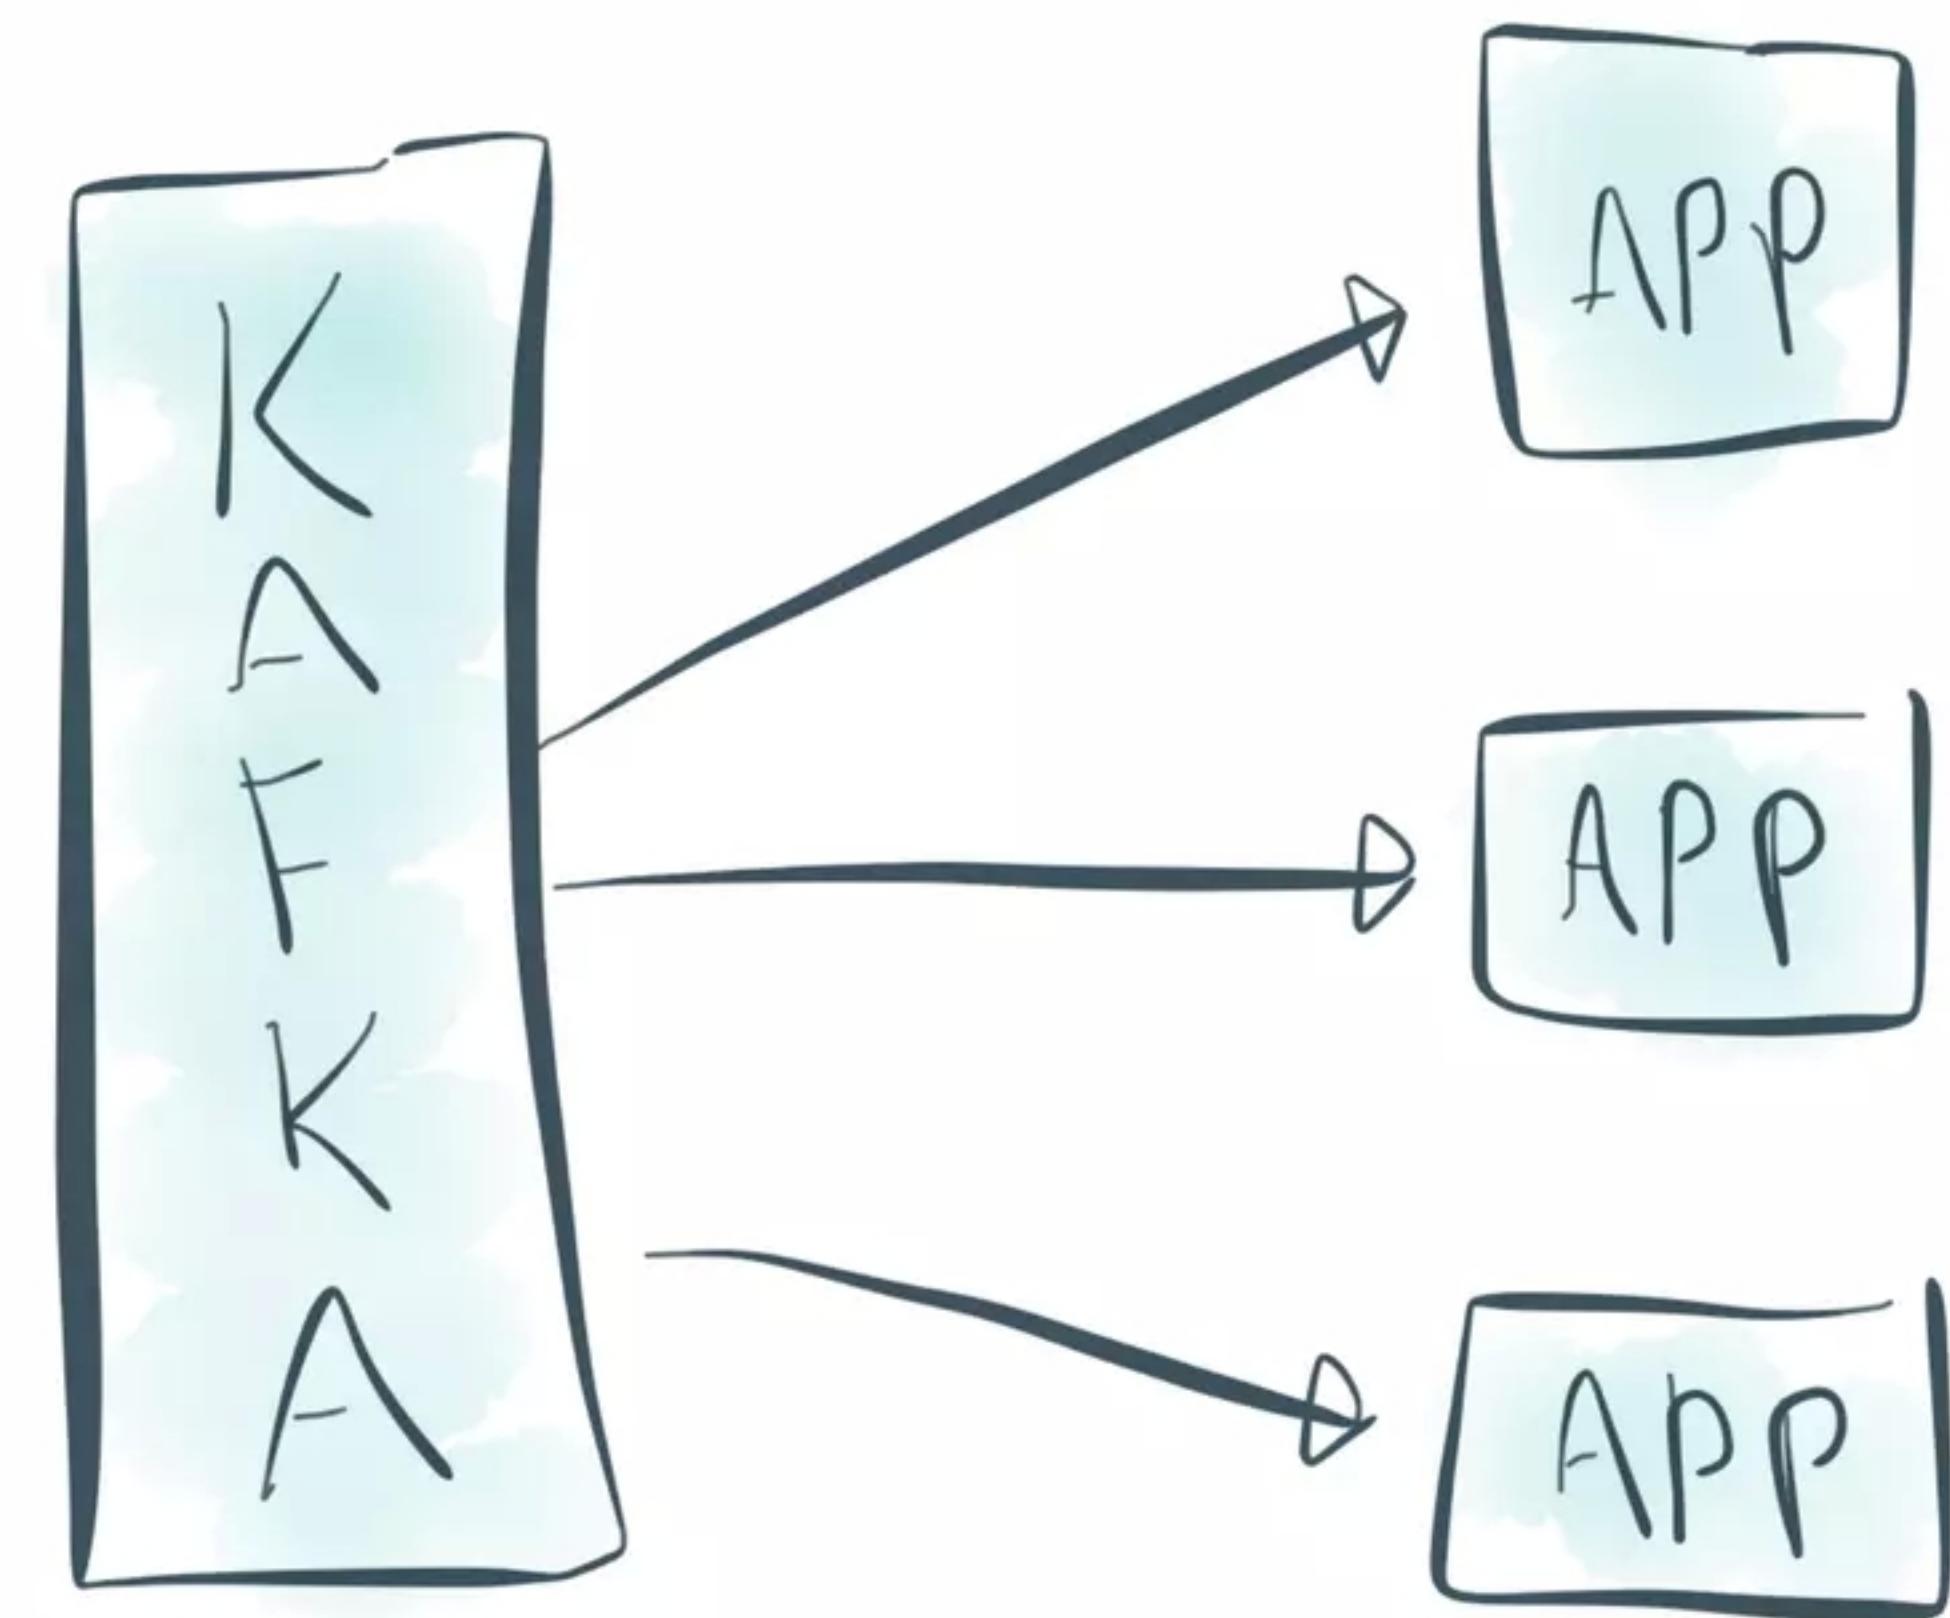
---

# Kafka

# PRODUCER

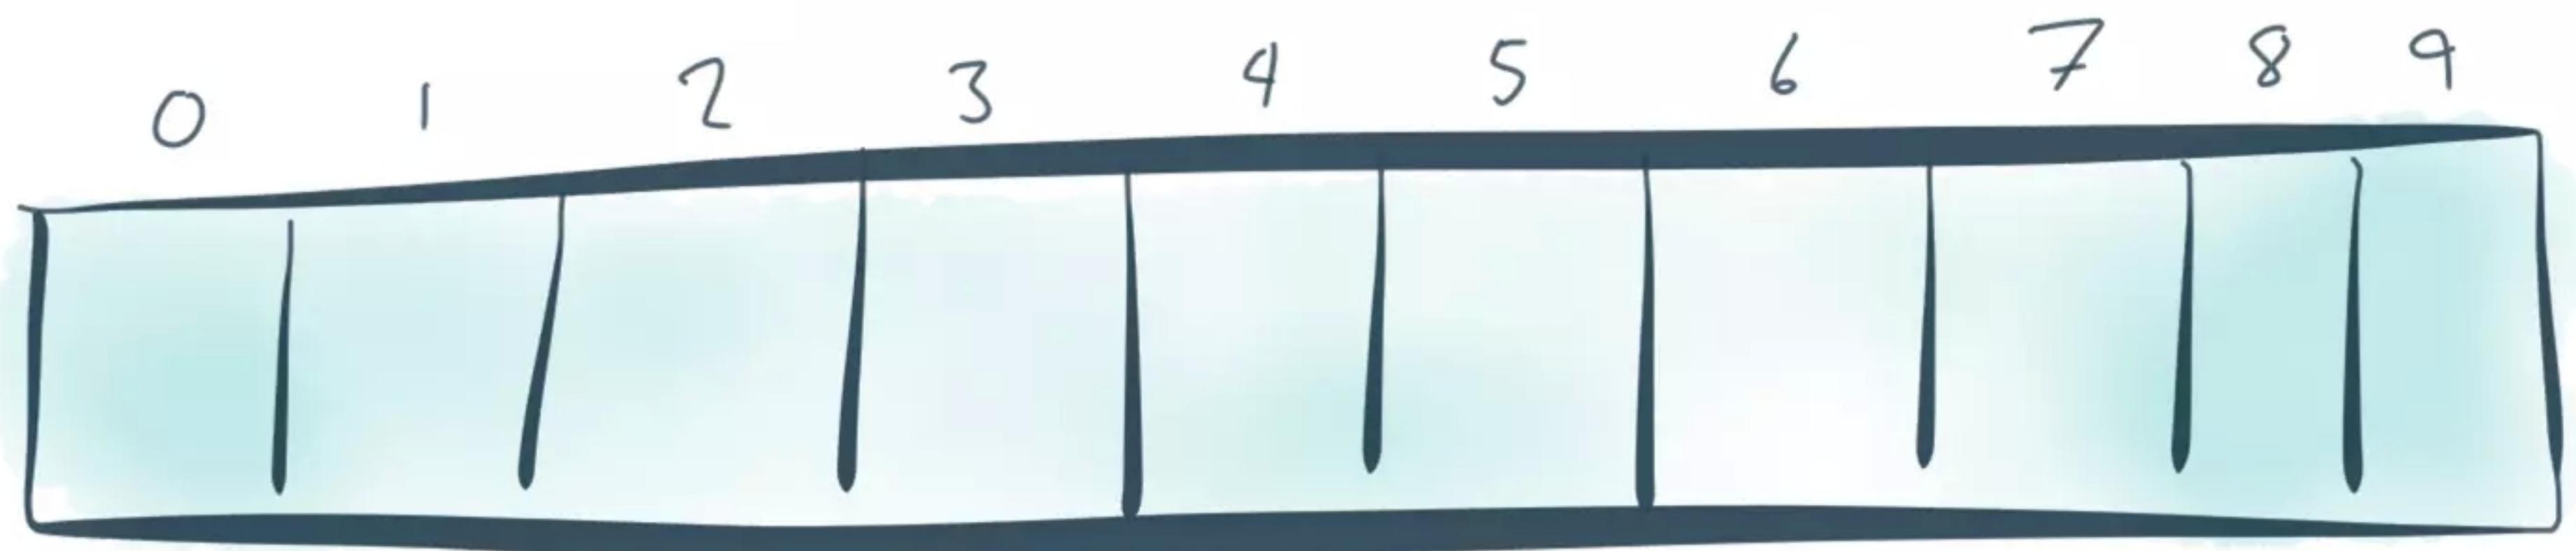


# CONSUMER



LOGS

NEXT  
WRITE



---

# Examples

---

# Yelp

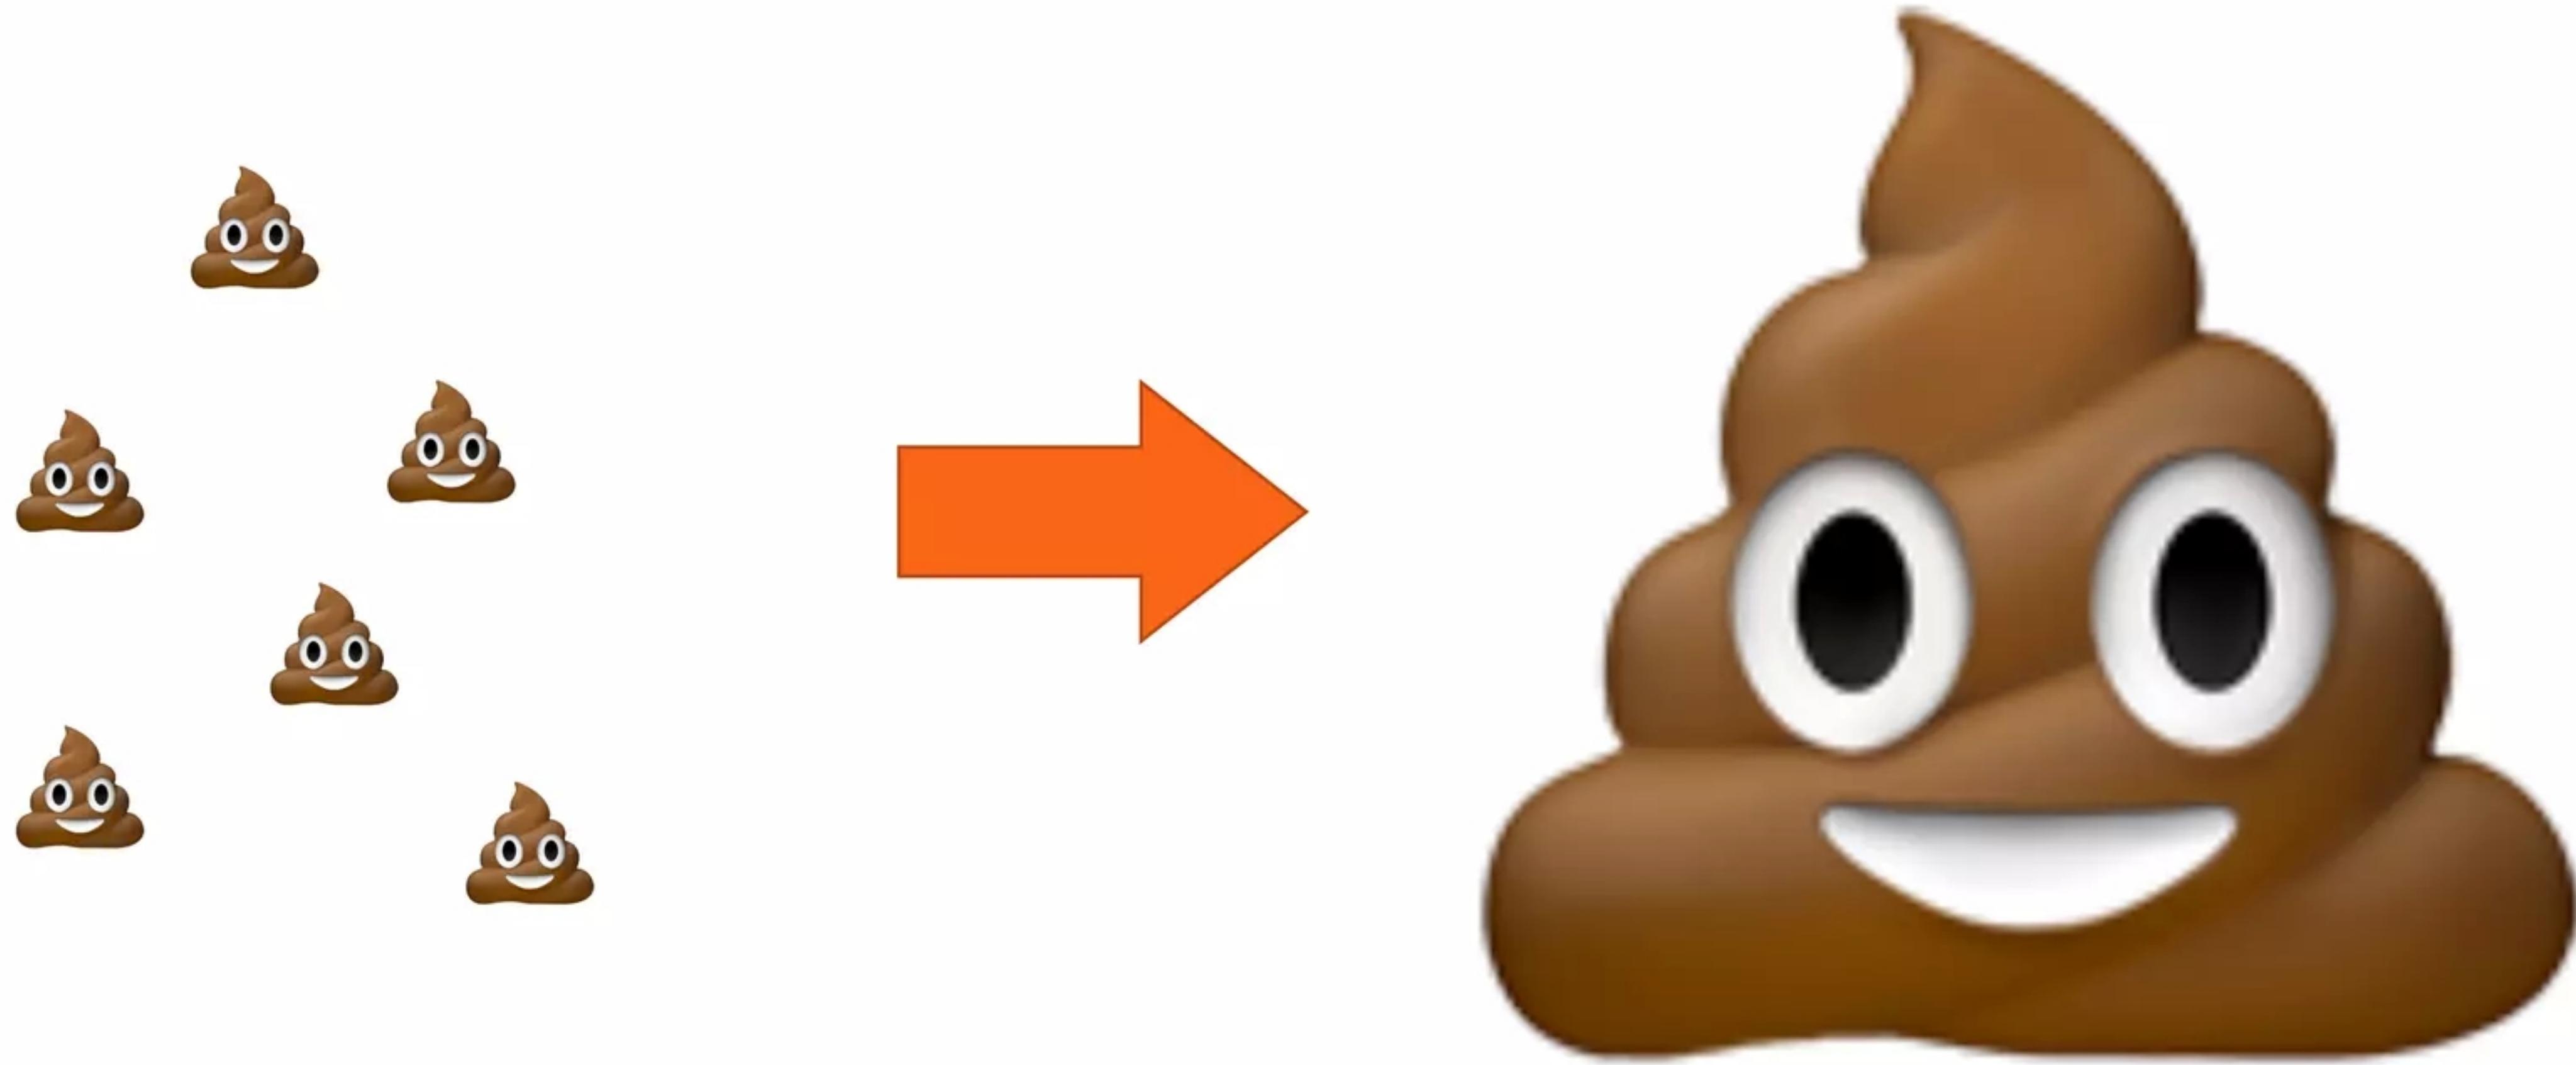
<https://engineeringblog.yelp.com/2016/07/billions-of-messages-a-day-yelps-real-time-data-pipeline.html>

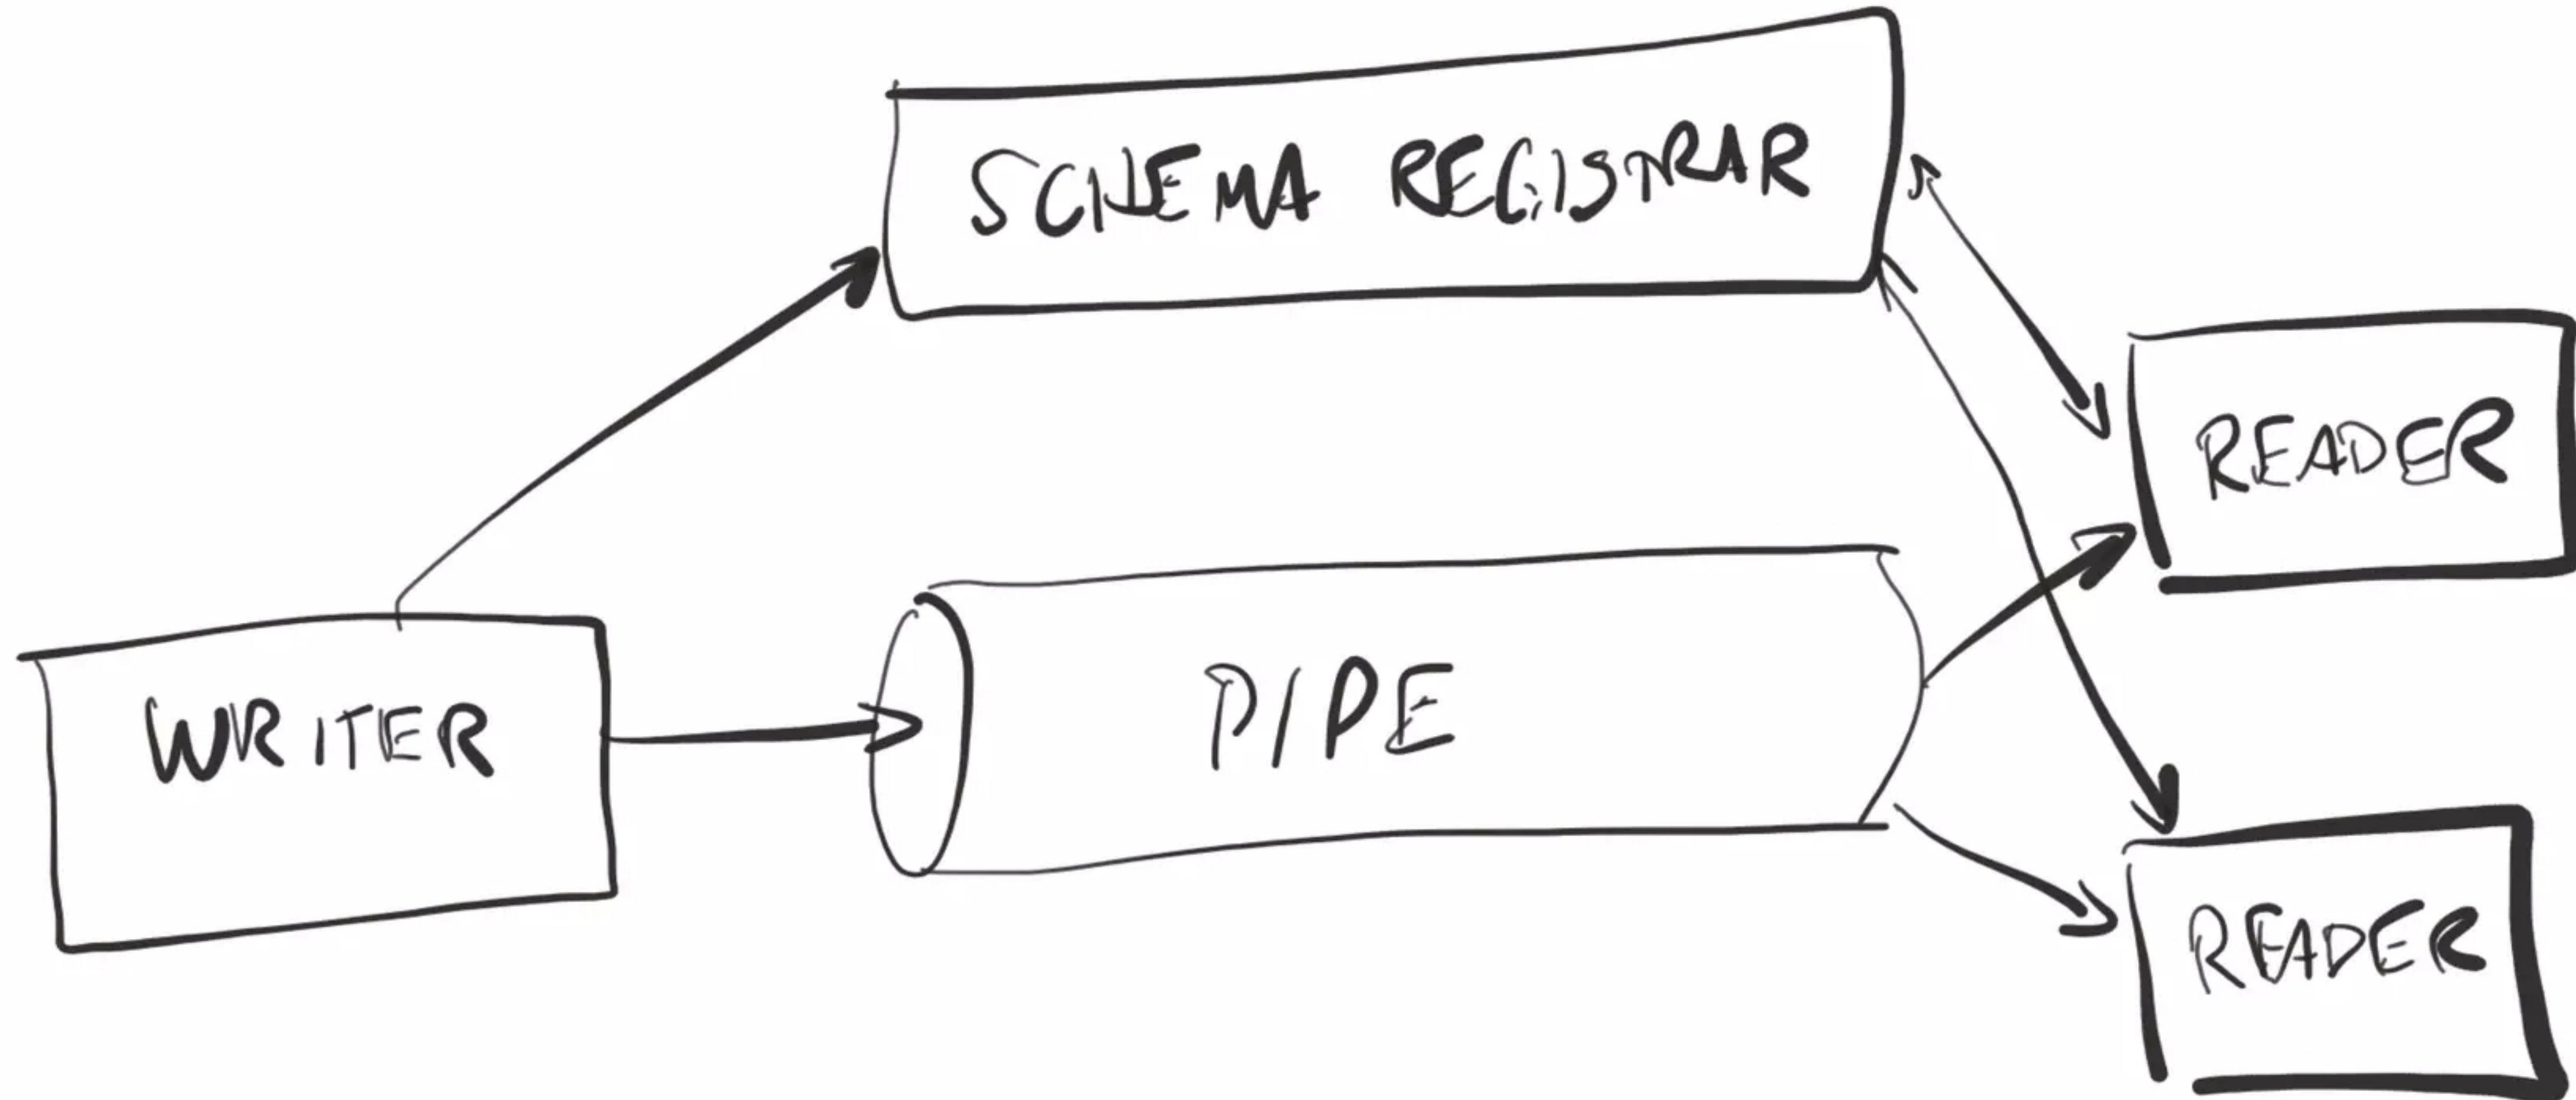
Justin C., Software Engineer  
(May, 2016)

# A Streaming Backbone

---

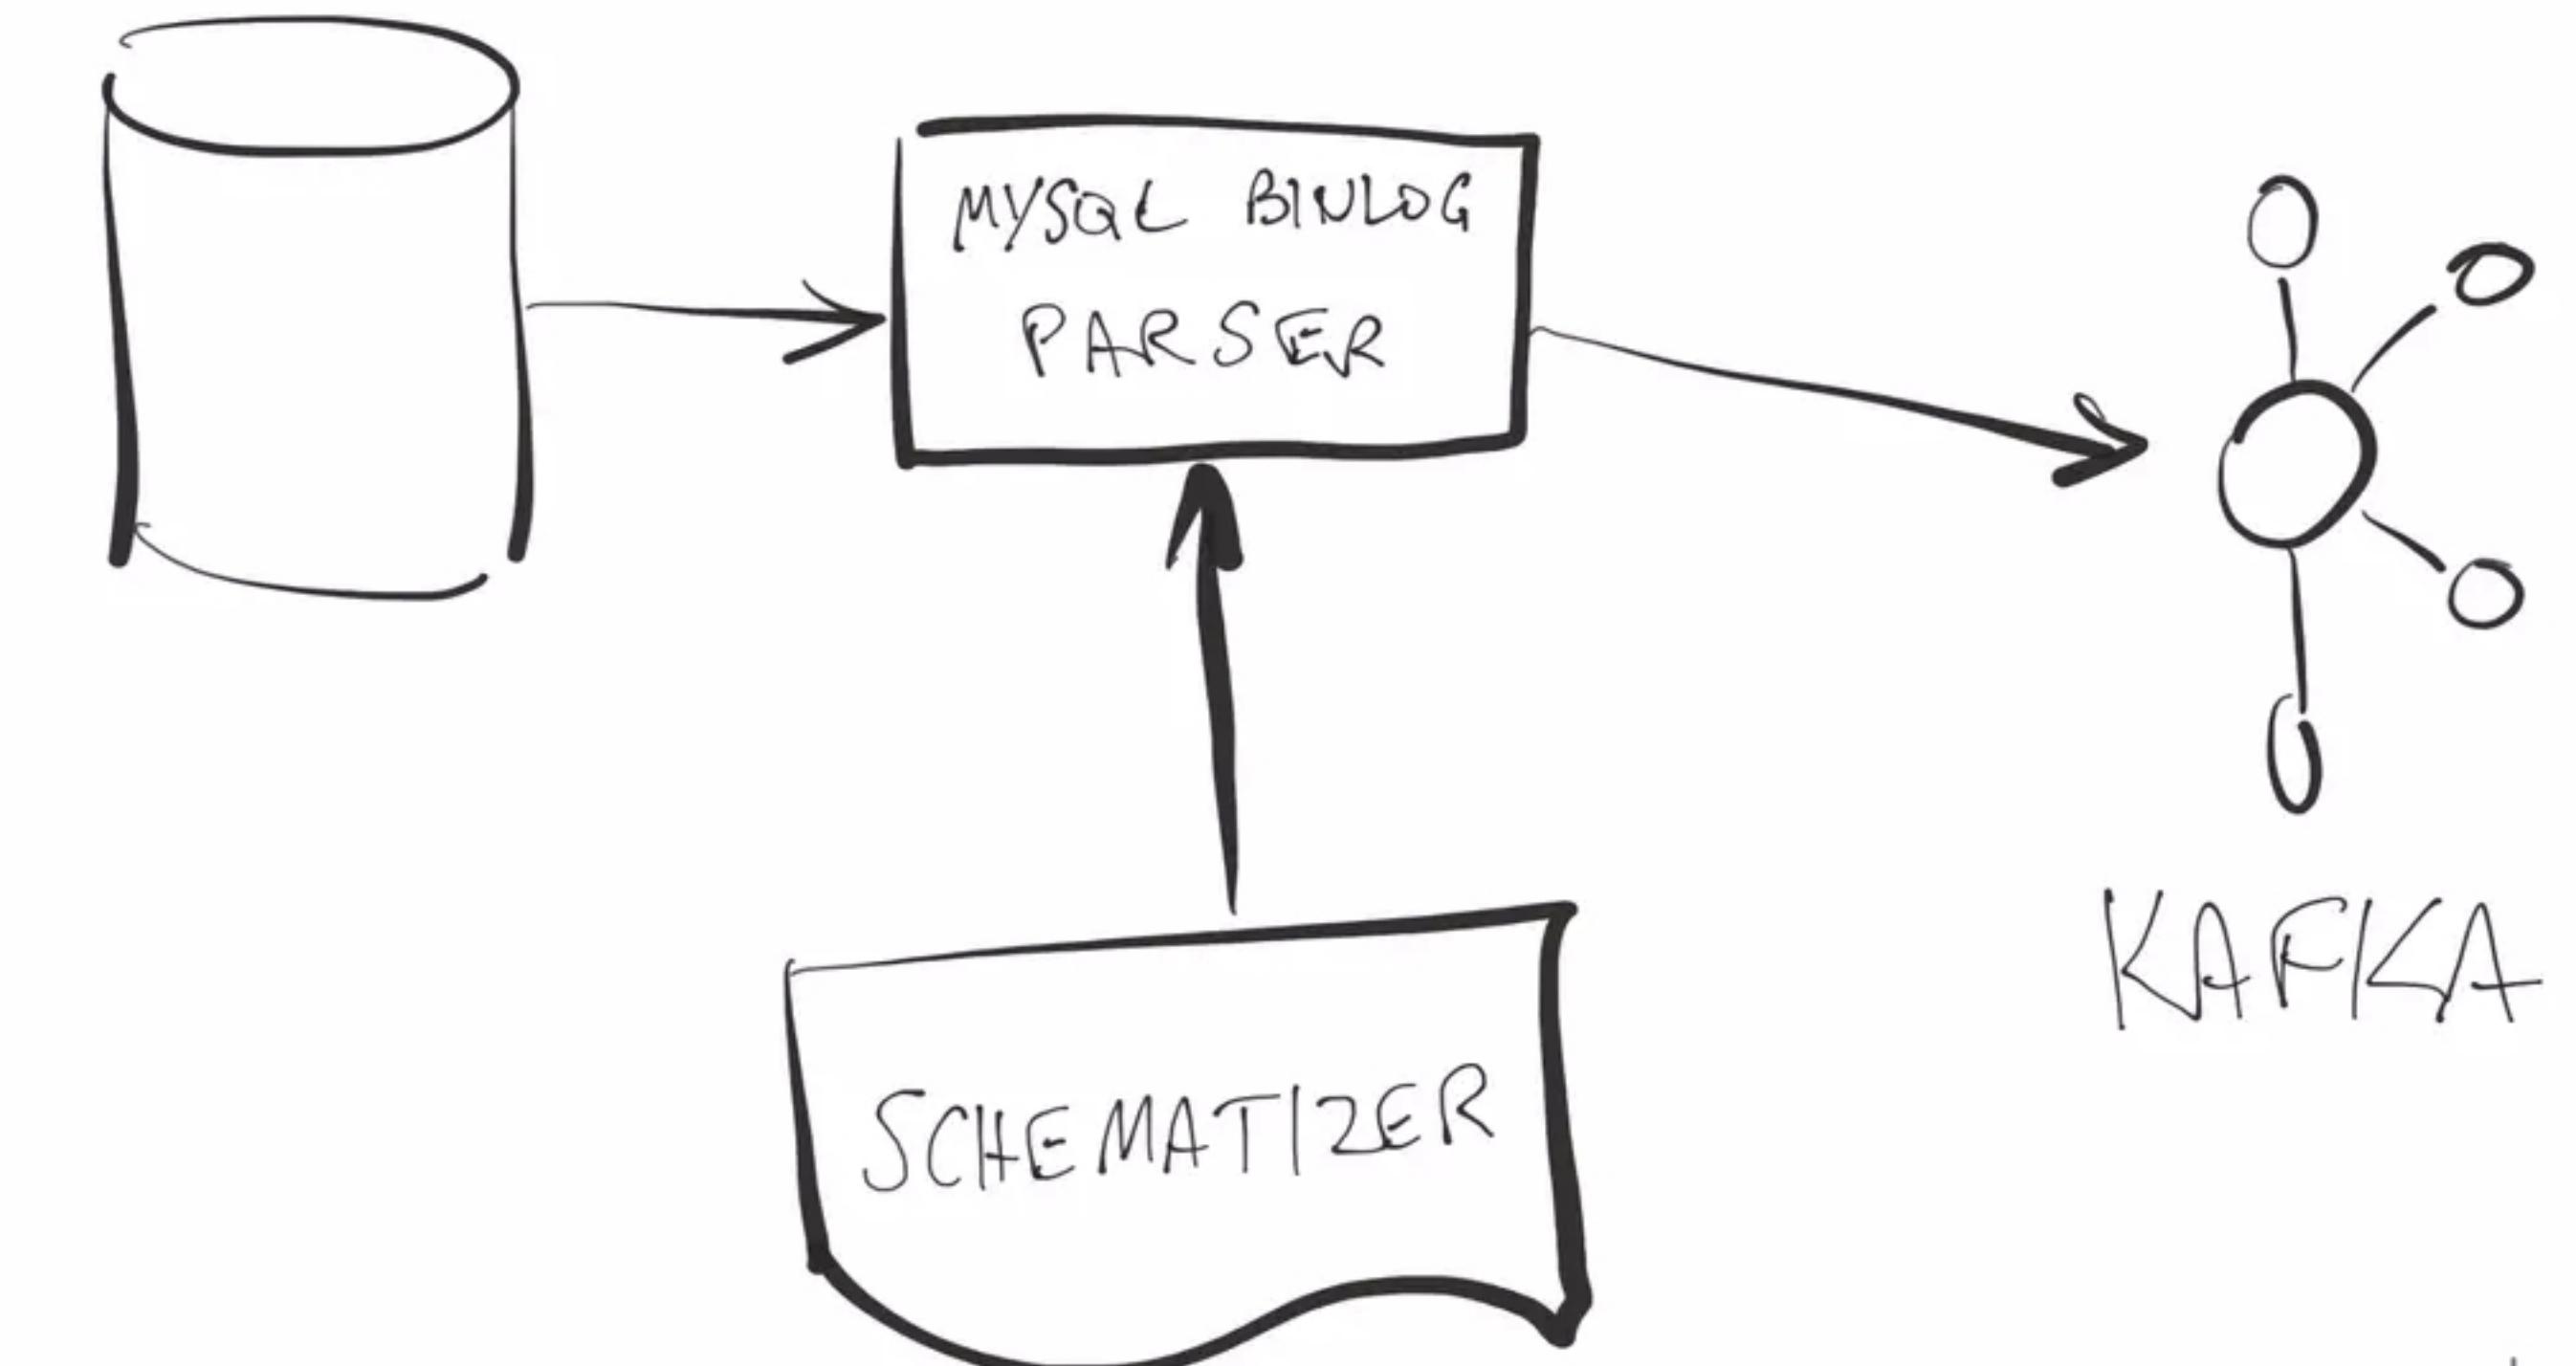
- 150 production services
- Metcalf's Law
- REST scales poorly
- N+1 problem and bulk data services





# Streaming Databases

- Stream-table duality
- But some services still want databases
- Individual services use MySQL
- Millions of rows across services

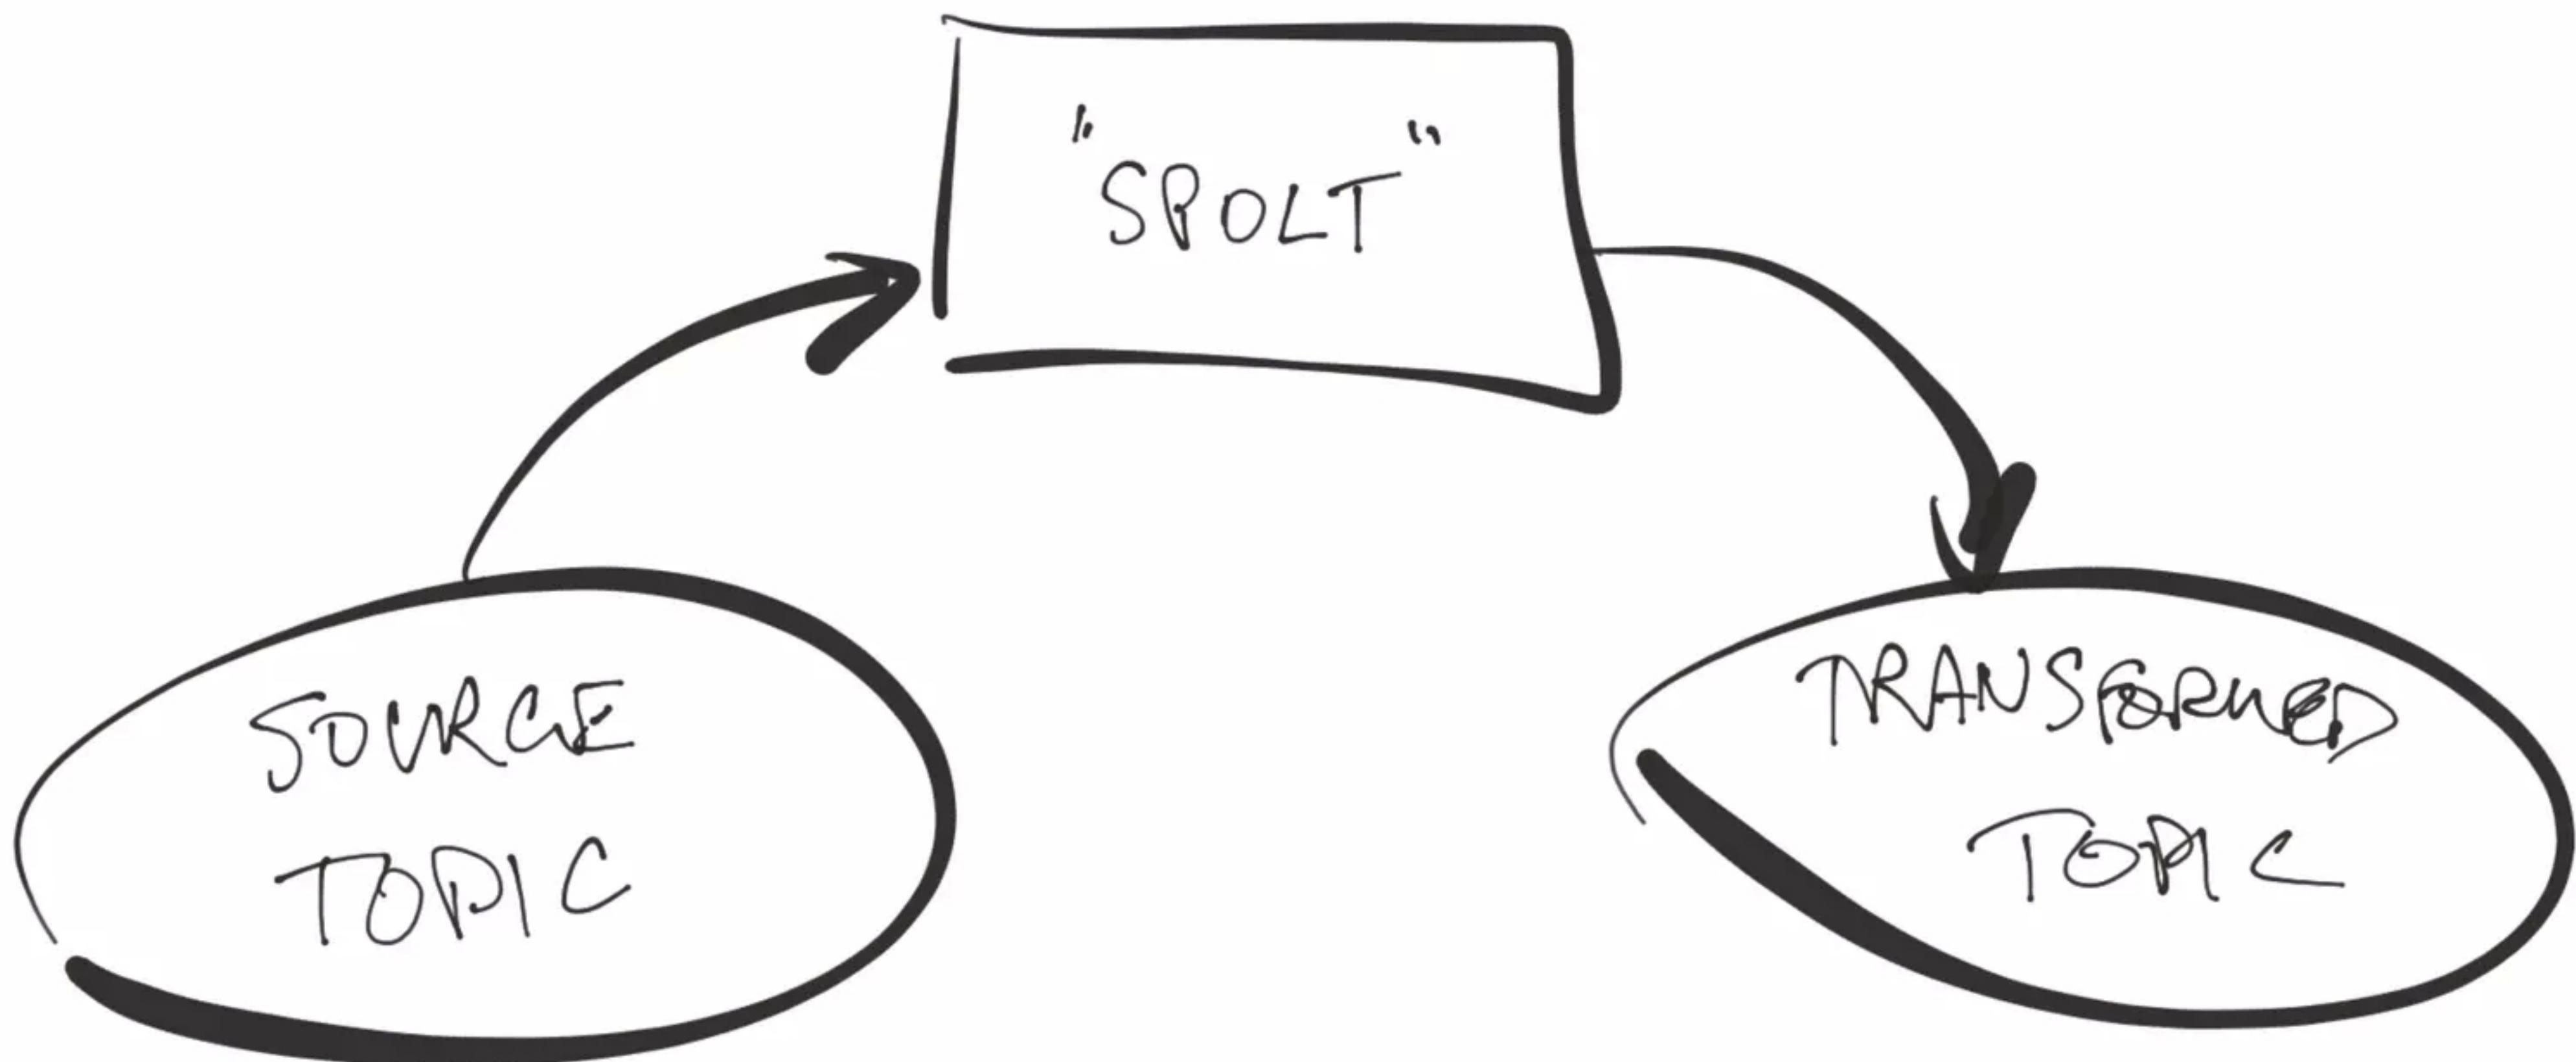


# Computing Streams

---

- In 2010, commitment to MapReduce was strong
- That fun was soon over
- Want to capitalize on schema translation tools
- Want near-real-time results
- Storm didn't support enough Python
- Wrote their own!

# PaaStorm Instance



# Yelp Conclusions

---

- Streaming pipeline ensures evolvability
- Stream/Table Duality
- Schema matters
- Yelp often ahead of the game!

---

# LINE Corp

<https://engineering.linecorp.com/en/blog/detail/80>

# Goals

---

- Unified data delivery pipeline between systems
- Replace a legacy, bespoke system for processing background tasks
- Had to pick a framework

## Option: Apache Samza™

---



Samza

# Option: Apache Kafka™ Streams

← → C ⌂ ⓘ https://www.confluent.io/blog/introducing-kafka-streams-stream-processing-made-simple



English | Deutsch



Product

Solutions

Services

Resources

About

Blog

Download



STREAM PROCESSING

## Introducing Kafka Streams: Stream Processing Made Simple



Jay Kreps

# Streaming API for Kafka

---

```
KStreamBuilder builder = new KStreamBuilder();

// Read the input Kafka topic into a KStream instance.
KStream<byte[], String> textLines = builder.stream("TextLinesTopic");

// Convert to upper case (:: is Java 8 syntax)
KStream<byte[], String> uppercasedWithMapValues = textLines.mapValues(String::toUpperCase);

// Write the results to a new Kafka topic called "UppercasedTextLinesTopic".
uppercasedWithMapValues.to("UppercasedTextLinesTopic");
KafkaStreams streams = new KafkaStreams(builder, streamsConfiguration);
streams.start();
```

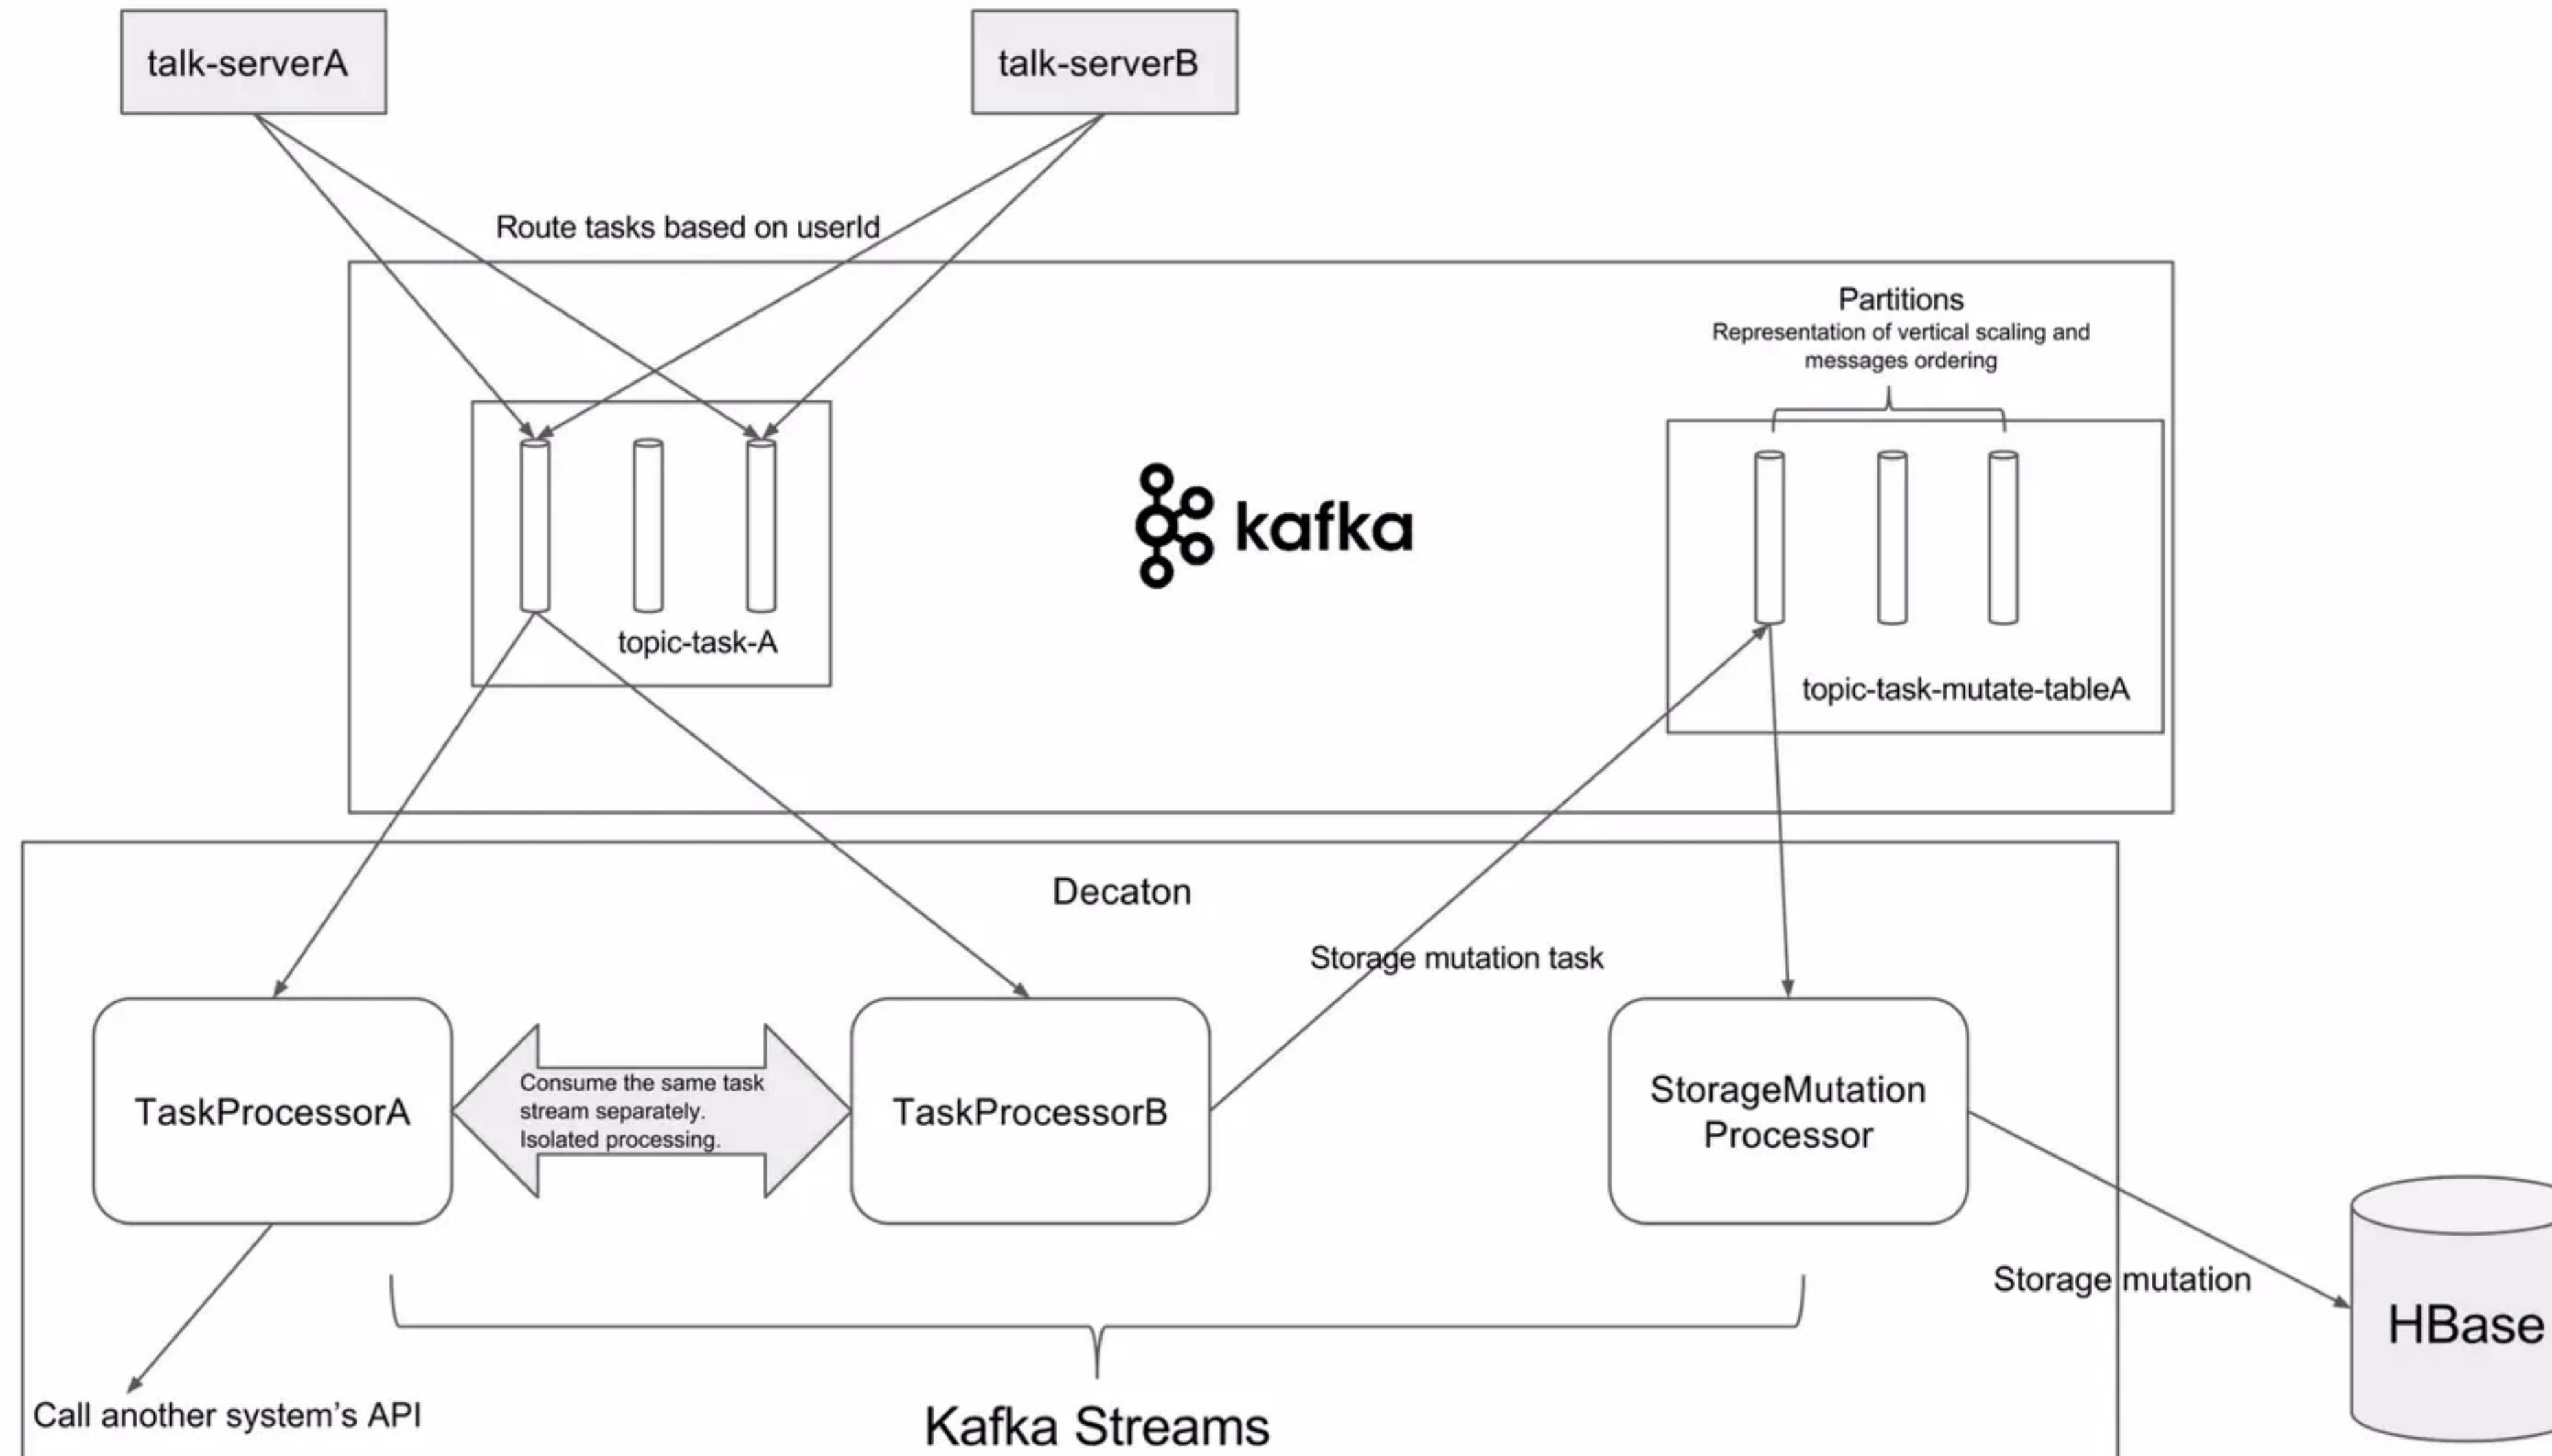
# Stateful Transformations

# Loopback Replicator

---

- One big input topic of all TalkOperations (1M messages/sec)
- Some services want a limited view of these
- Simple filtering

# Decaton



# LINE Conclusions

---

- Classical Kafka Streams implementation
- Better to “buy” than build

---

# Altspace VR



AltspaceVR

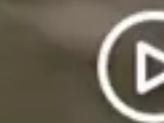
[Get AltspaceVR](#) [Events](#)

[Log In](#)

[Sign Up](#)

Be there together

[GET IT NOW](#)



[WATCH NOW](#)

Daydream

htc VIVE

oculus

SAMSUNG  
Gear VR

2D mode









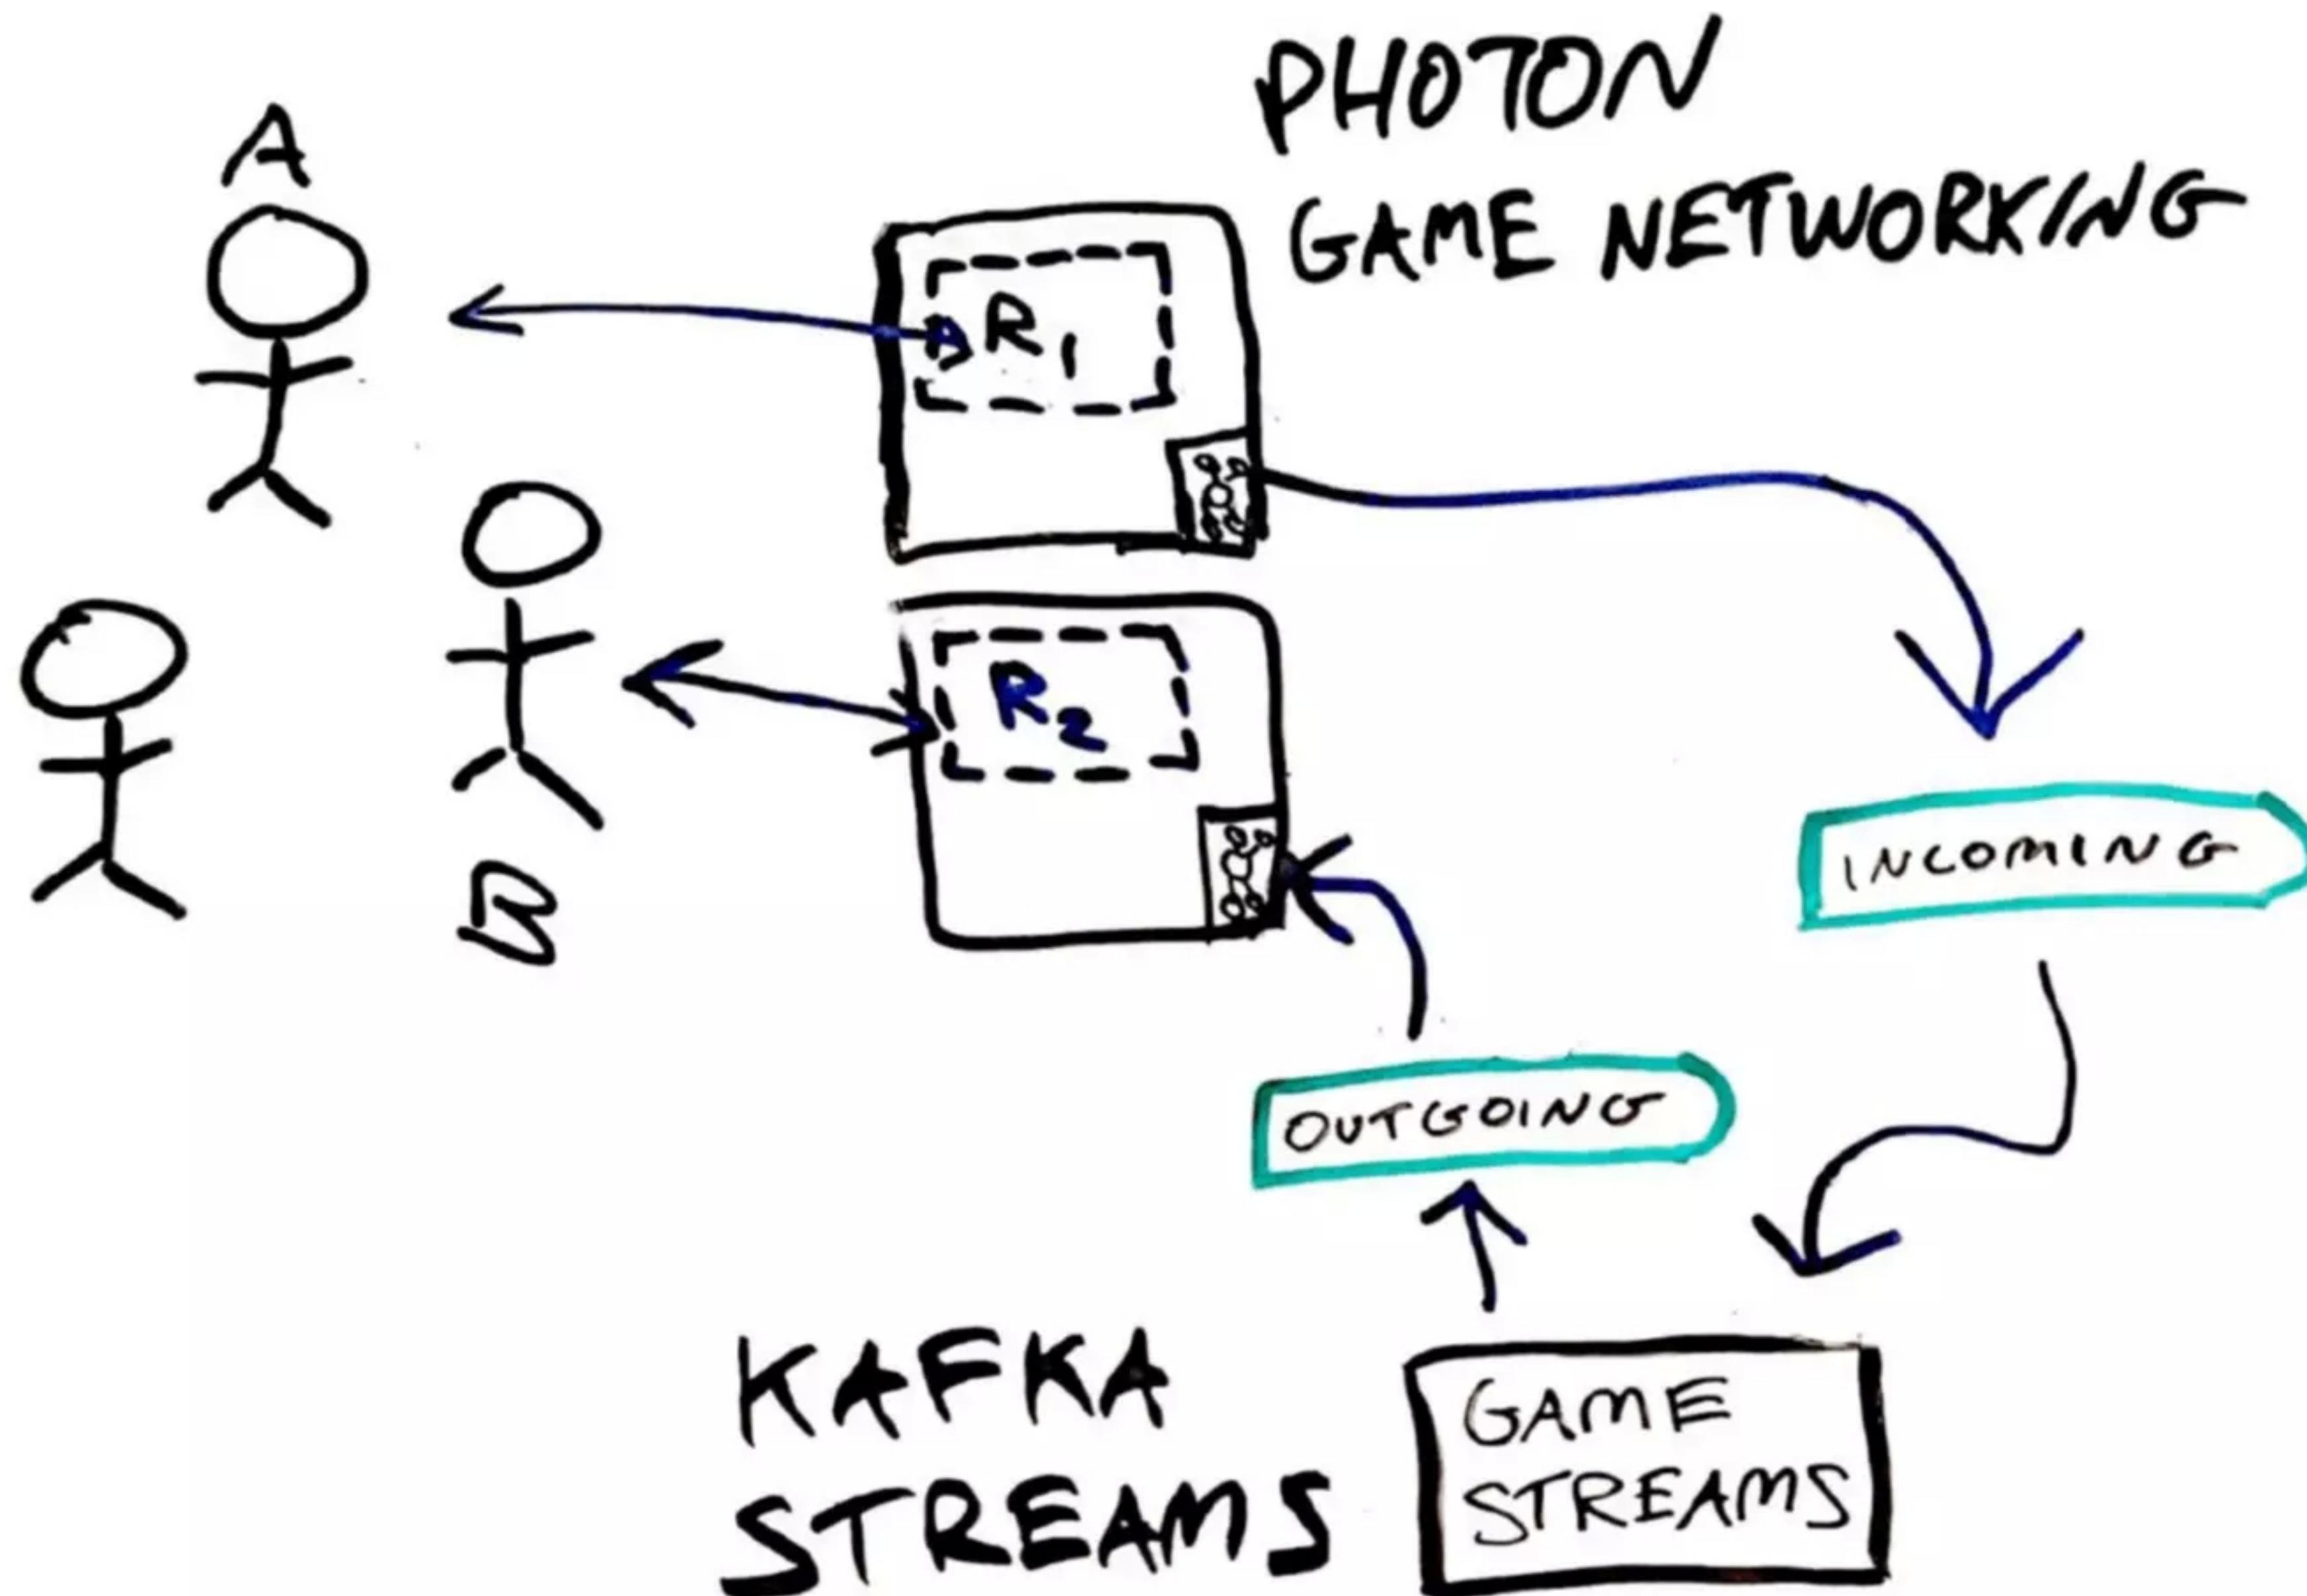
# REGGIE WATTS

LIVE IN VIRTUAL REALITY

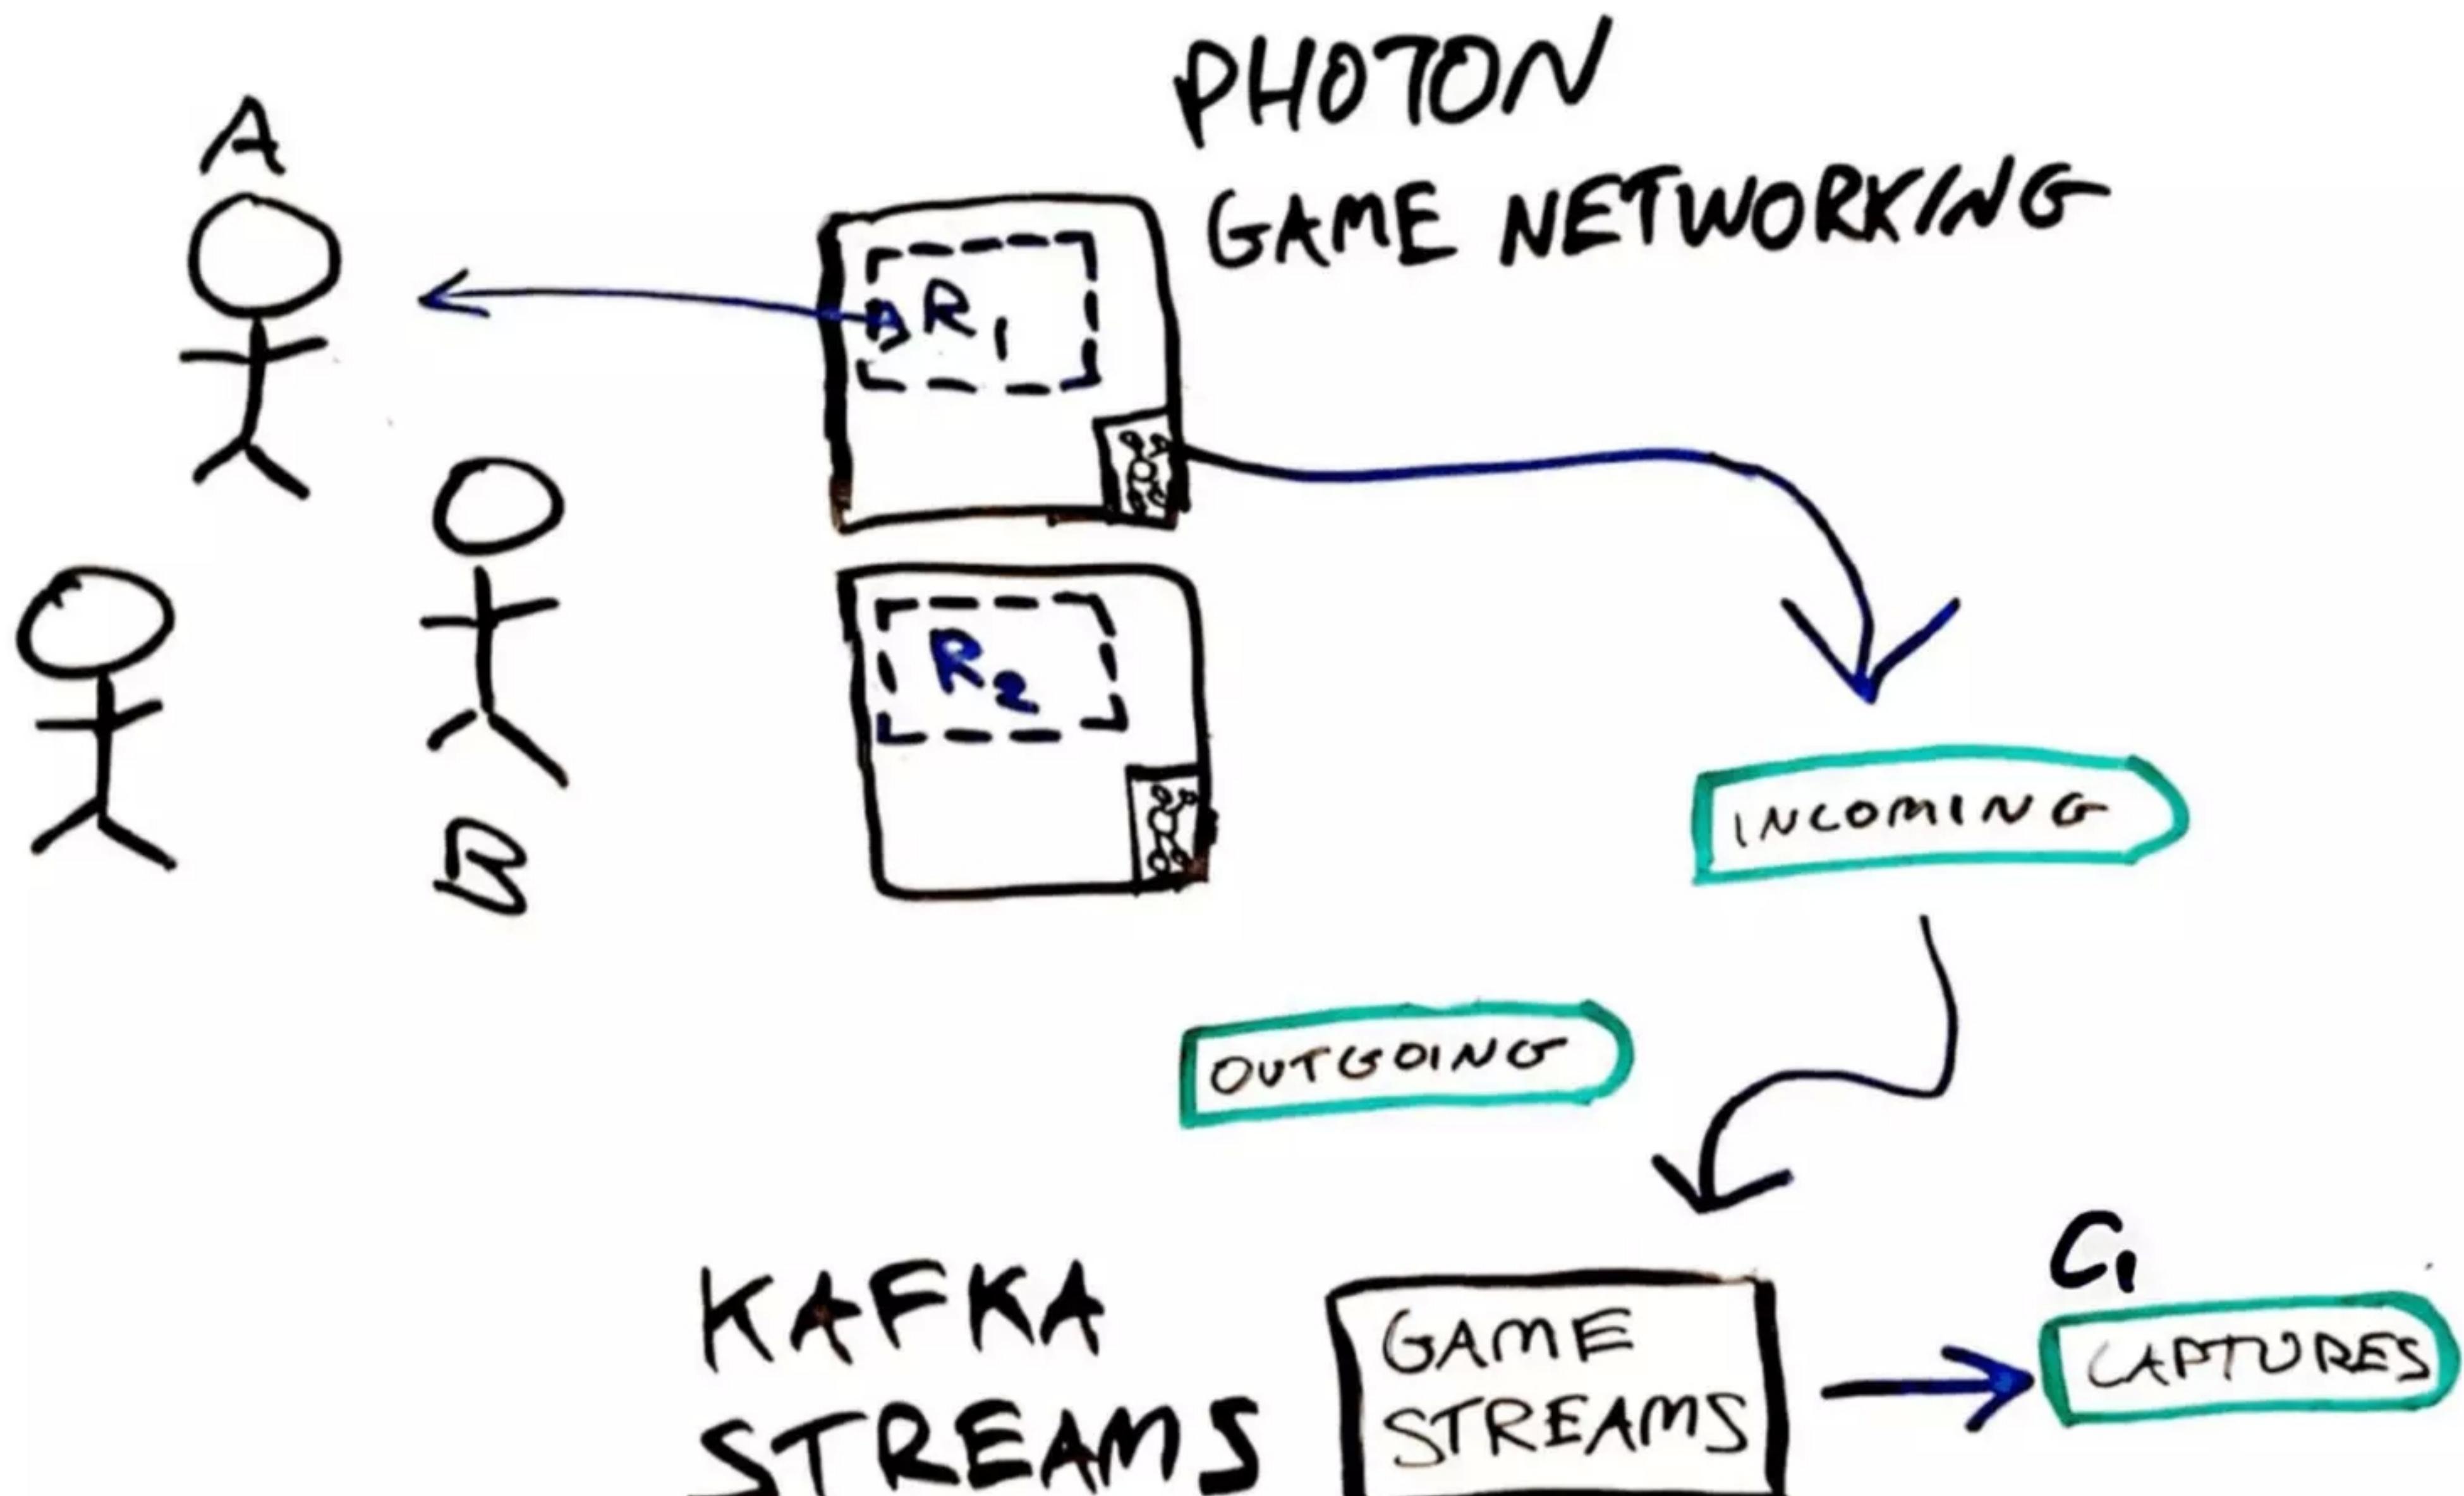
# Altspace Architecture

---

- Photon game engine
- Real-time play
- Capture for playback
- An entirely classical stream processing system



“Mirror User A to room R<sub>2</sub>”

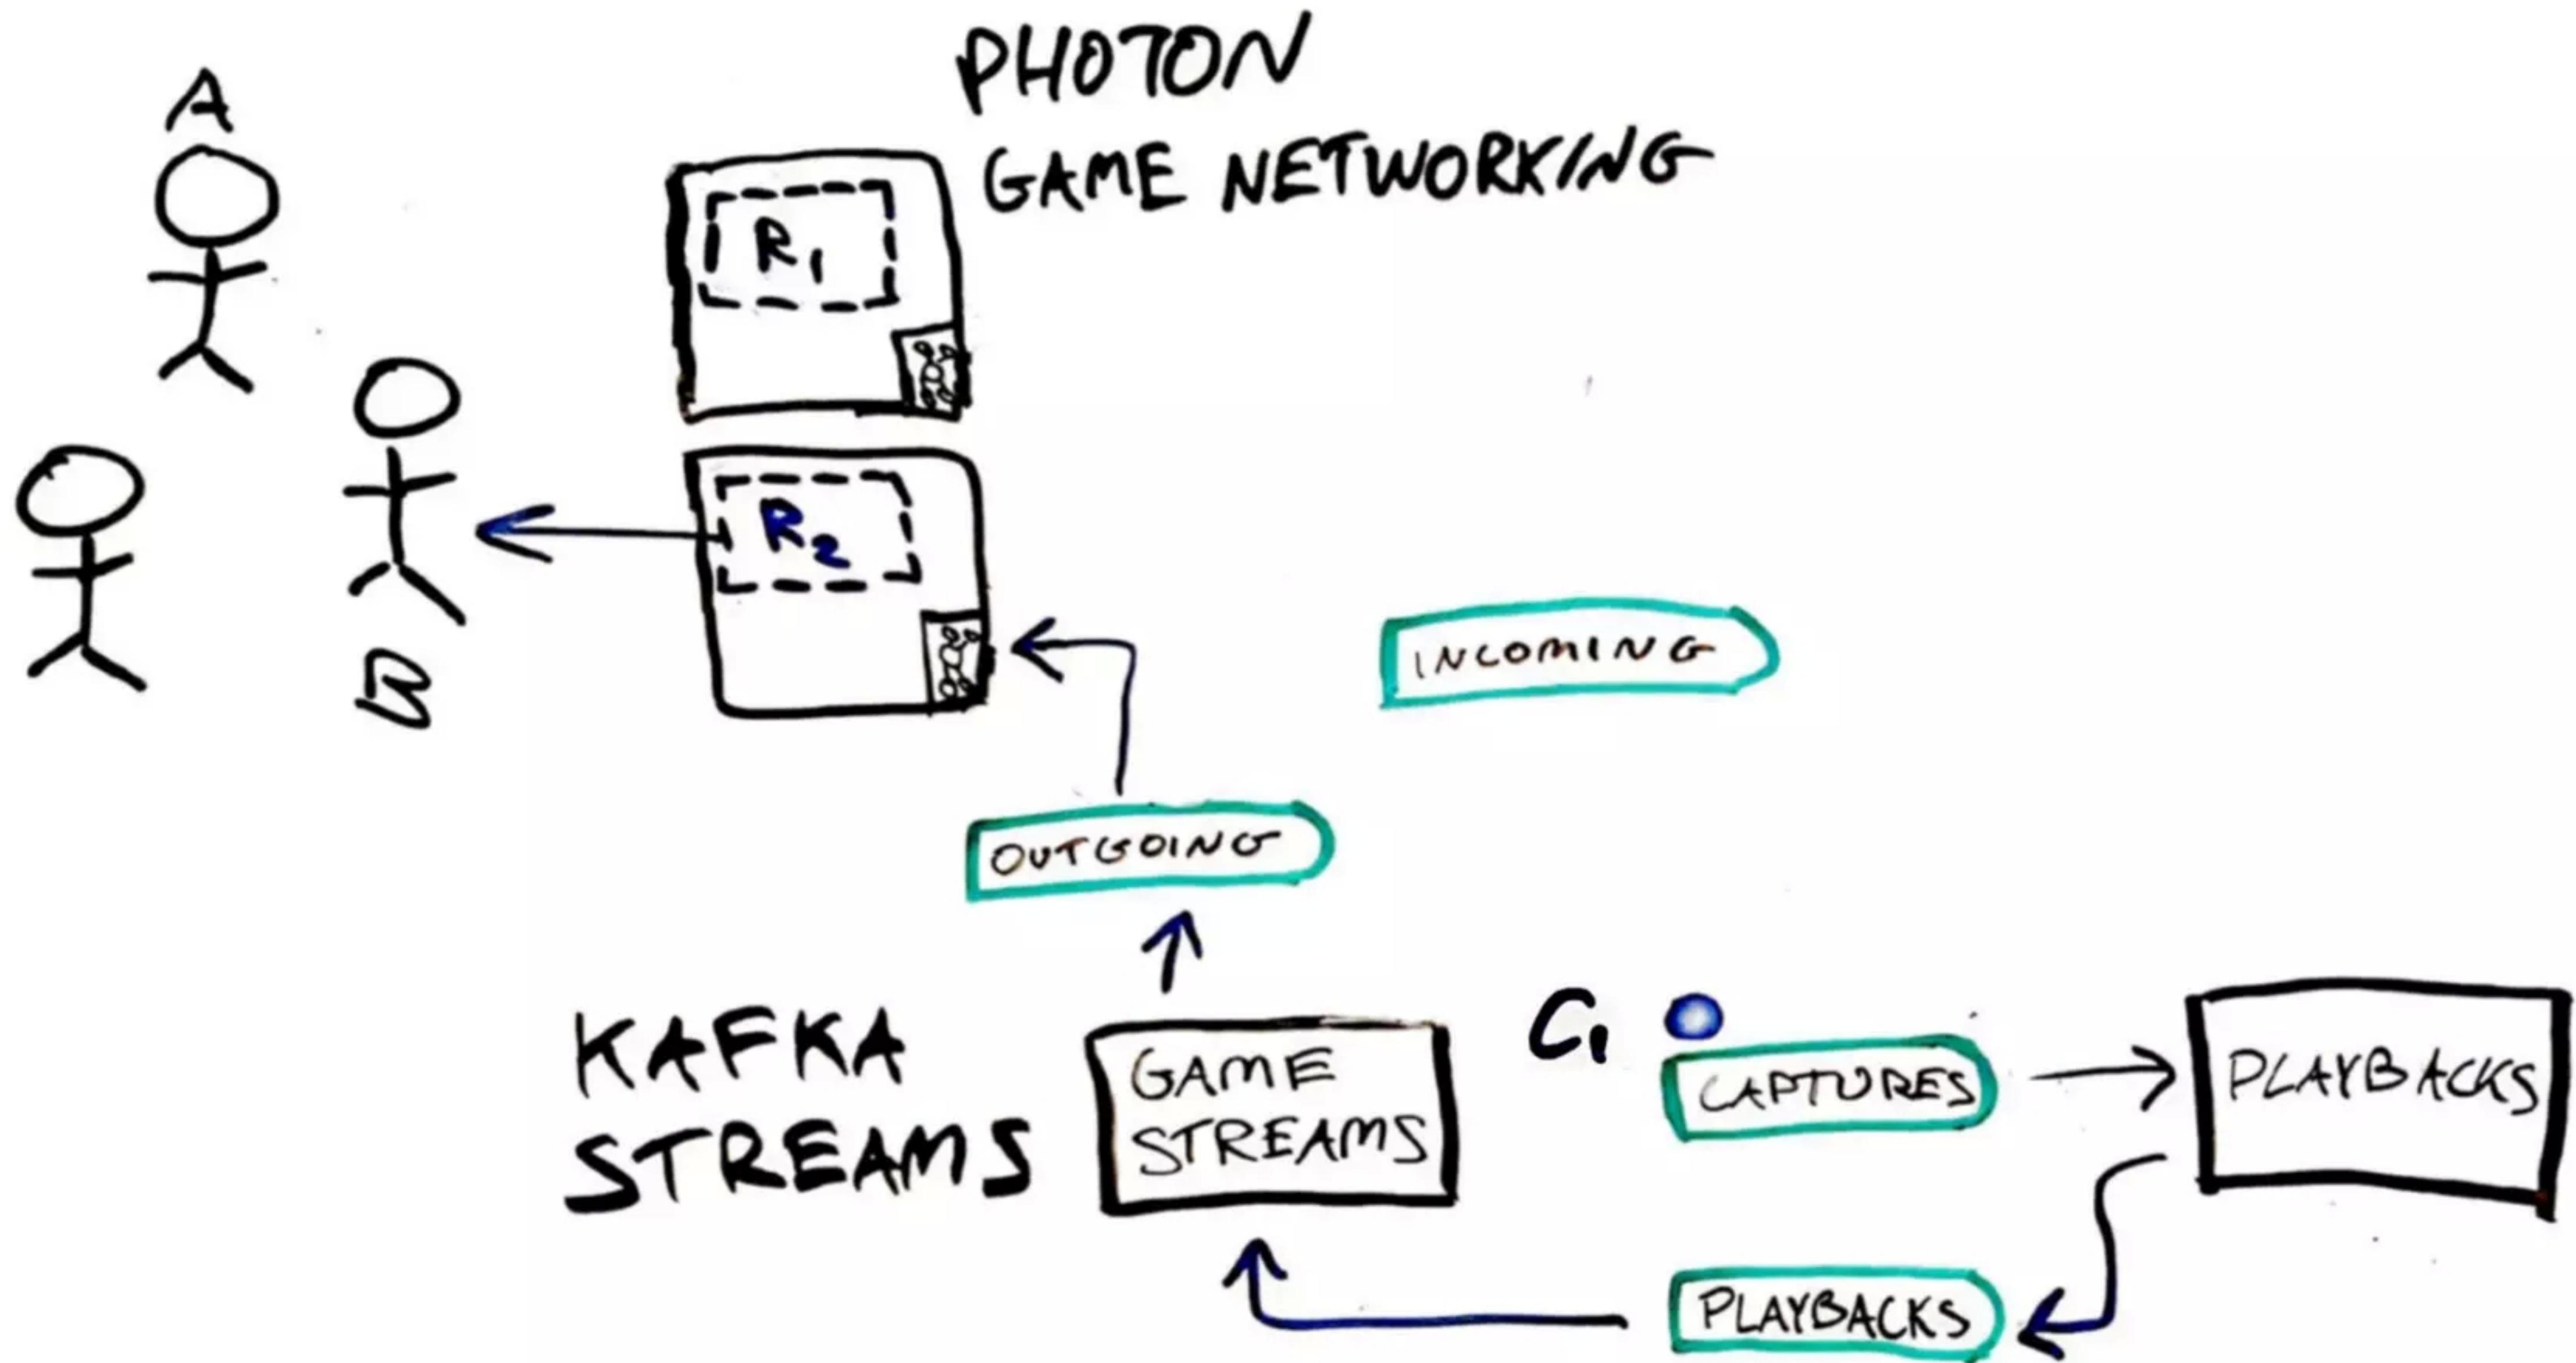


6 months later: "Capture User A"

# Playbacks

---

- Source capture data from Kafka or S3
- “Produce” it in a timed fashion



“Playback capture to room  $R_2$ ”

Source: Altspace



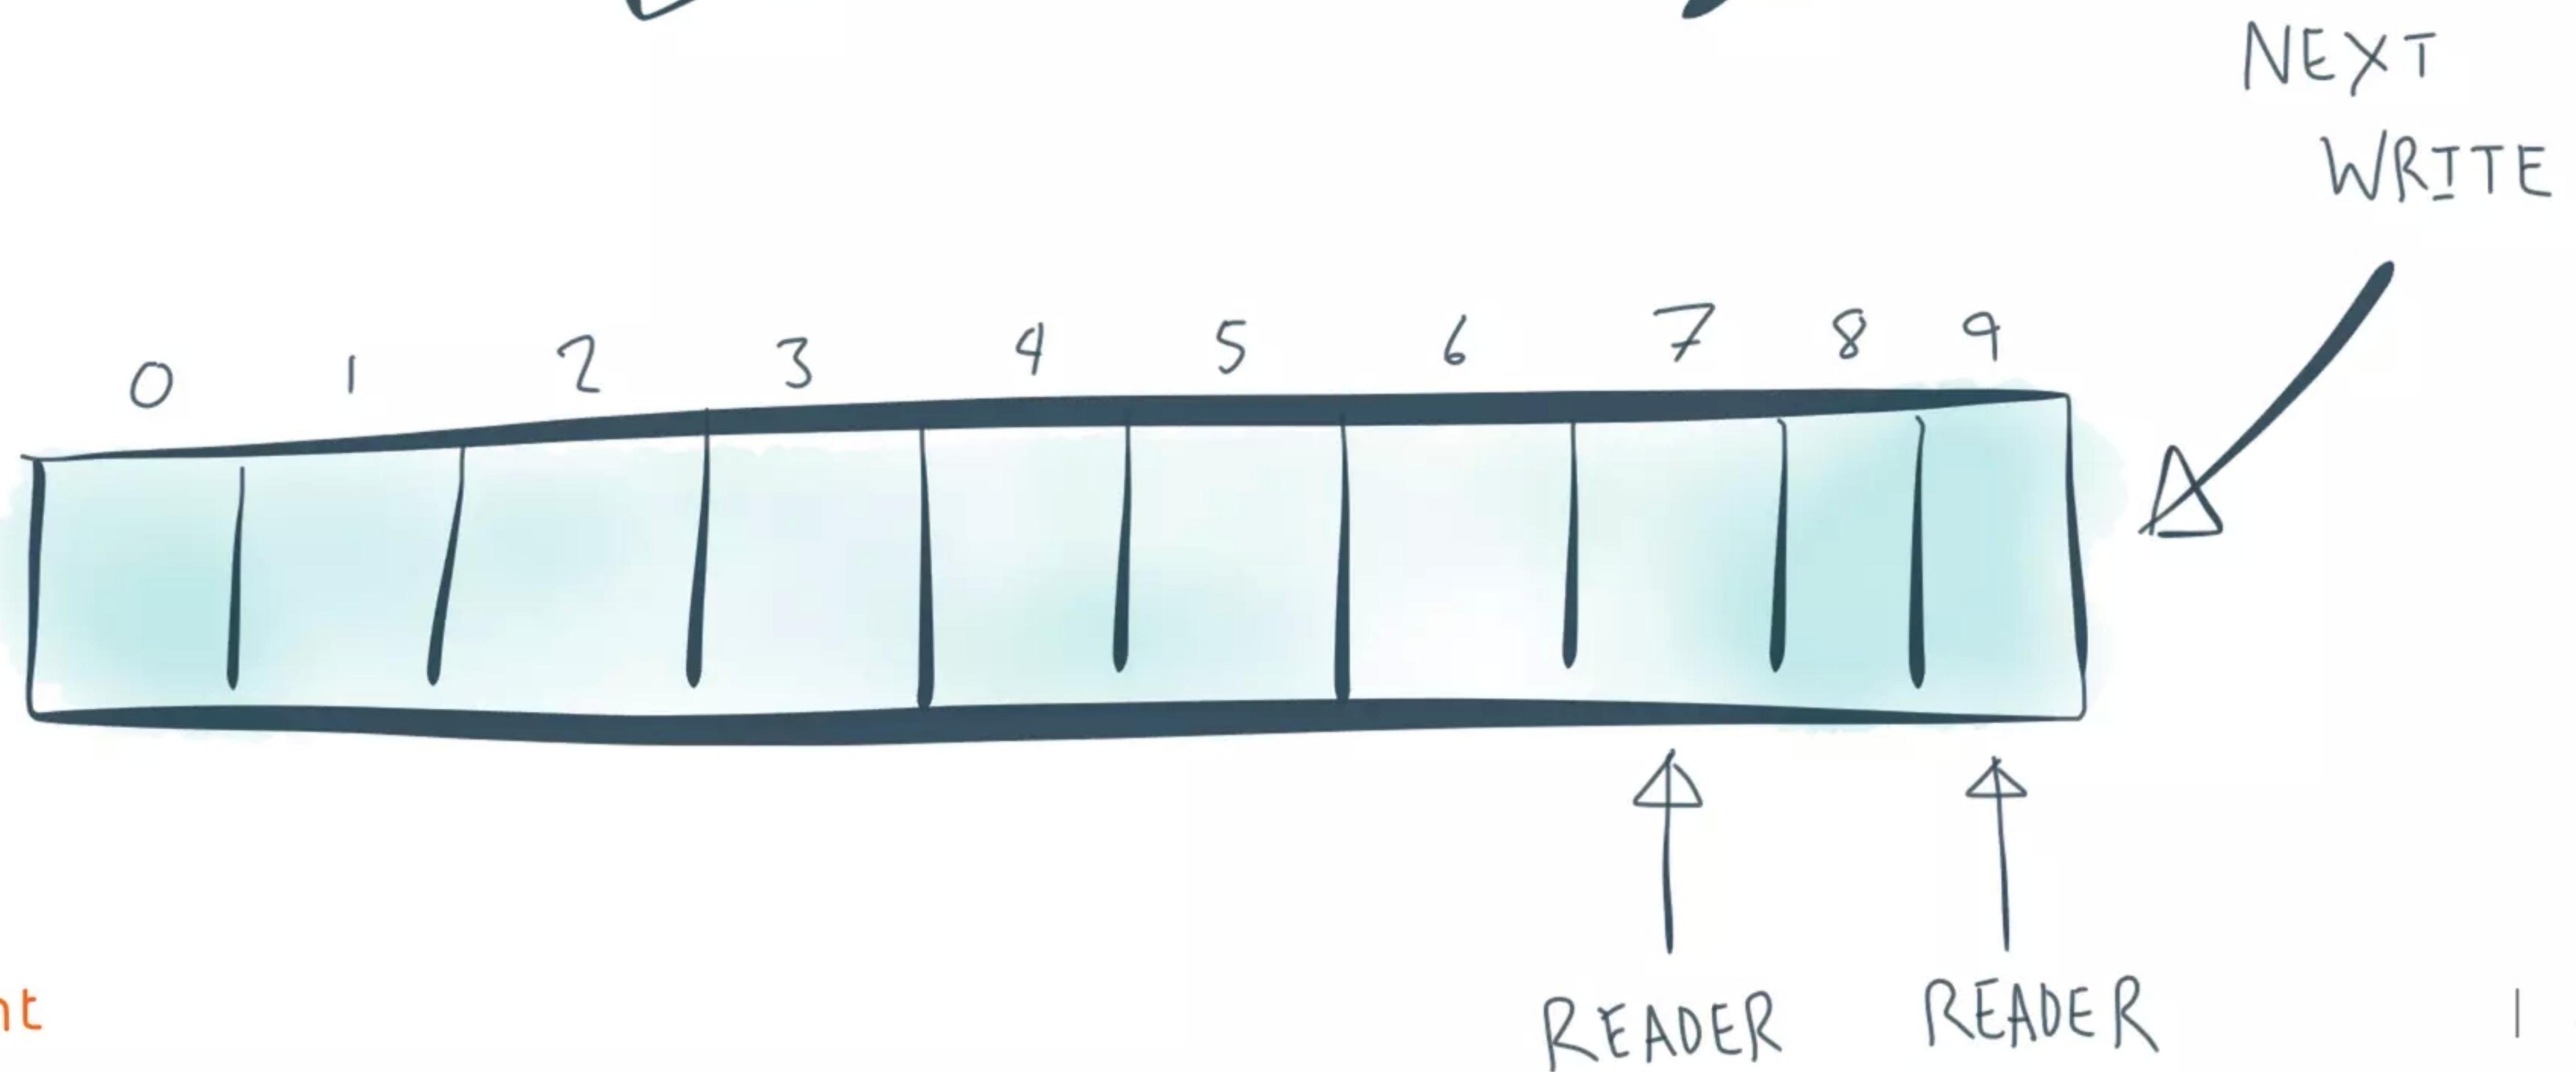
---

# The Dream Stream

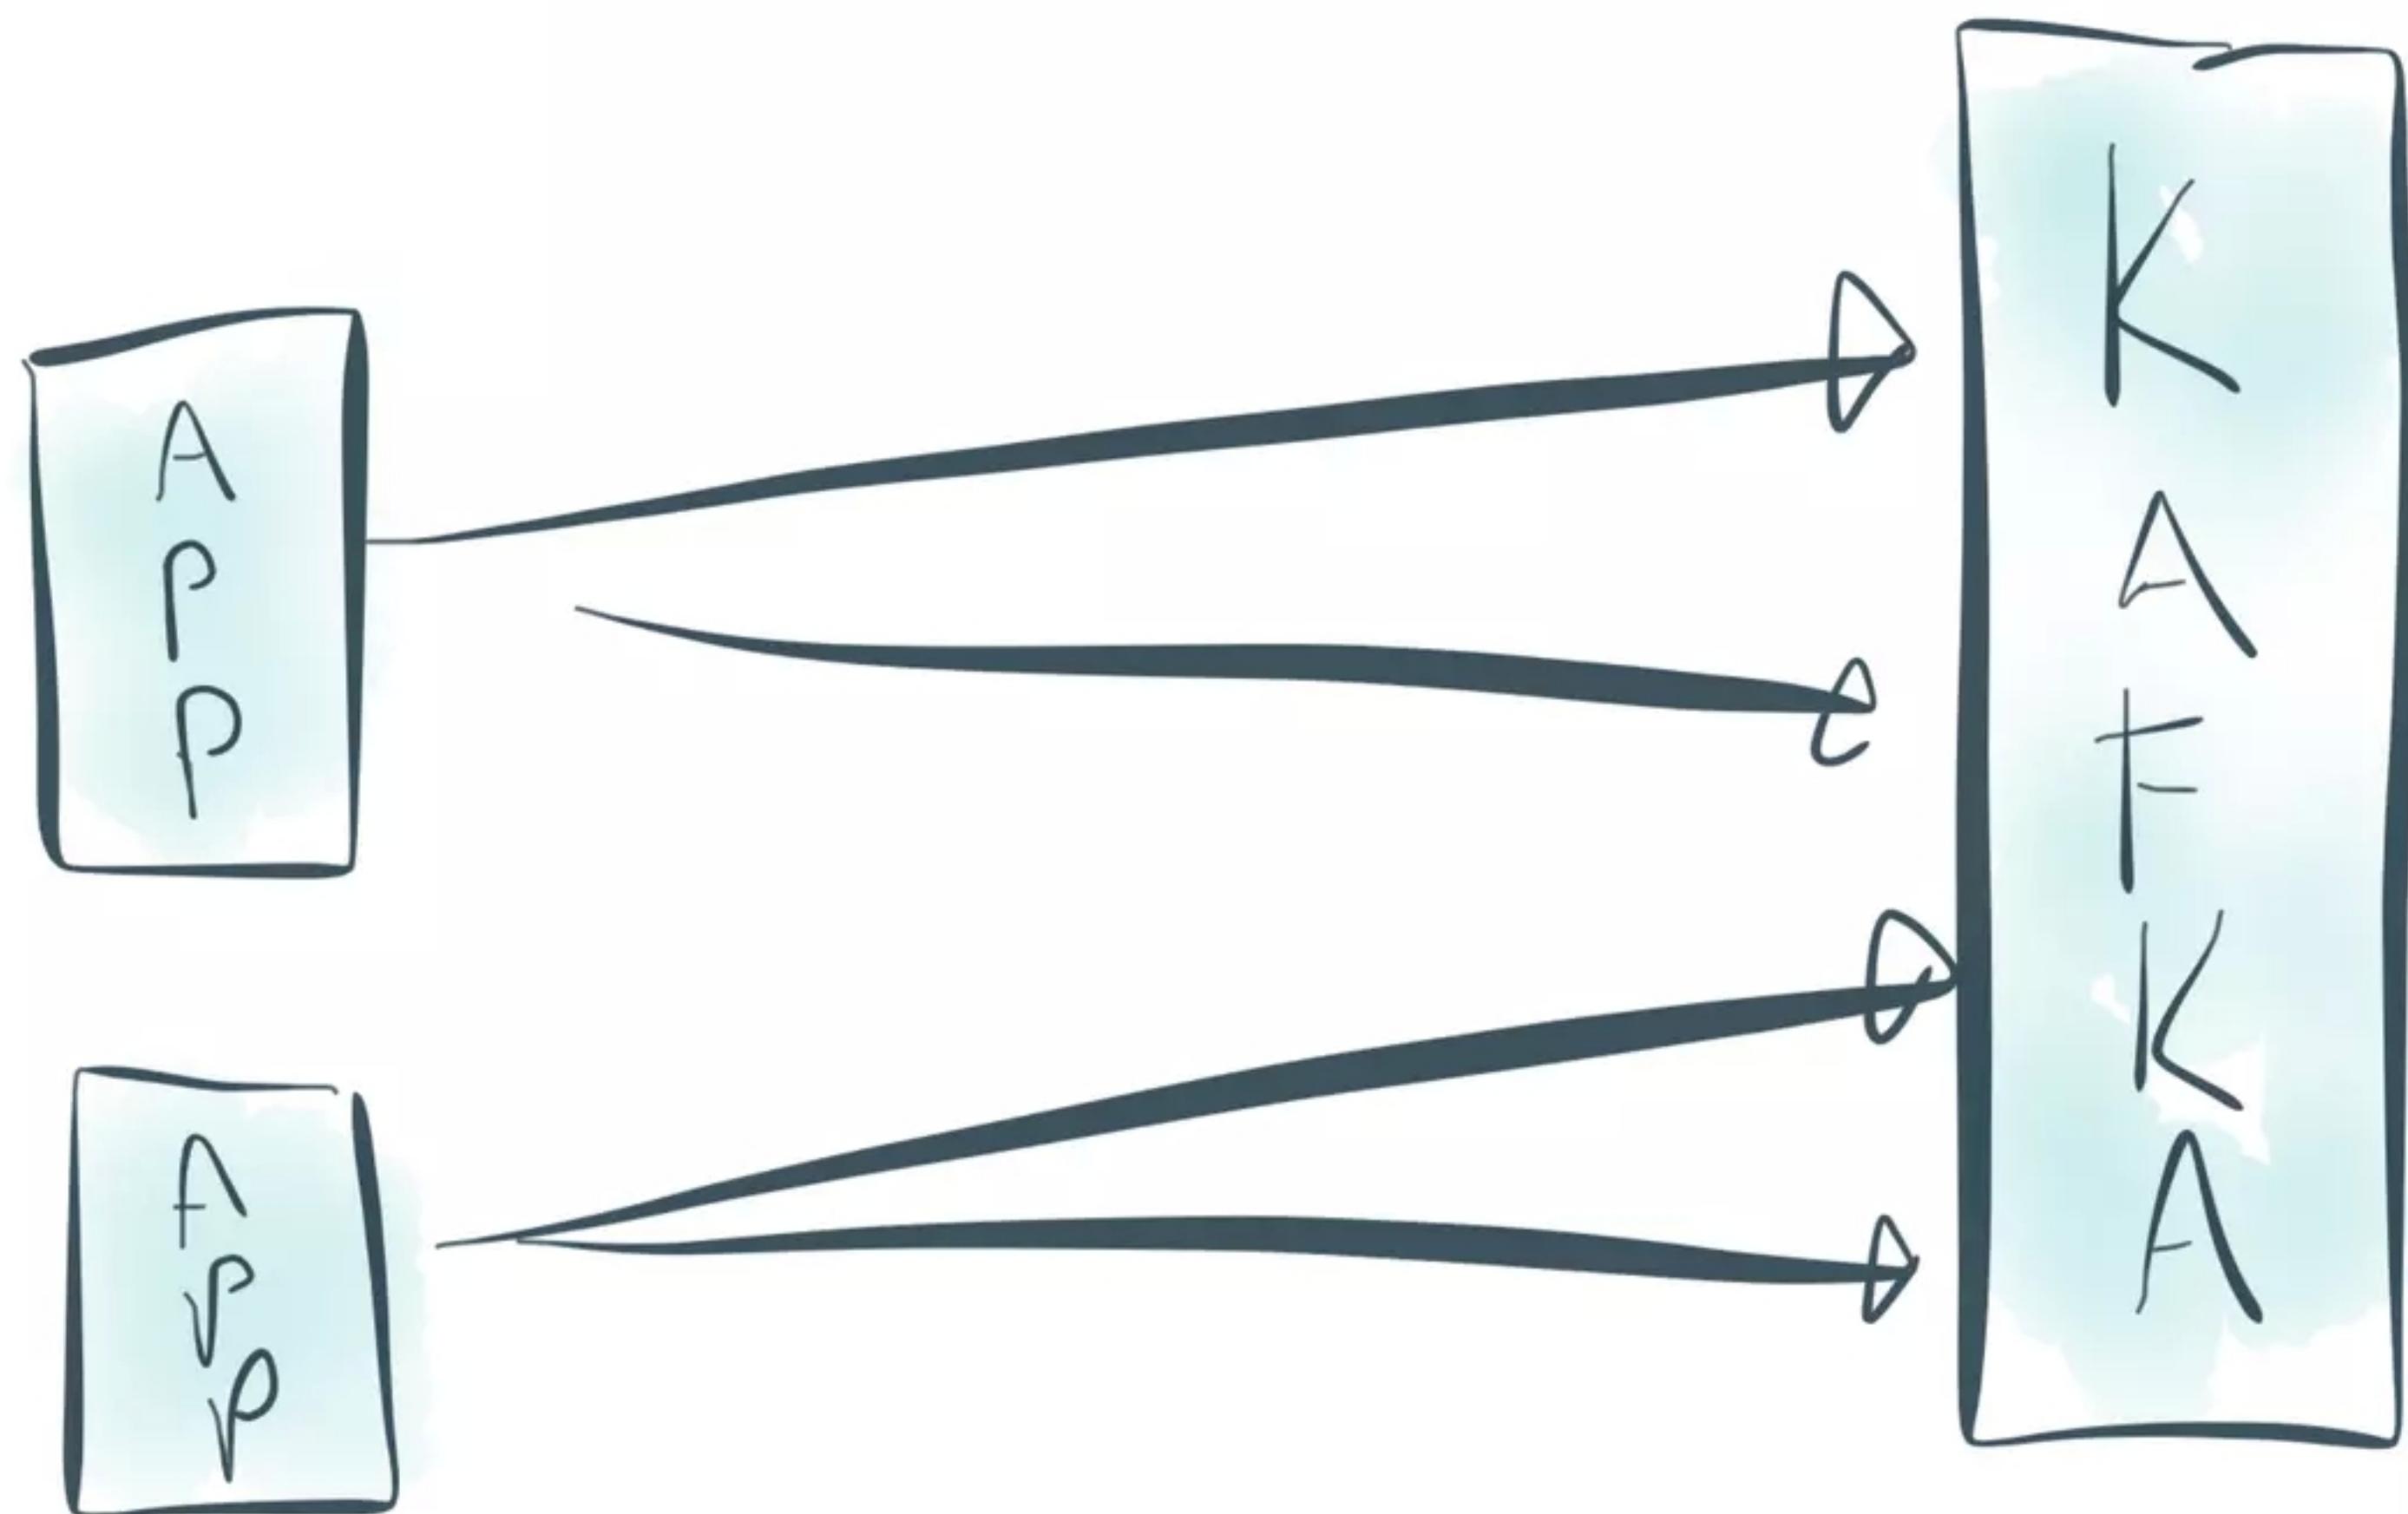
# STREAMING PLATFORM

- 1 PUB/SUB
- 2 STORE
- 3 PROCESS

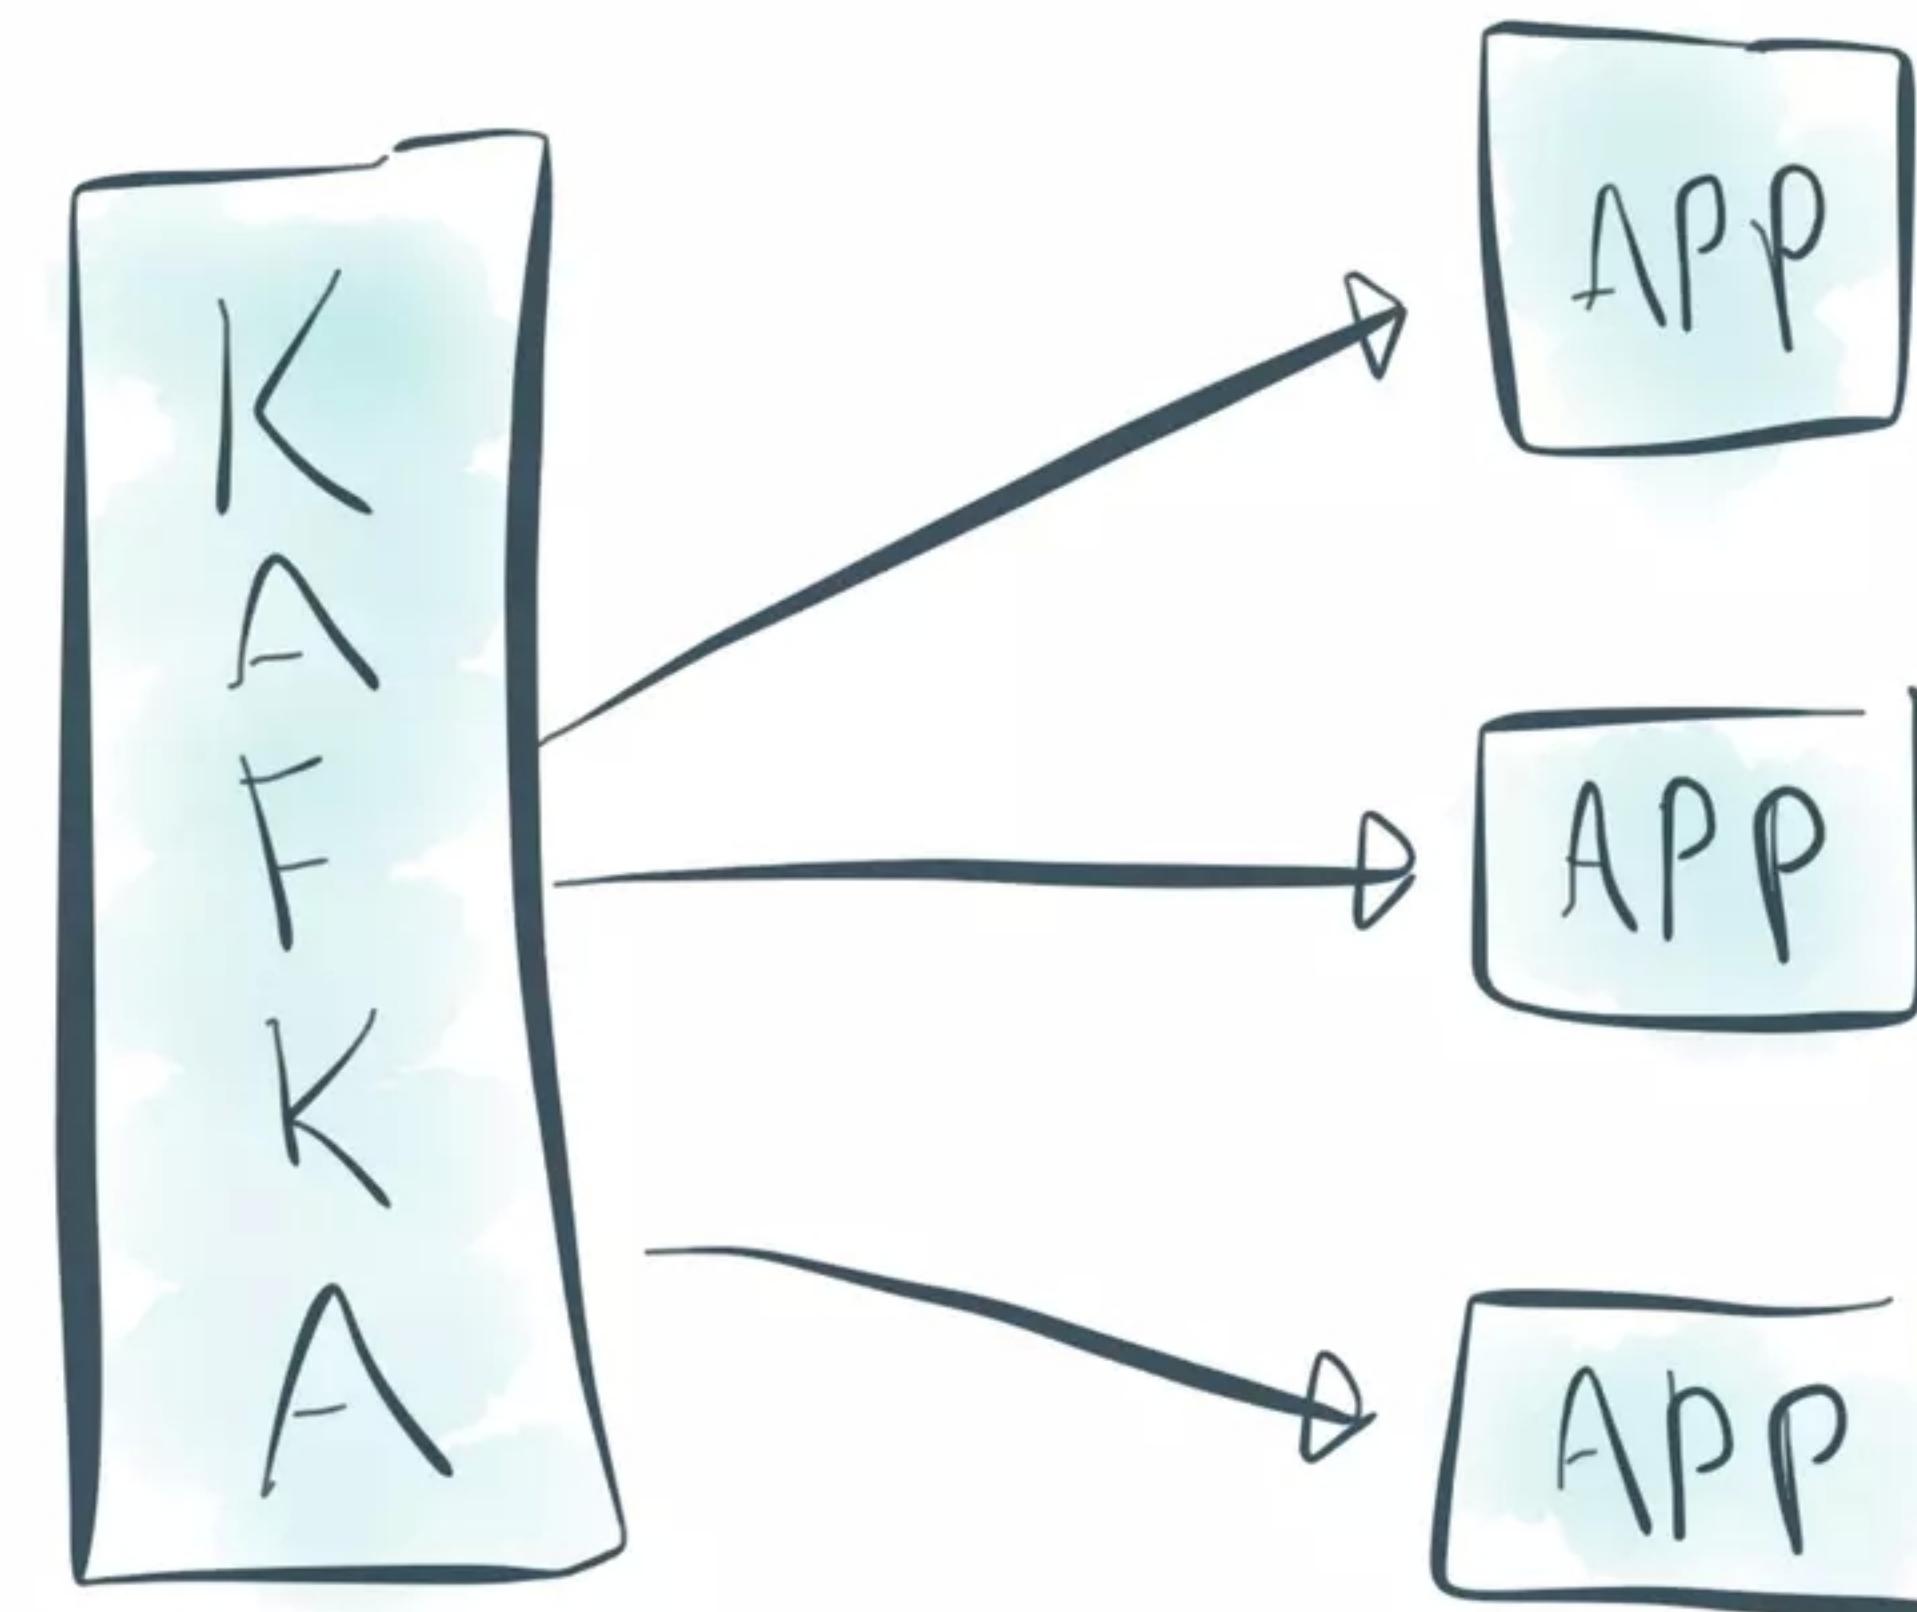
LOGS



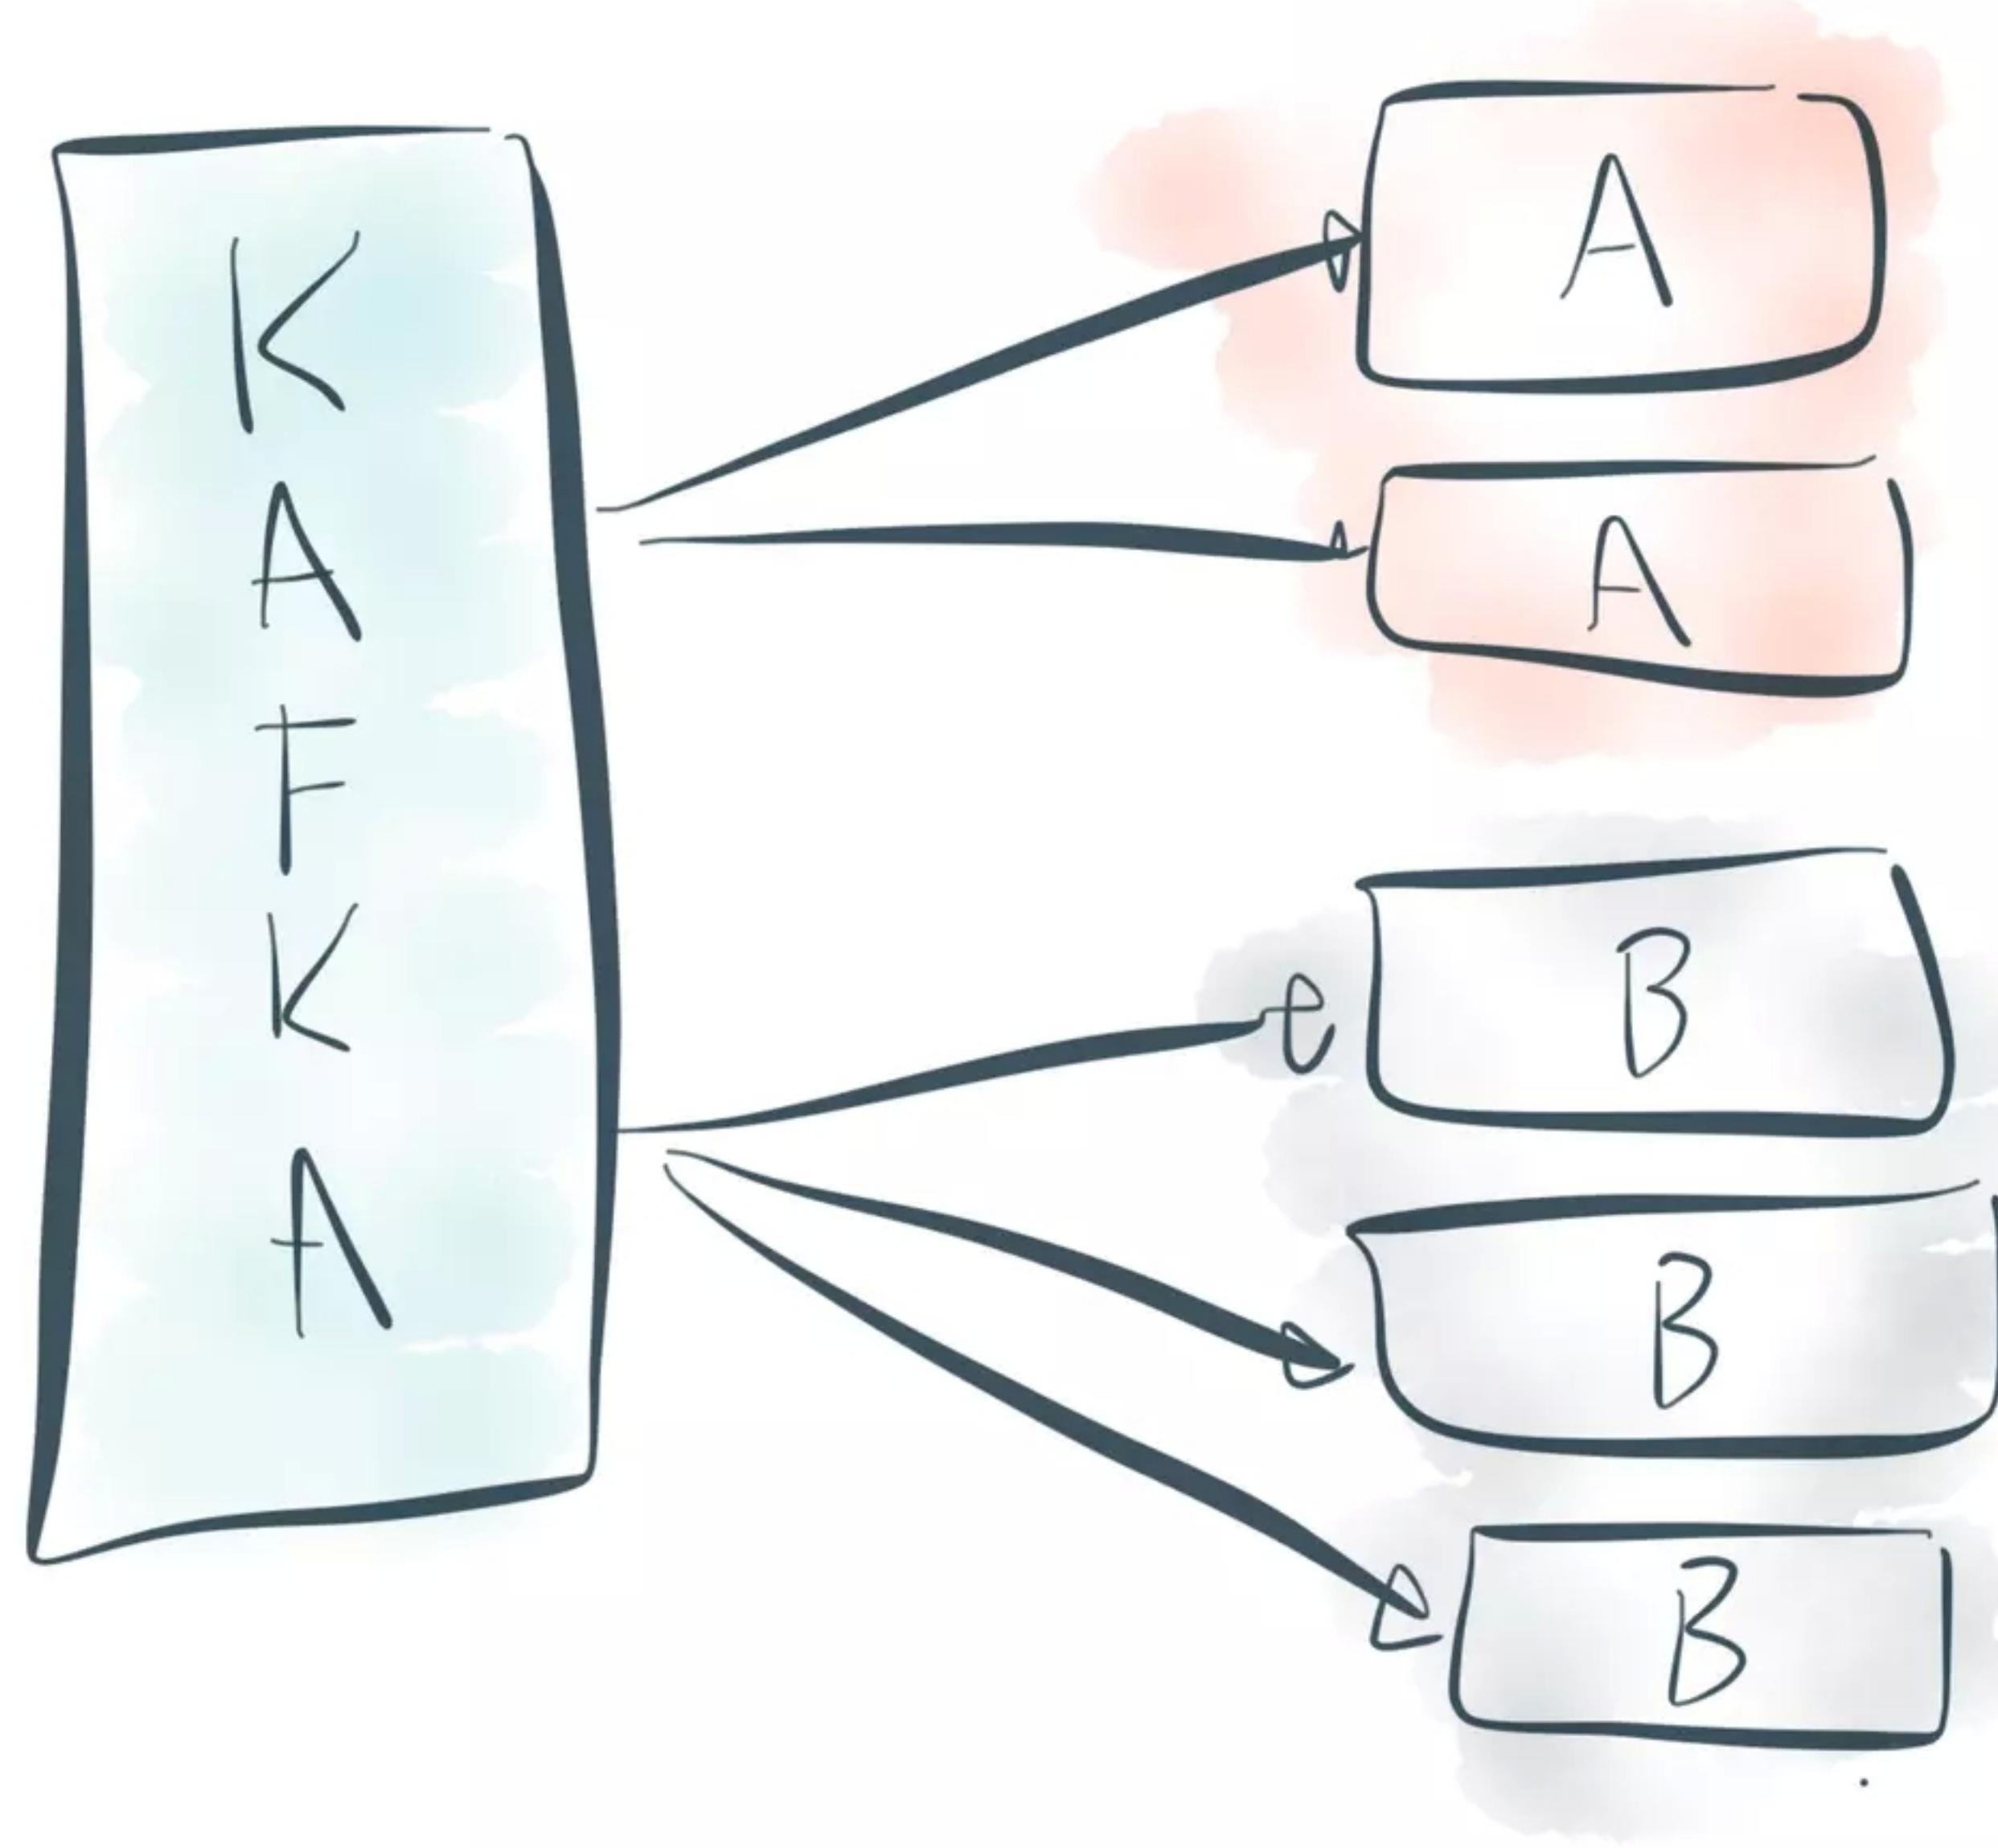
# PRODUCER



# CONSUMER



# GROUPS



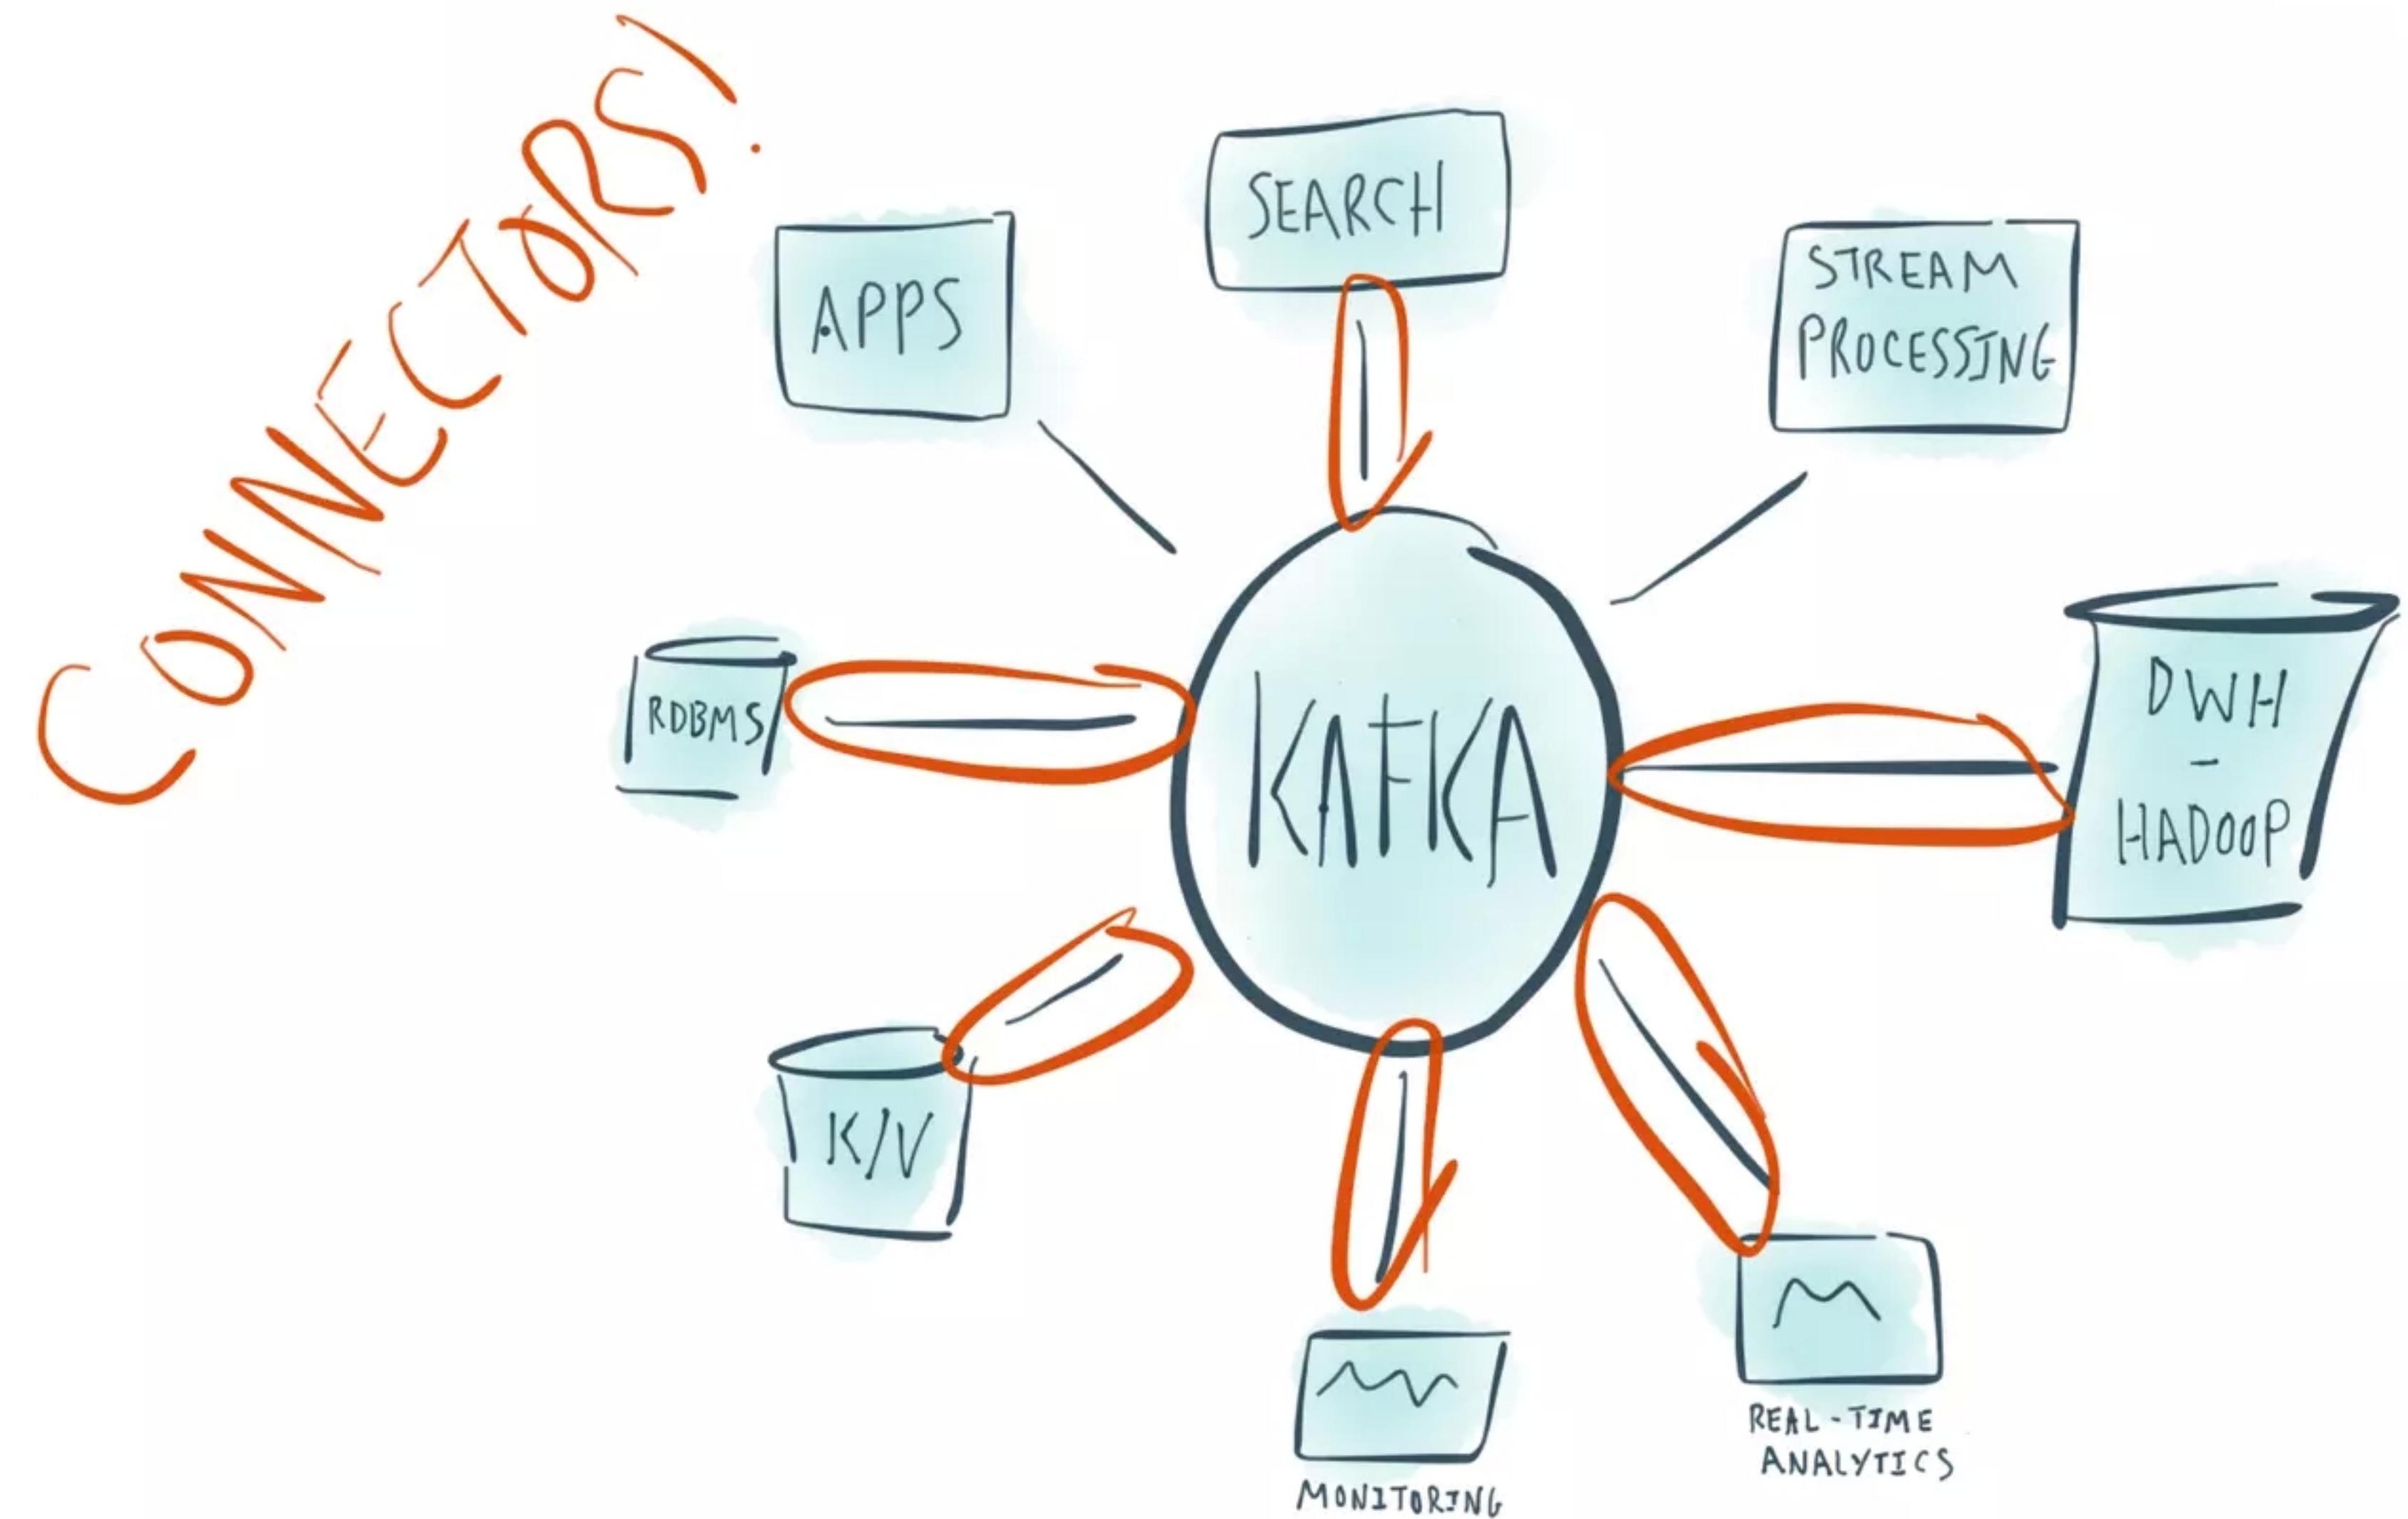
# KAFKA CONNECT



KAFKA CONNECT DOES  
THE HARD WORK

1. SCALE OUT
2. FAULT TOLERANCE
3. CENTRAL MANAGEMENT
4. SCHEMA PROPAGATION

# STREAMING PLATFORM



# STREAMING PLATFORM

- 1 PUB/SUB
- 2 STORE
- 3 PROCESS

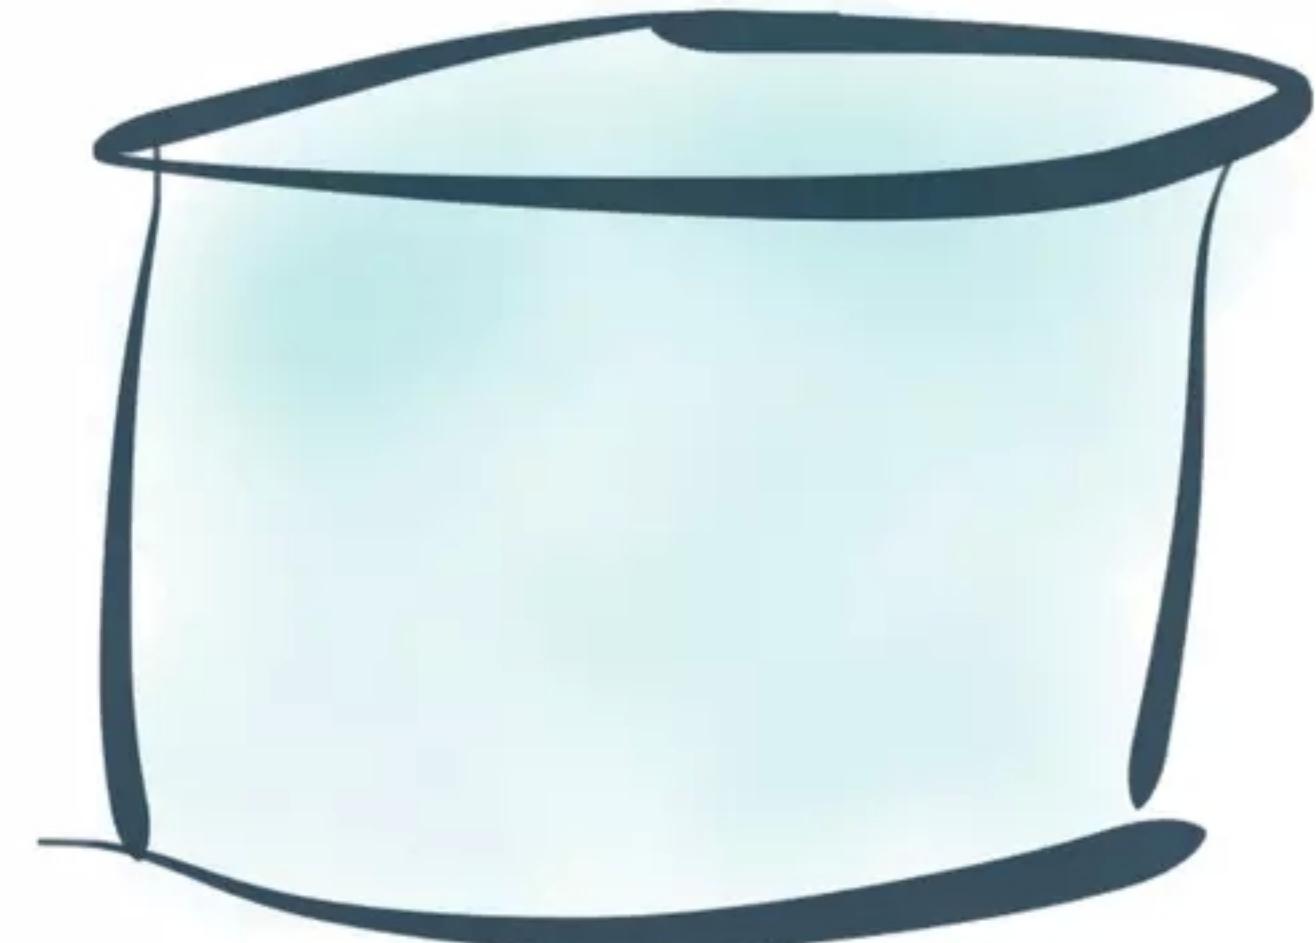
WHY  
STORE ?

# SCALABILITY OF A FILESYSTEM

- HUNDREDS OF MB/SEC THROUGHPUT
- MANY TB PER SERVER
- COMMODITY HARDWARE
- O(1) WRITES

# GUARANTEES OF A DATABASE

- STRICT ORDERING
- PERSISTENCE



# DISTRIBUTED BY DESIGN

- REPLICATION
- FAULT TOLERANCE
- PARTITIONING
- ELASTIC SCALING

# STREAMING PLATFORM

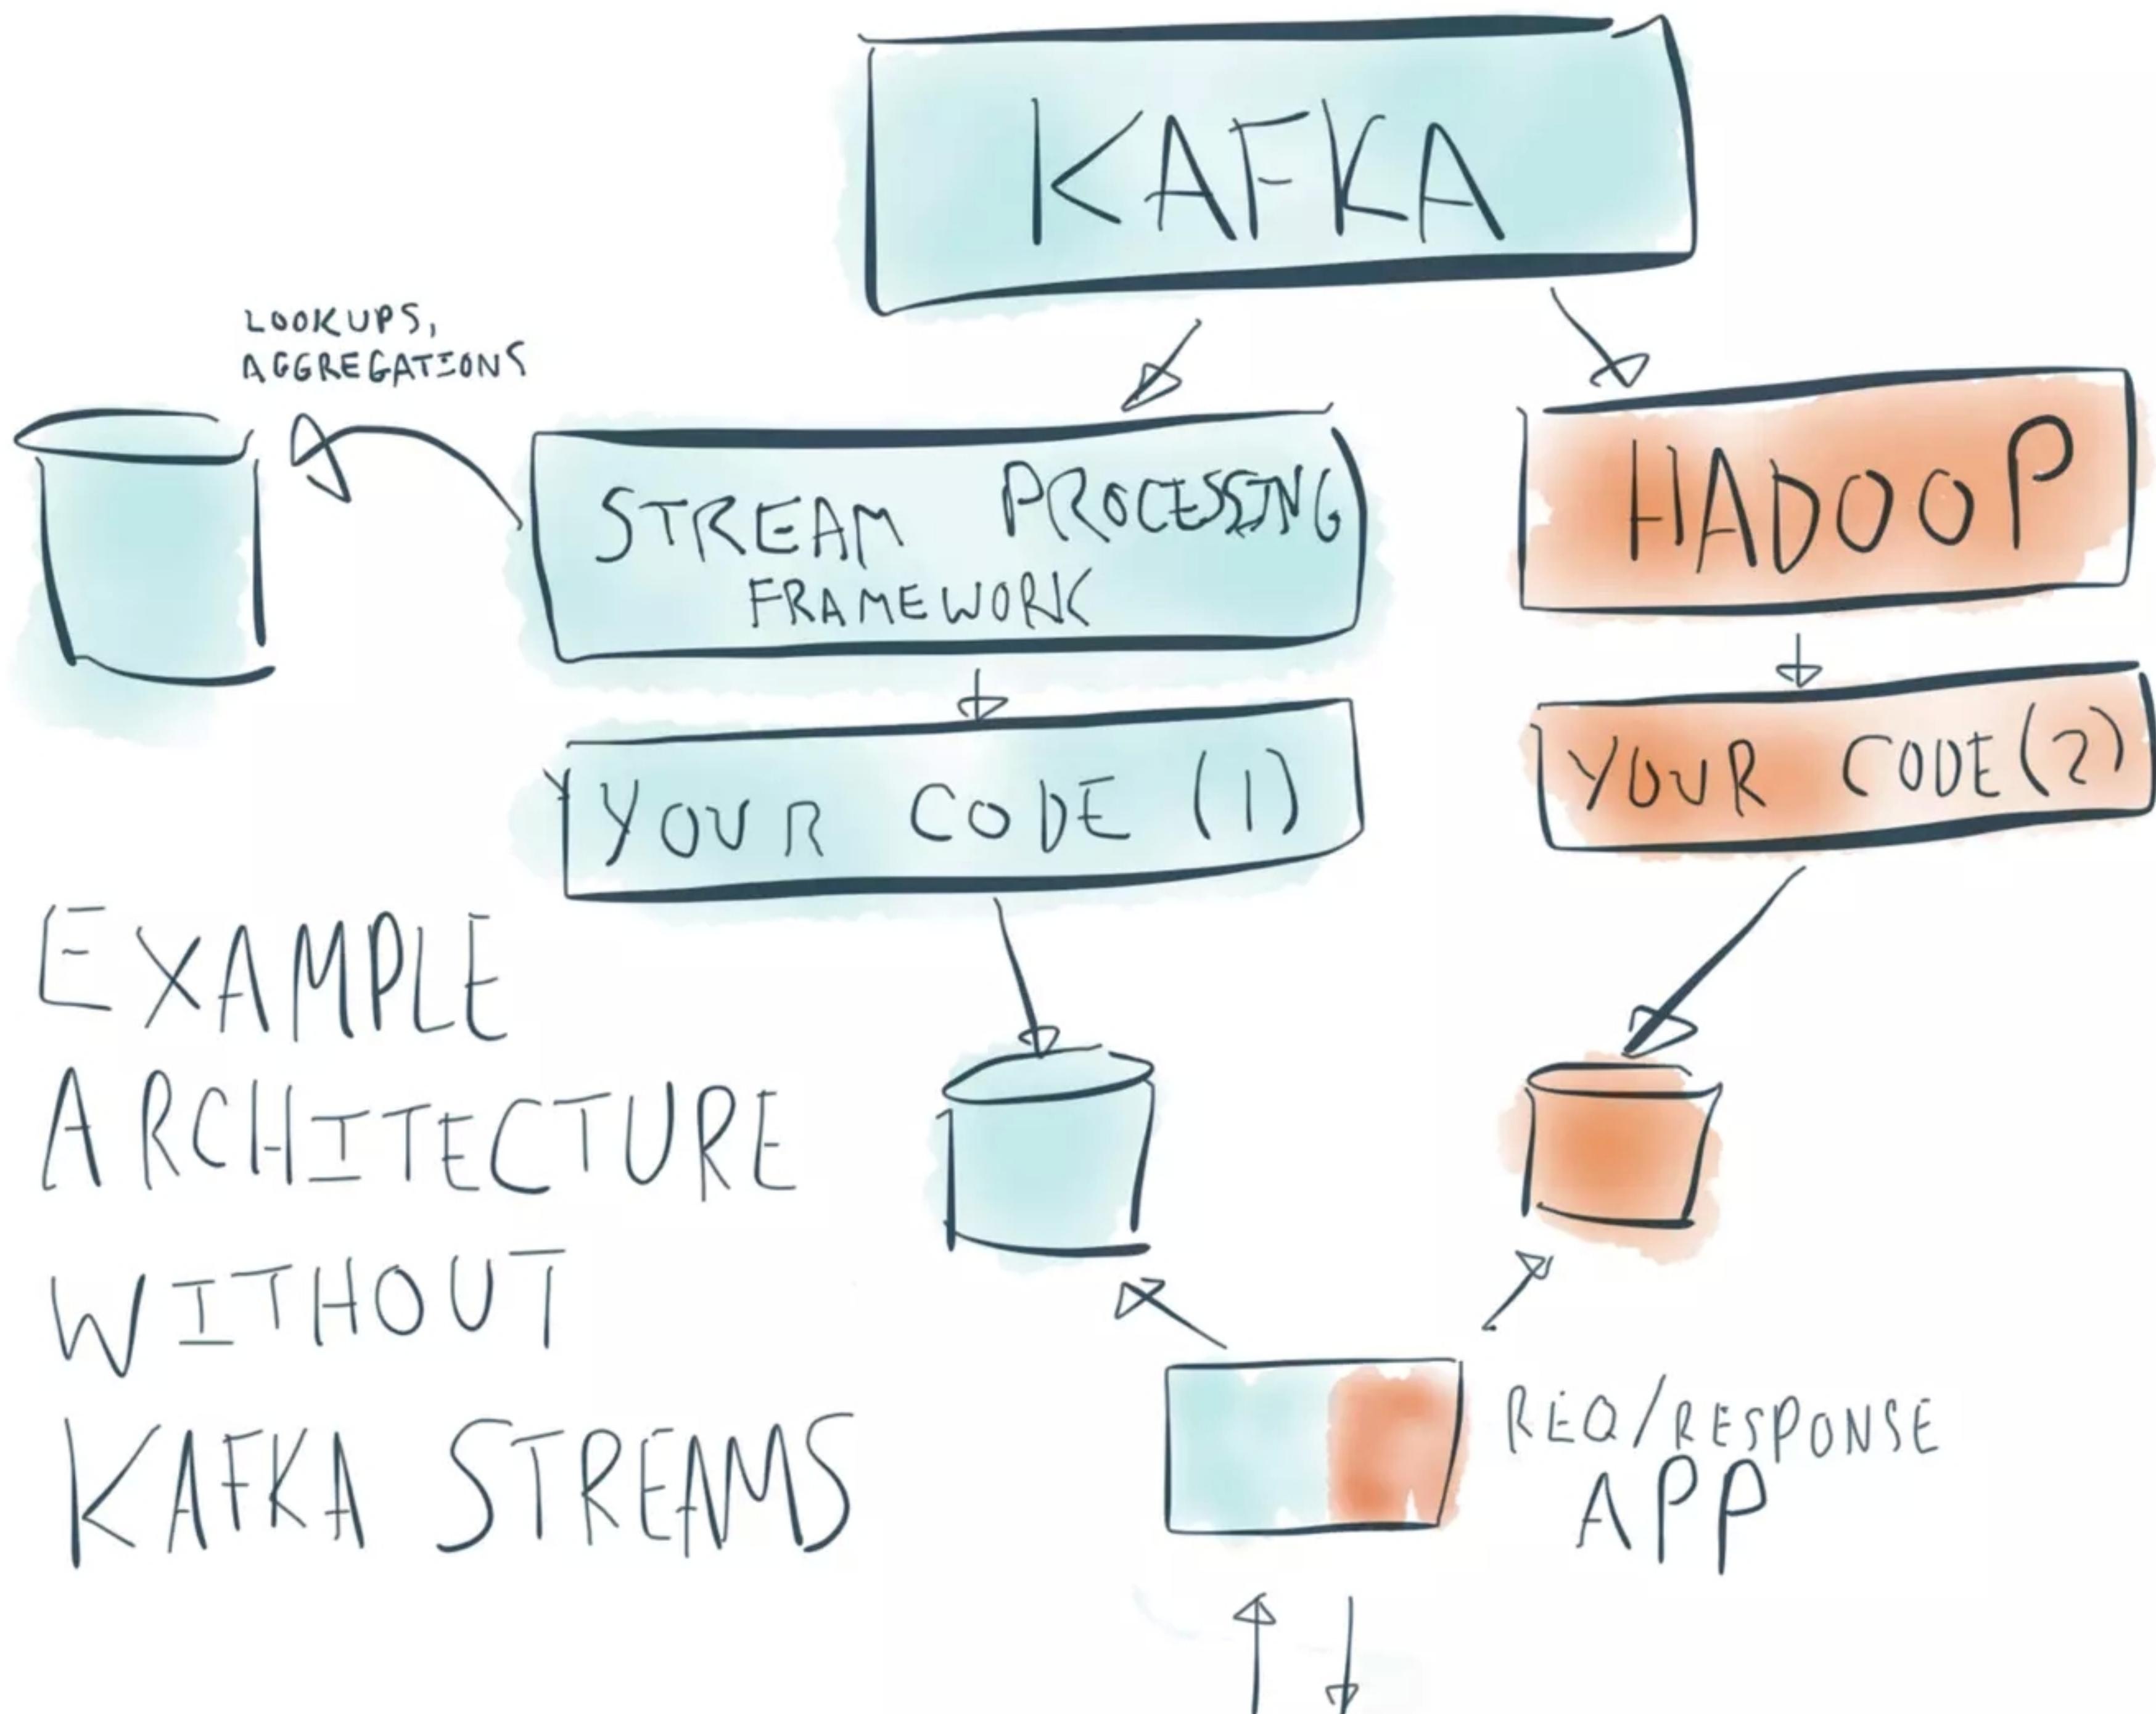
- 1 PUB/SUB
- 2 STORE
- 3 PROCESS

# KAFKA STREAMS

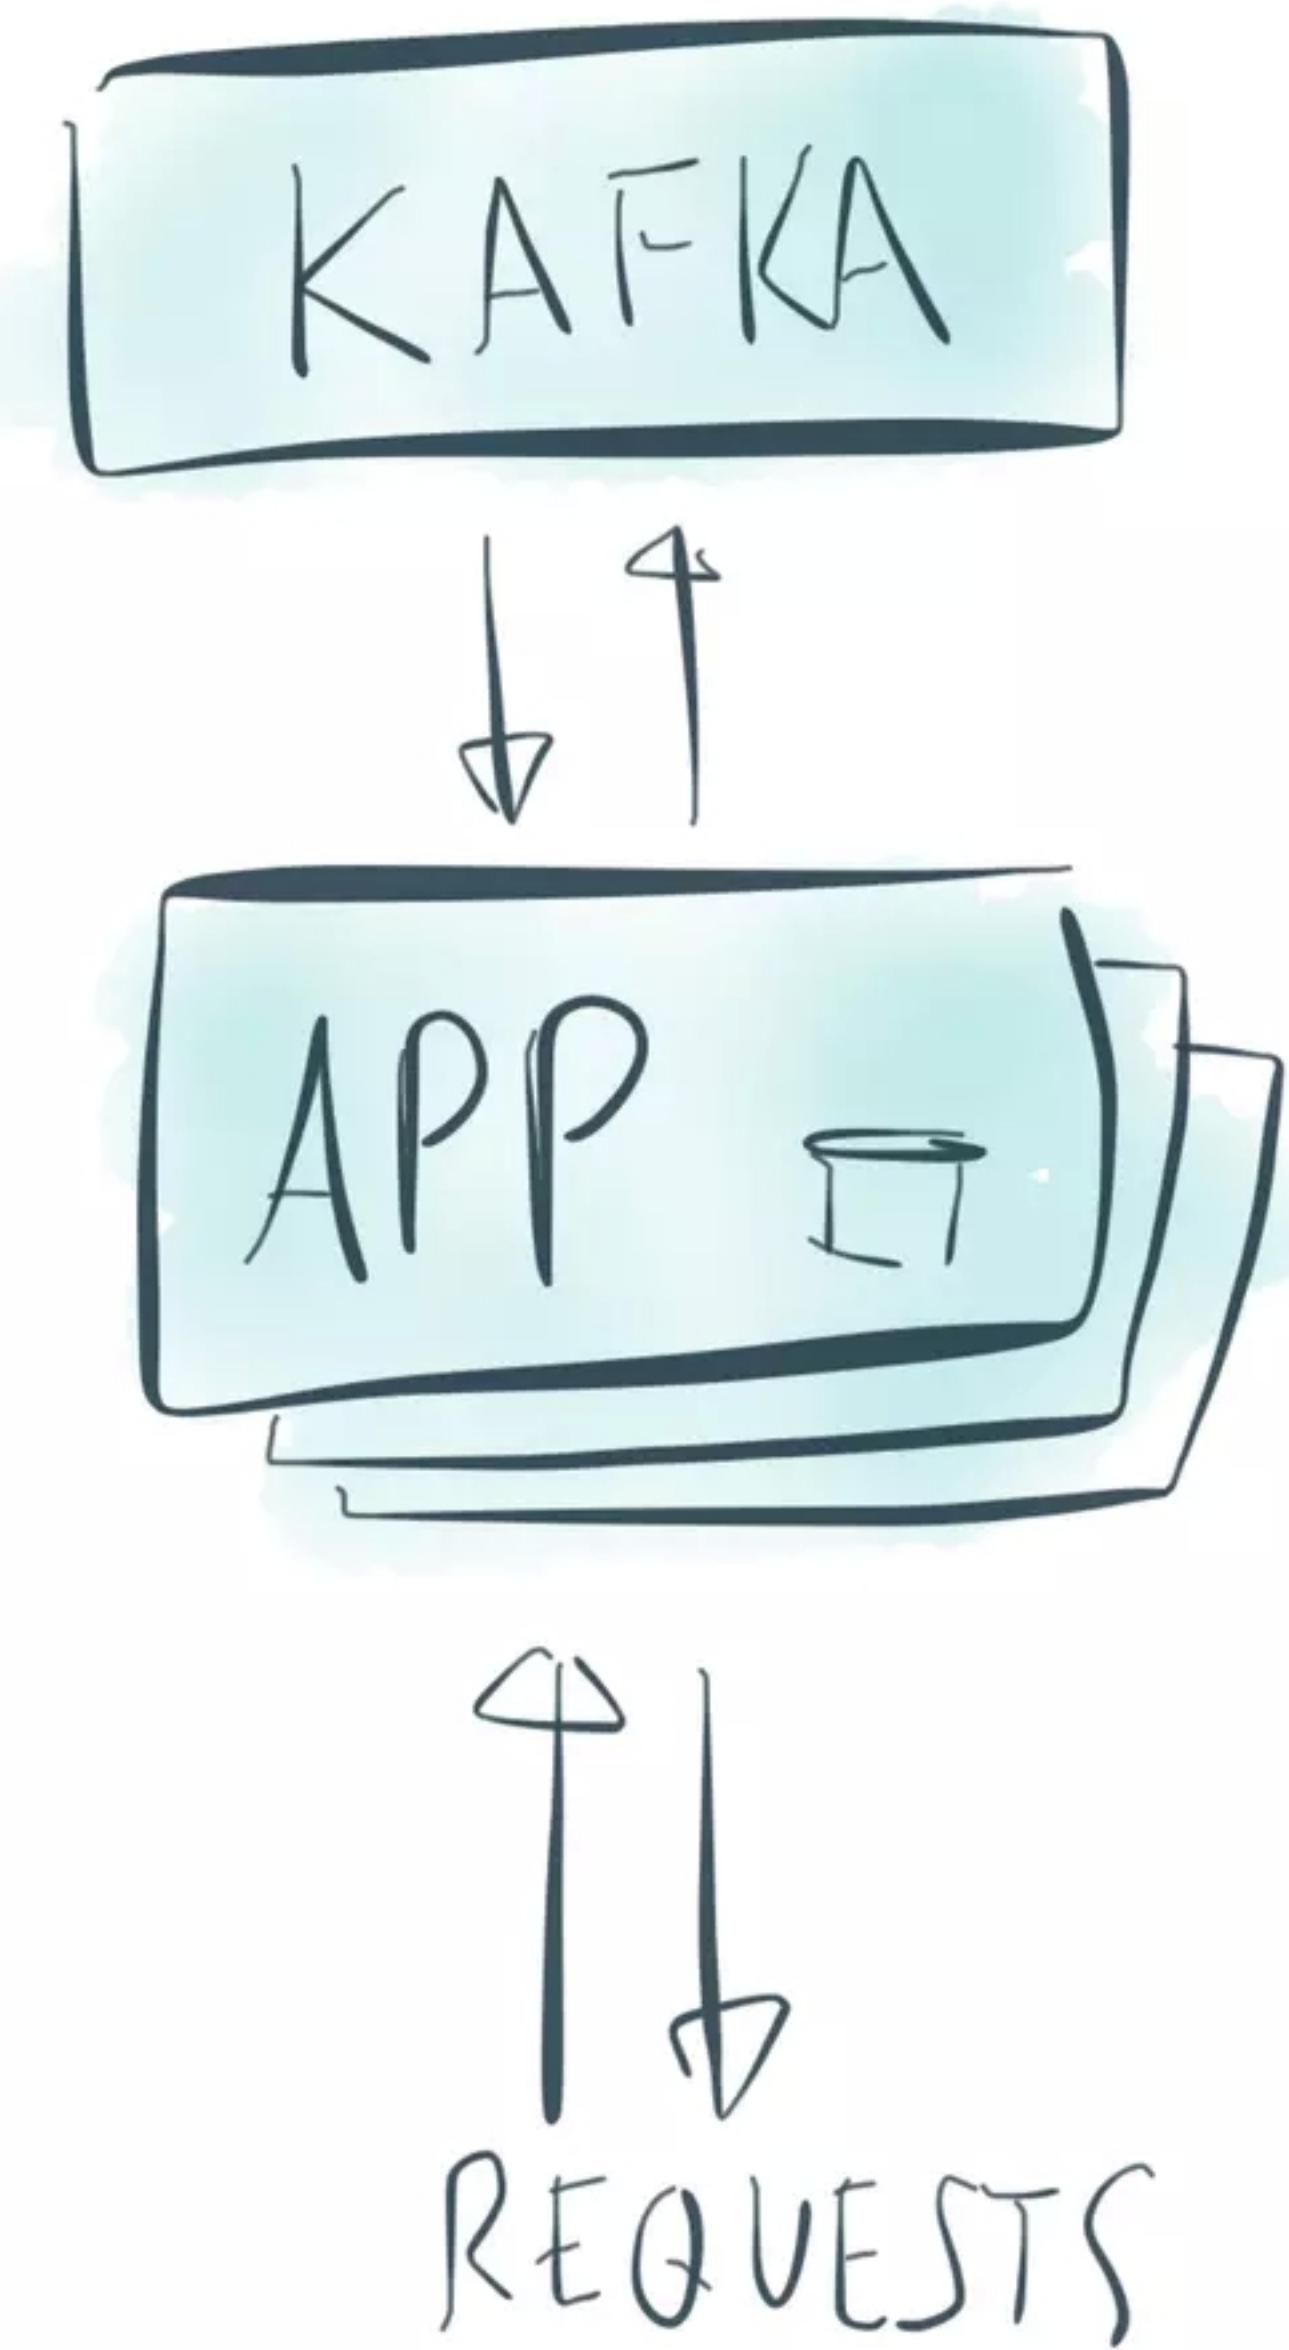
- SIMPLE LIBRARY
- CONVENIENT DSL
- DATAFLOW STYLE WINDOWING
- REPROCESSING
- NO MICROBATCH
- LOCAL STATE

# KEY OPERATIONS

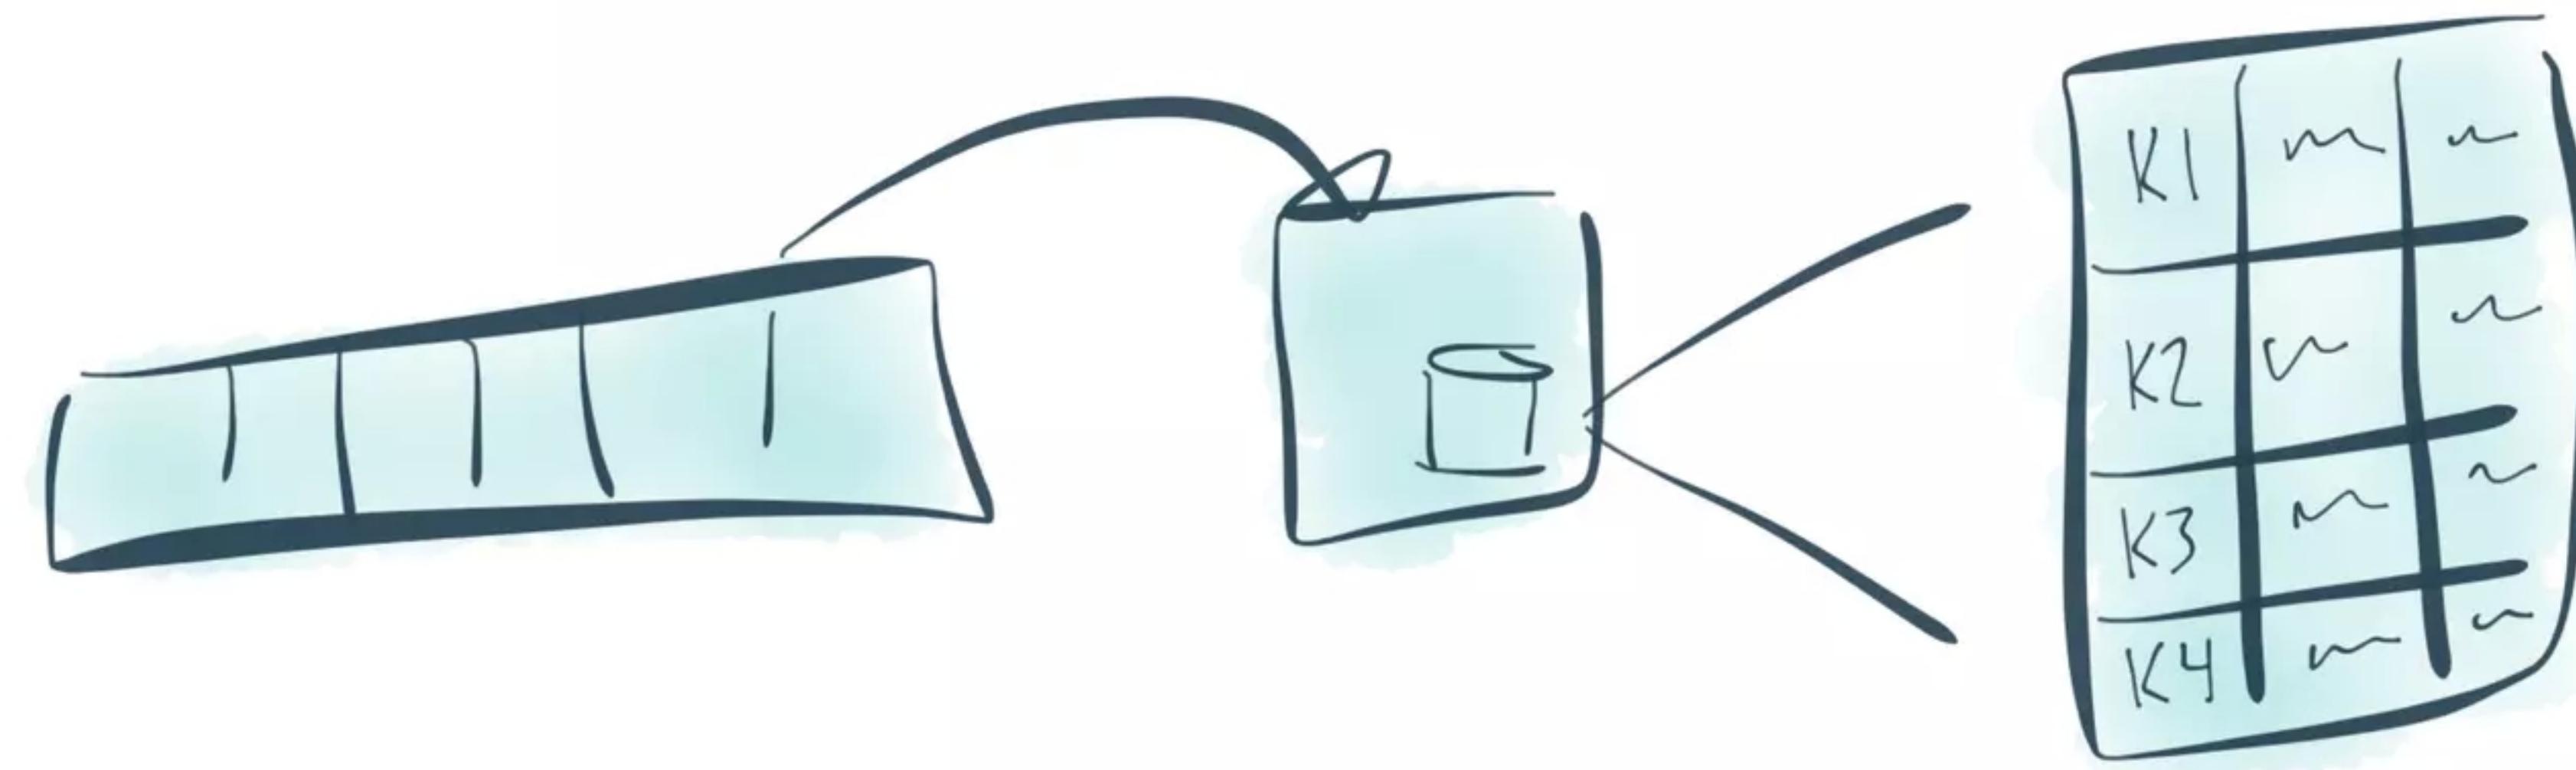
- MAP
- FILTER
- AGGREGATE (COUNT, SUM, ETC)
- JOIN



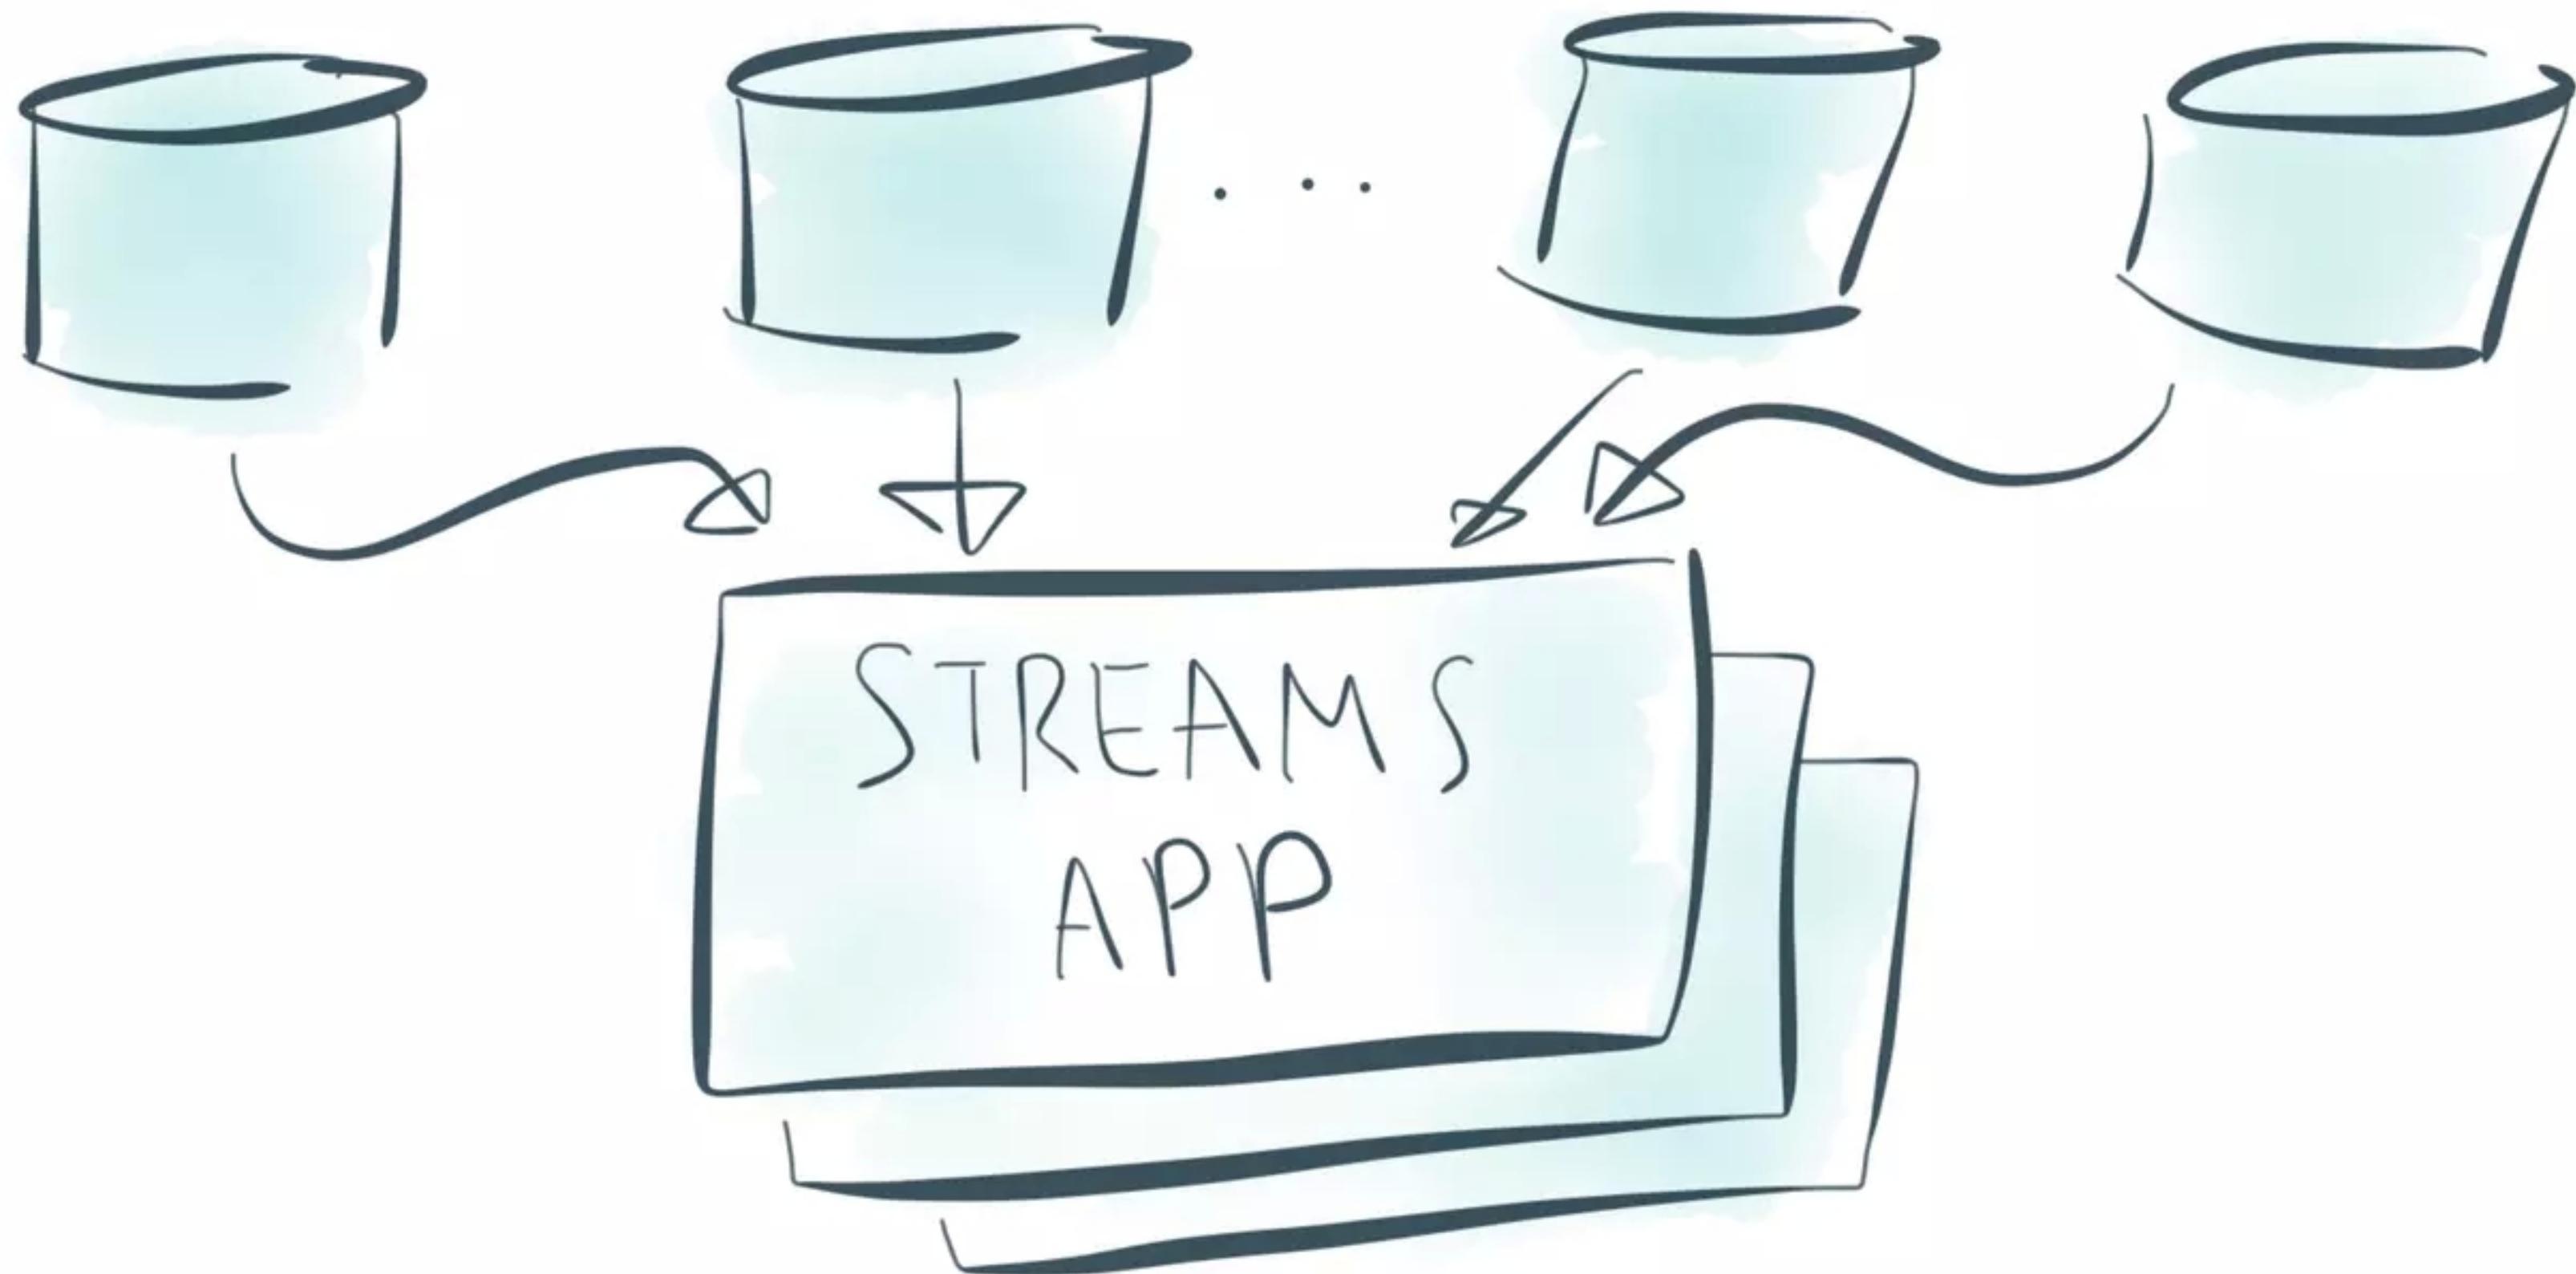
EXAMPLE  
ARCHITECTURE  
WITH  
KAFKA  
STREAMS



# UNIFY TABLES & STREAMS

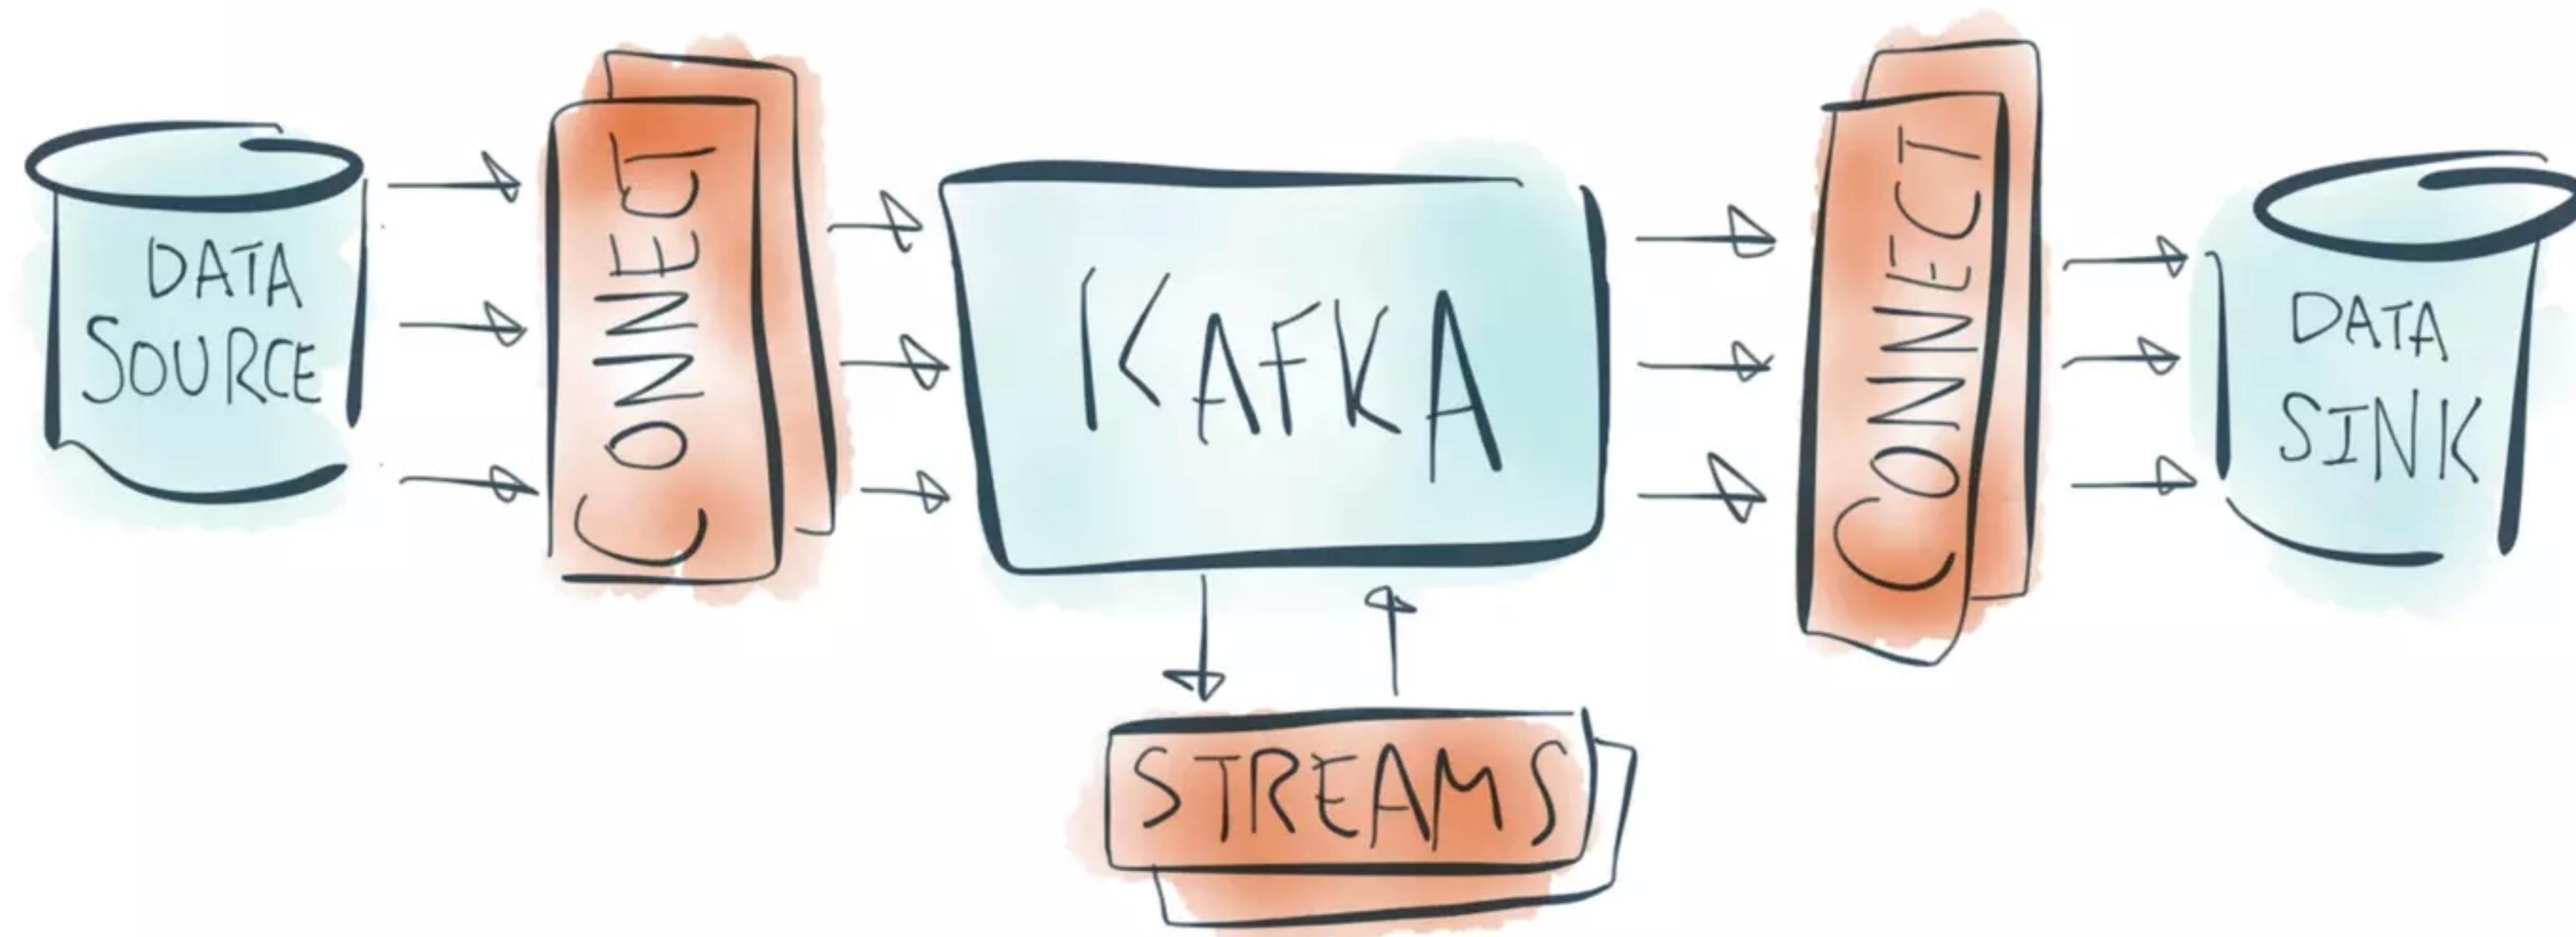


# EXAMPLE: CUSTOMER 360

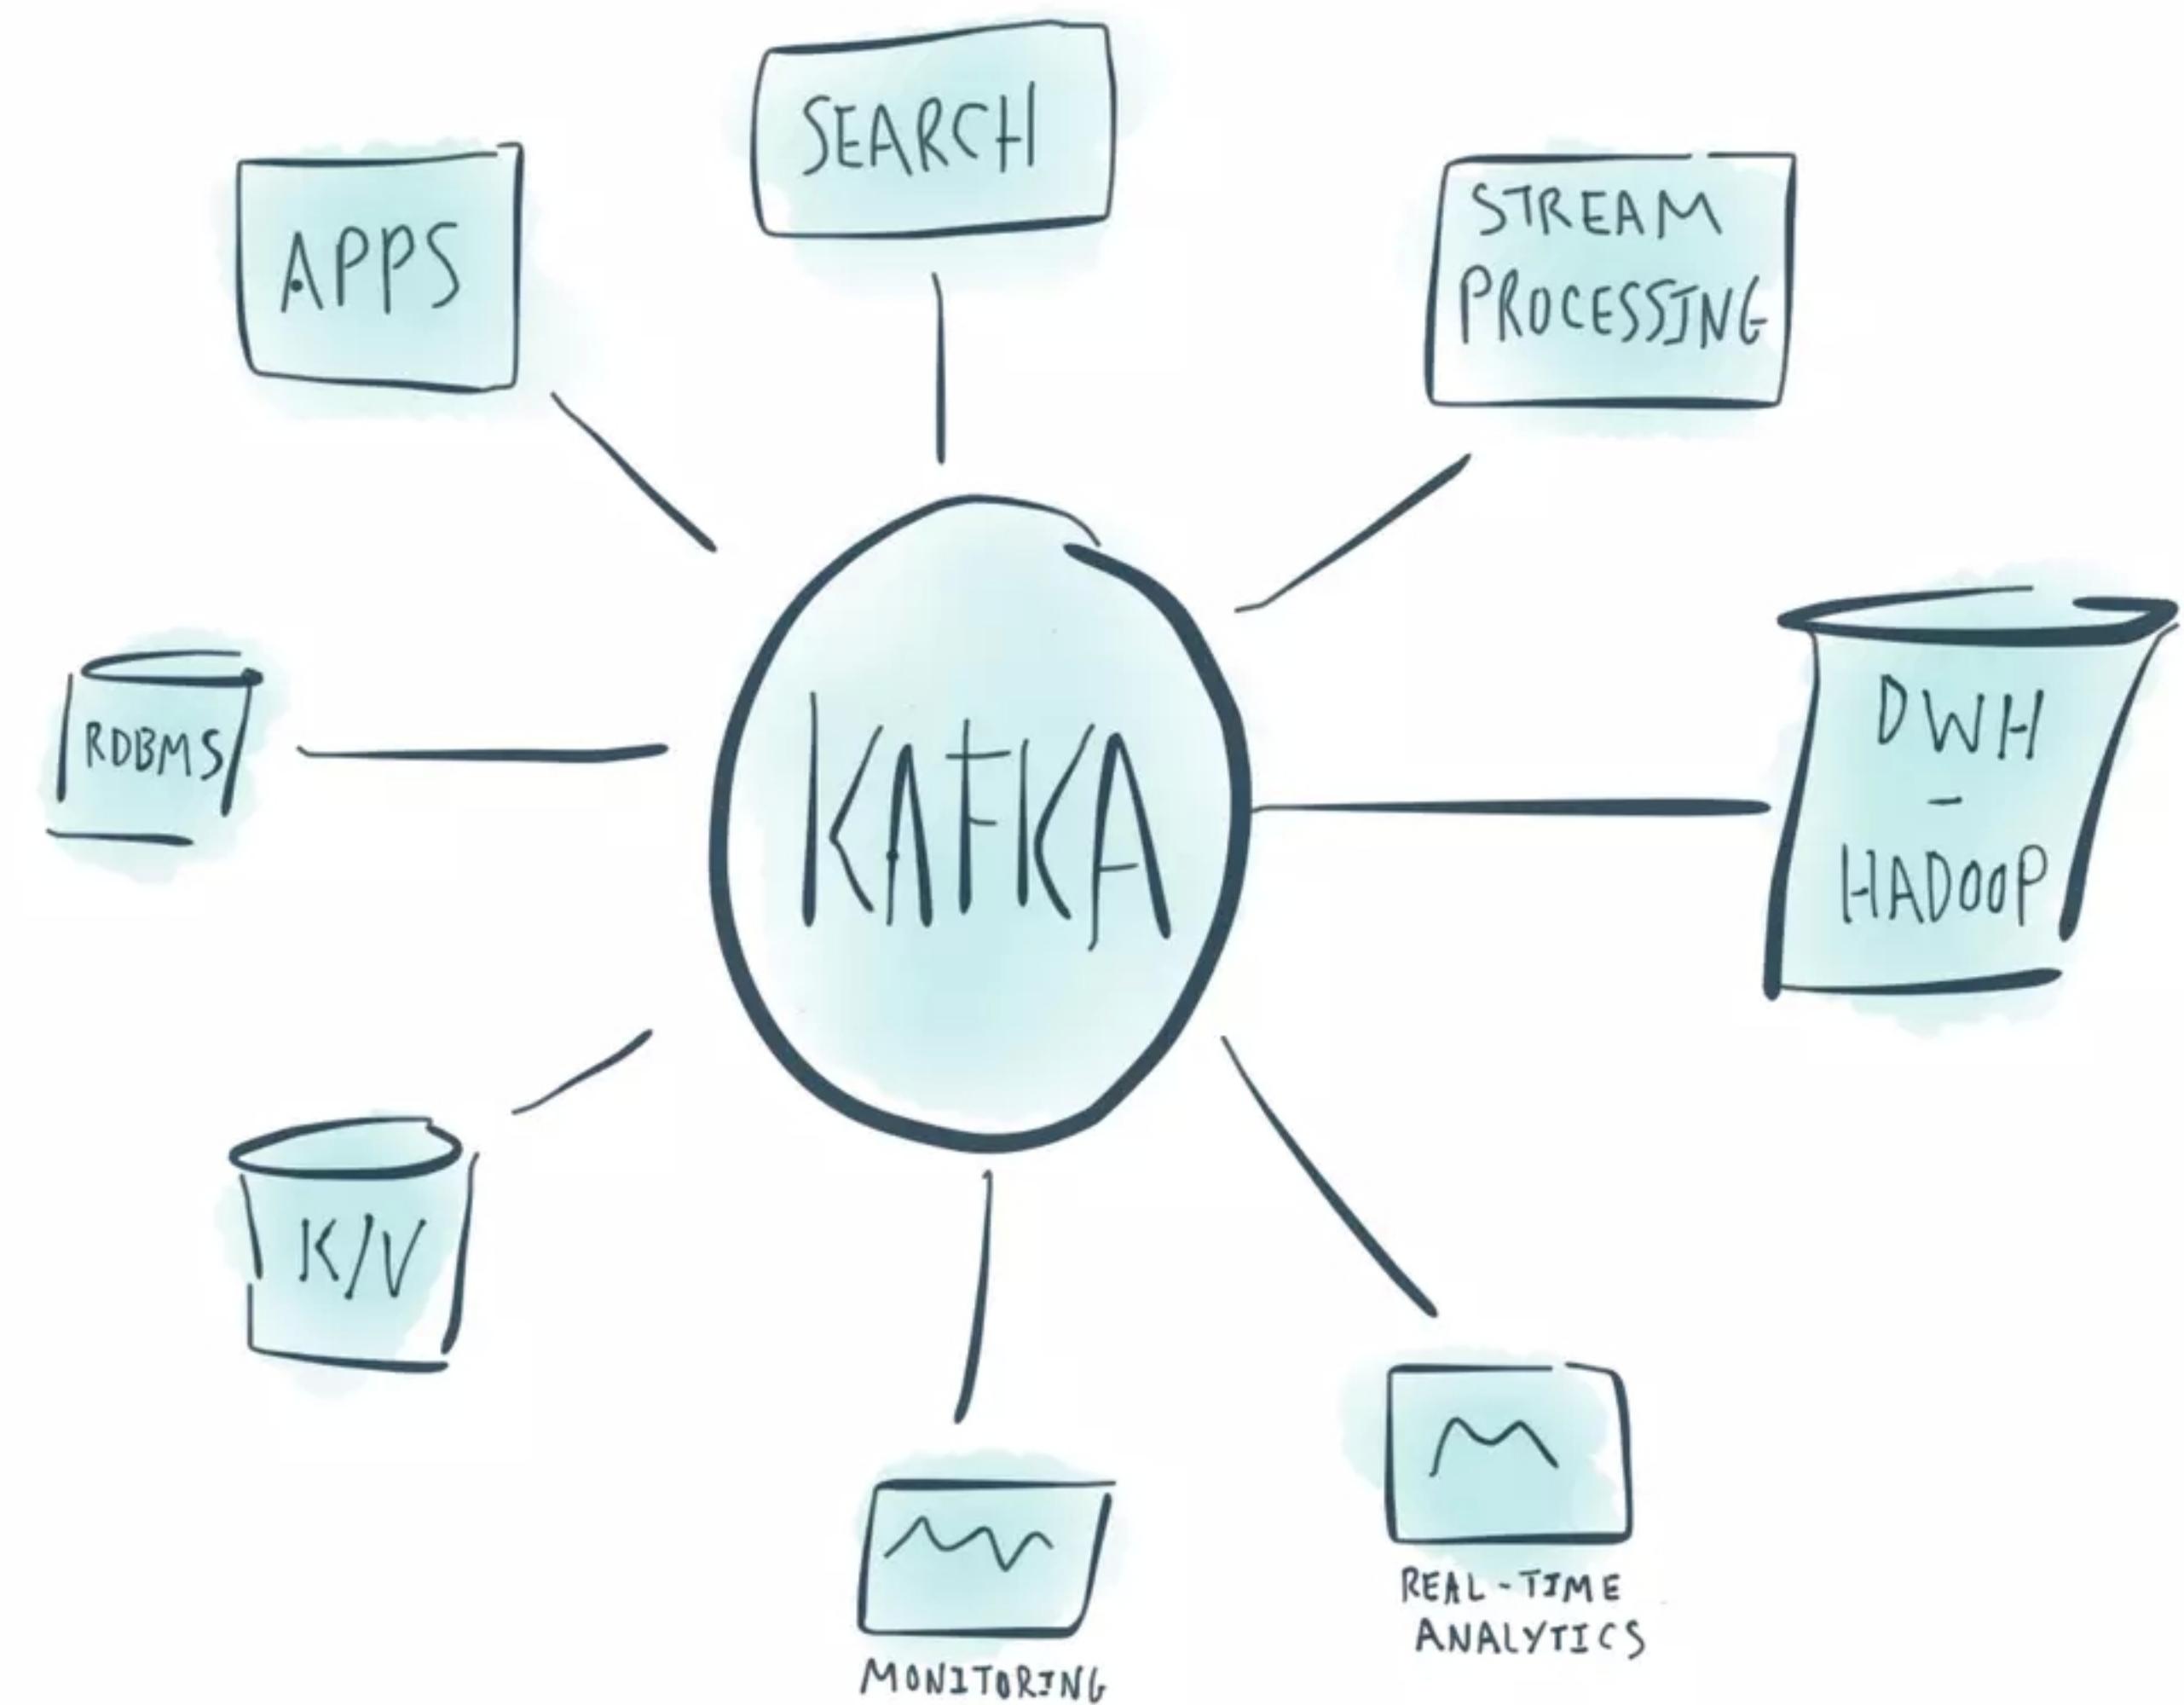


PPR  
LOOK UPS

# CONNECT + STREAMS



# STREAMING PLATFORM



# STREAMING PLATFORM ADOPTION

APP → APPS → PLATFORM

# Get Started with Apache Kafka Today!

Confluent Open Source 3.2

---

[Download ZIP](#)   [Download TAR](#)

[Install DEB](#)   [Install RPM](#)

**THE place to start with Apache Kafka!**

 Confluent Open Source

Apache Kafka packaged together with additional open source components to make streaming data better. Quickly get your Kafka environment running.

---

**Open source package**

Connectors   Clients

Schema Registry   REST Proxy

**Thoroughly tested and quality assured**

**More extensible developer experience**

**Easy upgrade path to Confluent Enterprise**



**Find your local Meetup Group**

<https://www.confluent.io/apache-kafka-meetups/>

**Join us in Slack**

<https://slackpass.io/confluentcommunity>



---



**Thank you**

@tlberglund