



Tips and Tricks for Apache Kafka™ Ops

Jen Snipes, Senior Customer Success Architect
<https://www.linkedin.com/in/jennifersnipes/>

Kat Grigg, Senior Customer Success Architect @katgrigg

June 2020

Agenda

- *What is Apache Kafka*
- *Kafka Operations*
 - *Configuring*
 - *Deploying*
 - *Maintaining*
 - *Monitoring*
 - *Debugging*

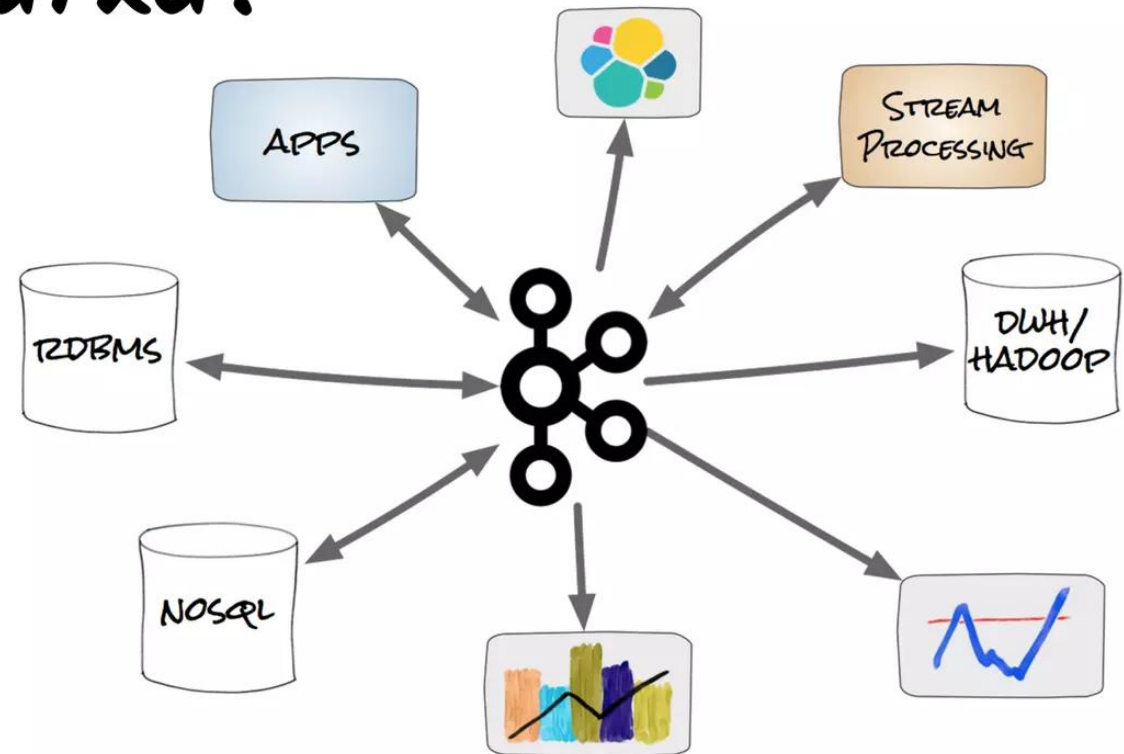
What is Apache Kafka?

Open Source

Real time Event Streaming Platform

Scalable, Distributed Log

Fault tolerant storage




Moving Data

Producers



Consumers

Connect

Streams



Log Structured Data Flow



Kafka Operations

5 Categories of Ops

Configuring

Deploying

Maintaining

Monitoring

Debugging

Configuring Kafka



Data Retention

- Disk space available?
- SLA on consuming?
- Keyed data? Cardinality?



Security

- Start with plaintext
- SSL and SASL? Choose one first
- Console clients are great for testing



Optimization

- **Throughput**
 - increase partitions, batch size, linger.ms
 - use acks 1
 - compression
 - parallelize consumers
- **Latency**
 - no compression
 - balance number of partitions
 - increase fetcher threads
 - producer linger.ms=0, consumer fetch.min.bytes=1
- **Durability**
 - acks = all, RF = 3, minISR = 2
 - leverage retries
 - unclean leader election set to false
 - disable auto commit



- **Availability**
 - minISR 1
 - unclean leader election true
 - increase recovery threads
 - low consumer session.timeout
 - optimize number of partitions

Deploying Kafka



Start Easy

- How do you run your database(s)?
- Do you know Docker well?
- Do you know Kubernetes well?
- If no, then don't use them
- Cloud an option?



One does not simply start Brokers

- JMX access (metrics reporter)
- Environment Variables
- Logging capacity
- Don't forget ZooKeeper
- File descriptors (ulimit)
- Disk deployment style



Maintaining Kafka

Updates and Upgrades



Rolling upgrade Tips

- [Gwen Shapira's Kafka Summit Talk 2019](#)
- Set protocols and message formats
- Be explicit
- Go slow, too fast and ISR sadness
- Upgrade often, point releases in prod



Rebalancing

- Lose/Add Broker
- Throttle early, adjust later
- Throttles don't affect in sync replicas
- Not sure? Batch the reassign
- See Monitoring



Updating Configs

Brokers

- Lots of [dynamic config](#) options!
- Sync config across the cluster

Topics

- [Overrides](#) are your friend
- Topic level unclean leader election



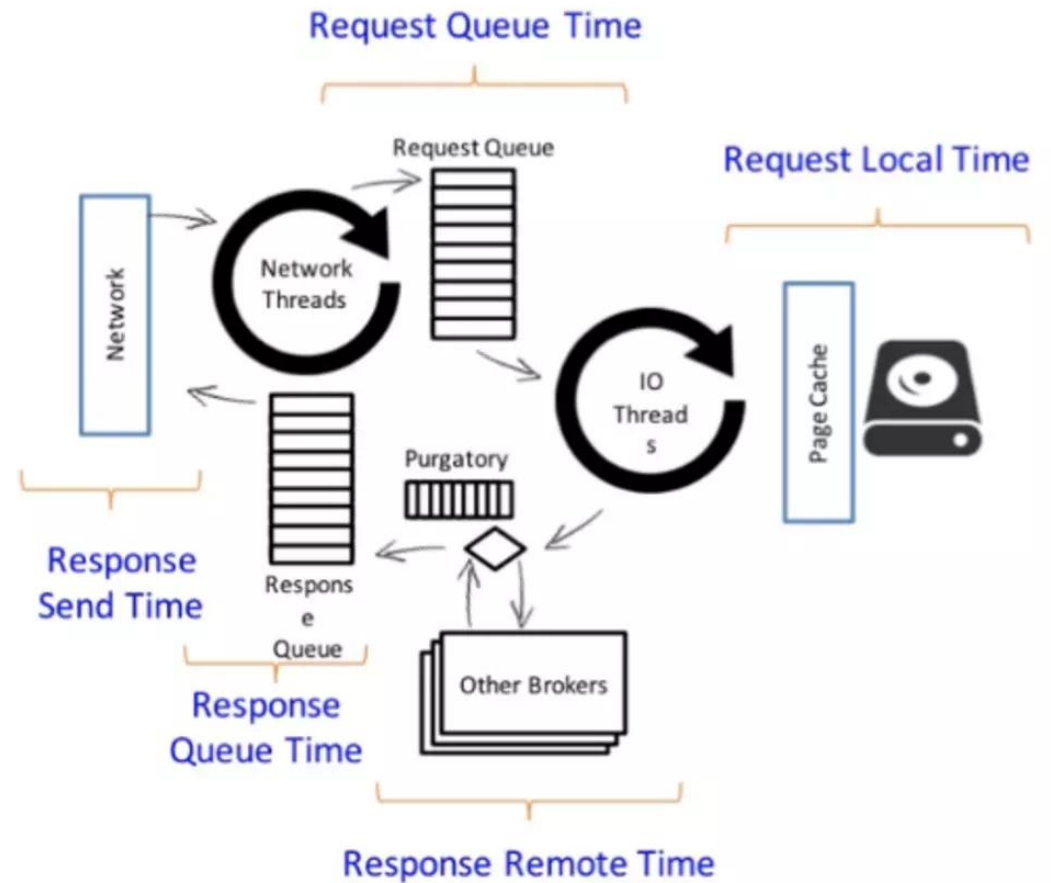
Monitoring Kafka

What do you watch?



Kafka Thread Model

- Monitor all of these hops with JMX
- Transparent way to see activity
- Especially watch request queue size
- Don't blindly bump thread count



Log Cleaner / Others

- Watch your log cleaner with JMX
- Bytes In/Out just makes sense
- Don't need to dashboard everything
- Alert don't automate (at least at first)



Broker Metrics

Top 5 Broker JMX metrics you should be watching

- UnderReplicatedPartitions
- RequestHandlerAvgIdlePercent
- NetworkProcessorAvgIdlePercent
- RequestQueueSize
- TotalTimeMs



Clients / Consumers

- batch sizes
- failed-authentication-rate
- io-ratio
- io-wait-ratio
- io-wait-time-ns-avg



Debugging Kafka

When there's trouble...



Logs about the Log

- Debugging? Head to the logs!
 - Defaults separates files reasonably
 - Splunk, ELK, etc.? filter by appender
 - Metadata files are important too
- **controller.log** -- written to by active controller
 - **state-change.log** -- updates received from and sent to controller
 - **server.log** -- basic broker operations (segment rolls, replica fetching, etc.)
 - **log-cleaner.log** -- cleaner thread (related to compaction)
 - **kafka-authorizer.log** -- secure connections only (set to DEBUG if you want successful connections)
 - **kafka-request.log** -- super verbose, will talk about more
 - **/var/lib/recovery-point-offset-checkpoint** – last offset flushed to disk
 - **/var/lib/cleaner-offset-checkpoint** – offset up to which the cleaner has cleaned (compacted topics only)
 - **/var/lib/replication-offset-checkpoint** – last committed offset

Request logging

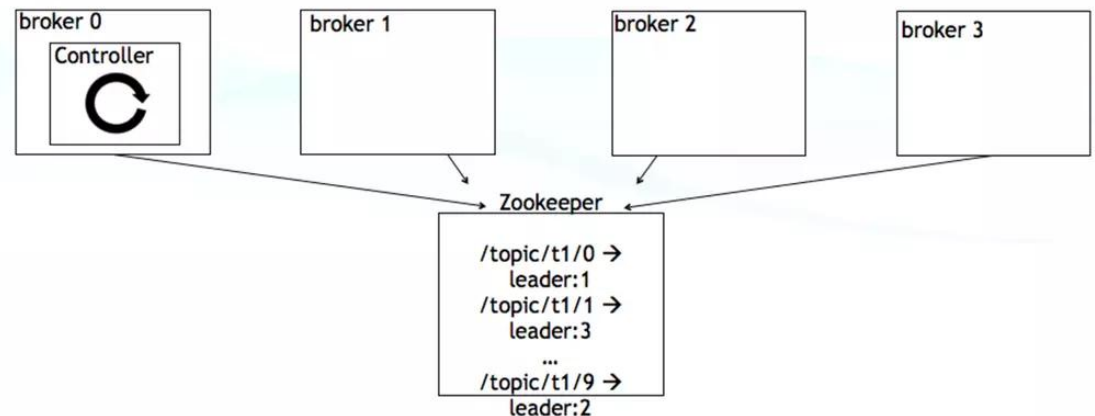
- kafka-request.log
- All information about all requests
- Verbose, not great for live debugging
- Use when no other options for root cause
- Rogue clients, single request slowdowns, etc.



Controller Issues

Partition state: stored in ZK, cached in controller

- [Kafka Controller Deep Dive](#)
- [Kafka Needs No Keeper](#)
- ISR updates aren't happening
- Replication stopped for no reason
- Broker refusing to become leader for extended period
- `rmmr /controller`



The Offsets Topic

- `__consumer_offsets` → Stores offsets and group metadata
- ***`bin/kafka-run-class kafka.tools.DumpLogSegments --offsets-decoder`***
- `offsets.retention.minutes` default 1 day (7 days on 2.0)
- Growing huge? Check the log cleaner health
- High Memory Usage on Brokers? Offsets are cached
- High CPU Usage on Brokers? Make sure this topic isn't huge
- Please replication factor 3

Restarting? Think about it...

- Restarts in distributed systems are tricky
- Have a reason
- Understand startup time
- Clean shutdowns, limited scope



References

- <https://www.confluent.io/resources/kafka-the-definitive-guide/>
- <https://www.confluent.io/kafka-summit-san-francisco-2019/please-upgrade-apache-kafka-now>
- <https://conferences.oreilly.com/strata/strata-ny-2018/public/schedule/detail/68957>
- <https://www.confluent.io/kafka-summit-san-francisco-2019/kafka-needs-no-keeper>
- <https://events.confluent.io/meetups>

**BLACK
LIVES
MATTER**