



Apache Kafka™ Fundamentals: What Every Software Engineer Should Know about Streams and Tables in Kafka

Michael G. Noll

Senior Technologist, Office of the CTO
@miguno

Kafka Meetup Zurich, Switzerland, Feb 2020

Event Streaming

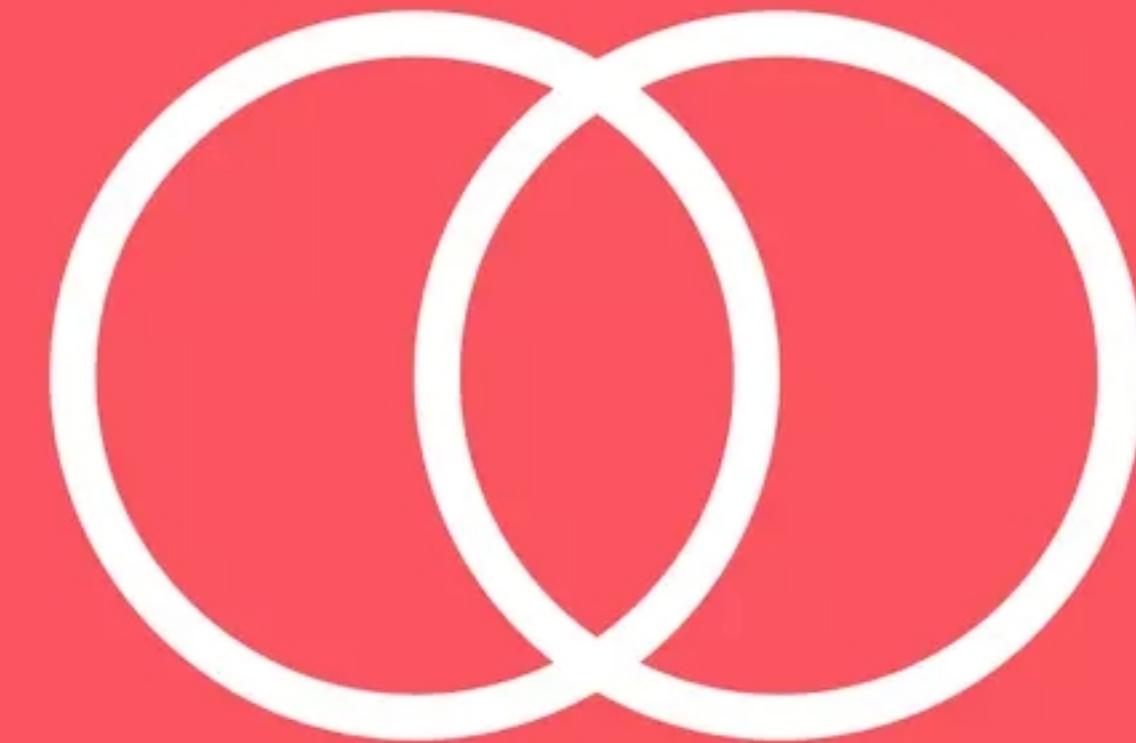
Streams



Tables

Event Streaming

Streams

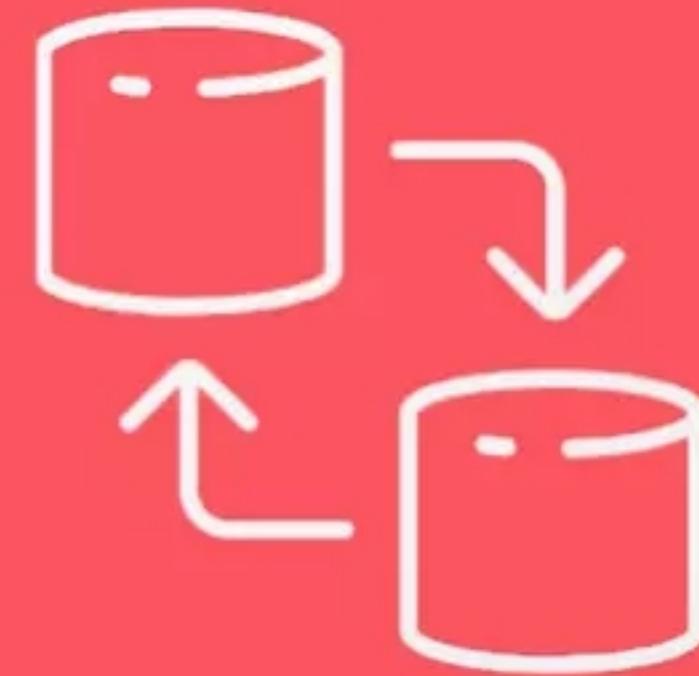


Tables

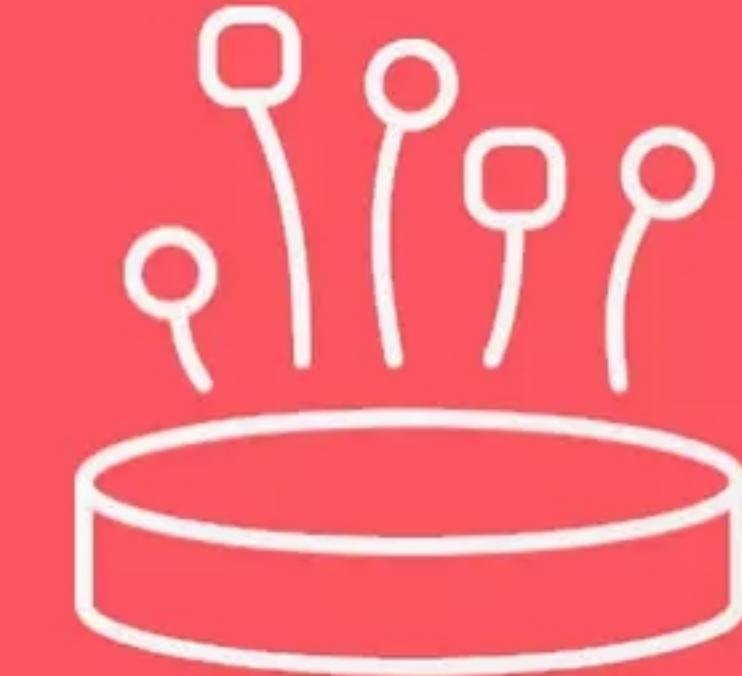
Events, Streams, Tables

A Primer

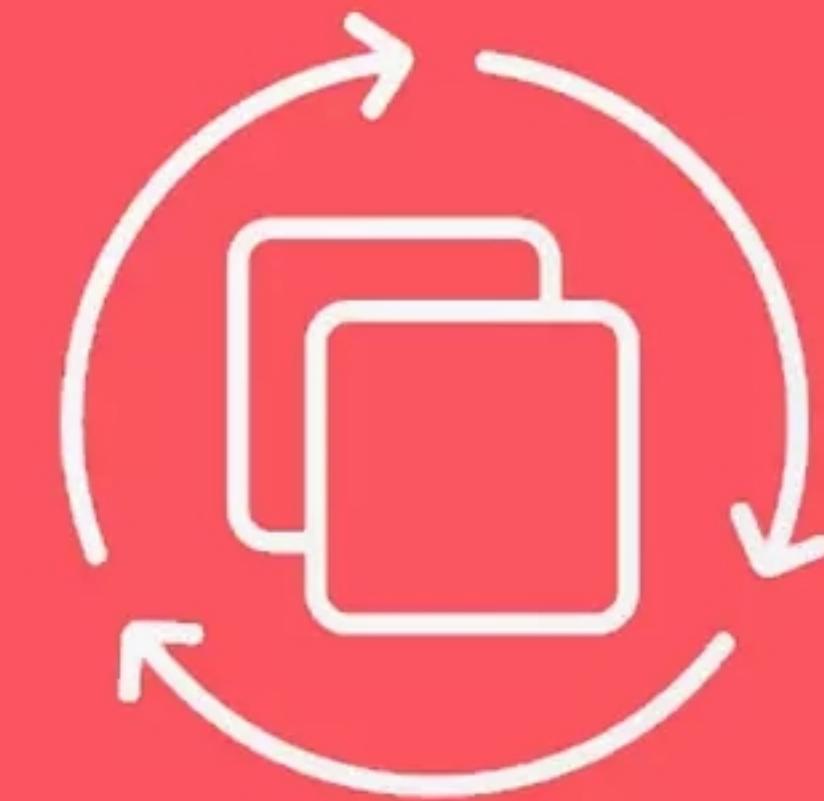
An Event Streaming Platform gives you three key functionalities



**Publish & Subscribe
to Events**



**Store
Events**



**Process & Analyze
Events**

An Event records the fact that something happened



A good
was sold



An invoice
was issued to
Michael



Alice paid
\$200 to Bob
on Feb 01, 2020
at 8:31am



A new customer
registered

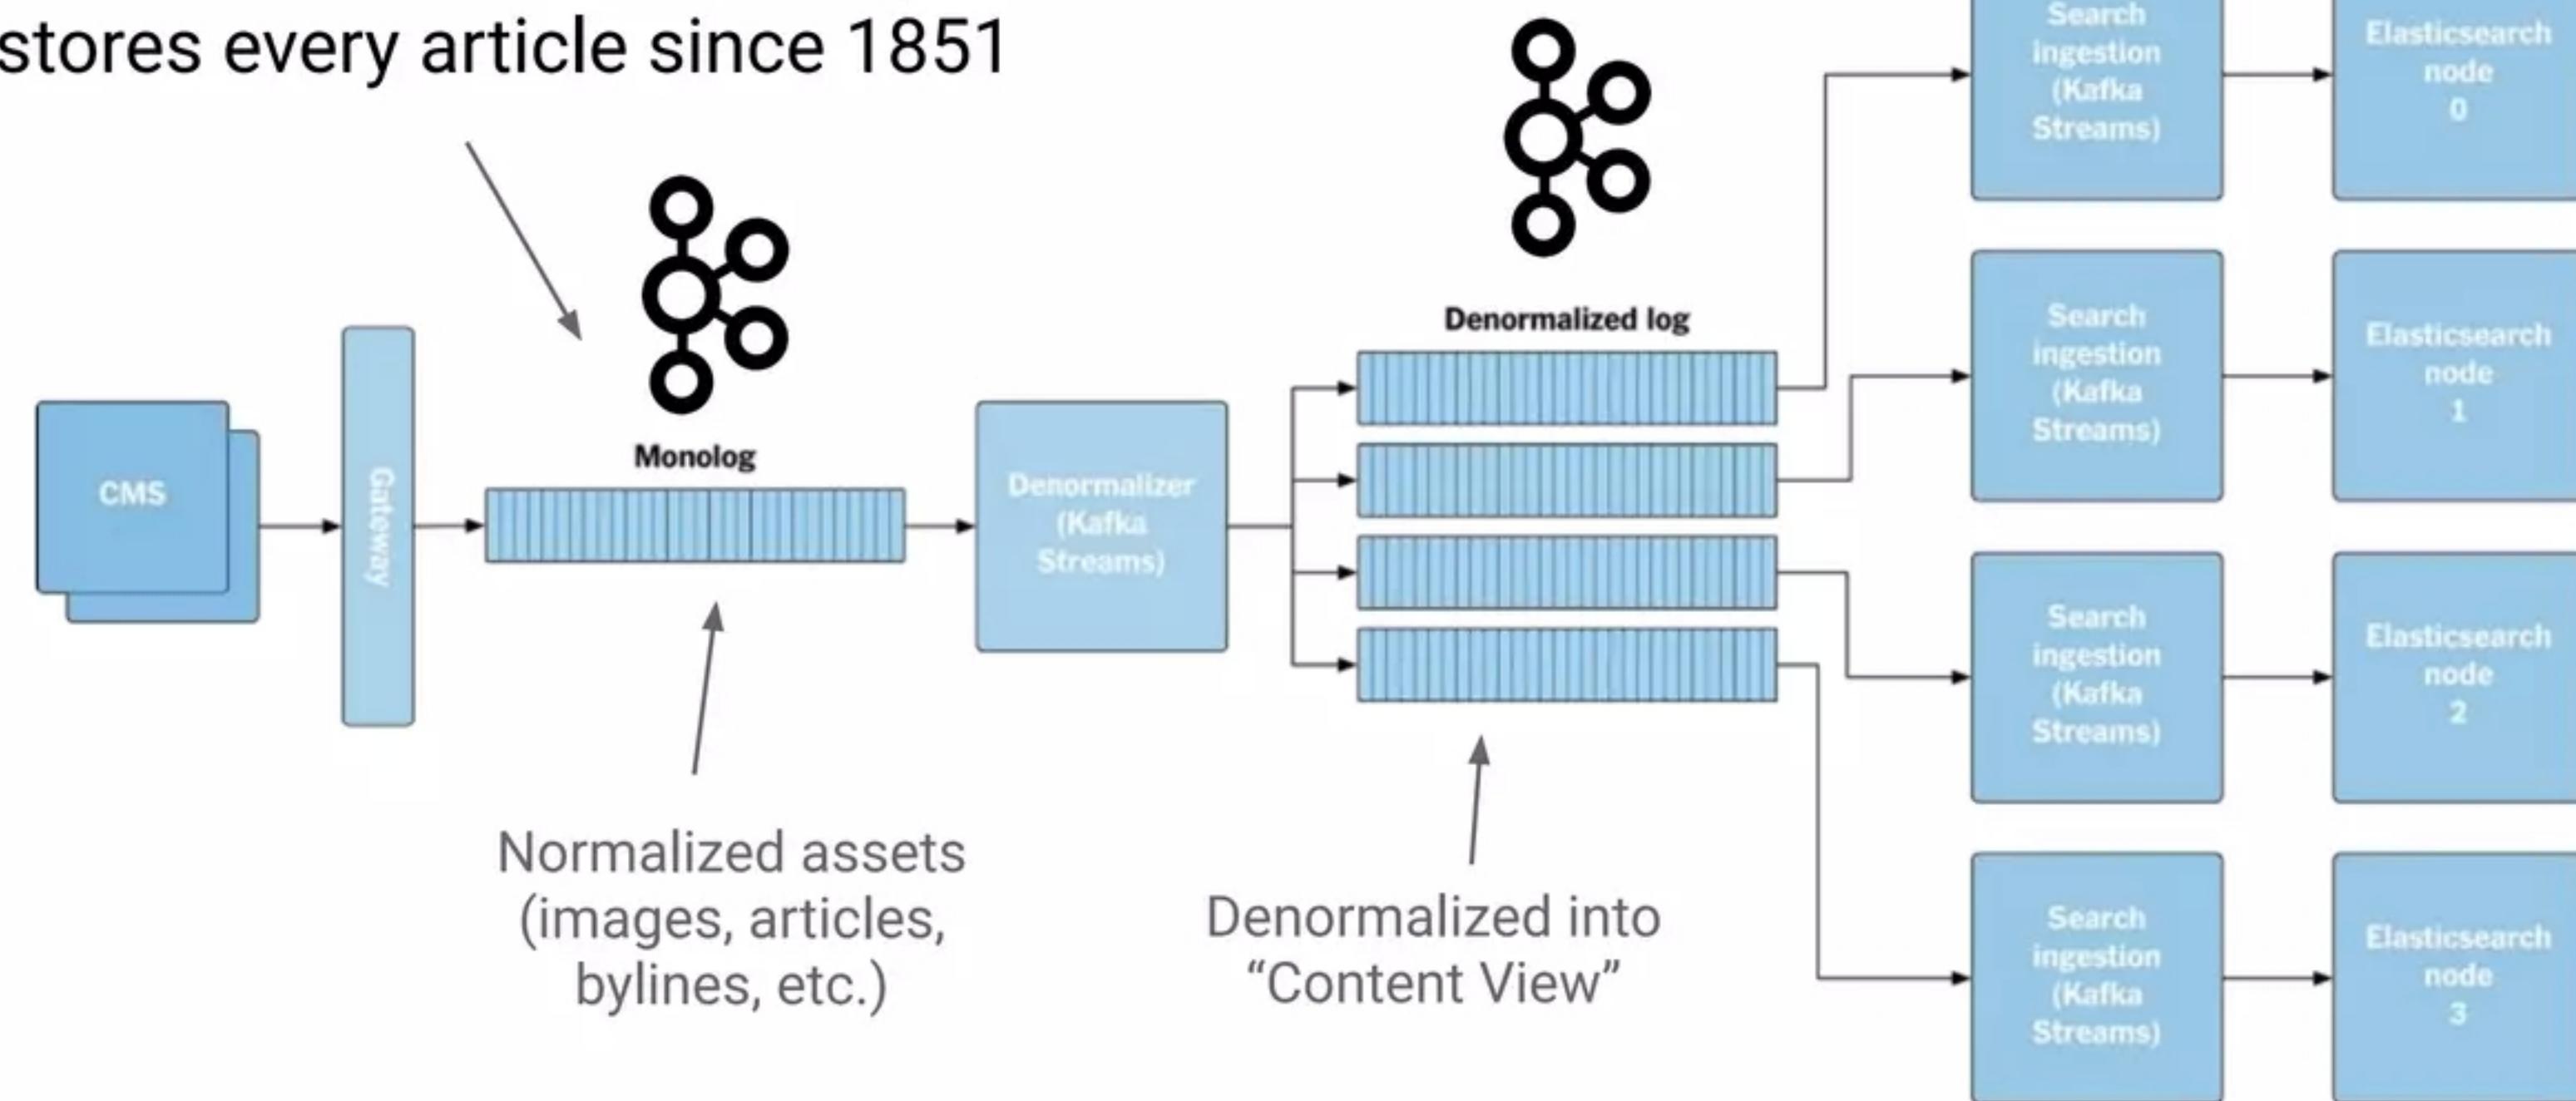
A Stream records history as a sequence of Events



Streams record history, even hundreds of years

The New York Times

Kafka = Source of Truth,
stores every article since 1851



**Streams
record history**



“The ledger of sales.”

**Tables
represent state**



“The sales totals.”

Streams record history

1. e4 e5
2. Nf3 Nc6
3. Bc4 Bc5
4. d3 Nf6
5. Nbd2

“The sequence of moves.”

Tables represent state



“The state of the board.”

**Streams = INSERT only
Immutable, append-only**

**Tables = INSERT, UPDATE, DELETE
Mutable, row key (event.key) identifies which row**

The key to mutability is ... the event.key!

	Stream	Table
Has unique key constraint?	No	Yes
First event with key bob arrives	INSERT	INSERT
Another event with key bob arrives	INSERT	UPDATE
Event with key bob and value null arrives	INSERT	DELETE
Event with key null arrives	INSERT	<ignored>

RDBMS analogy: A Stream is ~ a Table that has no unique key and is append-only.

Creating a table from a stream or topic

Table



```
KTable<String, String> table  
= builder.table("input-topic",  
    Consumed.with(Serdes.String(),  
        Serdes.String()));
```



```
CREATE TABLE myTable  
(username VARCHAR, location VARCHAR)  
WITH (KAFKA_TOPIC='input-topic',  
      VALUE_FORMAT='...', ...)
```

Stream (facts)



alice	Paris
bob	Sydney
alice	Rome
bob	Lima
alice	Berlin

bob	Lima
alice	Berlin

Table (dims)

alice	Berlin
bob	Lima

Streams record history

Tables represent state

--	--	--	--	--	--



table changes

aggregation
like SUM(), COUNT()

Aggregating a stream (COUNT example)

Table



```
KTable<String, Long> locationsPerUser  
= locationUpdatesStream  
    .groupBy((k, v) -> v.username)  
    .count();
```



```
CREATE TABLE locationsPerUser AS  
SELECT username, COUNT(*)  
FROM locationUpdatesStream  
GROUP BY username  
EMIT CHANGES;
```

Aggregating a stream (COUNT example)

Stream
(changelog)

Table

Stream alice | Paris

older data

newer data



```
KTable<String, Long> locationsPerUser  
= locationUpdatesStream  
    .groupBy((k, v) -> v.username)  
    .count();
```



```
CREATE TABLE locationsPerUser AS  
SELECT username, COUNT(*)  
FROM locationUpdatesStream  
GROUP BY username  
EMIT CHANGES;
```

Topics, Partitions, and Other Storage Fundamentals

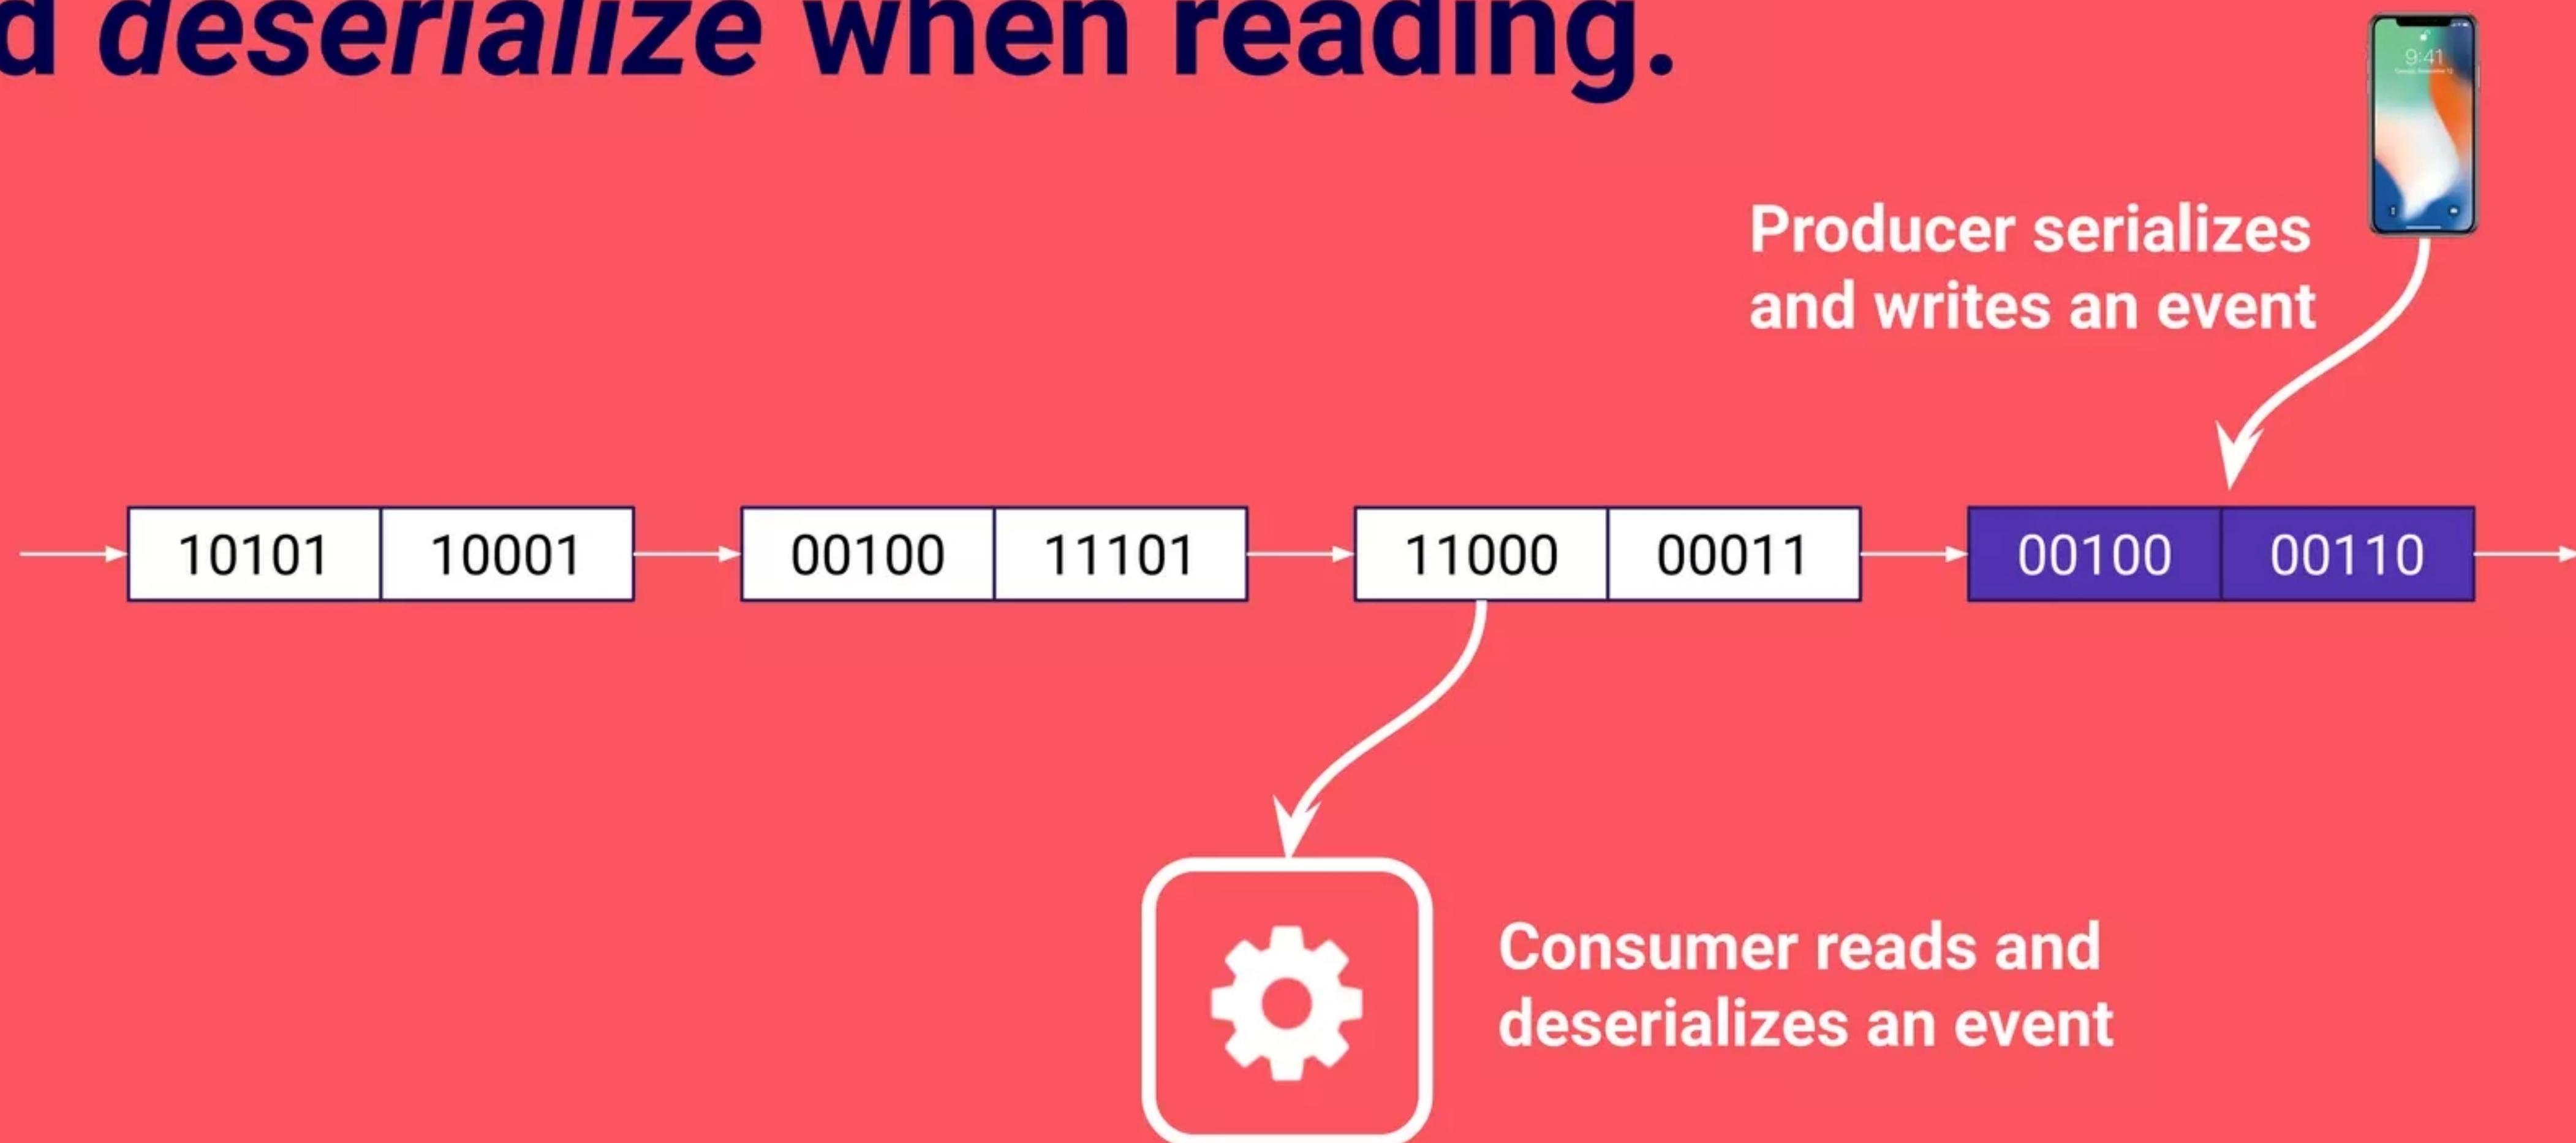
**A Topic
is where events are durably stored,
like a file in a distributed file system**



An unbounded sequence of serialized events,
where each event is represented as an
encoded key-value pair* or “Kafka message”.

Storage formats

Clients **serialize events to bytes when writing, and deserialize when reading.**



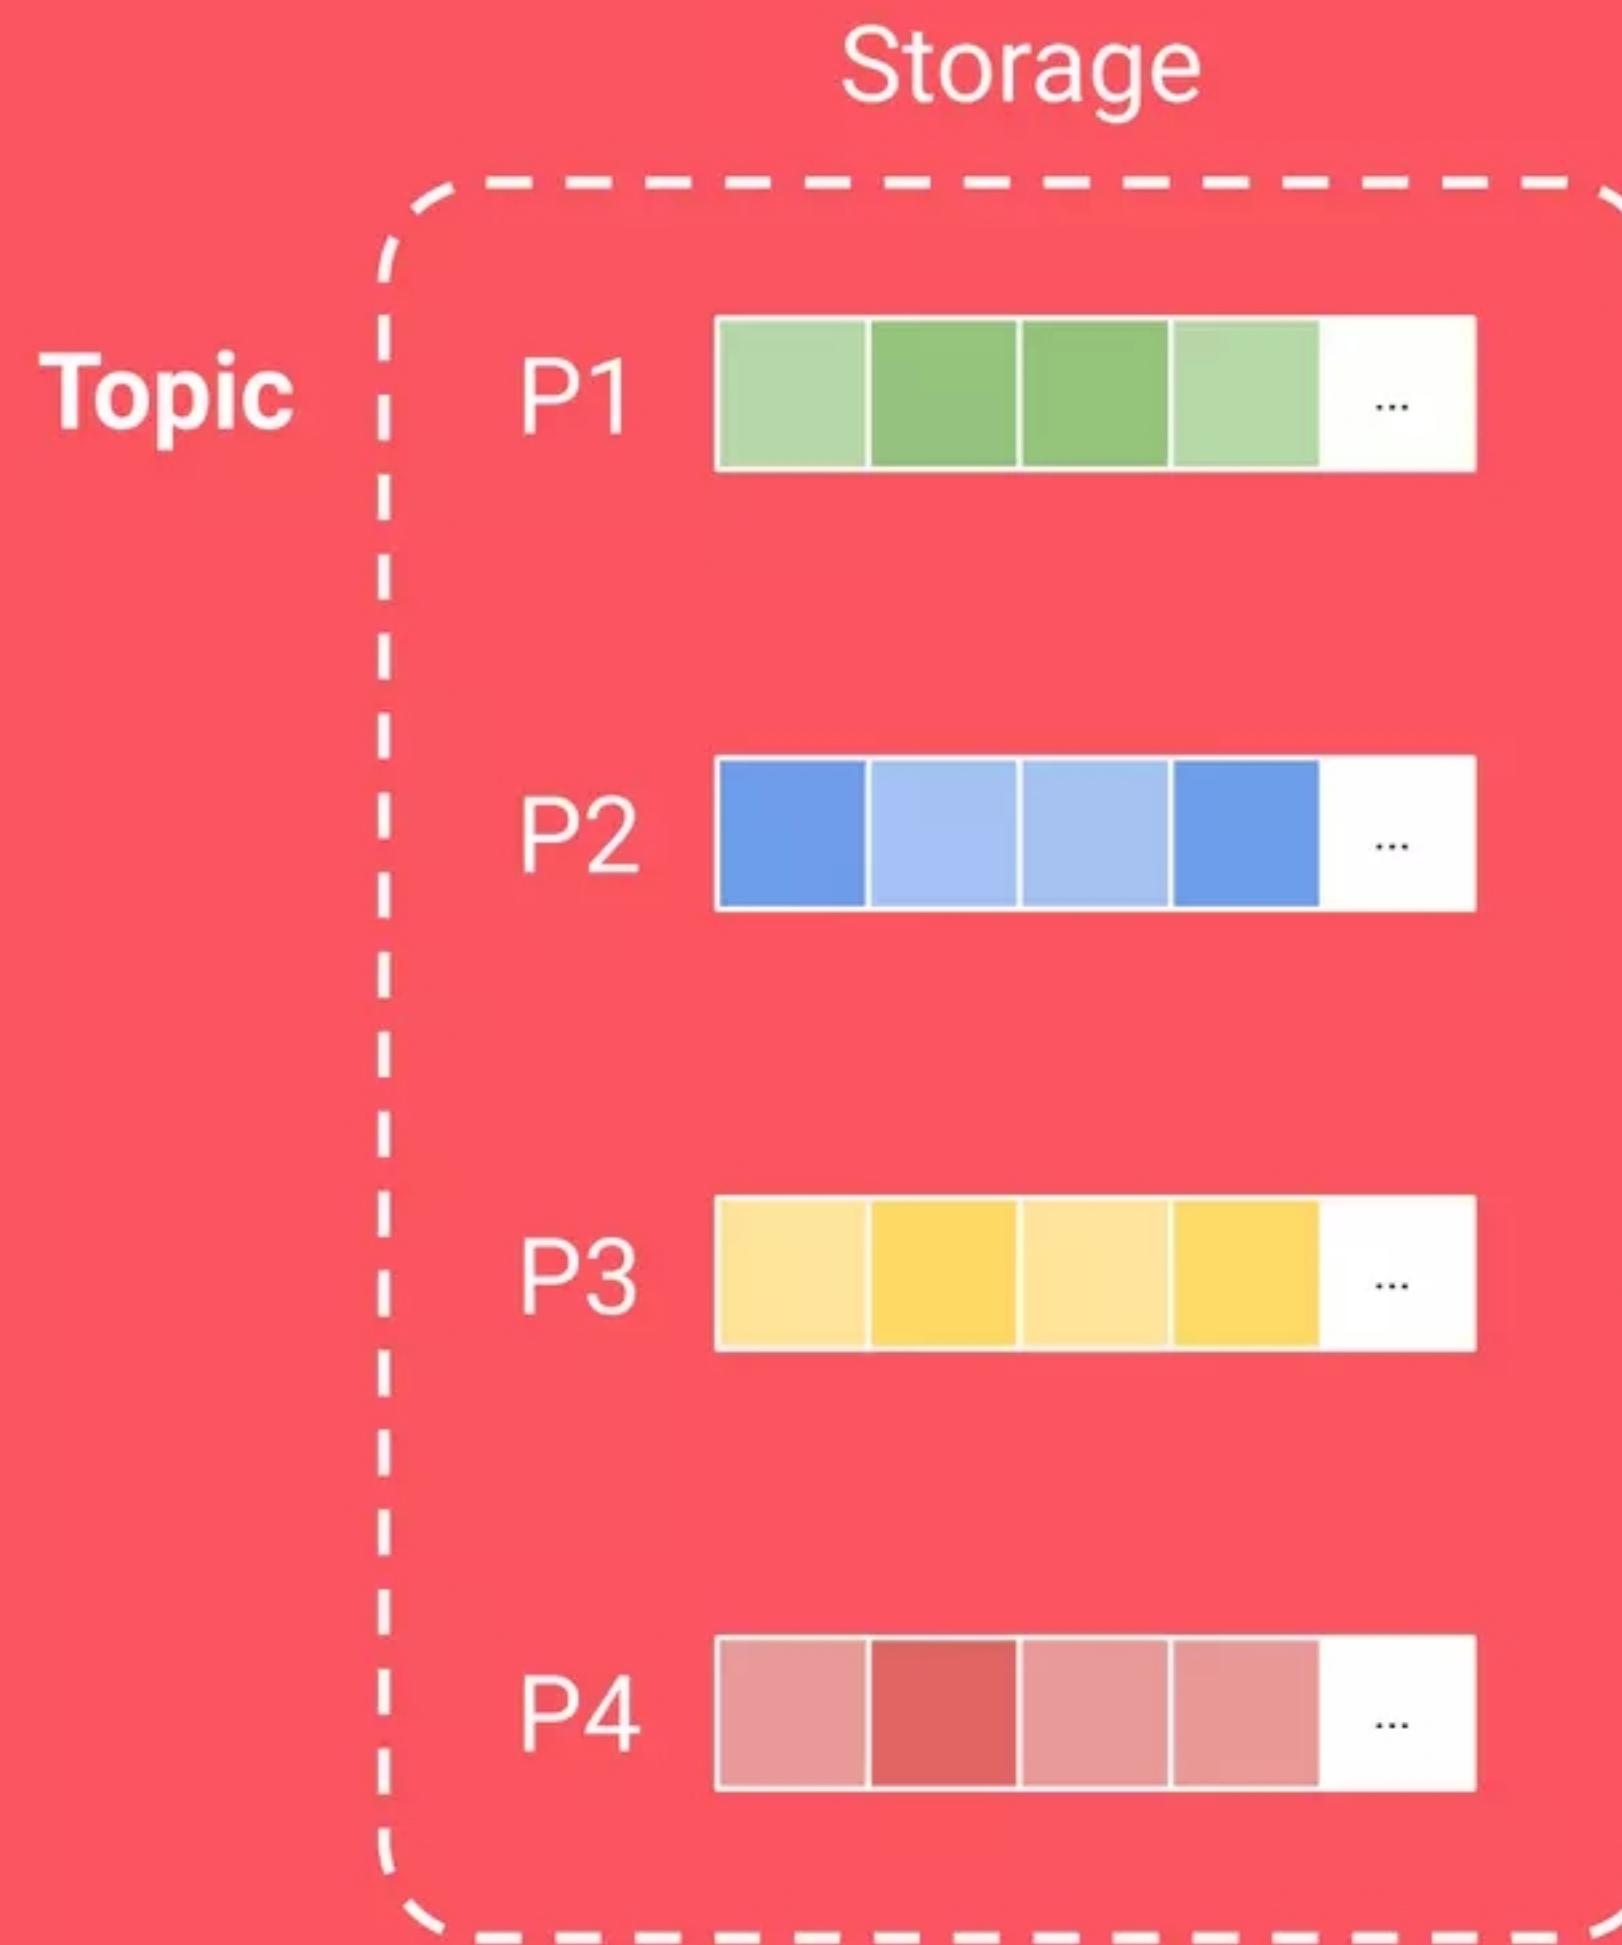
Storage formats

Brokers take no part in this. All they see is bytes.
Being this “dumb” is actually really smart!



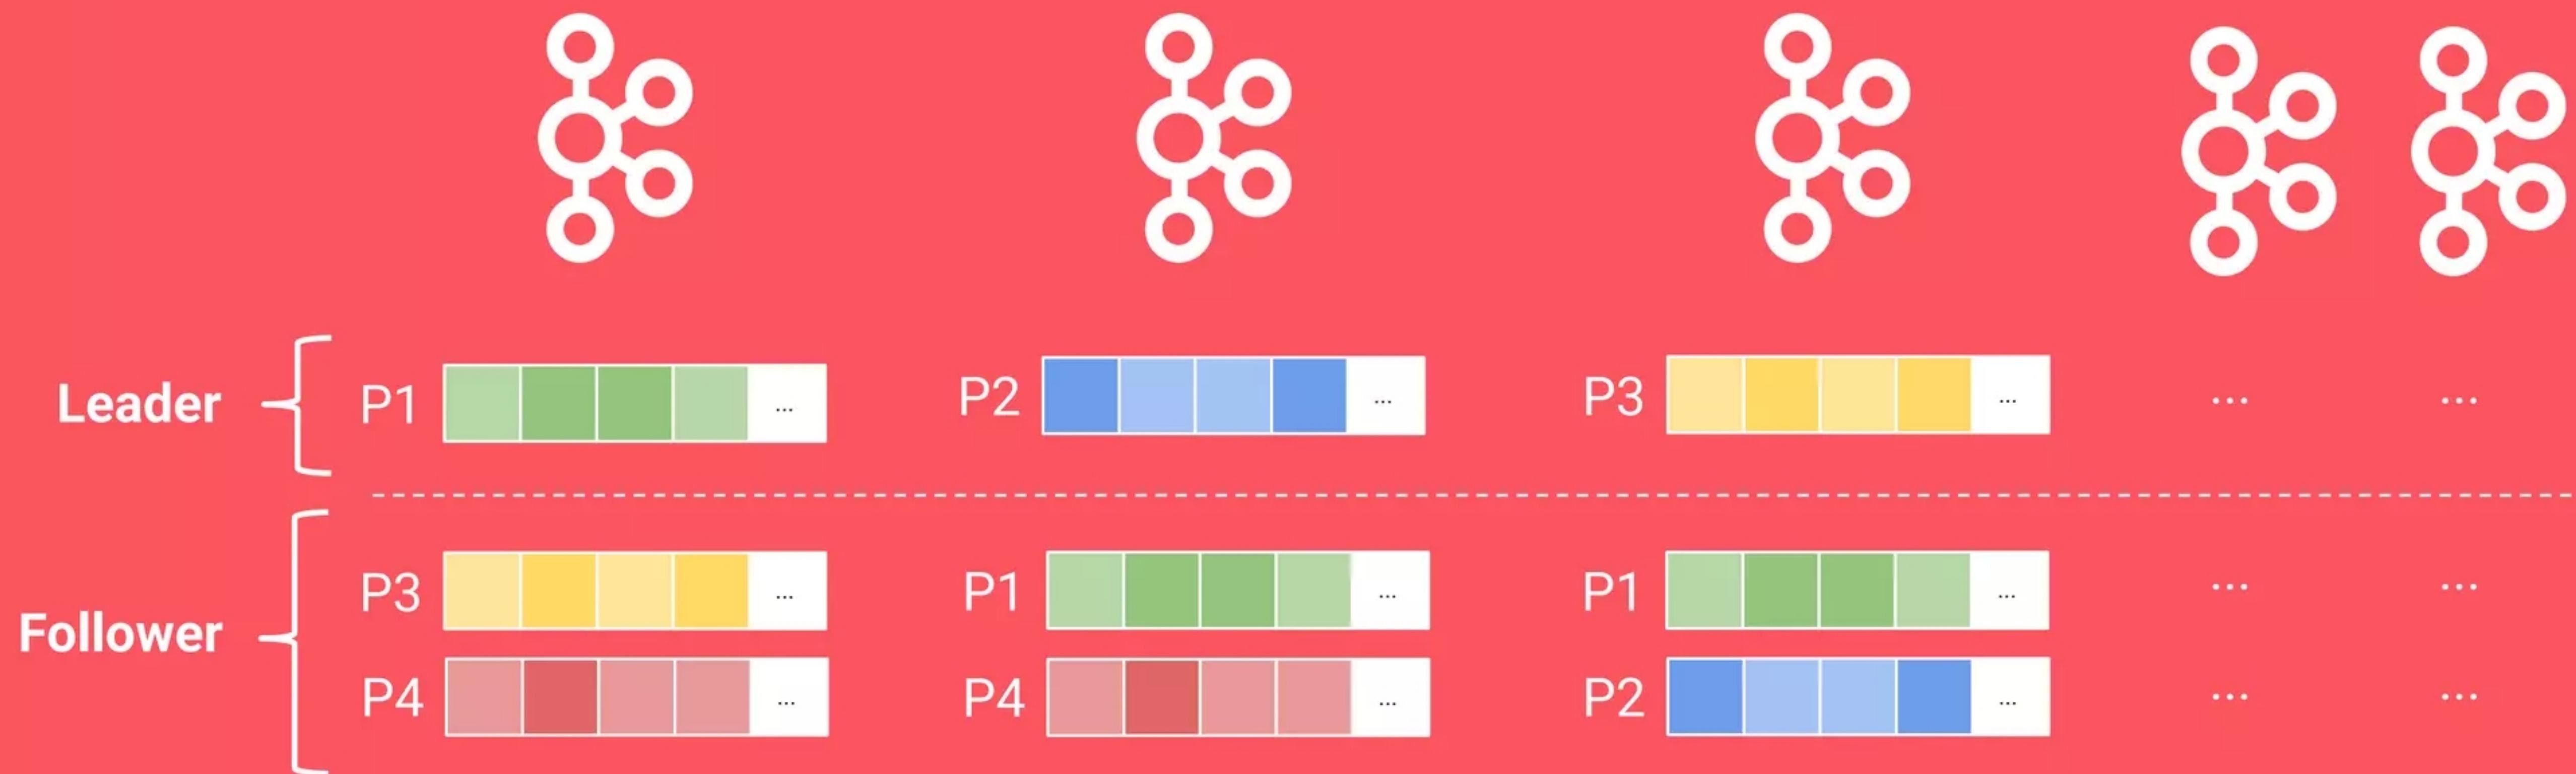
Storage is partitioned

And ‘partitions’ are the most important concept in Kafka



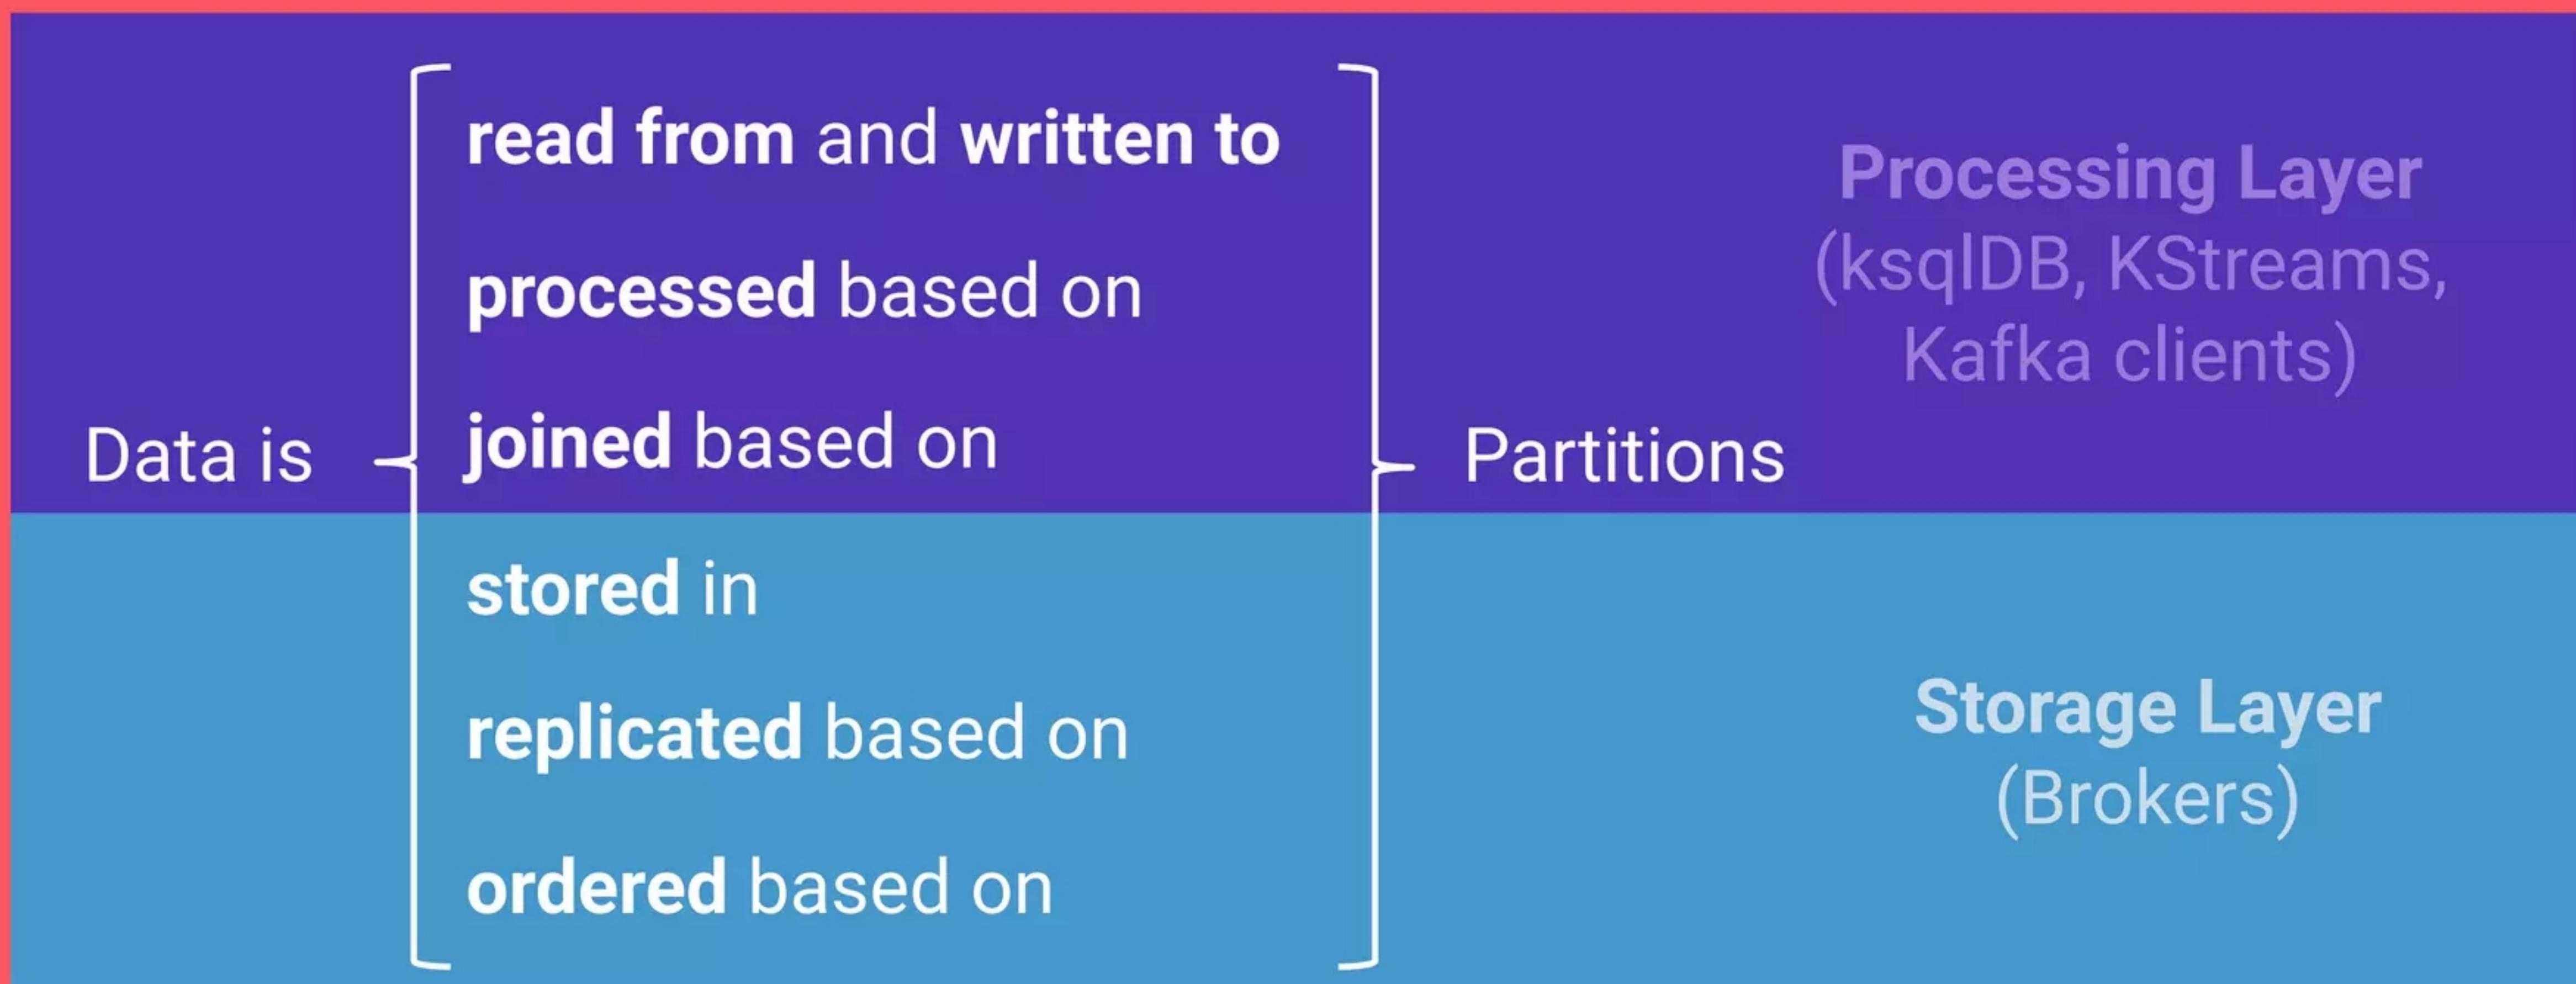
Partitions are spread across Kafka brokers

A broker can be leader and/or follower of 1+ partitions

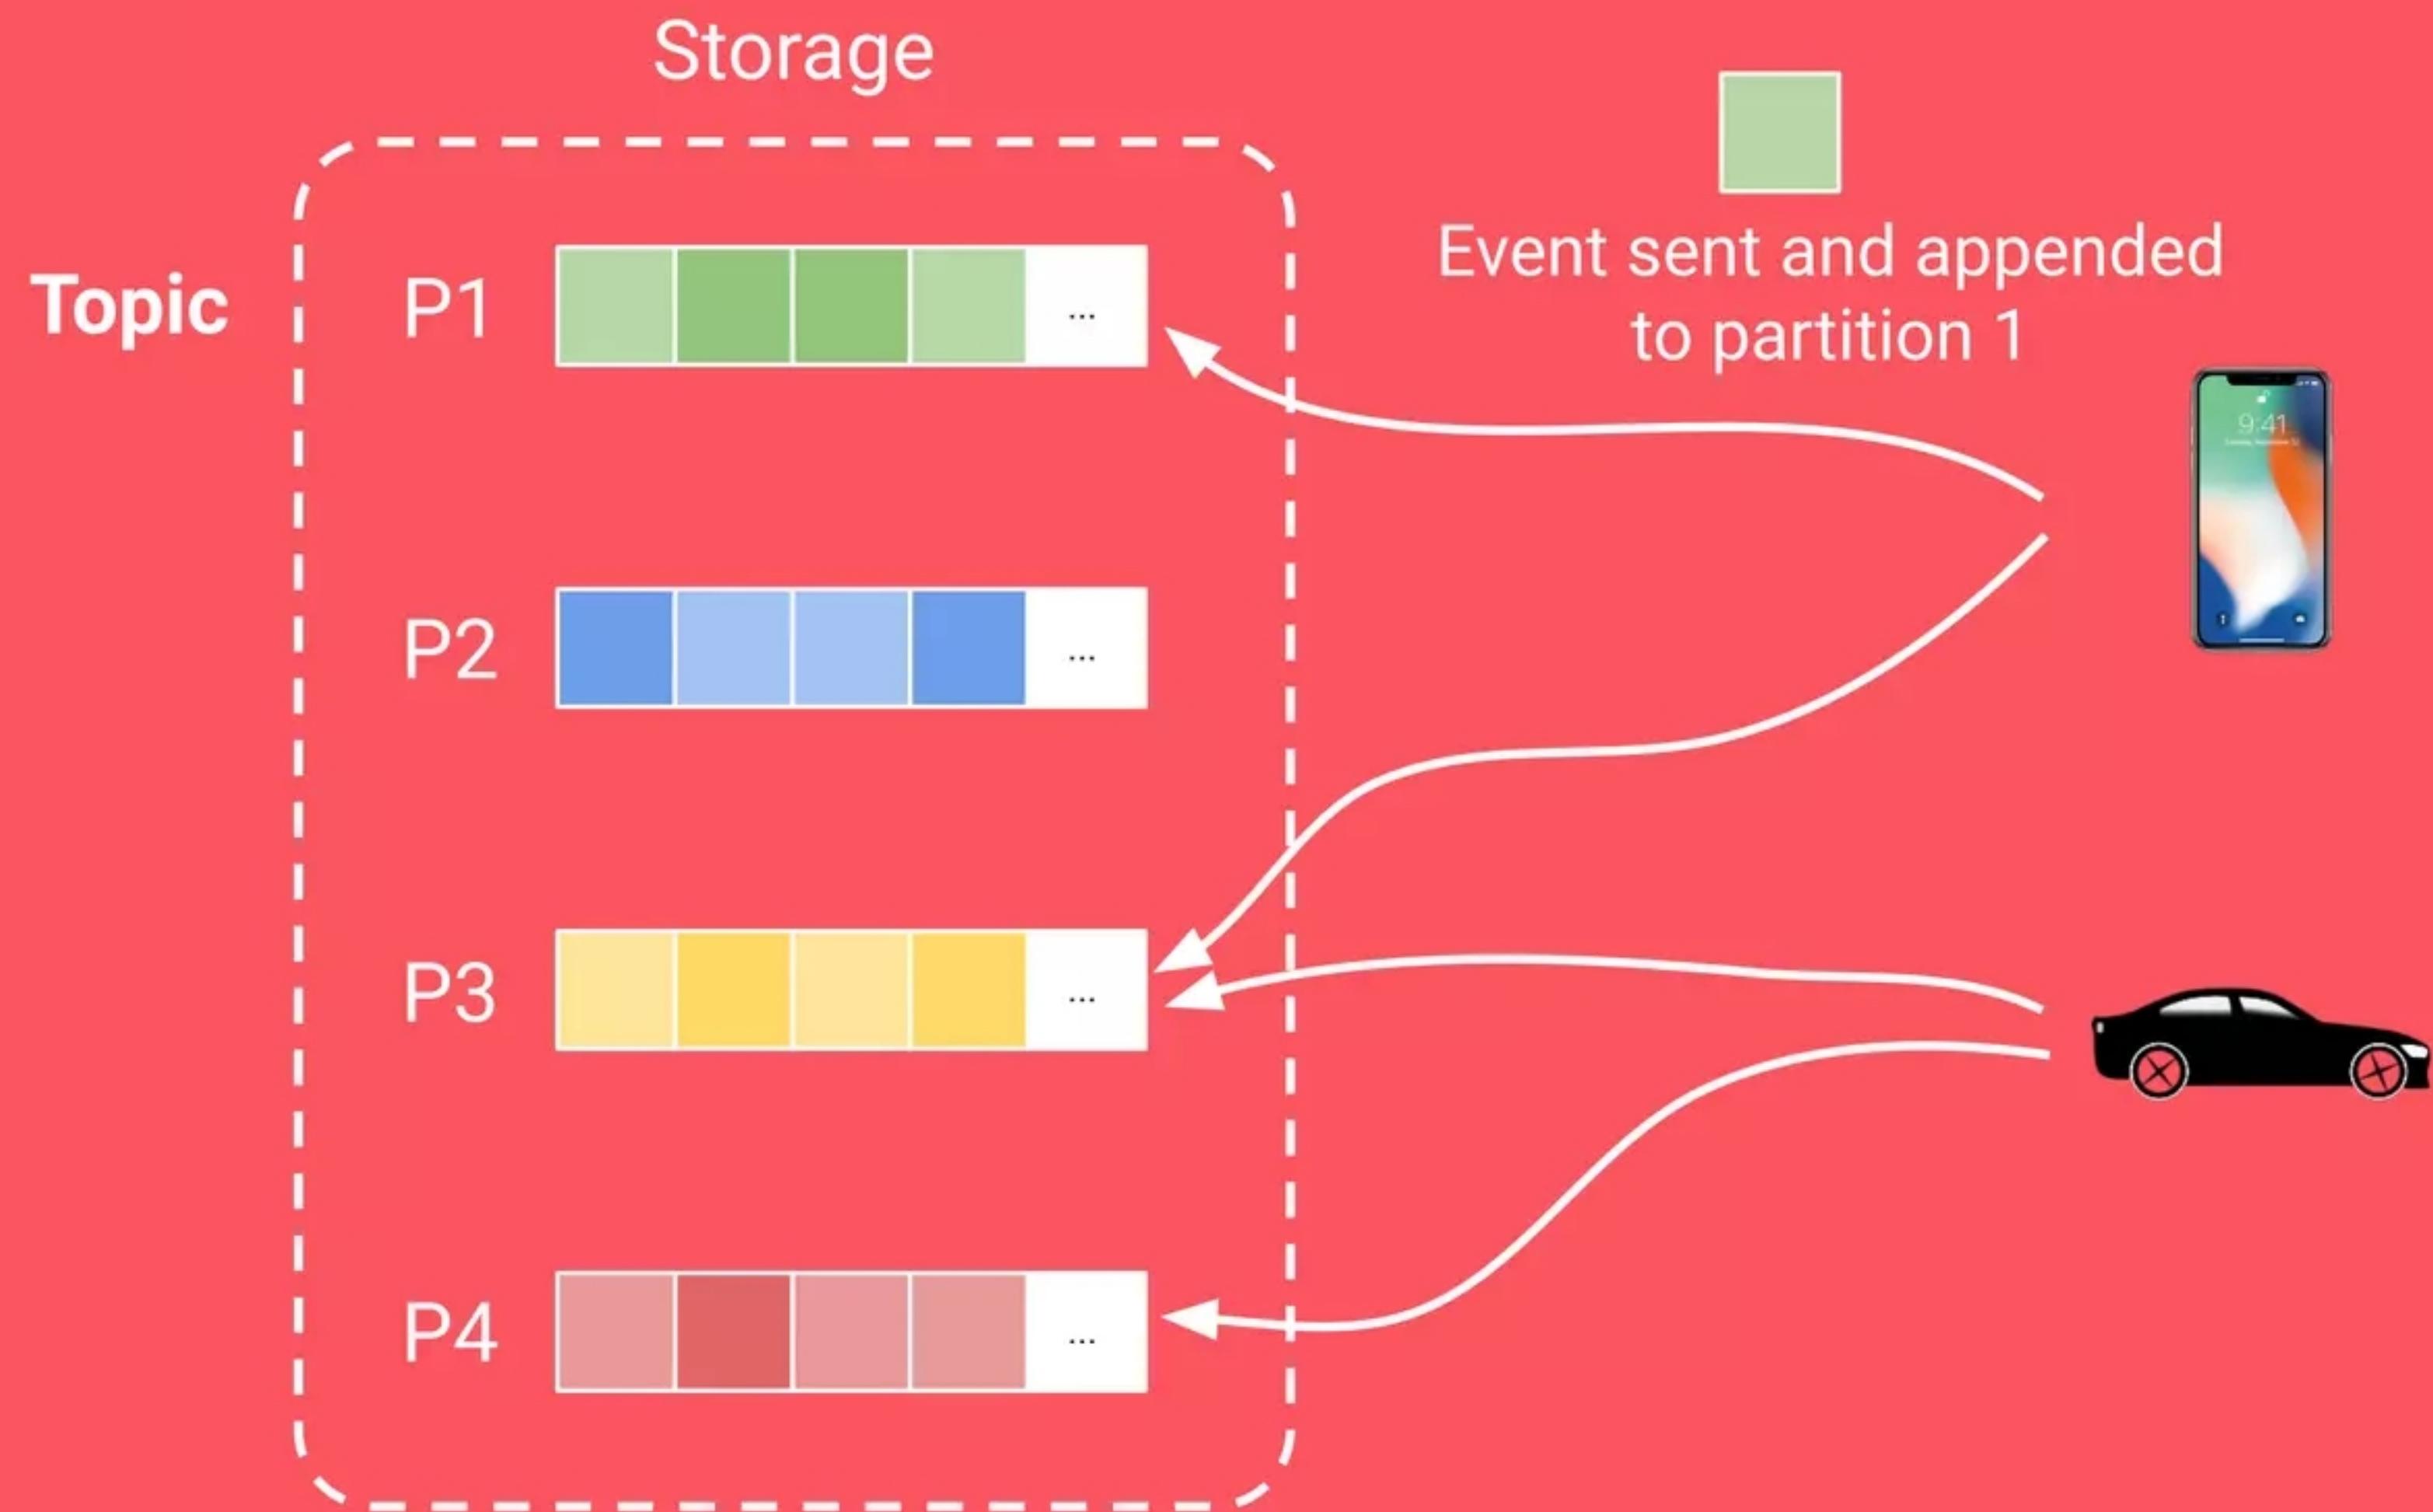


Partitions play a central role in Kafka

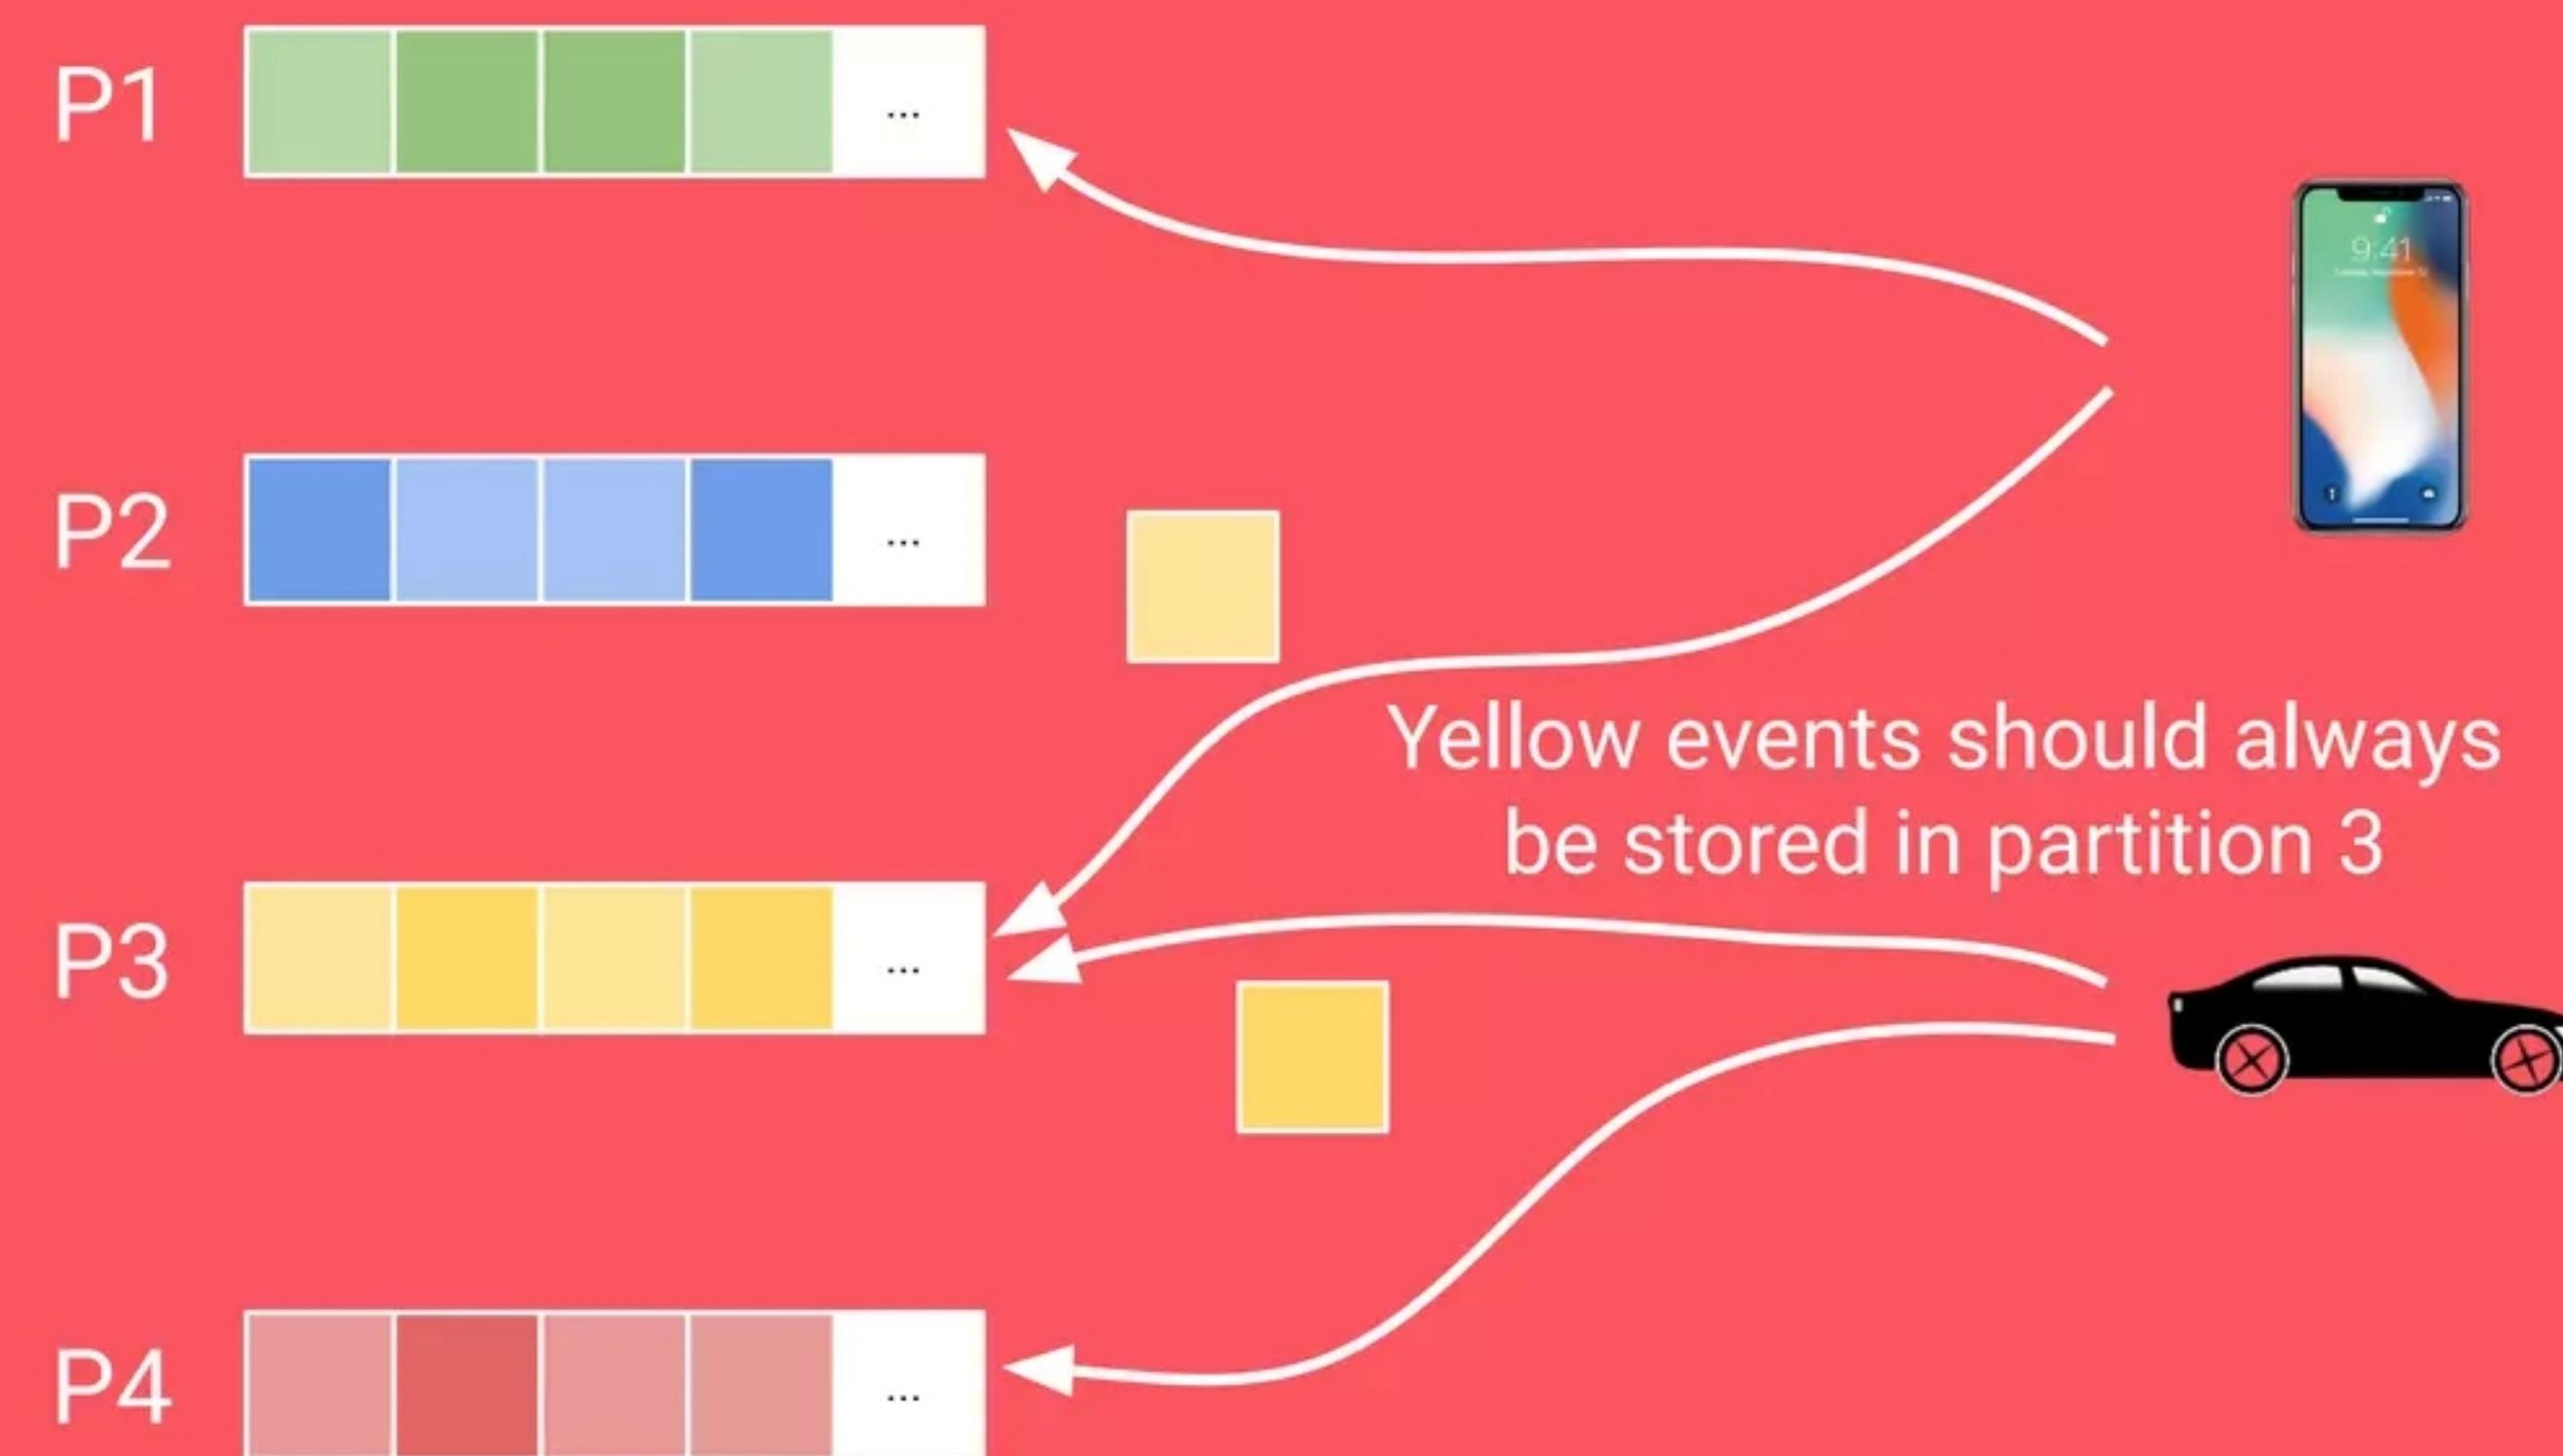
Topics are partitioned. Partitions enable **scalability, elasticity, fault tolerance**.



Event producers determine event partitioning through a partitioning function $f(\text{event.key}, \text{event.value})^*$



**Events with same key should be in same partition
to ensure proper ordering of related events
to ensure processing by Consumers returns expected results**



Top causes for same key in *different* partitions

1. Topic configuration

You increased the number of partitions in a topic

2. Producer configuration

A producer uses a custom partitioner

→ Be careful in this situation!

Processing Fundamentals with ksqlDB and Kafka Streams

From Topics to Streams and Tables

Topics

live in the Kafka *storage layer*, ‘filesystem’ (Brokers)

Streams and Tables

live in the Kafka *processing layer* (ksqlDB, KStreams)

Topics vs. Streams and Tables

Processing Layer
(ksqldb, KStreams)

Table

plus aggregation

Event Stream



plus schema (serdes)

Storage Layer
(Brokers)

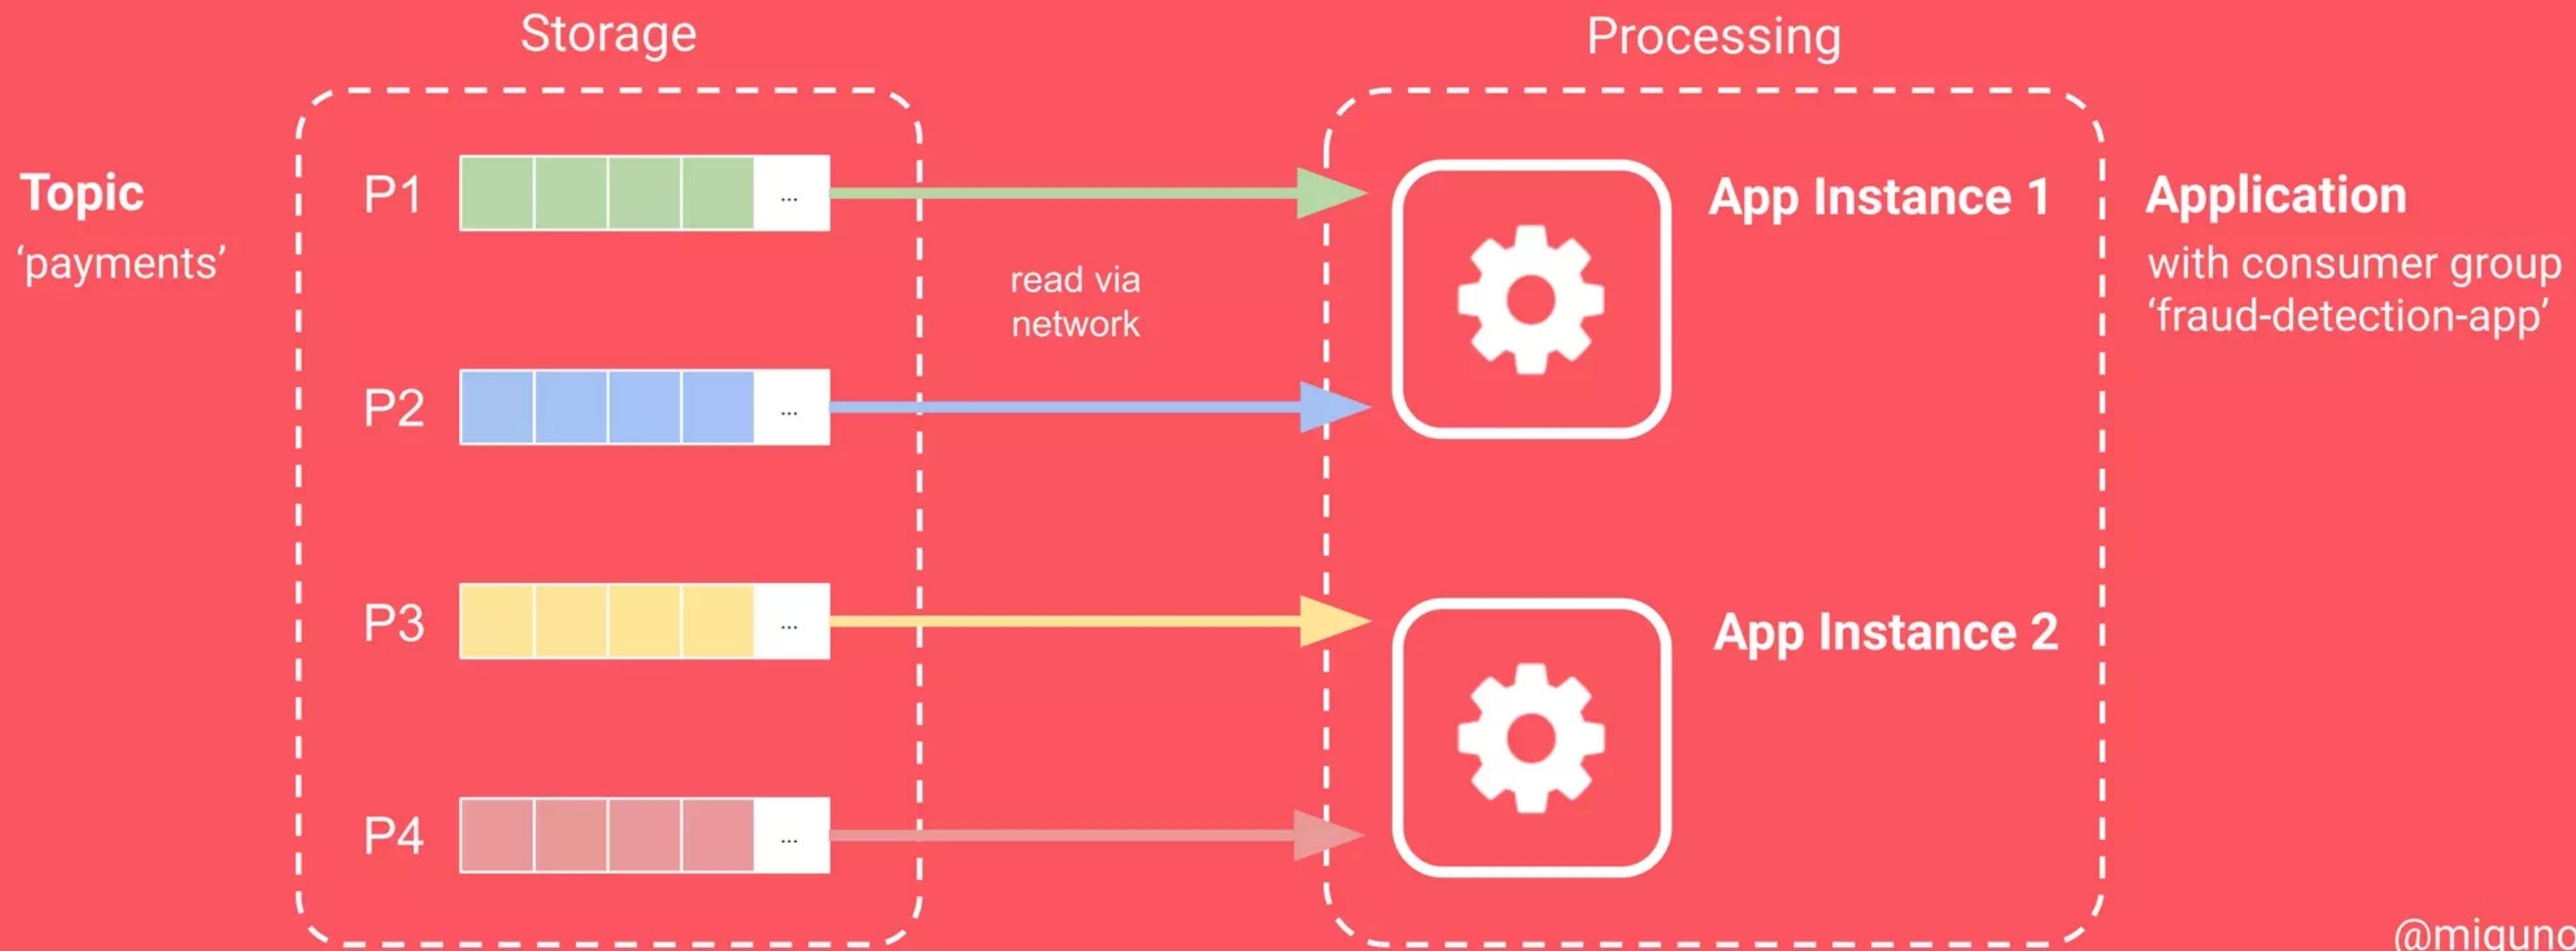
Topic



alice	2
bob	1

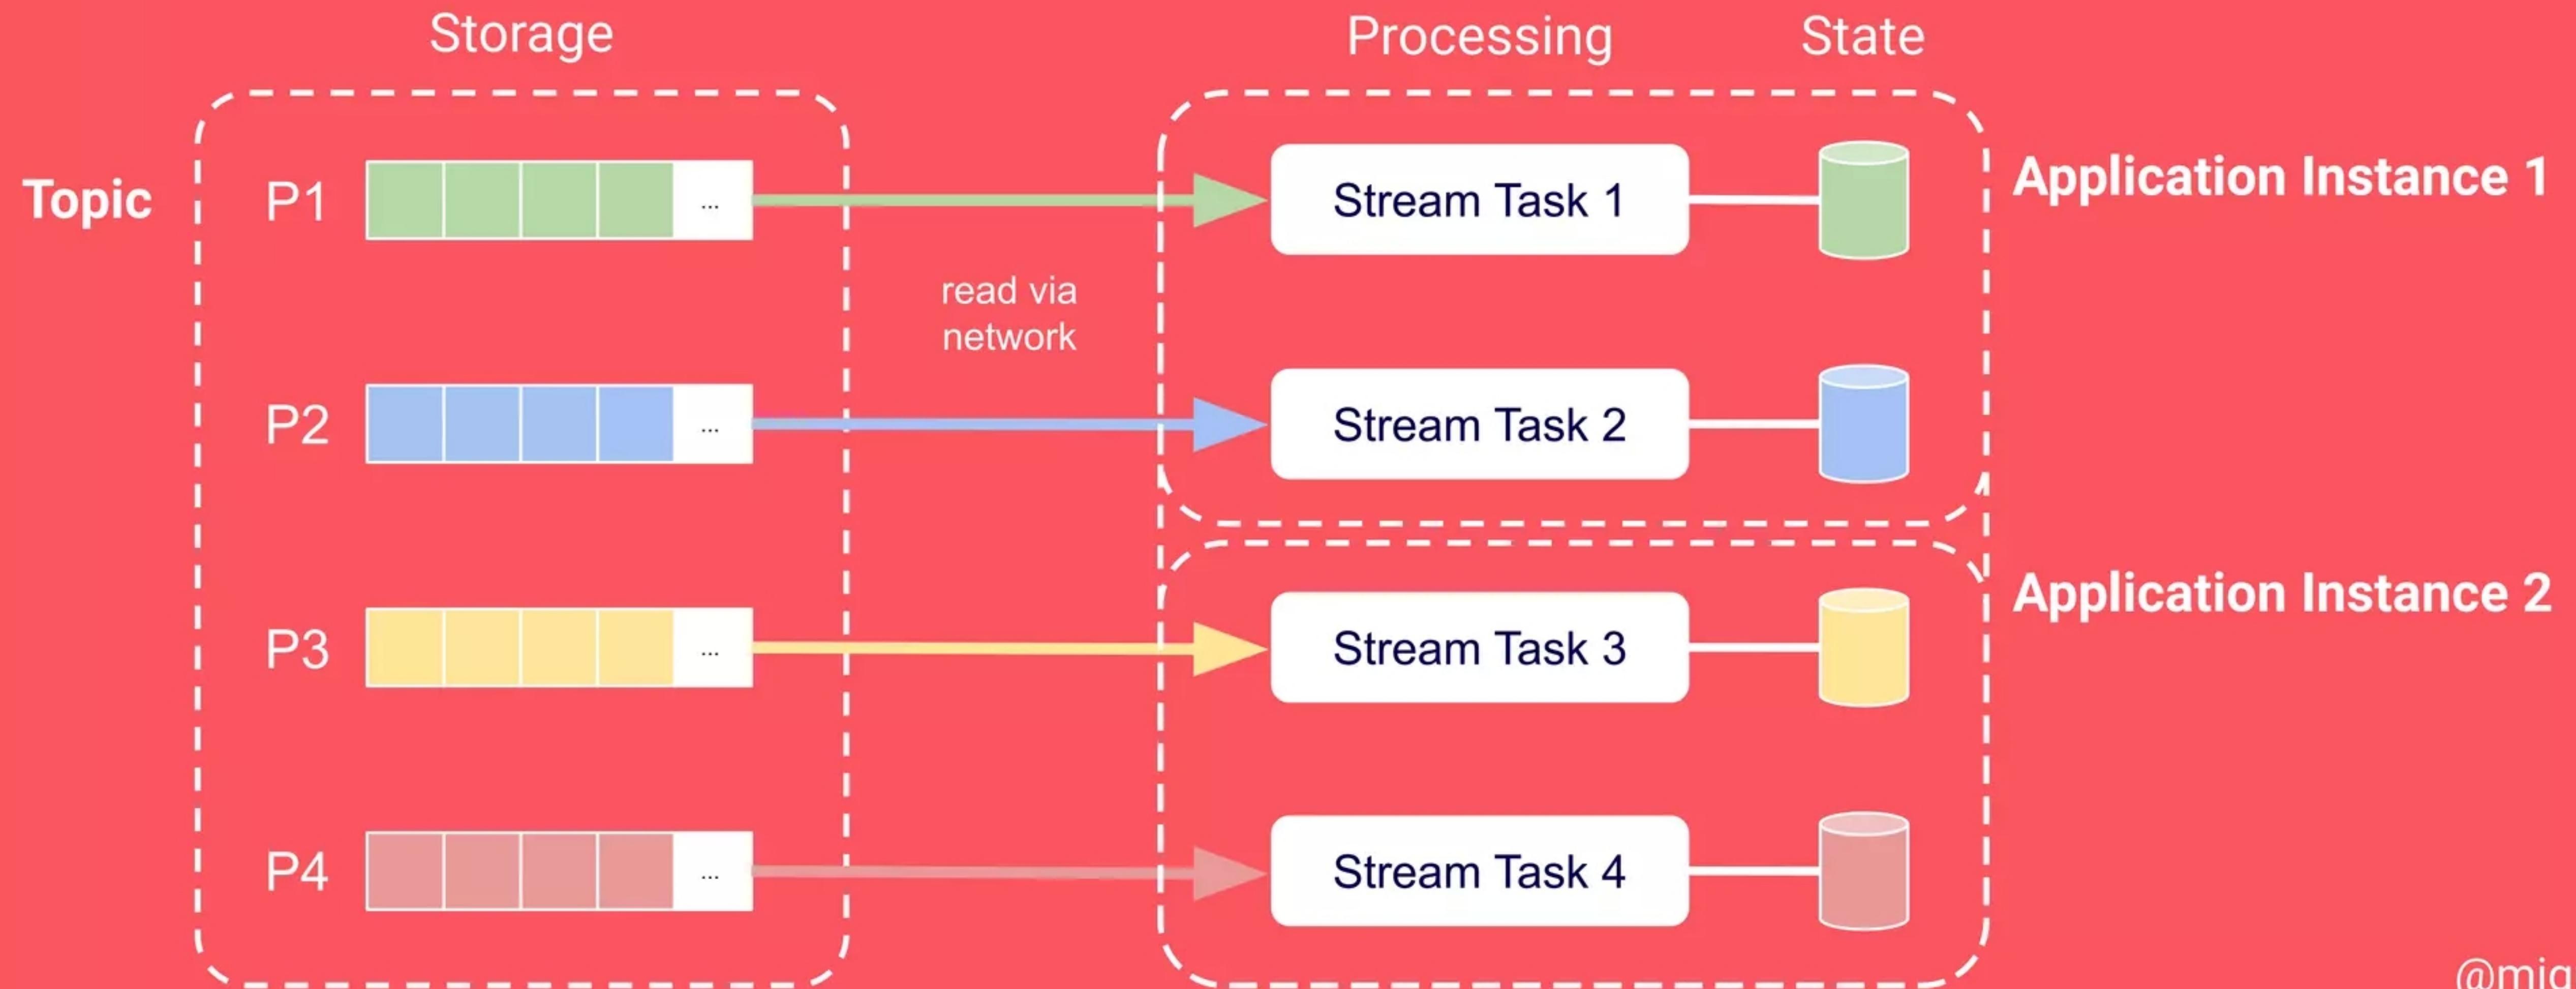
Kafka Processing

Data is processed per-partition

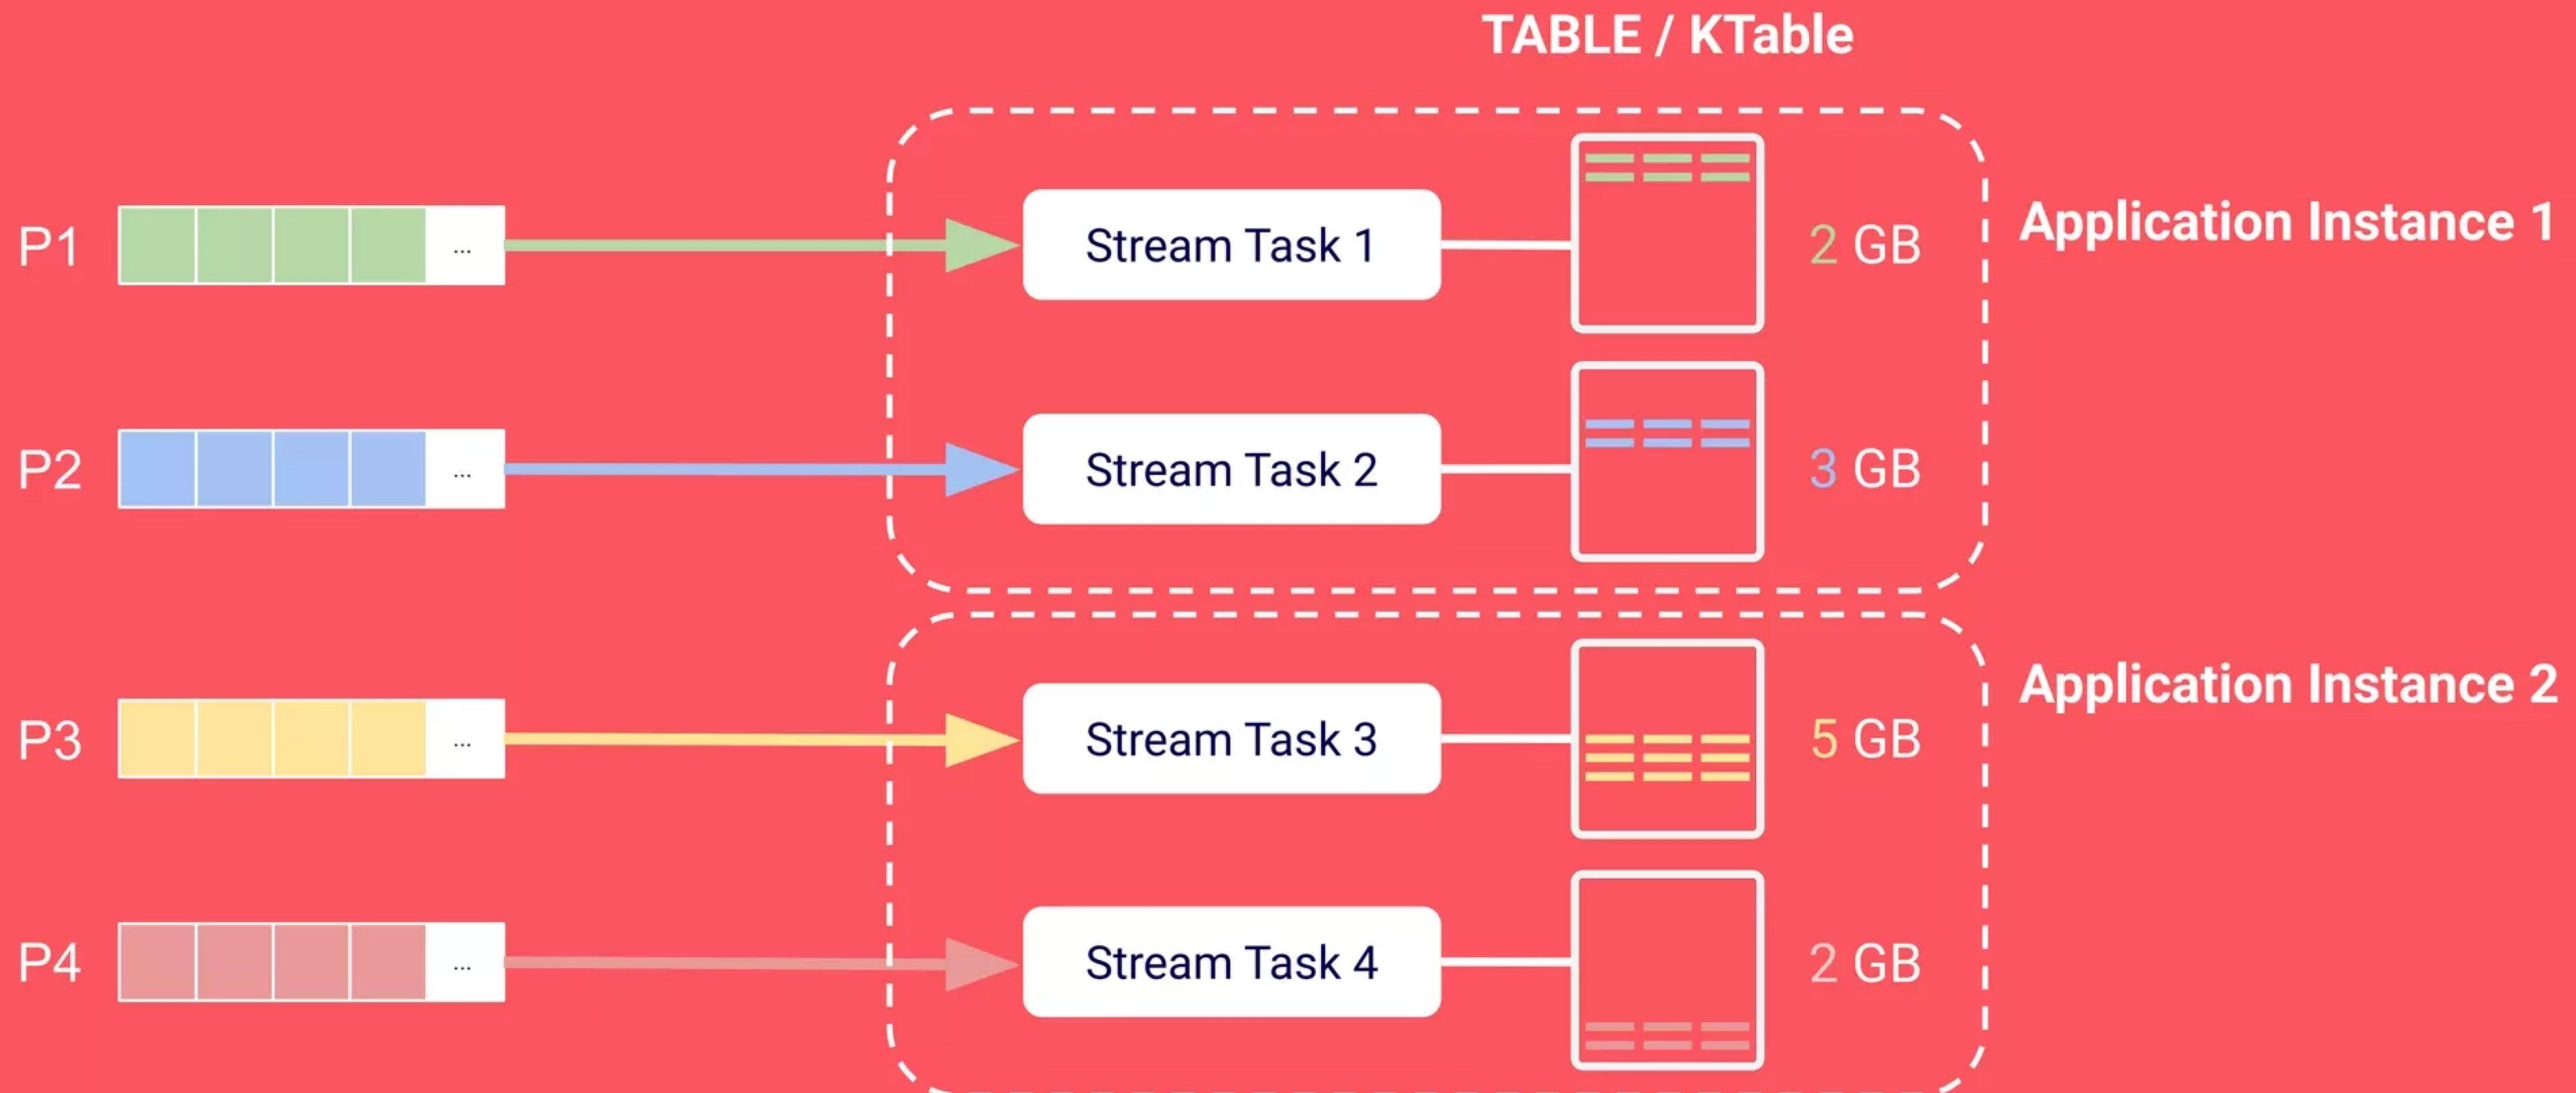


Kafka Processing

Stream tasks are the unit of parallelism

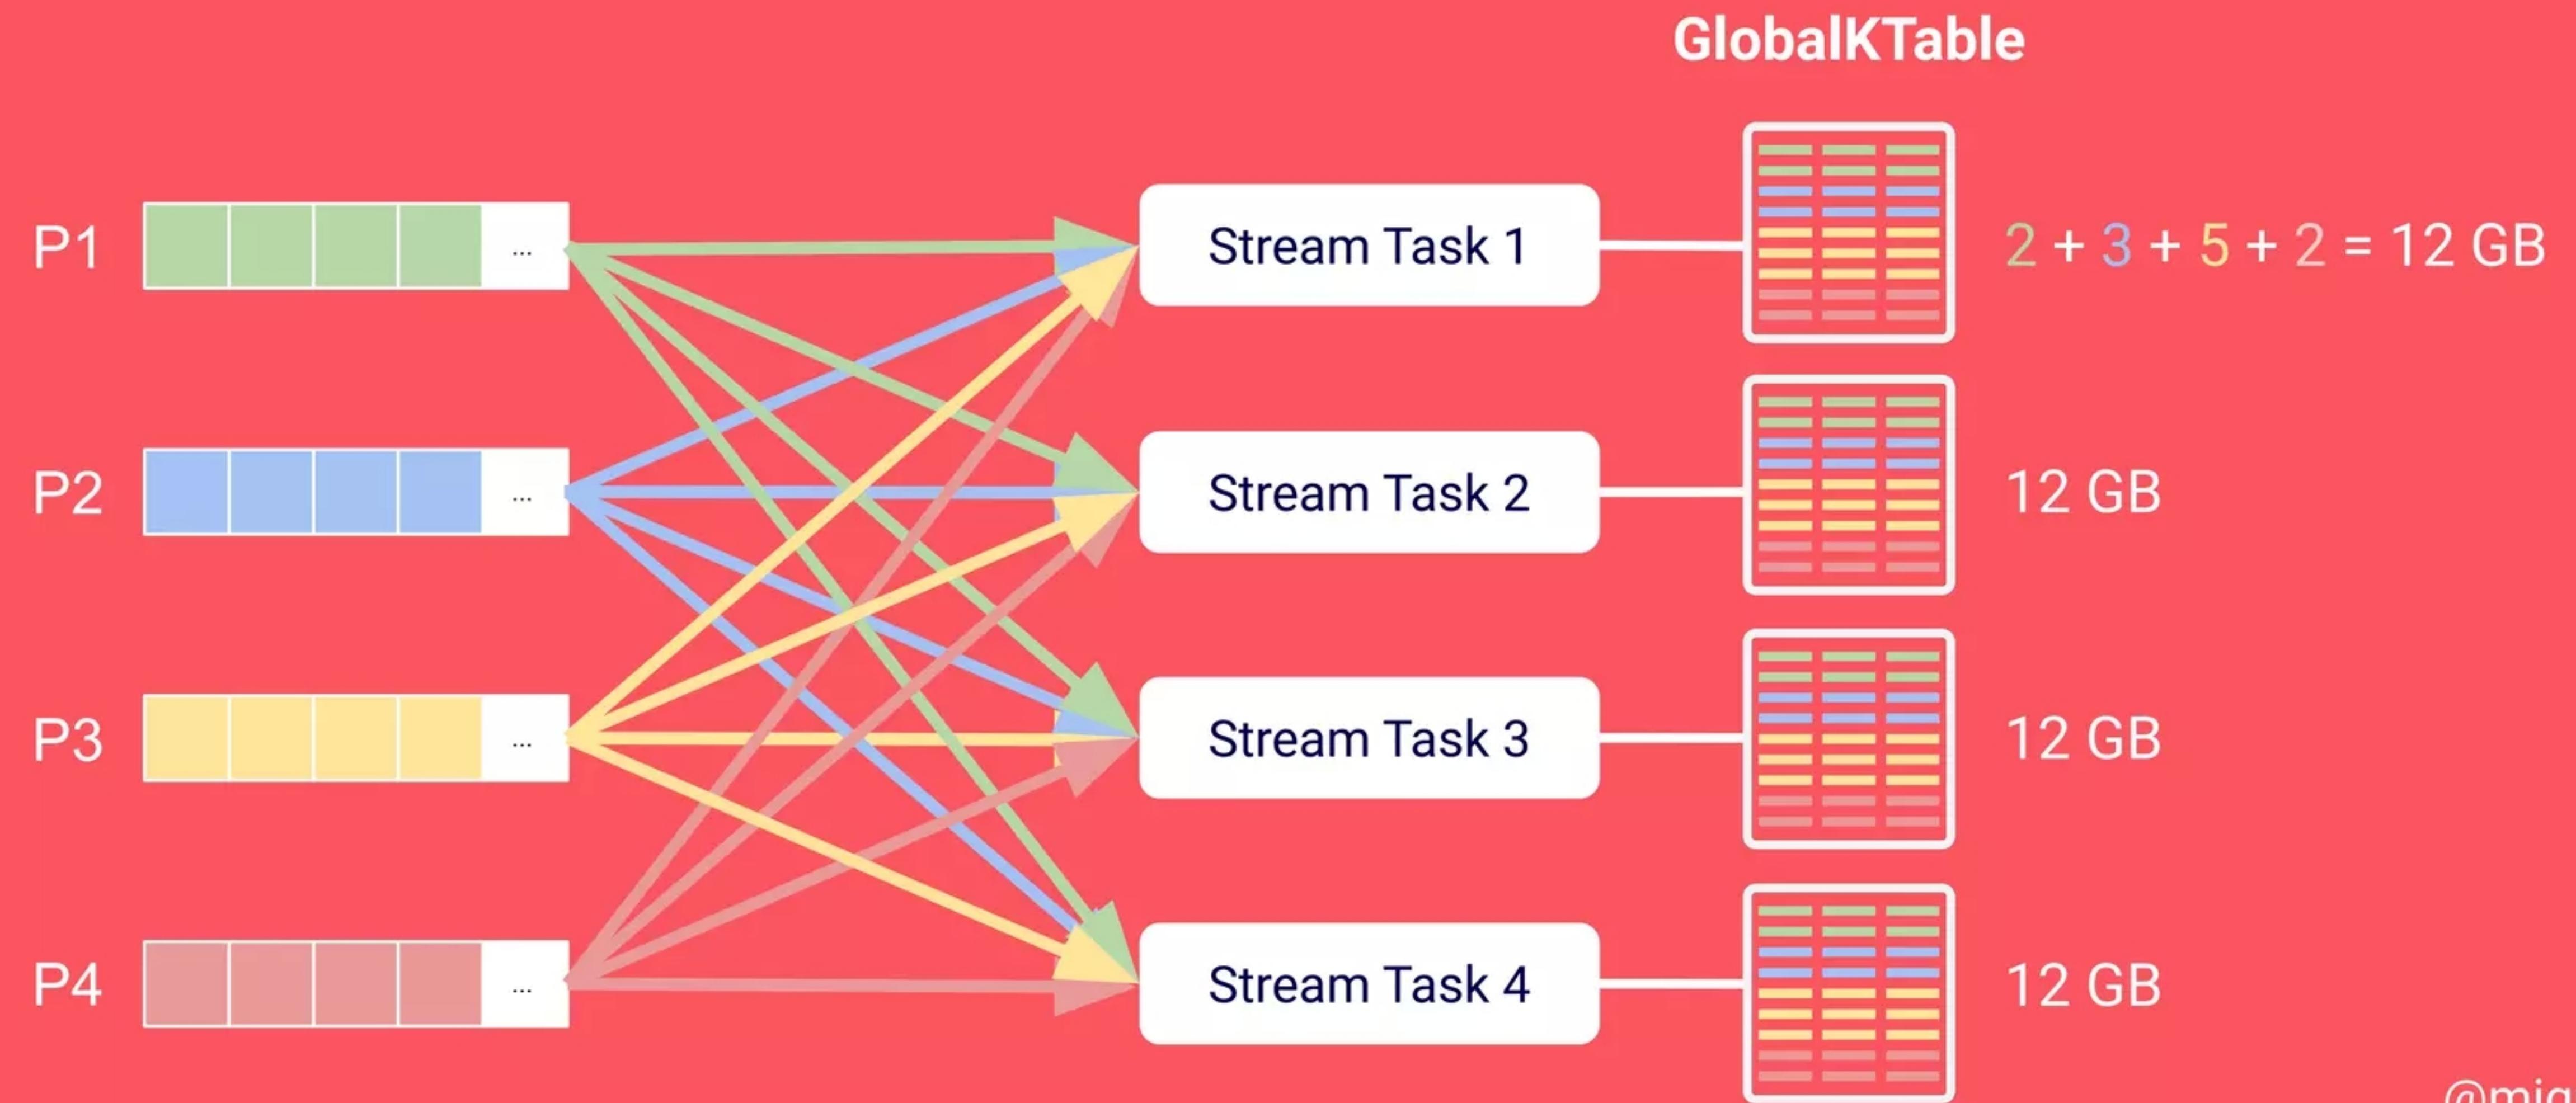


Streams and Tables are partitioned, too



Global Tables give *complete* data to every task

Great for joins without re-keying the input or to broadcast info



Topics vs. Streams and Tables

Concept	Schema	Partitioned	Unbounded	Ordering Guarantee***	Mutable	Single Value Per Key	Fault Tolerant
Storage Layer							
Topic	No* (raw bytes)	Yes	Yes	Yes	No	No	Yes
Processing Layer							
Stream	Yes	Yes	Yes	Yes	No	No	Yes
Table	Yes	Yes	No**	No	Yes	Yes	Yes
Global Table	Yes	No	No**	No	Yes	Yes	Yes

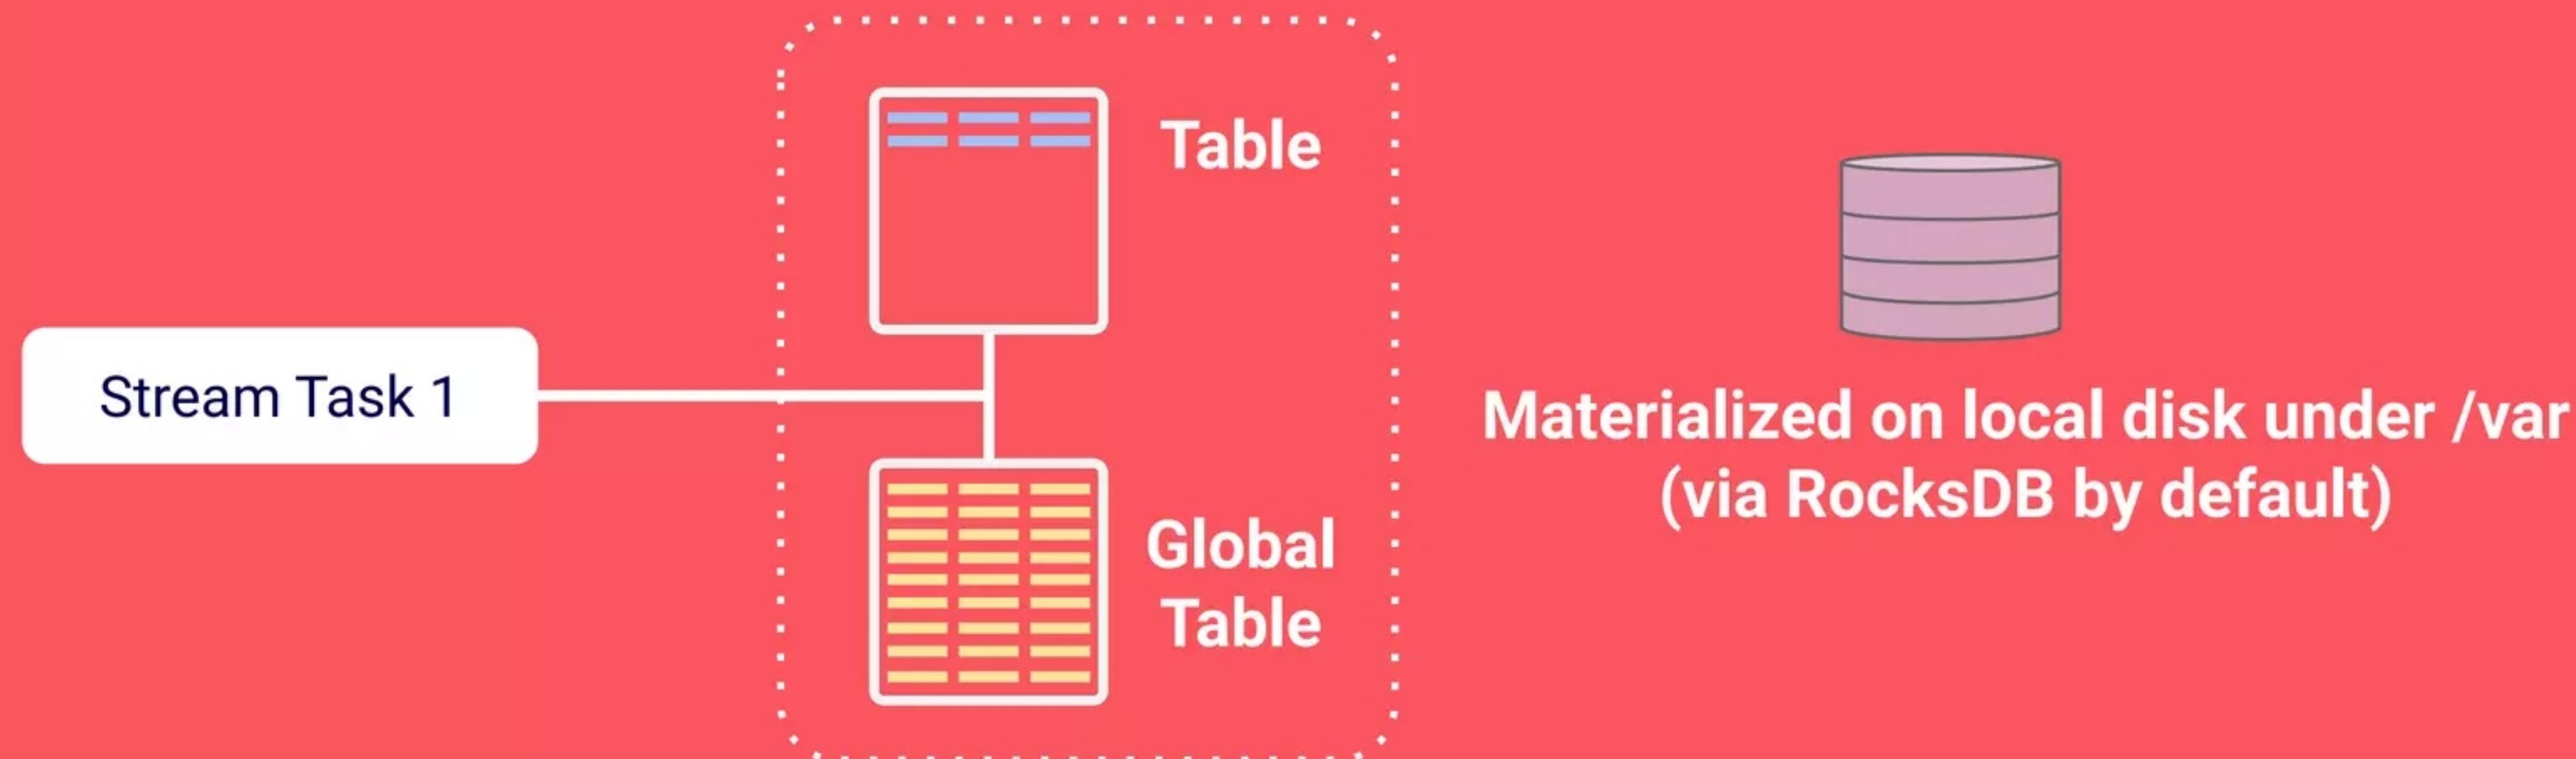
*When using Schema Registry with Avro, then a schema is technically available for the data in a topic.

**Generally speaking the answer is Yes but, in practice, tables are almost always bounded due to finite key space.

***Ordering per partition, not across the multiple partitions in a topic.

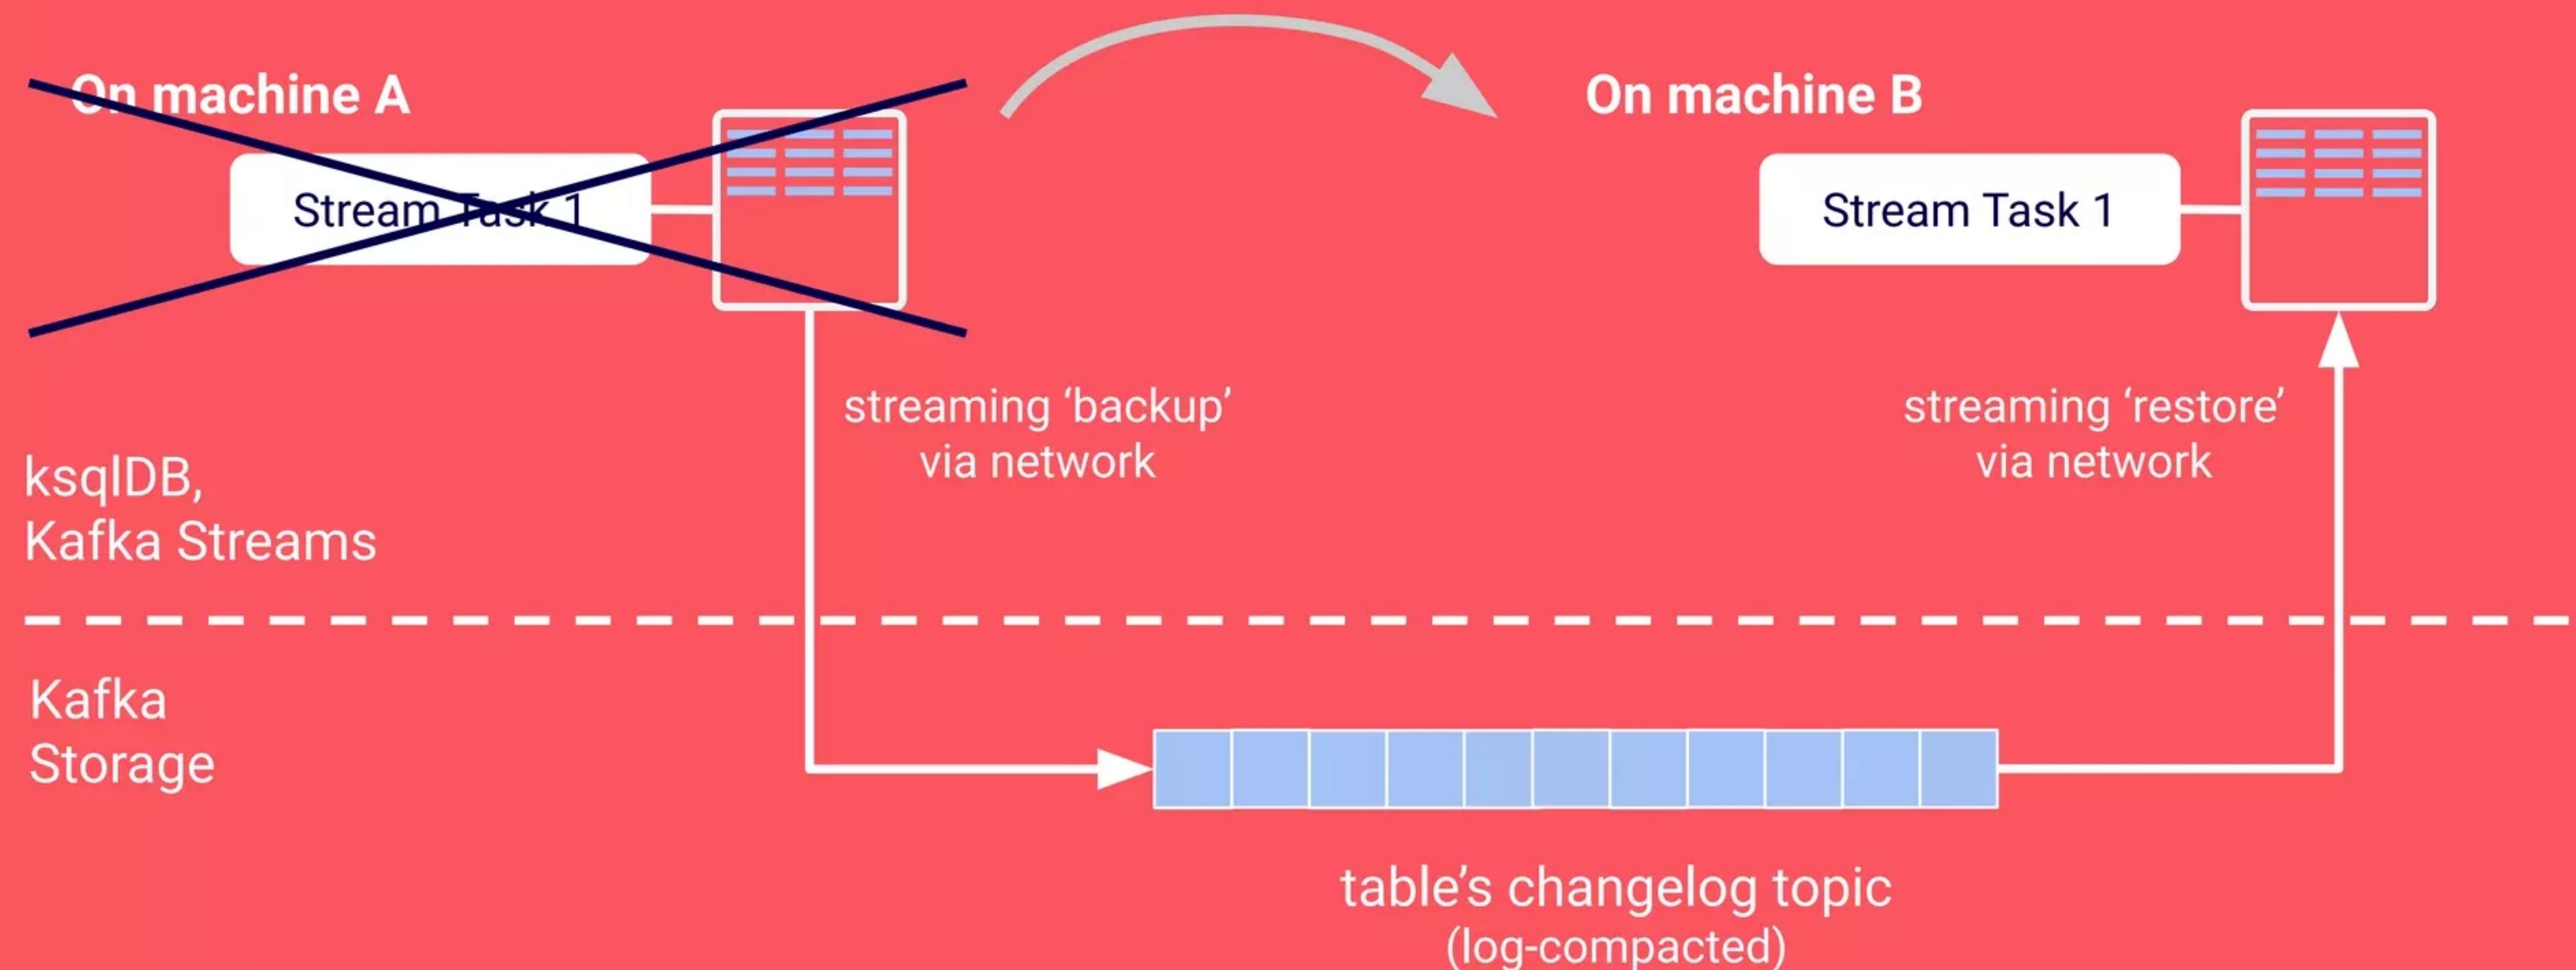
Elasticity, Fault Tolerance, and Other Advanced Concepts

Tables are materialized on local disk so that processing is fast, and tables can be larger than RAM. Enables large-scale state. Saves \$\$\$ on cloud instances.



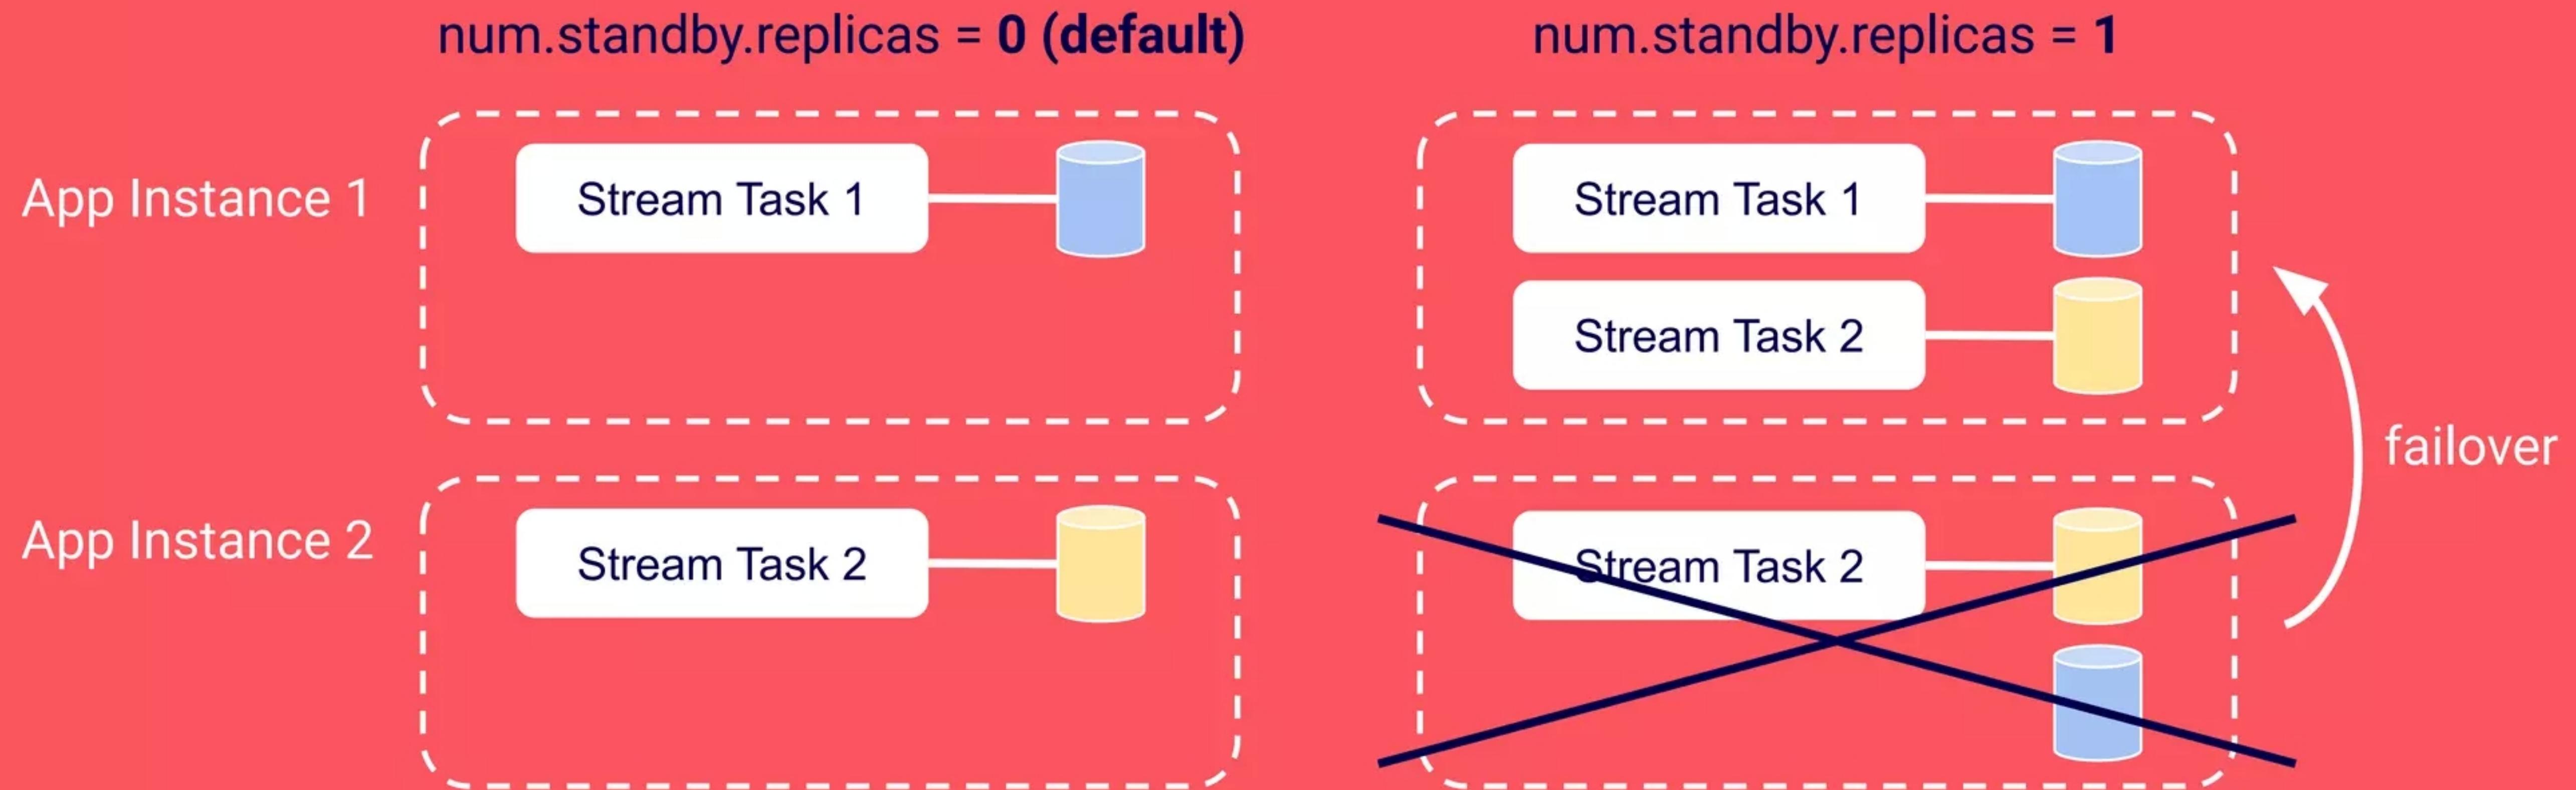
**But machines and containers can be lost!
How can we make this fault-tolerant?**

Tables and other state are made fault-tolerant by storing their data durably, remotely in Kafka topics



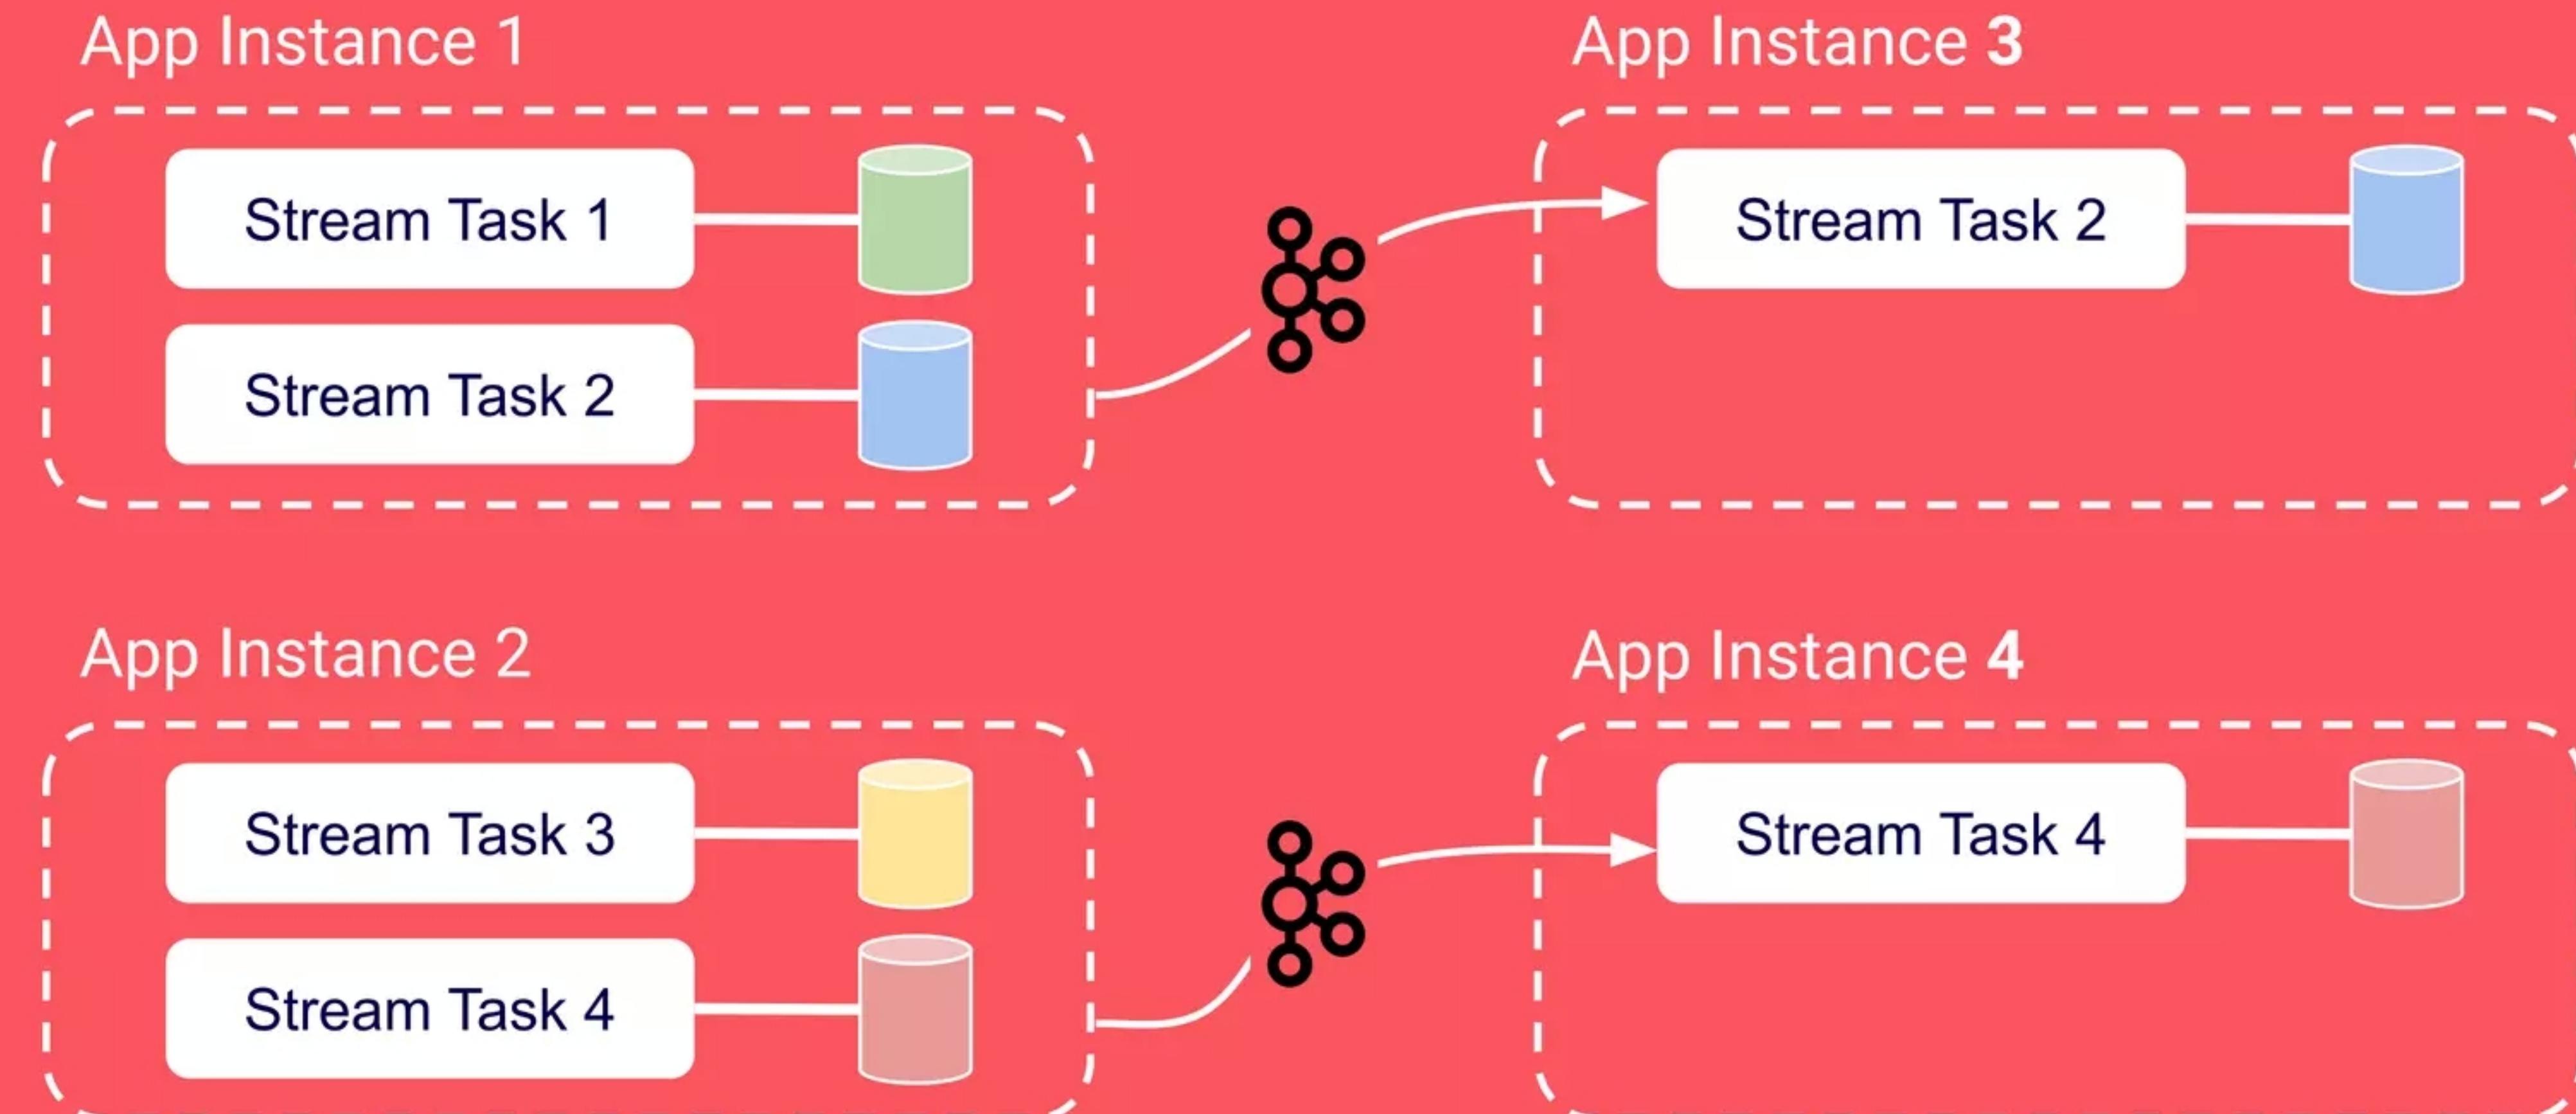
Standby Replicas speed up application recovery

App instances can optionally maintain copies of another instance's table/state data to minimize failover times



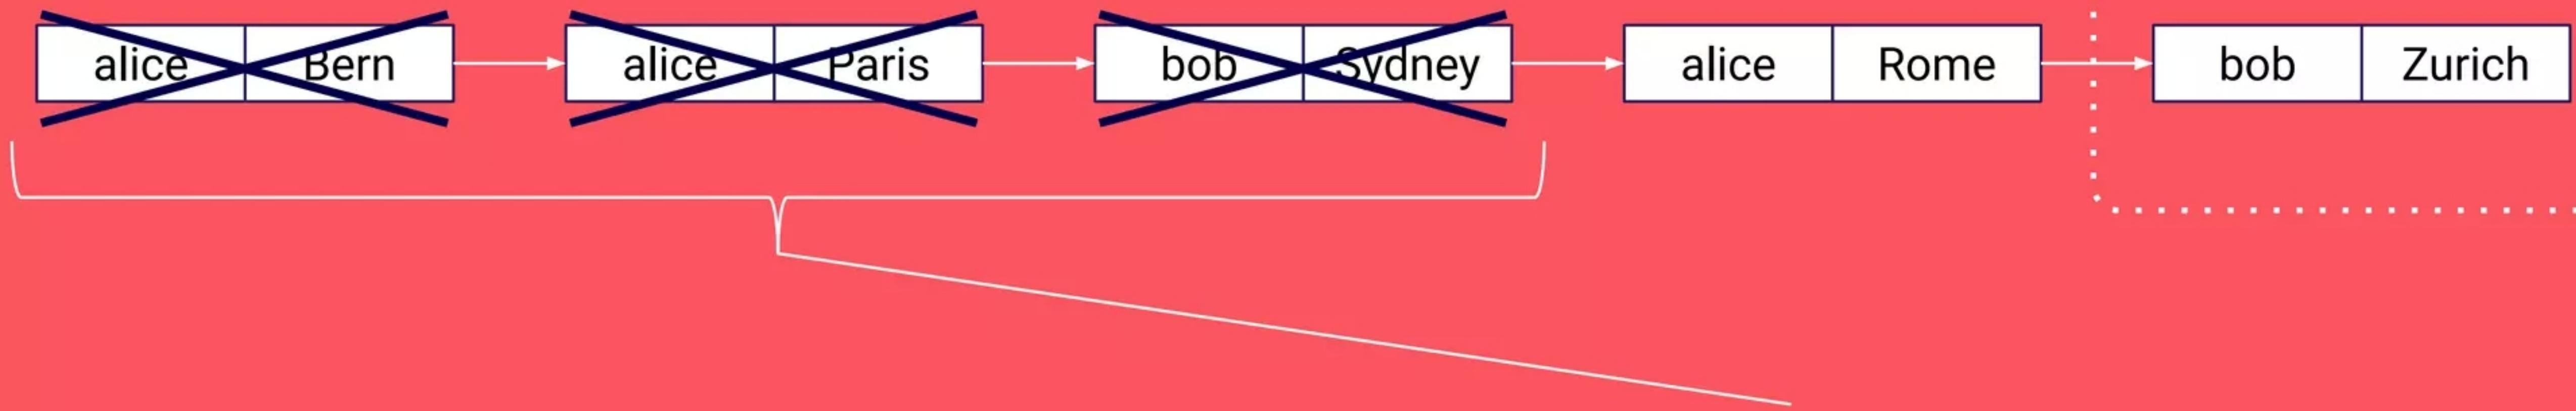
Elastic scaling

Stream tasks are migrated, including tables/state (via Kafka)



Tables <-> topic log-compaction

A table's underlying topic is **compacted by default** to save Kafka storage space
to speed up failover and recovery for processing



Note: Compaction intentionally removes part of a table's history. If you need the full history and don't have the historic data elsewhere, consider disabling compaction.

TL;DR for Log Compaction

Have a Stream?

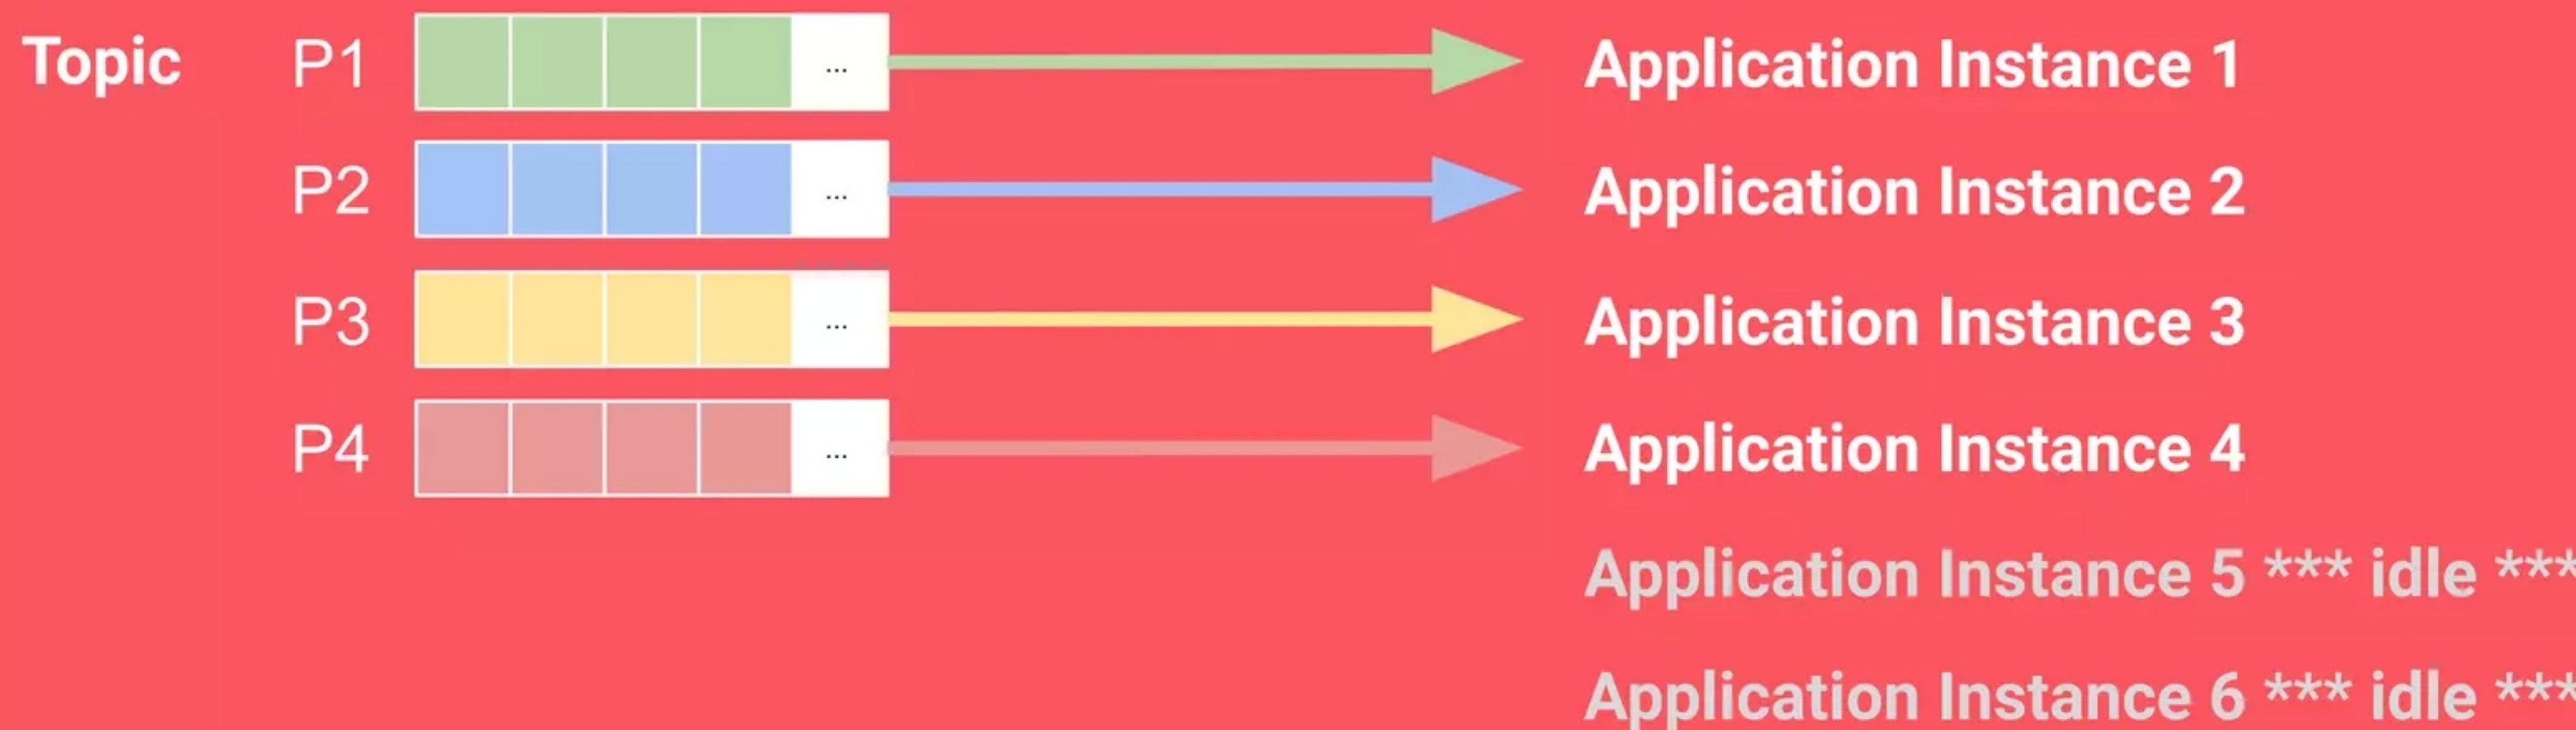
→ **Disable log compaction for its topic (= default)**

Have a Table?

→ **Enable log compaction for its topic (= default)**

Disable only when needed, see previous slide

Max processing parallelism = #input partitions



- Need higher parallelism? Increase the original topic's partition count.
- Higher parallelism for just one use case? Derive a new topic from the original with higher partition count. Lower its retention to save storage.

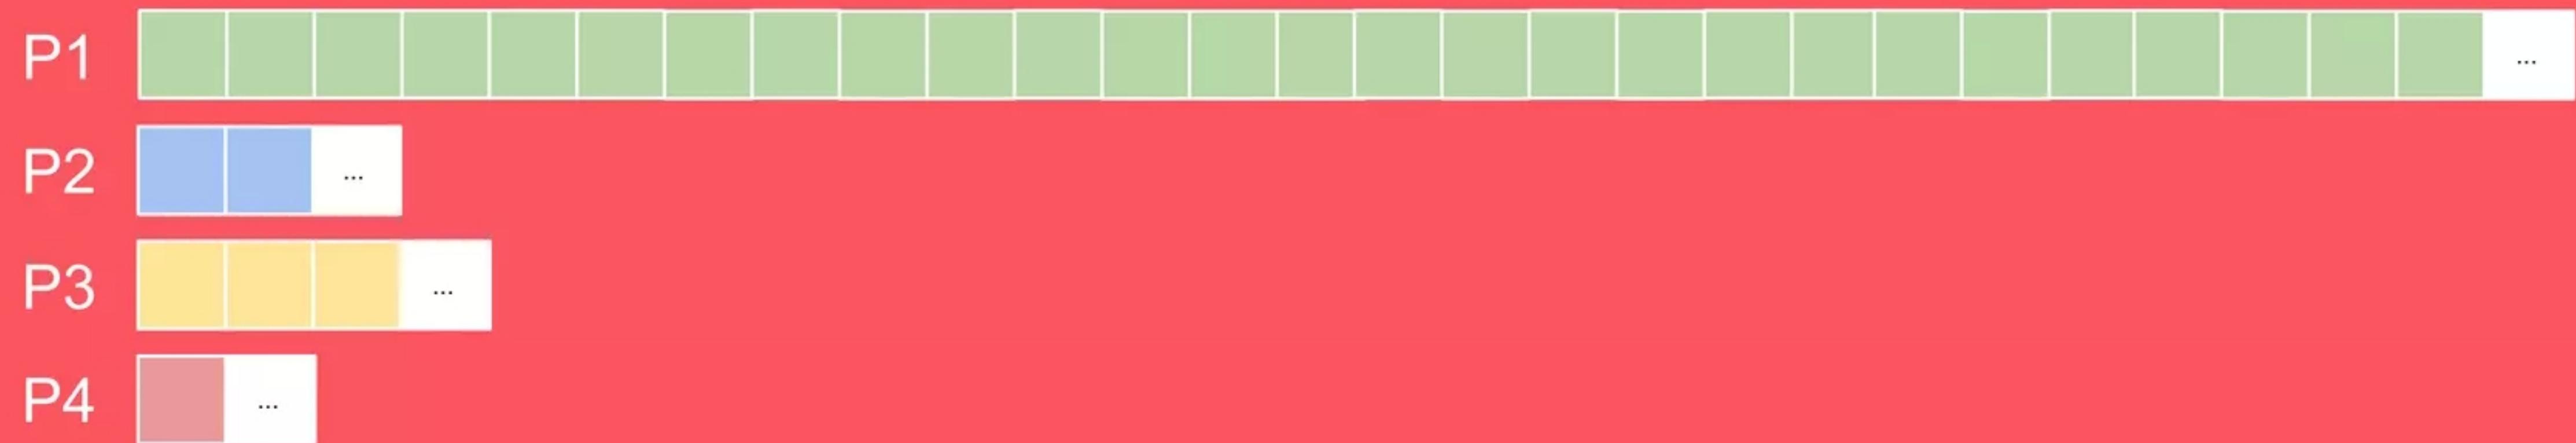
How to increase number of partitions when needed

ksqlDB example: statement below creates a new stream with the desired number of partitions.

```
CREATE STREAM products_repartitioned  
    WITH (PARTITIONS=30) AS  
        SELECT * FROM products  
        EMIT CHANGES
```

'Hot' partitions can be problematic, often caused by

1. Events not evenly distributed across partitions
2. Events evenly distributed but certain events take longer to process



Strategies to address hot partitions include

- 1a. Ingress: Find better partitioning function $f(\text{event.key})$ for producers
- 1b. Storage: Re-partition data into new topic if you can't change the original
2. Scale processing vertically, e.g. more powerful CPU instances

How to re-partition your data when needed

ksqlDB example: statement below creates a new stream with changed number of partitions and a new field as event.key (so that its data is now correctly co-partitioned for joining)

```
CREATE STREAM products_repartitioned
  WITH (PARTITIONS=42) AS
    SELECT * FROM products
    PARTITION BY product_id
    EMIT CHANGES
```

Capacity Planning and Sizing

Sorry, not covered here! I recommend:

ksqlDB Performance Tuning for Fun and Profit, by Nick Dearden

<https://kafka-summit.org/sessions/ksql-performance-tuning-fun-profit/>

ksqlDB documentation

<https://docs.ksqldb.io/en/latest/operate-and-deploy/capacity-planning/>

Kafka Streams documentation

<https://docs.confluent.io/current/streams/sizing.html>

Where to go from here?

My blog series on Kafka Fundamentals to learn more

<https://www.confluent.io/blog/kafka-streams-tables-part-1-event-streaming>

Confluent Cloud, the fastest & easiest way to get started with Kafka

<https://www.confluent.io/confluent-cloud/>

Confluent Platform, if you want to operate Kafka yourself

<https://www.confluent.io/product/confluent-platform/>



THANK YOU

@miguno
michael@confluent.io



cnfl.io/meetups



cnfl.io/blog



cnfl.io/slack